

Automated Formal Verification of Event Orderings in Parallel Hardware and Software



Yatin A. Manerkar (Princeton University)



Memory Consistency Models (MCMs)


- MCMs specify ordering rules which constrain values that can be read by load instructions in parallel programs
 - Sequential Consistency [Lamport IEEE Tran. 1979]
- ISA-level MCMs defined in terms of acyclicity, irreflexivity, etc. of relational patterns
 - [Shasha & Snir TOPLAS 1988] [Alglave et al. TOPLAS 2014]

Message passing (mp) litmus test

Core 0	Core 1
(i1) [x] ← 1	(i3) r1 ← [y]
(i2) [y] ← 1	(i4) r2 ← [x]

Under SC: Forbid r1=1, r2=0

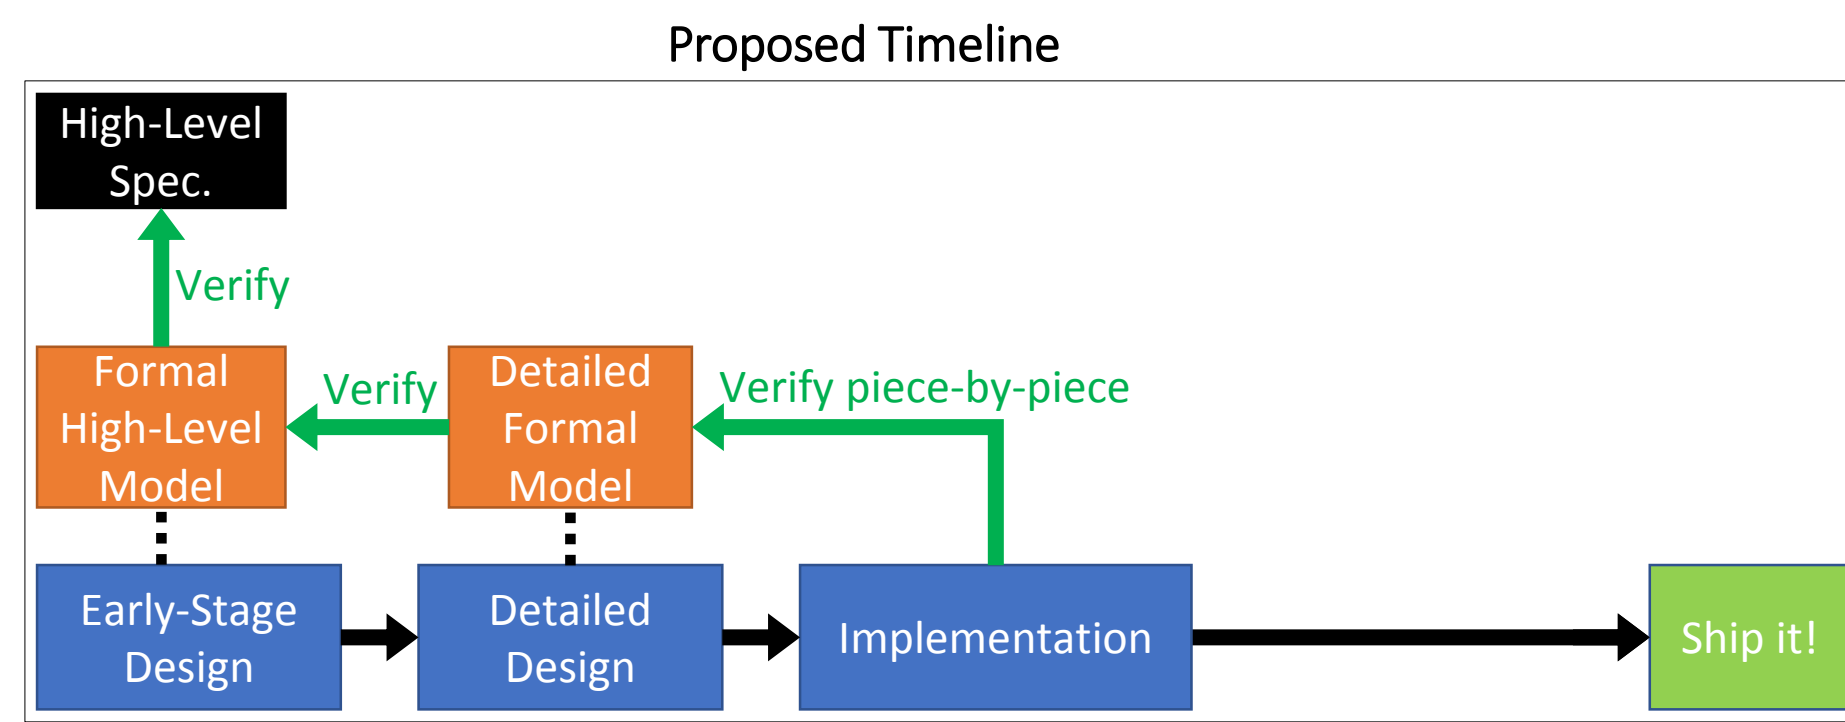
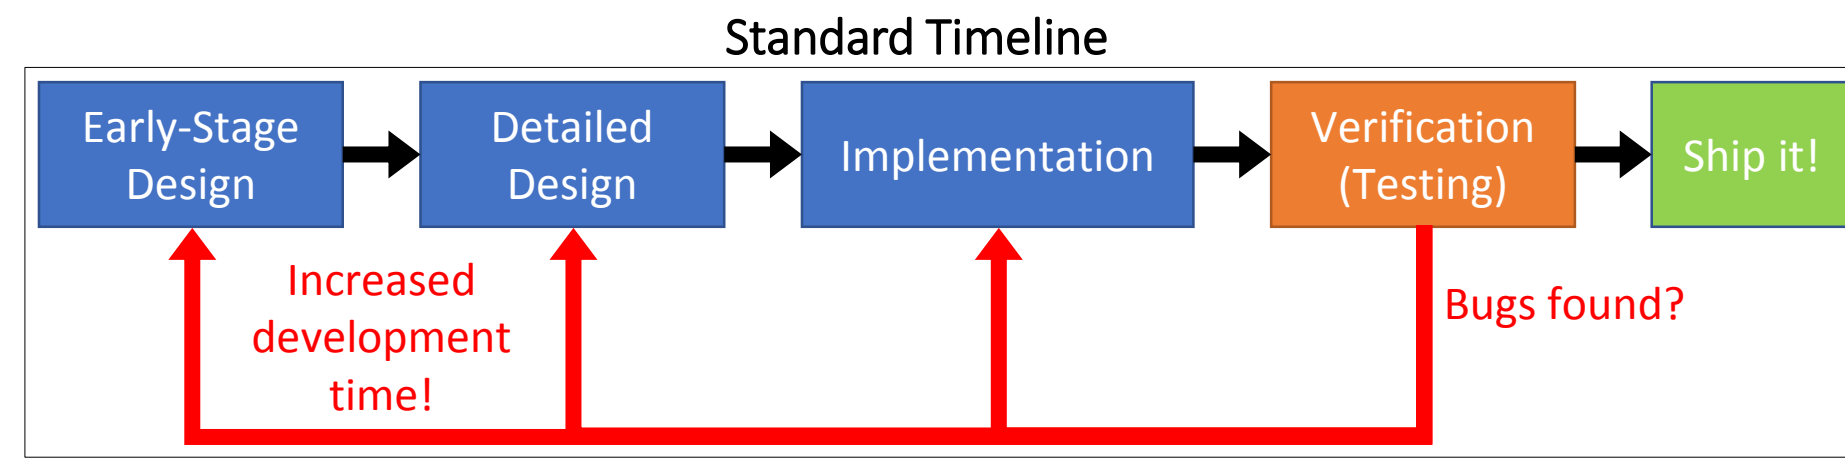
ISA-Level MCM Specification of SC
acyclic (po U co U rf U fr)



Nodes: Instructions
Edges: ISA-level relations

Bringing Verification Earlier in the Design Timeline

- Verification costs now dominate total hardware design cost [Foster DAC 2014]
- Formal verification early in the design timeline can catch bugs quicker, and allow models to evolve with the system!



Prior Work: Microarchitectural MCM Modelling


- Microarchitectural ordering specifications defined as set of μ spec axioms [Lustig, et al. ASPLOS 2016]
- Microarchitectural executions are μ hb (microarchitectural happens-before) graphs

μ spec Microarchitectural Ordering Specification

Axiom "Fetch_is_FIFO":

```
... EdgeExists ((i1, IF), (i2, IF))  
=> AddEdge ((i1, EX), (i2, EX)).  
...
```

μ hb graph of mp on 3-stage processor

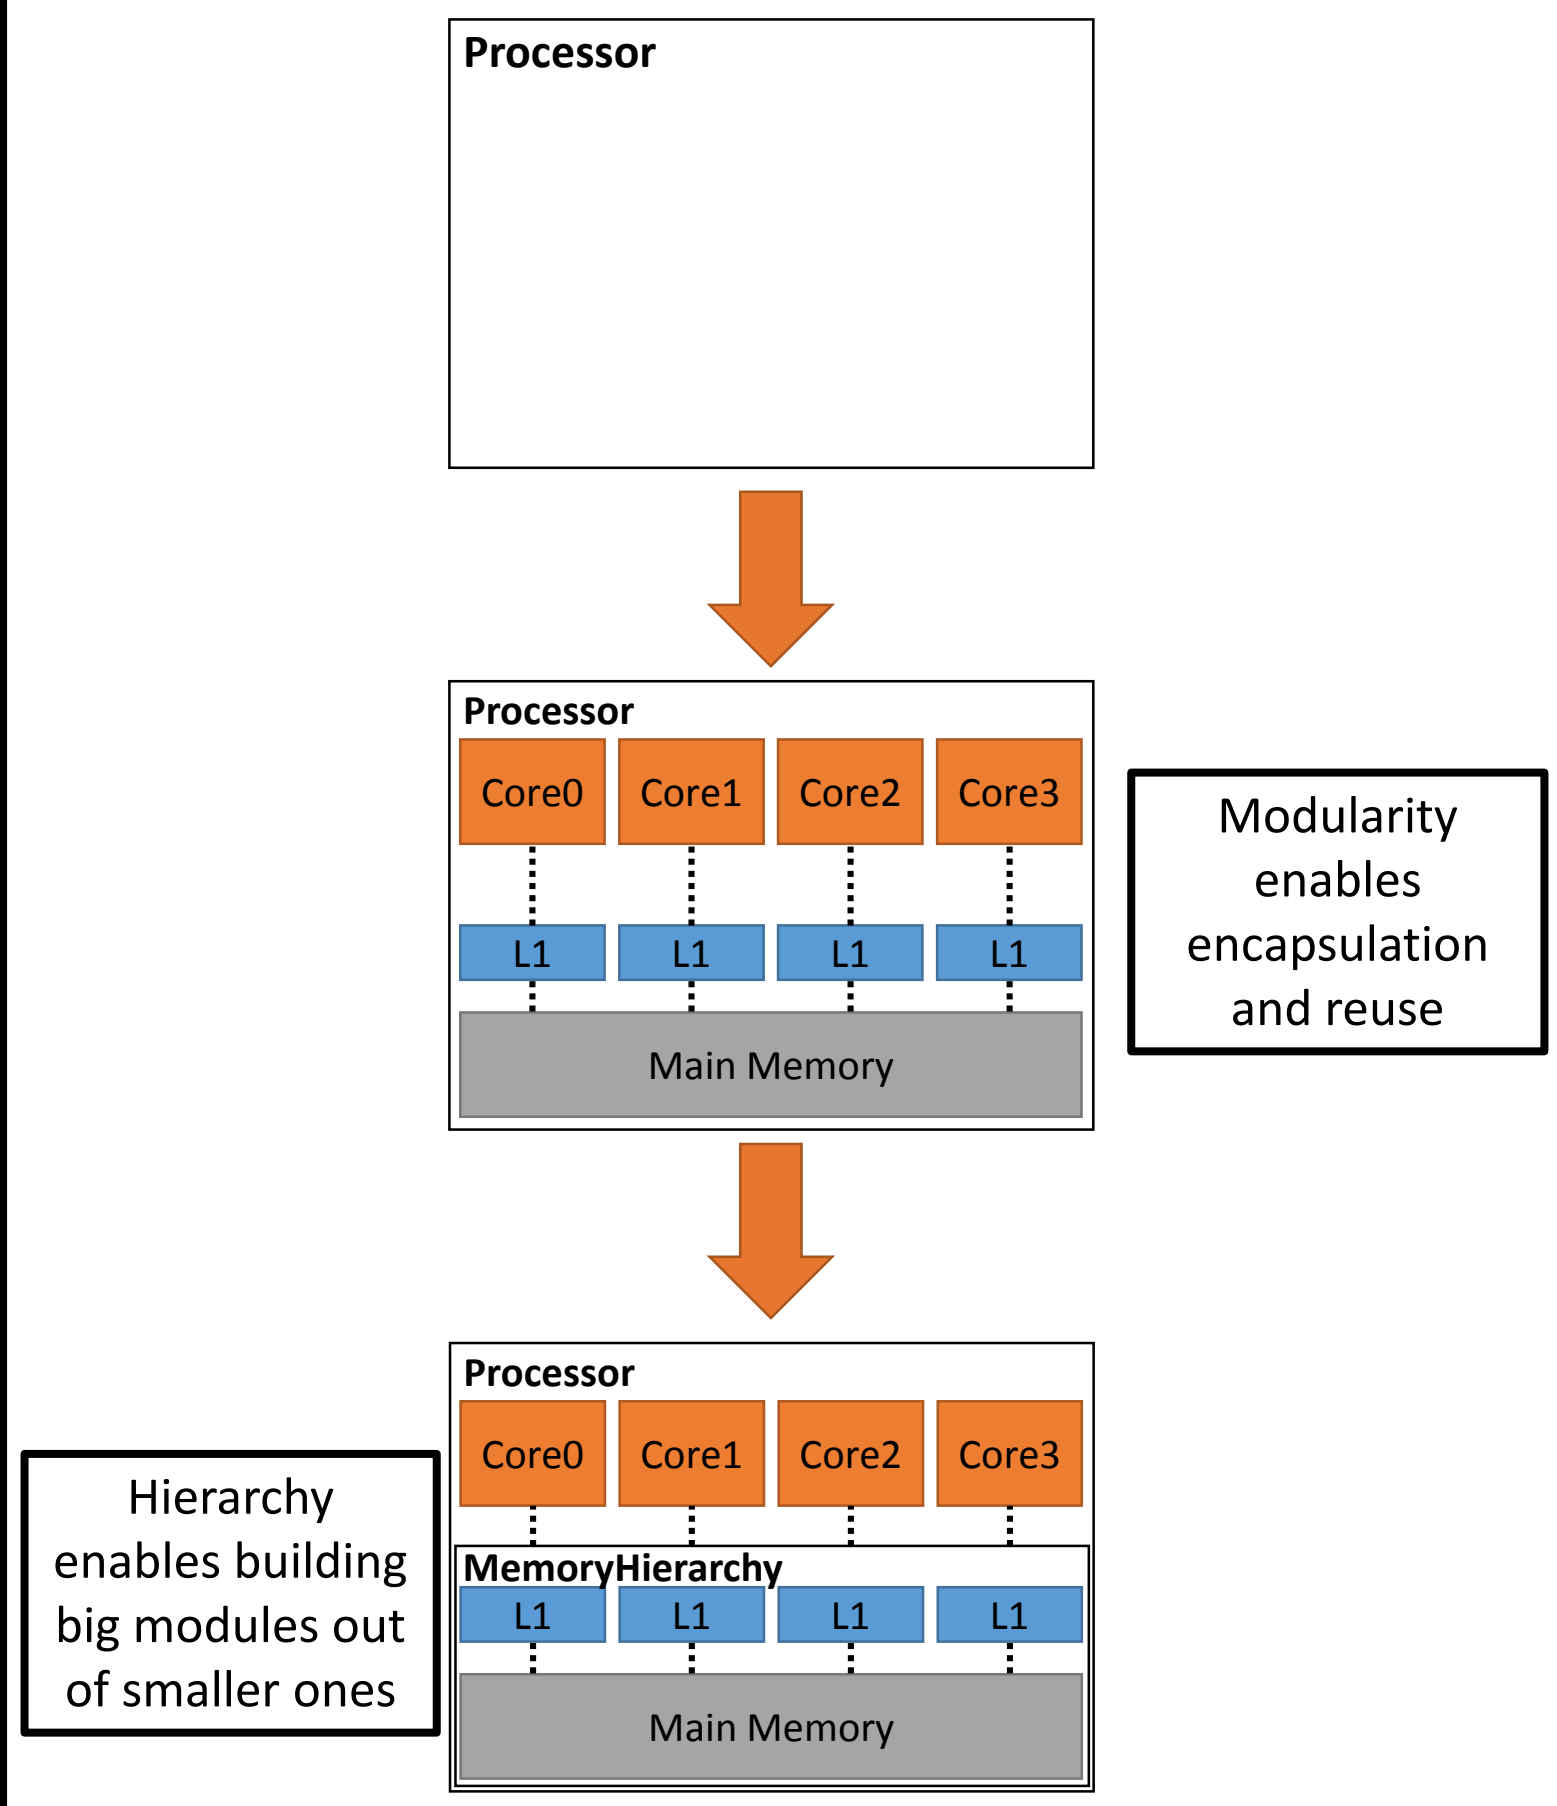


Nodes: Instr. sub-events
Edges: Happens-before order

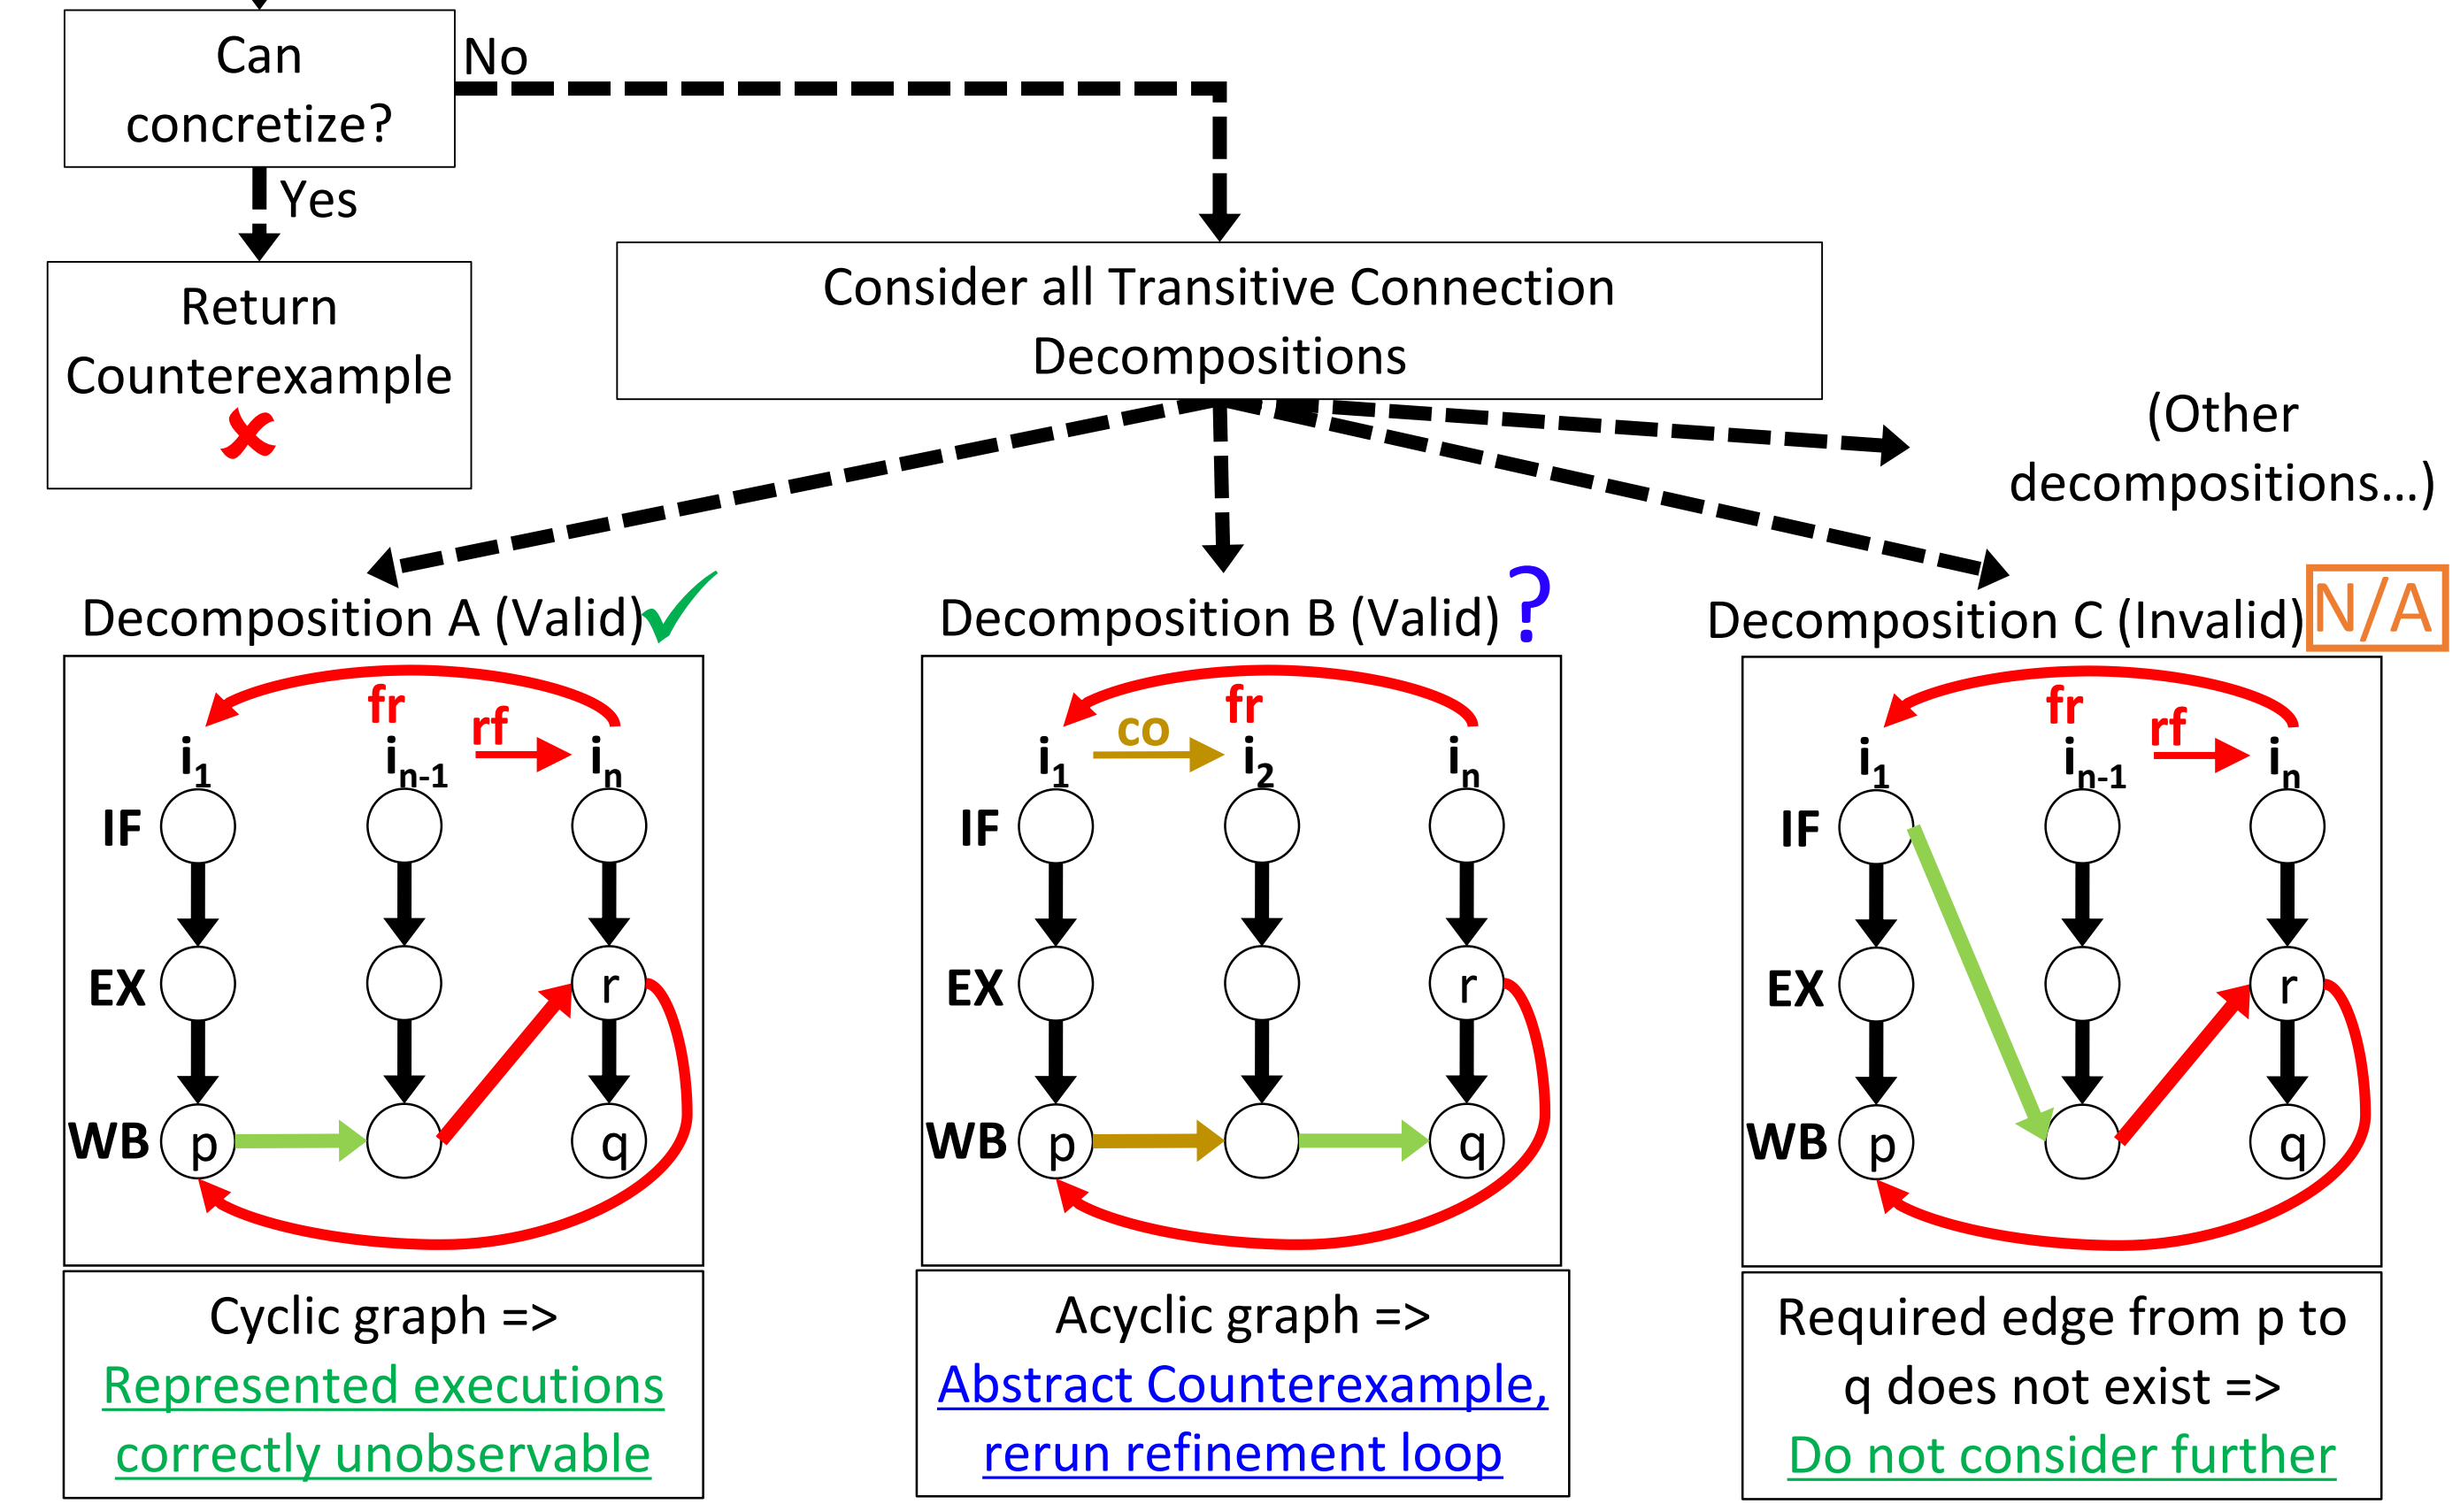


Cyclic graph => Unobservable
Acyclic graph => Observable

RealityCheck: The Need for Modularity

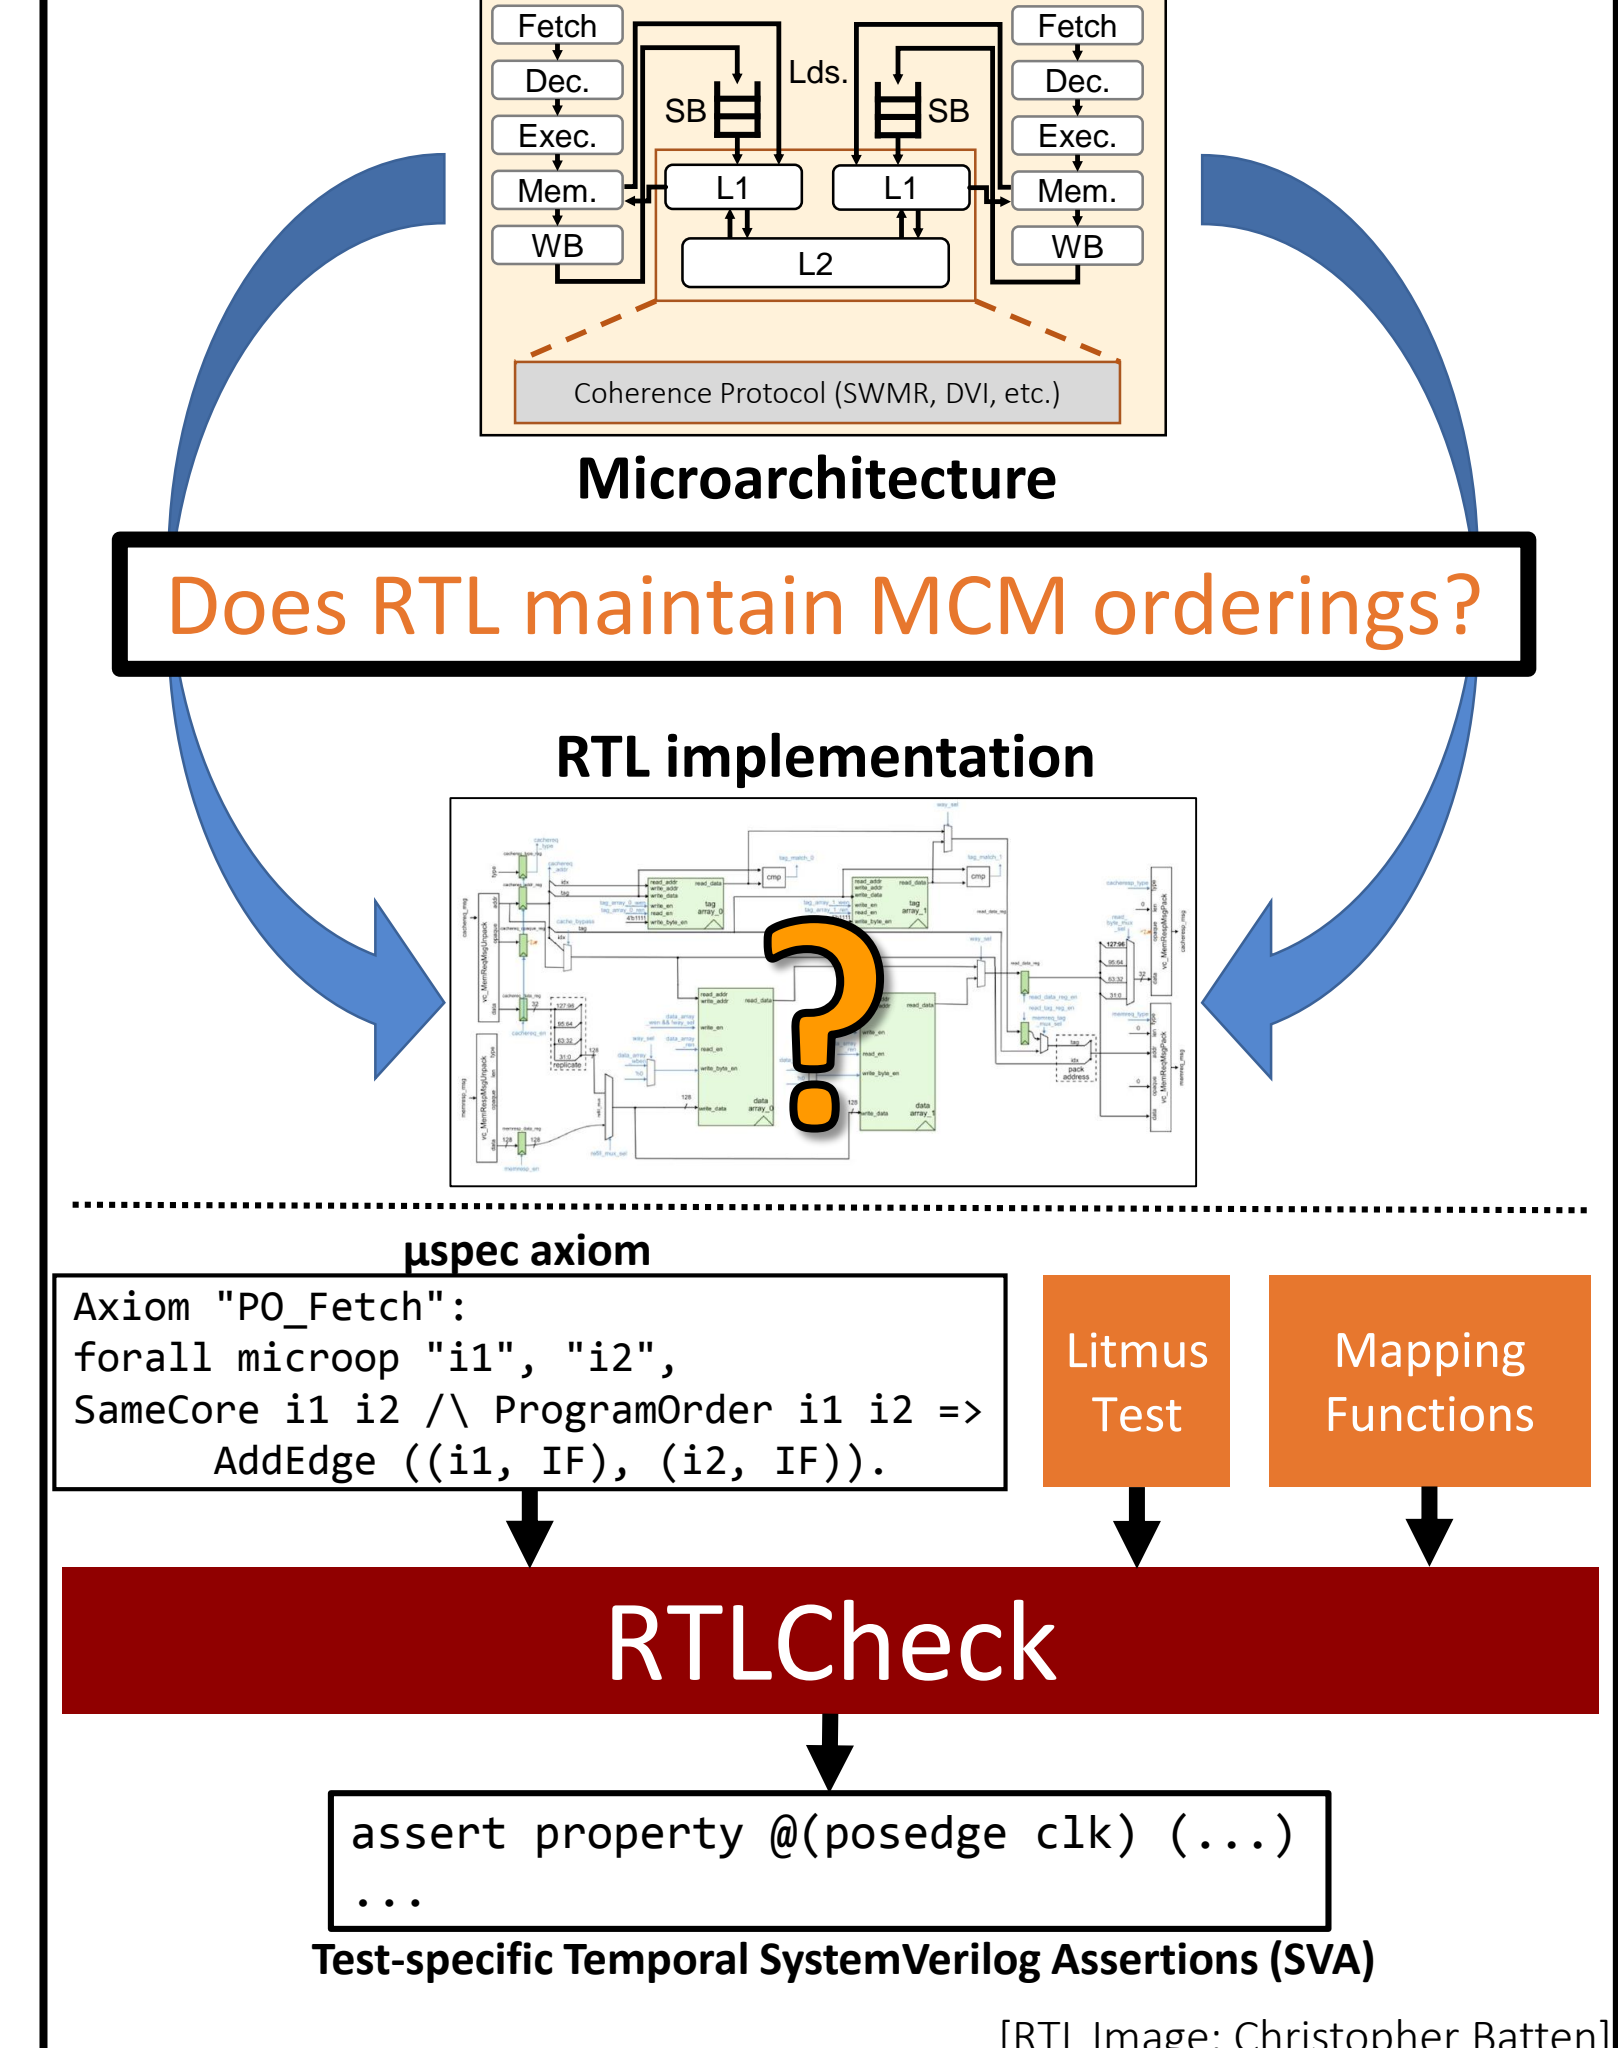
- Flat verification is at odds with realities of design process
- Teams develop components independently and later connect them together



PipeProof: Automated All-Program Microarchitectural MCM Correctness Proofs



RTLCheck: Going Down to Verilog



RTLCheck: The μ spec/SVA Mismatch

- Tricky to translate from μ spec to SVA correctly
 - Situation further complicated by SVA verifiers (e.g. Cadence JasperGold) not fully supporting SVA spec
 - Example:
- Axiom "Read_Returns_0":
- Load returns 0 => Load is BeforeAllWrites
- In SVA, this is represented as:
- ```
assume property (a); // Load returns 0
assert property (b); // Load is BeforeAllWrites
//The above is equivalent to...
assert property ((always a) implies (always b));
```
- SVA is based on temporal logic (G = always, F = eventually):  
$$G a \rightarrow G b = (\sim(G a)) \vee G b = (F \sim a) \vee G b$$
  - Assumptions introduce liveness (through use of F)
    - Expensive to check! [Cerny et al. 2010]
  - SVA verifiers **approximate**: only check assumptions until current state
    - Property is then easier to check...
    - ...but **assumptions no longer correctly enforced!**
  - RTLCheck Solution**: Generate properties that handle all test outcomes

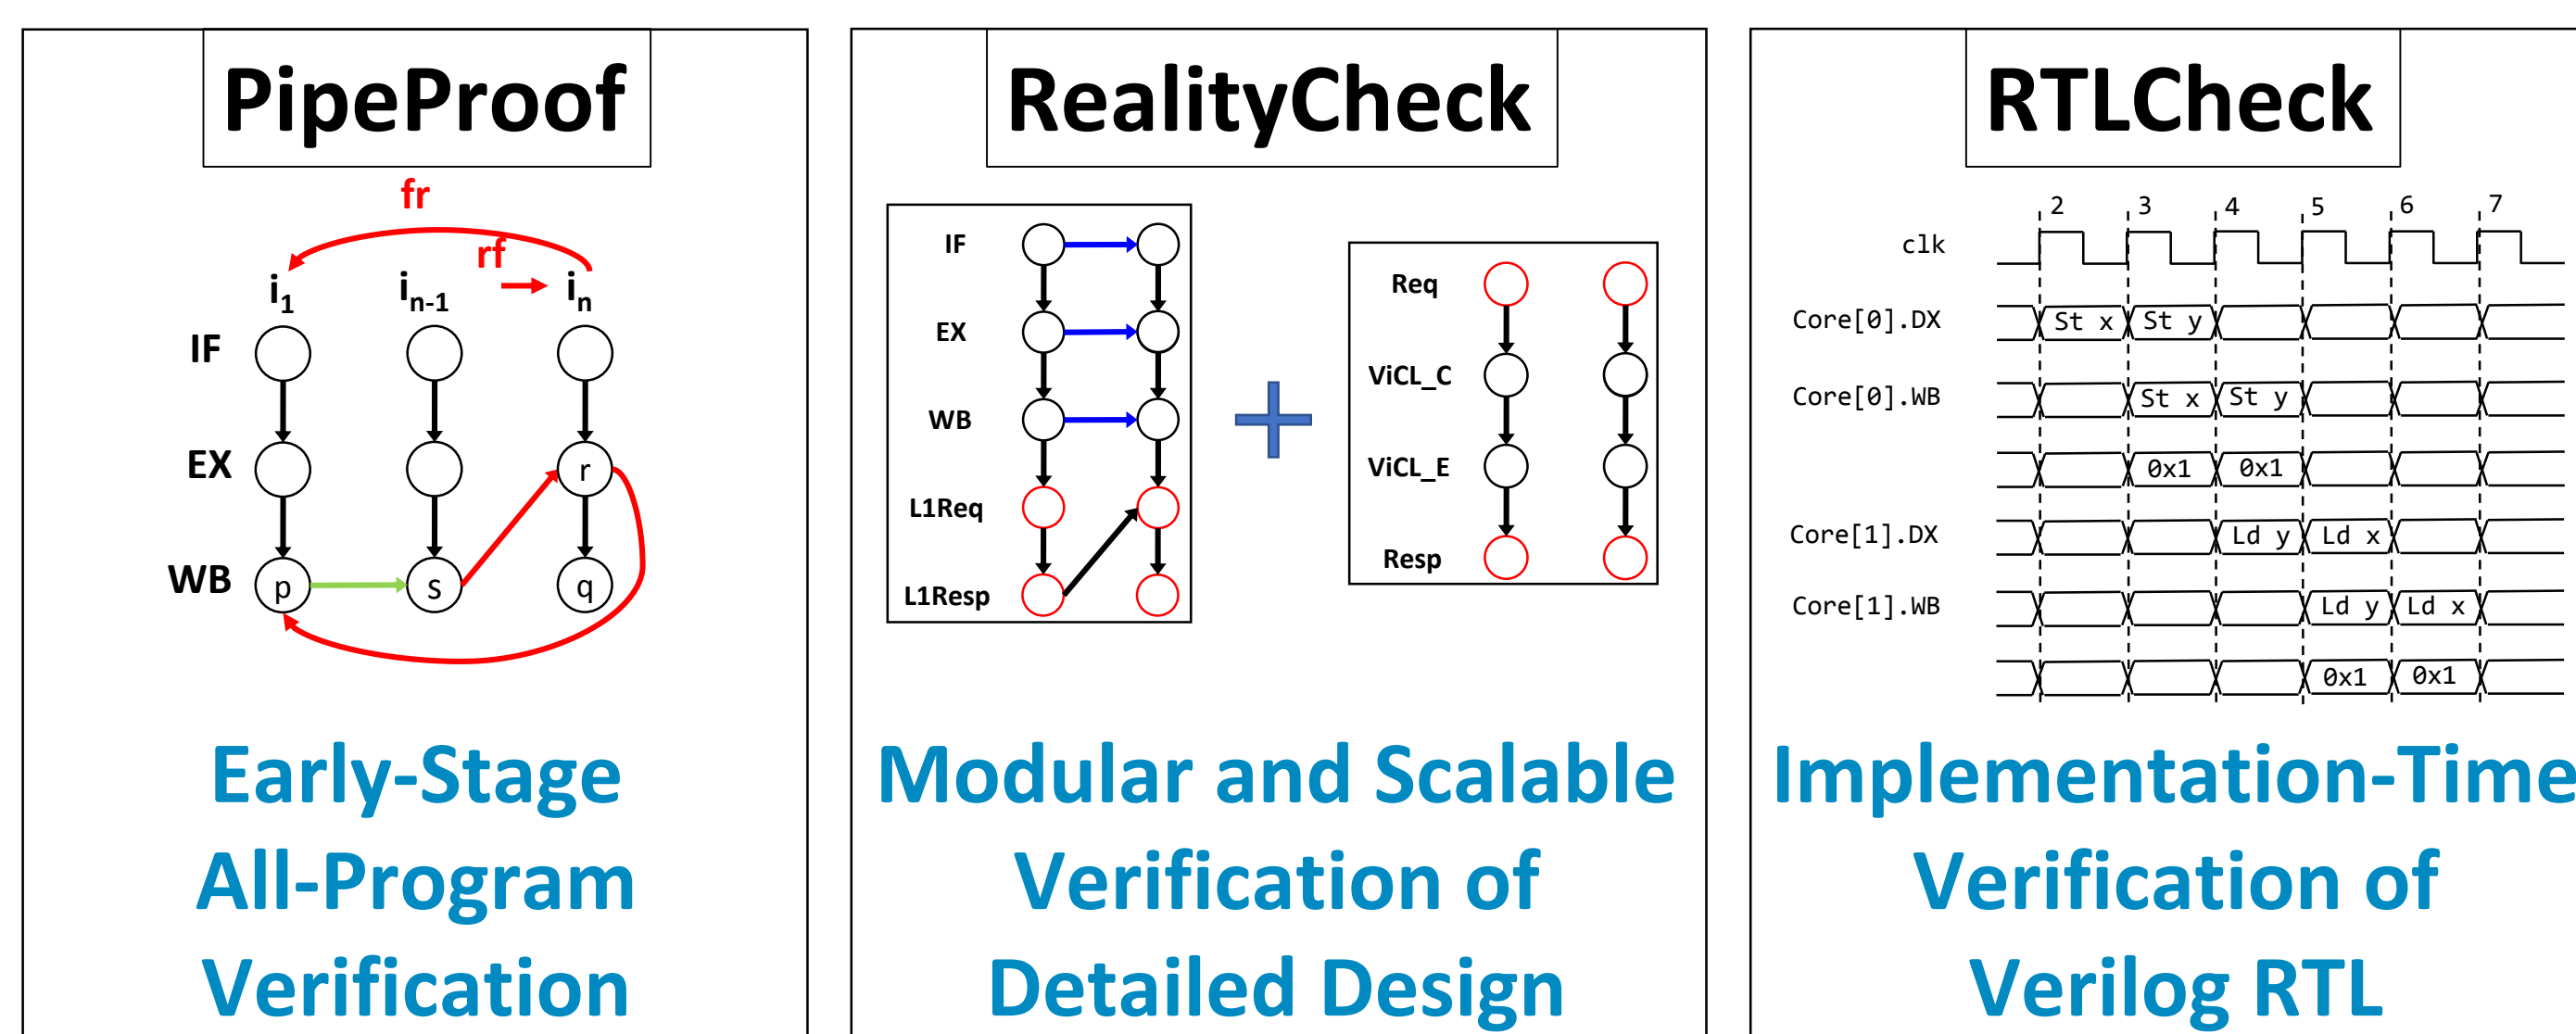
Note:  $\mu$ spec axiom abstracted for brevity

## Nuts, Bolts, and Results

- All tools provide automated formal verification
- Core to approach**: happens-before graphs checked for cycles using efficient SMT solvers
  - RealityCheck uses Z3 solver from Microsoft Research
- RTLCheck utilises commercial SVA prover (Cadence JasperGold) to check generated SVA assertions
- Some Results**:
  - PipeProof can prove correct simple microarchitectures implementing SC and TSO for all programs in < 1 hour
  - JasperGold proved 89% of all RTLCheck-generated properties for 56 litmus tests in 11 hrs/test for open-source RISC-V V-scale processor
  - RTLCheck discovered a bug in the memory implementation of V-scale (reported to developers)
  - Discovered two counterexamples to "trailing-sync" compiler mapping from C11 to Power and ARMv7 [Manerkar et al. CoRR 2016] using Trippel [Trippel et al. ASPLOS 2017]

## Different Tools for Different Points in the Verification Timeline

- PipeProof**: Early-stage flat formal verification of design across all possible programs against ISA-level MCM
- RealityCheck**: Mid-stage hierarchical formal verification of detailed design against ISA-level MCM (bounded)
- RTLCheck**: Late-stage formal verification of Verilog implementation against microarchitectural orderings (bounded)



## Conclusions

- Verification now dominates total hardware design cost
- Bringing verification **earlier** in design timeline can find bugs **quicker** and **reduce** overall development time
- Memory consistency models (MCMs) specify ordering rules which constrain values read by loads in parallel programs
- My thesis**: automated formal methodologies and tools for verification of MCM implementations
- PipeProof**: Ensure that design respects MCM for all possible programs before RTL is written!
- RealityCheck**: Modular, scalable verification capable of tackling large designs
- RTLCheck**: Checking early-stage microarchitectural ordering specifications on real processor implementations

Papers and code available at  
[www.cs.princeton.edu/~manerkar](http://www.cs.princeton.edu/~manerkar)