

# Automated Full-Stack Memory Model Verification with the Check suite

**Yatin Manerkar**

Princeton University

ARM Cambridge, July 20<sup>th</sup>, 2018

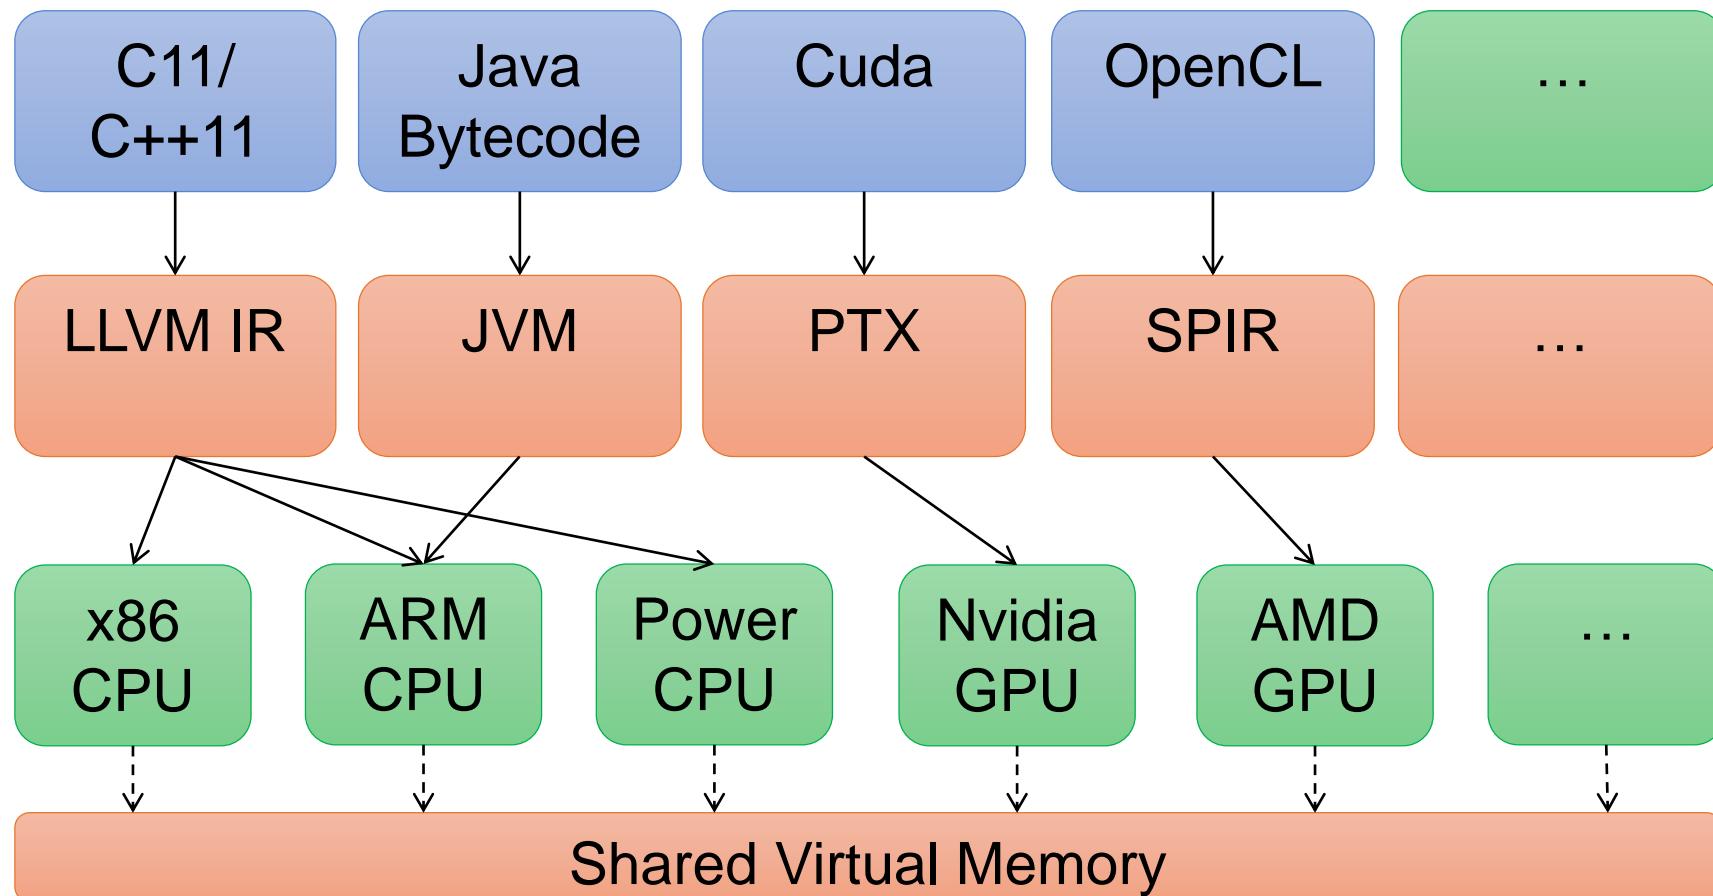
<http://check.cs.princeton.edu/>



# What are Memory (Consistency) Models?

## Memory Consistency Models (MCMs)

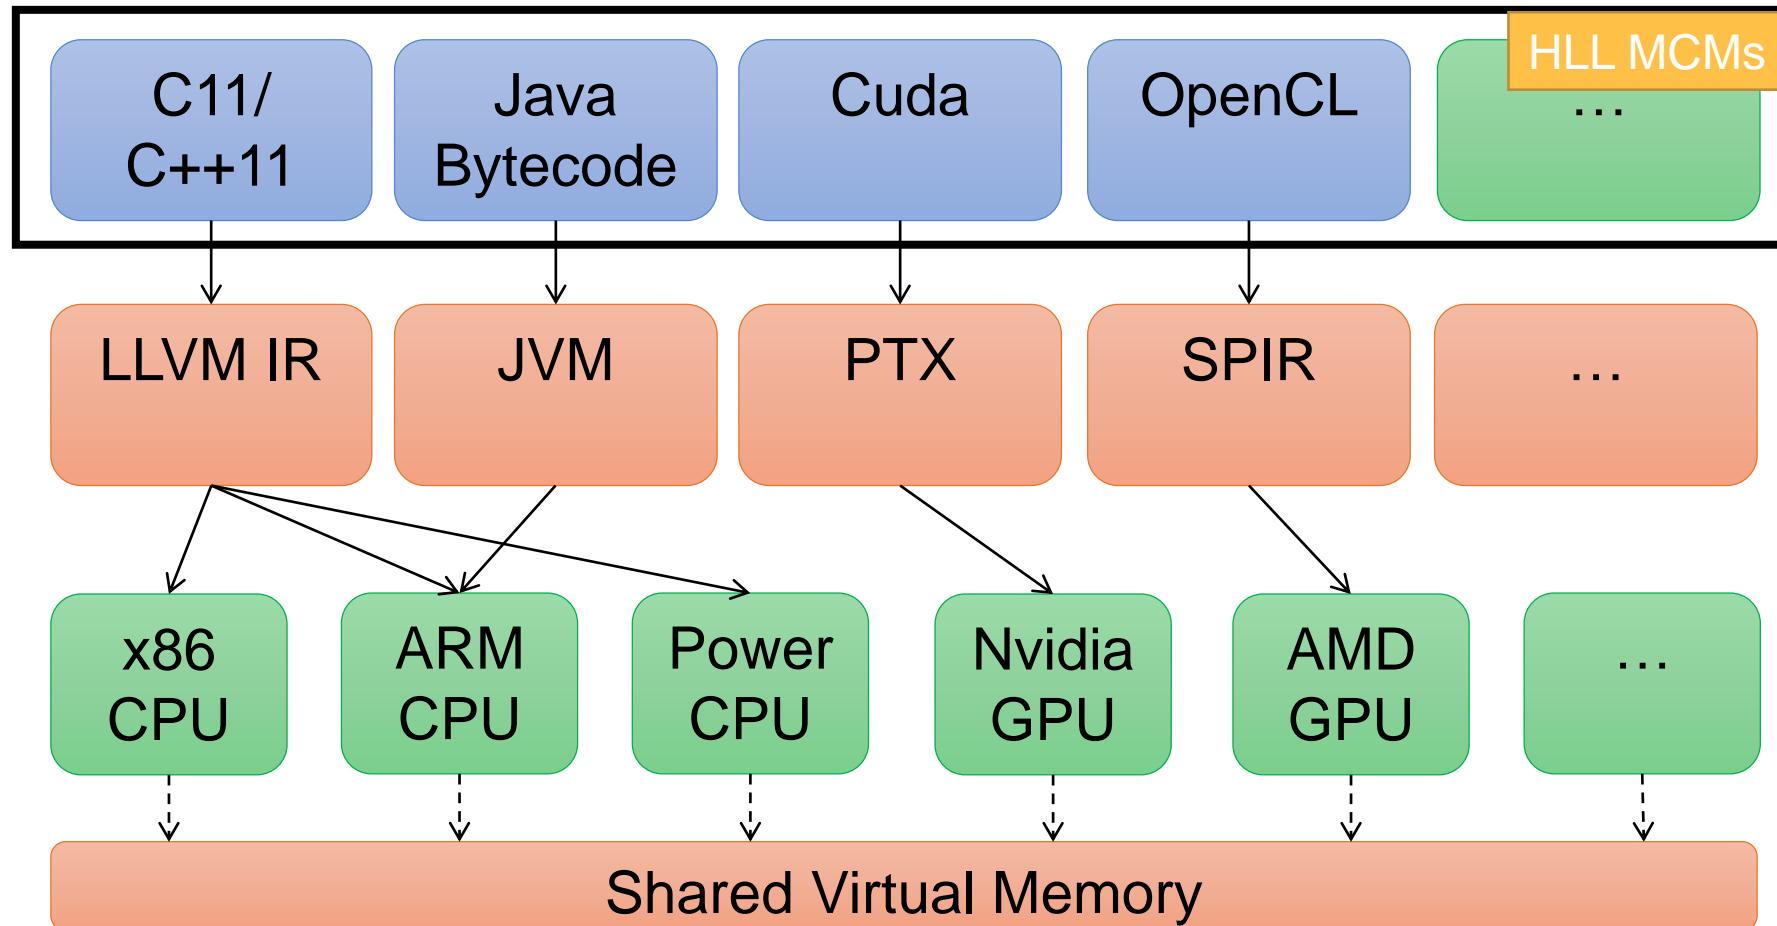
Specify rules and guarantees about the ordering and visibility of accesses to shared memory [Sorin et al., 2011].



# What are Memory (Consistency) Models?

## Memory Consistency Models (MCMs)

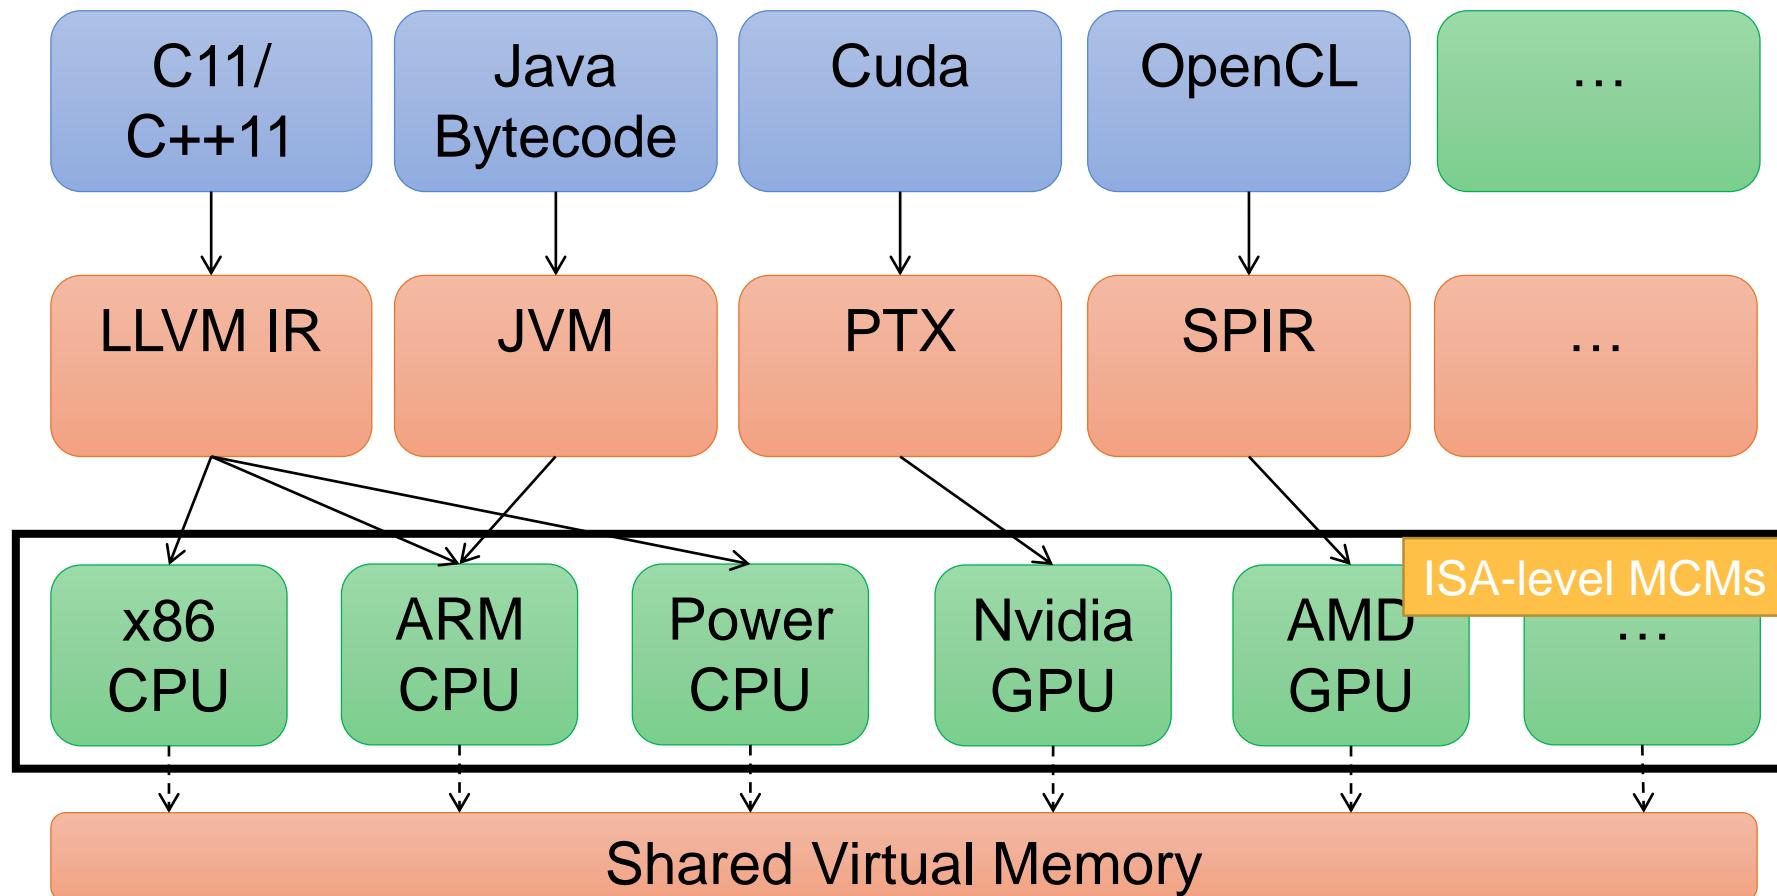
Specify rules and guarantees about the ordering and visibility of accesses to shared memory [Sorin et al., 2011].



# What are Memory (Consistency) Models?

## Memory Consistency Models (MCMs)

Specify rules and guarantees about the ordering and visibility of accesses to shared memory [Sorin et al., 2011].



# Sequential Consistency (SC) - Interleaving Model

- Defined by [Lamport 1979], execution is the same as if:
  - (R1) Memory ops of each processor appear in program order
  - (R2) Memory ops of all processors were executed in some total order  
(load reads the value of last store to its address in the total order)

Program (mp litmus test) (all addrs initially 0)		Legal Executions						Illegal Outcome
Core 0	Core 1	x=1	r1=y	x=1	x=1	r1=y	r1=y	
x=1	r1=y	y=1	r2=x	r1=y	r1=y	x=1	x=1	r1=1
y=1	r2=x	r1=y	x=1	r2=x	y=1	r2=x	y=1	r2=0
		r2=x	y=1	y=1	r2=x	y=1	r2=x	
		r1=1	r1=0					
		r2=1	r2=0					
						r1=0 r2=1		



# Sequential Consistency (SC) - Interleaving Model

- Defined by [Lamport 1979], execution is the same as if:
  - (R1) Memory ops of each processor appear in program order
  - (R2) Memory ops of all processors were executed in some total order  
(load reads the value of last store to its address in the total order)

Program (mp litmus test) (all addrs initially 0)		Legal Executions						Illegal Outcome
Core 0	Core 1	x=1	r1=y	x=1	x=1	r1=y	r1=y	
x=1	r1=y	y=1	r2=x	r1=y	r1=y	x=1	x=1	
y=1	r2=x	r1=y	x=1	r2=x	y=1	r2=x	y=1	
		r2=x	y=1	y=1	r2=x	y=1	r2=x	
								r1=1 r2=0
		r1=1	r1=0					
		r2=1	r2=0					
								r1=0 r2=1



# Hardware Implements Weak Memory Models

- Most processors don't implement SC
  - x86: Total Store Order (TSO): Relaxes Write->Read ordering
  - ARMv8 and Power relax more orderings
- Compilation to weak memory ISAs must maintain ordering guarantees
  - [Owens et al. TPHOLS 2009], [Batty et al. POPL 2011, POPL 2012], [Wickerson et al. OOPSLA 2015], ...

## C11 Source Code

```
atomic<int> x = 0;  
atomic<int> y = 0;
```

Thread 0	Thread 1
x = 1;	r1 = y;
y = 1;	r2 = x;

**C11 Forbids: r1 = 1, r2 = 0**



# Hardware Implements Weak Memory Models

- Most processors don't implement SC
  - x86: Total Store Order (TSO): Relaxes Write->Read ordering
  - ARMv8 and Power relax more orderings
- Compilation to weak memory ISAs must maintain ordering guarantees
  - [Owens et al. TPHOLS 2009], [Batty et al. POPL 2011, POPL 2012], [Wickerson et al. OOPSLA 2015], ...

## C11 Source Code

```
atomic<int> x = 0;  
atomic<int> y = 0;
```

Thread 0	Thread 1
x = 1;	r1 = y;
y = 1;	r2 = x;
<b>C11 Forbids: r1 = 1, r2 = 0</b>	



# Hardware Implements Weak Memory Models

- Most processors don't implement SC
  - x86: Total Store Order (TSO): Relaxes Write->Read ordering
  - ARMv8 and Power relax more orderings
- Compilation to weak memory ISAs must maintain ordering guarantees
  - [Owens et al. TPHOLS 2009], [Batty et al. POPL 2011, POPL 2012], [Wickerson et al. OOPSLA 2015], ...

## C11 Source Code

```
atomic<int> x = 0;  
atomic<int> y = 0;
```

Thread 0	Thread 1
x = 1;	r1 = y;
y = 1;	r2 = x;
<b>C11 Forbids: r1 = 1, r2 = 0</b>	

Compile

## ARMv8 Assembly Language

Initially, [x] = [y] = 0

Core 0	Core 1
stl #1, [x]	lda r1, [y]
stl #1, [y]	lda r2, [x]
<b>ARMv8 forbids: r1 = 1, r2 = 0</b>	



# Hardware Implements Weak Memory Models

- Most processors don't implement SC
  - x86: Total Store Order (TSO): Relaxes Write->Read ordering
  - ARMv8 and Power relax more orderings

**Is the ARMv8 hardware correctly implementing  
the ARMv8 MCM?**

atomic<int> y = 0;		Initially, [x] = [y] = 0	
Thread 0	Thread 1	Core 0	Core 1
x = 1;	r1 = y;	stl #1, [x]	lda r1, [y]
y = 1;	r2 = x;	stl #1, [y]	lda r2, [x]
<b>C11 Forbids: r1 = 1, r2 = 0</b>		<b>ARMv8 forbids: r1 = 1, r2 = 0</b>	

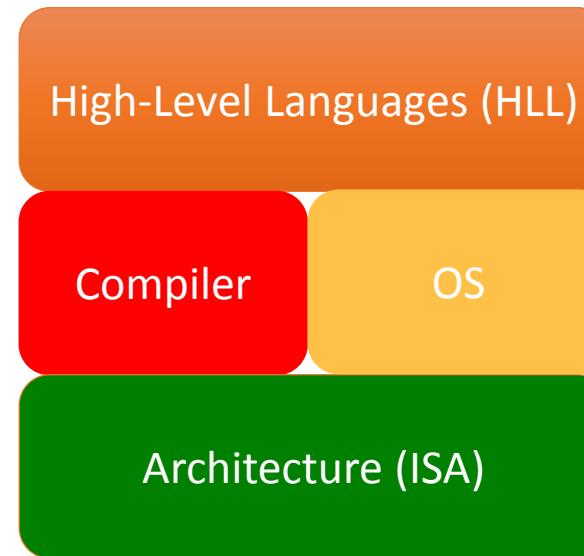
Compile →



# MCM Verification is a Full-Stack Problem!

[Batty et al. POPL 2011, POPL 2012]  
[Algave et al. TOPLAS 2014]  
[Wickerson et al. OOPSLA 2015]

...



Is compiler maintaining  
HLL guarantees?

Is the ISA-level MCM  
formally defined?

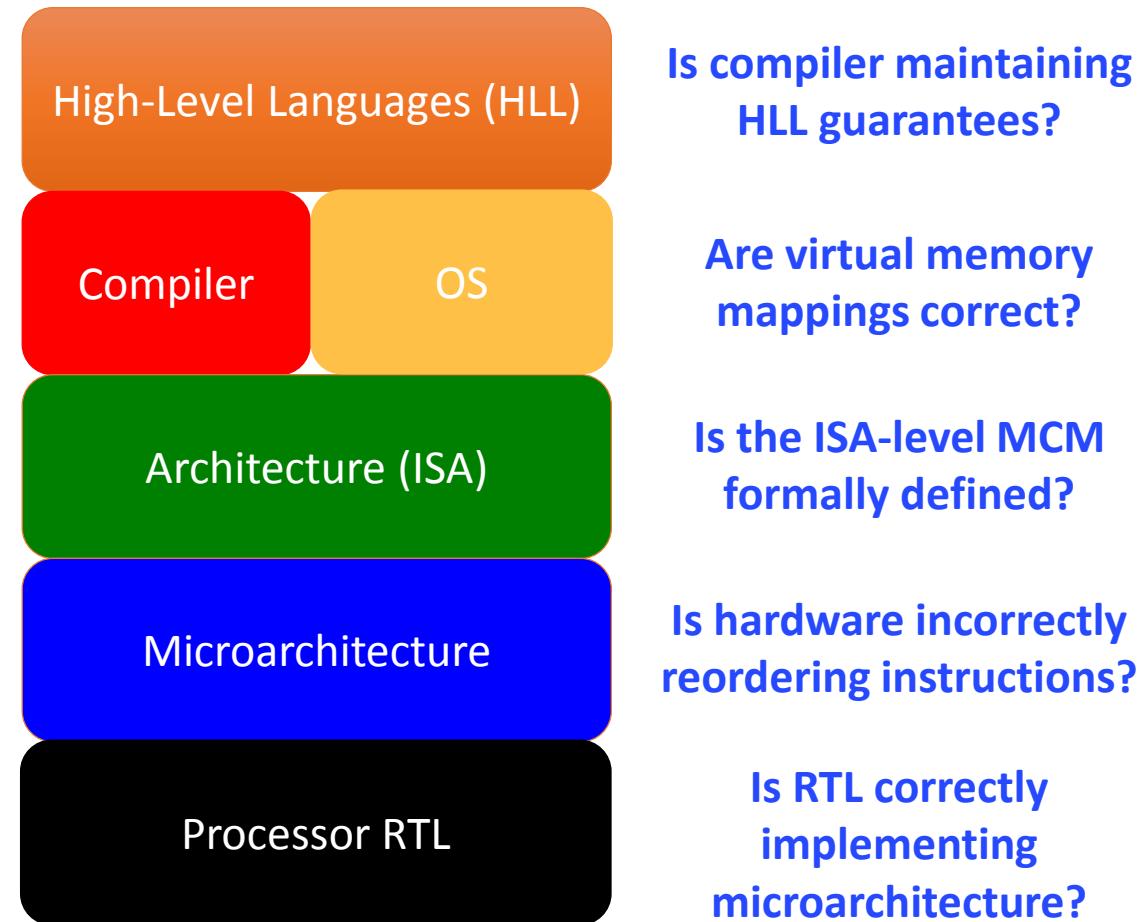
- Each layer has responsibilities for ensuring correct MCM operation
- Need MCM checking tools at all layers of the computing stack!



# MCM Verification is a Full-Stack Problem!

[Batty et al. POPL 2011, POPL 2012]  
[Alglave et al. TOPLAS 2014]  
[Wickerson et al. OOPSLA 2015]

...



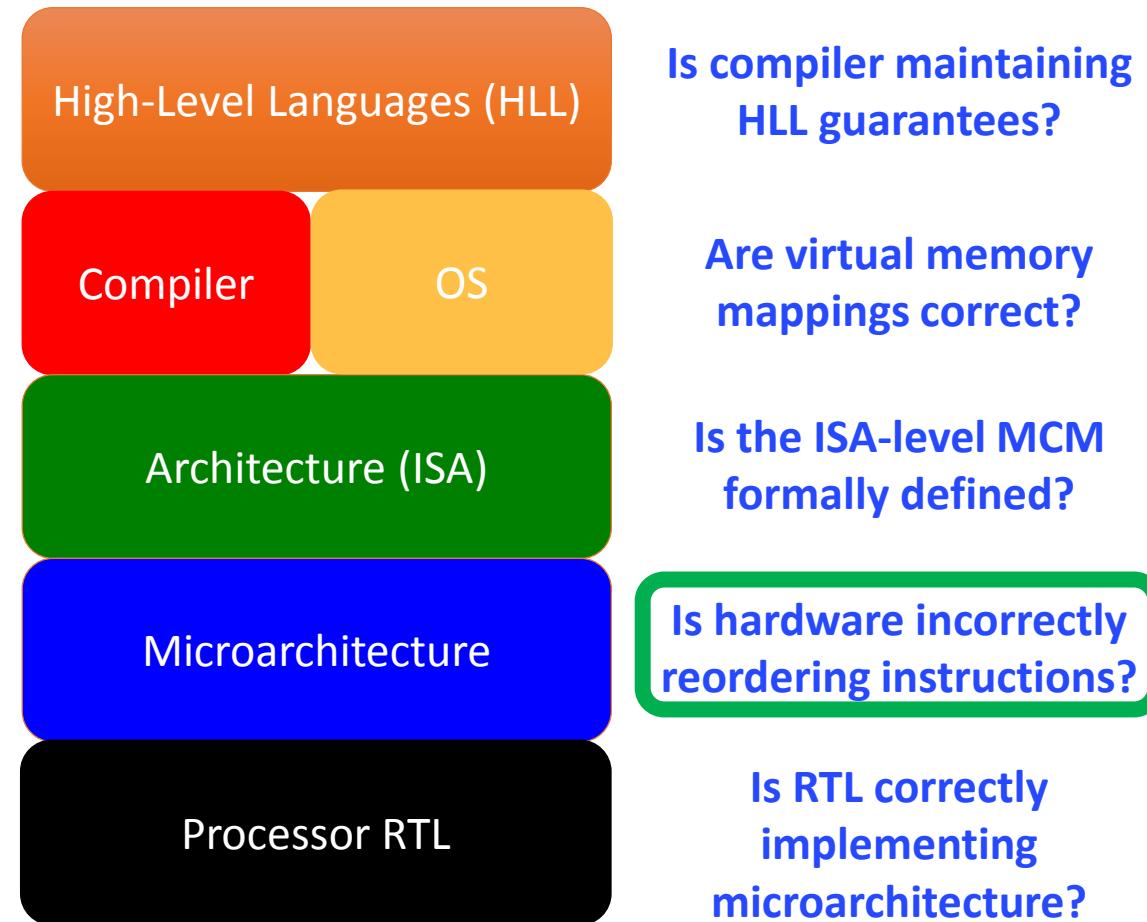
- Each layer has responsibilities for ensuring correct MCM operation
- Need MCM checking tools at all layers of the computing stack!



# MCM Verification is a Full-Stack Problem!

[Batty et al. POPL 2011, POPL 2012]  
[Alglave et al. TOPLAS 2014]  
[Wickerson et al. OOPSLA 2015]

...



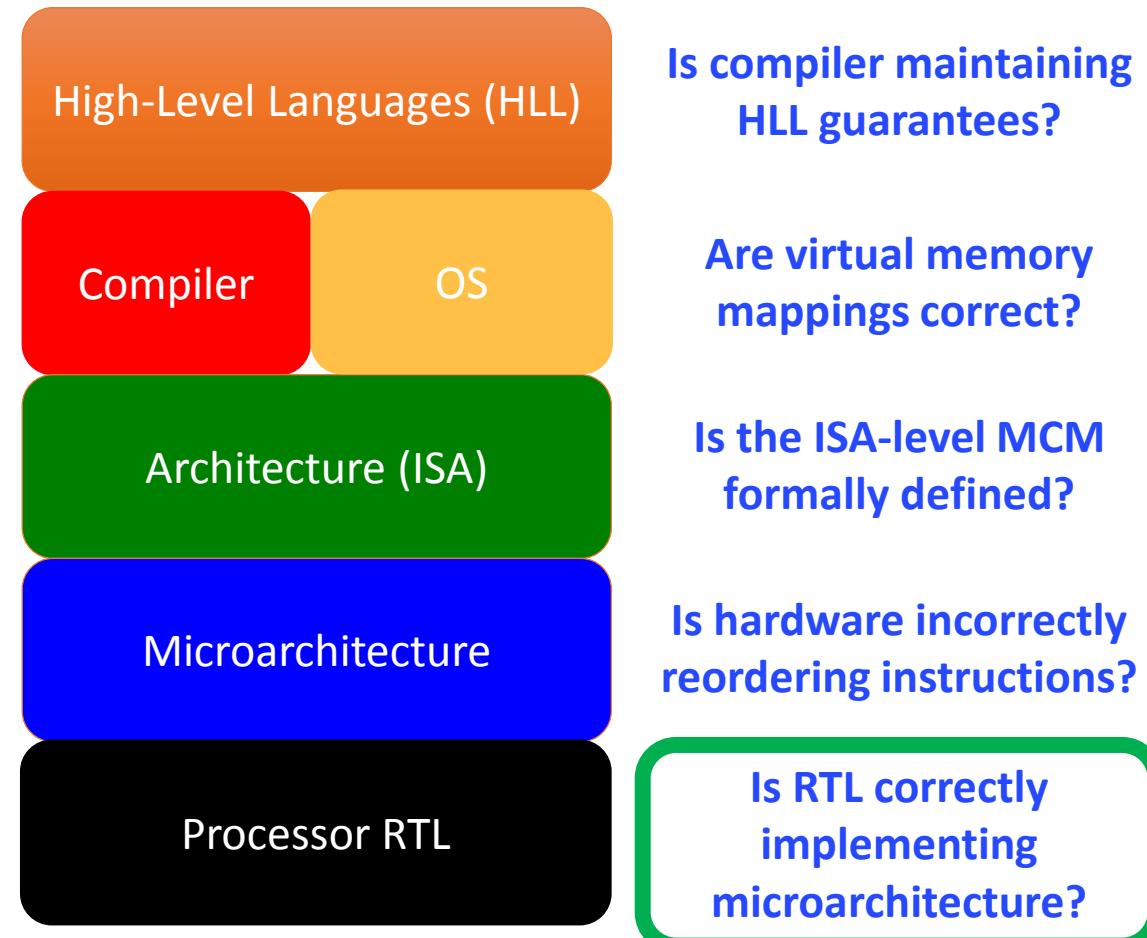
- Each layer has responsibilities for ensuring correct MCM operation
- Need MCM checking tools at all layers of the computing stack!



# MCM Verification is a Full-Stack Problem!

[Batty et al. POPL 2011, POPL 2012]  
[Alglave et al. TOPLAS 2014]  
[Wickerson et al. OOPSLA 2015]

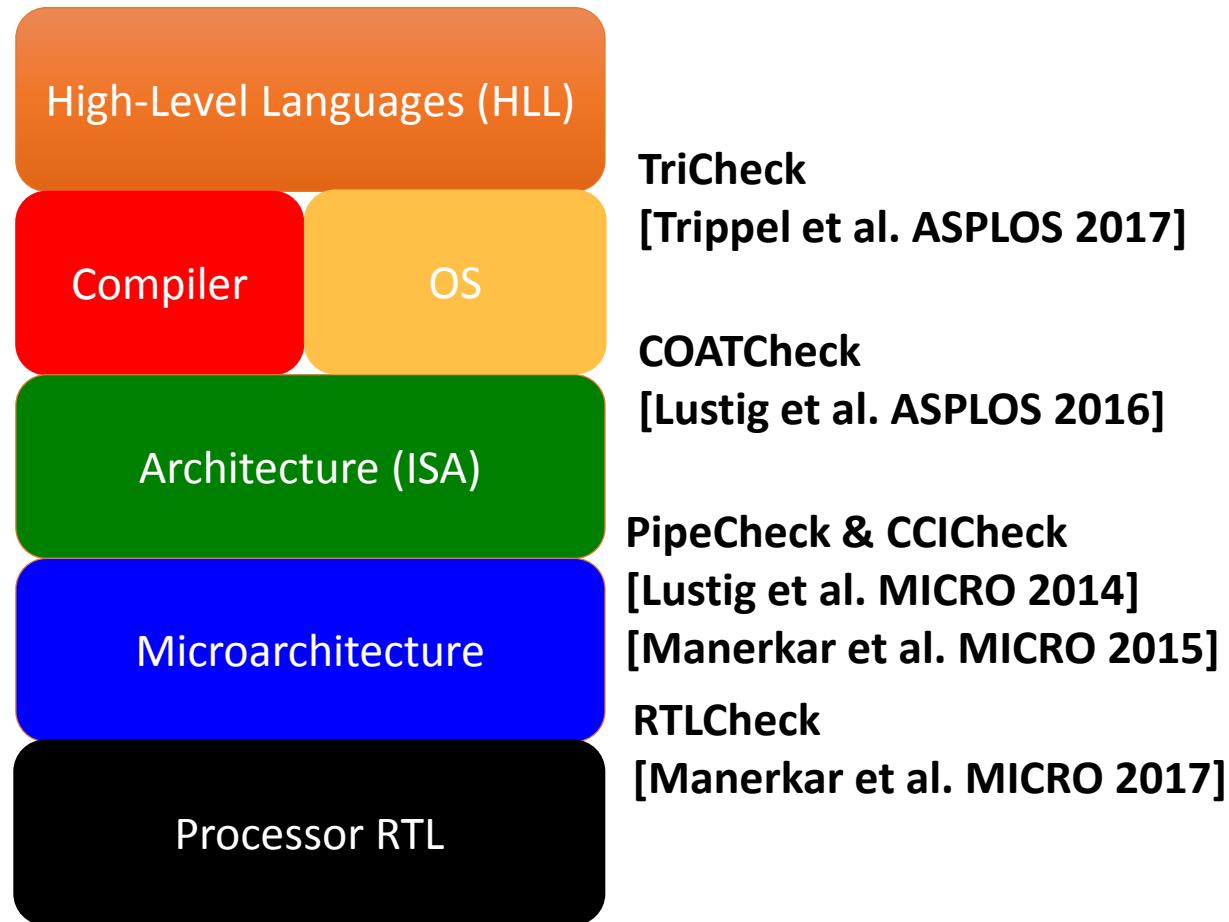
...



- Each layer has responsibilities for ensuring correct MCM operation
- Need MCM checking tools at all layers of the computing stack!



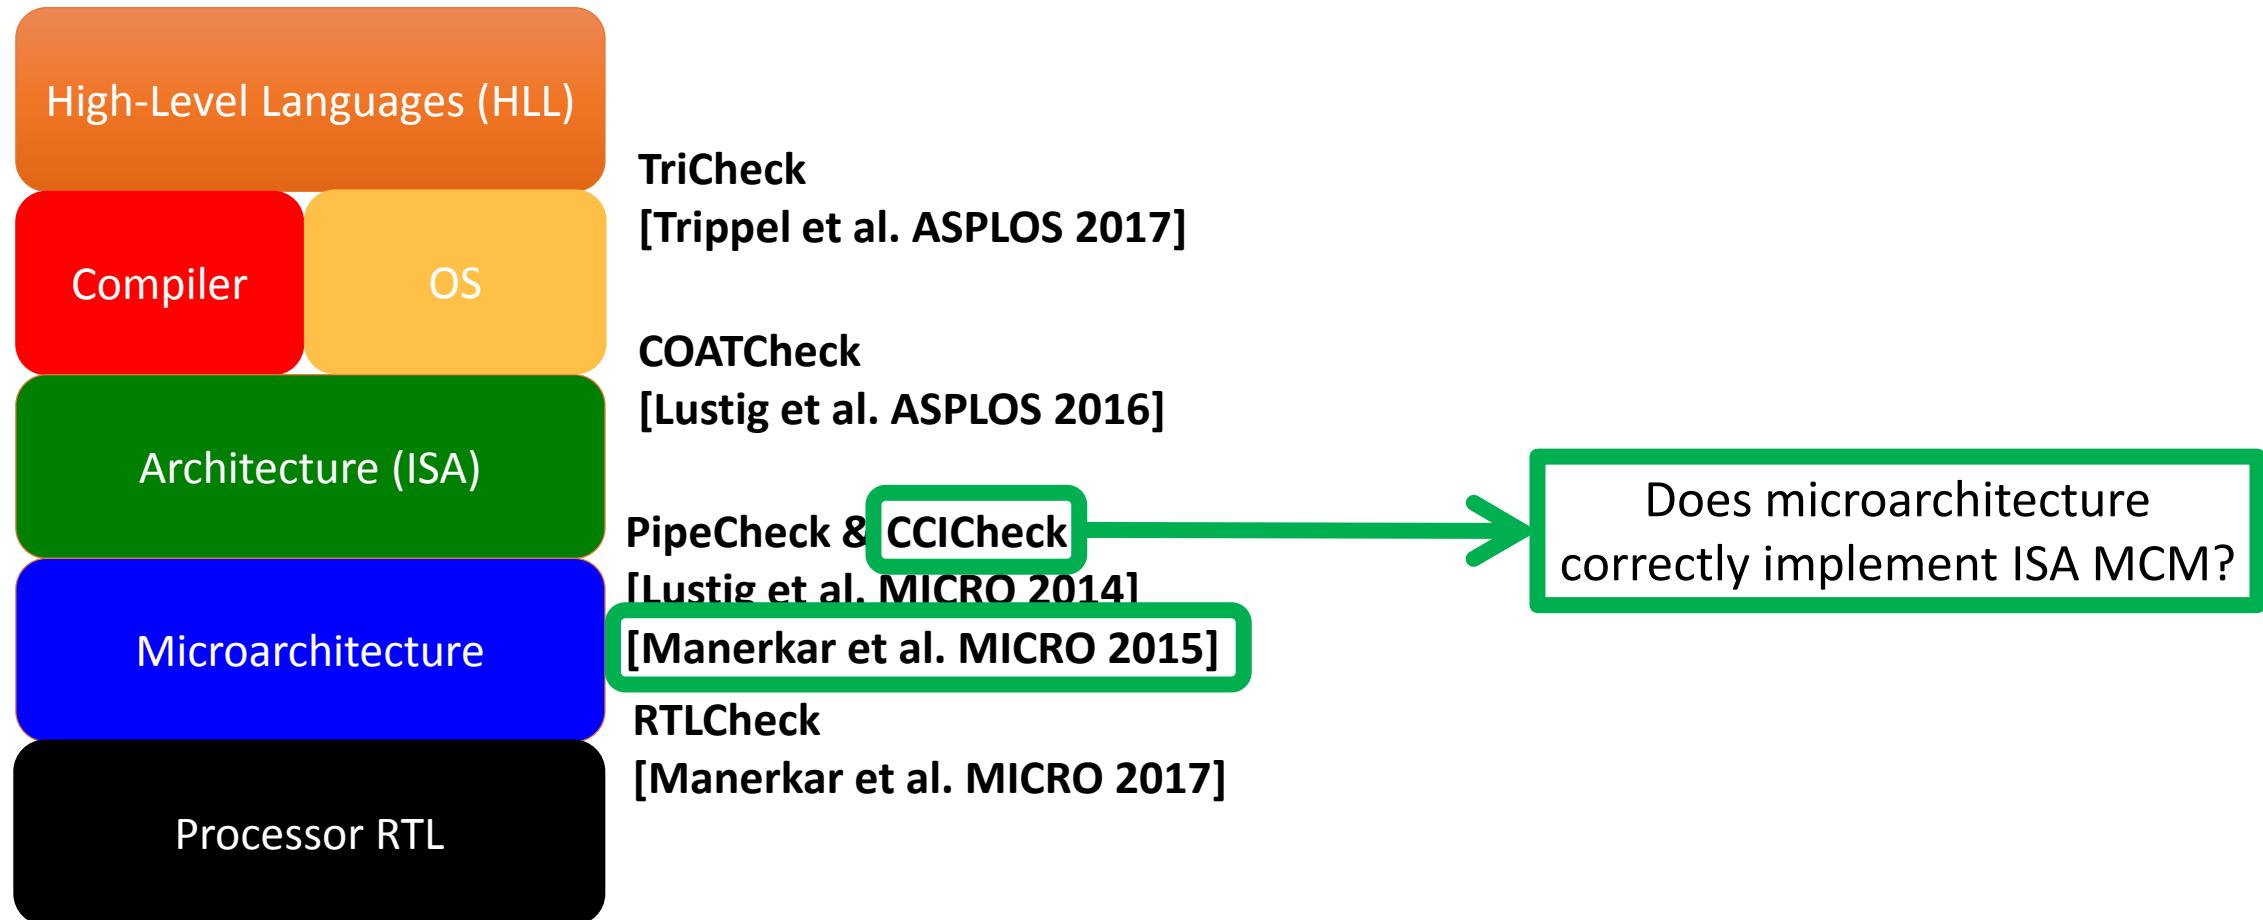
# Check Suite: Full-Stack Automated MCM Analysis



- Suite of tools at various levels of computing stack
- **Automated Full-Stack MCM checking** across litmus test suites



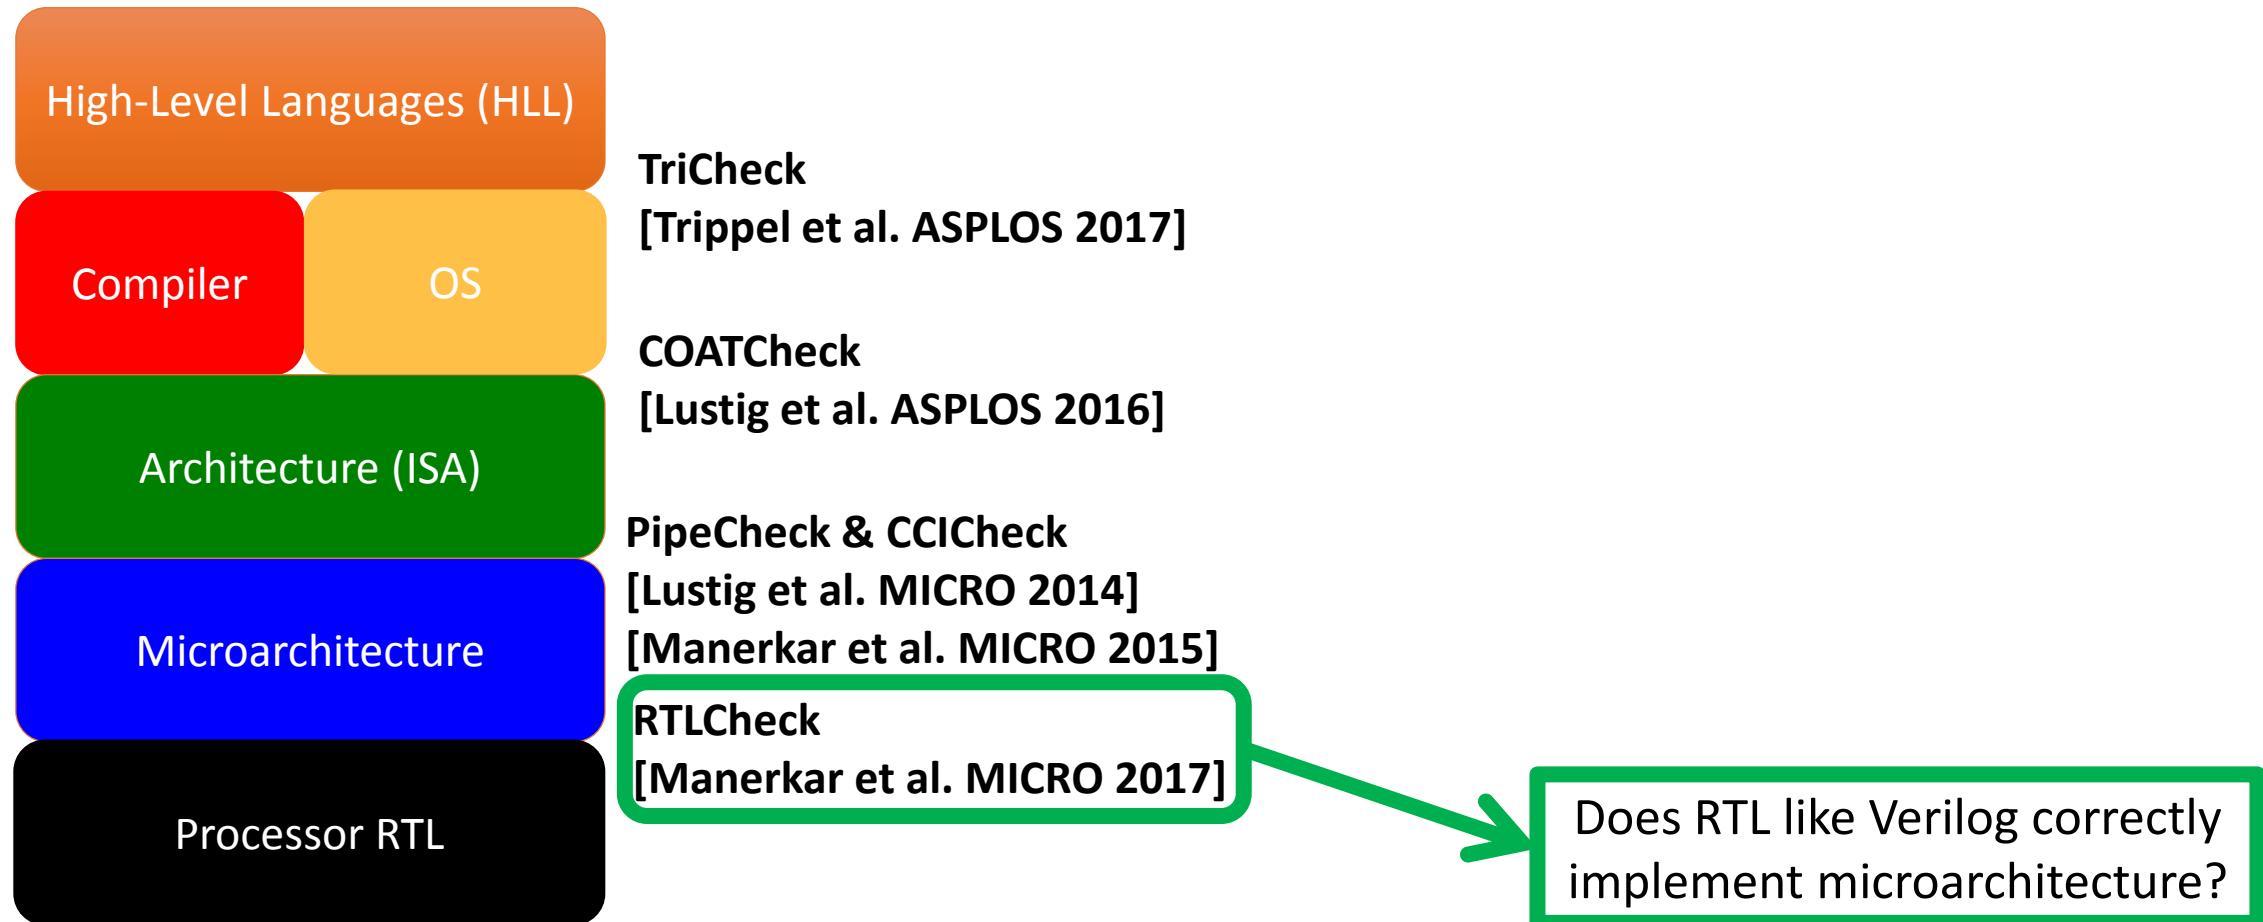
# Check Suite: Full-Stack Automated MCM Analysis



- Suite of tools at various levels of computing stack
- **Automated Full-Stack MCM checking** across litmus test suites



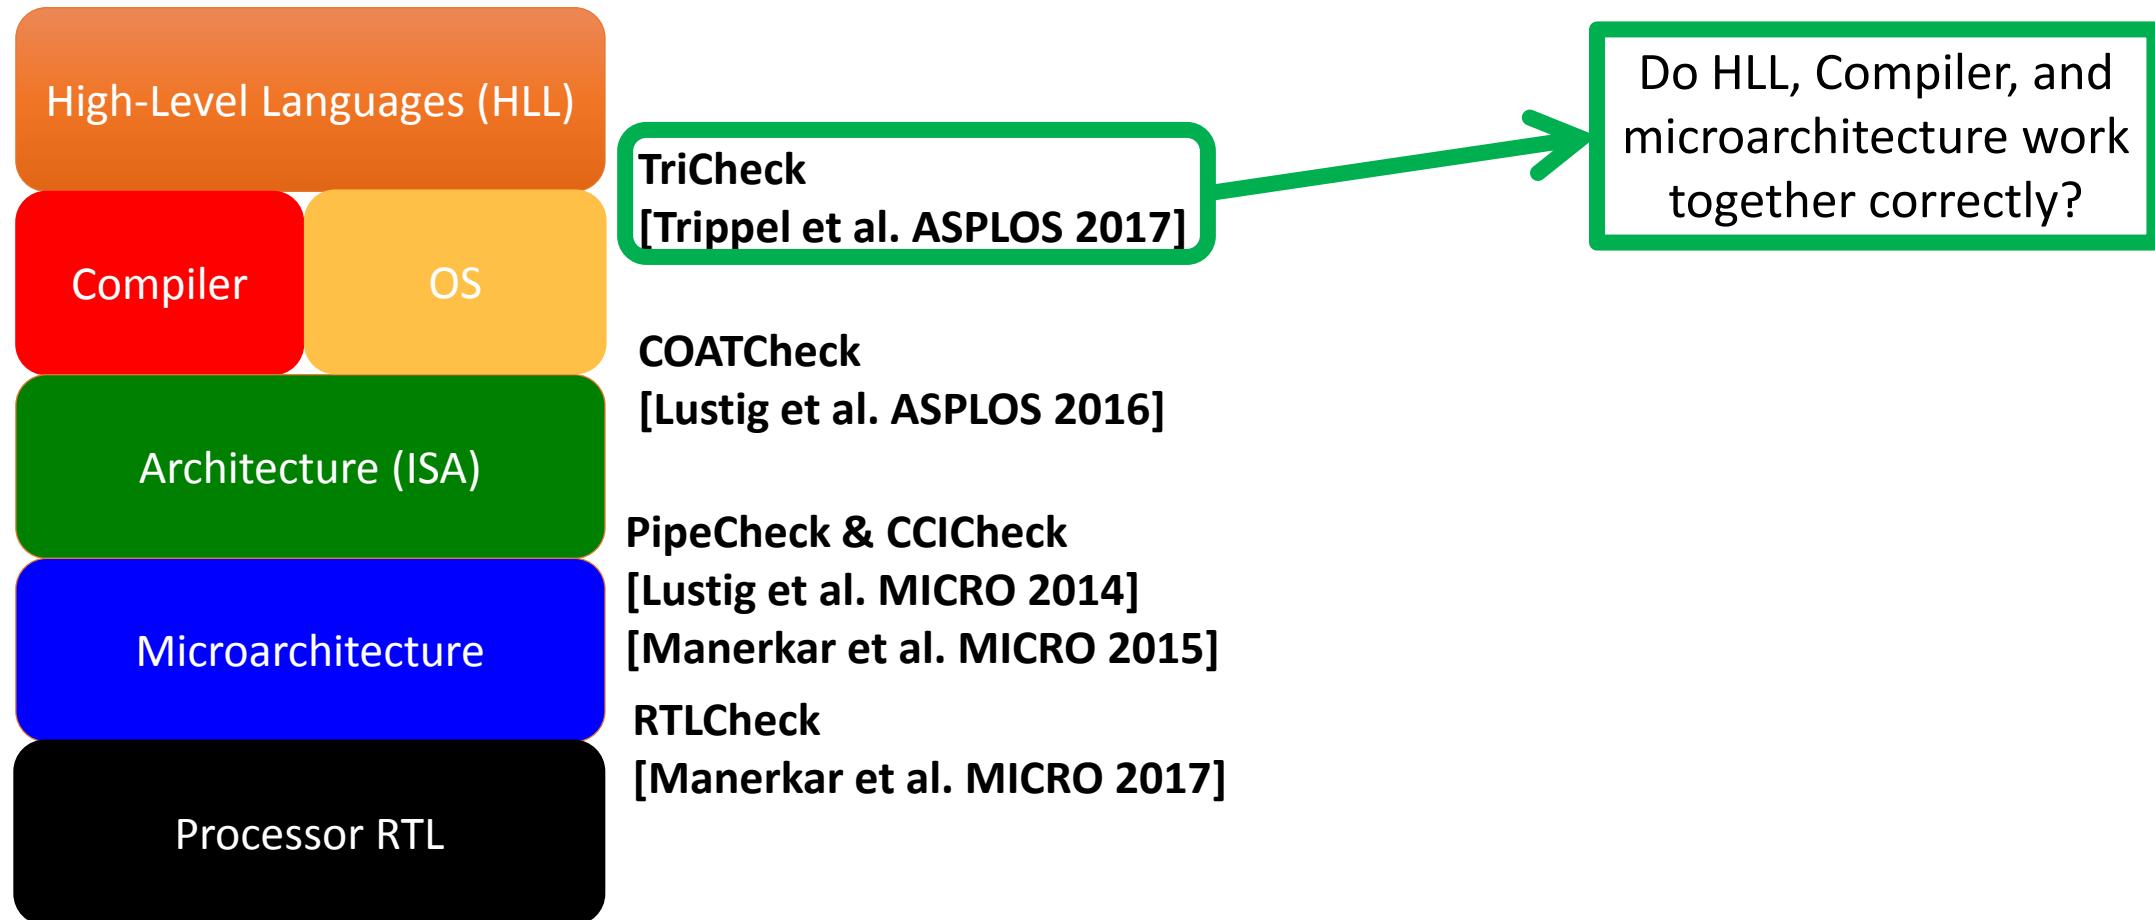
# Check Suite: Full-Stack Automated MCM Analysis



- Suite of tools at various levels of computing stack
- **Automated Full-Stack MCM checking** across litmus test suites



# Check Suite: Full-Stack Automated MCM Analysis



- Suite of tools at various levels of computing stack
- **Automated Full-Stack MCM checking** across litmus test suites



# Check Suite: Full-Stack Automated MCM Analysis

High-Level Languages (HLL)

Compiler

OS

Architecture (ISA)

Microarchitecture

Processor RTL

**TriCheck**

[Trippel et al. ASPLOS 2017]

**COATCheck**

[Lustig et al. ASPLOS 2016]

**PipeCheck & CCICheck**

[Lustig et al. MICRO 2014]

[Manerkar et al. MICRO 2015]

**RTLCheck**

[Manerkar et al. MICRO 2017]

So far, tools have found bugs in:

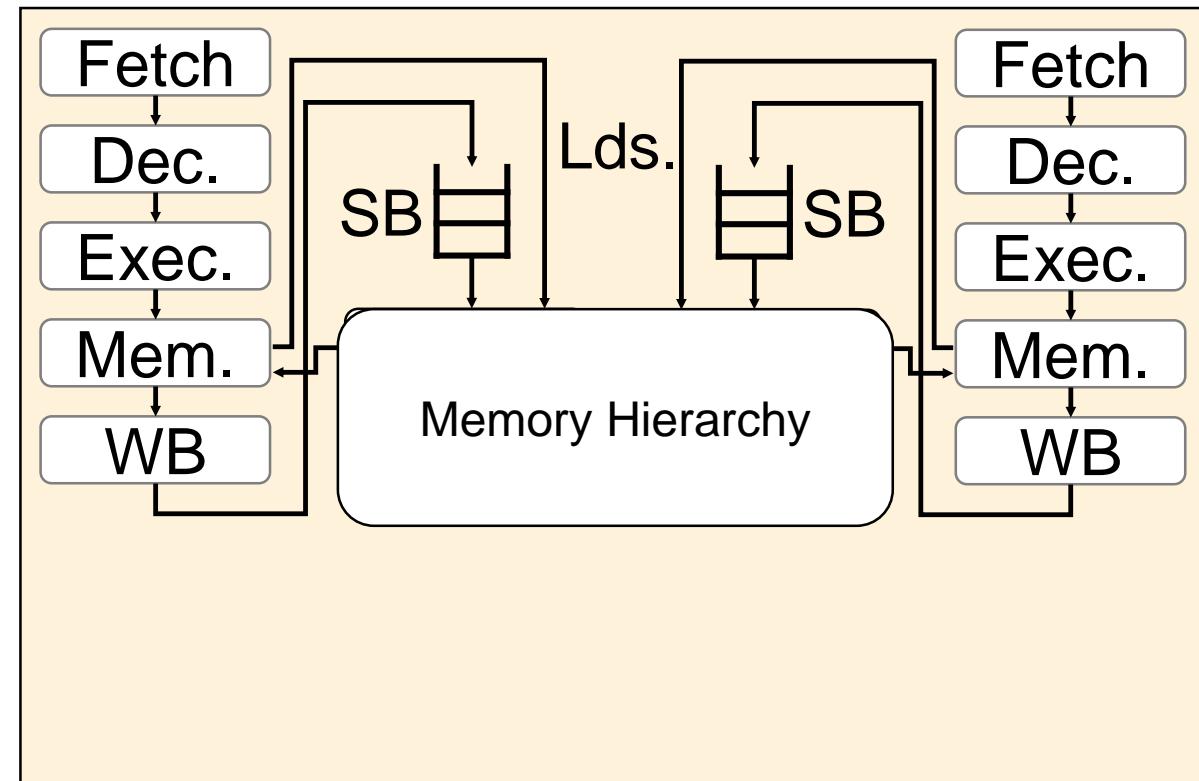
- Widely-used **gem5** Research simulator
- Cache coherence paper (**TSO-CC**)
- **IBM XL C++ compiler** (fixed in v13.1.5)
- In-design commercial processors
- **RISC-V draft ISA** specification
- Compiler mapping proofs
- **C11** memory model
- Open-source processor **RTL**

- Suite of tools at various levels of computing stack
- **Automated Full-Stack MCM checking** across litmus test suites



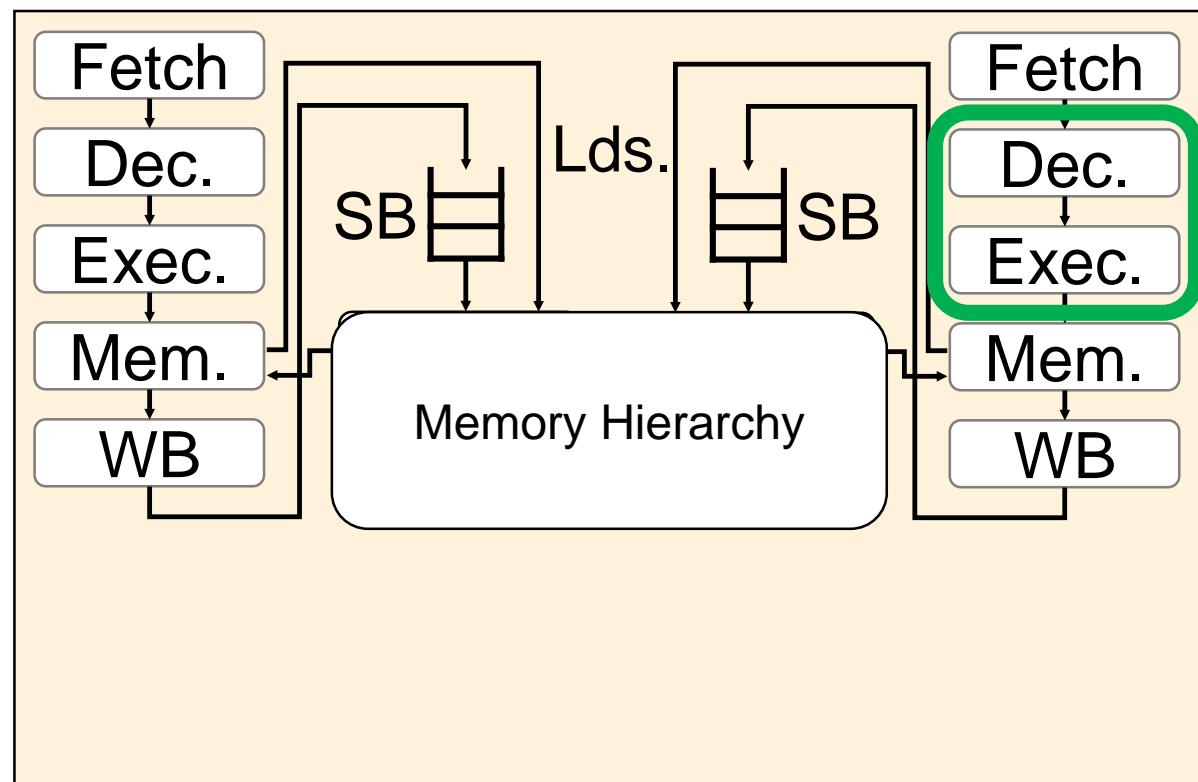
# Modelling Microarchitecture: Going below the ISA

- Hardware enforces consistency model using smaller localized orderings
  - In-order fetch/decode/execute...
  - Orderings enforced by memory hierarchy
  - ...and many more



# Modelling Microarchitecture: Going below the ISA

- Hardware enforces consistency model using smaller localized orderings
  - In-order fetch/decode/execute...
  - Orderings enforced by memory hierarchy
  - ...and many more



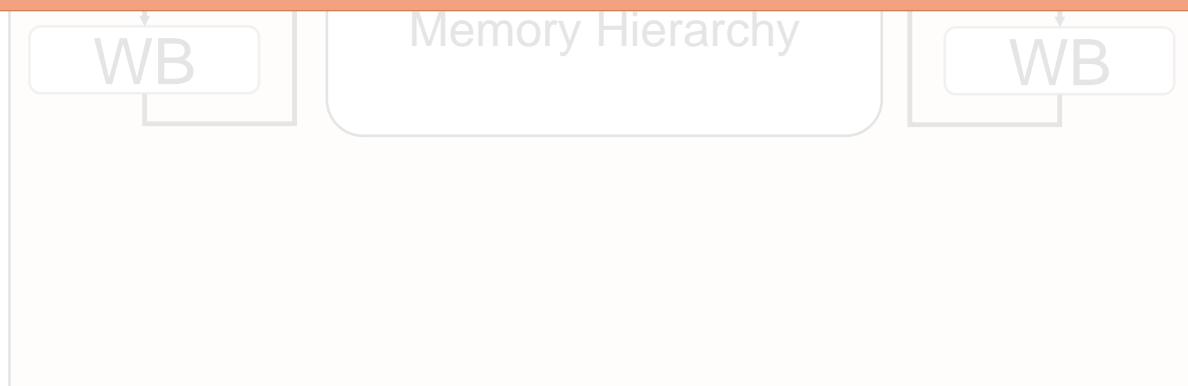
Pipeline stages  
may be FIFO to  
ensure in-order  
execution



# Modelling Microarchitecture: Going below the ISA

- Hardware enforces consistency model using smaller localized orderings
  - In-order fetch/decode/execute...

**Do individual orderings correctly work together  
to satisfy consistency model?**



Pipeline stages  
may be FIFO to  
ensure in-order  
execution



# Microarchitectural Consistency Checking

Microarchitecture in  $\mu$ spec DSL

**Axiom "Decode\_is\_FIFO":**

```
... EdgeExists ((i1, Decode), (i2, Decode))
=> AddEdge ((i1, Execute), (i2, Execute)).
```

**Axiom "PO\_Fetch":**

```
... SameCore i1 i2 /\ ProgramOrder i1 i2 =>
AddEdge ((i1, Fetch), (i2, Fetch)).
```



Litmus Test

Core 0	Core 1
(i1) $[x] \leftarrow 1$	(i3) $r1 \leftarrow [y]$
(i2) $[y] \leftarrow 1$	(i4) $r2 \leftarrow [x]$
Under SC: Forbid $r1=1, r2=0$	



# Microarchitectural Consistency Checking

Microarchitecture in µspec DSL

**Axiom "Decode\_is\_FIFO":**

```
... EdgeExists ((i1, Decode), (i2, Decode))  
      => AddEdge ((i1, Execute), (i2, Execute)).
```

**Axiom "PO\_Fetch":**

```
... SameCore i1 i2 /\ ProgramOrder i1 i2 =>  
      AddEdge ((i1, Fetch), (i2, Fetch)).
```

Each **axiom** specifies an ordering  
that µarch should respect

Core 0	Core 1
(i1) $[x] \leftarrow 1$	(i3) $r1 \leftarrow [y]$
(i2) $[y] \leftarrow 1$	(i4) $r2 \leftarrow [x]$
Under SC: Forbid $r1=1, r2=0$	



# Microarchitectural Consistency Checking

Microarchitecture in  $\mu$ spec DSL

**Axiom "Decode\_is\_FIFO":**

```
... EdgeExists ((i1, Decode), (i2, Decode))
=> AddEdge ((i1, Execute), (i2, Execute)).
```

**Axiom "PO\_Fetch":**

```
... SameCore i1 i2 /\ ProgramOrder i1 i2 =>
AddEdge ((i1, Fetch), (i2, Fetch)).
```



Litmus Test

Core 0	Core 1
(i1) $[x] \leftarrow 1$	(i3) $r1 \leftarrow [y]$
(i2) $[y] \leftarrow 1$	(i4) $r2 \leftarrow [x]$
Under SC: Forbid $r1=1, r2=0$	



# Microarchitectural Consistency Checking

Microarchitecture in  $\mu$ spec DSL

**Axiom "Decode\_is\_FIFO":**

```
... EdgeExists ((i1, Decode), (i2, Decode))  
=> AddEdge ((i1, Execute), (i2, Execute)).
```

**Axiom "PO\_Fetch":**

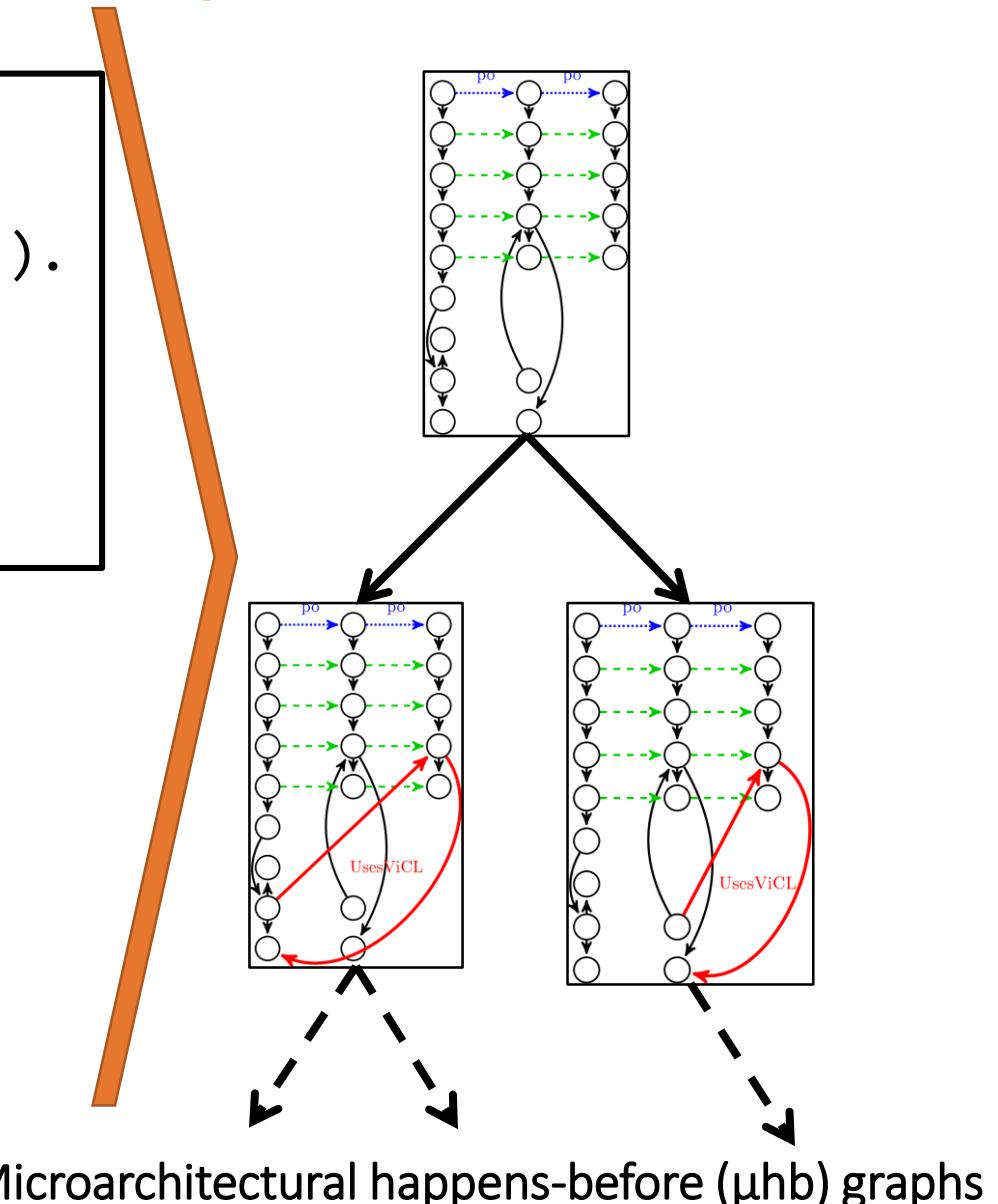
```
... SameCore i1 i2 /\ ProgramOrder i1 i2 =>  
AddEdge ((i1, Fetch), (i2, Fetch)).
```



Litmus Test

Core 0	Core 1
(i1) $[x] \leftarrow 1$	(i3) $r1 \leftarrow [y]$
(i2) $[y] \leftarrow 1$	(i4) $r2 \leftarrow [x]$

Under SC: Forbid  $r1=1, r2=0$



# Microarchitectural Consistency Checking

Microarchitecture in  $\mu$ spec DSL

**Axiom "Decode\_is\_FIFO":**

```
... EdgeExists ((i1, Decode), (i2, Decode))  
=> AddEdge ((i1, Execute), (i2, Execute)).
```

**Axiom "PO\_Fetch":**

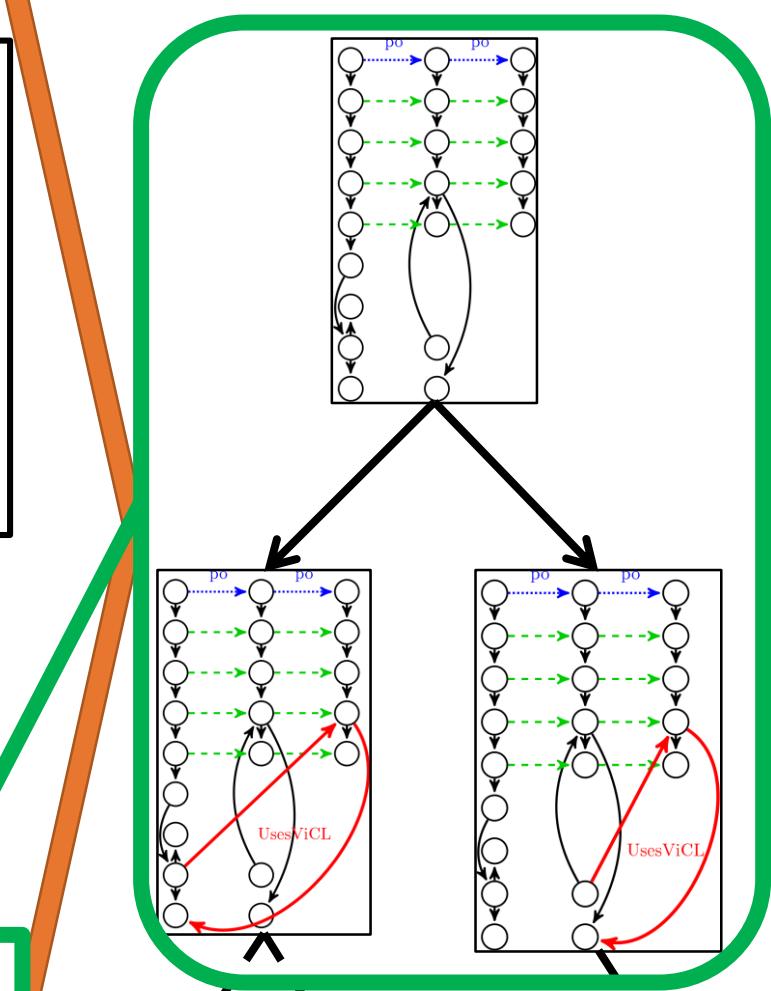
```
... SameCore i1 i2 /\ ProgramOrder i1 i2 =>  
AddEdge ((i1, Fetch), (i2, Fetch)).
```



Litmus Test

Core 0	Core 1
(i1) $[x] \leftarrow 1$	(i3) $r1 \leftarrow [y]$

Microarch. verification checks that  
combination of axioms satisfies MCM



Architectural happens-before ( $\mu$ hb) graphs



# PipeCheck: Executions as $\mu$ hb Graphs [Lustig et al. MICRO 2014]

Core 0

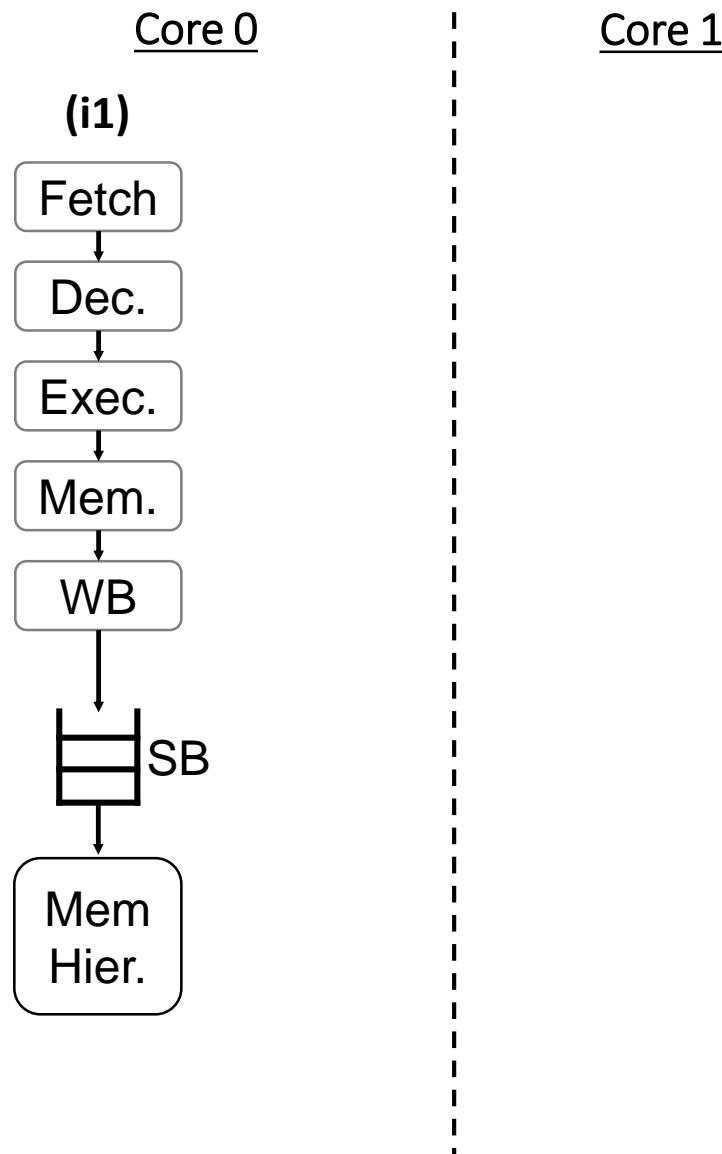
Core 1

Litmus Test **mp**

<b>Core 0</b>	<b>Core 1</b>
(i1) St $[x] \leftarrow 1$	(i3) Ld $r1 \leftarrow [y]$
(i2) St $[y] \leftarrow 1$	(i4) Ld $r2 \leftarrow [x]$
Under TSO: Forbid $r1=1, r2=0$	



# PipeCheck: Executions as $\mu$ hb Graphs [Lustig et al. MICRO 2014]



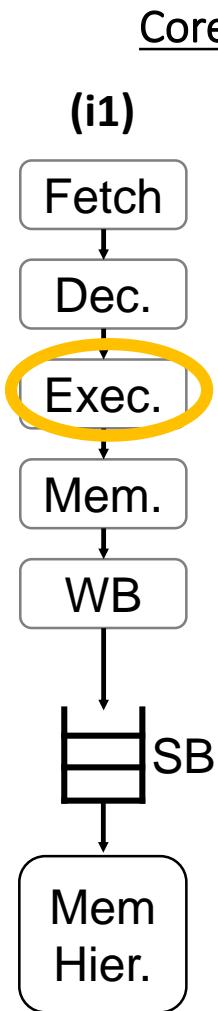
Core 1

Litmus Test **mp**

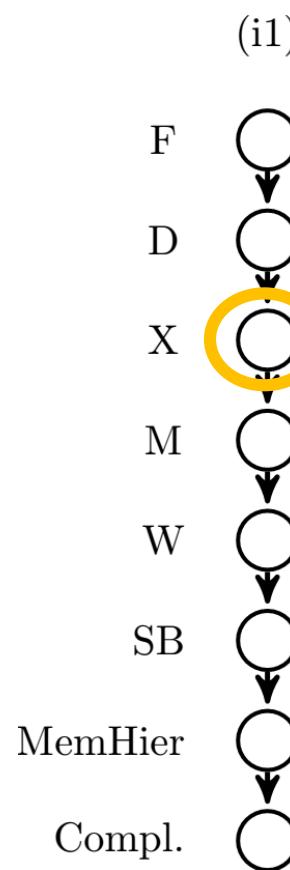
Core 0	Core 1
(i1) St $[x] \leftarrow 1$	(i3) Ld $r1 \leftarrow [y]$
(i2) St $[y] \leftarrow 1$	(i4) Ld $r2 \leftarrow [x]$
Under TSO: Forbid $r1=1, r2=0$	



# PipeCheck: Executions as $\mu$ hb Graphs [Lustig et al. MICRO 2014]



Core 1

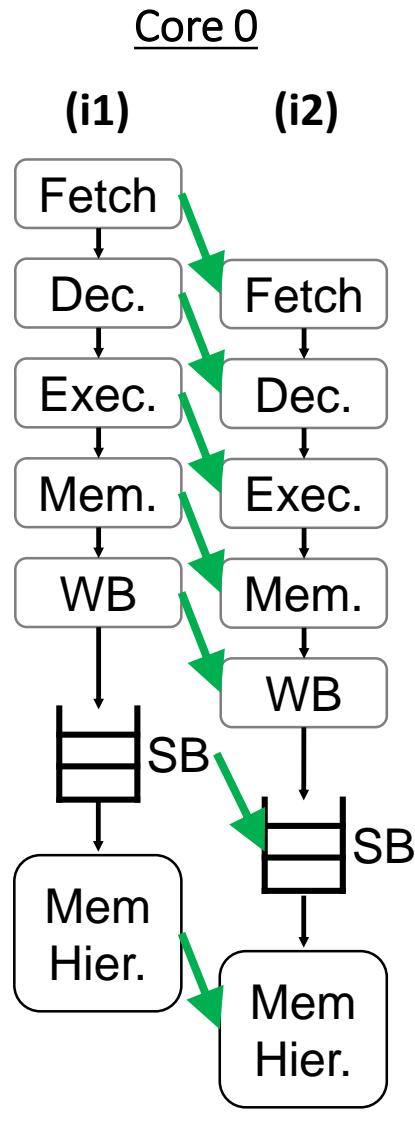


Litmus Test **mp**

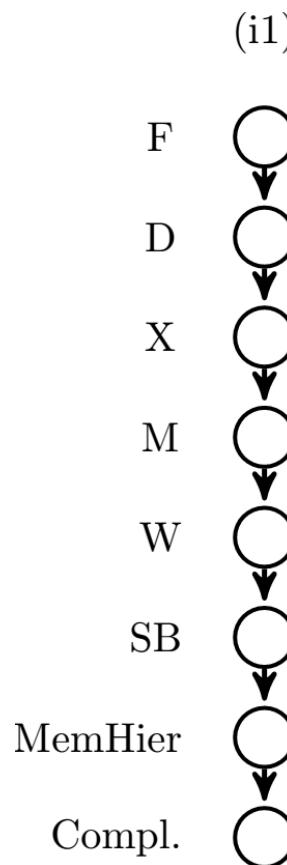
<b>Core 0</b>	<b>Core 1</b>
(i1) St $[x] \leftarrow 1$	(i3) Ld $r1 \leftarrow [y]$
(i2) St $[y] \leftarrow 1$	(i4) Ld $r2 \leftarrow [x]$
Under TSO: Forbid $r1=1, r2=0$	



# PipeCheck: Executions as $\mu$ hb Graphs [Lustig et al. MICRO 2014]



Core 1

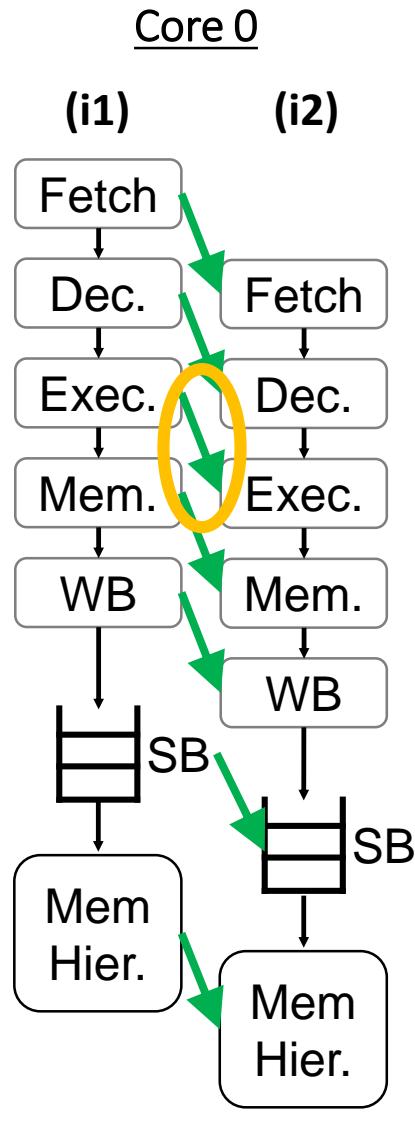


Litmus Test **mp**

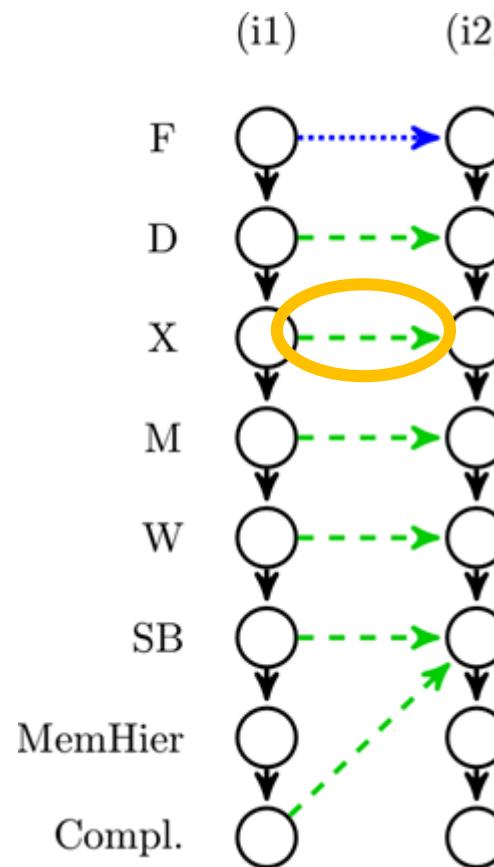
<b>Core 0</b>	<b>Core 1</b>
(i1) St $[x] \leftarrow 1$	(i3) Ld $r1 \leftarrow [y]$
(i2) St $[y] \leftarrow 1$	(i4) Ld $r2 \leftarrow [x]$
Under TSO: Forbid $r1=1, r2=0$	



# PipeCheck: Executions as $\mu$ hb Graphs [Lustig et al. MICRO 2014]



Core 1

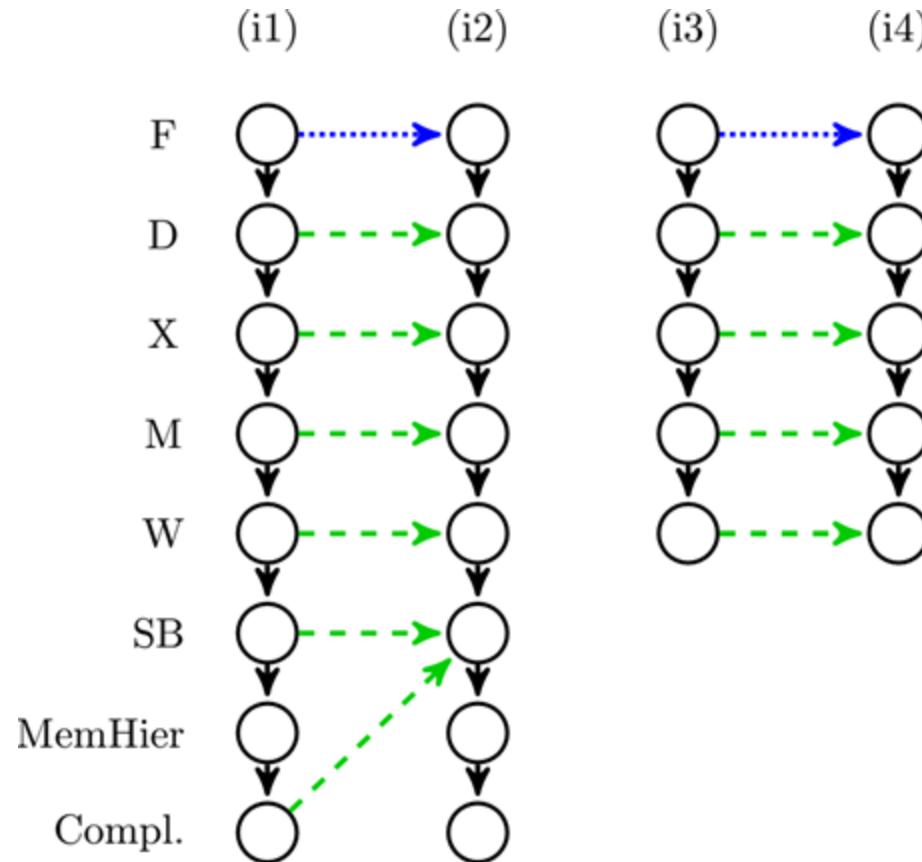
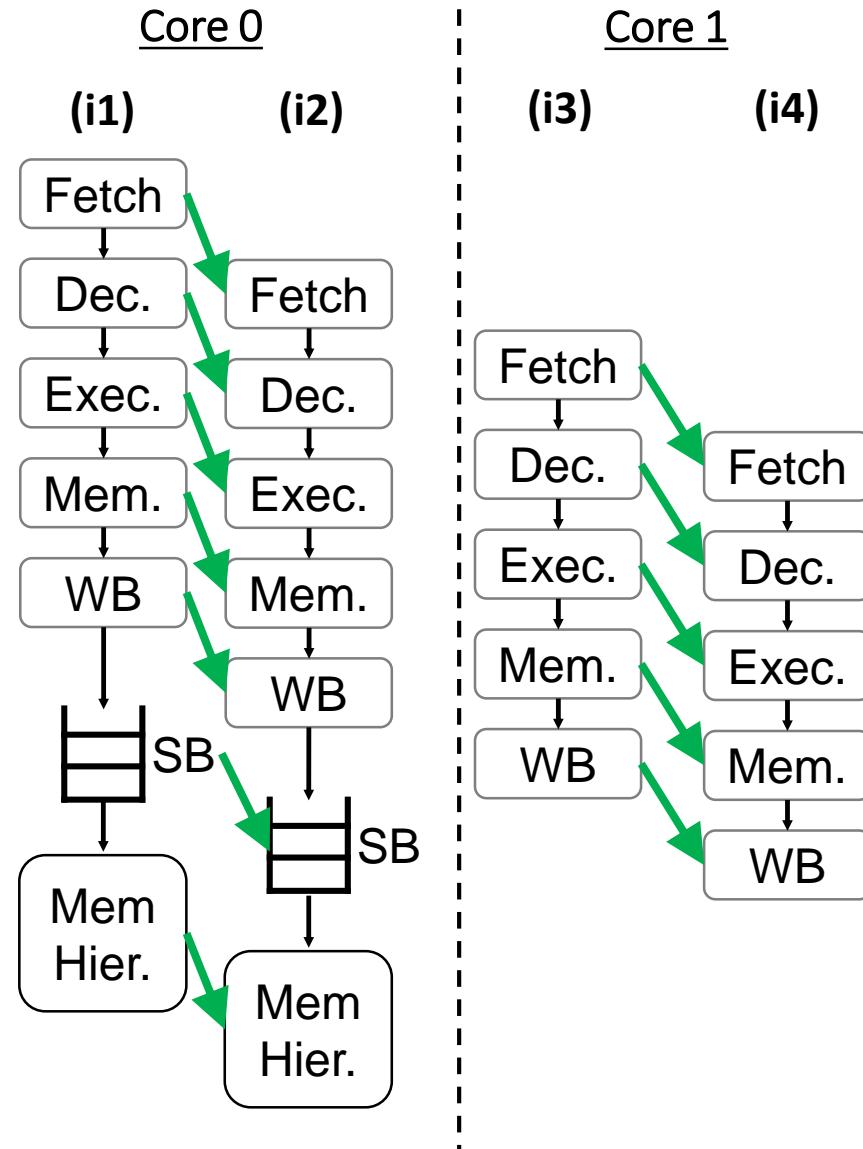


Litmus Test **mp**

Core 0	Core 1
(i1) St $[x] \leftarrow 1$	(i3) Ld $r1 \leftarrow [y]$
(i2) St $[y] \leftarrow 1$	(i4) Ld $r2 \leftarrow [x]$
Under TSO: Forbid $r1=1, r2=0$	



# PipeCheck: Executions as μhb Graphs [Lustig et al. MICRO 2014]

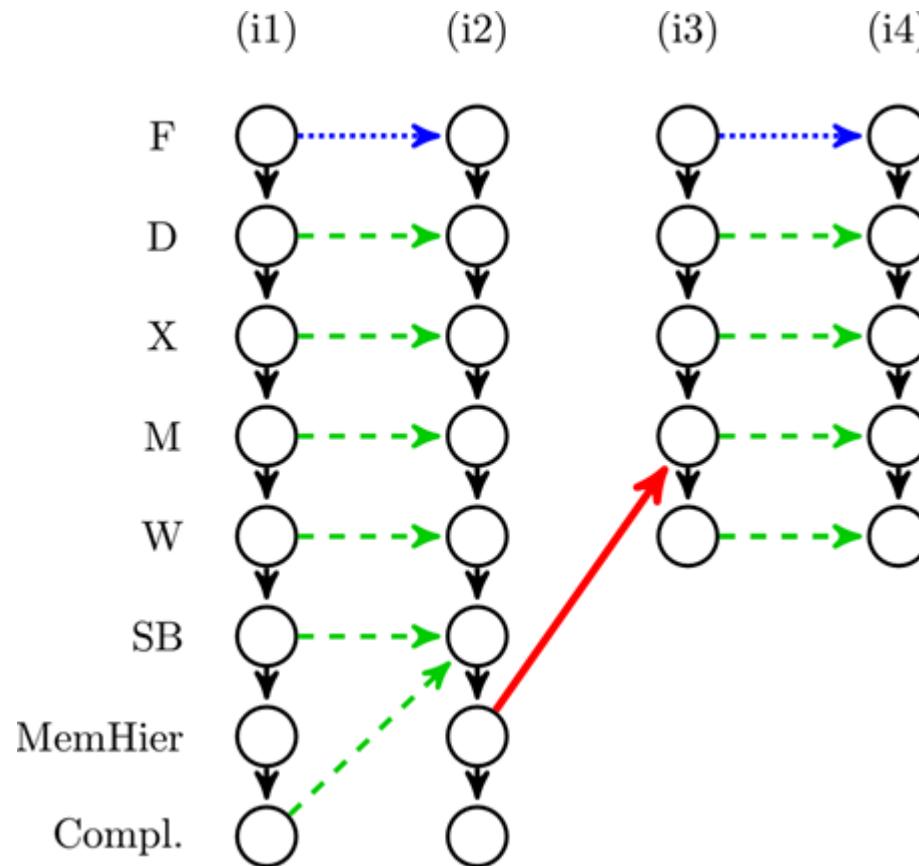
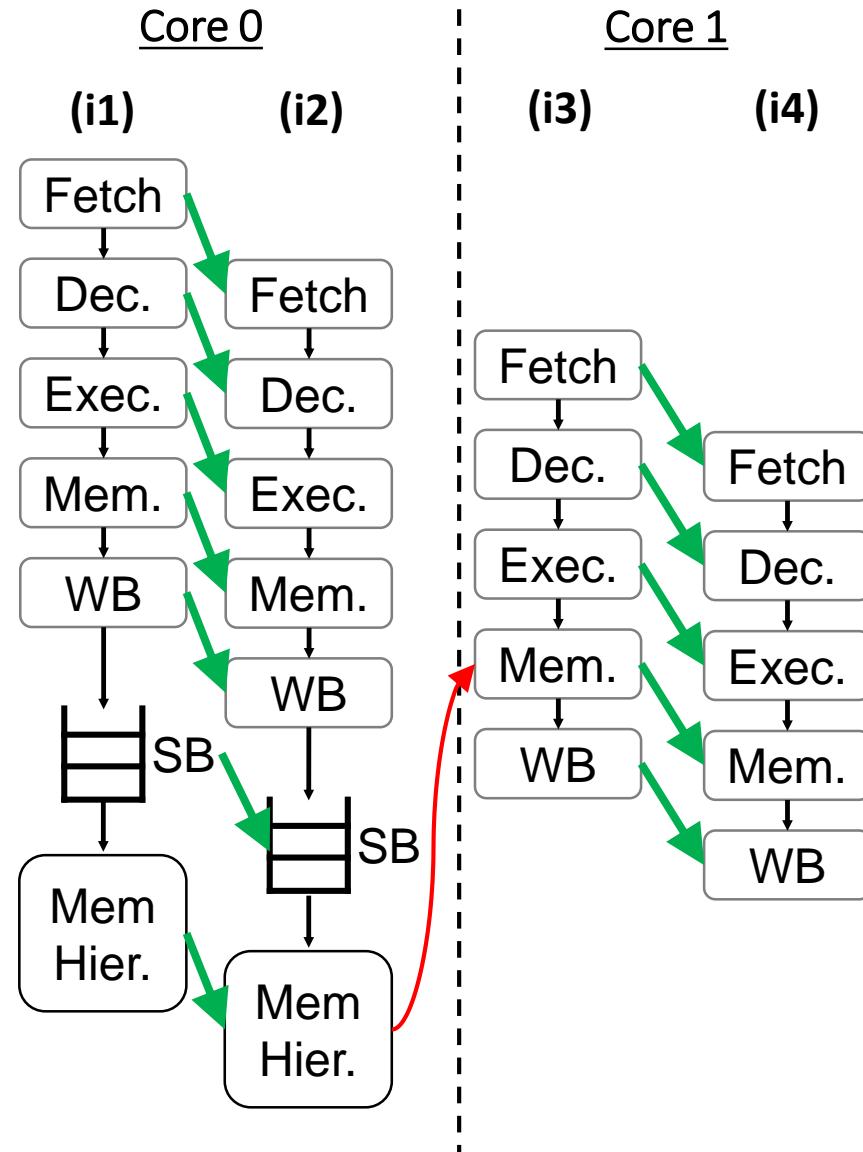


Litmus Test **mp**

Core 0	Core 1
(i1) St $[x] \leftarrow 1$	(i3) Ld $r_1 \leftarrow [y]$
(i2) St $[y] \leftarrow 1$	(i4) Ld $r_2 \leftarrow [x]$
Under TSO: Forbid $r_1=1, r_2=0$	



# PipeCheck: Executions as $\mu$ hb Graphs [Lustig et al. MICRO 2014]



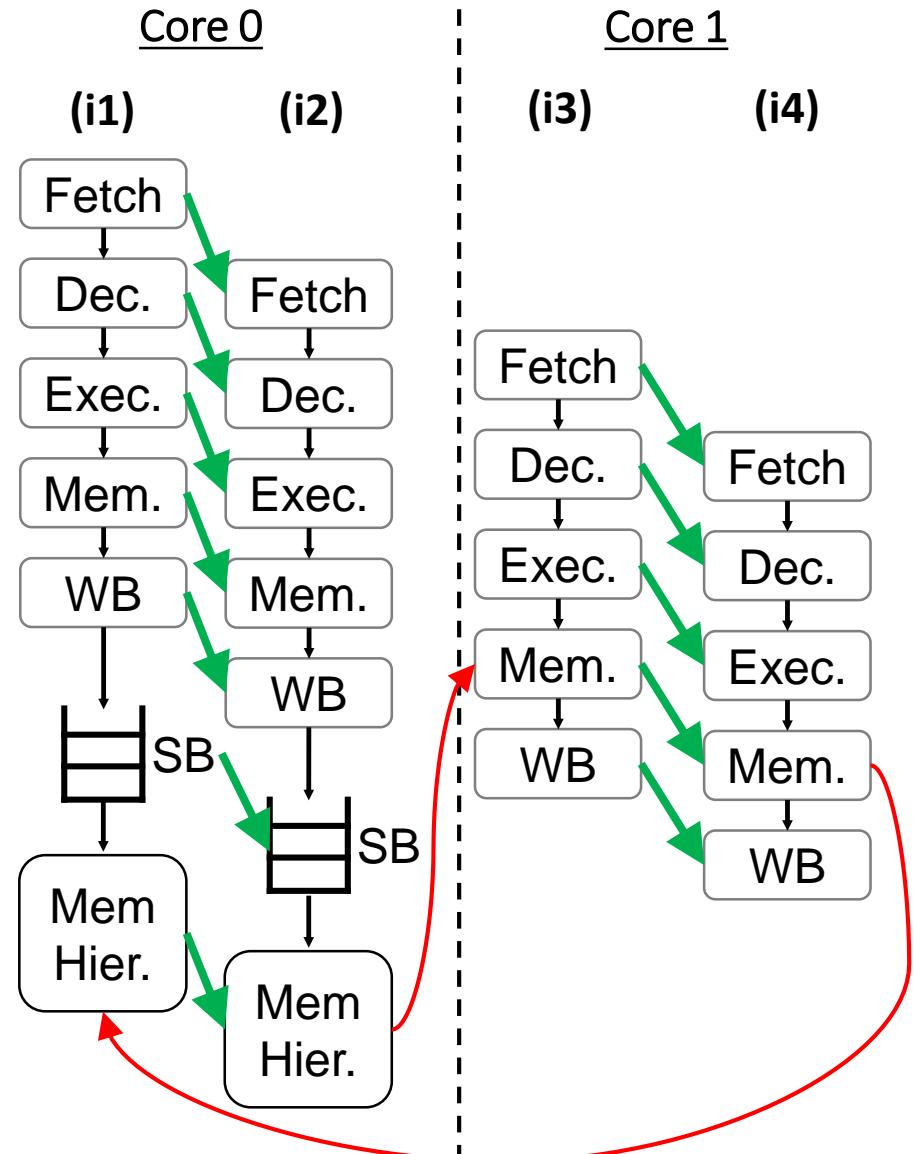
Litmus Test **mp**

Core 0	Core 1
(i1) St $[x] \leftarrow 1$	(i3) Ld $r_1 \leftarrow [y]$
(i2) St $[y] \leftarrow 1$	(i4) Ld $r_2 \leftarrow [x]$

Under TSO: Forbid  $r_1=1, r_2=0$

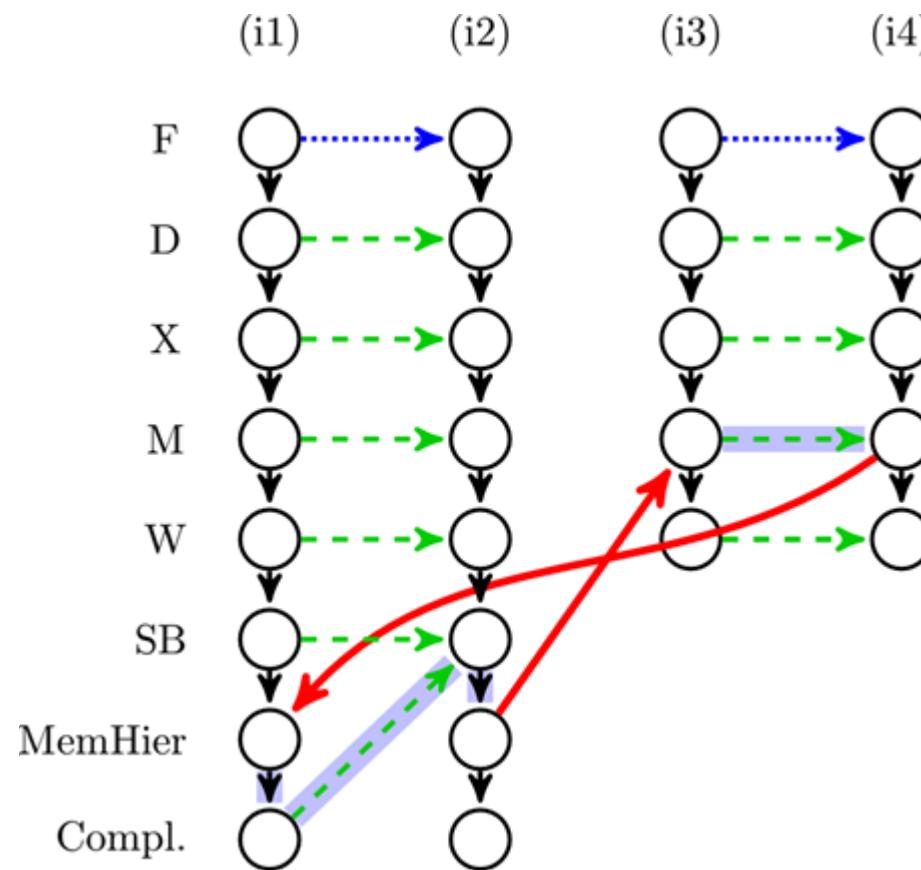


# PipeCheck: Executions as μhb Graphs [Lustig et al. MICRO 2014]



Core 1

(i3)      (i4)



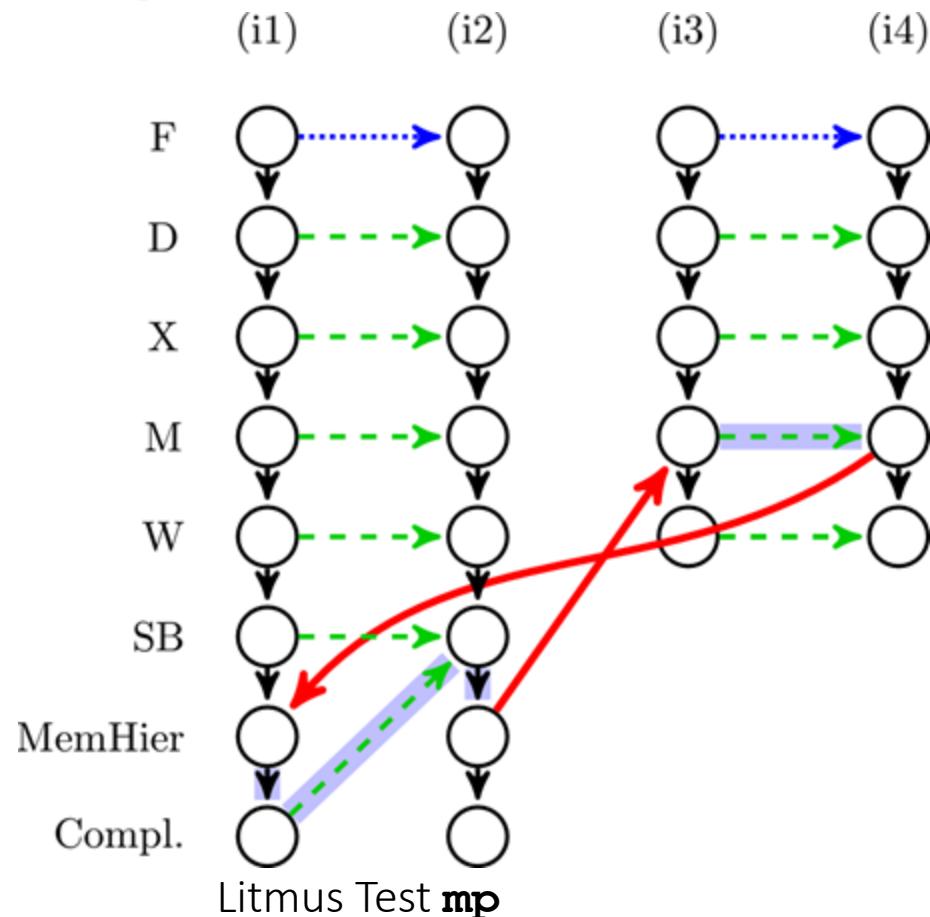
Litmus Test **mp**

Core 0	Core 1
(i1) St $[x] \leftarrow 1$	(i3) Ld $r1 \leftarrow [y]$
(i2) St $[y] \leftarrow 1$	(i4) Ld $r2 \leftarrow [x]$

Under TSO: Forbid  $r1=1, r2=0$



# PipeCheck: Microarchitectural Correctness

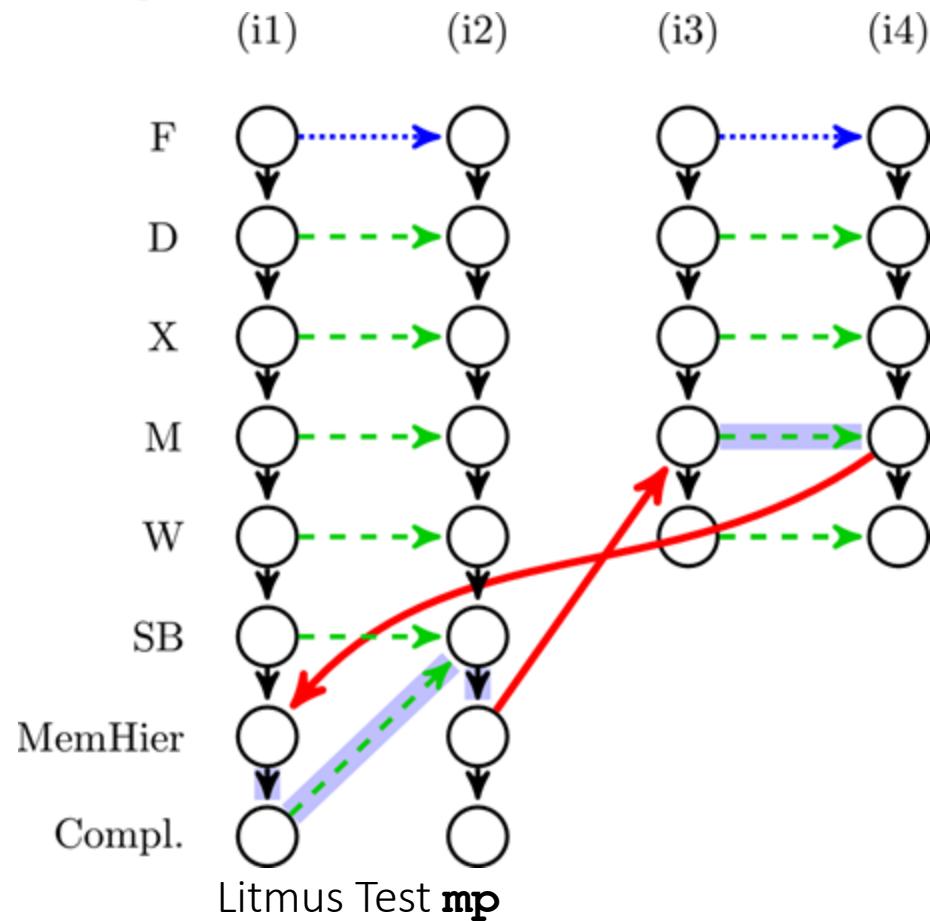


- Cycle in  $\mu$ hb graph => event has to happen before itself (impossible)
- **Cyclic** graph → **unobservable** on  $\mu$ arch
- **Acyclic** graph → **observable** on  $\mu$ arch
- Exhaustively enumerate and check all possible execs of litmus test on  $\mu$ arch
  - Implemented using fast SMT solvers
  - Compare against ISA-level outcome from **herd** [Alglave et al. TOPLAS 2014]

Core 0	Core 1
(i1) $[x] \leftarrow 1$	(i3) $r1 \leftarrow [y]$
(i2) $[y] \leftarrow 1$	(i4) $r2 \leftarrow [x]$
Under SC: Forbid $r1=1, r2=0$	



# PipeCheck: Microarchitectural Correctness



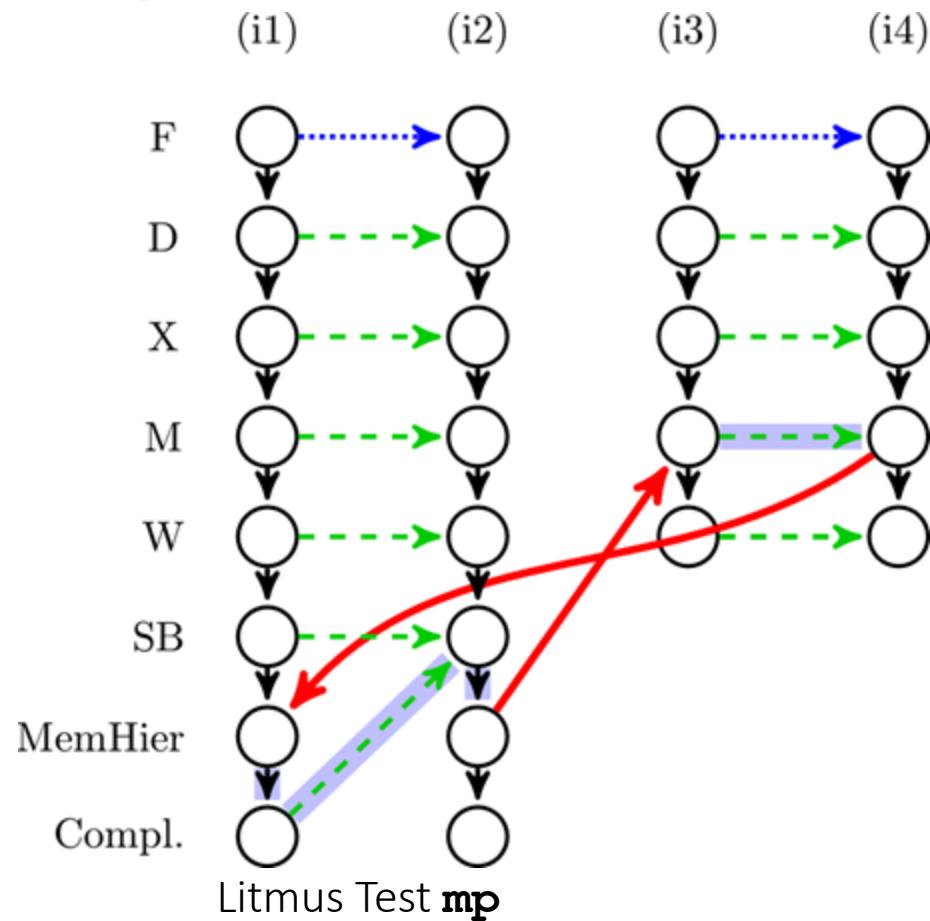
Core 0	Core 1
(i1) $[x] \leftarrow 1$	(i3) $r1 \leftarrow [y]$
(i2) $[y] \leftarrow 1$	(i4) $r2 \leftarrow [x]$
Under SC: Forbid $r1=1, r2=0$	

- Cycle in  $\mu$ hb graph => event has to happen before itself (impossible)
- **Cyclic** graph → **unobservable** on  $\mu$ arch
- **Acyclic** graph → **observable** on  $\mu$ arch
- Exhaustively enumerate and check all possible execs of litmus test on  $\mu$ arch
  - Implemented using fast SMT solvers
  - Compare against ISA-level outcome from **herd** [Alglave et al. TOPLAS 2014]

ISA-Level Outcome	Observable ( $\geq 1$ Graph Acyclic)	Not Observable (All Graphs Cyclic)
Allowed	OK	OK (stricter than necessary)
Forbidden	Consistency violation!	OK



# PipeCheck: Microarchitectural Correctness



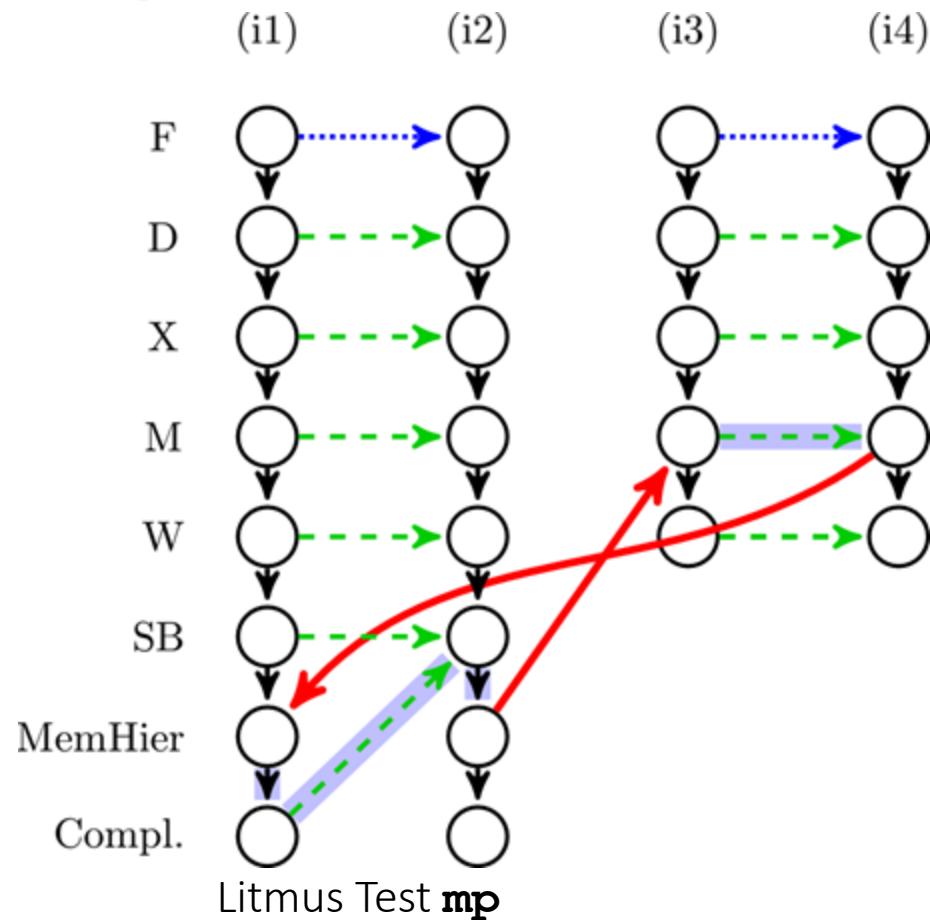
Core 0	Core 1
(i1) [x] ← 1	(i3) r1 ← [y]
(i2) [y] ← 1	(i4) r2 ← [x]
Under SC: Forbid r1=1, r2=0	

- Cycle in  $\mu$ hb graph => event has to happen before itself (impossible)
- **Cyclic** graph → **unobservable** on  $\mu$ arch
- **Acyclic** graph → **observable** on  $\mu$ arch
- Exhaustively enumerate and check all possible execs of litmus test on  $\mu$ arch
  - Implemented using fast SMT solvers
  - Compare against ISA-level outcome from **herd** [Alglave et al. TOPLAS 2014]

ISA-Level Outcome	Observable ( $\geq 1$ Graph Acyclic)	Not Observable (All Graphs Cyclic)
Allowed	OK	OK (stricter than necessary)
Forbidden	Consistency violation!	OK



# PipeCheck: Microarchitectural Correctness



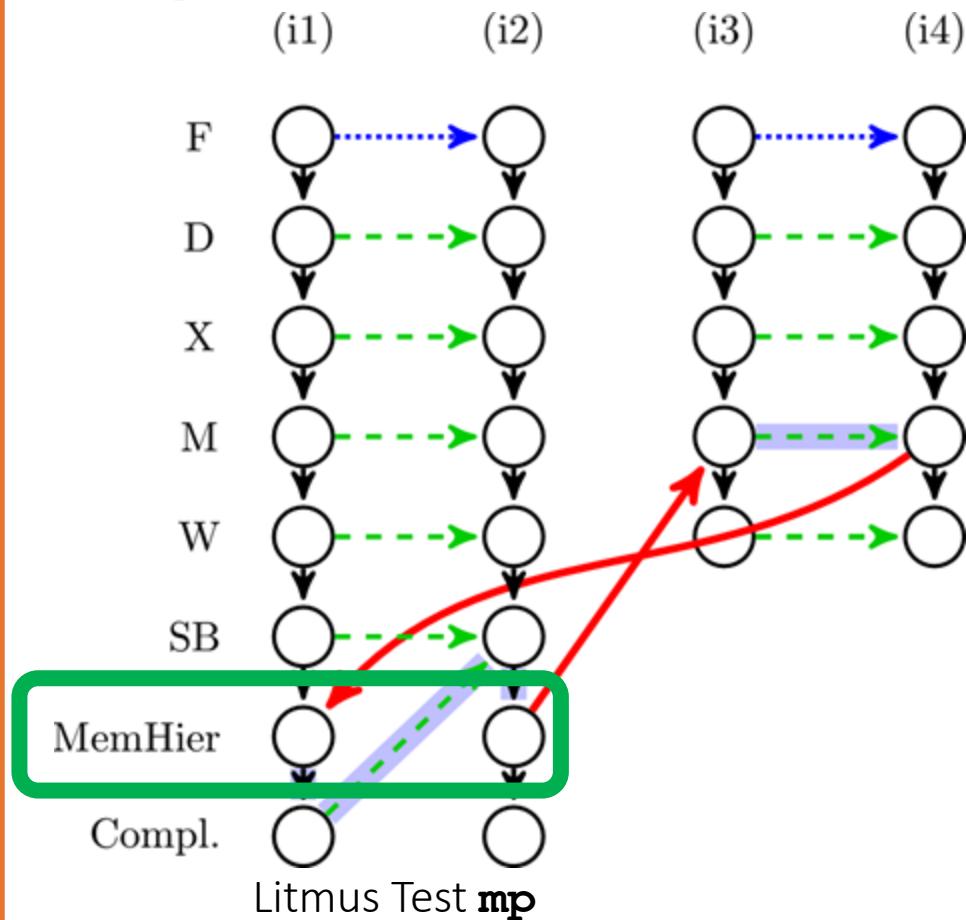
Core 0	Core 1
(i1) $[x] \leftarrow 1$	(i3) $r1 \leftarrow [y]$
(i2) $[y] \leftarrow 1$	(i4) $r2 \leftarrow [x]$
Under SC: Forbid $r1=1, r2=0$	

- Cycle in  $\mu$ hb graph => event has to happen before itself (impossible)
- **Cyclic** graph → **unobservable** on  $\mu$ arch
- **Acyclic** graph → **observable** on  $\mu$ arch
- Exhaustively enumerate and check all possible execs of litmus test on  $\mu$ arch
  - Implemented using fast SMT solvers
  - Compare against ISA-level outcome from **herd** [Alglave et al. TOPLAS 2014]

ISA-Level Outcome	Observable ( $\geq 1$ Graph Acyclic)	Not Observable (All Graphs Cyclic)
Allowed	OK	OK (stricter than necessary)
Forbidden	Consistency violation!	OK



# PipeCheck: Microarchitectural Correctness



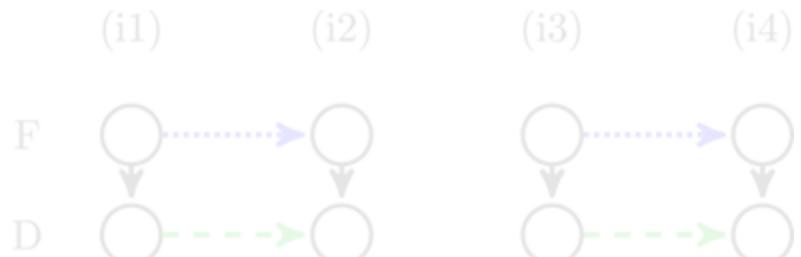
Core 0	Core 1
(i1) $[x] \leftarrow 1$	(i3) $r1 \leftarrow [y]$
(i2) $[y] \leftarrow 1$	(i4) $r2 \leftarrow [x]$
Under SC: Forbid $r1=1, r2=0$	

- Cycle in  $\mu$ hb graph => event has to happen before itself (impossible)
- **Cyclic** graph → **unobservable** on  $\mu$ arch
- **Acyclic** graph → **observable** on  $\mu$ arch
- Exhaustively enumerate and check all possible execs of litmus test on  $\mu$ arch
  - Implemented using fast SMT solvers
  - Compare against ISA-level outcome from **herd** [Alglave et al. TOPLAS 2014]

ISA-Level Outcome	Observable ( $\geq 1$ Graph Acyclic)	Not Observable (All Graphs Cyclic)
Allowed	OK	OK (stricter than necessary)
Forbidden	Consistency violation!	OK



# PipeCheck: Microarchitectural Correctness



- Cycle in  $\mu$ hb graph => event has to happen before itself (impossible)
- Cyclic graph → unobservable on uarch

Abstracted memory hierarchy prevents  
verification of complex coherence issues!

Compr.		Litmus Test	mp
Core 0	Core 1		
(i1) $[x] \leftarrow 1$	(i3) $r1 \leftarrow [y]$		
(i2) $[y] \leftarrow 1$	(i4) $r2 \leftarrow [x]$		
Under SC: Forbid r1=1, r2=0			

ICFP [Bogomolov et al., TOPLAS 2014]

ISA-Level Outcome	Observable ( $\geq 1$ Graph Acyclic)	Not Observable (All Graphs Cyclic)
Allowed	OK	OK (stricter than necessary)
Forbidden	Consistency violation!	OK



# CCICheck: Coherence vs Consistency

High-Level Languages (HLL)

Compiler

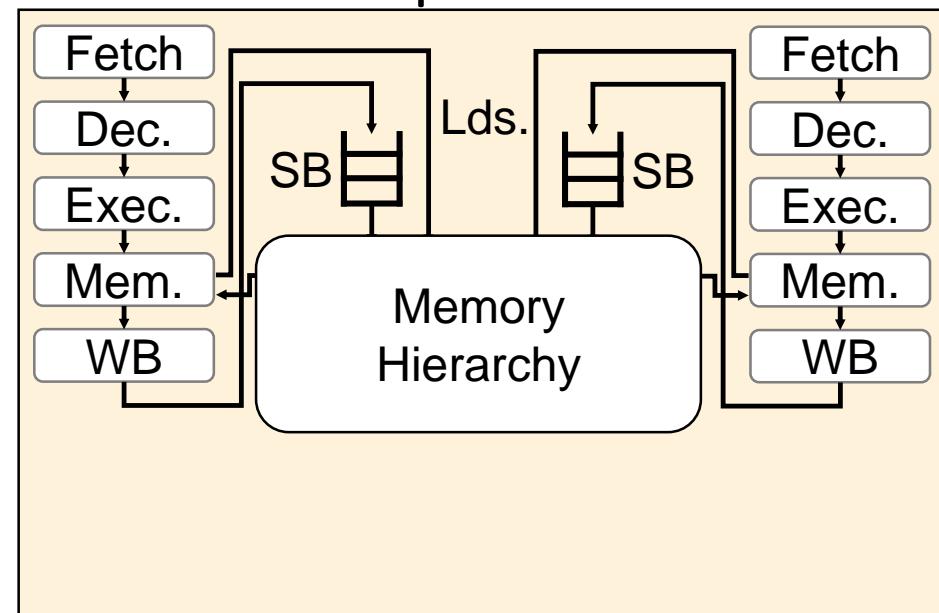
OS

Architecture (ISA)

Microarchitecture

Processor RTL

- Memory hierarchy is a collection of caches
  - Coherence protocols ensure that all caches agree on the value of any variable
- **CCICheck [Manerkar et al. MICRO 2015]** shows that consistency verification often cannot simply treat memory hierarchy abstractly
  - Nominated for Best Paper at MICRO 2015



# CCICheck: Coherence vs Consistency

High-Level Languages (HLL)

Compiler

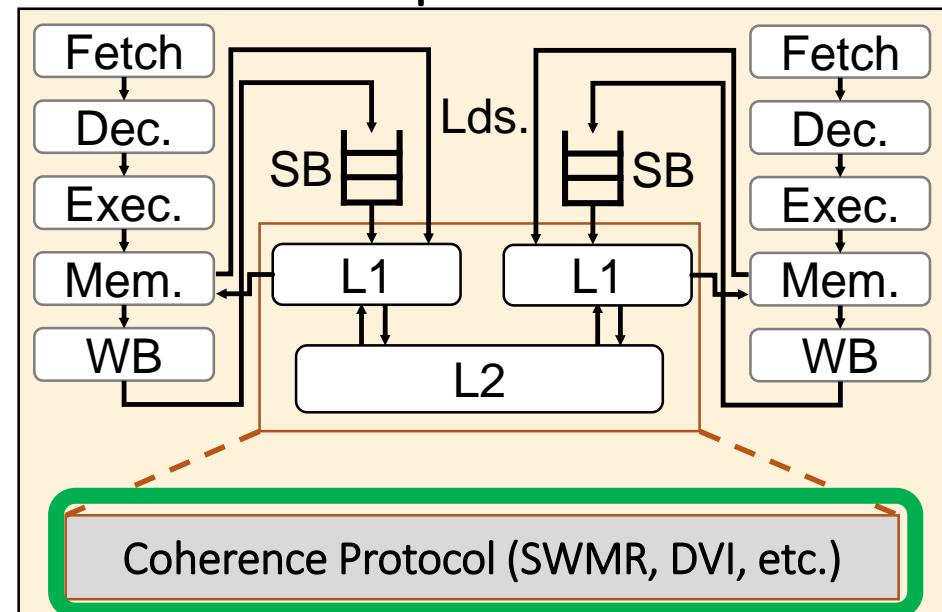
OS

Architecture (ISA)

Microarchitecture

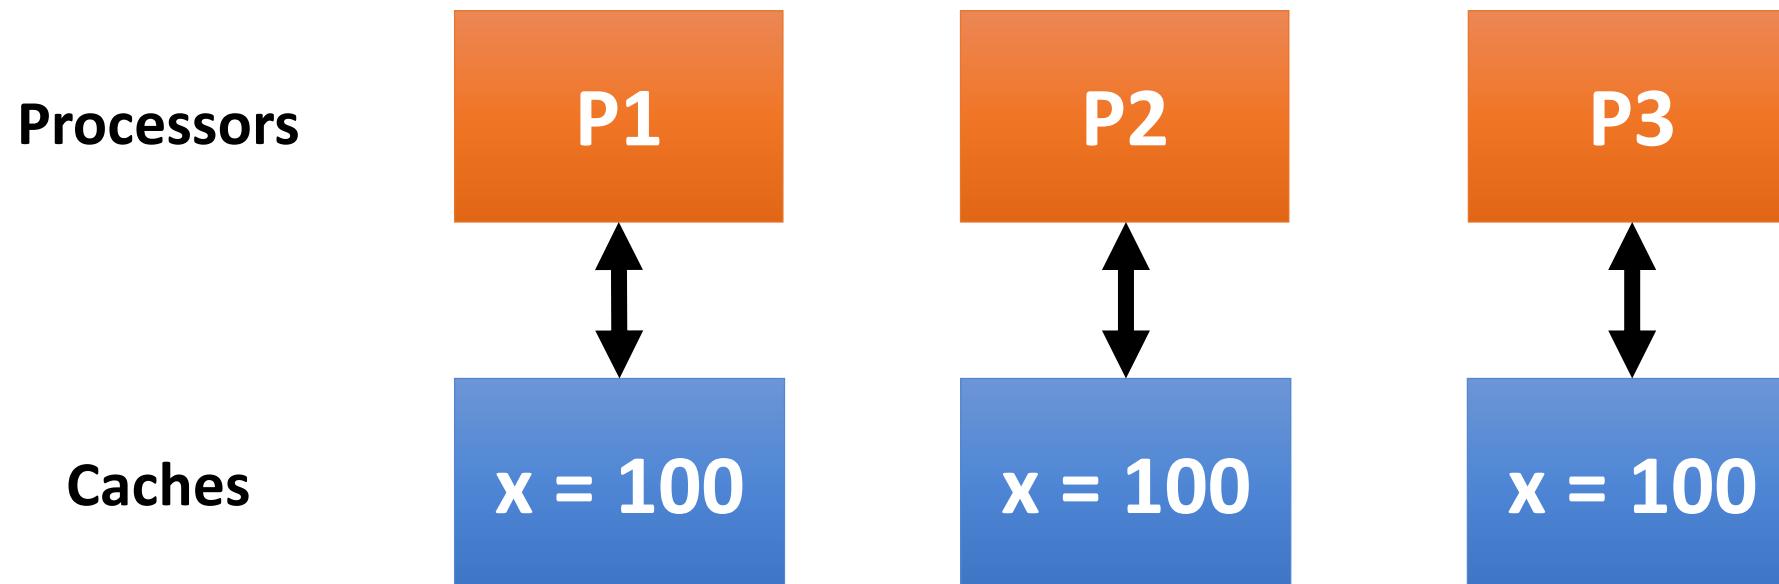
Processor RTL

- Memory hierarchy is a collection of caches
  - Coherence protocols ensure that all caches agree on the value of any variable
- **CCICheck [Manerkar et al. MICRO 2015]** shows that consistency verification often cannot simply treat memory hierarchy abstractly
  - Nominated for Best Paper at MICRO 2015



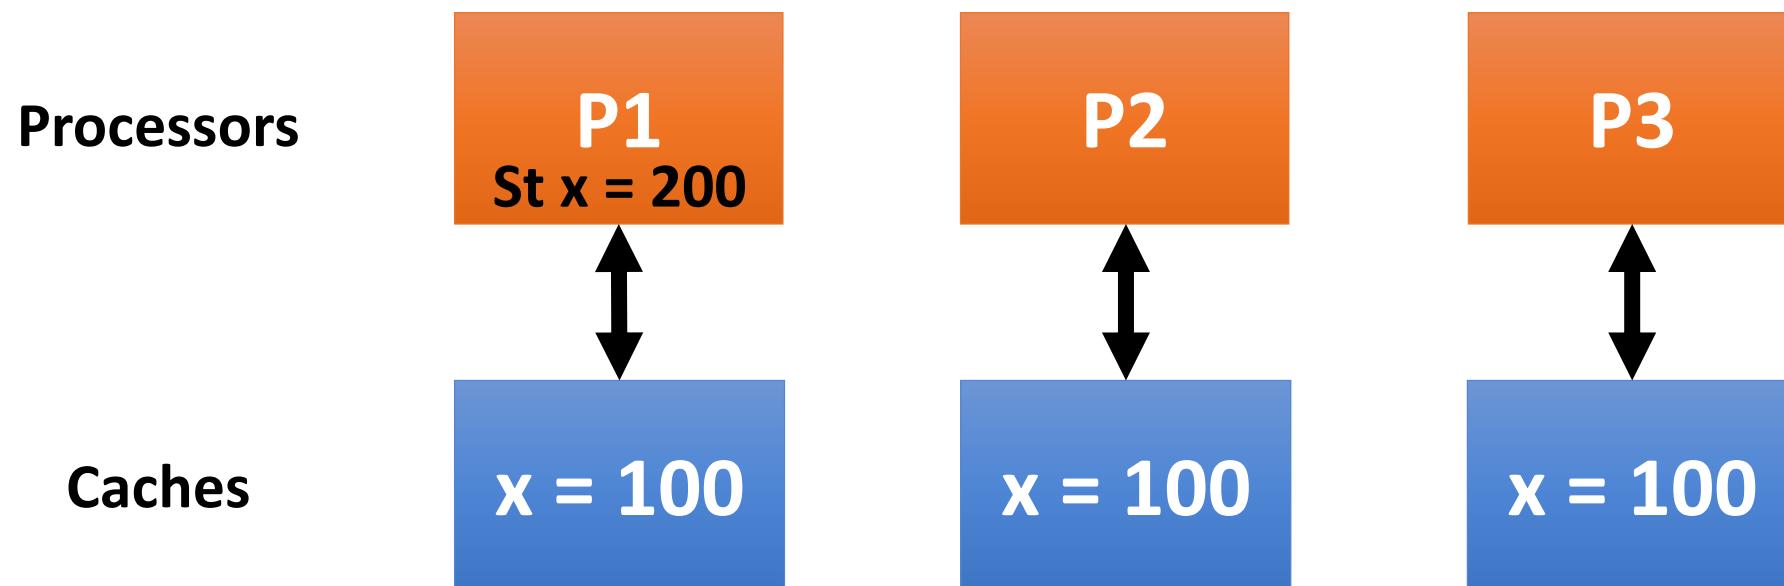
# Coherence Protocol Example

- If P1 updates the value of x to 200, the stale value of x in other processors must be **invalidated**
- If P3 wants to subsequently read/write x, it must request the new value
- **SWMR** = Single-Writer Multiple Readers, **DVI** = Data Value Invariant



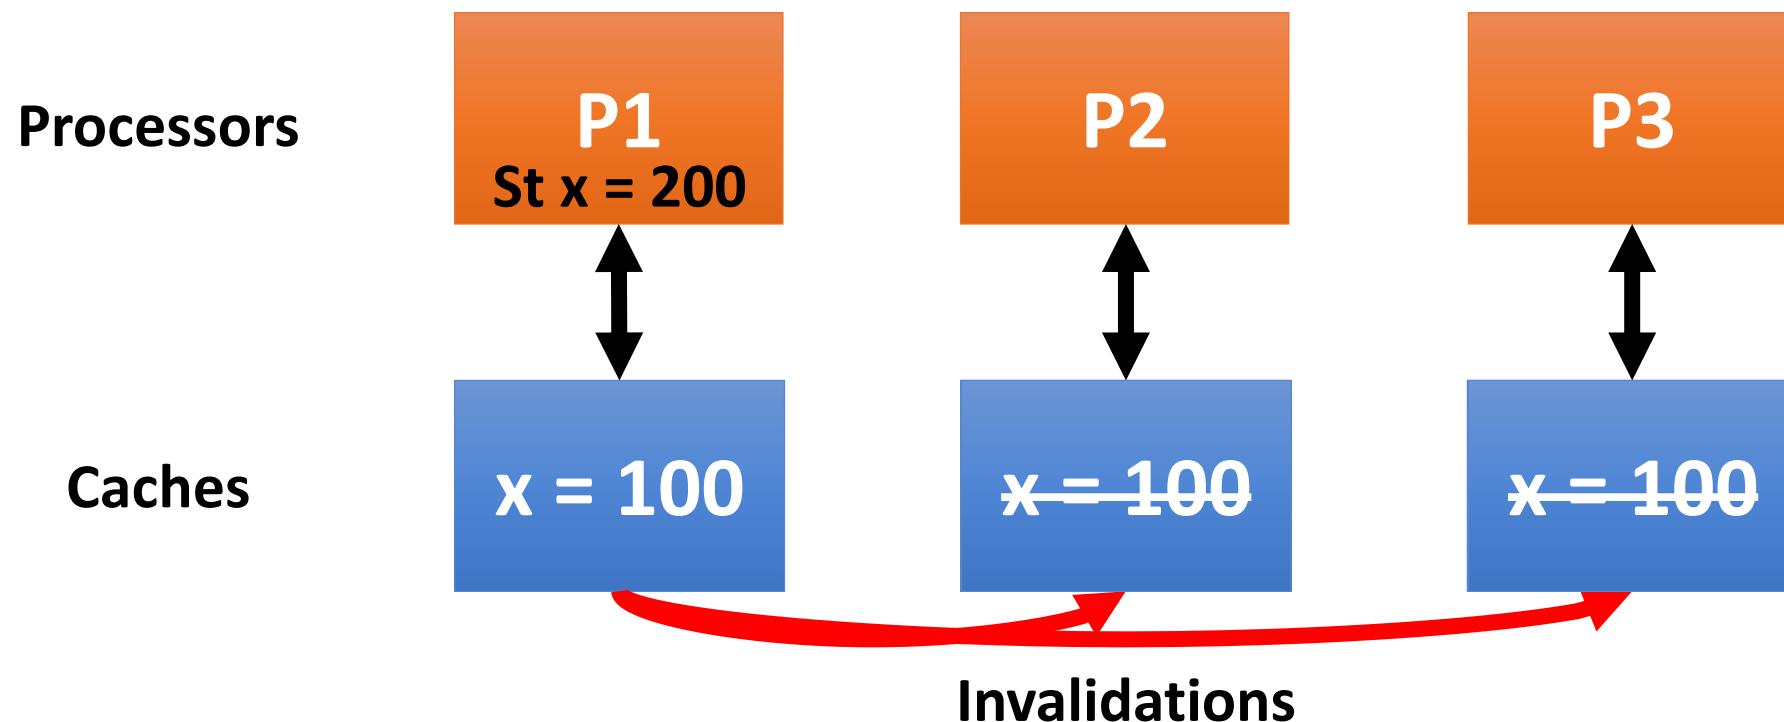
# Coherence Protocol Example

- If P1 updates the value of x to 200, the stale value of x in other processors must be **invalidated**
- If P3 wants to subsequently read/write x, it must request the new value
- **SWMR** = Single-Writer Multiple Readers, **DVI** = Data Value Invariant



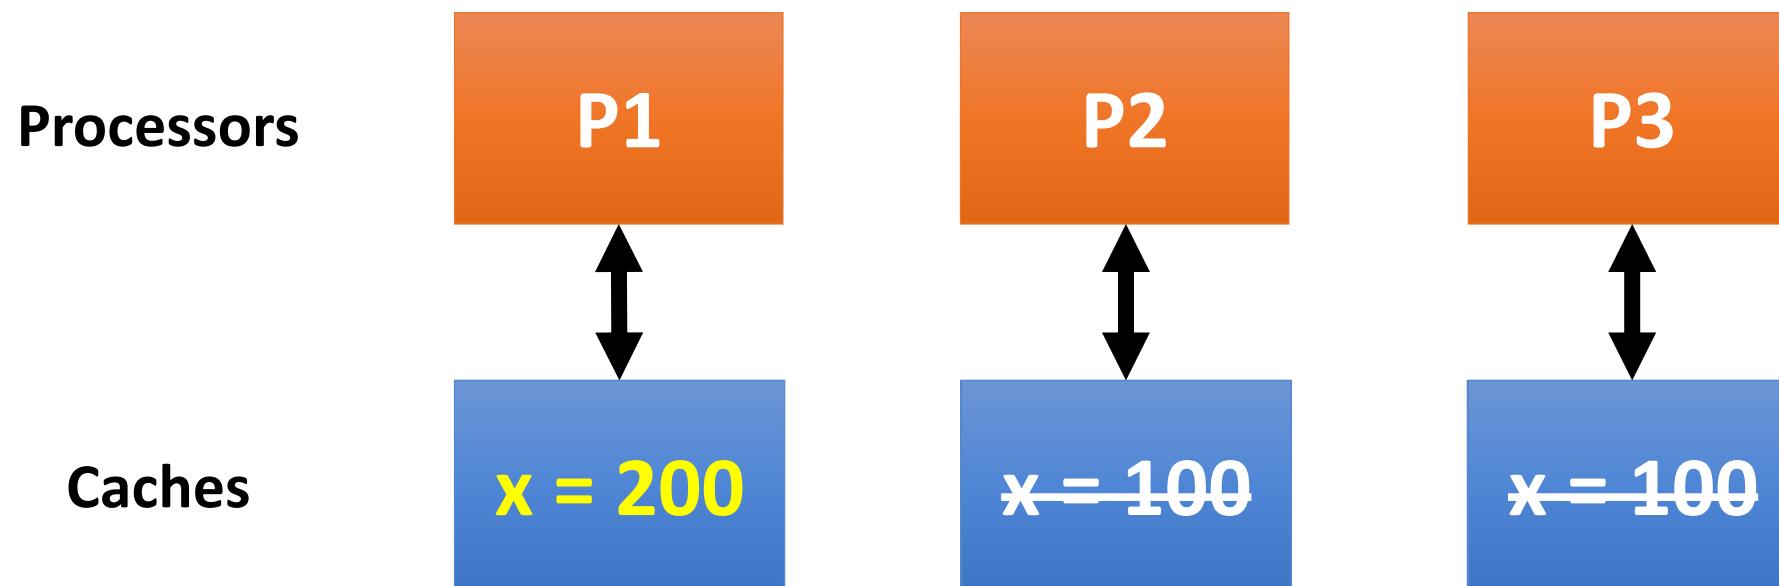
# Coherence Protocol Example

- If P1 updates the value of  $x$  to 200, the stale value of  $x$  in other processors must be **invalidated**
- If P3 wants to subsequently read/write  $x$ , it must request the new value
- **SWMR** = Single-Writer Multiple Readers, **DVI** = Data Value Invariant



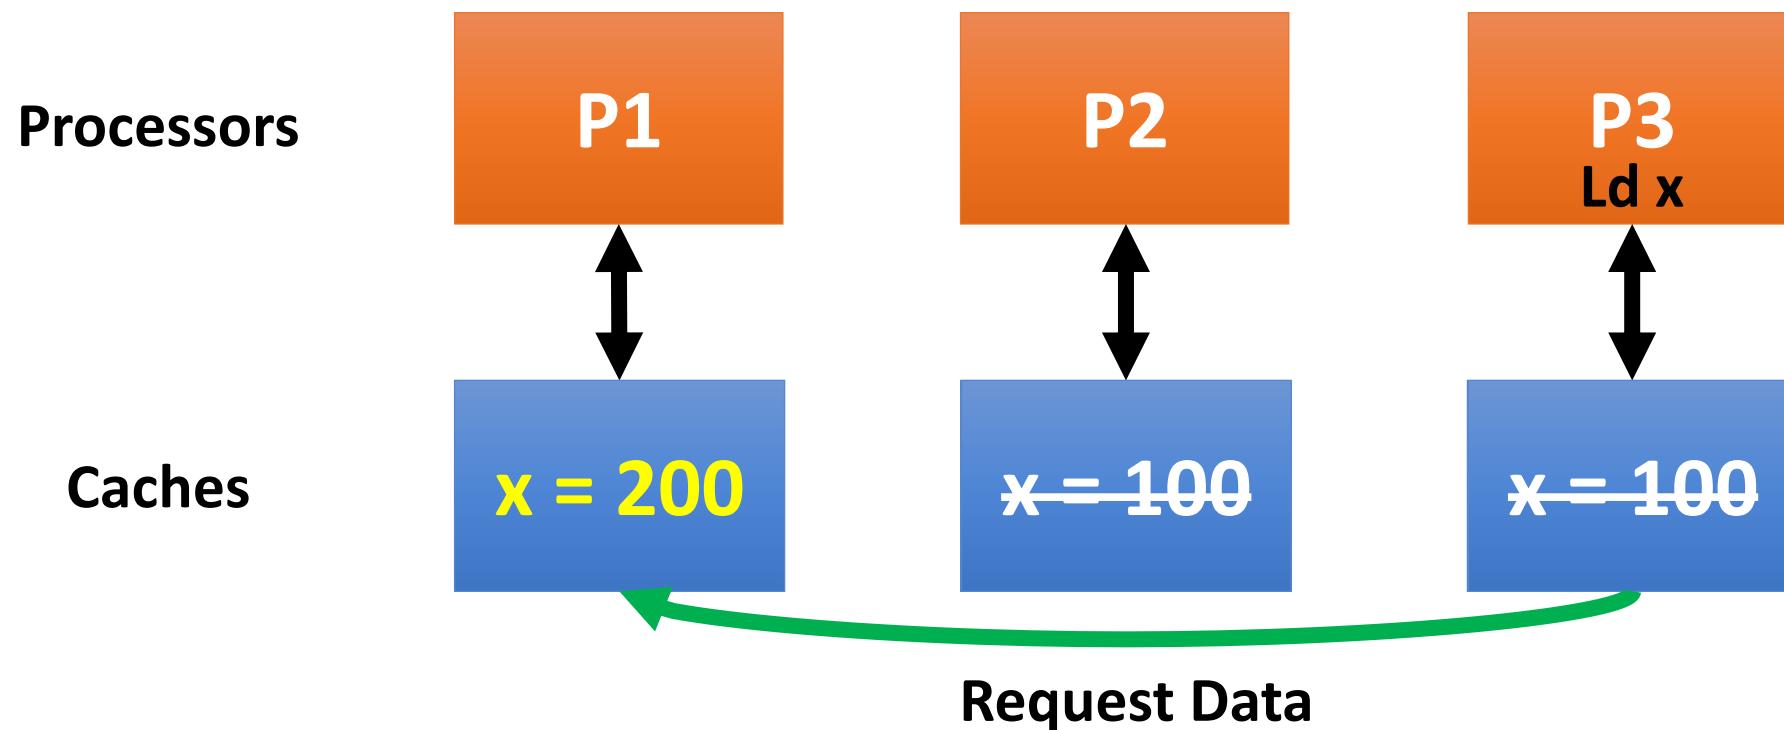
# Coherence Protocol Example

- If P1 updates the value of  $x$  to 200, the stale value of  $x$  in other processors must be **invalidated**
- If P3 wants to subsequently read/write  $x$ , it must request the new value
- **SWMR** = Single-Writer Multiple Readers, **DVI** = Data Value Invariant



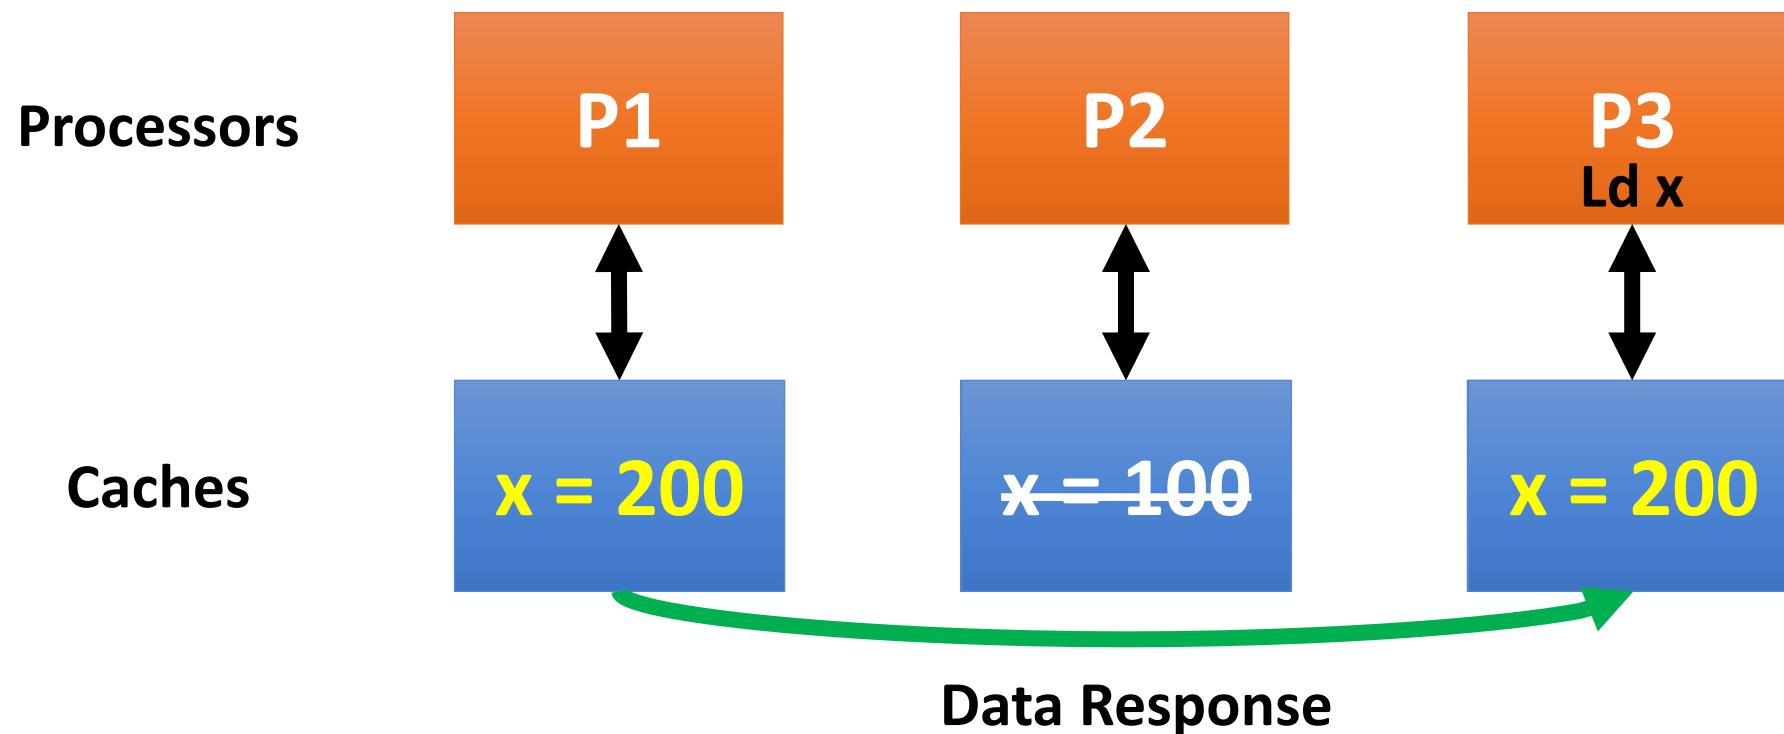
# Coherence Protocol Example

- If P1 updates the value of  $x$  to 200, the stale value of  $x$  in other processors must be **invalidated**
- If P3 wants to subsequently read/write  $x$ , it must request the new value
- **SWMR** = Single-Writer Multiple Readers, **DVI** = Data Value Invariant



# Coherence Protocol Example

- If P1 updates the value of  $x$  to 200, the stale value of  $x$  in other processors must be **invalidated**
- If P3 wants to subsequently read/write  $x$ , it must request the new value
- **SWMR** = Single-Writer Multiple Readers, **DVI** = Data Value Invariant



# Motivating Example – “Peekaboo” [Sorin et al. Primer 2011]

- Three optimizations: correct individually, but not in combination



# Motivating Example – “Peekaboo” [Sorin et al. Primer 2011]

- Three optimizations: correct individually, but not in combination
1. Prefetching



# Motivating Example – “Peekaboo” [Sorin et al. Primer 2011]

- Three optimizations: correct individually, but not in combination
  1. Prefetching
  2. Invalidation before use
    - Invalidation can arrive before data
    - Acknowledge Inv early rather than wait for data to arrive
    - But repeated inv before use → livelock [Kubiatowicz et al. ASPLOS 1992]



# Motivating Example – “Peekaboo” [Sorin et al. Primer 2011]

- Three optimizations: correct individually, but not in combination

1. Prefetching

2. Invalidation before use

- Invalidation can arrive before data
- Acknowledge Inv early rather than wait for data to arrive
- But repeated inv before use → livelock [Kubiatowicz et al. ASPLOS 1992]

3. Livelock avoidance: allow destination core to perform one operation on data when it arrives, even if already invalidated

[Sorin et al. Primer 2011]

- Does **not** break coherence
- Sometimes **intentionally** returns stale data



# Motivating Example – “Peekaboo”

- Consider **mp** with the livelock-avoidance mechanism:

Core 0	Core 1
(i1) $[x] \leftarrow 1$	(i3) $r1 \leftarrow [y]$
(i2) $[y] \leftarrow 1$	(i4) $r2 \leftarrow [x]$
Under SC: Forbid $r1=1, r2=0$	

## Optimizations:

1. Prefetching
2. Invalidation-before-use
3. Livelock avoidance

Core 0

x: Shared  
y: Modified

$[x] \leftarrow 1$   
 $[y] \leftarrow 1$

Core 1

x: Invalid  
y: Invalid

$r1 \leftarrow [y]$   
 $r2 \leftarrow [x]$



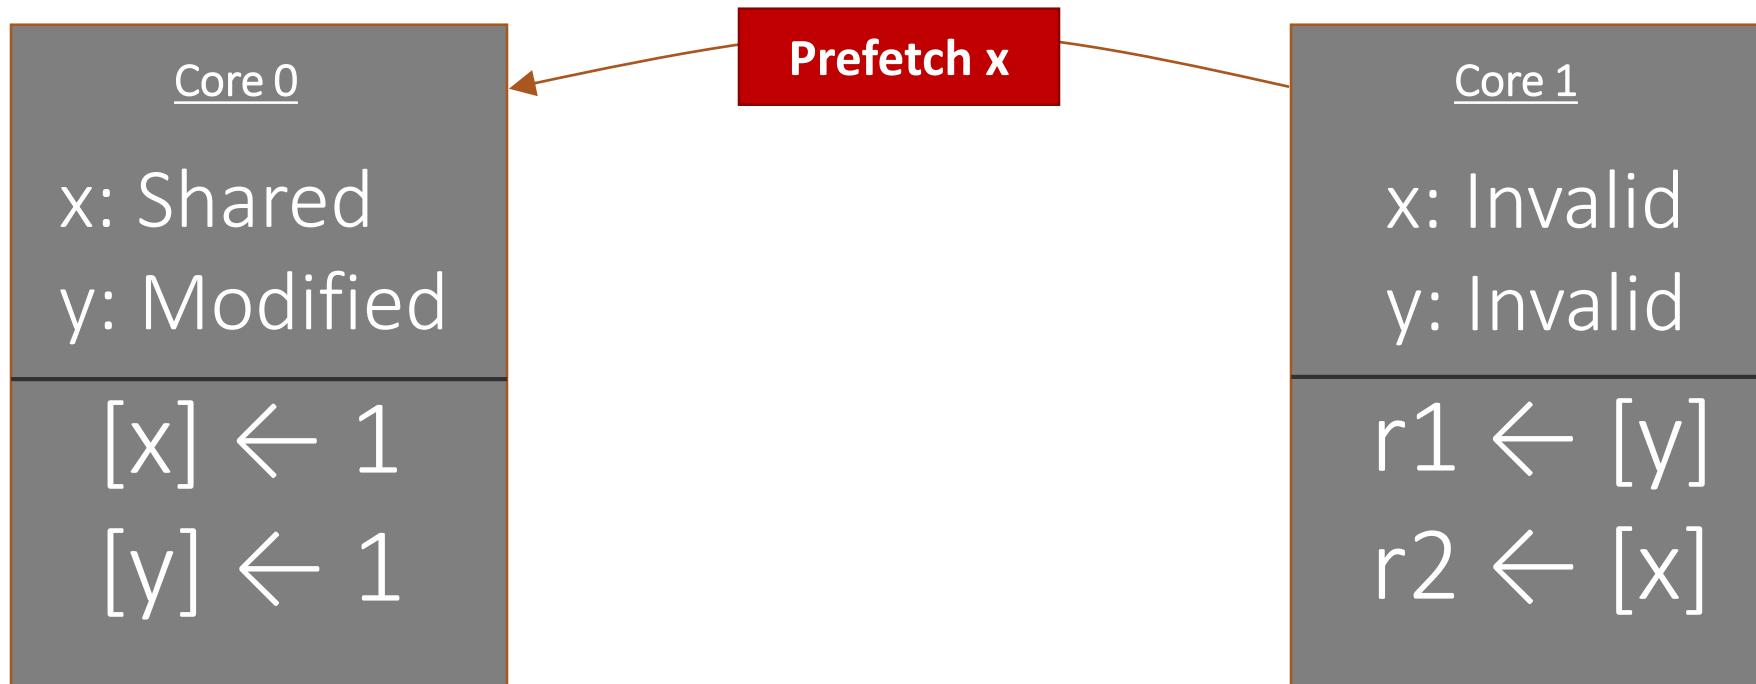
# Motivating Example – “Peekaboo”

- Consider **mp** with the livelock-avoidance mechanism:

Core 0	Core 1
(i1) $[x] \leftarrow 1$	(i3) $r1 \leftarrow [y]$
(i2) $[y] \leftarrow 1$	(i4) $r2 \leftarrow [x]$
Under SC: Forbid $r1=1, r2=0$	

## Optimizations:

1. Prefetching
2. Invalidation-before-use
3. Livelock avoidance



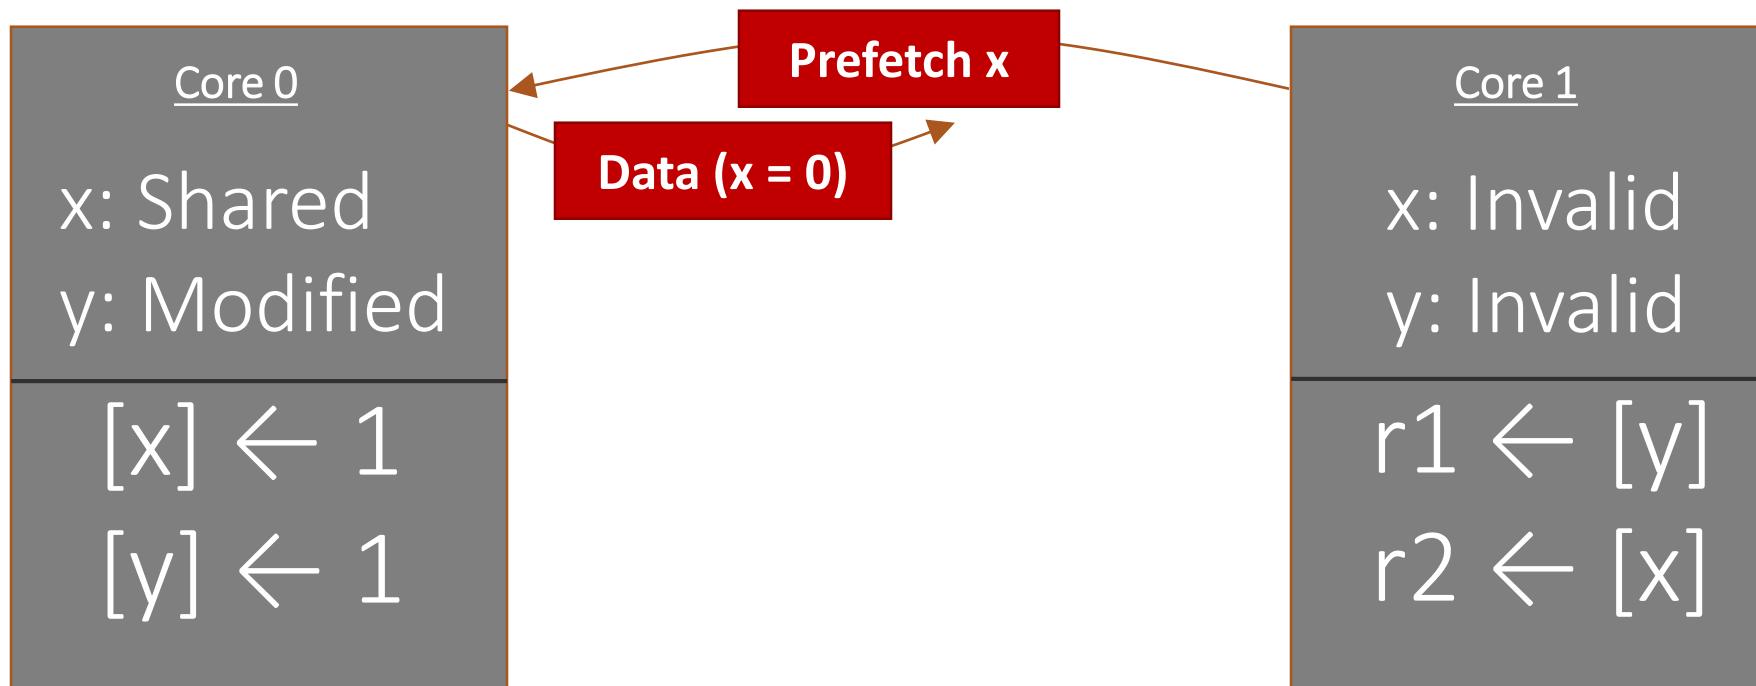
# Motivating Example – “Peekaboo”

- Consider **mp** with the livelock-avoidance mechanism:

Core 0	Core 1
(i1) $[x] \leftarrow 1$	(i3) $r1 \leftarrow [y]$
(i2) $[y] \leftarrow 1$	(i4) $r2 \leftarrow [x]$
Under SC: Forbid $r1=1, r2=0$	

## Optimizations:

1. Prefetching
2. Invalidation-before-use
3. Livelock avoidance



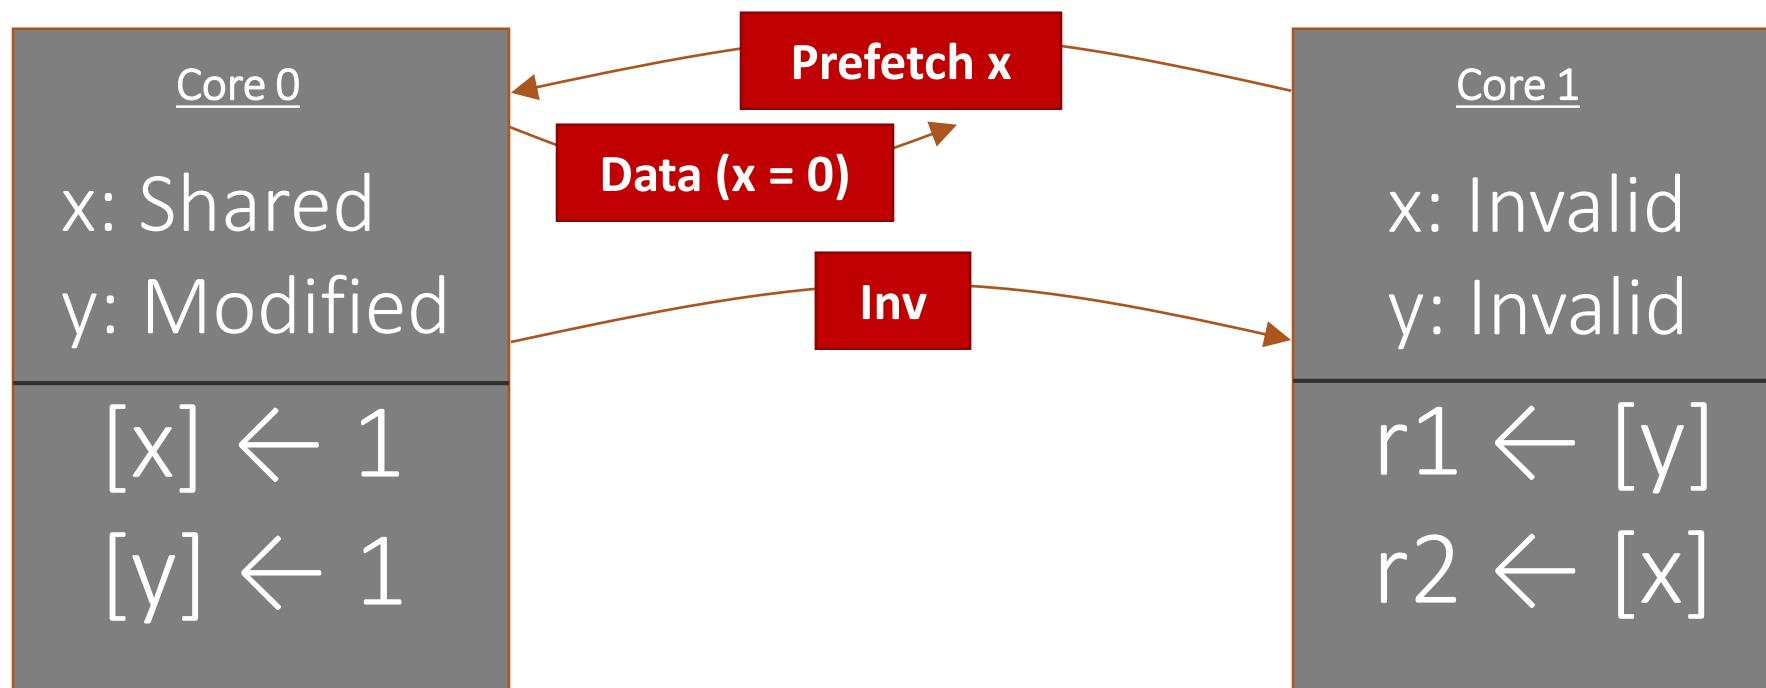
# Motivating Example – “Peekaboo”

- Consider **mp** with the livelock-avoidance mechanism:

Core 0	Core 1
(i1) $[x] \leftarrow 1$	(i3) $r1 \leftarrow [y]$
(i2) $[y] \leftarrow 1$	(i4) $r2 \leftarrow [x]$
Under SC: Forbid $r1=1, r2=0$	

## Optimizations:

1. Prefetching
2. Invalidation-before-use
3. Livelock avoidance



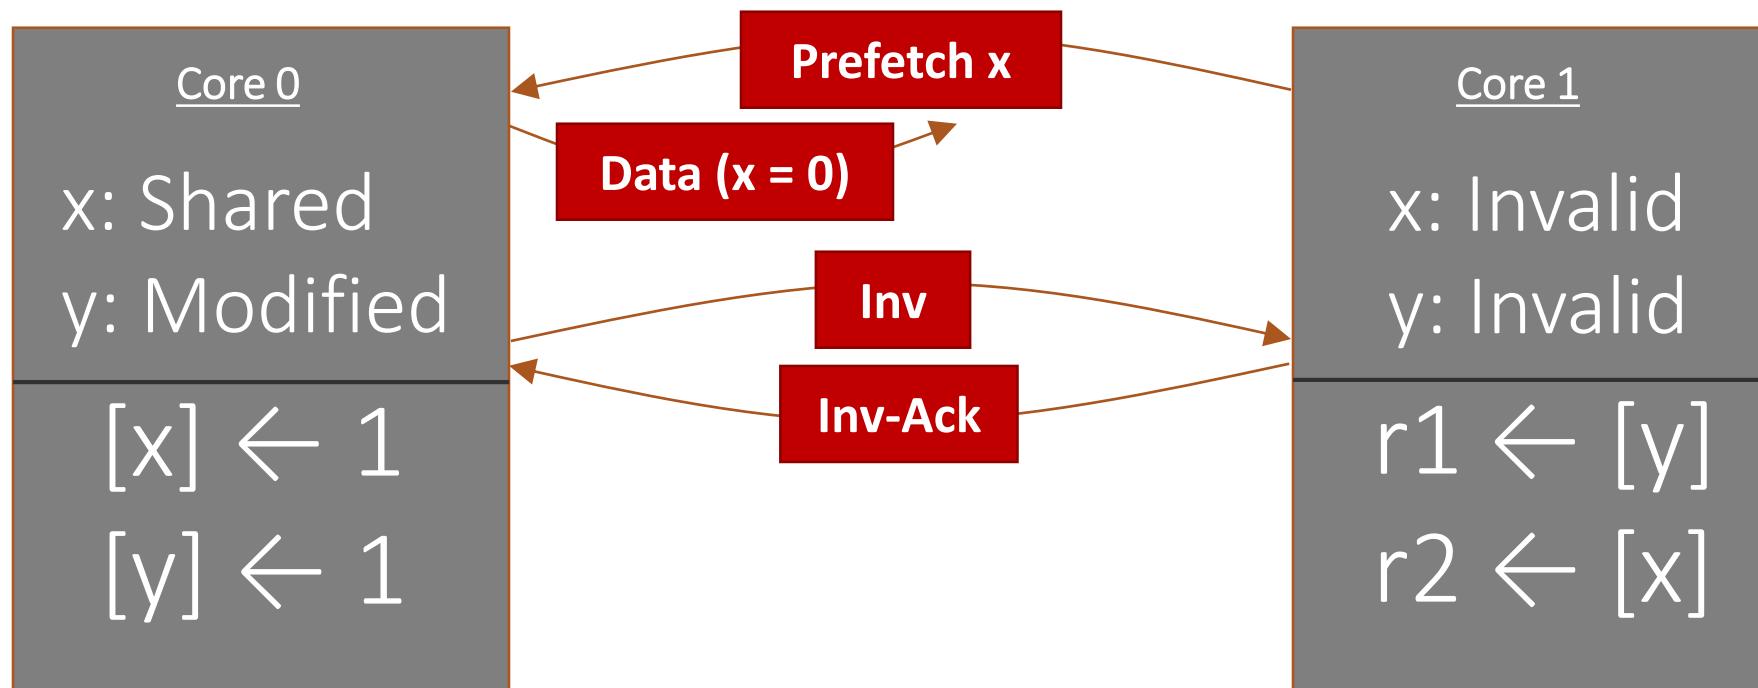
# Motivating Example – “Peekaboo”

- Consider **mp** with the livelock-avoidance mechanism:

Core 0	Core 1
(i1) $[x] \leftarrow 1$	(i3) $r1 \leftarrow [y]$
(i2) $[y] \leftarrow 1$	(i4) $r2 \leftarrow [x]$
Under SC: Forbid $r1=1, r2=0$	

## Optimizations:

1. Prefetching
2. Invalidation-before-use
3. Livelock avoidance



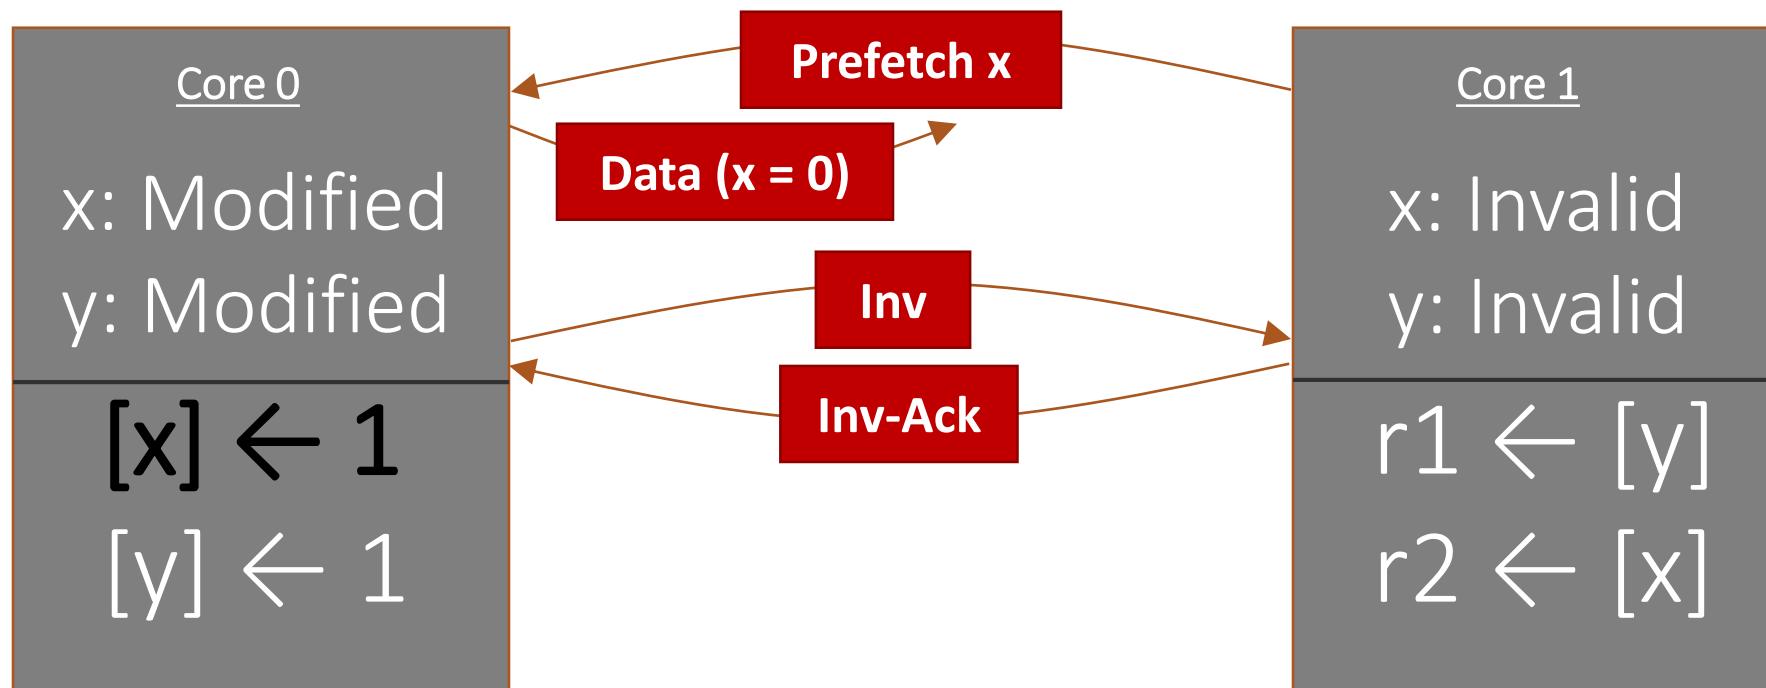
# Motivating Example – “Peekaboo”

- Consider **mp** with the livelock-avoidance mechanism:

Core 0	Core 1
(i1) $[x] \leftarrow 1$	(i3) $r1 \leftarrow [y]$
(i2) $[y] \leftarrow 1$	(i4) $r2 \leftarrow [x]$
Under SC: Forbid $r1=1, r2=0$	

## Optimizations:

1. Prefetching
2. Invalidation-before-use
3. Livelock avoidance



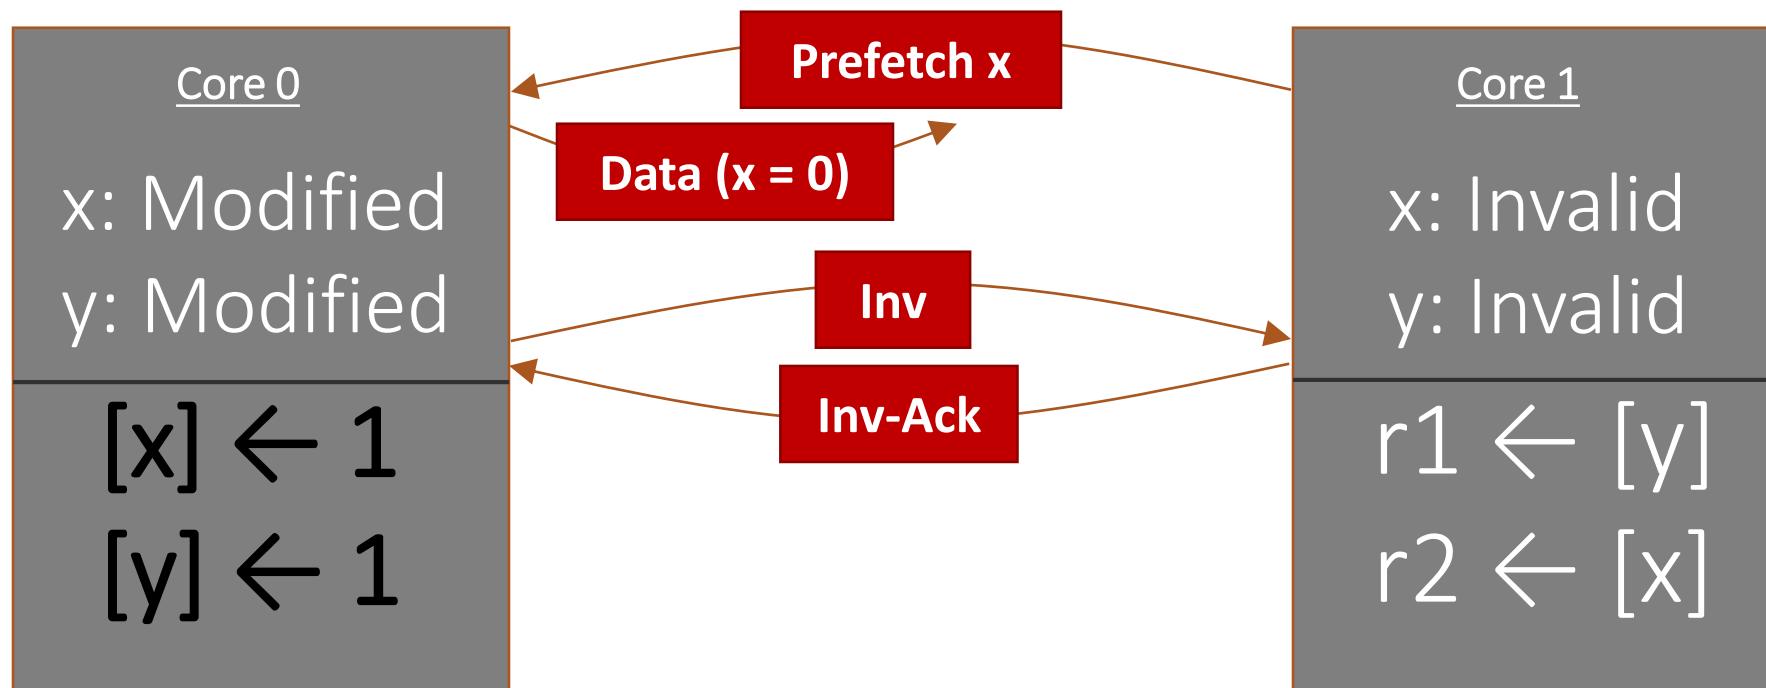
# Motivating Example – “Peekaboo”

- Consider **mp** with the livelock-avoidance mechanism:

Core 0	Core 1
(i1) $[x] \leftarrow 1$	(i3) $r1 \leftarrow [y]$
(i2) $[y] \leftarrow 1$	(i4) $r2 \leftarrow [x]$
Under SC: Forbid $r1=1, r2=0$	

## Optimizations:

1. Prefetching
2. Invalidation-before-use
3. Livelock avoidance



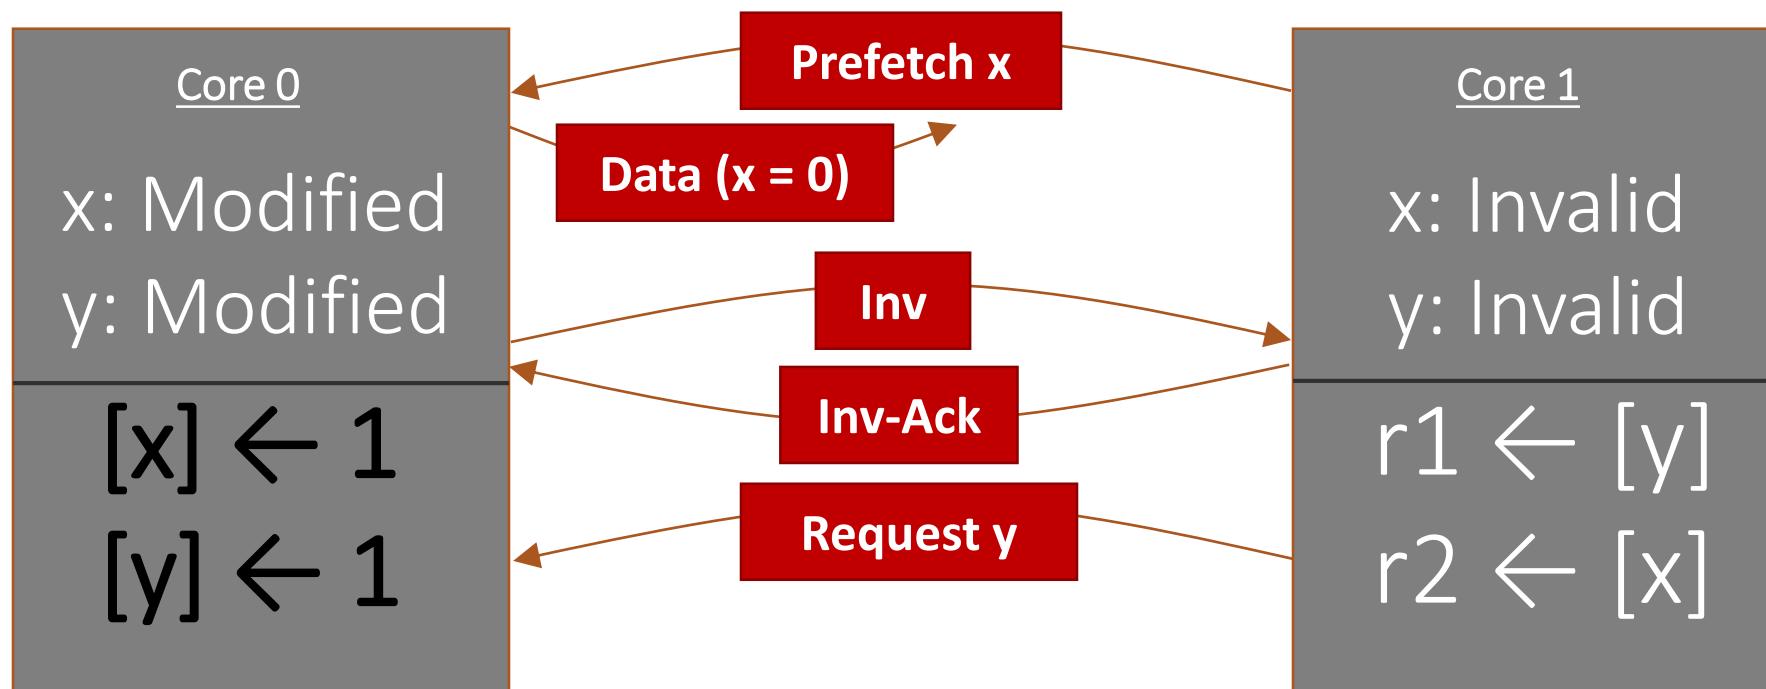
# Motivating Example – “Peekaboo”

- Consider **mp** with the livelock-avoidance mechanism:

Core 0	Core 1
(i1) $[x] \leftarrow 1$	(i3) $r1 \leftarrow [y]$
(i2) $[y] \leftarrow 1$	(i4) $r2 \leftarrow [x]$
Under SC: Forbid $r1=1, r2=0$	

## Optimizations:

1. Prefetching
2. Invalidation-before-use
3. Livelock avoidance



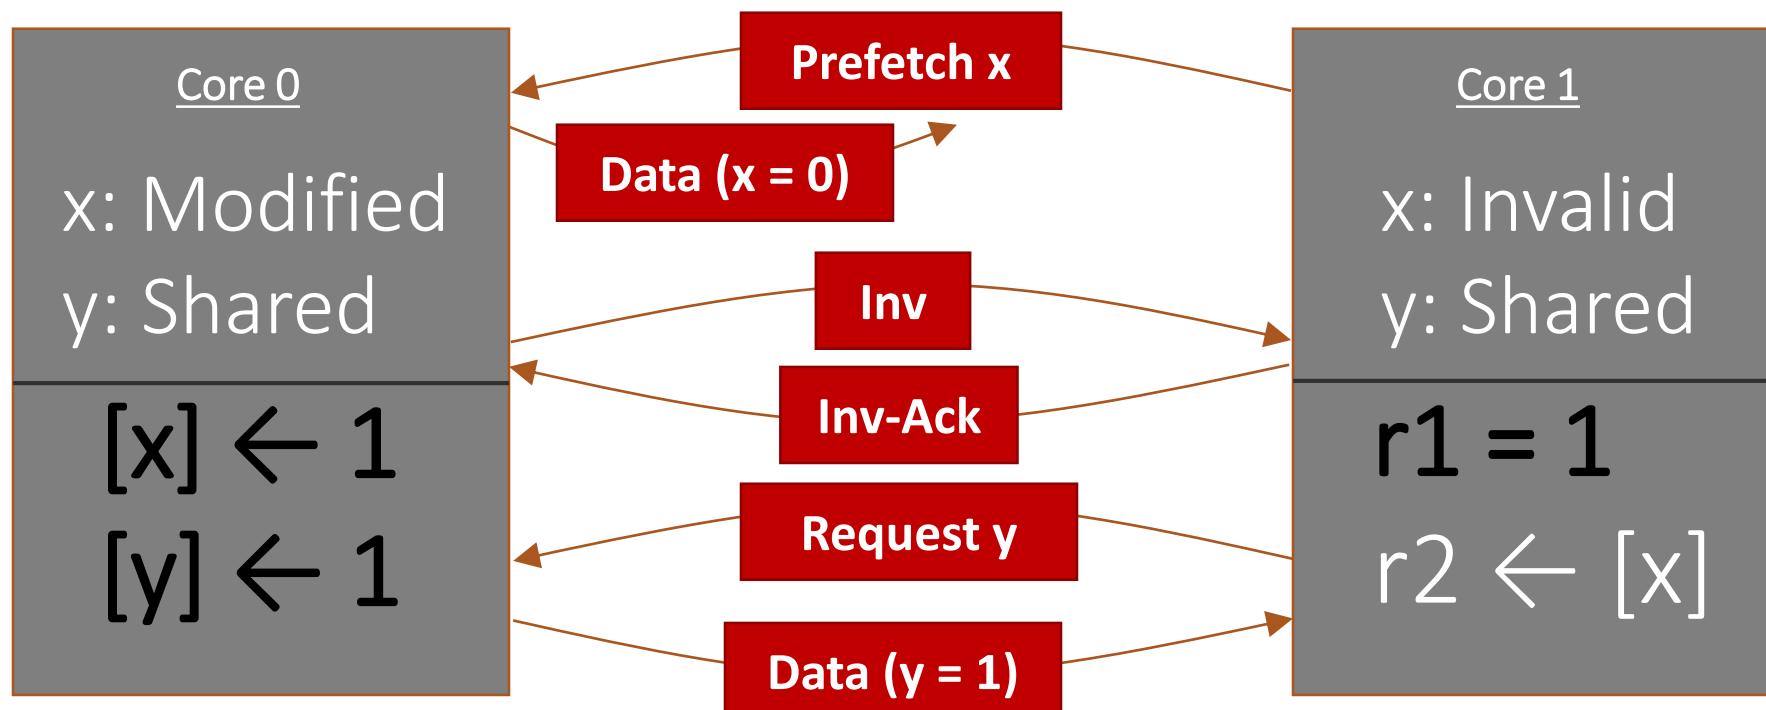
# Motivating Example – “Peekaboo”

- Consider **mp** with the livelock-avoidance mechanism:

Core 0	Core 1
(i1) $[x] \leftarrow 1$	(i3) $r1 \leftarrow [y]$
(i2) $[y] \leftarrow 1$	(i4) $r2 \leftarrow [x]$
Under SC: Forbid $r1=1, r2=0$	

## Optimizations:

1. Prefetching
2. Invalidation-before-use
3. Livelock avoidance



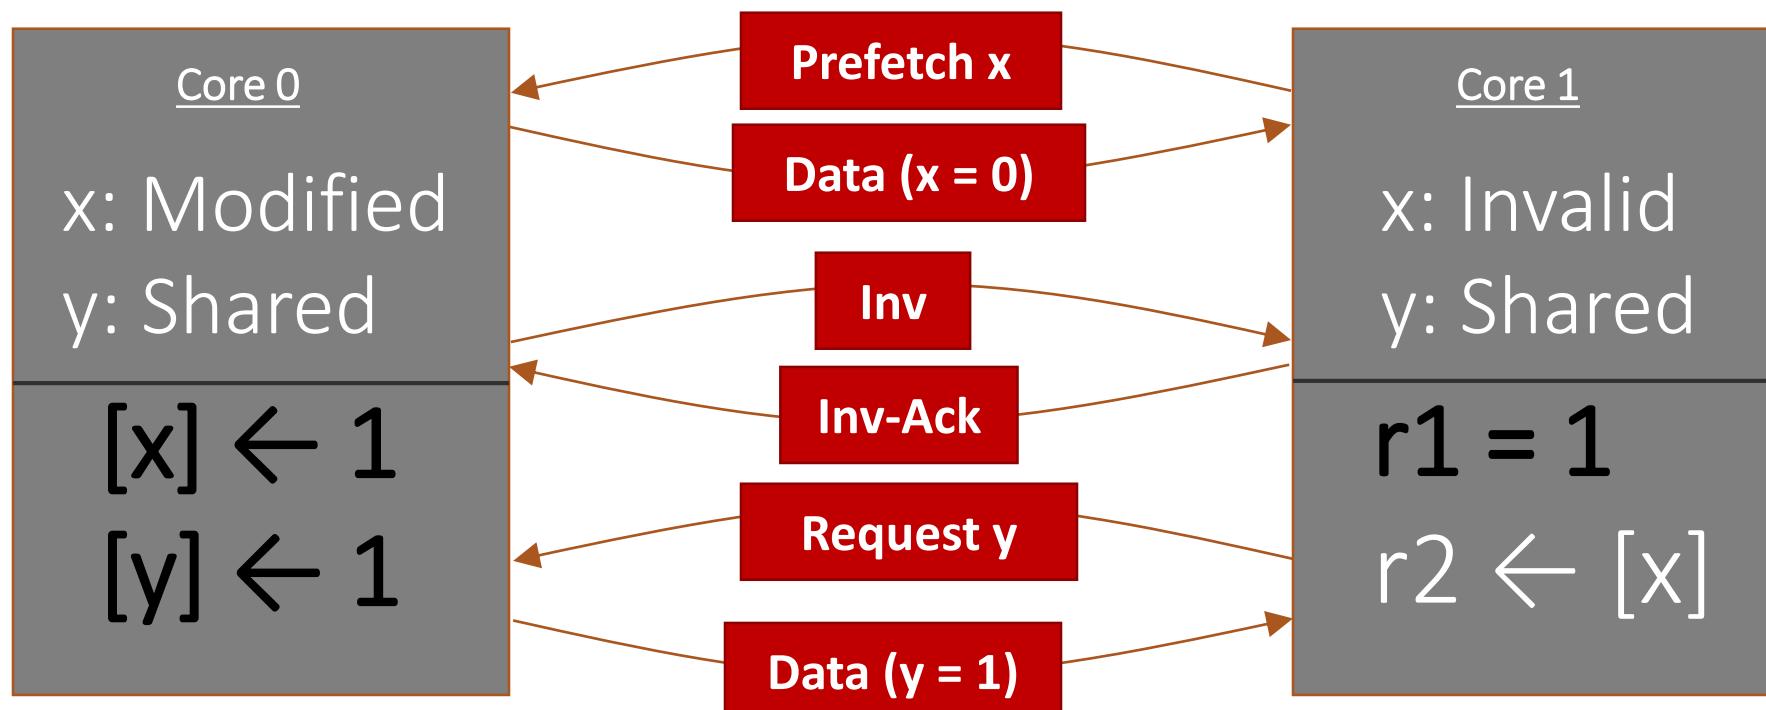
# Motivating Example – “Peekaboo”

- Consider **mp** with the livelock-avoidance mechanism:

Core 0	Core 1
(i1) $[x] \leftarrow 1$	(i3) $r1 \leftarrow [y]$
(i2) $[y] \leftarrow 1$	(i4) $r2 \leftarrow [x]$
Under SC: Forbid $r1=1, r2=0$	

## Optimizations:

1. Prefetching
2. Invalidation-before-use
3. Livelock avoidance



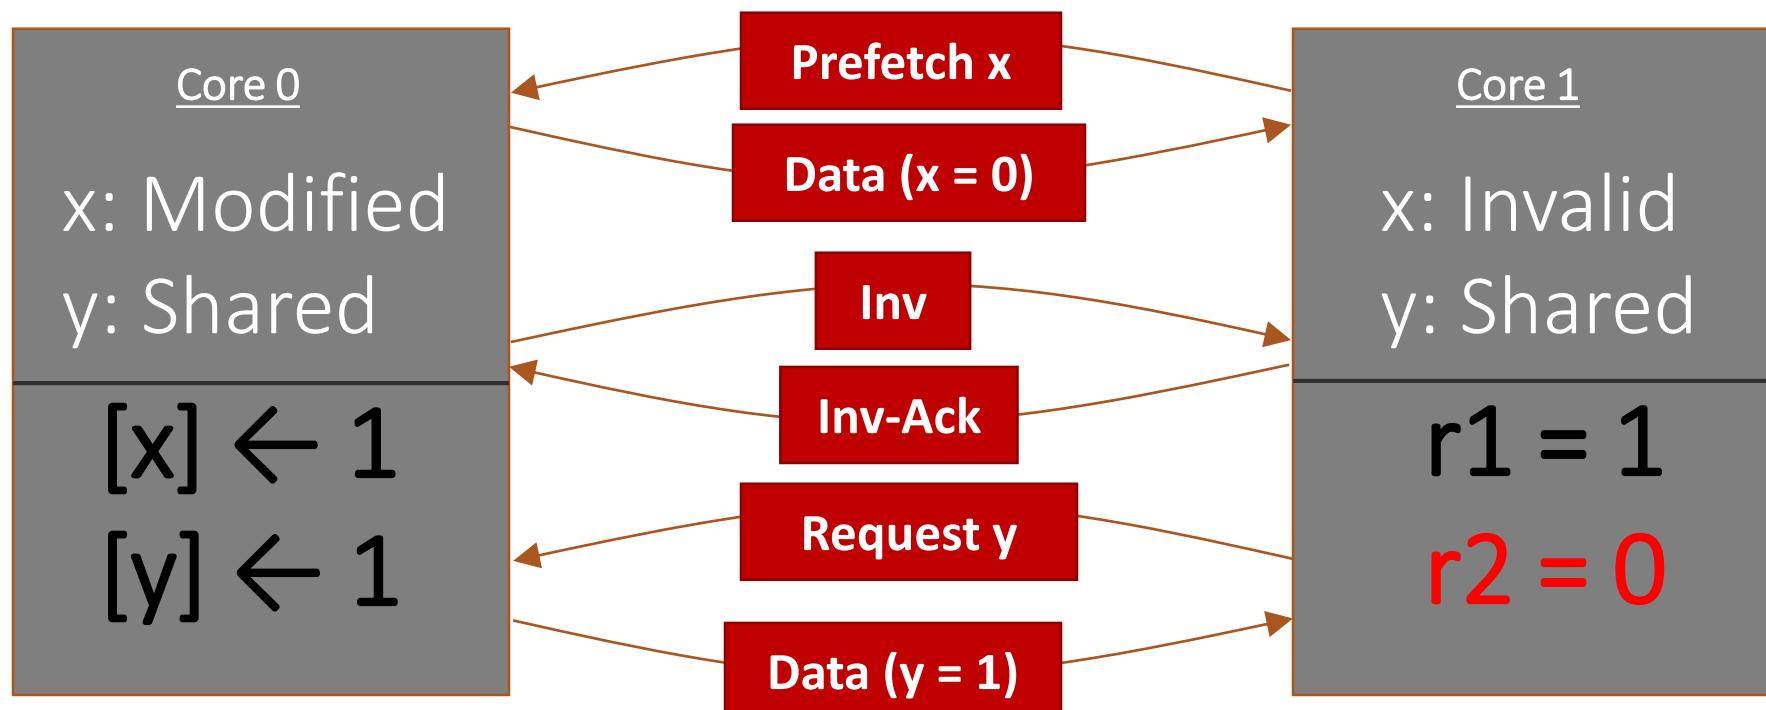
# Motivating Example – “Peekaboo”

- Consider **mp** with the livelock-avoidance mechanism:

Core 0	Core 1
(i1) $[x] \leftarrow 1$	(i3) $r1 \leftarrow [y]$
(i2) $[y] \leftarrow 1$	(i4) $r2 \leftarrow [x]$
Under SC: Forbid $r1=1, r2=0$	

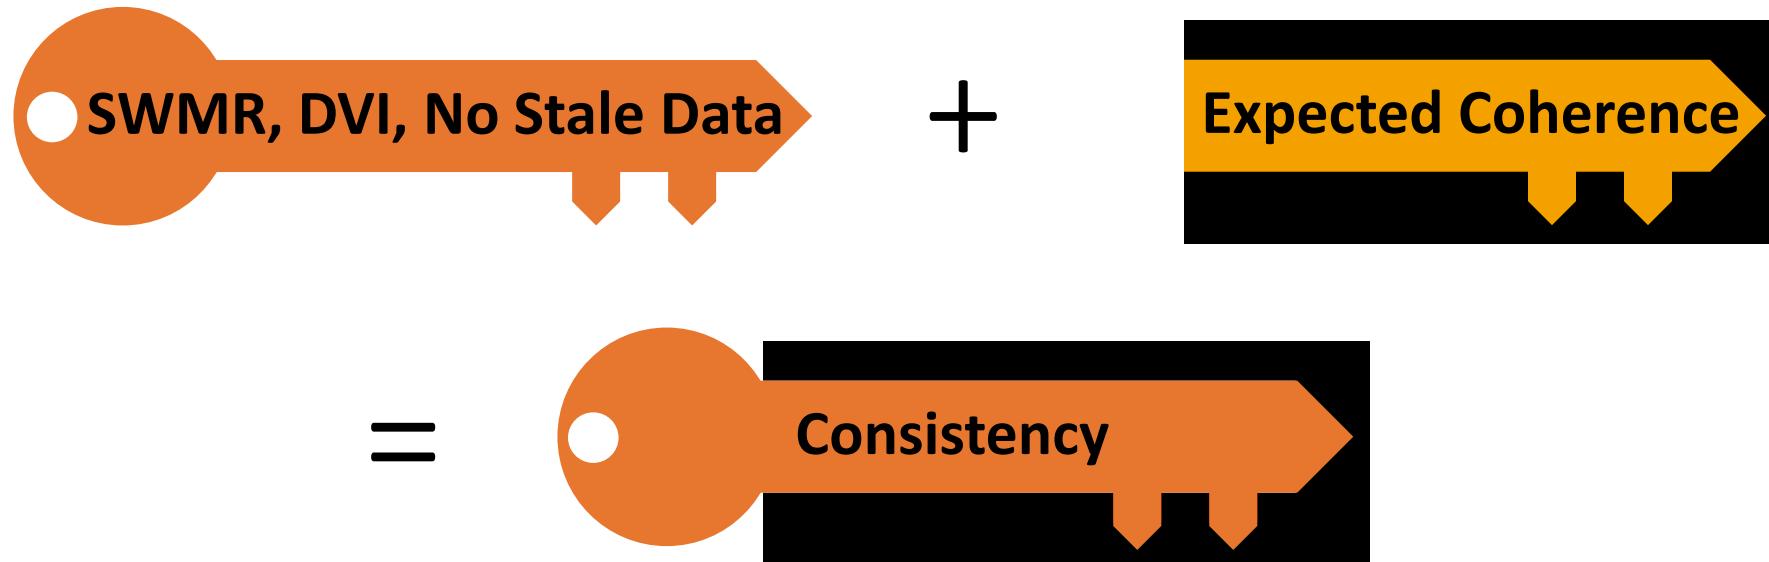
## Optimizations:

1. Prefetching
2. Invalidation-before-use
3. Livelock avoidance



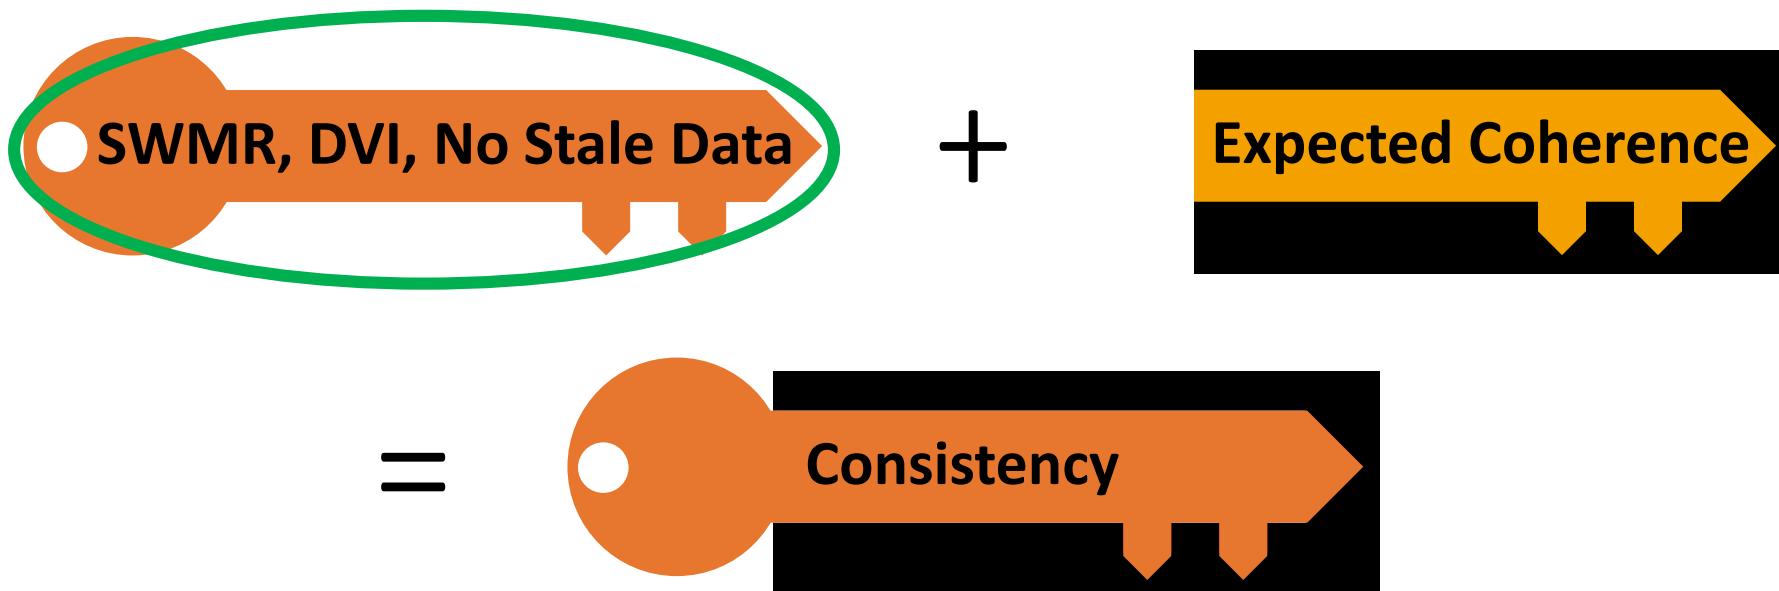
# The Coherence-Consistency Interface (CCI)

- CCI = coherence protocol guarantees to microarch. + orderings microarch. expects from coherence protocol



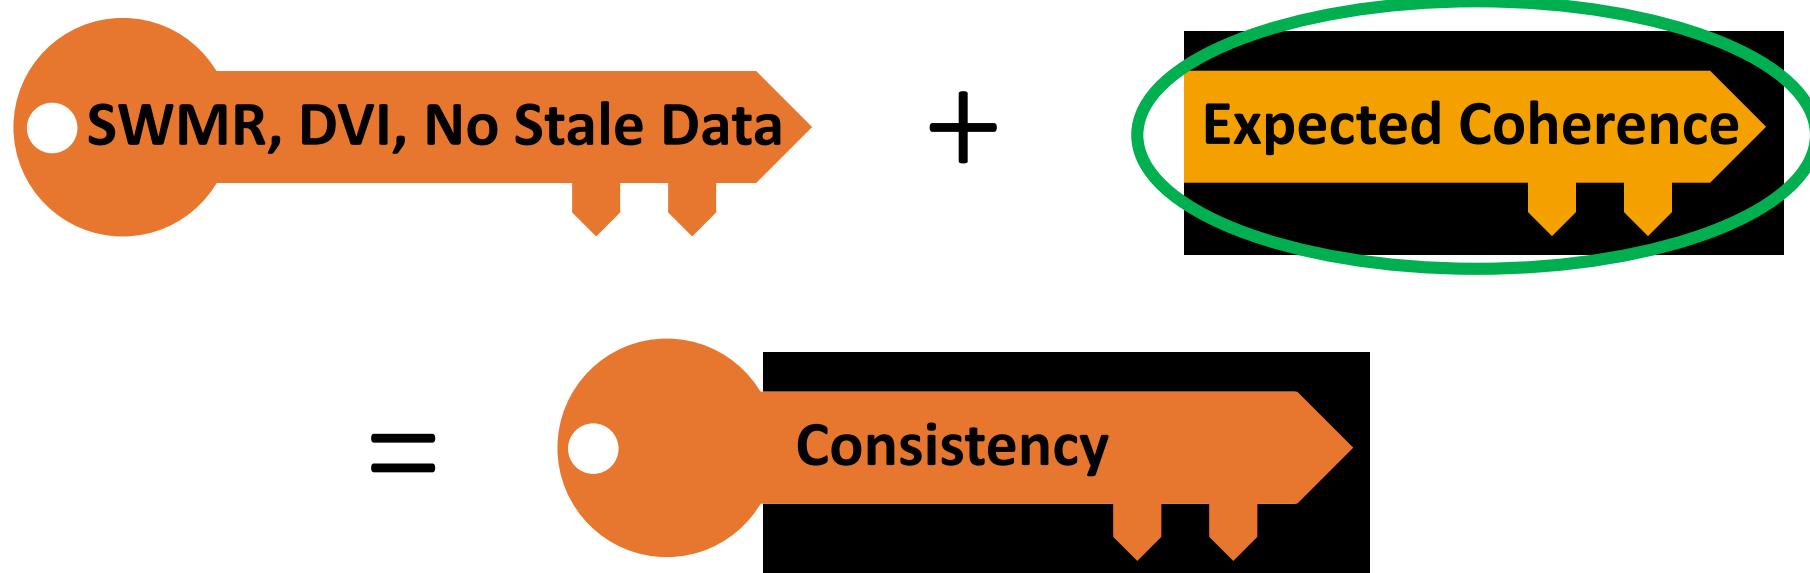
# The Coherence-Consistency Interface (CCI)

- CCI = coherence protocol guarantees to microarch. + orderings microarch. expects from coherence protocol



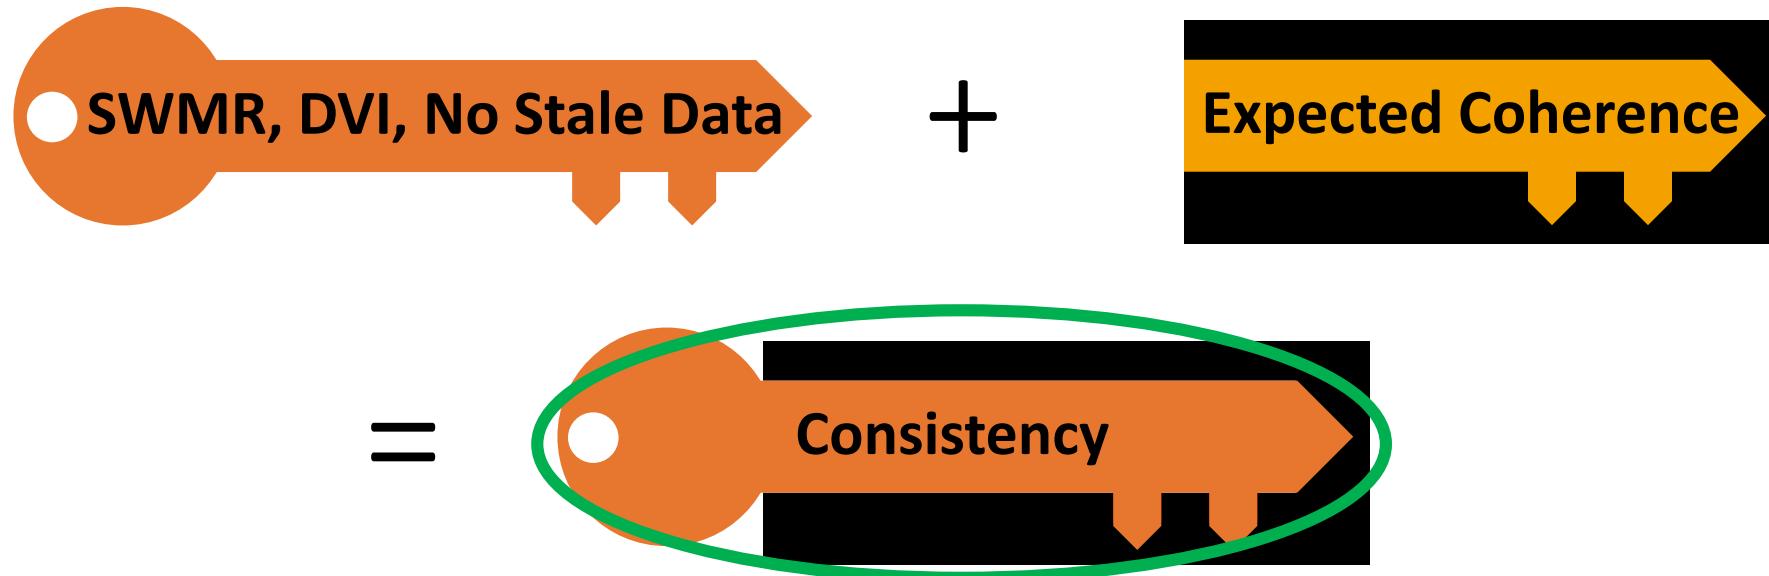
# The Coherence-Consistency Interface (CCI)

- CCI = coherence protocol guarantees to microarch. + orderings microarch. expects from coherence protocol



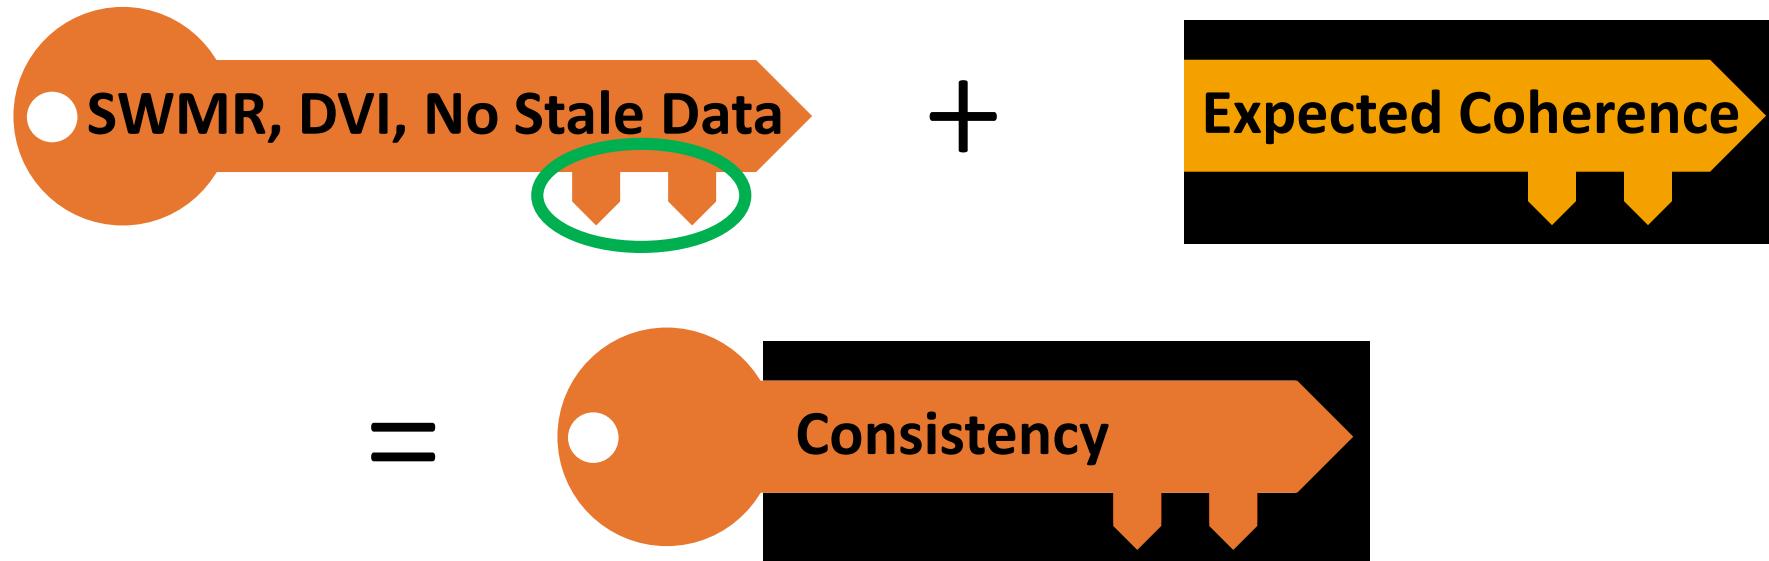
# The Coherence-Consistency Interface (CCI)

- CCI = coherence protocol guarantees to microarch. + orderings microarch. expects from coherence protocol



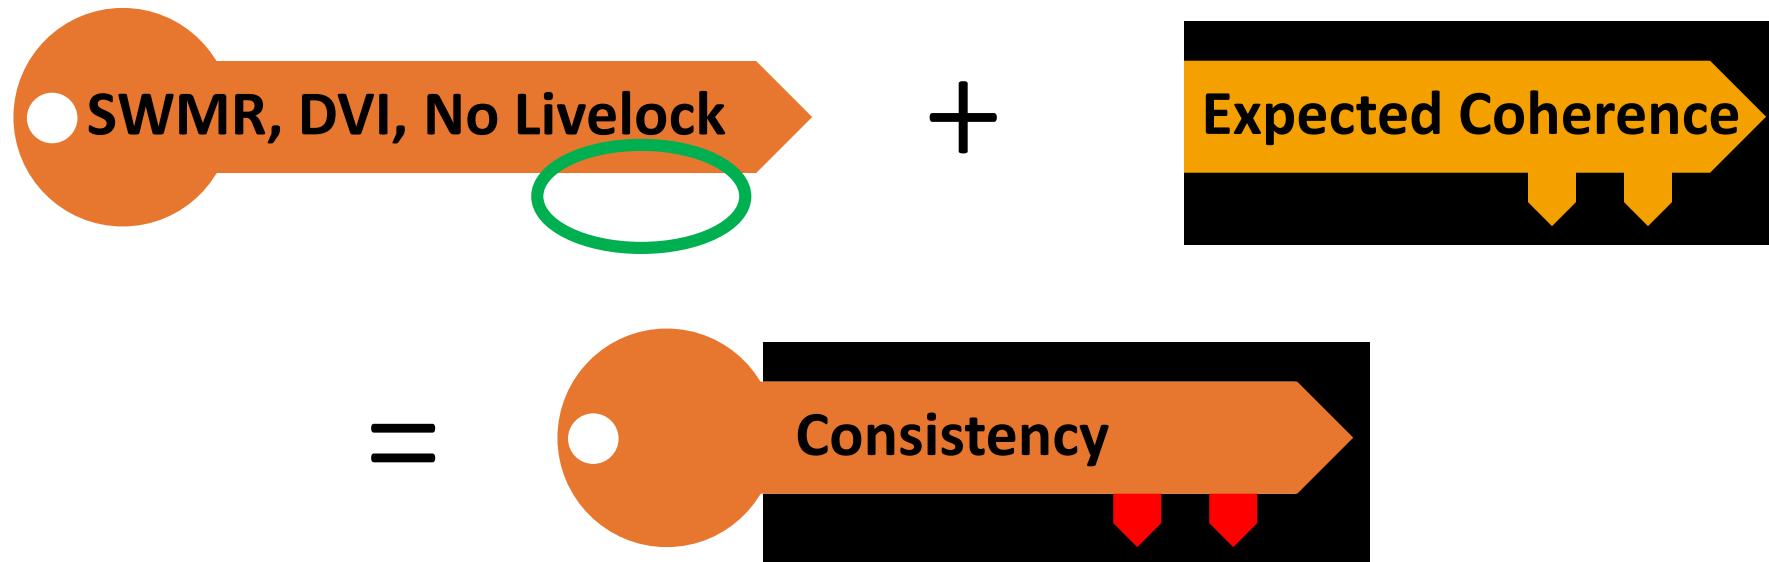
# The Coherence-Consistency Interface (CCI)

- CCI = coherence protocol guarantees to microarch. + orderings microarch. expects from coherence protocol



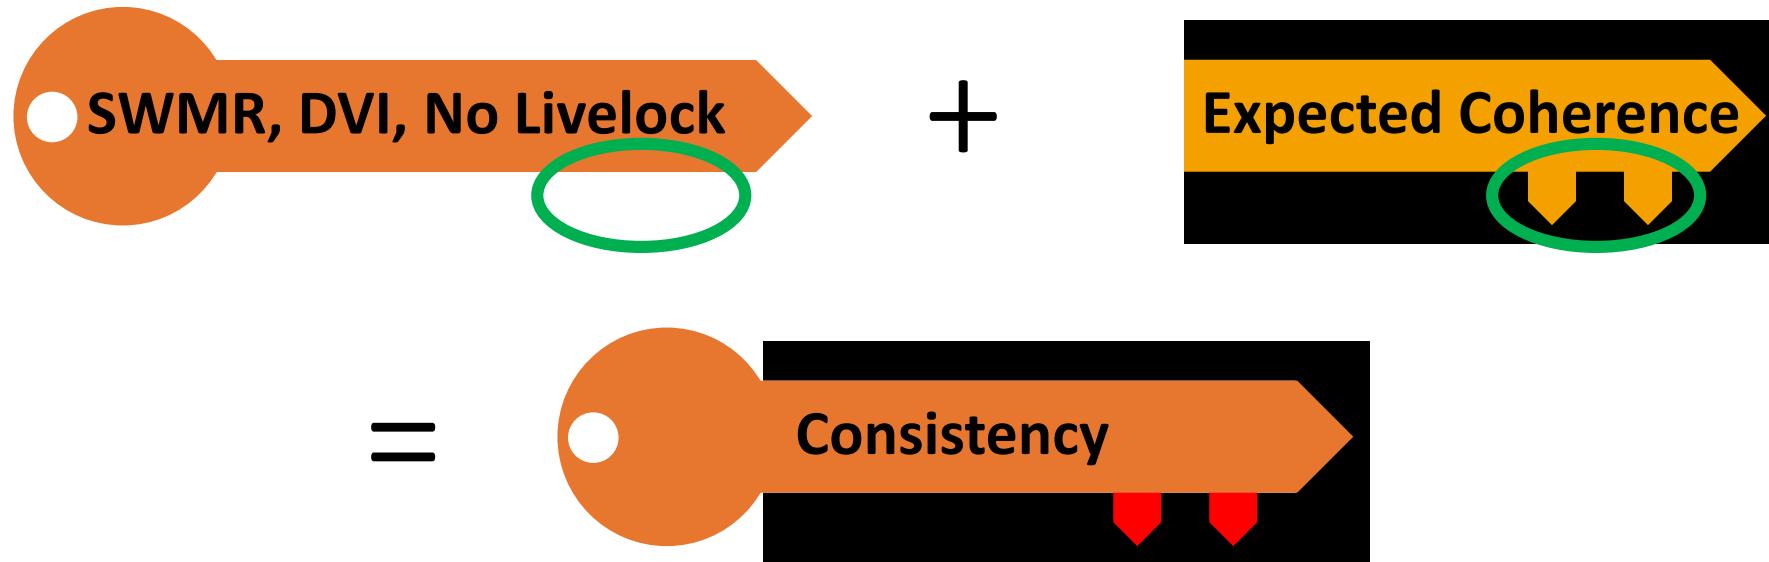
# The Coherence-Consistency Interface (CCI)

- CCI = coherence protocol guarantees to microarch. + orderings microarch. expects from coherence protocol



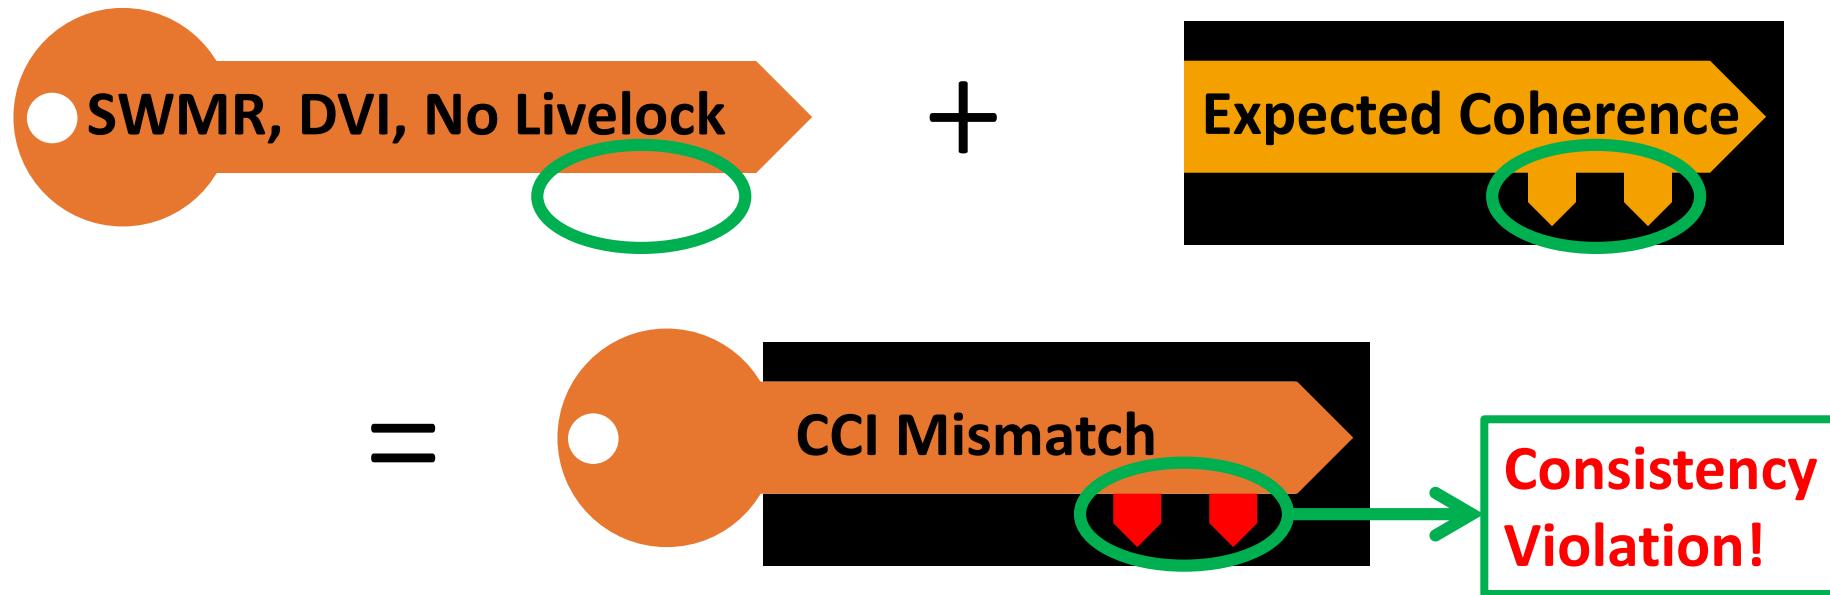
# The Coherence-Consistency Interface (CCI)

- CCI = coherence protocol guarantees to microarch. + orderings microarch. expects from coherence protocol



# The Coherence-Consistency Interface (CCI)

- CCI = coherence protocol guarantees to microarch. + orderings microarch. expects from coherence protocol

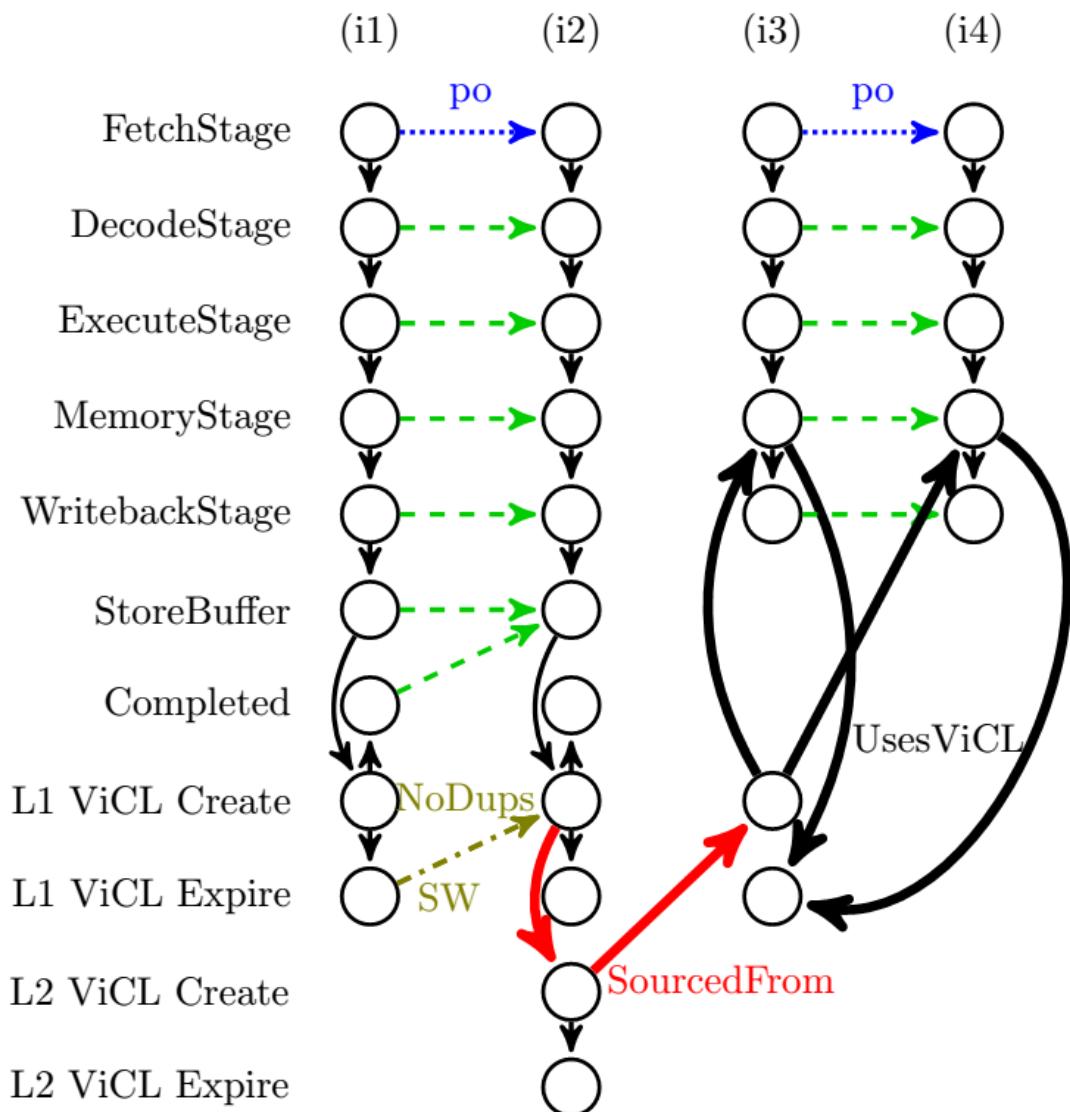


# ViCL: Value in Cache Lifetime

- Need a way to model cache occupancy and coherence events for:
  - Coherence protocol optimizations (eg: Peekaboo)
  - Partial incoherence and lazy coherence (GPUs, etc)
- A ViCL is a 4-tuple:  
 $(cache\_id, address, data\_value, generation\_id)$
- **cache\_id** and **generation\_id** uniquely identify each cache line
- A ViCL 4-tuple maps on to the period of time over which the cache line serves the data value for the address



# ViCLs in $\mu$ hb Graphs



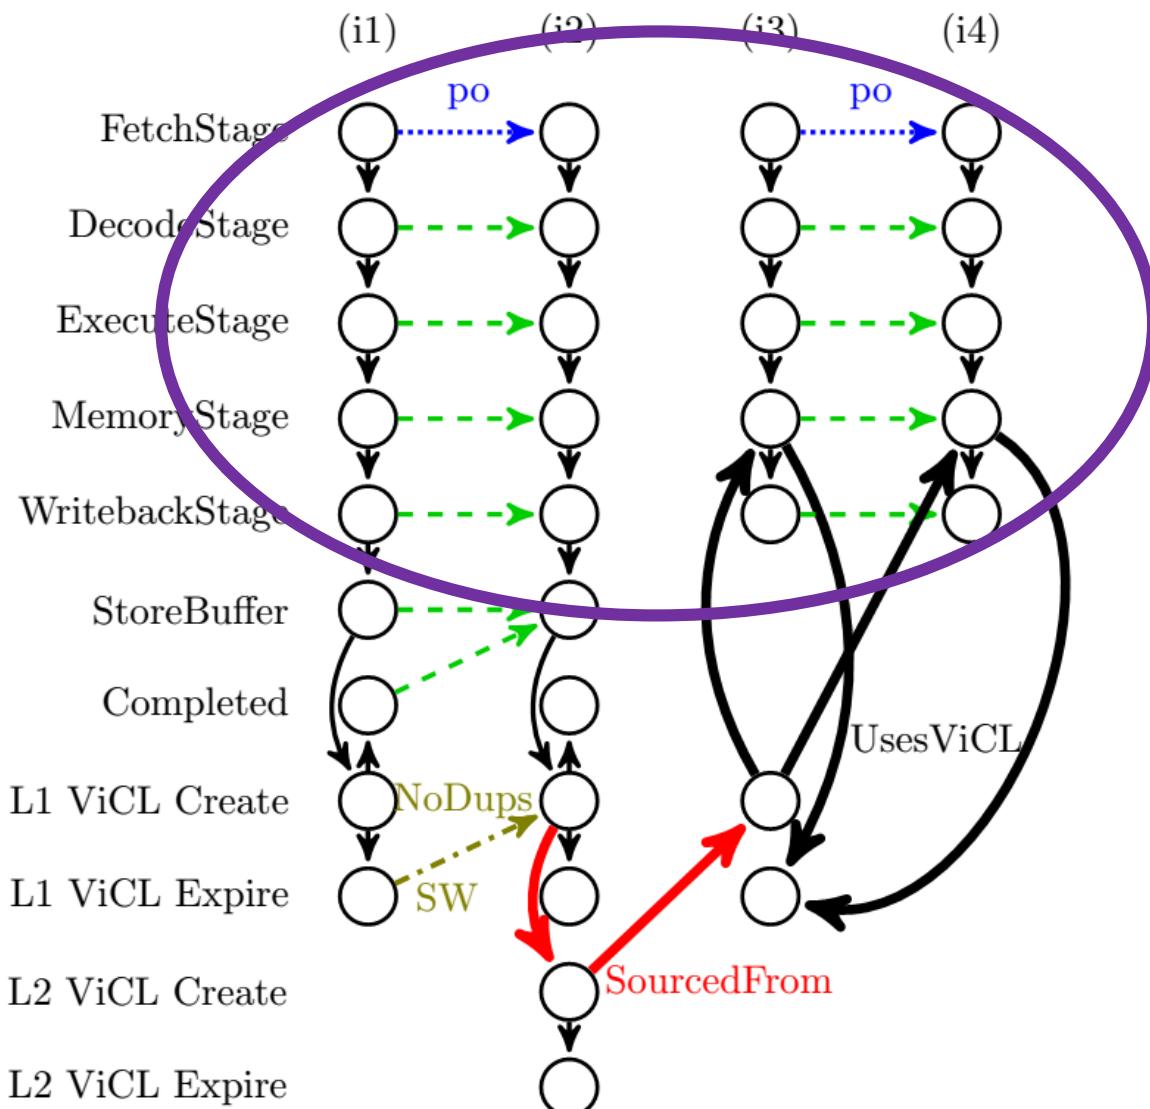
- ViCLs start at a **ViCL Create** event and end at a **ViCL Expire** event
  - Correspond to nodes in  $\mu$ hb graphs
  - Axioms over these nodes and edges enforce coherence and data movement orderings
- Use pipeline model from PipeCheck, but add ViCL nodes and edges

Litmus Test **co-mp**

Core 0	Core 1
(i1) St $[x] \leftarrow 1$	(i3) Ld $r1 \leftarrow [x]$
(i2) St $[x] \leftarrow 2$	(i4) Ld $r2 \leftarrow [x]$
In TSO: $r1=2, r2=2$ Allowed	



# ViCLs in $\mu$ hb Graphs

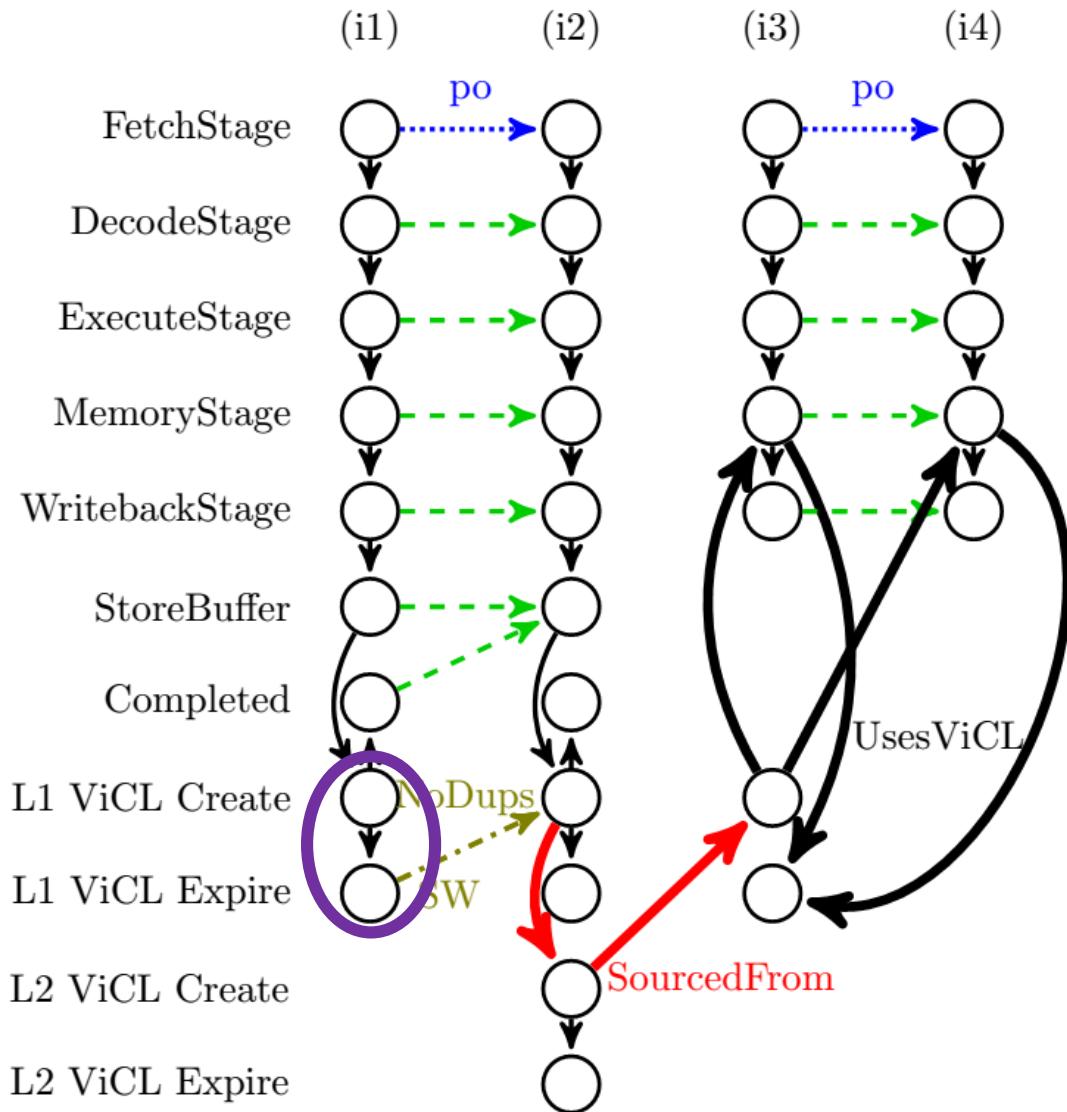


- ViCLs start at a **ViCL Create** event and end at a **ViCL Expire** event
  - Correspond to nodes in  $\mu$ hb graphs
  - Axioms over these nodes and edges enforce coherence and data movement orderings
- Use pipeline model from PipeCheck, but add ViCL nodes and edges

Litmus Test **co-mp**

Core 0	Core 1
(i1) St $[x] \leftarrow 1$	(i3) Ld r1 $\leftarrow [x]$
(i2) St $[x] \leftarrow 2$	(i4) Ld r2 $\leftarrow [x]$
In TSO: r1=2, r2=2 Allowed	

# ViCLs in $\mu$ hb Graphs

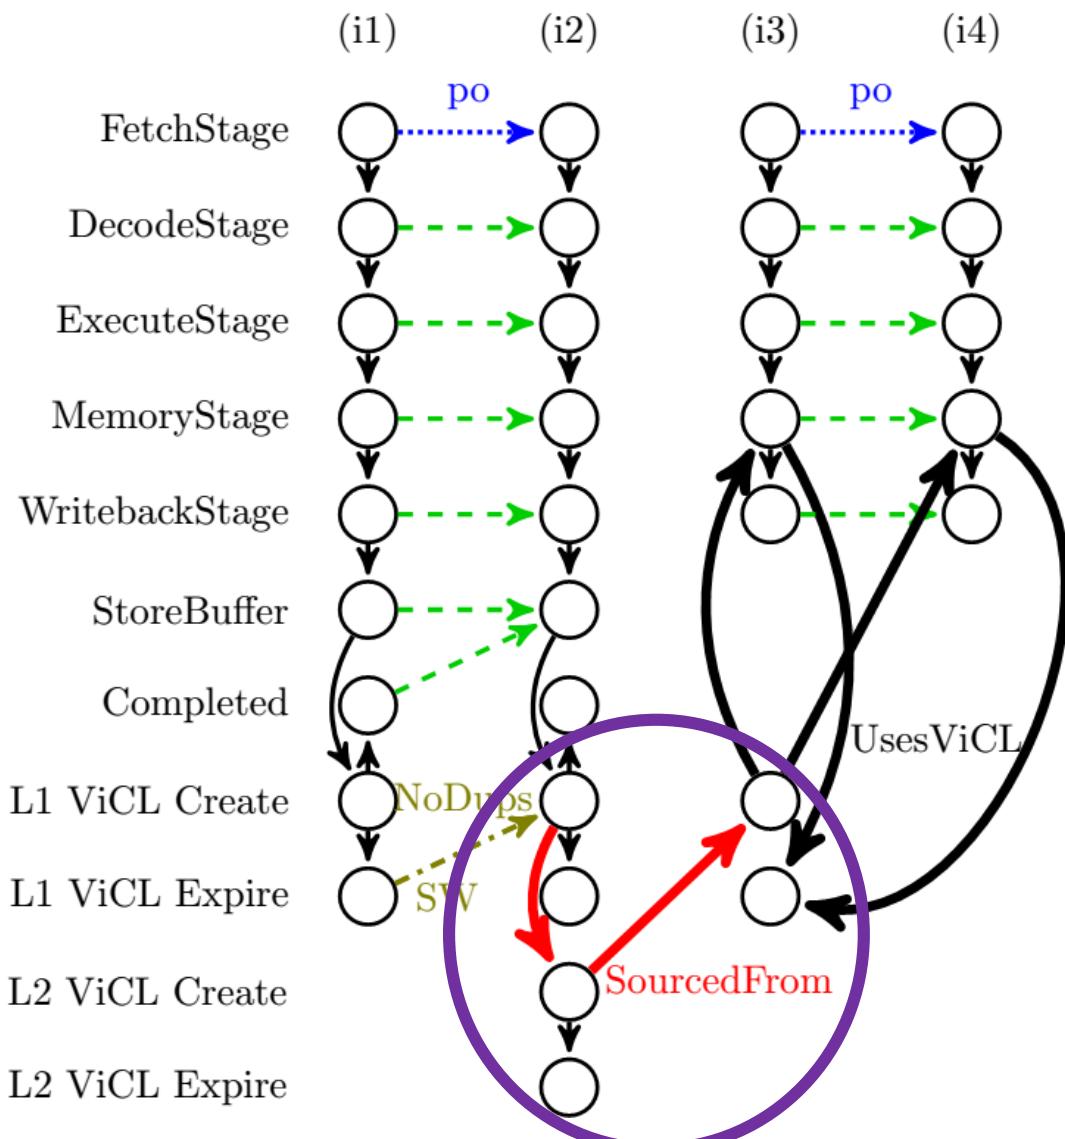


- ViCLs start at a **ViCL Create** event and end at a **ViCL Expire** event
  - Correspond to nodes in  $\mu$ hb graphs
  - Axioms over these nodes and edges enforce coherence and data movement orderings
- Use pipeline model from PipeCheck, but add ViCL nodes and edges

Litmus Test **co-mp**

Core 0	Core 1
(i1) St $[x] \leftarrow 1$	(i3) Ld $r1 \leftarrow [x]$
(i2) St $[x] \leftarrow 2$	(i4) Ld $r2 \leftarrow [x]$
In TSO: $r1=2, r2=2$ Allowed	

# ViCLs in $\mu$ hb Graphs



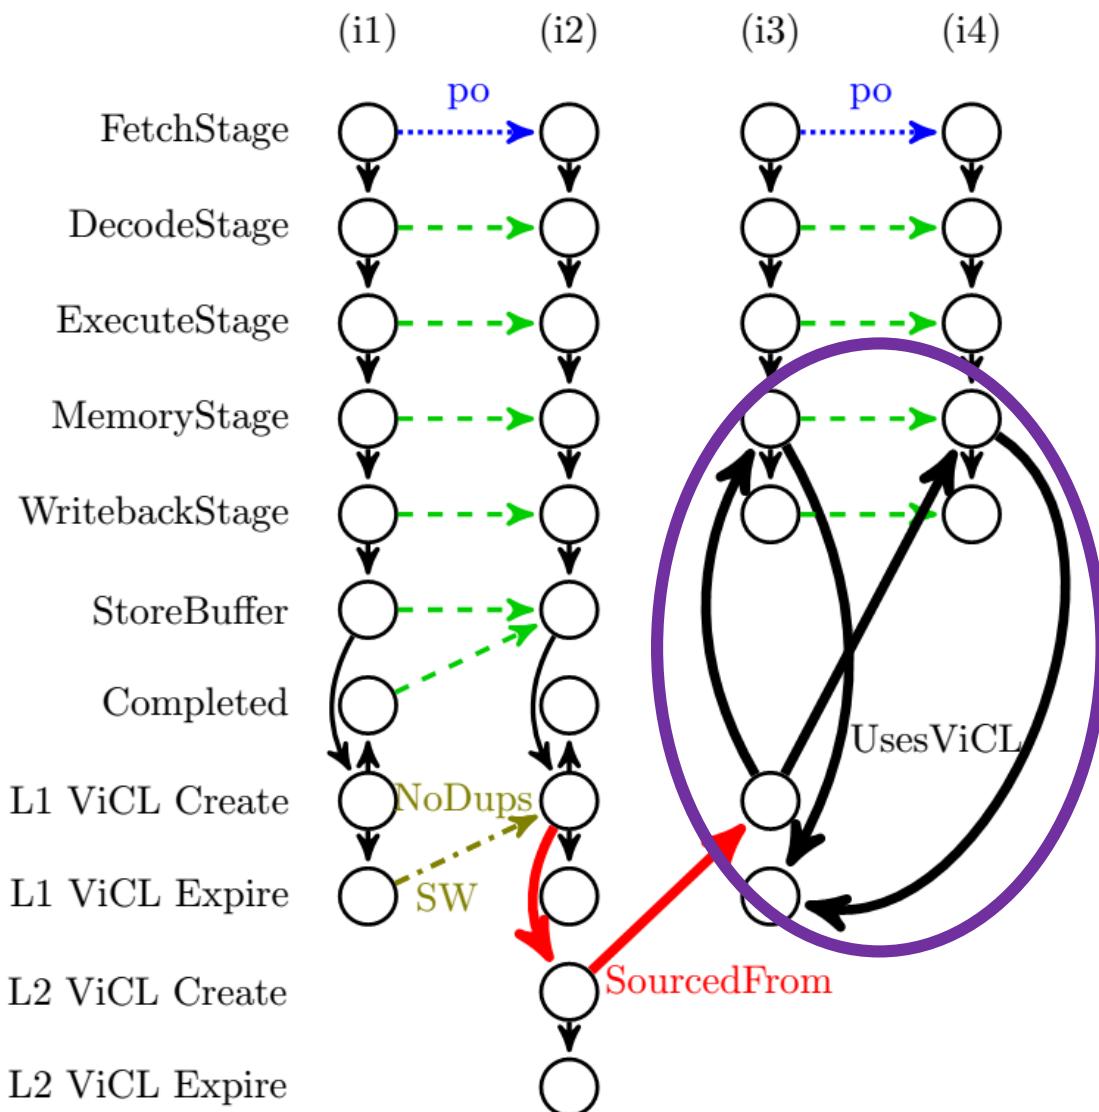
- ViCLs start at a **ViCL Create** event and end at a **ViCL Expire** event
  - Correspond to nodes in  $\mu$ hb graphs
  - Axioms over these nodes and edges enforce coherence and data movement orderings
- Use pipeline model from PipeCheck, but add ViCL nodes and edges

Litmus Test **co-mp**

Core 0	Core 1
(i1) St $[x] \leftarrow 1$	(i3) Ld $r1 \leftarrow [x]$
(i2) St $[x] \leftarrow 2$	(i4) Ld $r2 \leftarrow [x]$

In TSO: r1=2, r2=2 Allowed

# ViCLs in $\mu$ hb Graphs



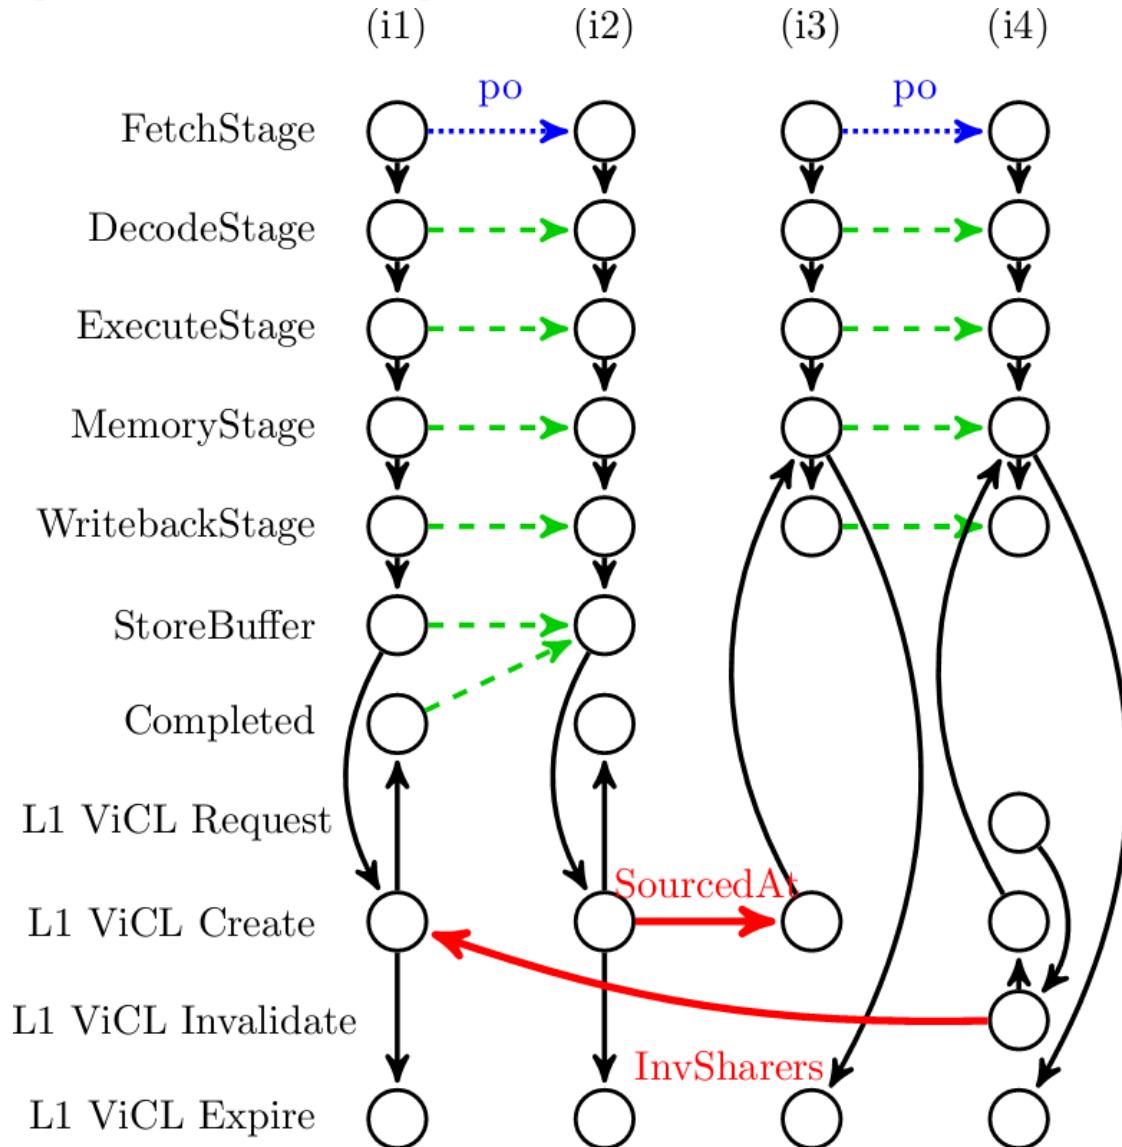
- ViCLs start at a **ViCL Create** event and end at a **ViCL Expire** event
  - Correspond to nodes in  $\mu$ hb graphs
  - Axioms over these nodes and edges enforce coherence and data movement orderings
- Use pipeline model from PipeCheck, but add ViCL nodes and edges

Litmus Test **co-mp**

Core 0	Core 1
(i1) St $[x] \leftarrow 1$	(i3) Ld r1 $\leftarrow [x]$
(i2) St $[x] \leftarrow 2$	(i4) Ld r2 $\leftarrow [x]$

In TSO: r1=2, r2=2 Allowed

# $\mu$ hb Graph for the Peekaboo Problem



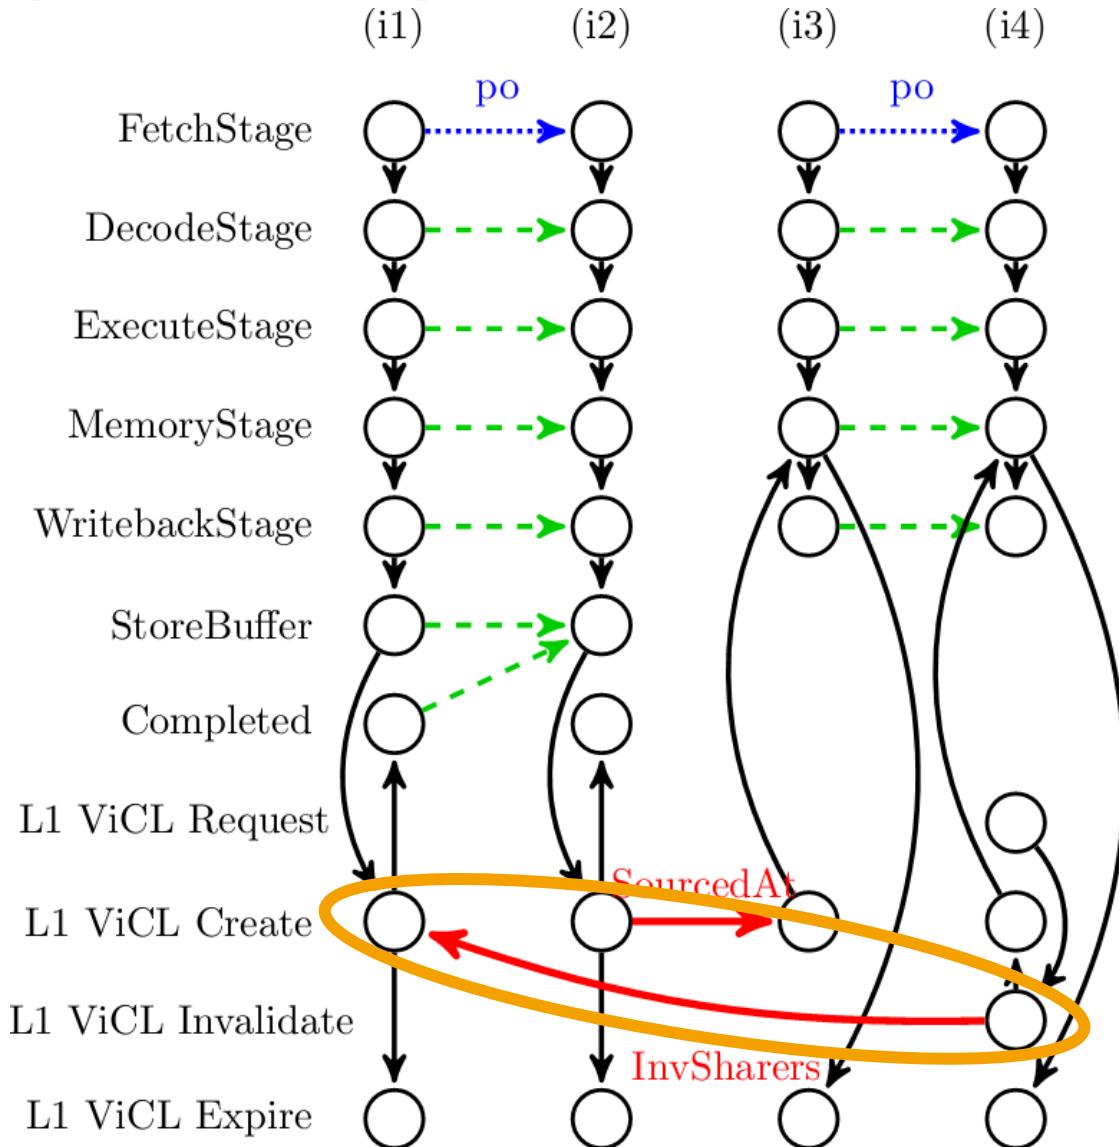
- Additional nodes represent ViCL requests and invalidations
- **Solution:** Invalidated data only usable if accessing load/store is oldest in program order at time of request [Sorin et al. Primer 2011]
- TSO-CC protocol [Elver and Nagarajan HPCA 2014] was vulnerable to variant of Peekaboo!
  - Now fixed

Core 0	Core 1
(i1) $[x] \leftarrow 1$	(i3) $r1 \leftarrow [y]$
(i2) $[y] \leftarrow 1$	(i4) $r2 \leftarrow [x]$

Under SC: Forbid  $r1=1, r2=0$



# $\mu$ hb Graph for the Peekaboo Problem



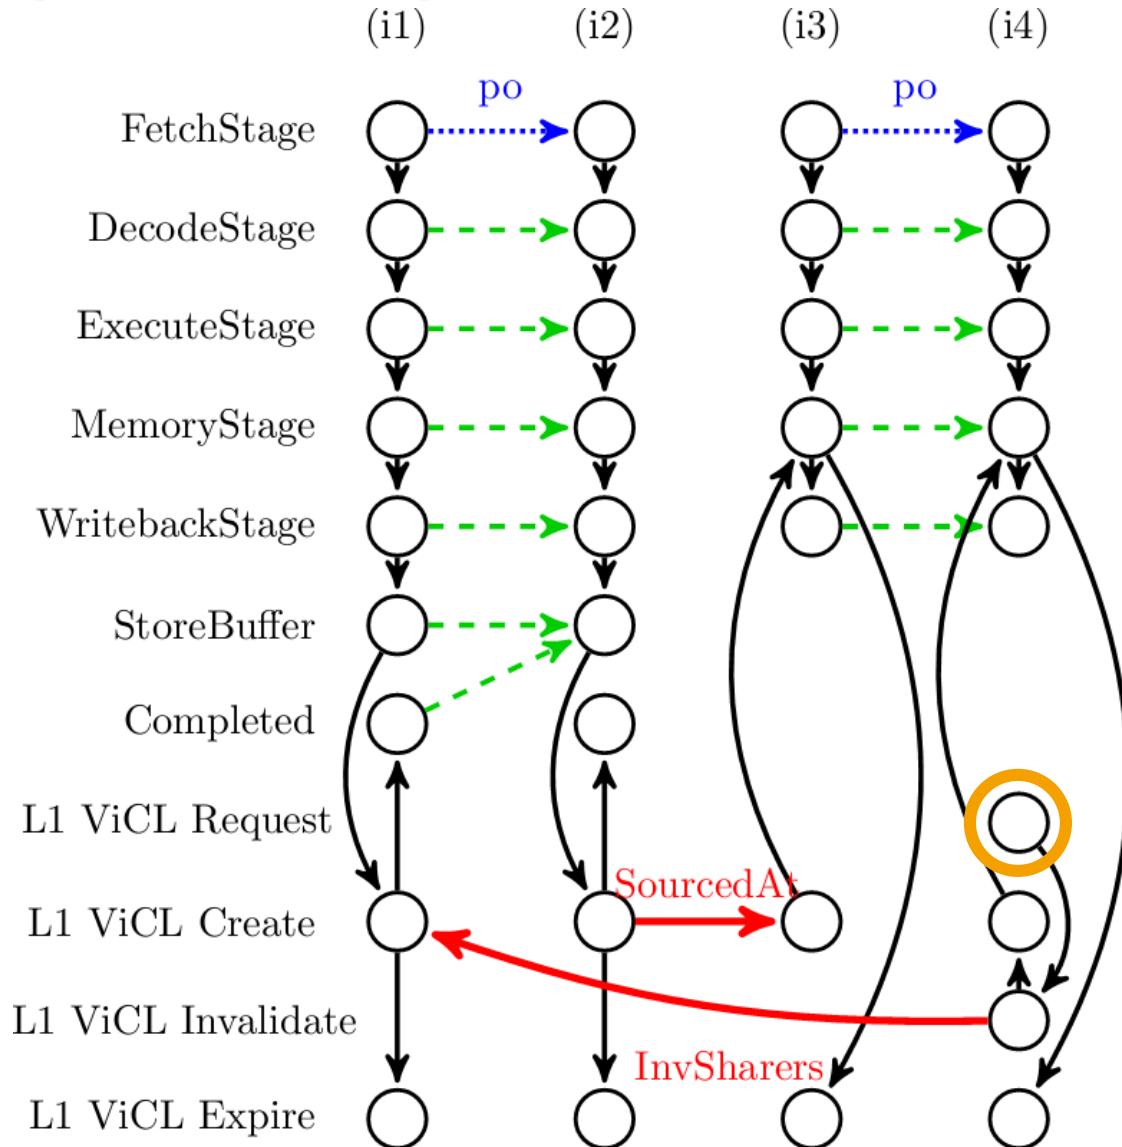
- Additional nodes represent ViCL requests and invalidations
- **Solution:** Invalidated data only usable if accessing load/store is oldest in program order at time of request [Sorin et al. Primer 2011]
- TSO-CC protocol [Elver and Nagarajan HPCA 2014] was vulnerable to variant of Peekaboo!
  - Now fixed

Core 0	Core 1
(i1) $[x] \leftarrow 1$	(i3) $r1 \leftarrow [y]$
(i2) $[y] \leftarrow 1$	(i4) $r2 \leftarrow [x]$

Under SC: Forbid  $r1=1, r2=0$



# $\mu$ hb Graph for the Peekaboo Problem

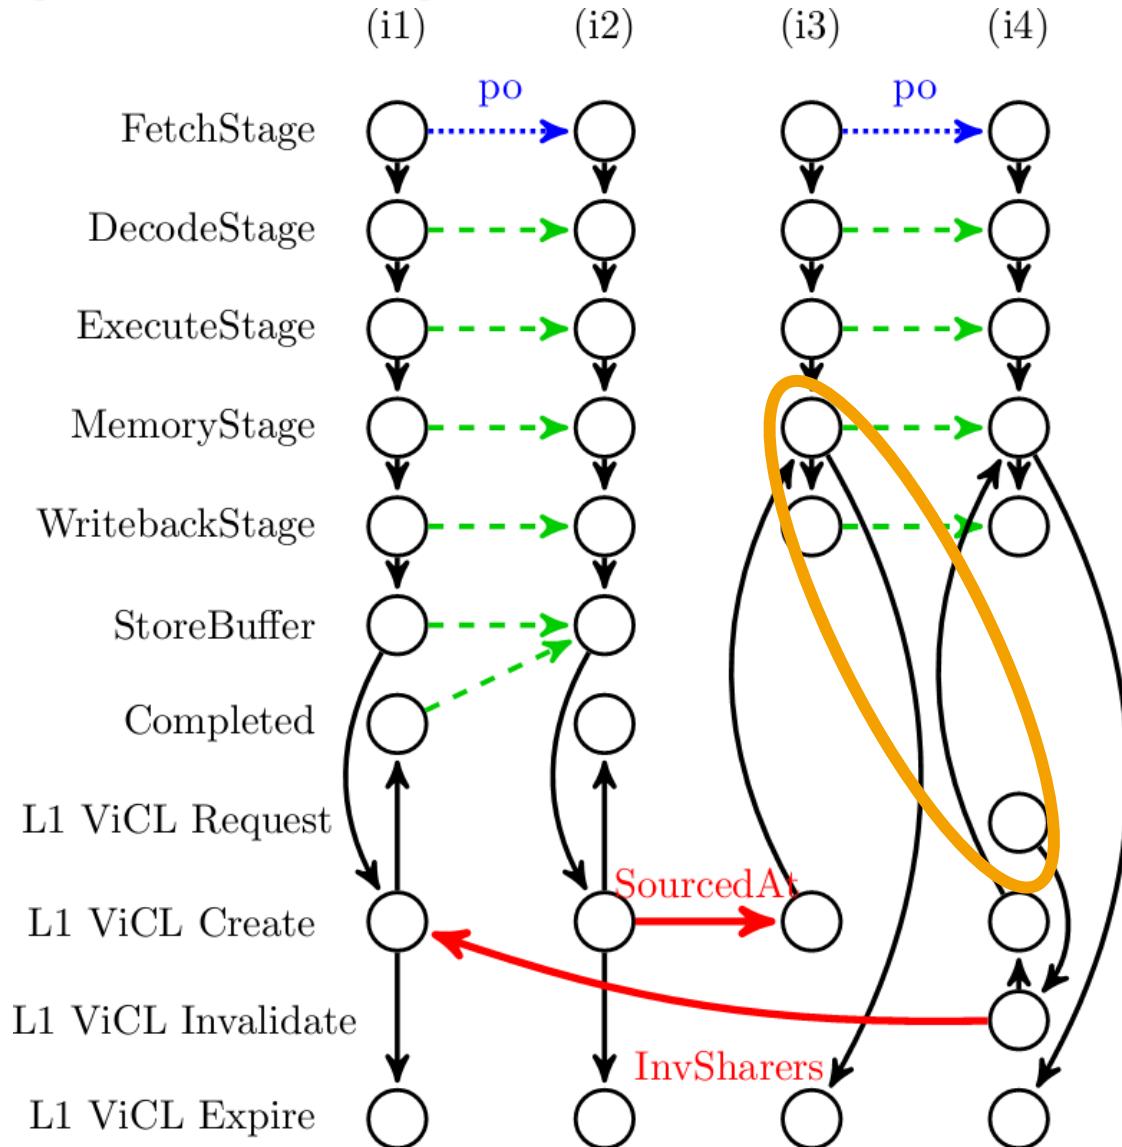


- Additional nodes represent ViCL requests and invalidations
- **Solution:** Invalidated data only usable if accessing load/store is oldest in program order at time of request [Sorin et al. Primer 2011]
- TSO-CC protocol [Elver and Nagarajan HPCA 2014] was vulnerable to variant of Peekaboo!
  - Now fixed

Core 0	Core 1
(i1) $[x] \leftarrow 1$	(i3) $r1 \leftarrow [y]$
(i2) $[y] \leftarrow 1$	(i4) $r2 \leftarrow [x]$

Under SC: Forbid  $r1=1, r2=0$

# $\mu$ hb Graph for the Peekaboo Problem



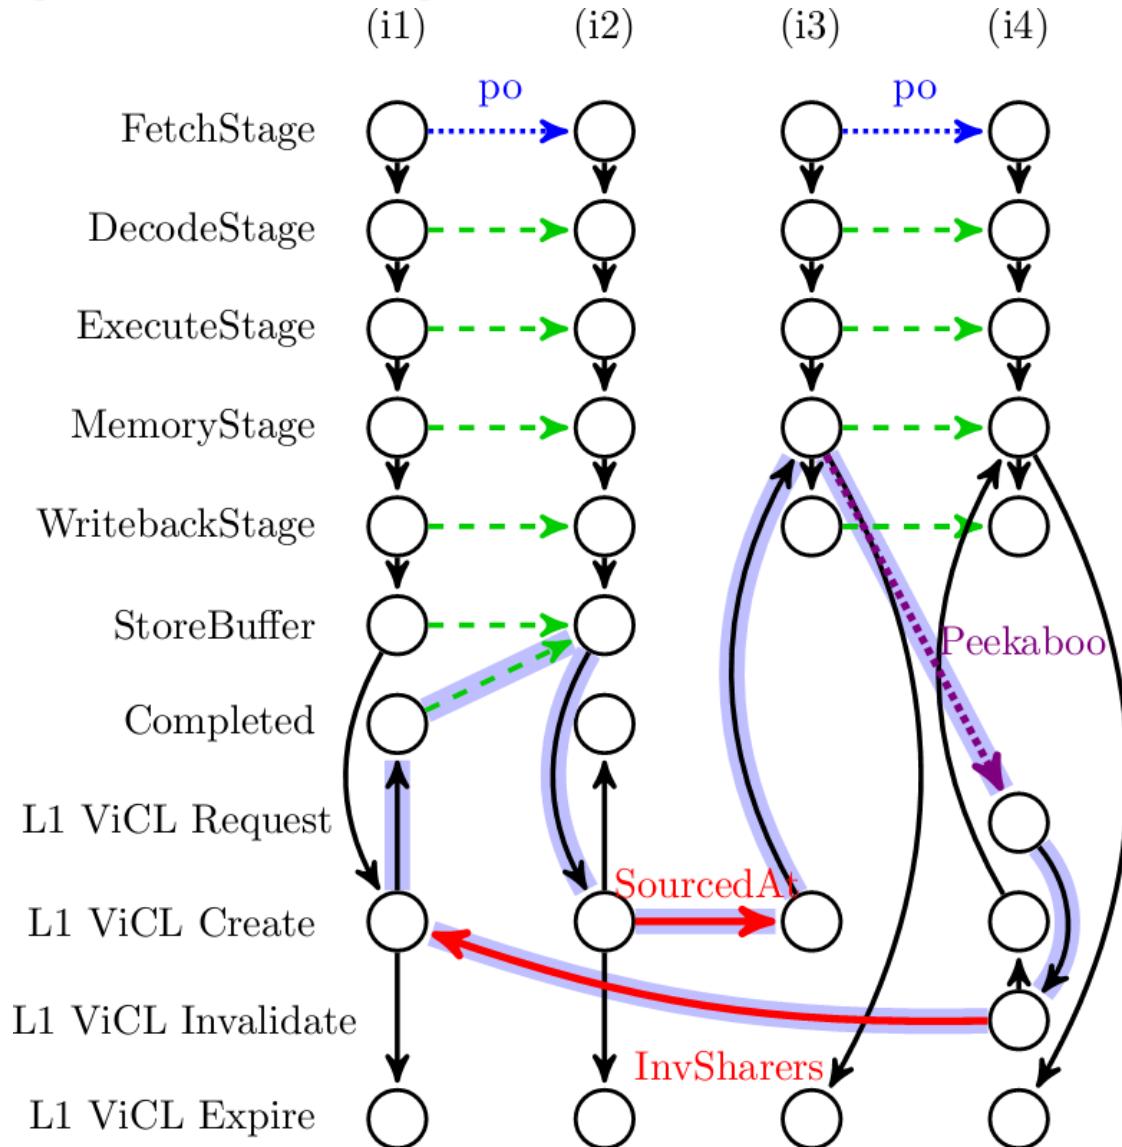
- Additional nodes represent ViCL requests and invalidations
- **Solution:** Invalidated data only usable if accessing load/store is oldest in program order at time of request [Sorin et al. Primer 2011]
- TSO-CC protocol [Elver and Nagarajan HPCA 2014] was vulnerable to variant of Peekaboo!
  - Now fixed

Core 0	Core 1
(i1) $[x] \leftarrow 1$	(i3) $r1 \leftarrow [y]$
(i2) $[y] \leftarrow 1$	(i4) $r2 \leftarrow [x]$

Under SC: Forbid  $r1=1, r2=0$



# $\mu$ hb Graph for the Peekaboo Problem



- Additional nodes represent ViCL requests and invalidations
- **Solution:** Invalidated data only usable if accessing load/store is oldest in program order at time of request [Sorin et al. Primer 2011]
- TSO-CC protocol [Elver and Nagarajan HPCA 2014] was vulnerable to variant of Peekaboo!
  - Now fixed

Core 0	Core 1
(i1) $[x] \leftarrow 1$	(i3) $r1 \leftarrow [y]$
(i2) $[y] \leftarrow 1$	(i4) $r2 \leftarrow [x]$

Under SC: Forbid  $r1=1, r2=0$

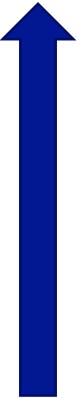
# CCICheck Takeaways

- Coherence & consistency often closely coupled in implementations
- In such cases, coherence & consistency cannot be verified separately
- **CCICheck: CCI-aware microarchitectural MCM checking**
  - Uses ViCL (Value in Cache Lifetime) abstraction
- Discovered bug in TSO-CC lazy coherence protocol

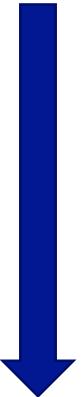


# ISA-level MCMs in the Hardware-Software Stack

High-Level Languages (HLLs)



New ISA-level MCM

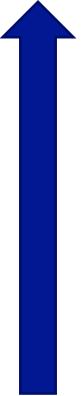


Hardware



# ISA-level MCMs in the Hardware-Software Stack

High-Level Languages (HLLs)



New ISA-level MCM

Which orderings  
must be guaranteed  
by hardware?



Hardware



# ISA-level MCMs in the Hardware-Software Stack

High-Level Languages (HLLs)

Which orderings does  
the compiler need to  
enforce?

New ISA-level MCM

Which orderings  
must be guaranteed  
by hardware?

Hardware



# ISA-level MCMs in the Hardware-Software Stack

High-Level Languages (HLLs)

Which orderings does

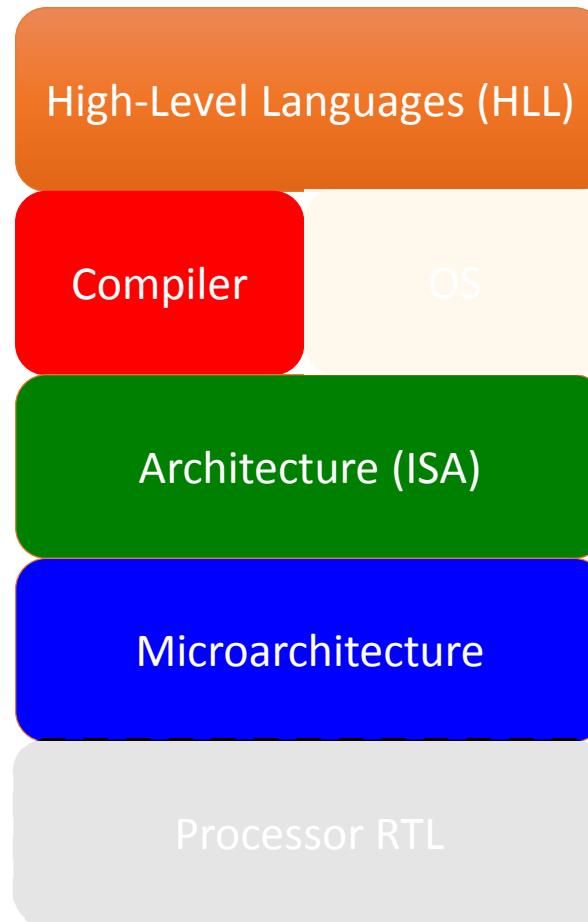
**TriCheck checks that HLL, compiler, ISA, and  
hardware align on MCM requirements**

Which orderings  
must be guaranteed  
by hardware?

Hardware



# TriCheck: Layers of the Stack are Intertwined



- ISA-level MCMs should allow microarchitectural optimizations but also be compatible with HLLs
- **TriCheck [Trippel et al. ASPLOS 2017]** enables holistic analysis of HLL memory model, ISA-level MCM, compiler mappings, and microarchitectures
  - **Mapping:** translation of HLL synchronization primitives to one or more assembly language instructions
- Also useful for checking HLL compiler mappings to ISA-level MCMs
- Selected as one of 12 “*Top Picks of Comp. Arch. Conferences*” for 2017



# TriCheck: Comparing HLL to Microarchitecture

HLL  
Model  
e.g. C11

HLL Litmus  
Test Variants

HLL to ISA  
Compiler  
Mapping

$\mu$ spec  
Microarch.  
Model

Four Primary Inputs



# TriCheck: Comparing HLL to Microarchitecture

HLL  
Model  
e.g. C11

HLL Litmus  
Test Variants

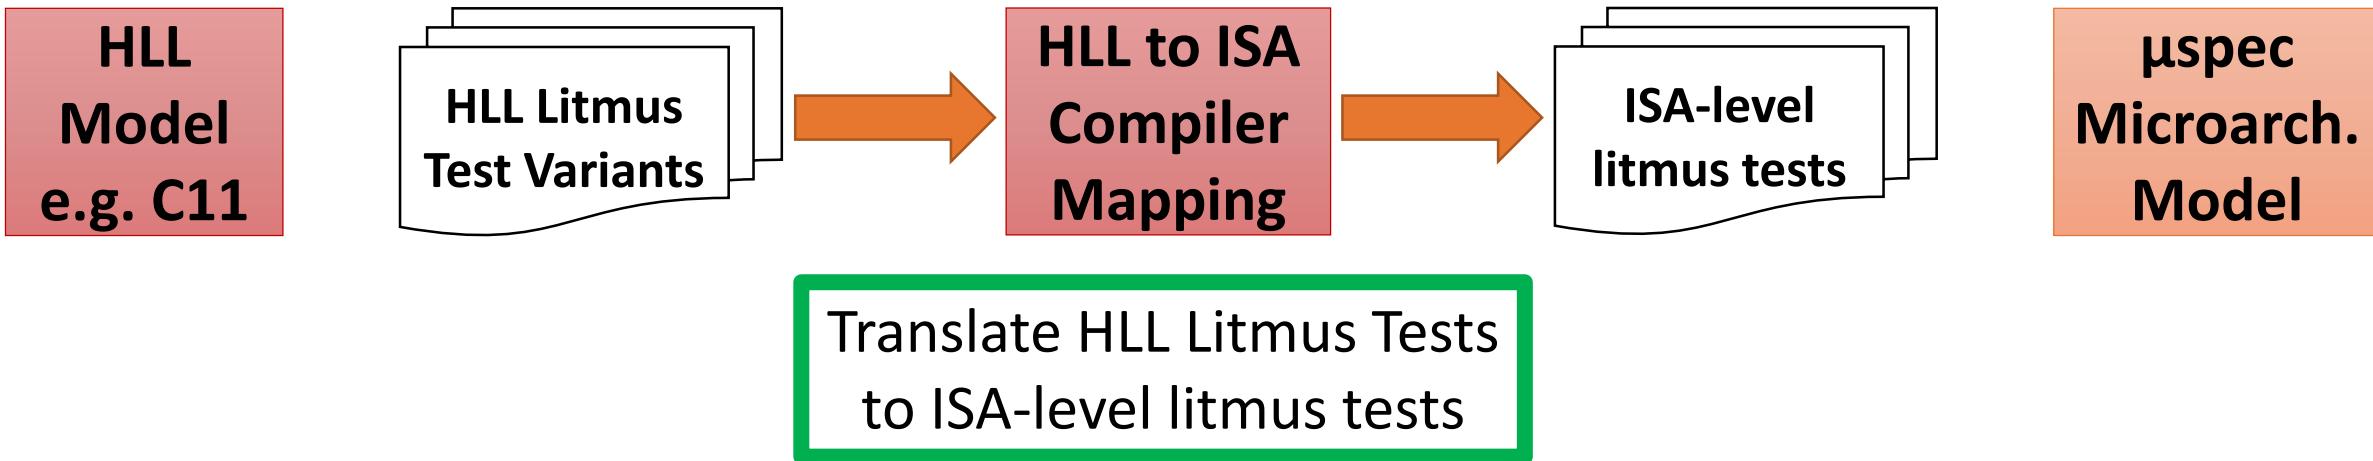
HLL to ISA  
Compiler  
Mapping

$\mu$ spec  
Microarch.  
Model

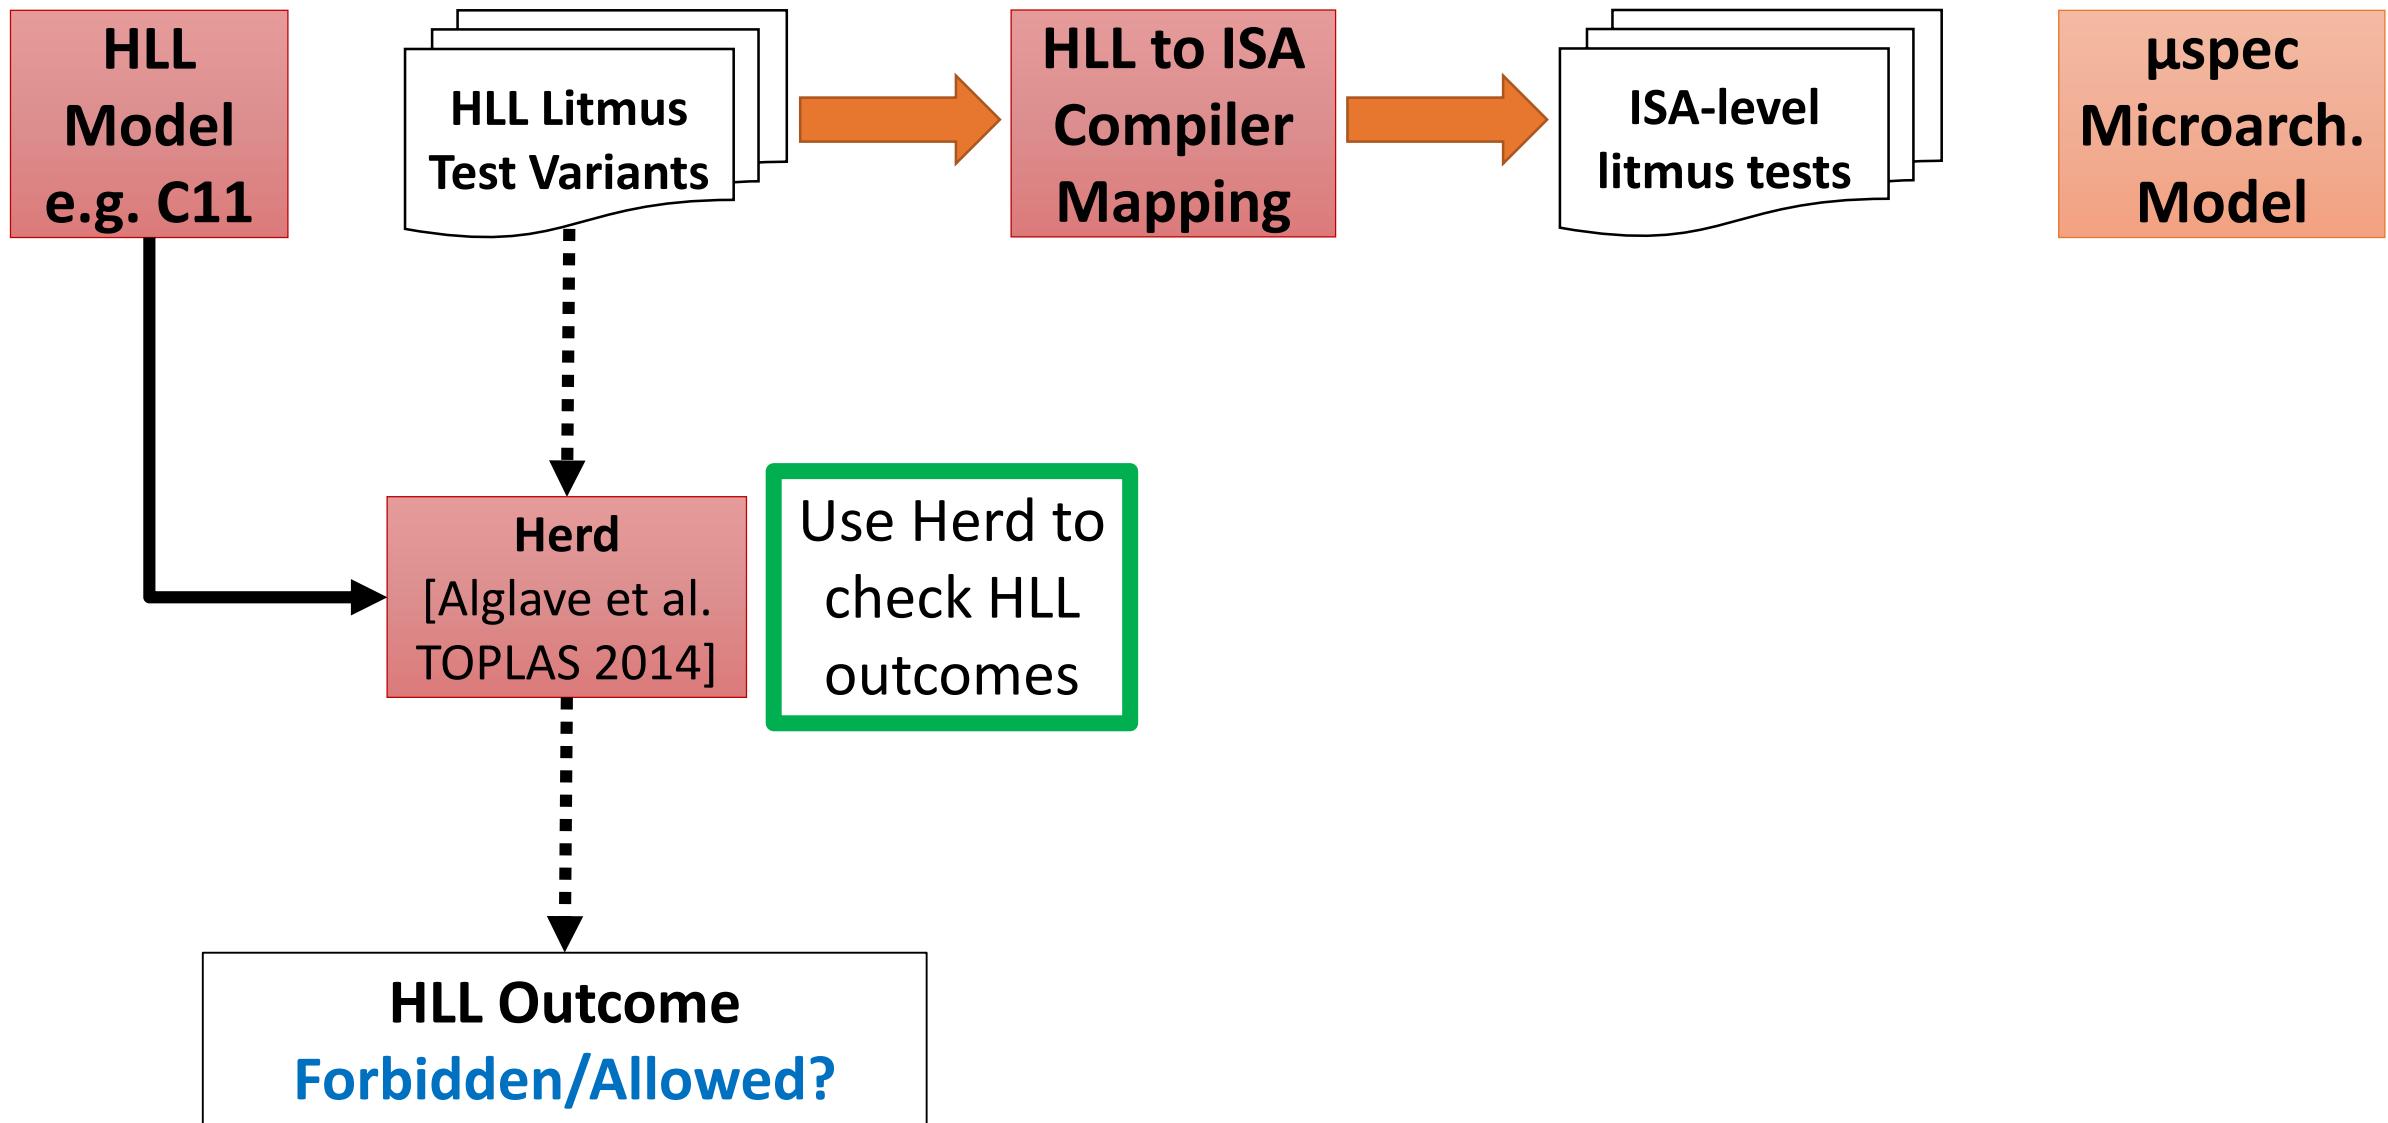
Examine all C11  
**memory\_order**  
combinations  
**(release, acquire,  
relaxed, seq\_cst)**  
for HLL litmus tests



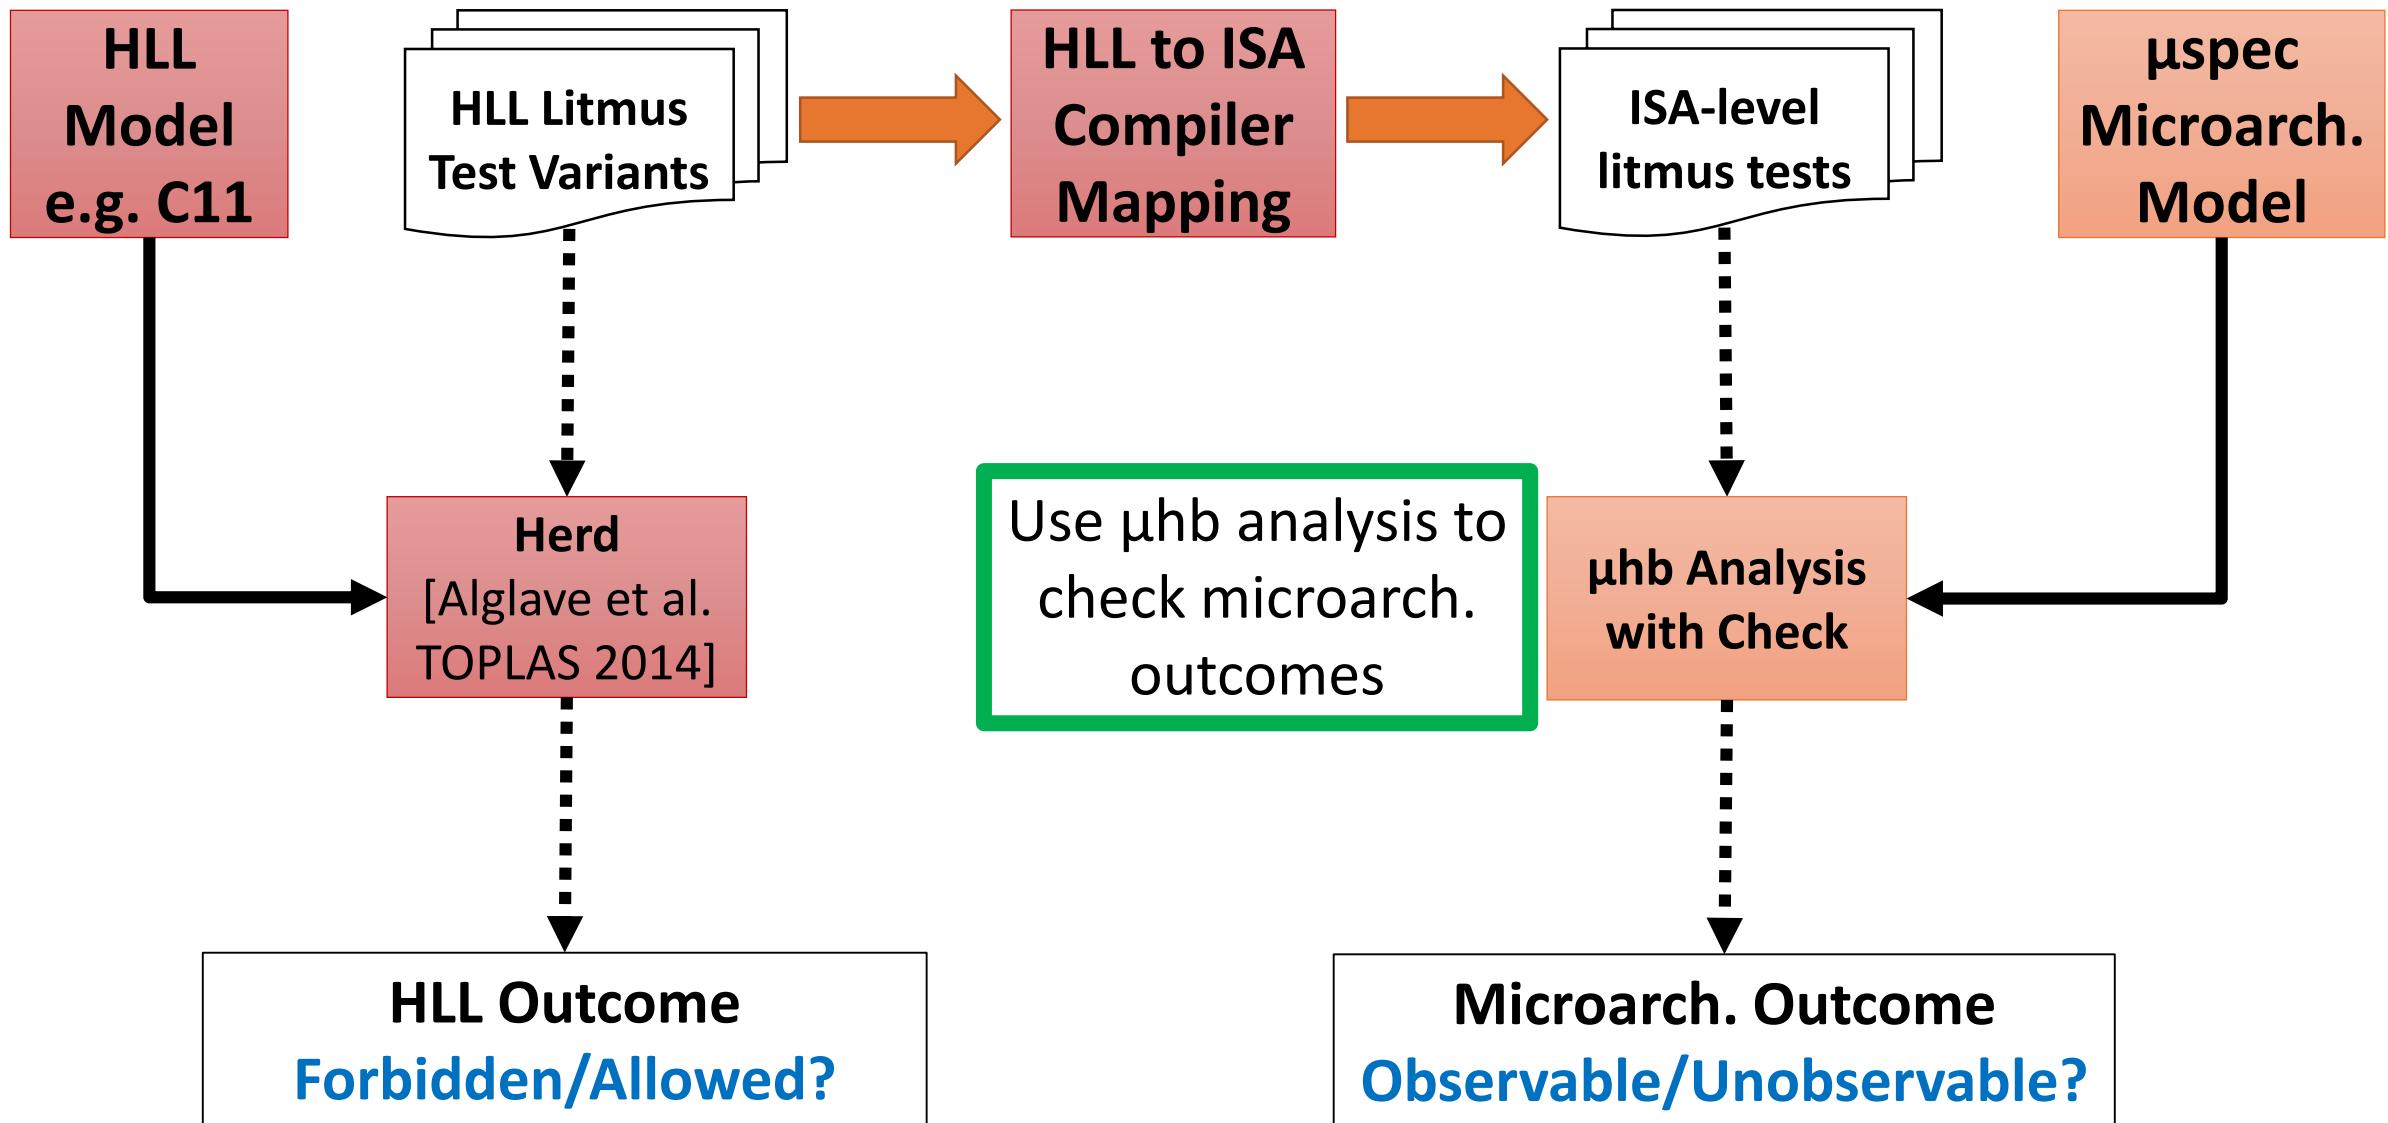
# TriCheck: Comparing HLL to Microarchitecture



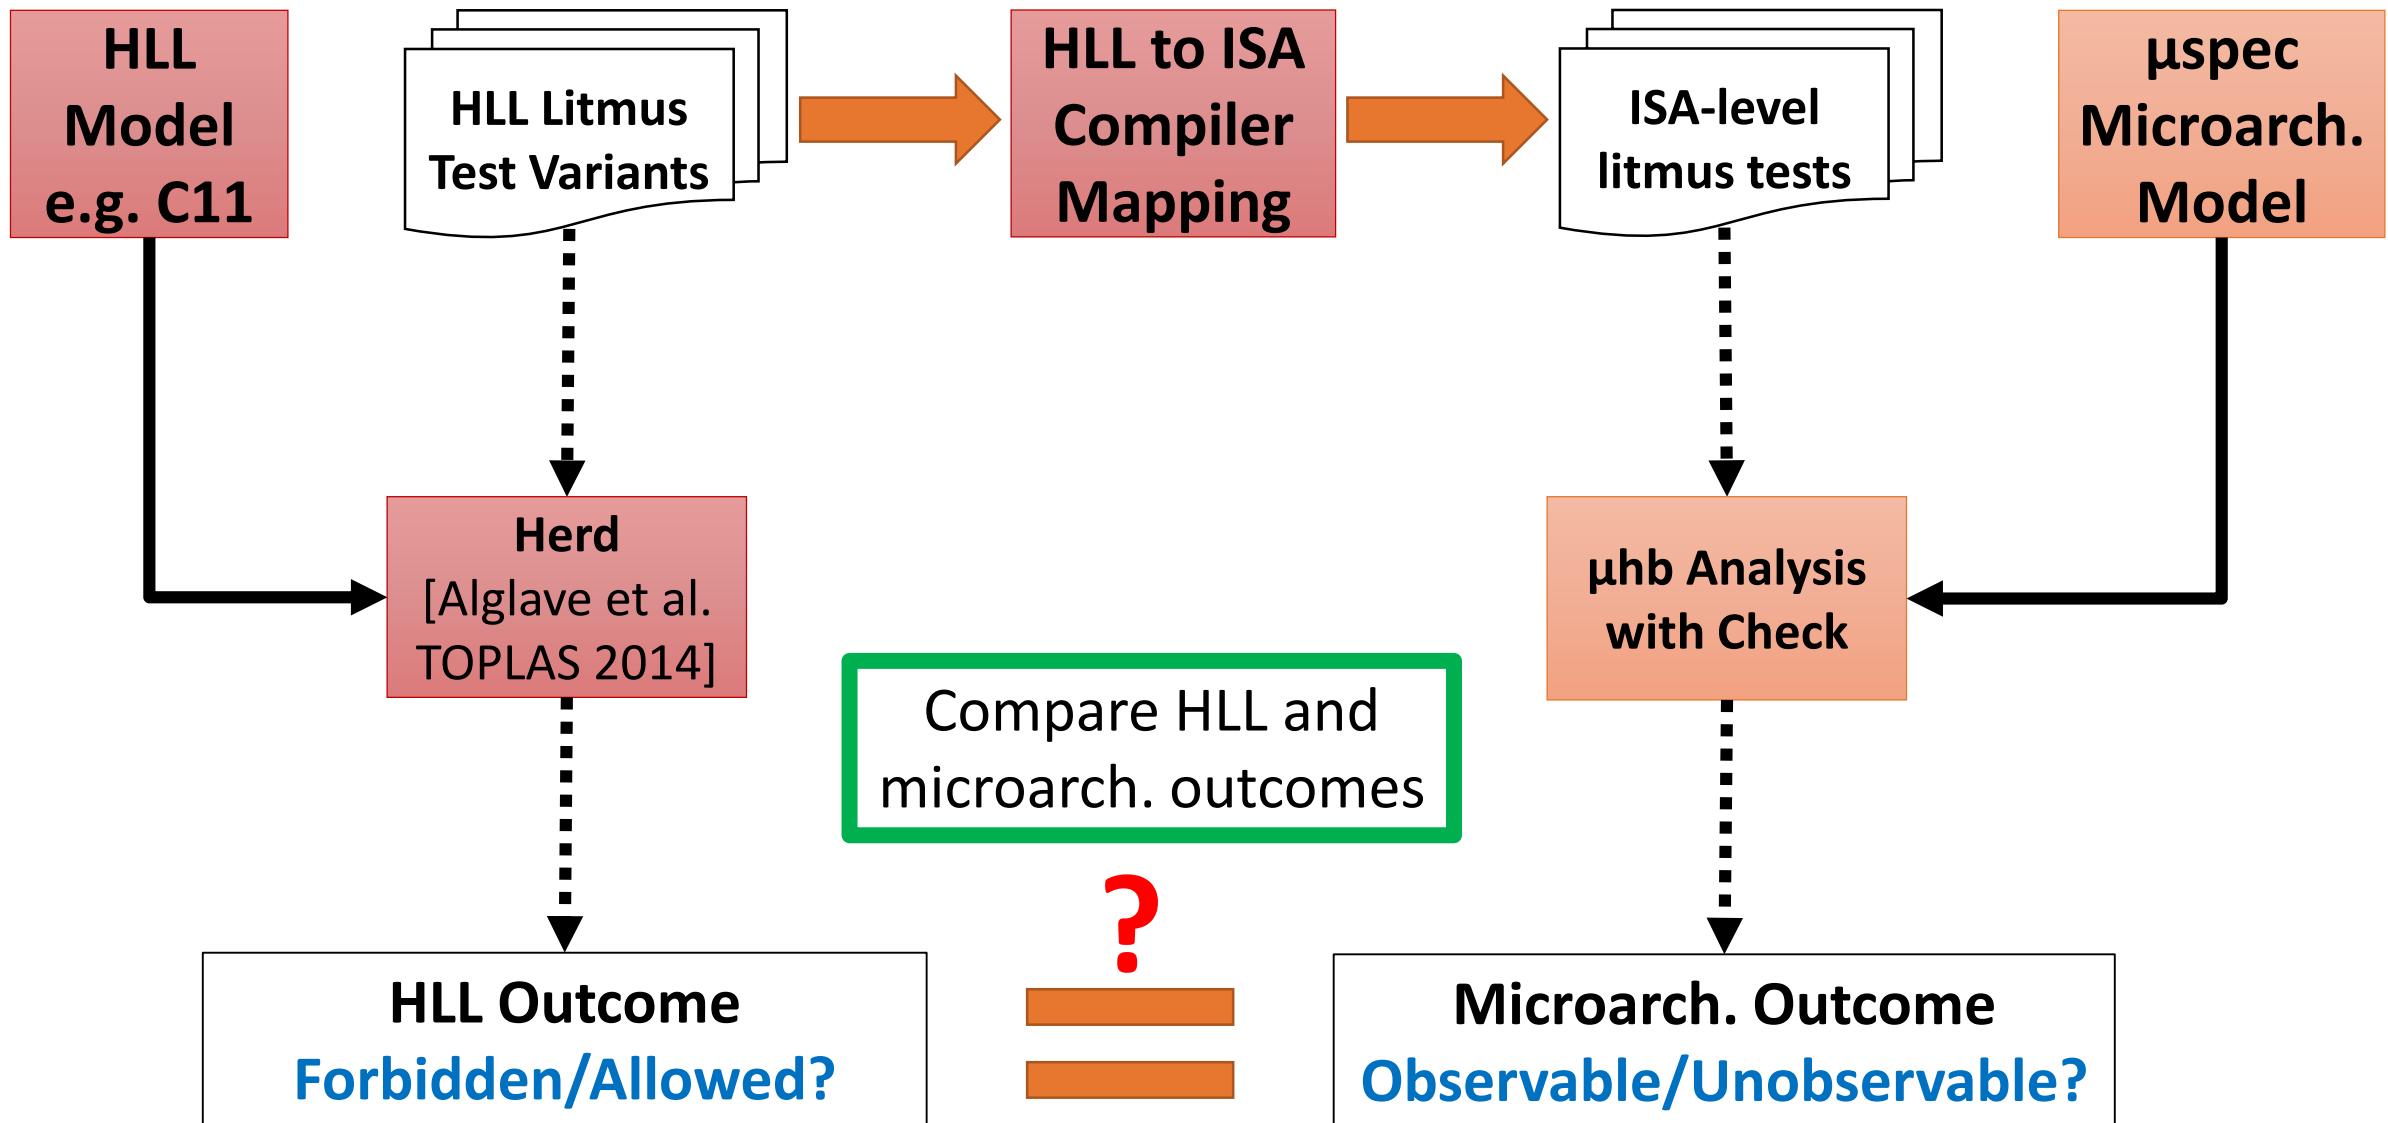
# TriCheck: Comparing HLL to Microarchitecture



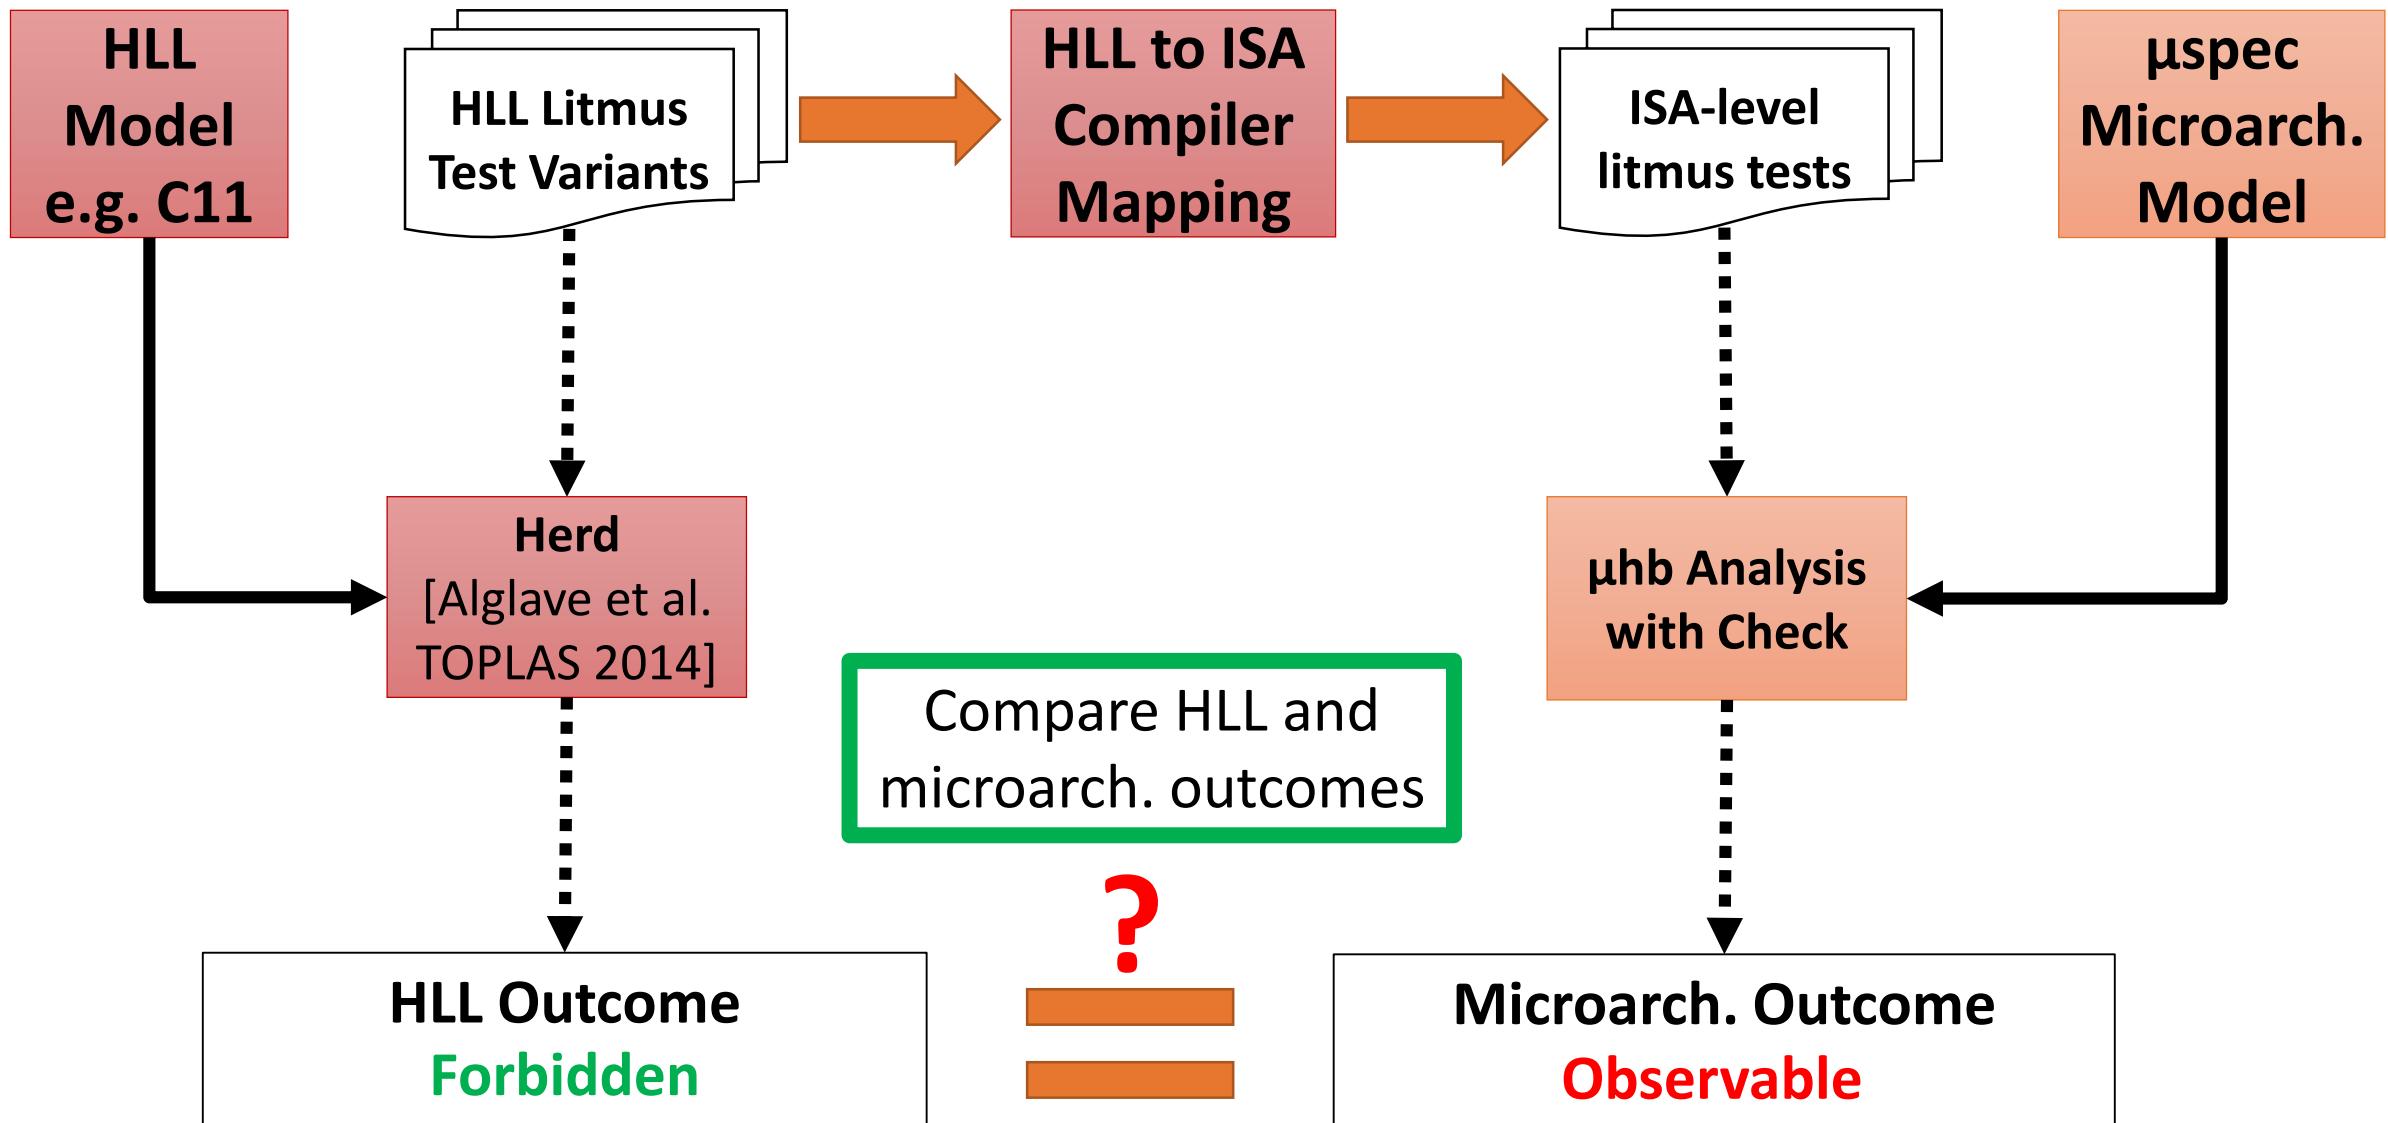
# TriCheck: Comparing HLL to Microarchitecture



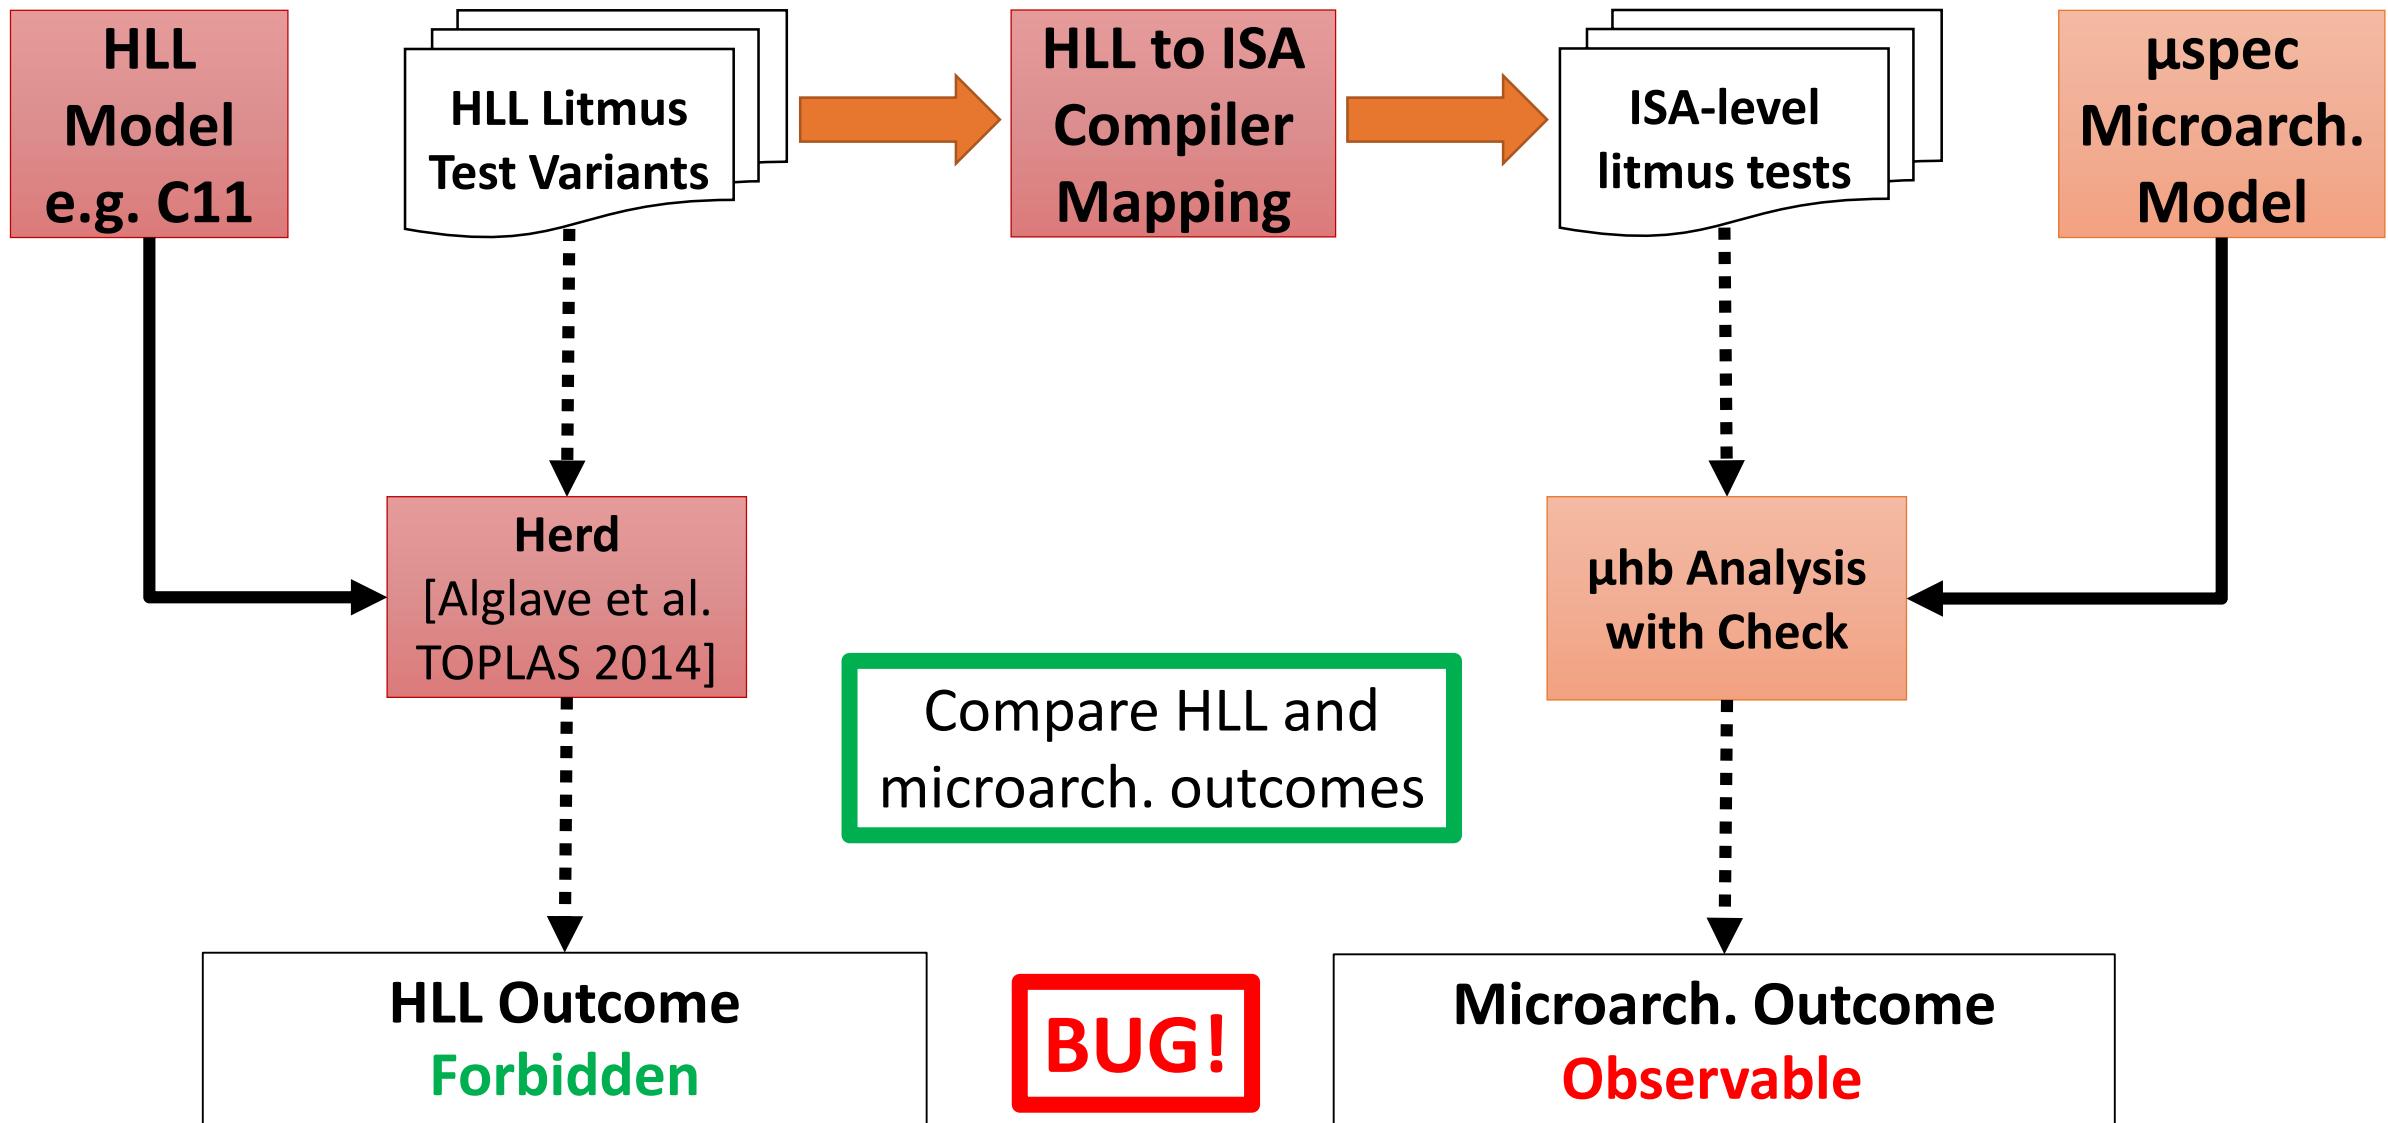
# TriCheck: Comparing HLL to Microarchitecture



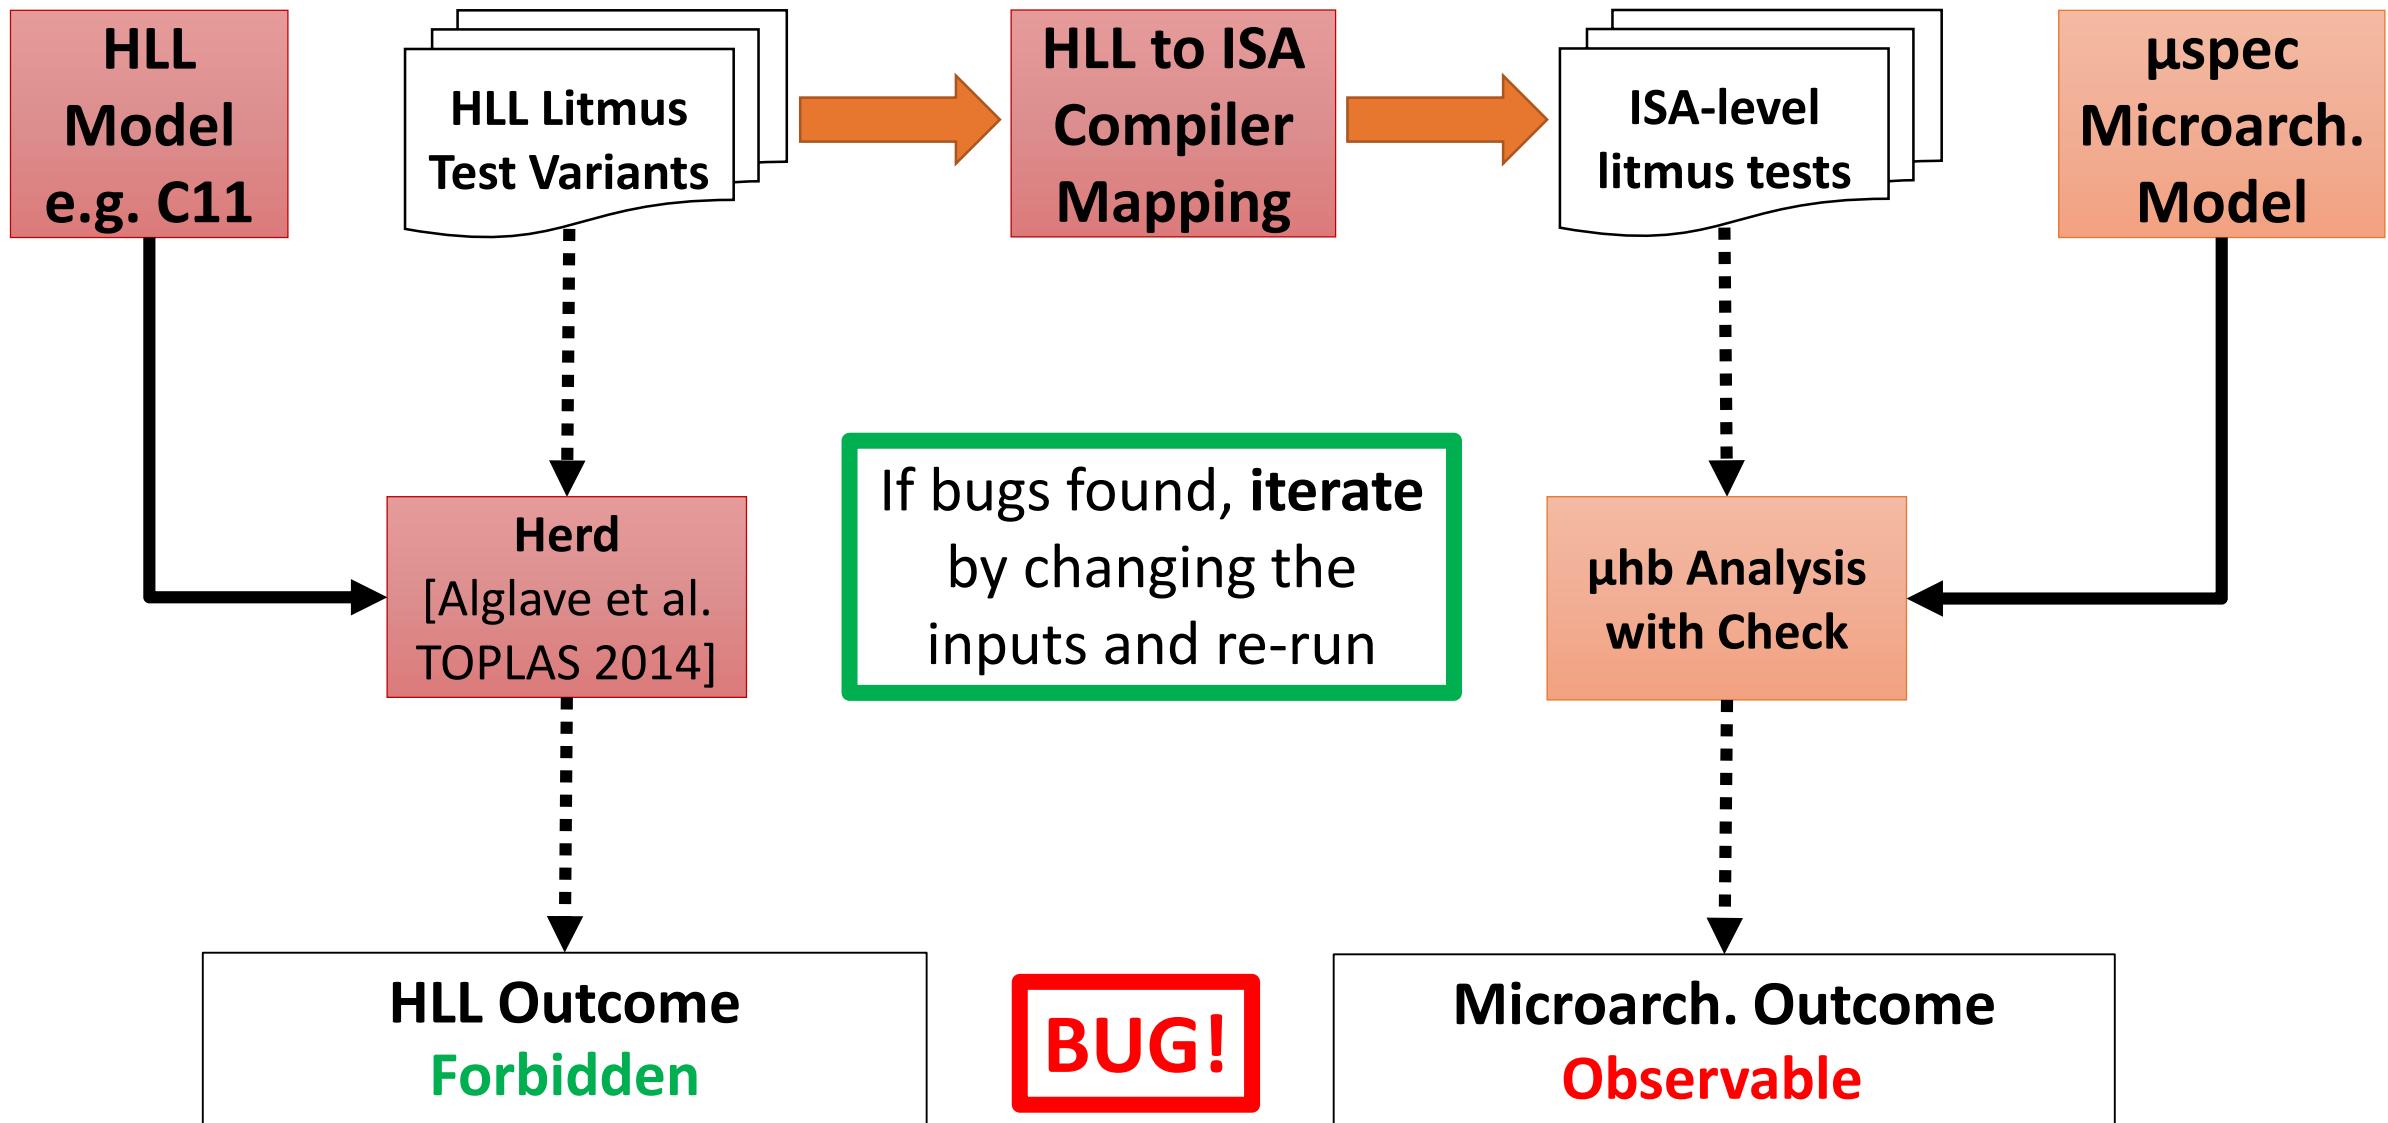
# TriCheck: Comparing HLL to Microarchitecture



# TriCheck: Comparing HLL to Microarchitecture



# TriCheck: Comparing HLL to Microarchitecture



# Using TriCheck for ISA MCM Design: RISC-V

- Ran TriCheck on draft RISC-V ISA MCM with
  - C11 HLL MCM [Batty et al. POPL 2011] [Batty et al. POPL 2016]
  - Compiler mappings based on RISC-V manual
  - Variety of microarchitectures that relaxed various memory orderings
    - **All legal according to draft RISC-V spec**
    - Ranging from SC microarchitecture to one with reorderings allowed by ARM/Power
- Draft RISC-V MCM for Base ISA **incapable** of correctly compiling C11:
  - C11 outcome forbidden, but **impossible** to forbid on hardware
  - RISC-V fences **too weak** to restore orderings that implementations could relax

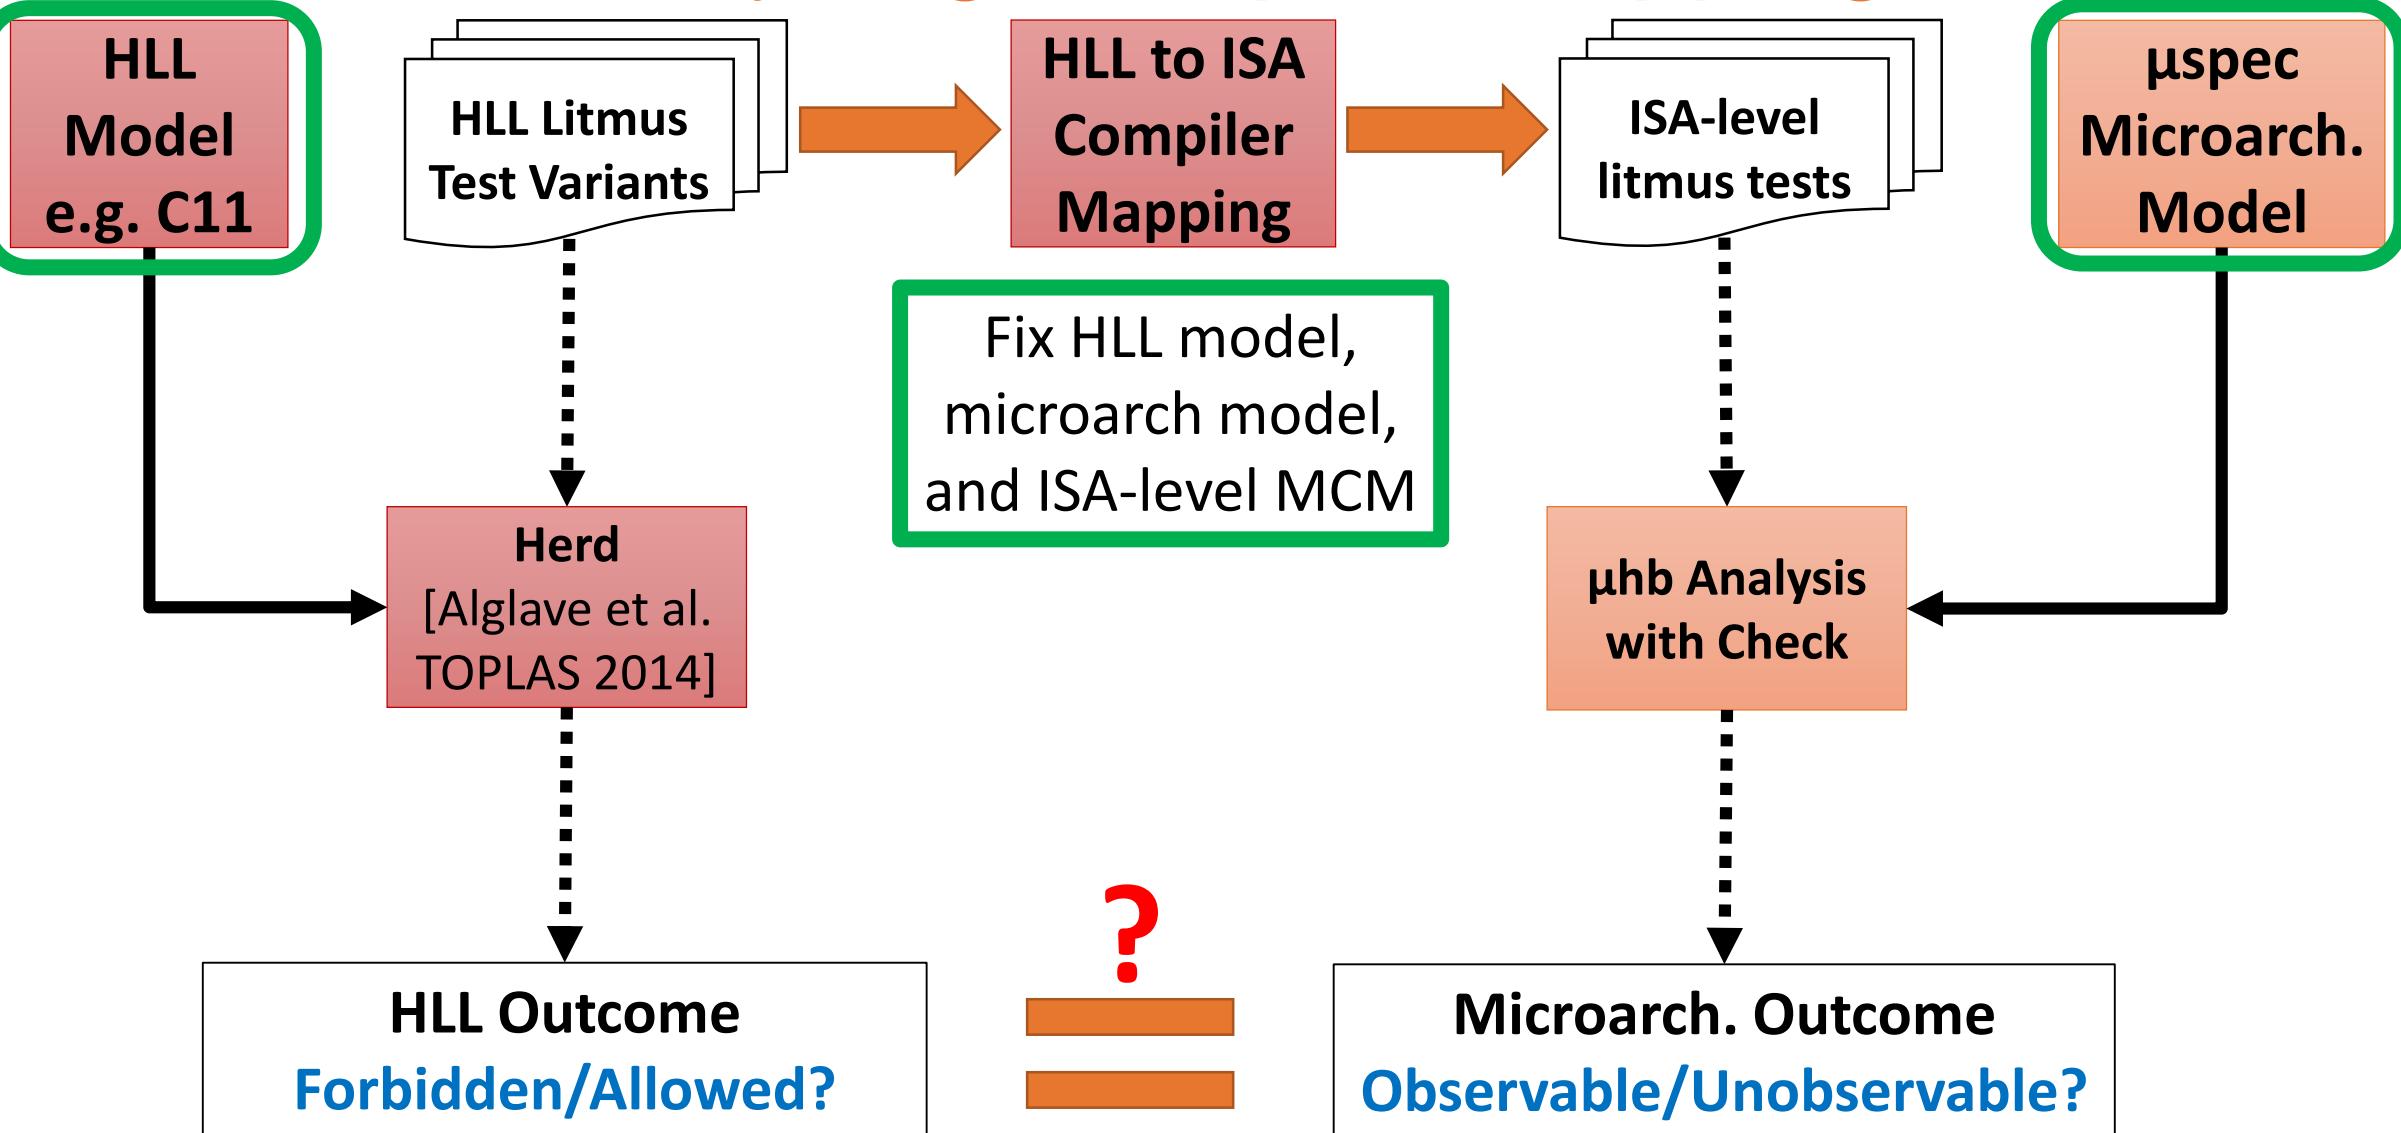


# Current RISC-V Status

- In response to our findings, RISC-V Memory Model Working Group was formed (we are members)
  - Mandate to create an MCM for RISC-V that satisfies community needs
- **Working Group has developed an MCM proposal that fixes the aforementioned bugs (and other issues)**
- **MCM proposal recently passed the 45-day public feedback period!**
  - Well on its way to being included in the next version of the RISC-V ISA spec



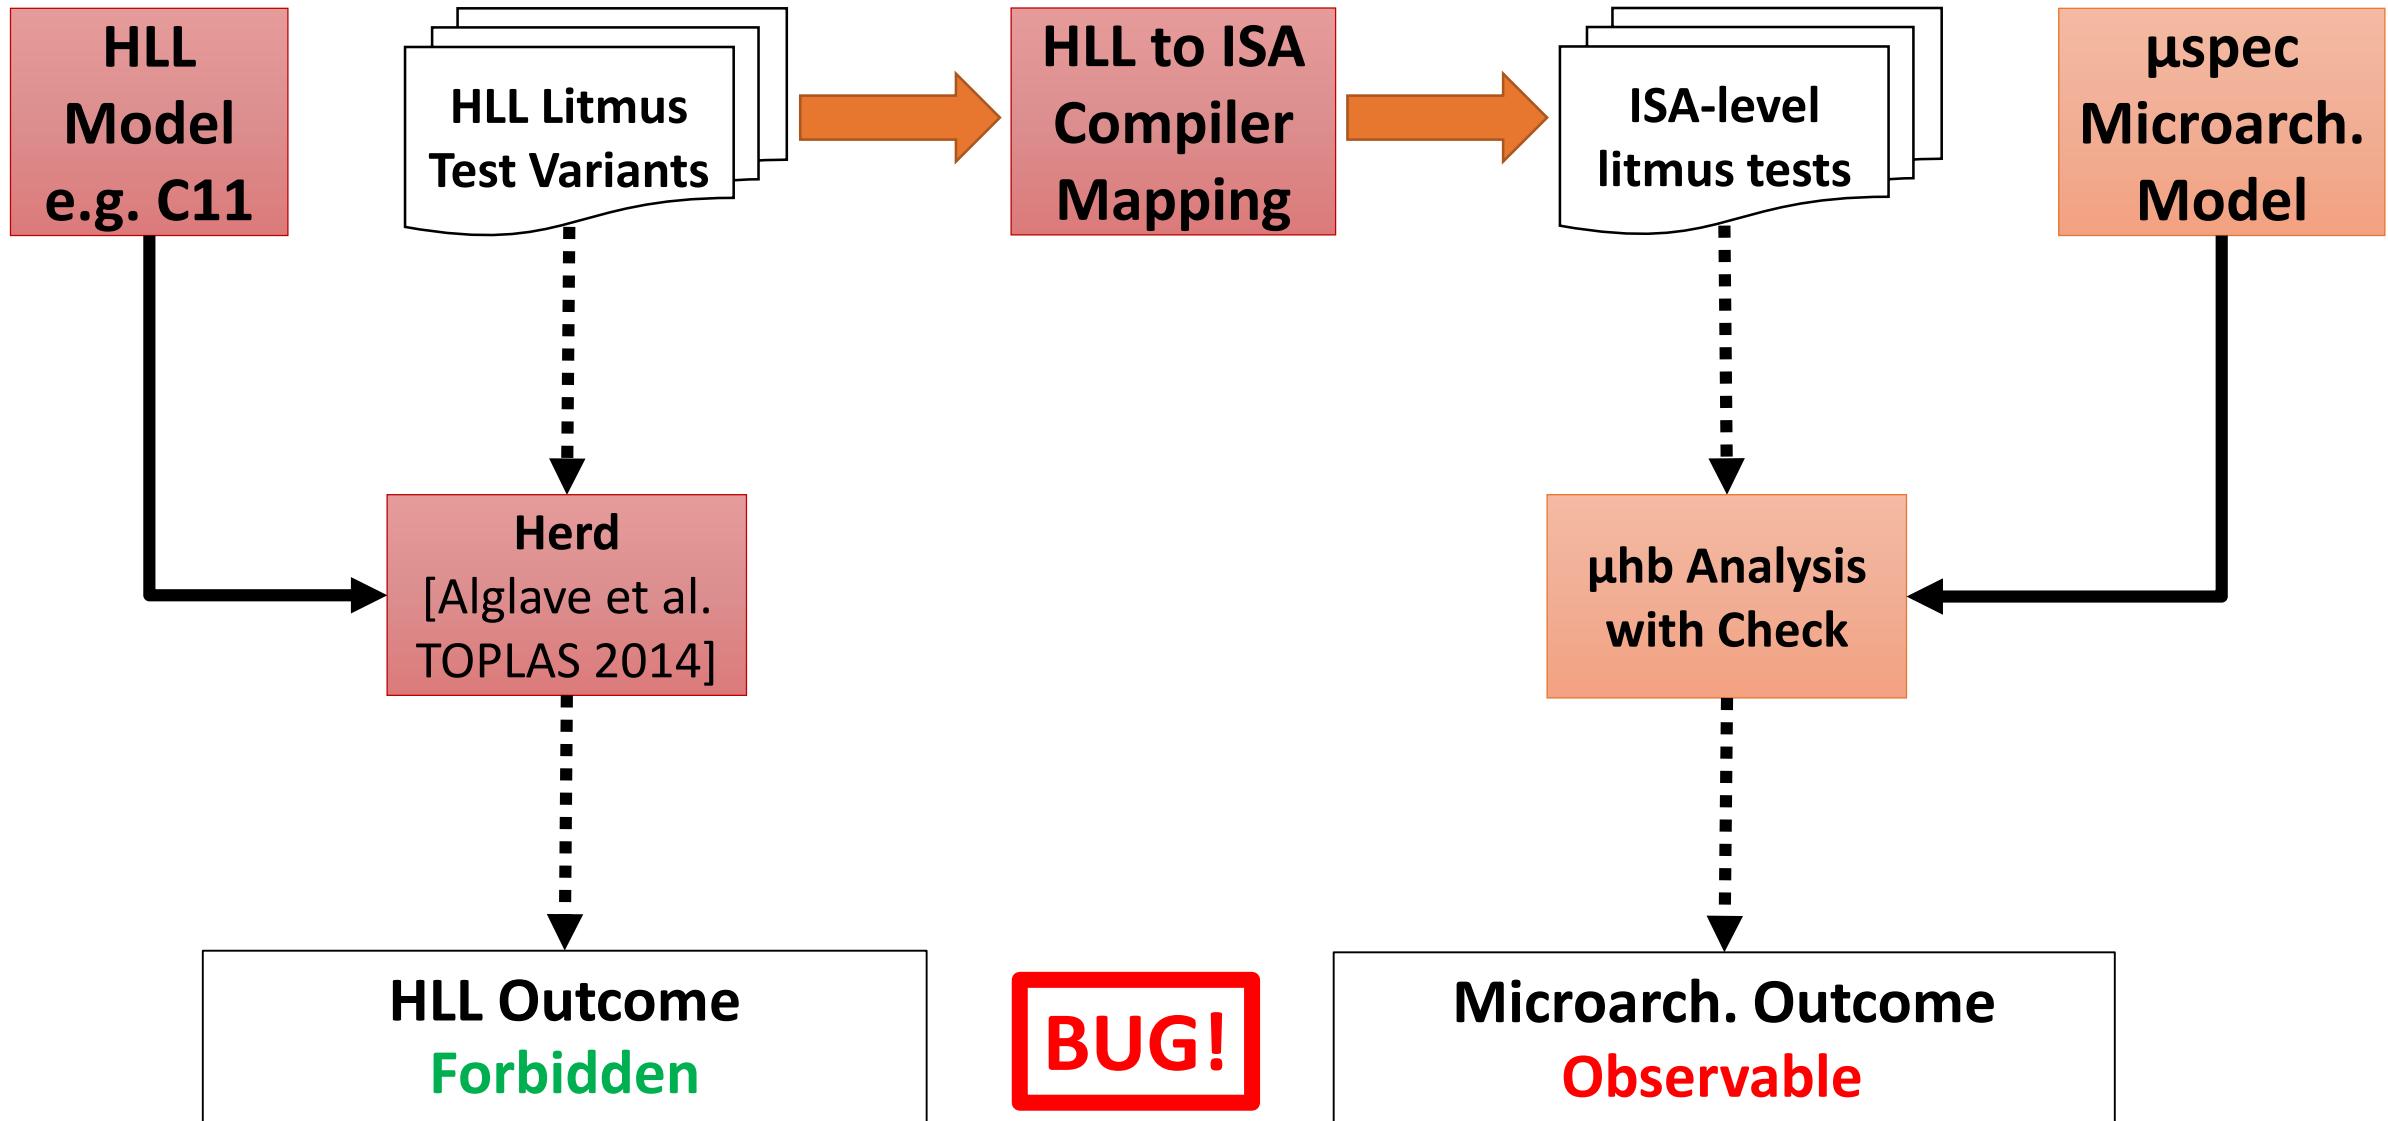
# TriCheck: Analysing Compiler Mappings



?



# TriCheck: Analysing Compiler Mappings



# Checking C11 Mappings to ARMv7/Power

- Ran TriCheck on microarch. with reordering similar to ARMv7/Power
  - Utilised “trailing-sync” compiler mapping [Batty et al. POPL 2012]
  - Discovered 2 cases where C11 outcome **forbidden**, but **allowed** by hardware!
  - Deduced that the mapping must be flawed
- Mapping was supposedly proven correct [Batty et al. POPL 2012]
  - Traced the loophole in the proof [Manerkar et al. CoRR’16]
- **Problem: C11 model slightly too strong for mappings**
  - C11 has happens-before (*hb*) ordering and total order on all SC accesses (*sc*)
  - *hb* and *sc* orders must agree with each other
  - **Trailing-sync mapping does not guarantee this for our counterexamples**



# Current state of C11

- “Leading-sync” mapping [McKenney and Silvera 2011]
  - Counterexample discovered concurrently to us [Lahav et al. PLDI 2017]
- **Both mappings currently broken**
- Possible solutions under discussion by C11 memory model committee:
  - RC11 [Lahav et al. PLDI 2017]: remove req. that *sc* and *hb* orders agree
    - Current mappings work, but reduces intuition in an already complicated C11 model
  - Adding extra fences to mappings
    - low performance, requires recompilation, counterexample pattern not common

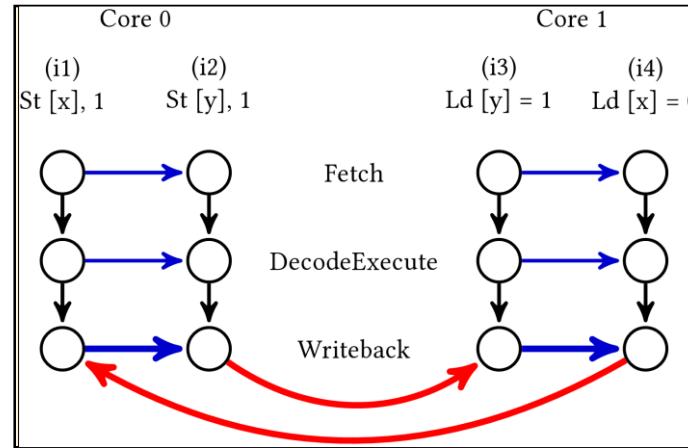


# TriCheck Takeaways

- Both HLL memory models and microarchitectural optimizations influence the design of ISA-level MCMs
- **TriCheck** enables holistic analysis of HLL memory model, ISA-level MCM, compiler mappings, and microarchitectural implementations
- TriCheck discovered numerous issues with draft RISC-V MCM
  - Influenced the design of the new RISC-V MCM
- Discovered two counterexamples to C11 -> ARMv7/Power compiler mappings
  - Mappings were previously “proven” correct; isolated flaw in proof



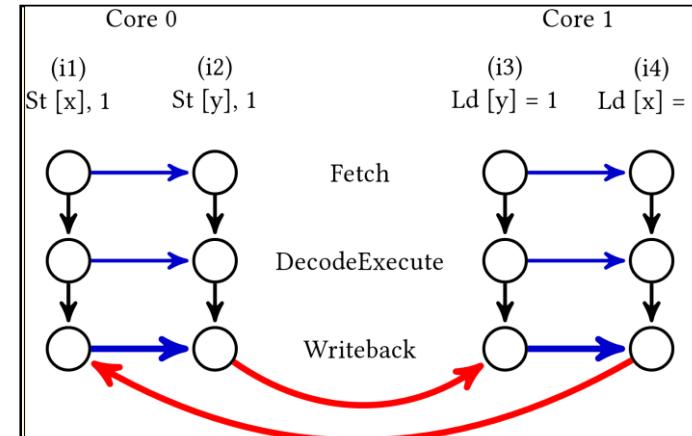
# Memory Consistency Checking for RTL



**Microarchitecture Checking**



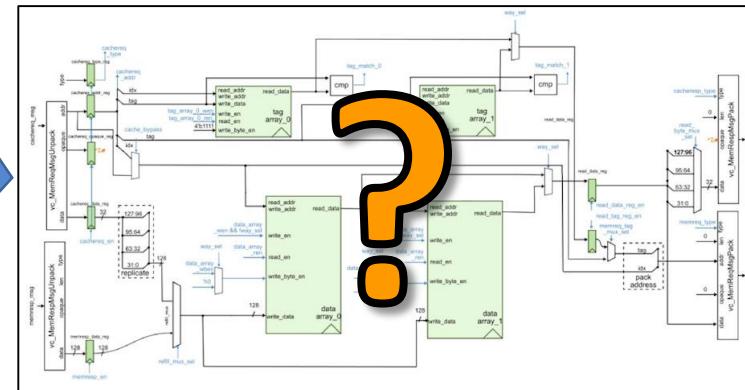
# Memory Consistency Checking for RTL



Microarchitecture Checking

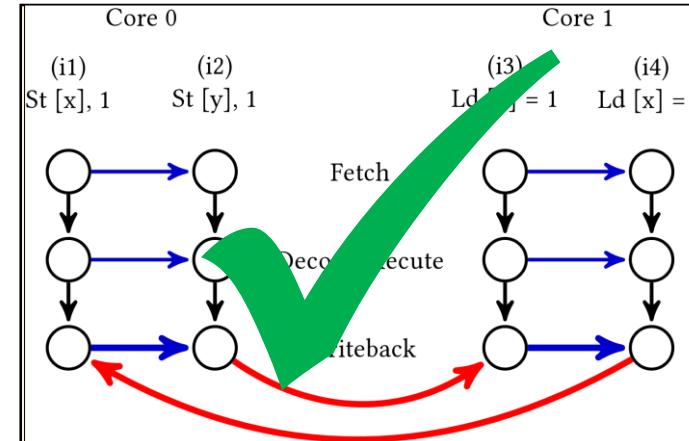
How to ensure RTL maintains orderings?

RTL implementation



[RTL Image: Christopher Batten]

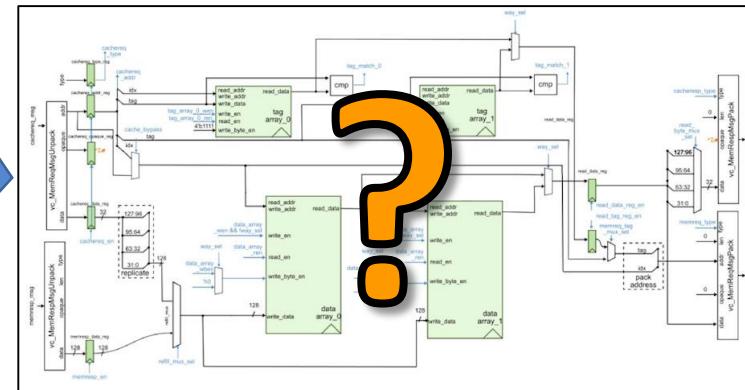
# Memory Consistency Checking for RTL



Microarchitecture Checking

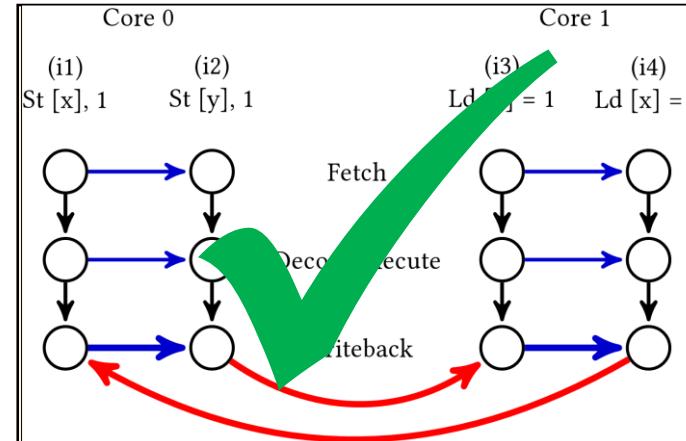
How to ensure RTL maintains orderings?

RTL implementation



[RTL Image: Christopher Batten]

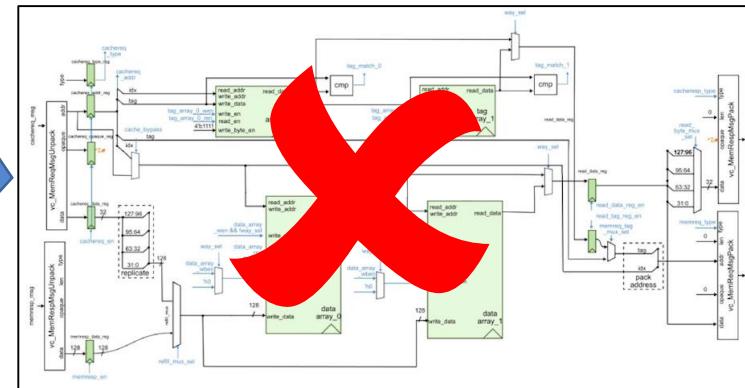
# Memory Consistency Checking for RTL



Microarchitecture Checking

How to ensure RTL maintains orderings?

RTL implementation



[RTL Image: Christopher Batten]

# RTLCheck: Checking RTL Implementations

High-Level Languages (HLL)

Compiler

OS

Architecture (ISA)

Microarchitecture

Processor RTL

- **RTLCheck [Manerkar et al. MICRO 2017]** enables checking microarchitectural axioms against an implementation’s Verilog RTL for litmus test suites
- This helps ensure that the RTL maintains orderings required for consistency
- Selected as an Honorable Mention from the “*Top Picks of Comp. Arch. Conferences*” for 2017



# RTL Verification is Maturing...

- ...but usually ignores memory consistency!
- Often use SystemVerilog Assertions (SVA)



# RTL Verification is Maturing...

- ...but usually ignores memory consistency!
- Often use SystemVerilog Assertions (SVA)

**ISA-Formal [Reid et al. CAV 2016]**

-Instr. Operational Semantics

**No MCM verification**



# RTL Verification is Maturing...

- ...but usually ignores memory consistency!
- Often use SystemVerilog Assertions (SVA)

**ISA-Formal [Reid et al. CAV 2016]**

-Instr. Operational Semantics

**No MCM verification**

**DOGReL [Stewart et al. DIFTS 2014]**

-Memory subsystem transactions

**No multicore MCM verification (?)**



# RTL Verification is Maturing...

- ...but usually ignores memory consistency!
- Often use SystemVerilog Assertions (SVA)

**ISA-Formal [Reid et al. CAV 2016]**

-Instr. Operational Semantics

**No MCM verification**

**DOGReL [Stewart et al. DIFTS 2014]**

-Memory subsystem transactions

**No multicore MCM verification (?)**

**Kami**

**[Vijayaraghavan et al. CAV 2015] [Choi et al. ICFP 2017]**

-MCM correctness for all programs, but...

**Needs Bluespec design and manual proofs!**



# RTL Verification is Maturing...

- ...but usually ignores memory consistency!
- Often use SystemVerilog Assertions (SVA)

**Lack of automated memory  
consistency verification at RTL!**

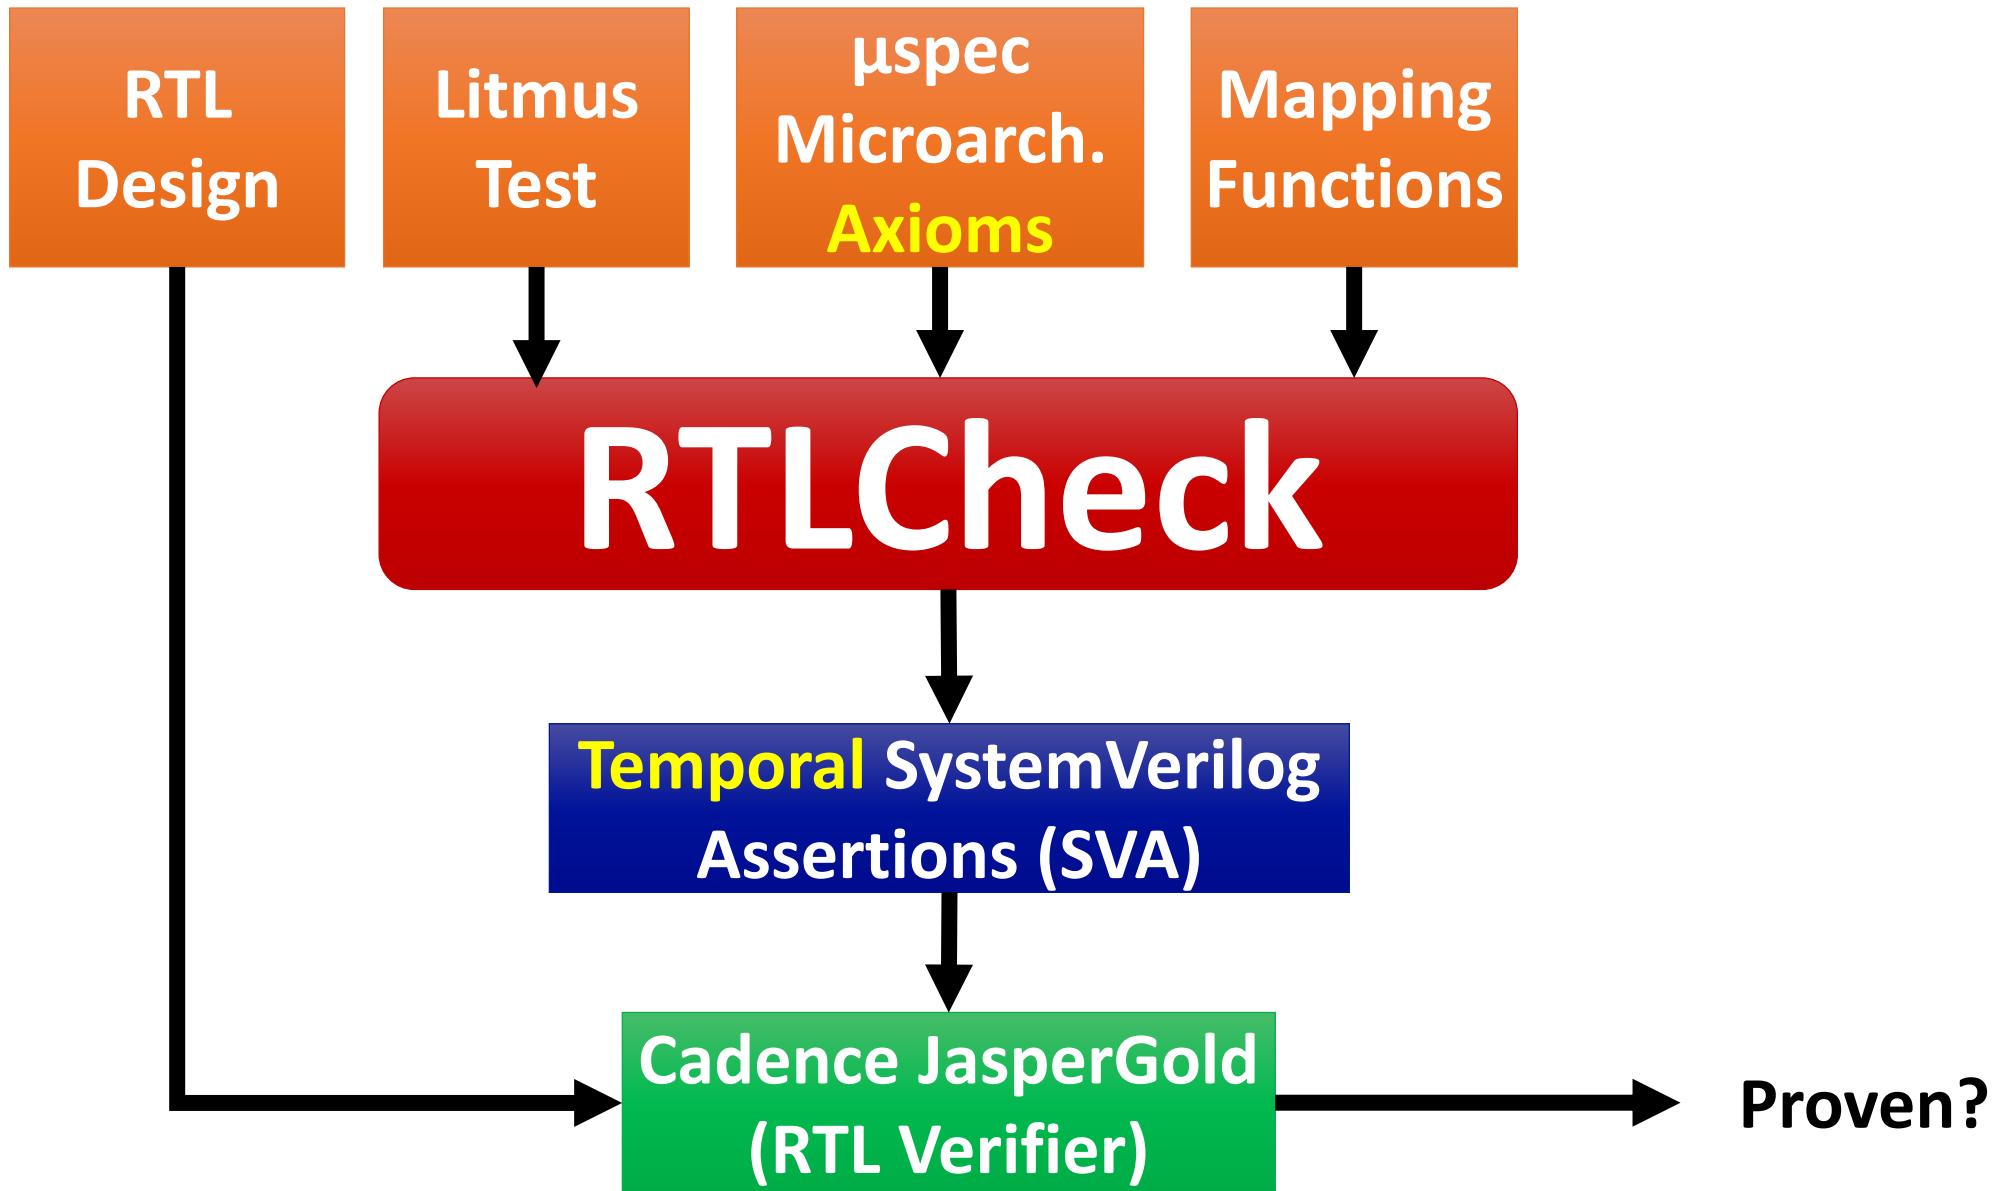
[Vijayaraghavan et al. CAV 2015] [Choi et al. ICFP 2017]

-MCM correctness for all programs, but...

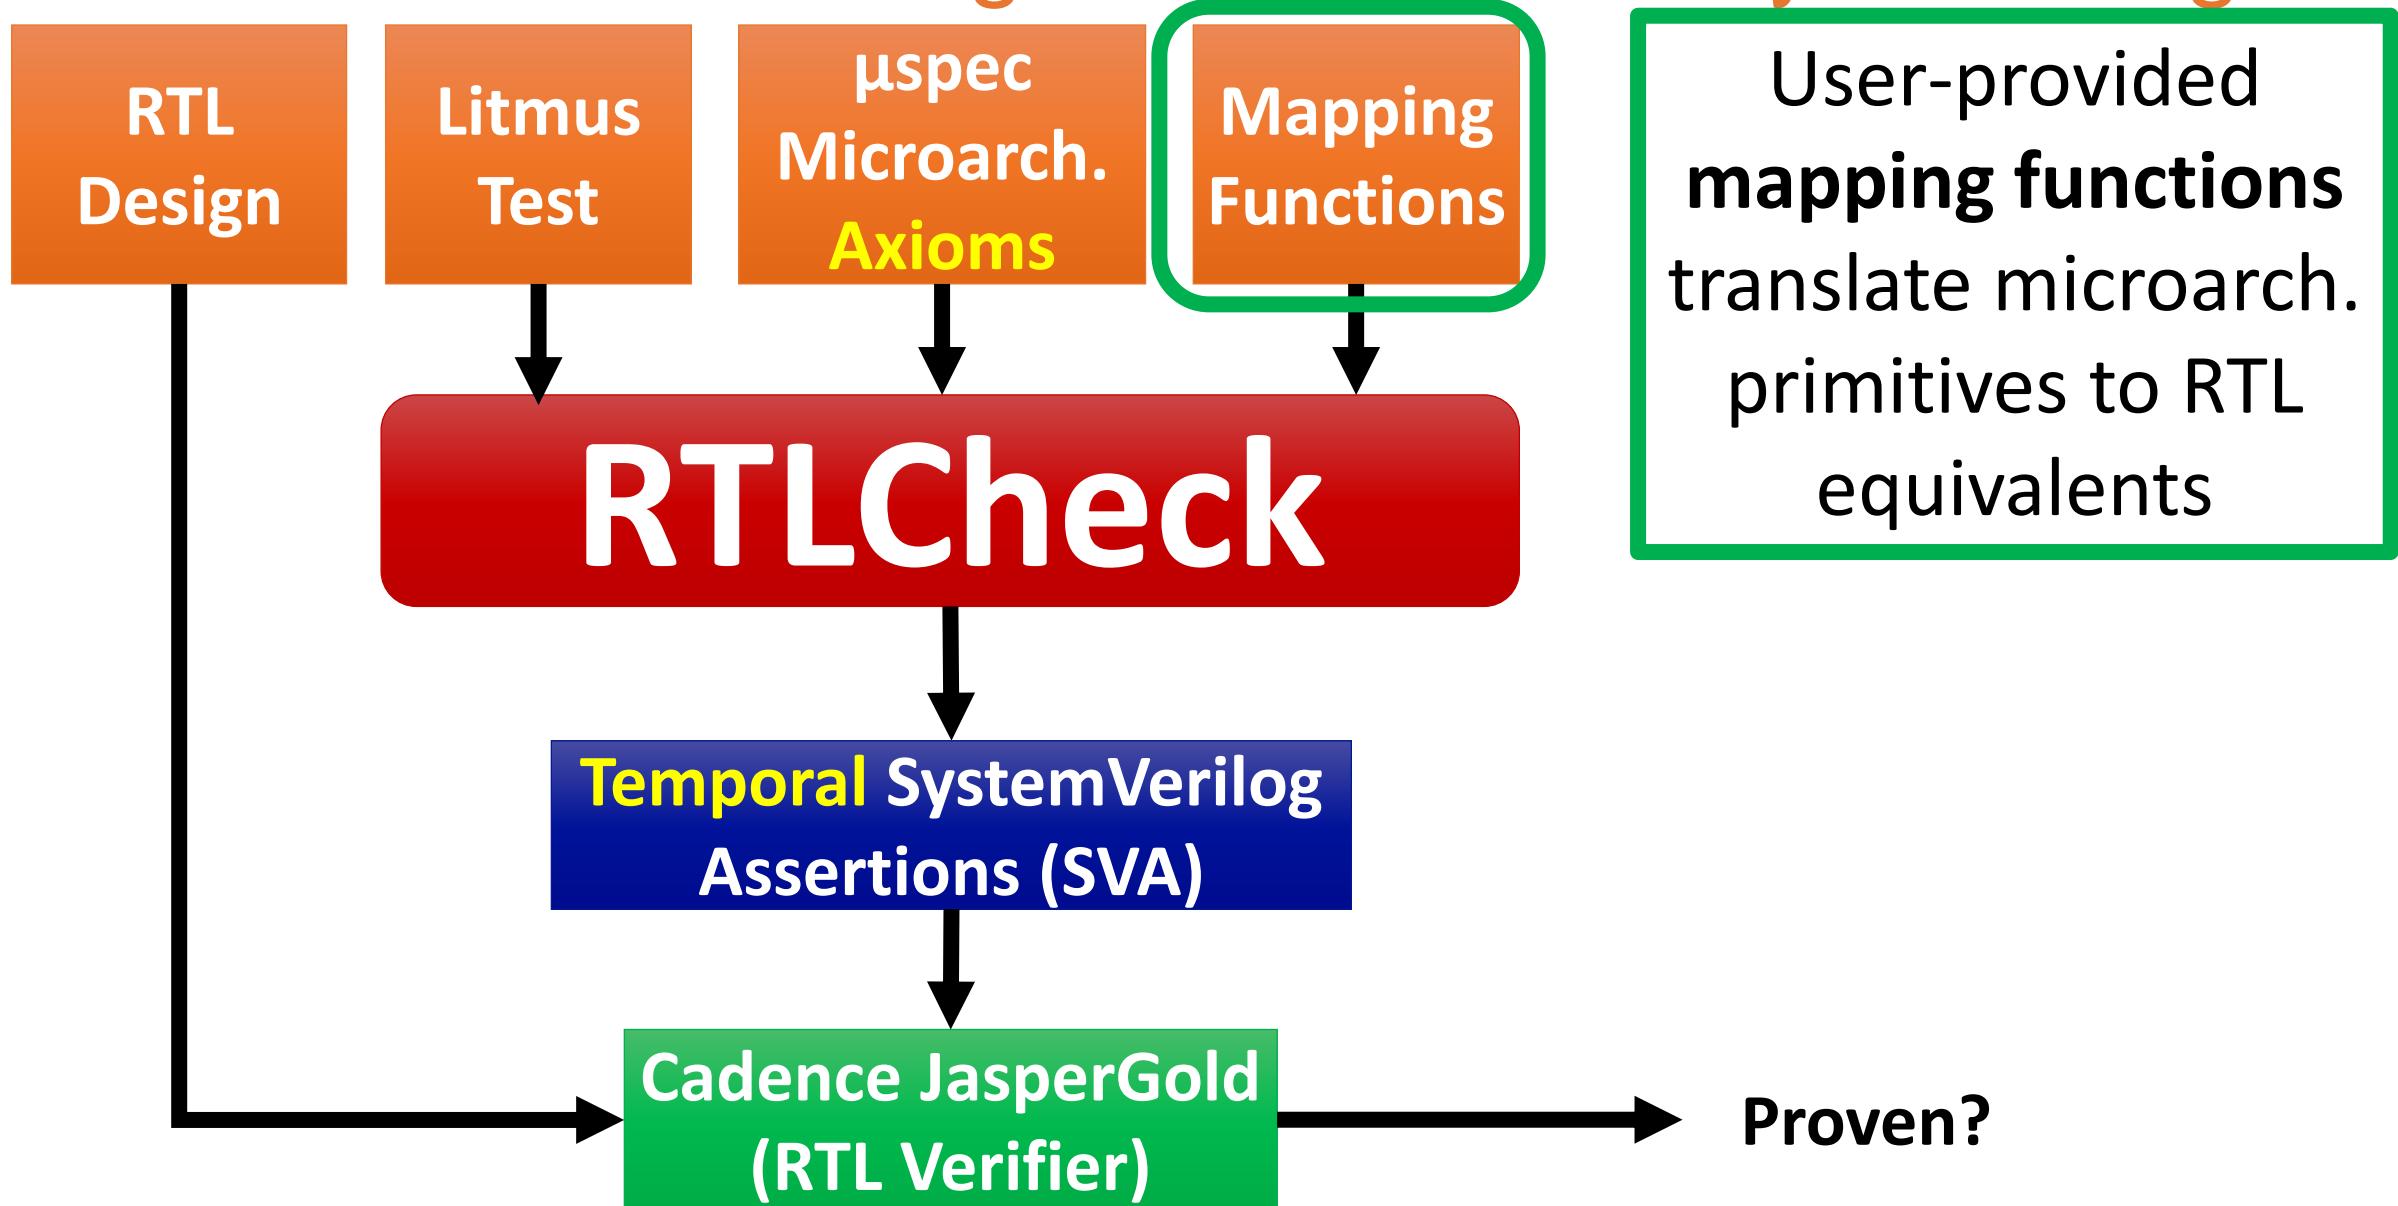
**Needs Bluespec design and manual proofs!**



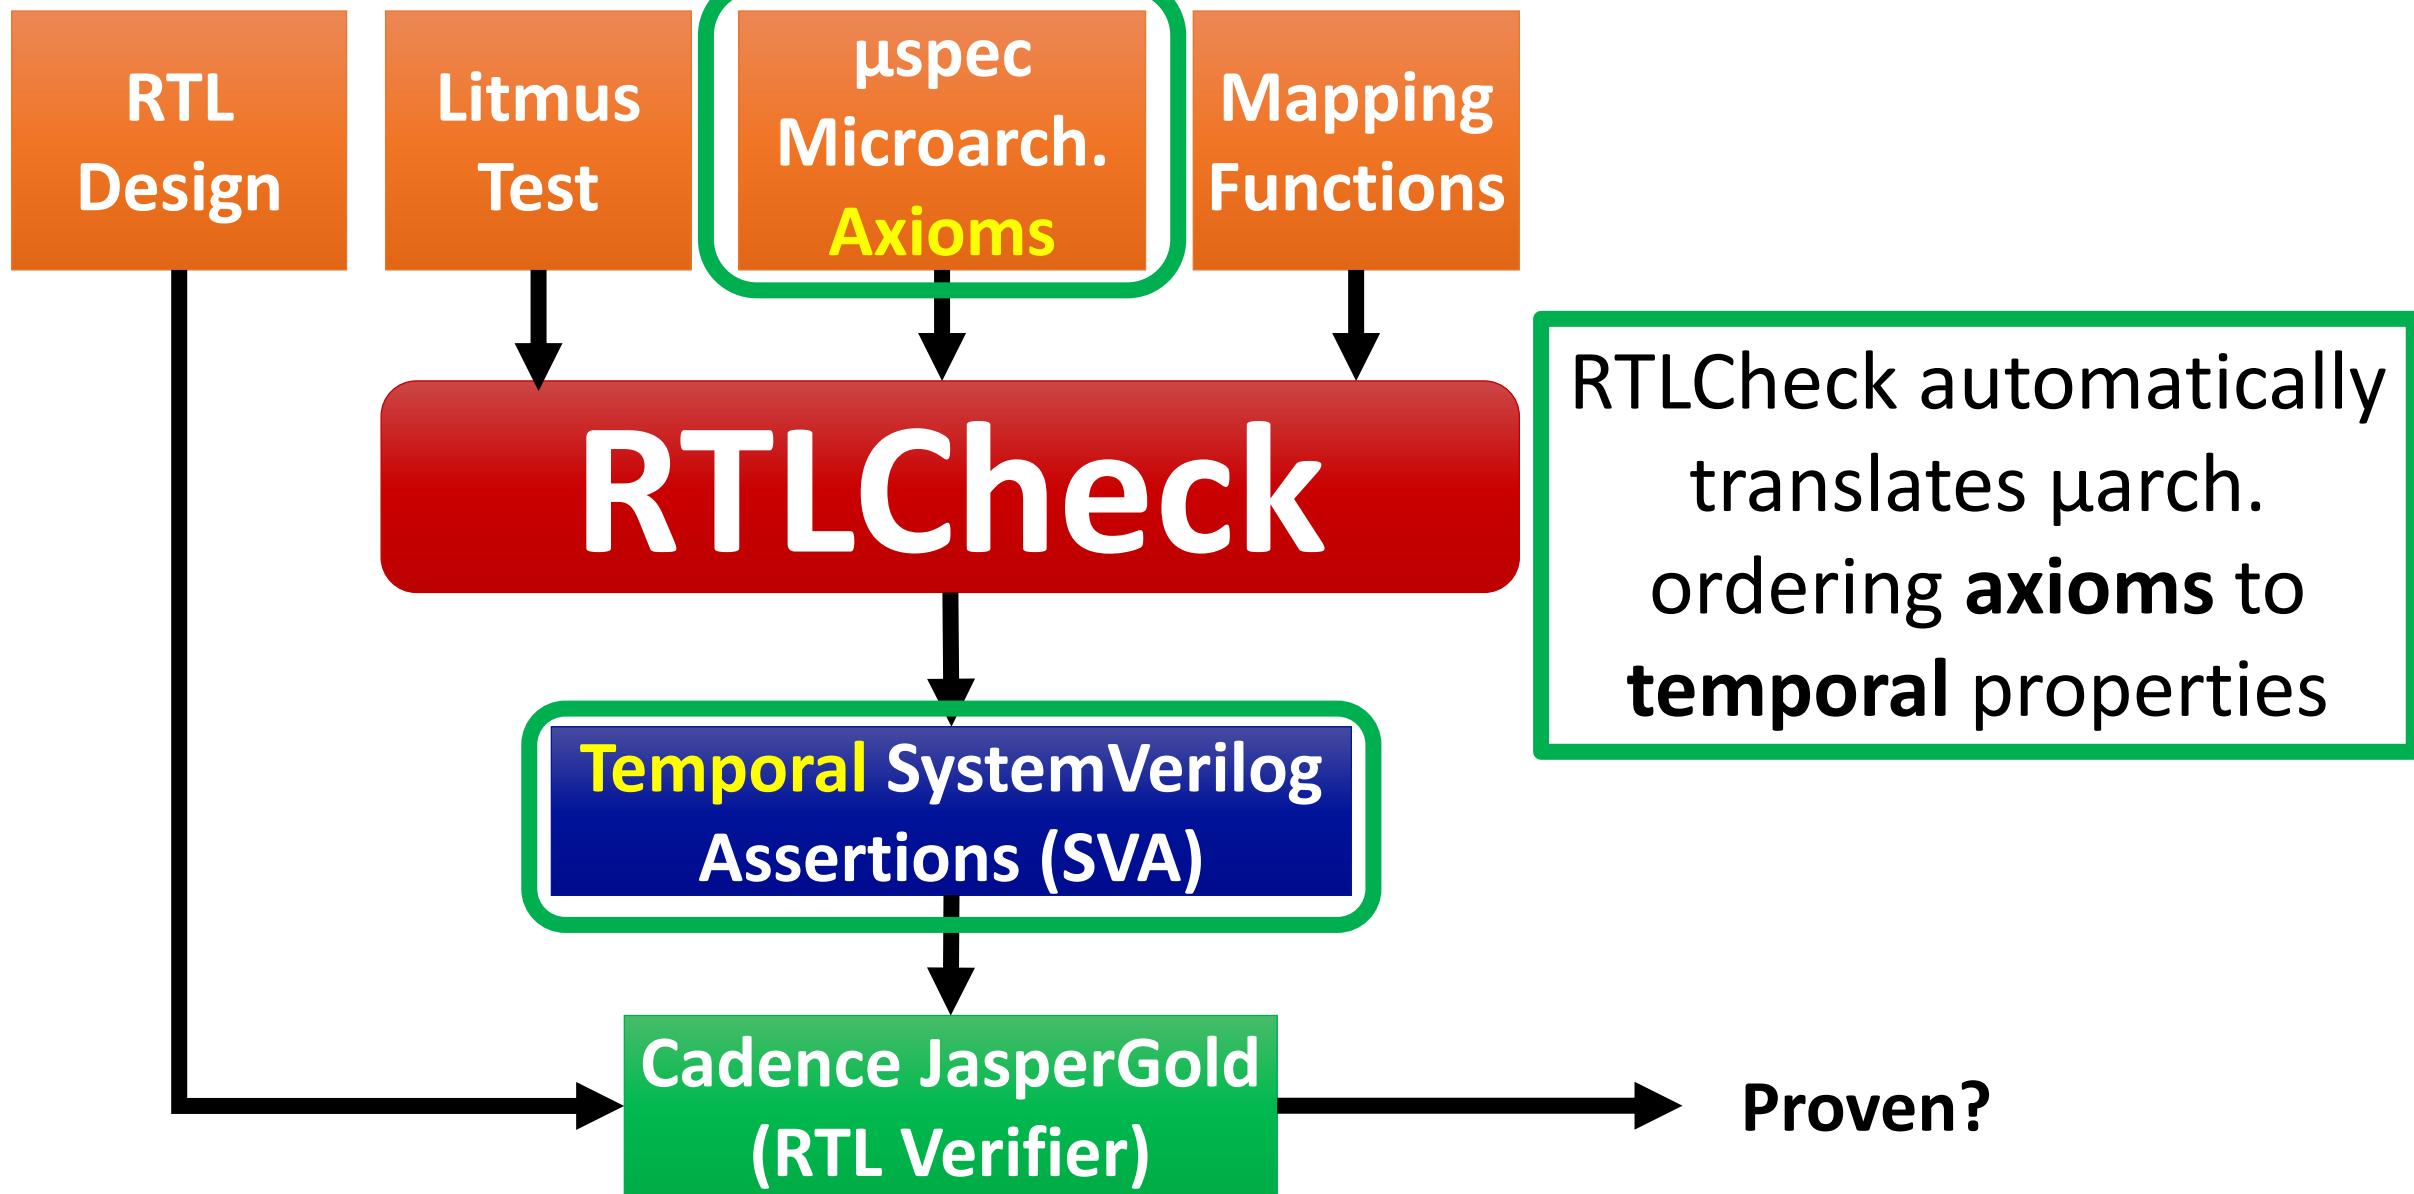
# RTLCheck: Checking RTL Consistency Orderings



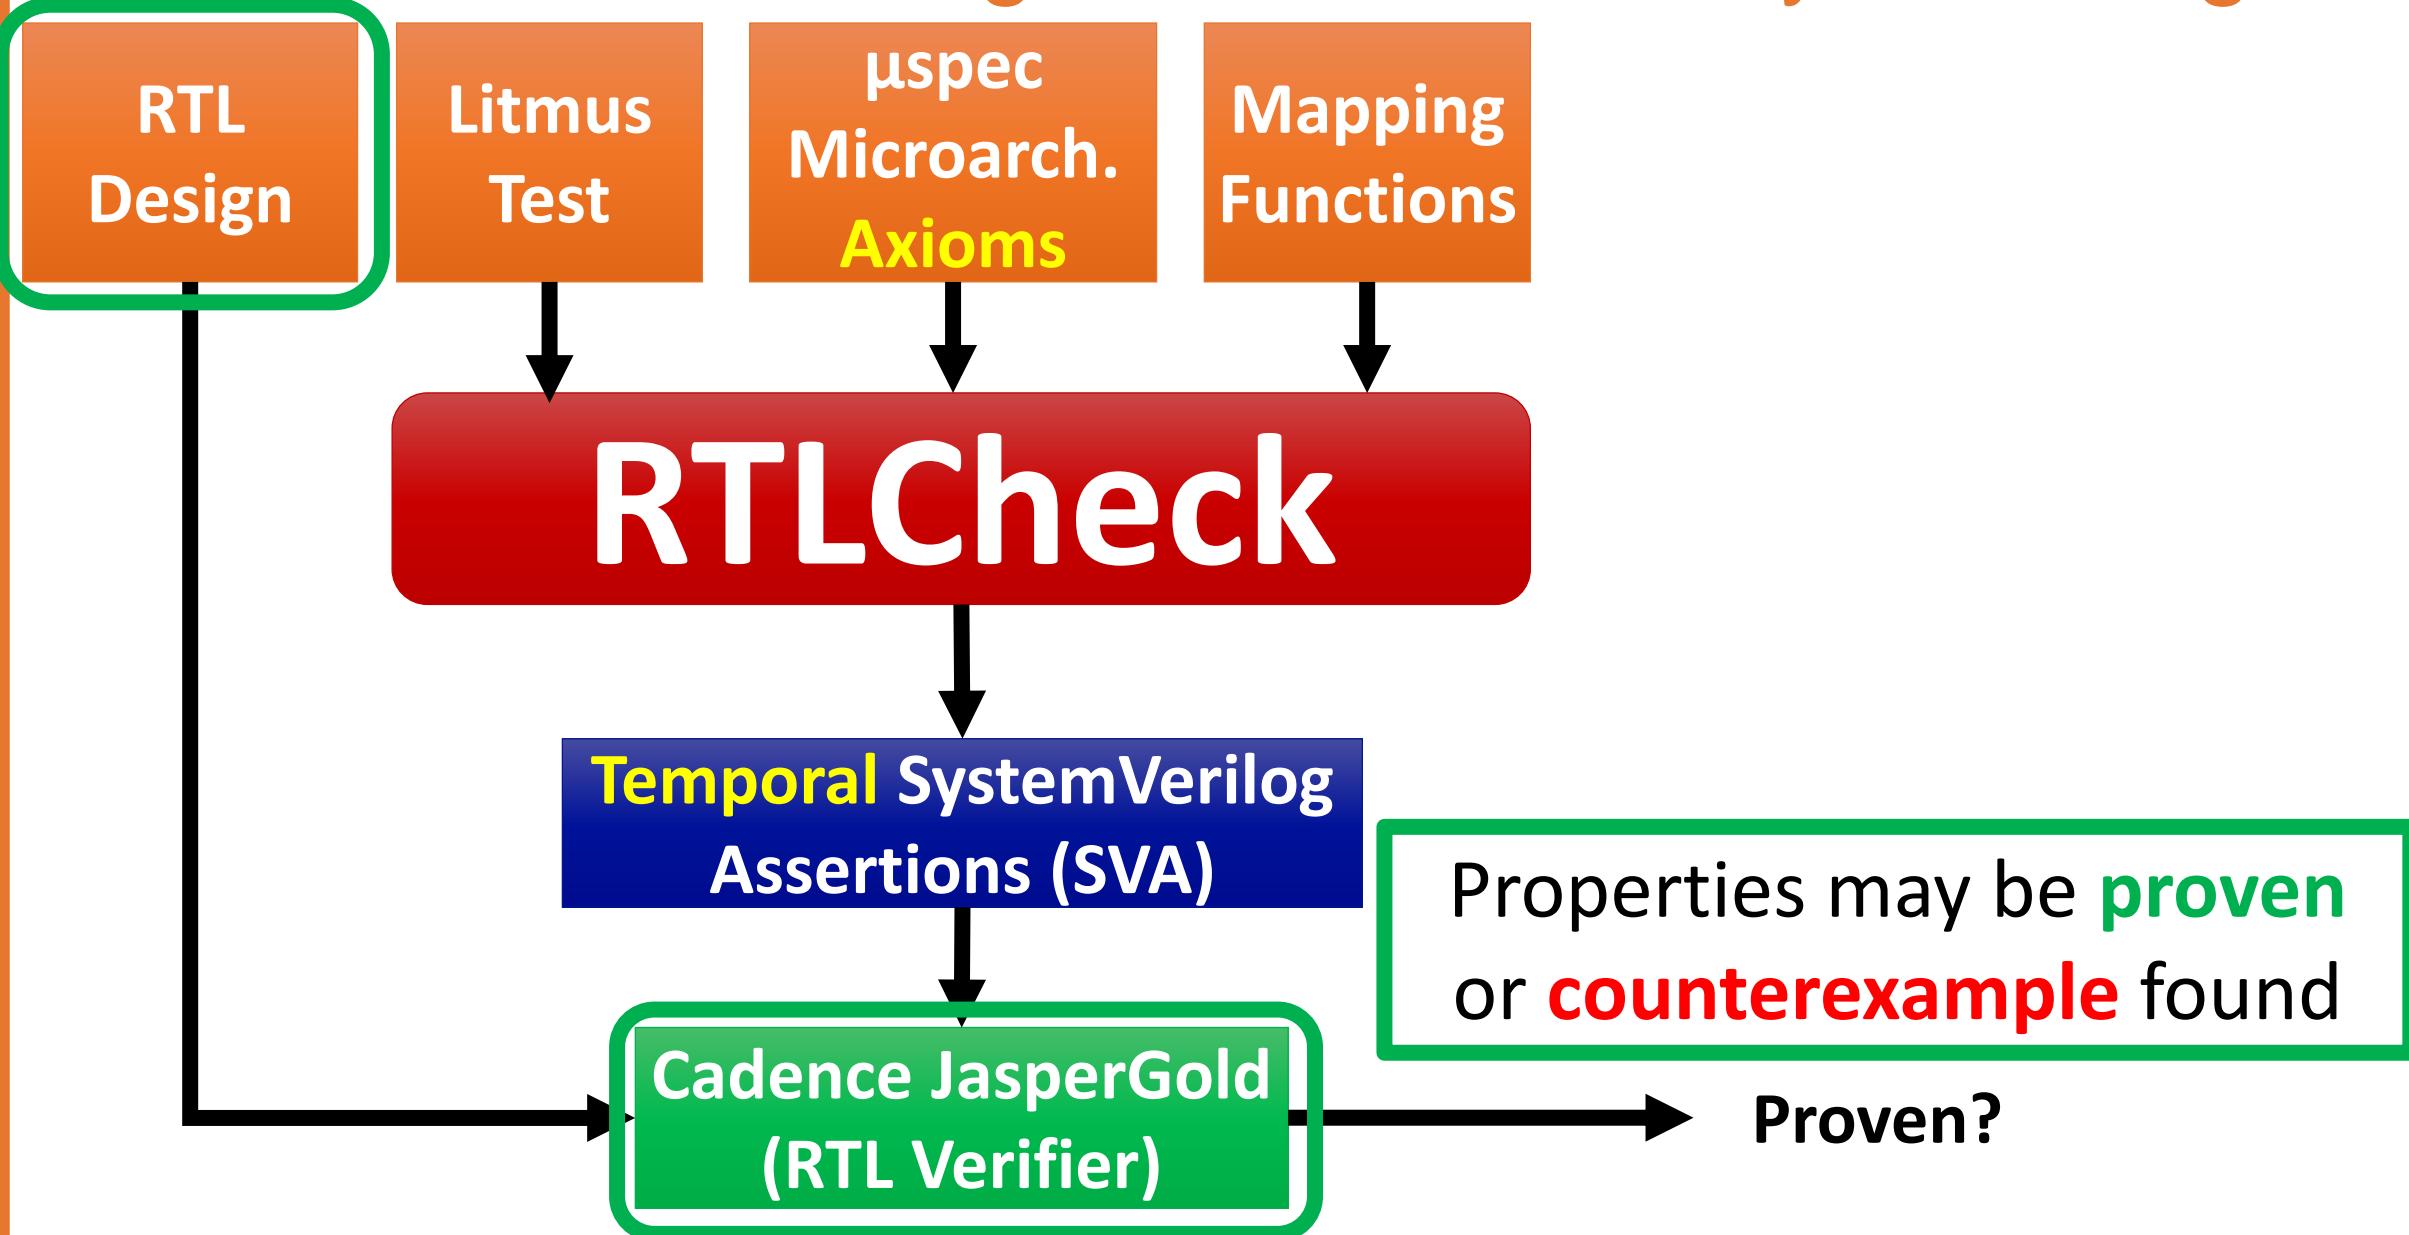
# RTLCheck: Checking RTL Consistency Orderings



# RTLCheck: Checking RTL Consistency Orderings



# RTLCheck: Checking RTL Consistency Orderings



# Meaning can be Lost in Translation!

小心地滑



# Meaning can be Lost in Translation!

小心地滑

(Caution: Slippery Floor)



# Meaning can be Lost in Translation!



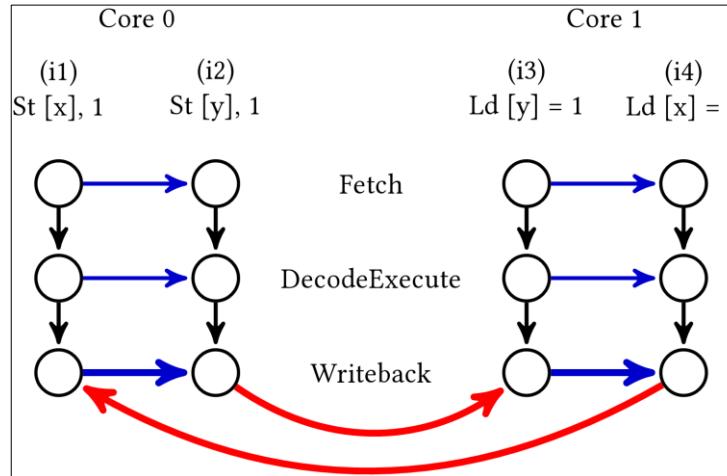
[Image: Barbara Younger]

[Inspiration: Tae Jun Ham]



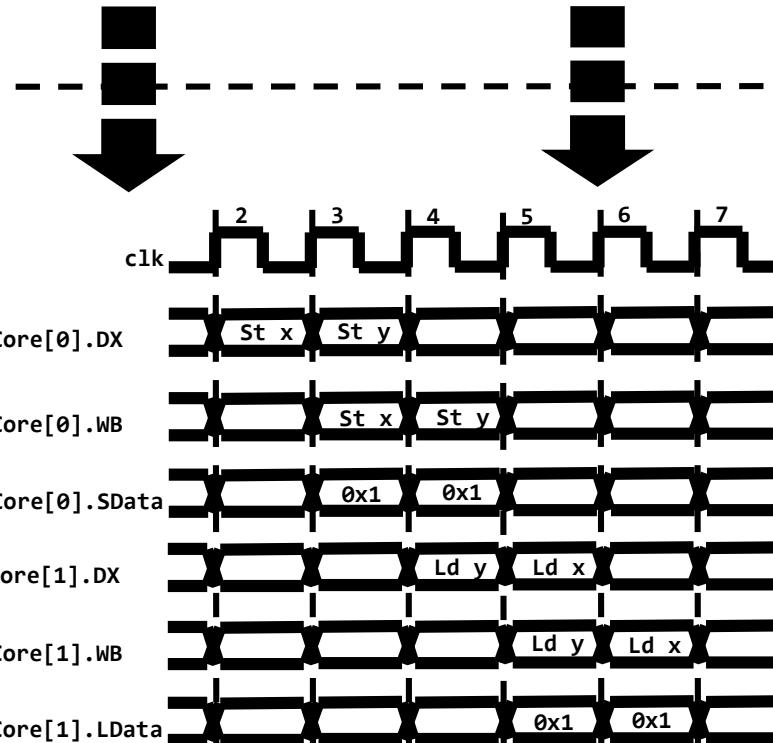
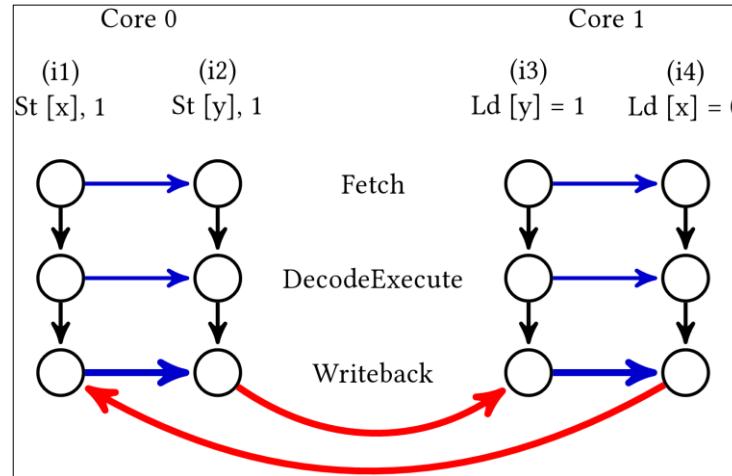
# RTLCheck: Checking Consistency at RTL

**Axiomatic  
Microarch.  
Analysis**



# RTLCheck: Checking Consistency at RTL

Axiomatic  
Microarch.  
Analysis

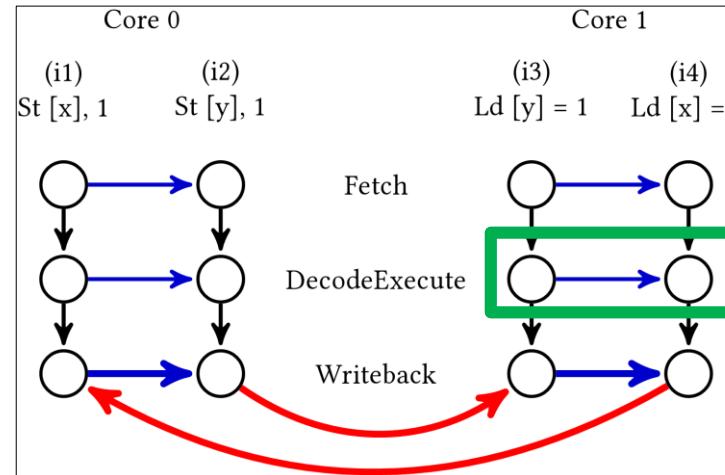


Temporal  
RTL Verification  
(SVA, etc)



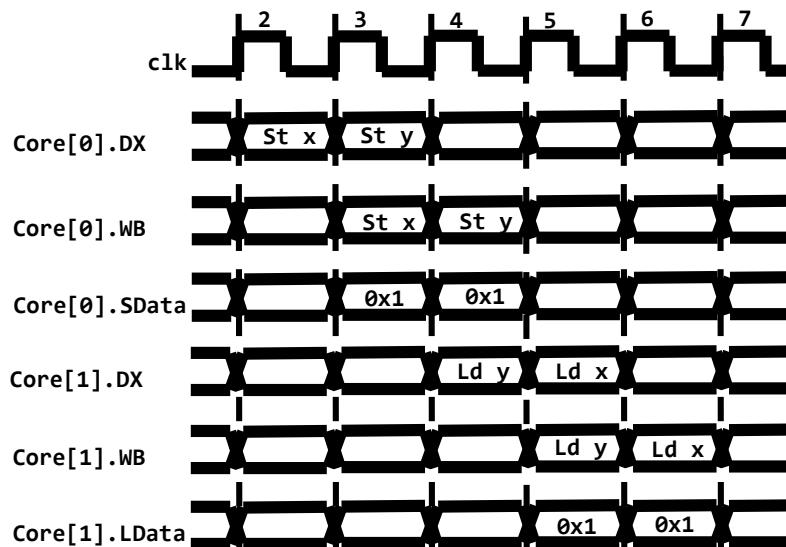
# RTLCheck: Checking Consistency at RTL

Axiomatic  
Microarch.  
Analysis



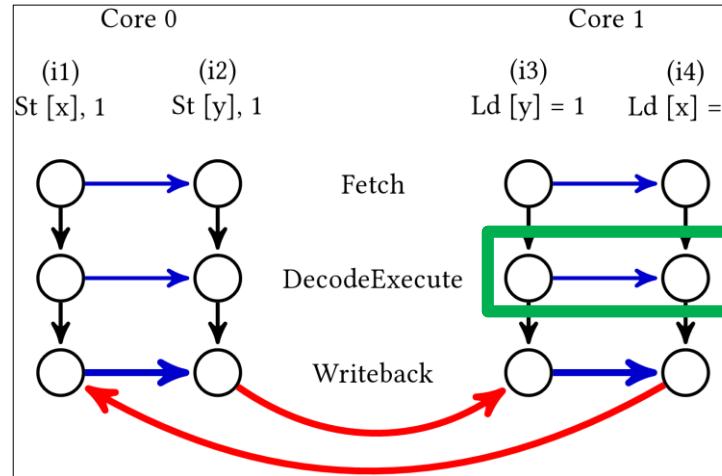
Abstract nodes  
and happens-  
before edges

Temporal  
RTL Verification  
(SVA, etc)



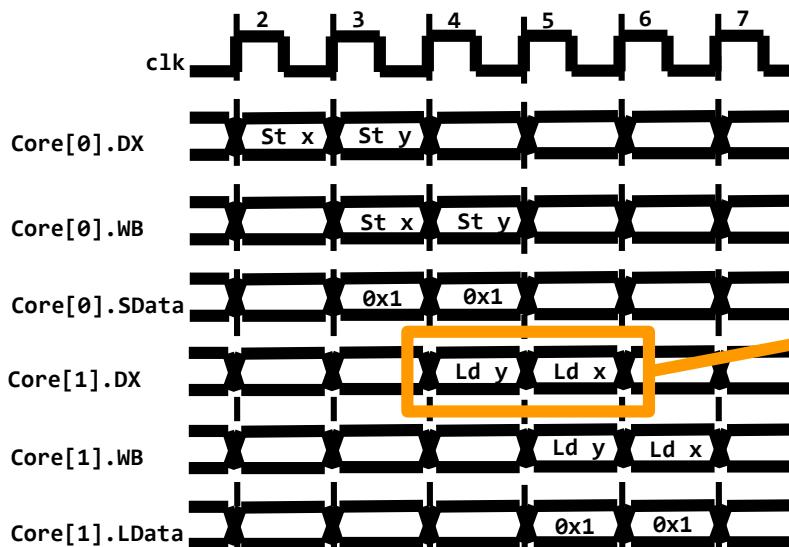
# RTLCheck: Checking Consistency at RTL

Axiomatic  
Microarch.  
Analysis



Abstract nodes  
and happens-  
before edges

Temporal  
RTL Verification  
(SVA, etc)

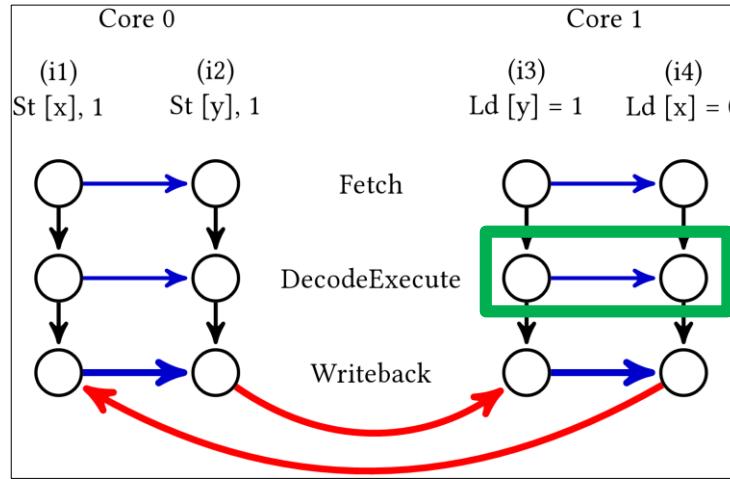


Concrete  
signals and  
clock cycles



# RTLCheck: Checking Consistency at RTL

Axiomatic  
Microarch.  
Analysis



Abstract nodes  
and happens-  
before edges

**Axiomatic/Temporal Mismatch!**

Temporal  
RTL Verification  
(SVA, etc)



Concrete  
signals and  
clock cycles



# Outcome Filtering in Axiomatic Analysis

- **Outcome Filtering:** Restrict test outcome to one particular outcome
  - Allows for more efficient verification
- Axiomatic models make outcome filtering easy

mp (Message Passing)

Core 0	Core 1
(i1) $x = 1;$	(i3) $r1 = y;$
(i2) $y = 1;$	(i4) $r2 = x;$



# Outcome Filtering in Axiomatic Analysis

- **Outcome Filtering:** Restrict test outcome to one particular outcome
  - Allows for more efficient verification
- Axiomatic models make outcome filtering easy

mp (Message Passing)

Core 0	Core 1
(i1) $x = 1;$	(i3) $r1 = y;$
(i2) $y = 1;$	(i4) $r2 = x;$
<b>Outcome: <math>r1 = 1, r2 = 1</math></b>	

Execution examined as a whole,  
so outcome can be enforced!



# Outcome Filtering in Axiomatic Analysis

- **Outcome Filtering:** Restrict test outcome to one particular outcome
  - Allows for more efficient verification
- Axiomatic models make outcome filtering easy

mp (Message Passing)

Core 0	Core 1
(i1) $x = 1;$	(i3) $r1 = y;$
(i2) $y = 1,$	(i4) $r2 = x;$
<b>Outcome: <math>r1 = 1, r2 = 1</math></b>	

Execution examined as a whole,  
so outcome can be enforced!



# Outcome Filtering in Axiomatic Analysis

- **Outcome Filtering:** Restrict test outcome to one particular outcome
  - Allows for more efficient verification
- Axiomatic models make outcome filtering easy

mp (Message Passing)

Core 0	Core 1
(i1) $x = 1,$	(i3) $r1 = y;$
(i2) $y = 1,$	(i4) $r2 = x;$
<b>Outcome: <math>r1 = 1, r2 = 1</math></b>	

Execution examined as a whole,  
so outcome can be enforced!



# Outcome Filtering in Temporal Verification

- Filtering executions by outcome requires expensive global analysis
  - Not done by many SVA verifiers, including JasperGold!

mp

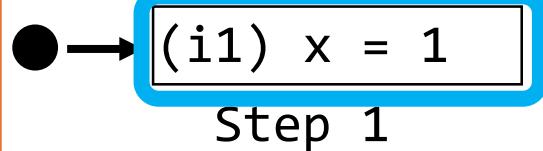
Core 0	Core 1
(i1) $x = 1;$	(i3) $r1 = y;$
(i2) $y = 1;$	(i4) $r2 = x;$
<b>Is <math>r1 = 1, r2 = 0</math> possible?</b>	



# Outcome Filtering in Temporal Verification

- Filtering executions by outcome requires expensive global analysis
  - Not done by many SVA verifiers, including JasperGold!

mp	
Core 0	Core 1
(i1) $x = 1;$	(i3) $r1 = y;$
(i2) $y = 1;$	(i4) $r2 = x;$
<b>Is <math>r1 = 1, r2 = 0</math> possible?</b>	

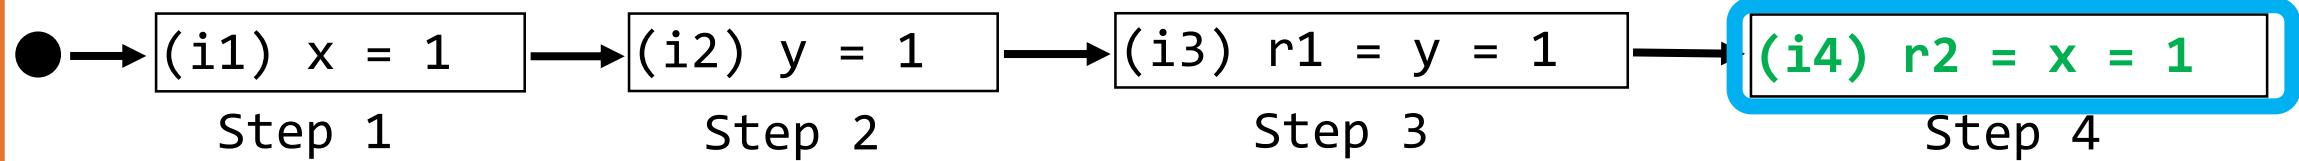


# Outcome Filtering in Temporal Verification

- Filtering executions by outcome requires expensive global analysis
  - Not done by many SVA verifiers, including JasperGold!

mp

Core 0	Core 1
(i1) $x = 1;$	(i3) $r1 = y;$
(i2) $y = 1;$	(i4) $r2 = x;$
Is $r1 = 1, r2 = 0$ possible?	

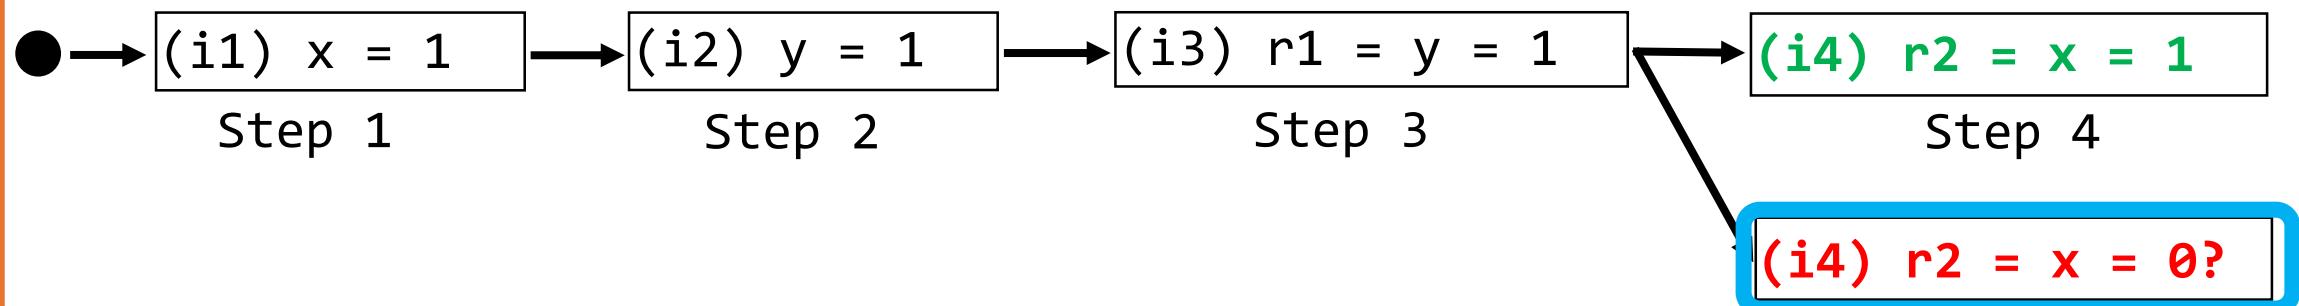


# Outcome Filtering in Temporal Verification

- Filtering executions by outcome requires expensive global analysis
  - Not done by many SVA verifiers, including JasperGold!

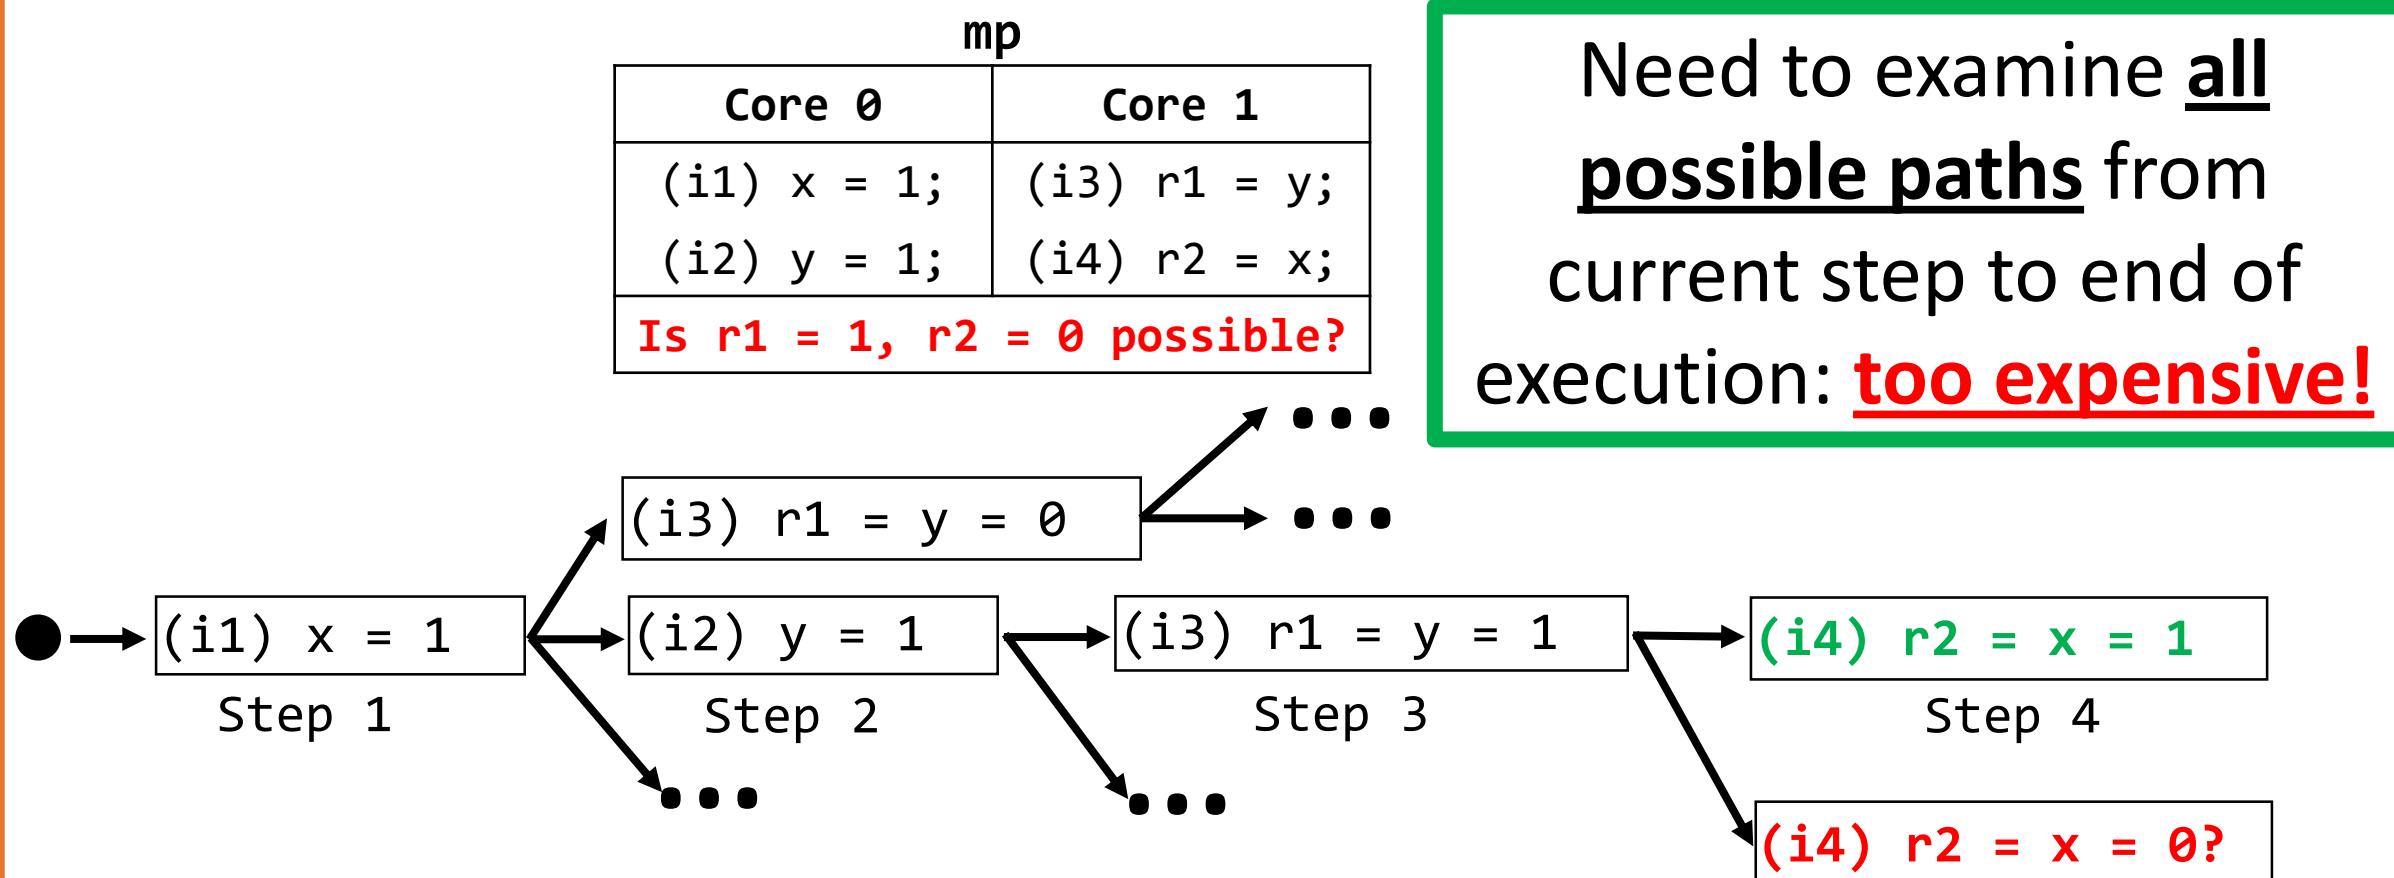
**mp**

Core 0	Core 1
(i1) $x = 1;$	(i3) $r1 = y;$
(i2) $y = 1;$	(i4) $r2 = x;$
Is $r1 = 1, r2 = 0$ possible?	



# Outcome Filtering in Temporal Verification

- Filtering executions by outcome requires expensive global analysis
  - Not done by many SVA verifiers, including JasperGold!

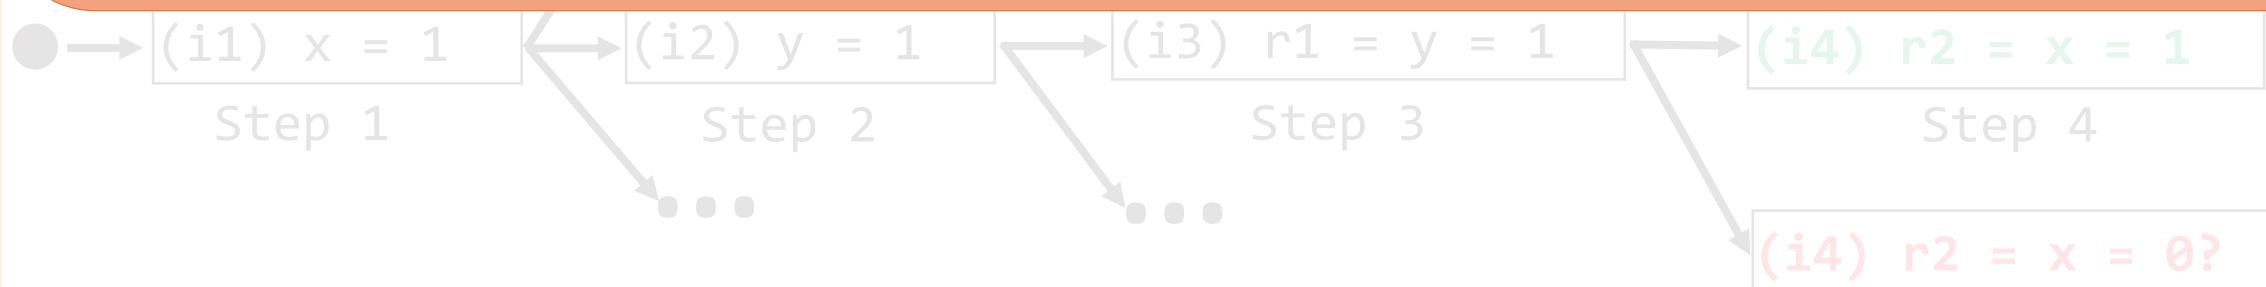


# Outcome Filtering in Temporal Verification

- Filtering executions by outcome requires expensive global analysis
  - Not done by many SVA verifiers, including JasperGold!

**SVA Verifier Approximation:** Only check if constraints hold up to current step

**Makes Outcome Filtering impossible!**



# $\mu$ spec Analysis Uses Outcome Filtering

mp	
Core 0	Core 1
(i1) $x = 1;$	(i3) $r1 = y;$
(i2) $y = 1;$	(i4) $r2 = x;$
<b>SC Forbids: <math>r1 = 1, r2 = 0</math></b>	

Axiom "*Read\_Values*":

Every load either reads **BeforeAllWrites** OR reads **FromLatestWrite**



# $\mu$ spec Analysis Uses Outcome Filtering

mp	
Core 0	Core 1
(i1) $x = 1;$	(i3) $r1 = y;$
(i2) $y = 1;$	(i4) $r2 = x;$
<b>SC Forbids: <math>r1 = 1, r2 = 0</math></b>	

Axiom "*Read\_Values*":

Every load either reads **BeforeAllWrites** OR reads **FromLatestWrite**



# $\mu$ spec Analysis Uses Outcome Filtering

mp	
Core 0	Core 1
(i1) $x = 1;$	(i3) $r1 = y;$
(i2) $y = 1;$	(i4) $r2 = x;$
<b>SC Forbids: <math>r1 = 1, r2 = 0</math></b>	

Axiom "*Read\_Values*":

Every load either reads **BeforeAllWrites** OR reads **FromLatestWrite**

**No write for load  
to read from!**



# $\mu$ spec Analysis Uses Outcome Filtering

mp	
Core 0	Core 1
(i1) $x = 1;$	(i3) $r1 = y;$
(i2) $y = 1;$	(i4) $r2 = x;$
<b>SC Forbids: <math>r1 = 1, r2 = 0</math></b>	

Axiom "*Read\_Values*":

~~Every load either reads BeforeAllWrites OR reads FromLatestWrite~~

**Outcome Filtering leads to simpler axioms!**



# Temporal Outcome Filtering Fails!

**Filtered Read\_Values:**

Unless Load returns non-zero value,

Load happens before all stores to its address

mp	
Core 0	Core 1
(i1) $x = 1;$	(i3) $r1 = y;$
(i2) $y = 1;$	(i4) $r2 = x;$
<b>SC Forbids: <math>r1 = 1, r2 = 0</math></b>	

Time (cycles) →

clk

Core[0].Commit

Core[0].SData

Core[1].Commit

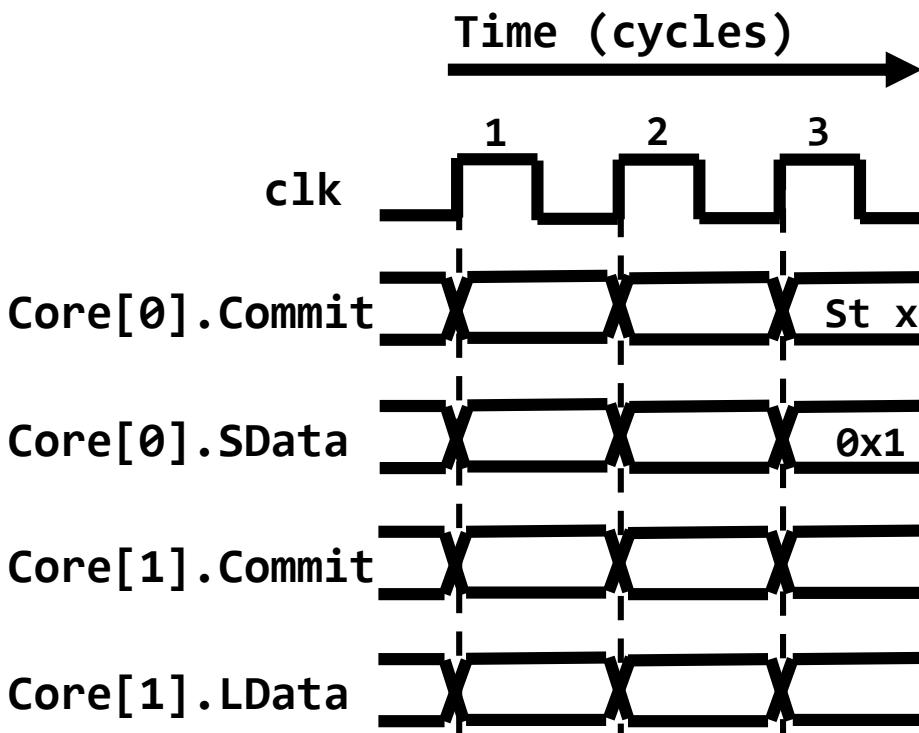
Core[1].LData

# Temporal Outcome Filtering Fails!

**Filtered Read\_Values:**

Unless Load returns non-zero value,

Load happens before all stores to its address



mp	
Core 0	Core 1
(i1) x = 1;	(i3) r1 = y;
(i2) y = 1;	(i4) r2 = x;
<b>SC Forbids: r1 = 1, r2 = 0</b>	

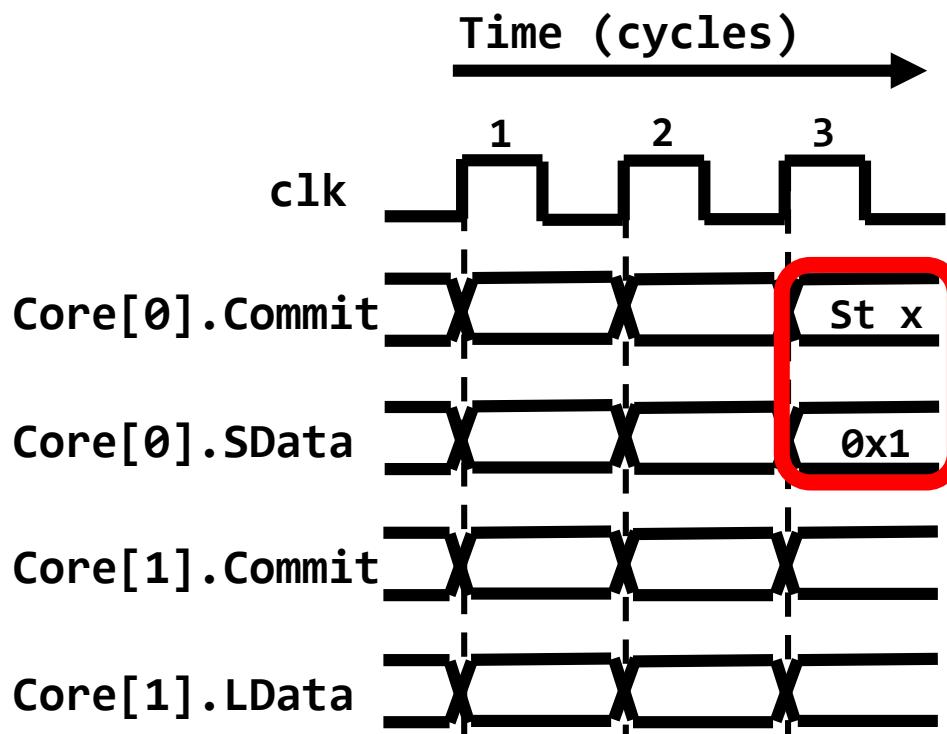
**After 3 cycles:**

# Temporal Outcome Filtering Fails!

**Filtered Read\_Values:**

Unless Load returns non-zero value,

Load happens before all stores to its address



mp	
Core 0	Core 1
(i1) x = 1;	(i3) r1 = y;
(i2) y = 1;	(i4) r2 = x;
<b>SC Forbids: r1 = 1, r2 = 0</b>	

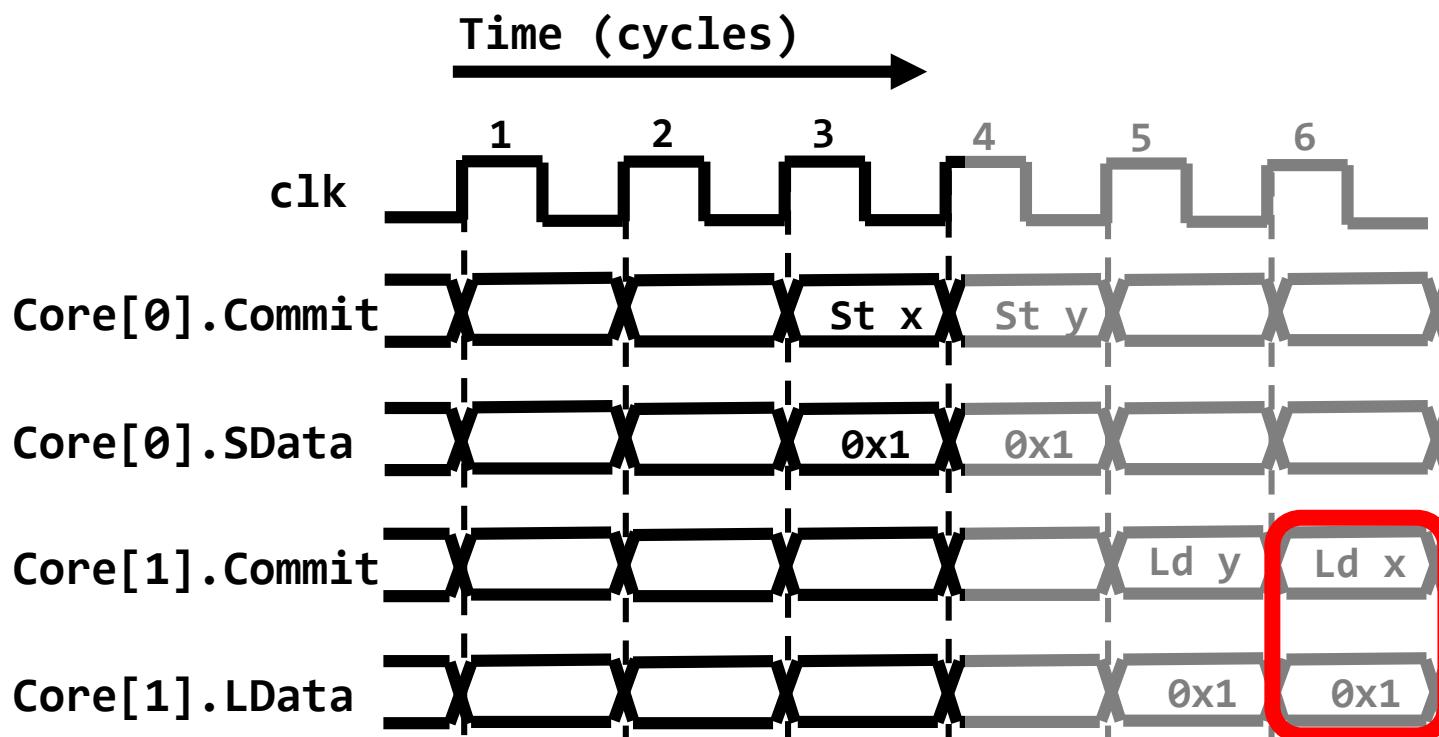
**After 3 cycles:**  
Store happens before load!  
**Property Violated?**

# Temporal Outcome Filtering Fails!

**Filtered Read\_Values:**

Unless Load returns non-zero value,

Load happens before all stores to its address



mp	
Core 0	Core 1
(i1) x = 1;	(i3) r1 = y;
(i2) y = 1;	(i4) r2 = x;
<b>SC Forbids: r1 = 1, r2 = 0</b>	

**After 3 cycles:**  
Store happens before load!  
**Property Violated?**

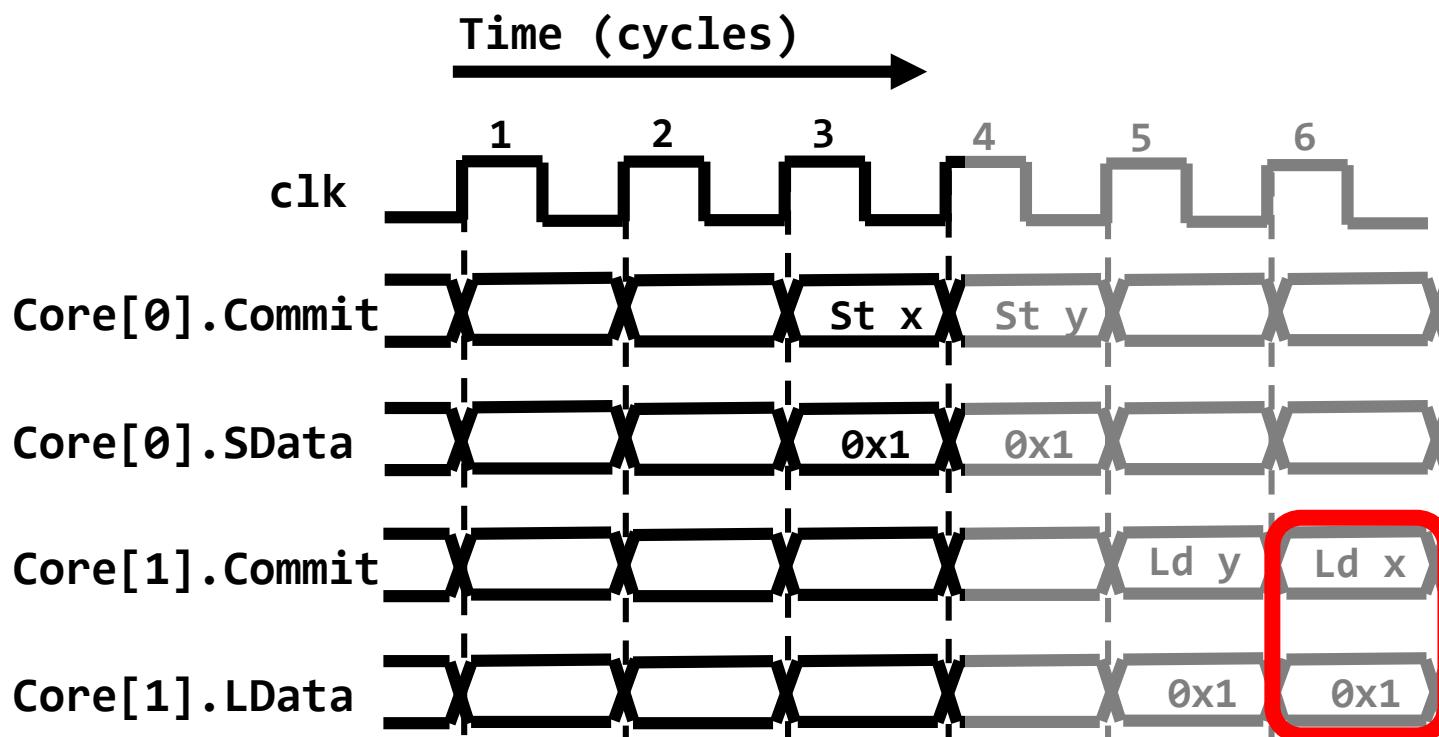
**After 6 cycles:**  
Load does not read 0  
**No Violation!**

# Temporal Outcome Filtering Fails!

**Filtered Read\_Values:**

Unless Load returns non-zero value,

Load happens before all stores to its address



mp	
Core 0	Core 1
(i1) $x = 1;$	(i3) $r1 = y;$
(i2) $y = 1;$	(i4) $r2 = x;$
<b>SC Forbids: <math>r1 = 1, r2 = 0</math></b>	

**After 3 cycles:**

Store happens before load!

**Property Violated?**

**After 6 cycles:**

Load does not read 0

**No Violation!**

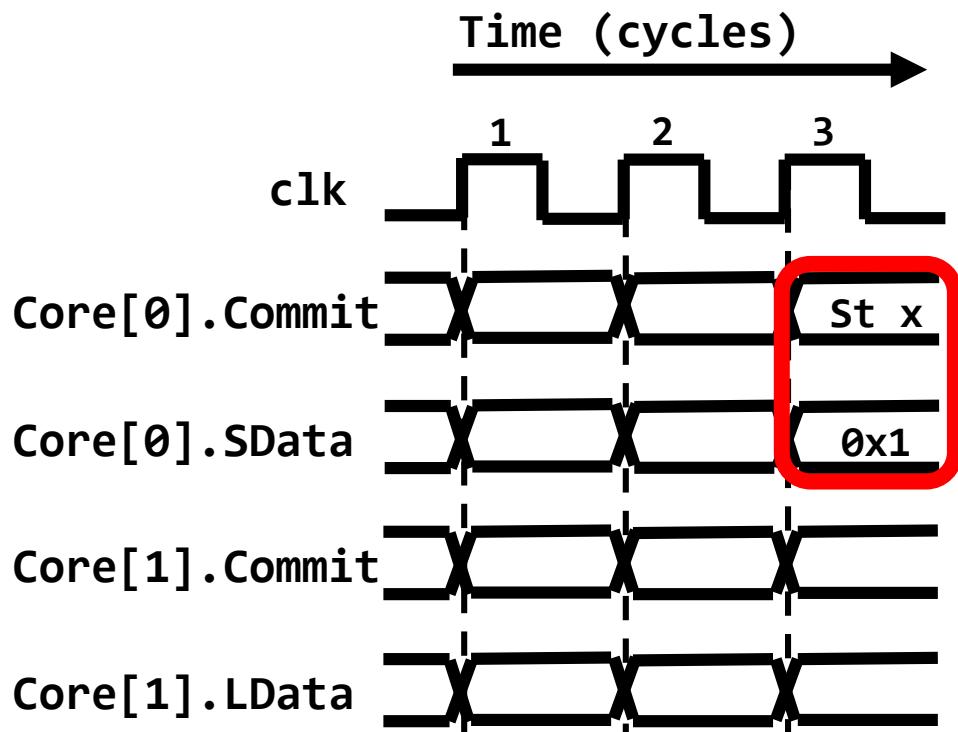
**But SVA verifiers don't check future cycles!**

# Temporal Outcome Filtering Fails!

**Filtered Read\_Values:**

Unless Load returns non-zero value,

Load happens before all stores to its address



Counterexample flagged despite hardware doing **nothing wrong!**

mp	
Core 0	Core 1
(i1) $x = 1;$	(i3) $r1 = y;$
(i2) $y = 1;$	(i4) $r2 = x;$
<b>SC Forbids: <math>r1 = 1, r2 = 0</math></b>	

**After 3 cycles:**

Store happens before load!

**Property Violated?**

**After 6 cycles:**

Load does not read 0

**No Violation!**

**But SVA verifiers don't check future cycles!**

# Solution: Load Value Constraints

- Don't simplify axioms; translate all cases
- Tag each case with appropriate ***load value constraints***
  - reflect the data constraints required for edge(s)

Core 0	Core 1
(i1) $x = 1;$	(i3) $r1 = y;$
(i2) $y = 1;$	(i4) $r2 = x;$
<b>SC Forbids: <math>r1 = 1, r2 = 0</math></b>	

Axiom "*Read\_Values*":

Every load either reads **BeforeAllWrites** OR reads **FromLatestWrite**

Property to check:

`mapNode(Ld x → St x, Ld x == 0) or mapNode(St x → Ld x, Ld x == 1);`



# Solution: Load Value Constraints

- Don't simplify axioms; translate all cases
- Tag each case with appropriate ***load value constraints***
  - reflect the data constraints required for edge(s)

Core 0	Core 1
(i1) $x = 1;$	(i3) $r1 = y;$
(i2) $y = 1;$	(i4) $r2 = x;$
<b>SC Forbids: <math>r1 = 1, r2 = 0</math></b>	

Axiom "*Read\_Values*":

Every load either reads **BeforeAllWrites** OR reads **FromLatestWrite**

Property to check:

`mapNode(Ld x → St x, Ld x == 0)` or `mapNode(St x → Ld x, Ld x == 1);`



# Solution: Load Value Constraints

- Don't simplify axioms; translate all cases
- Tag each case with appropriate ***load value constraints***
  - reflect the data constraints required for edge(s)

Core 0	Core 1
(i1) $x = 1;$	(i3) $r1 = y;$
(i2) $y = 1;$	(i4) $r2 = x;$
<b>SC Forbids: <math>r1 = 1, r2 = 0</math></b>	

Axiom "*Read\_Values*":

Every load either reads **BeforeAllWrites** OR reads **FromLatestWrite**

Property to check:

`mapNode(Ld x → St x, Ld x == 0) or mapNode(St x → Ld x, Ld x == 1);`



# Solution: Load Value Constraints

- Don't simplify axioms; translate all cases
- Tag each case with appropriate ***load value constraints***
  - reflect the data constraints required for edge(s)

Core 0	Core 1
(i1) $x = 1;$	(i3) $r1 = y;$
(i2) $y = 1;$	(i4) $r2 = x;$
<b>SC Forbids: <math>r1 = 1, r2 = 0</math></b>	

Axiom "*Read\_Values*":

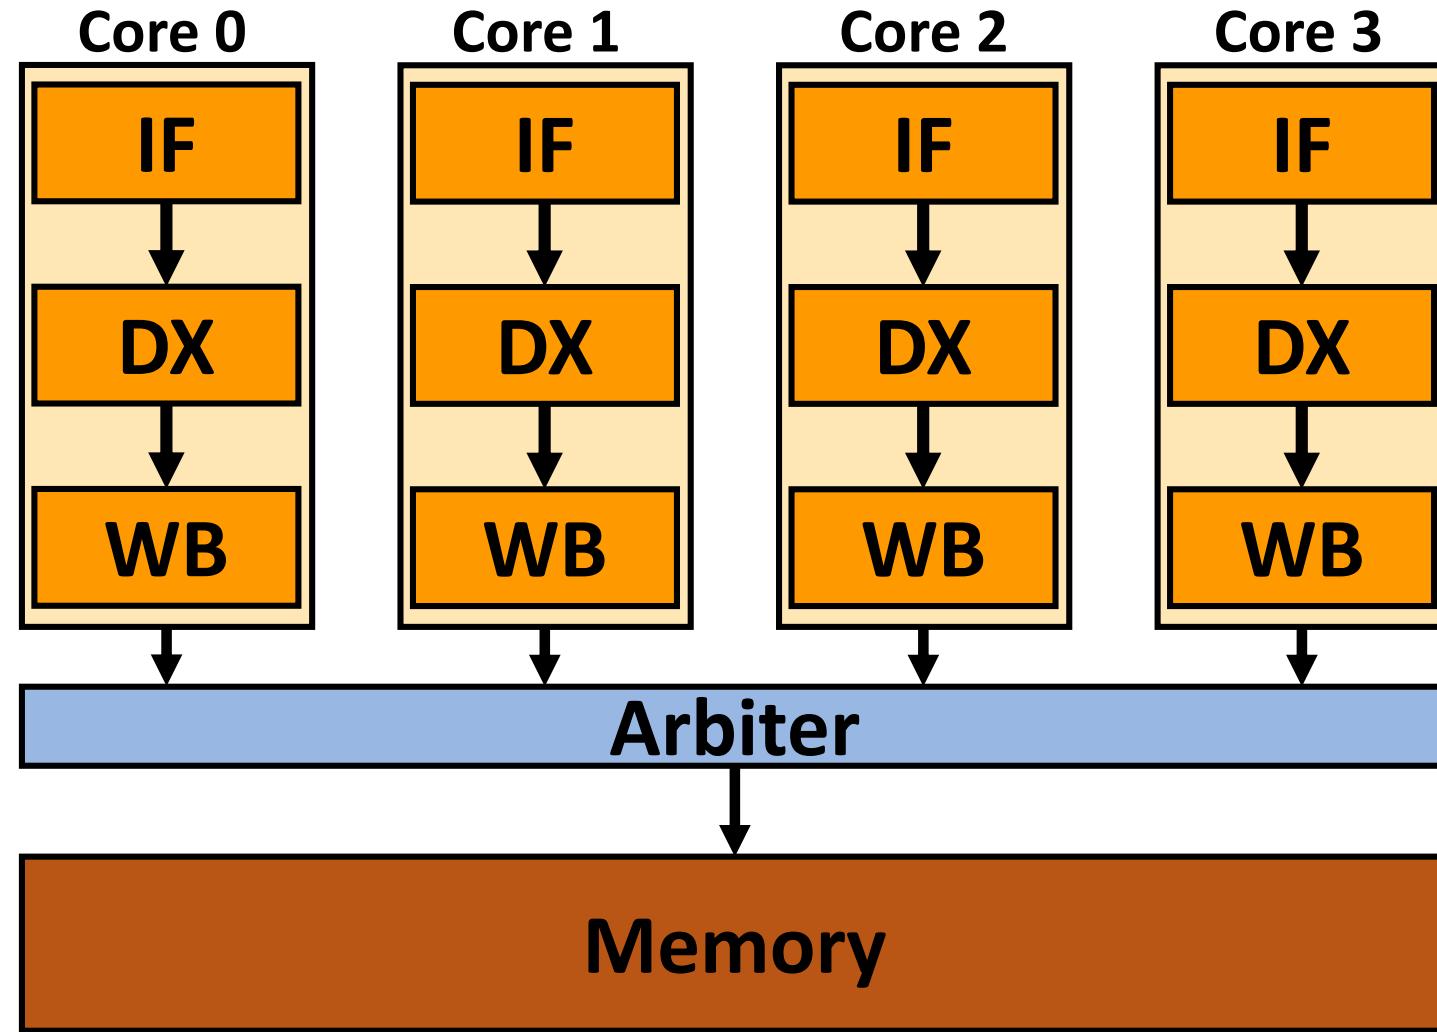
Every load either reads **BeforeAllWrites** **OR** reads **FromLatestWrite**

Property to check:

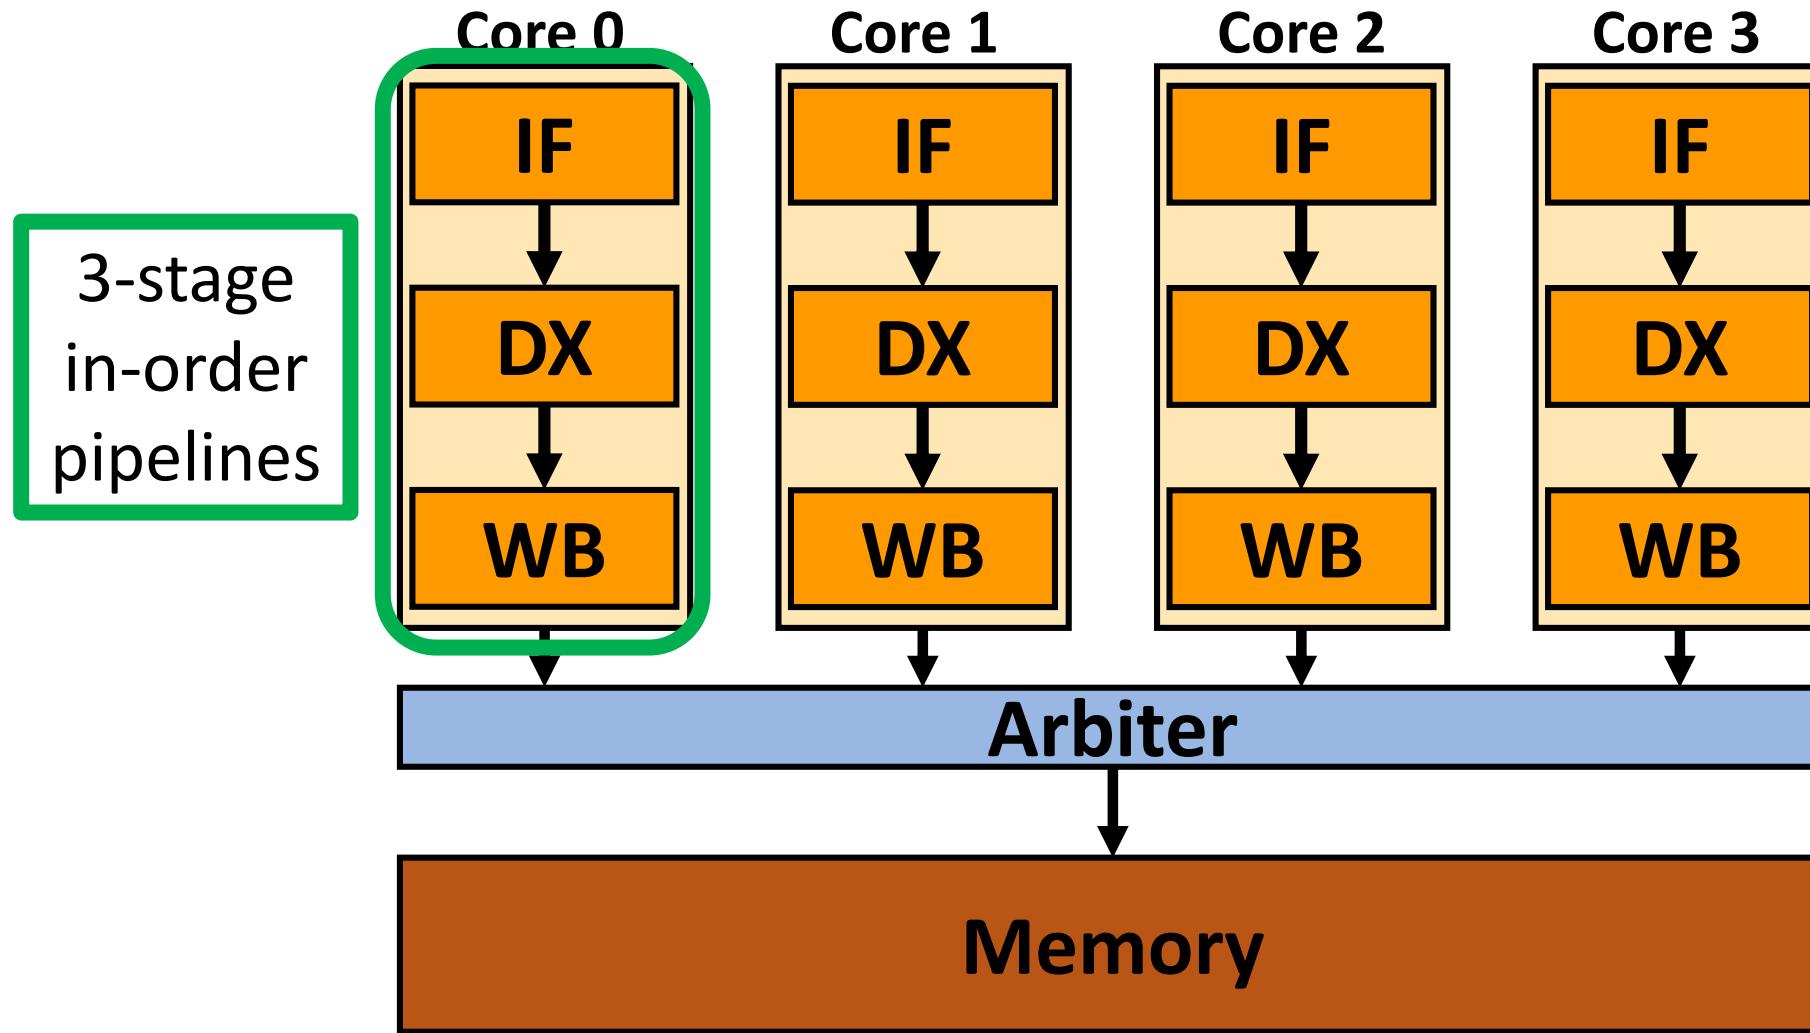
`mapNode(Ld x → St x, Ld x == 0) or mapNode(St x → Ld x, Ld x == 1);`



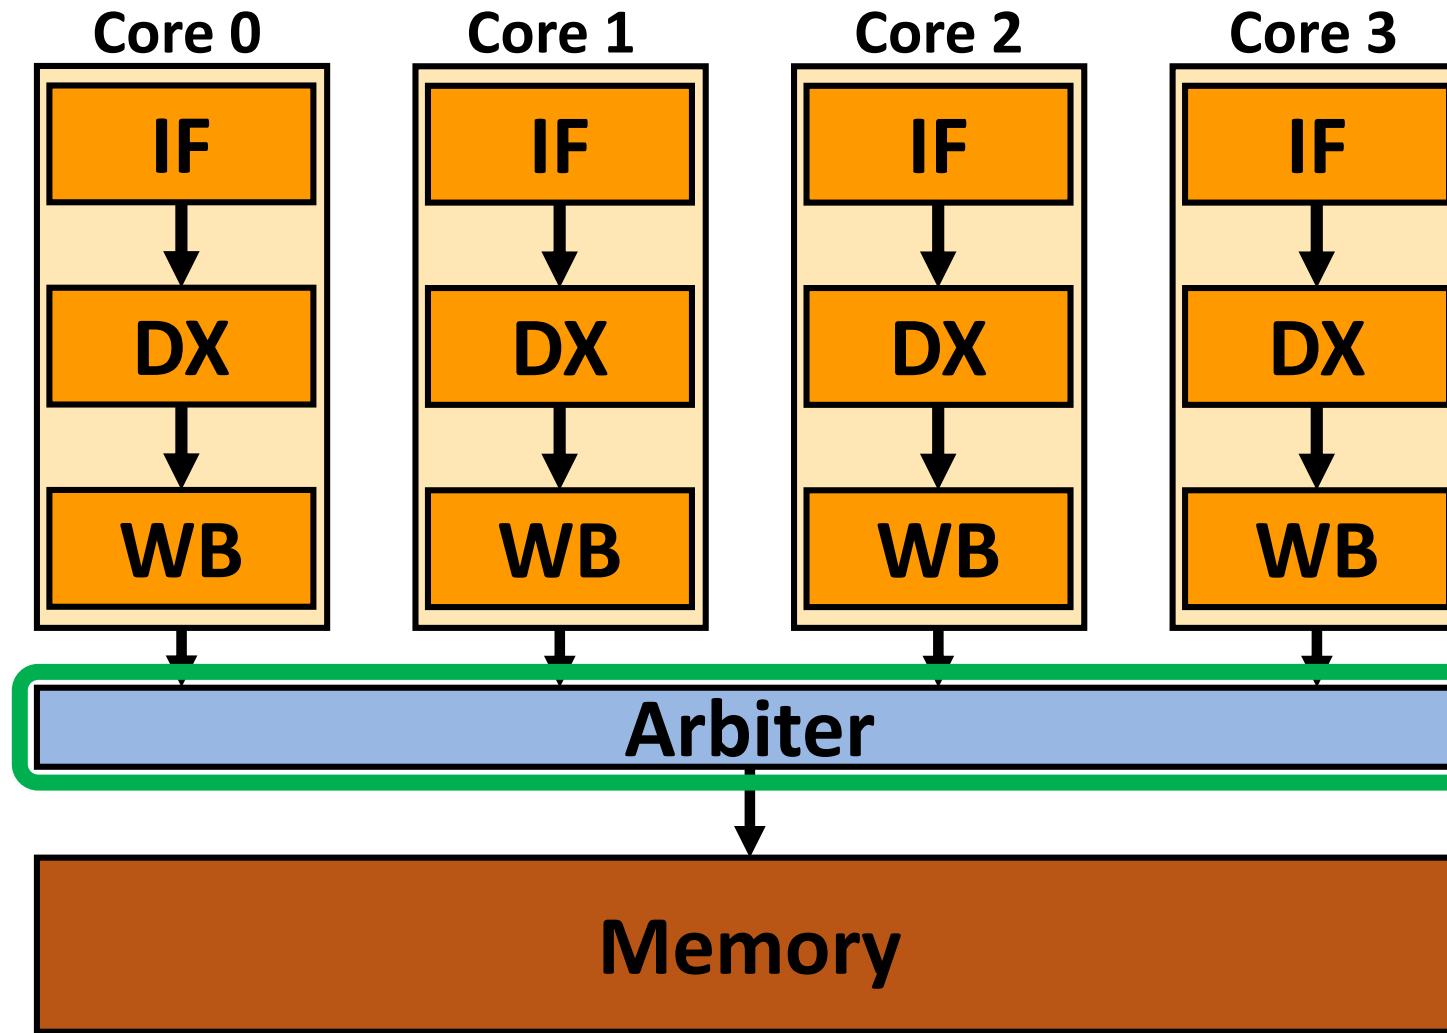
# Multi-V-scale: a Multicore Case Study



# Multi-V-scale: a Multicore Case Study

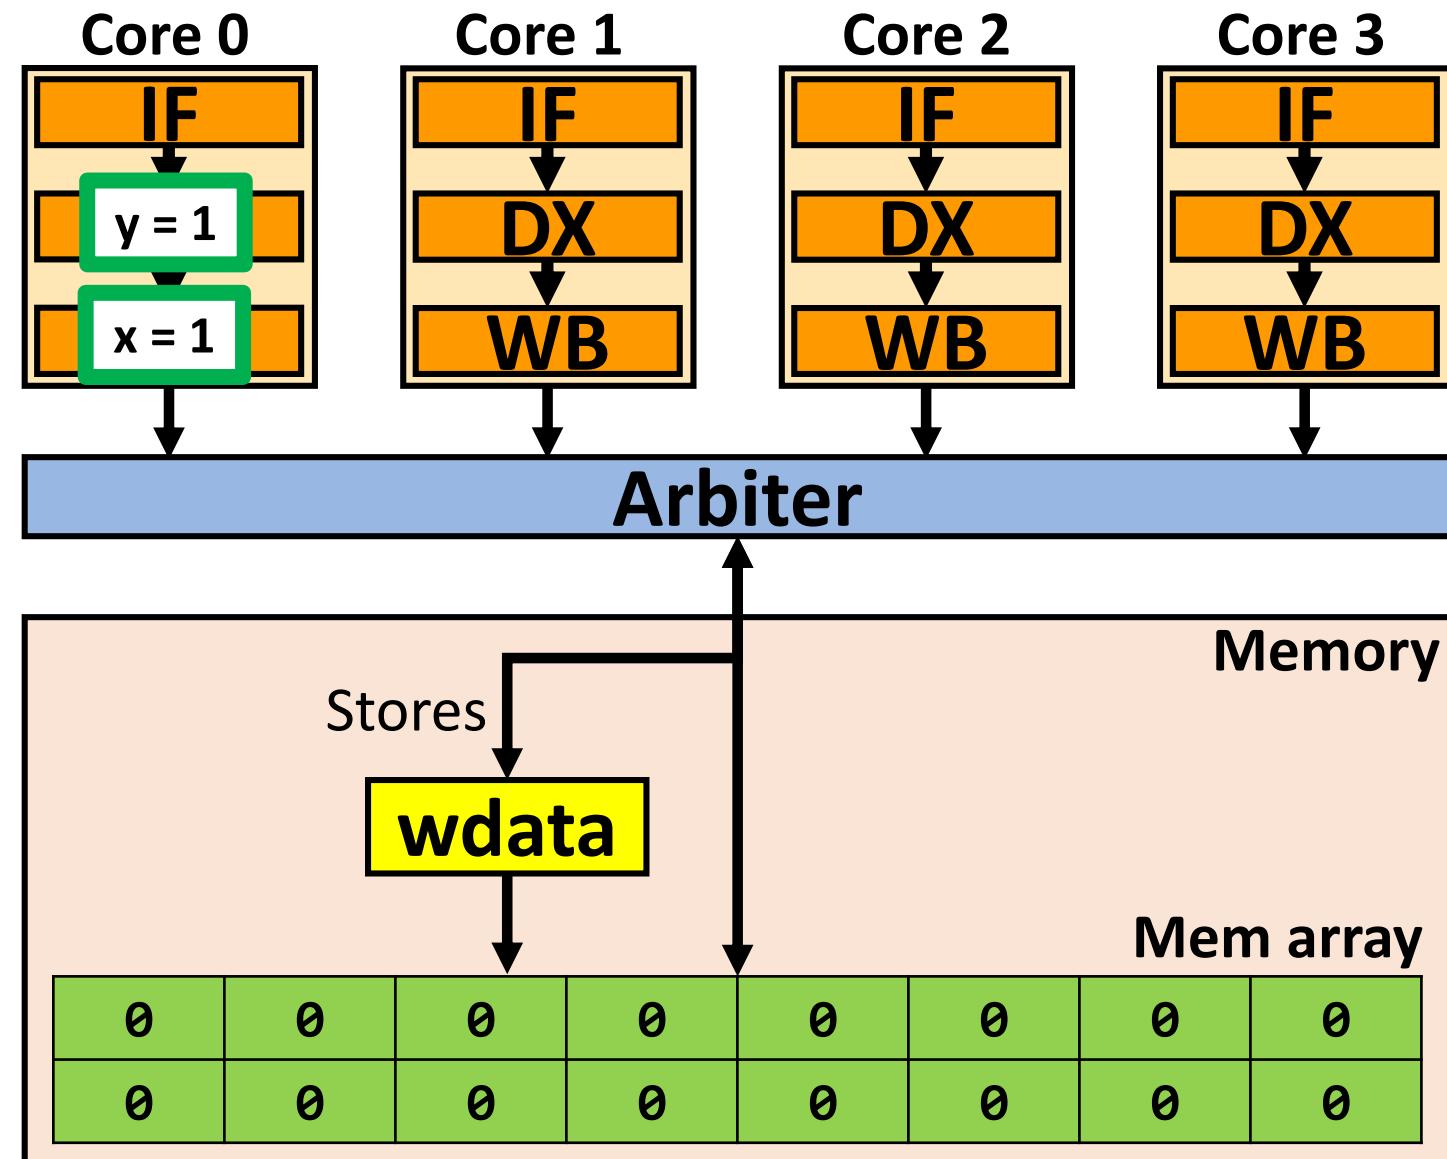


# Multi-V-scale: a Multicore Case Study



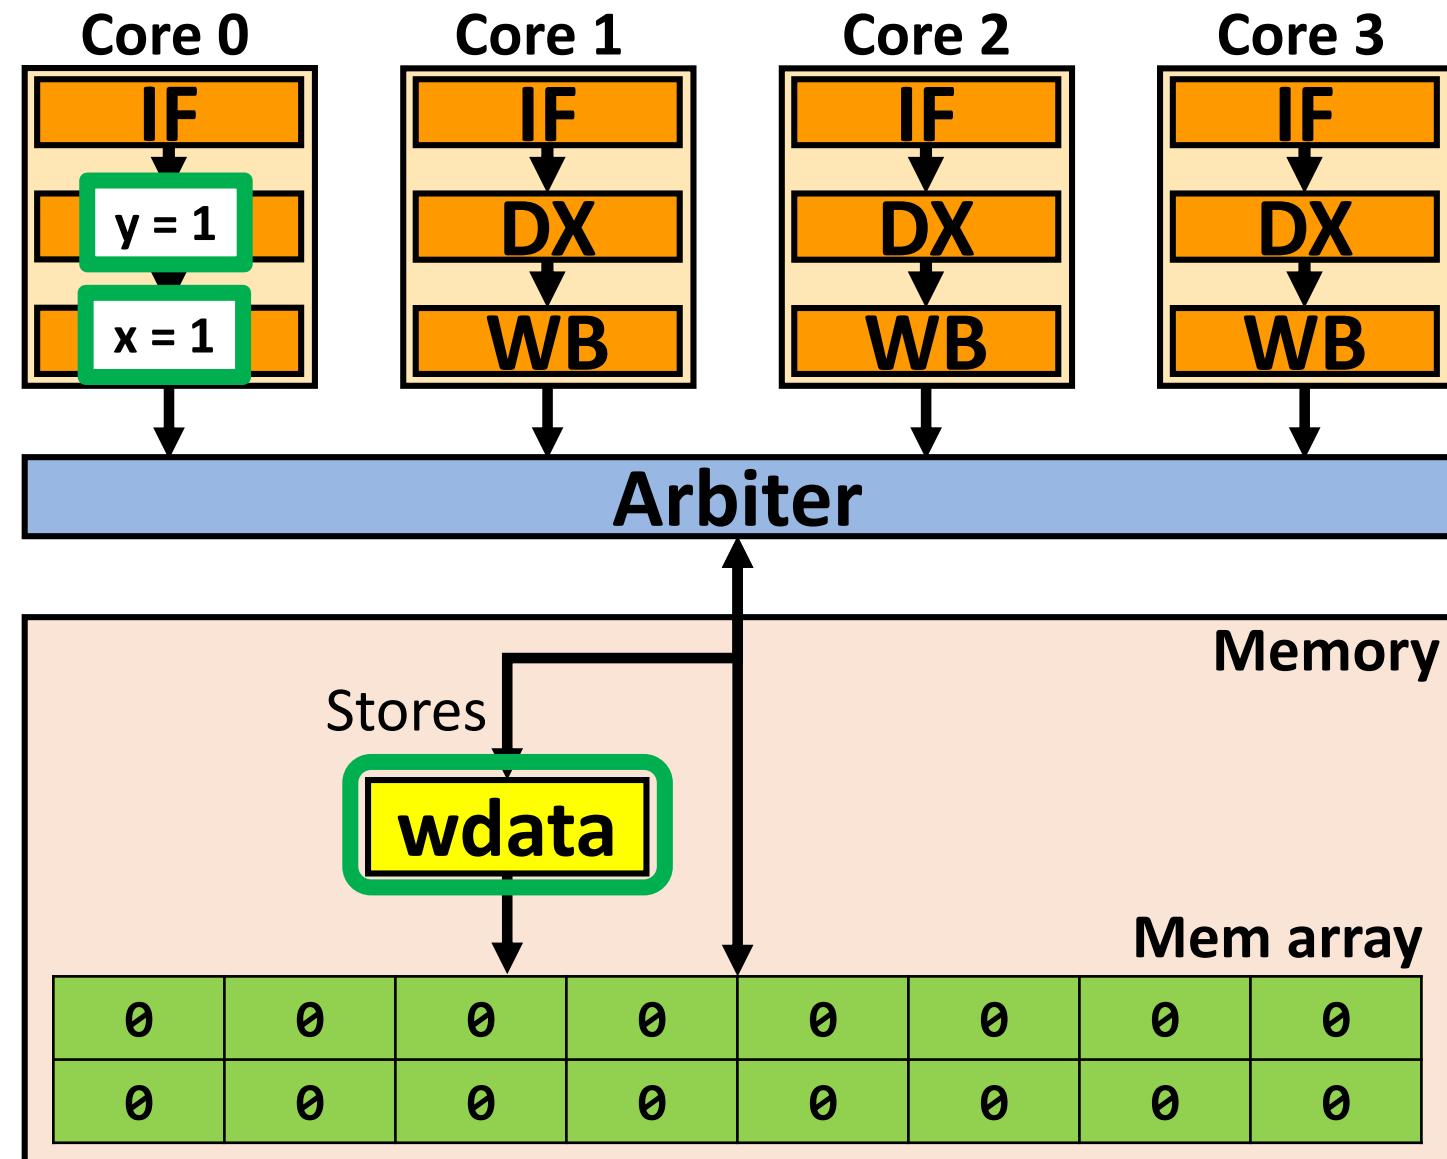
# Bug Discovered in V-scale

- V-scale memory internally writes stores to **wdata** register
- **wdata** pushed to memory when subsequent store occurs
- Akin to single-entry store buffer
- **When two stores are sent to memory in successive cycles, first of two stores is dropped by memory!**
- Fixed bug by eliminating **wdata**
- V-scale has since been deprecated by RISC-V Foundation



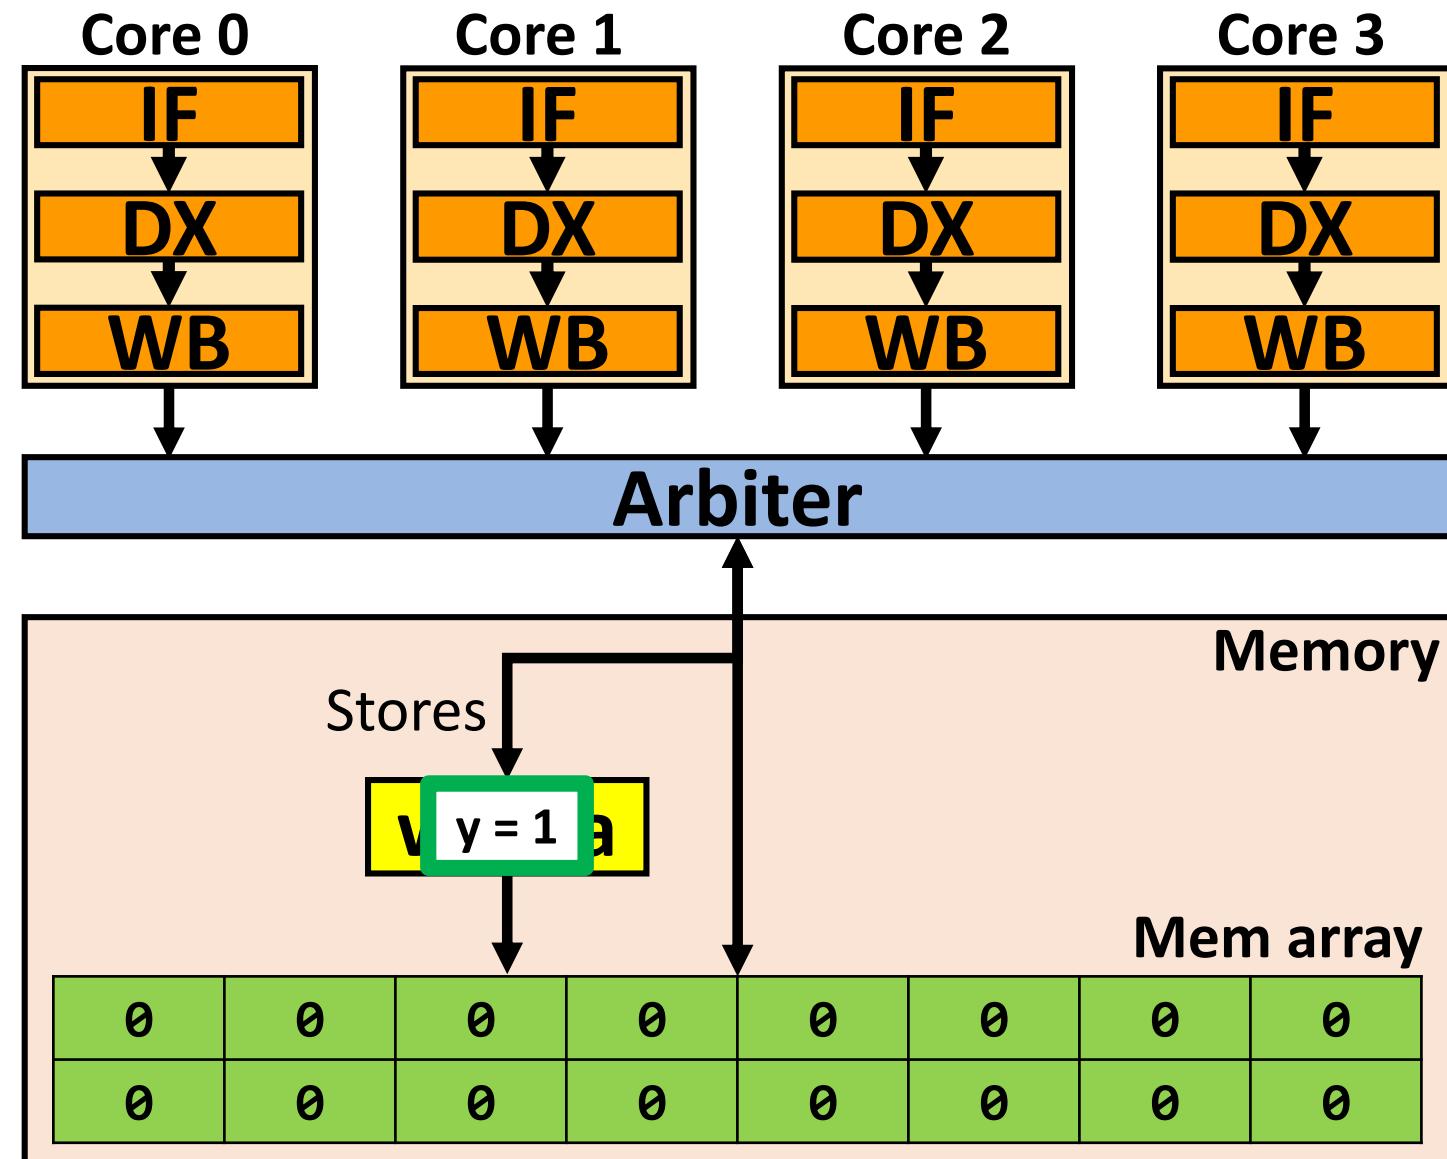
# Bug Discovered in V-scale

- V-scale memory internally writes stores to **wdata** register
- **wdata** pushed to memory when subsequent store occurs
- Akin to single-entry store buffer
- **When two stores are sent to memory in successive cycles, first of two stores is dropped by memory!**
- Fixed bug by eliminating **wdata**
- V-scale has since been deprecated by RISC-V Foundation



# Bug Discovered in V-scale

- V-scale memory internally writes stores to **wdata** register
- **wdata** pushed to memory when subsequent store occurs
- Akin to single-entry store buffer
- **When two stores are sent to memory in successive cycles, first of two stores is dropped by memory!**
- Fixed bug by eliminating **wdata**
- V-scale has since been deprecated by RISC-V Foundation

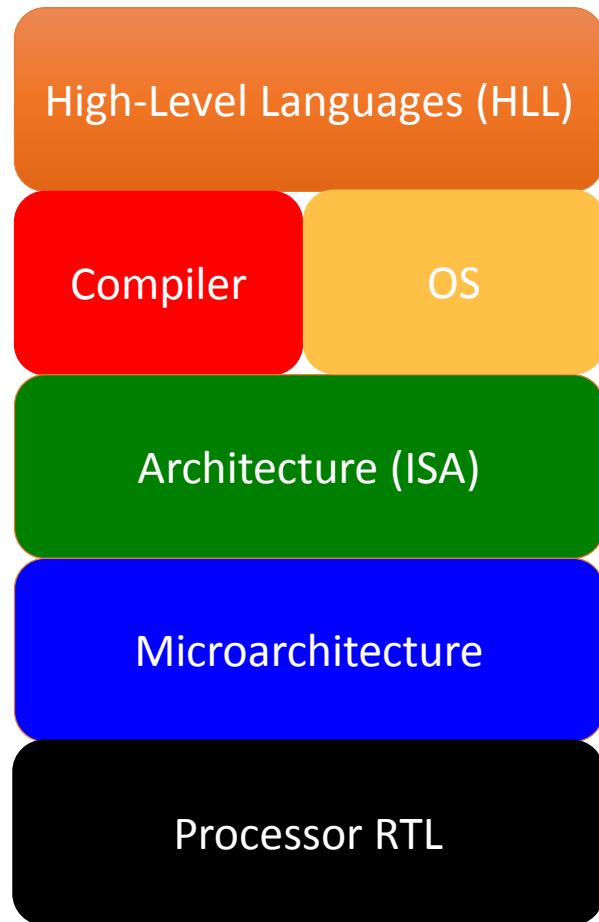


# RTLCheck Takeaways

- Microarchitectural models must be validated against RTL
- **RTLCheck**: Automated translation of **microarch. axioms** into equivalent temporal **SVA properties** for litmus test suites
  - Translation is complicated by the axiomatic-temporal mismatch
  - JasperGold was able to prove 90% of properties/test in 11 hours runtime
- Last piece of the Check suite; now have tools at all levels of the stack!



# Conclusion



- The Check suite provides automated full-stack MCM checking of implementations
- Litmus-test based verification to concentrate on error-prone cases
- Can check:
  - Implementation of HLL requirements
  - Virtual memory implementation
  - HLL Compiler mappings
  - Microarchitectural Orderings (including coherence)
  - and even RTL (Verilog)!
- All tools are open-source and publicly available!



# With Thanks to...

- Collaborators:

- Margaret Martonosi
- Daniel Lustig
- Caroline Trippel
- Michael Pellauer
- Aarti Gupta

- Funding:

- Princeton Wallace Memorial Honorific Fellowship
- STARnet C-FAR (Center for Future Architectures Research)
- JUMP ADA Center (Applications Driving Architectures)
- National Science Foundation



# Questions?

<http://www.cs.princeton.edu/~manerkar>

- **Yatin A. Manerkar**, Daniel Lustig, Margaret Martonosi, and Michael Pellauer. RTLCheck: Verifying the Memory Consistency of RTL Designs. The 50th Annual IEEE/ACM International Symposium on Microarchitecture (MICRO), October 2017.
- **Yatin A. Manerkar**, Caroline Trippel, Daniel Lustig, Michael Pellauer, and Margaret Martonosi. Counterexamples and Proof Loophole for the C/C++ to POWER and ARMv7 Trailing-Sync Compiler Mappings. CoRR abs/1611.01507, November 2016.
- Caroline Trippel, **Yatin A. Manerkar**, Daniel Lustig, Michael Pellauer, and Margaret Martonosi. TriCheck: Memory Model Verification at the Trisection of Software, Hardware, and ISA. The 22nd International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS), April 2017.
- **Yatin A. Manerkar**, Daniel Lustig, Michael Pellauer, and Margaret Martonosi. CCICheck: Using  $\mu$ hb Graphs to Verify the Coherence-Consistency Interface. The 48th Annual IEEE/ACM International Symposium on Microarchitecture (MICRO), December 2015.



<http://check.cs.princeton.edu/>

# Coherence and Consistency

- Most coherence protocols are not that simple!
  - Partial incoherence (e.g. GPUs) [Wickerson et al. OOPSLA 2016]
  - Lazy coherence (e.g. TSO-CC) [Elver and Nagarajan HPCA 2014]
- **CCI: Coherence-Consistency Interface**

Coherence

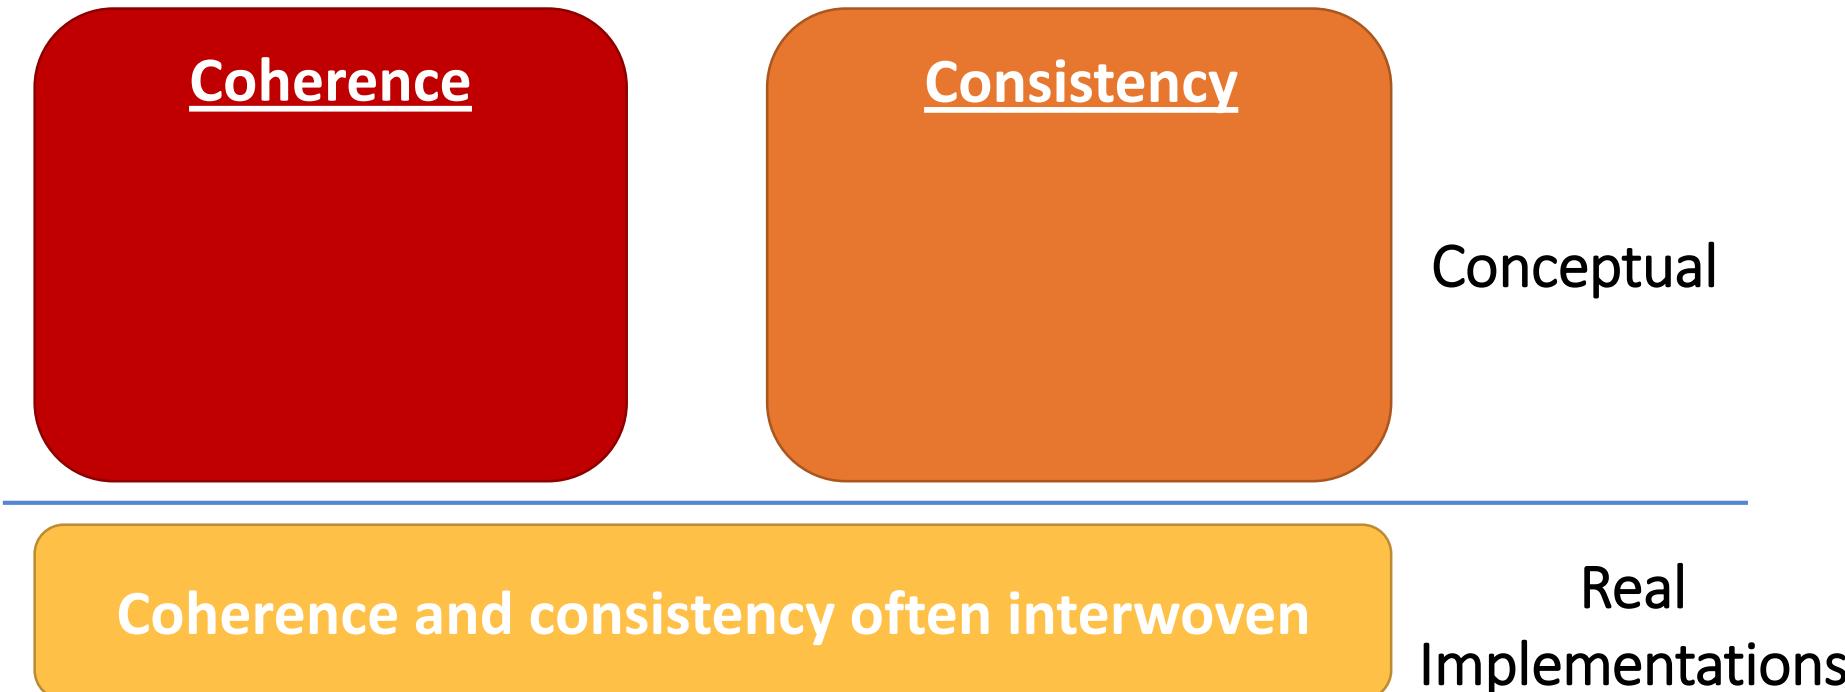
Consistency

Conceptual



# Coherence and Consistency

- Most coherence protocols are not that simple!
  - Partial incoherence (e.g. GPUs) [Wickerson et al. OOPSLA 2016]
  - Lazy coherence (e.g. TSO-CC) [Elver and Nagarajan HPCA 2014]
- **CCI: Coherence-Consistency Interface**



# Coherence and Consistency

- Most coherence protocols are not that simple!
  - Partial incoherence (e.g. GPUs) [Wickerson et al. OOPSLA 2016]
  - Lazy coherence (e.g. TSO-CC) [Elver and Nagarajan HPCA 2014]
- CCI: Coherence-Consistency Interface

## Coherence

Verifiers can't ignore consistency implications!

## Consistency

Verifiers can't assume abstract coherence/memory hierarchy!

Conceptual

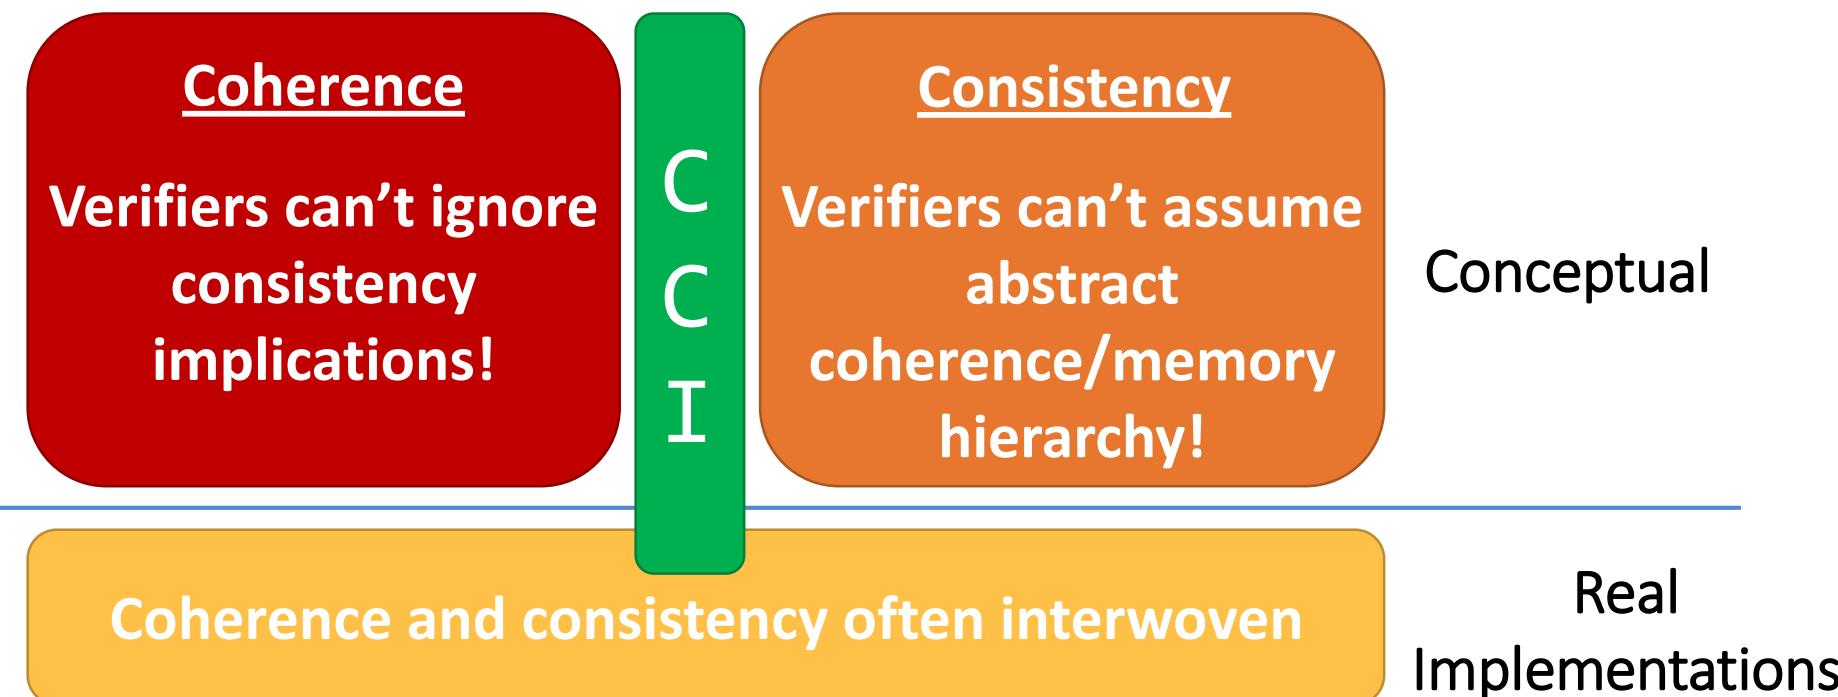
Coherence and consistency often interwoven

Real  
Implementations



# Coherence and Consistency

- Most coherence protocols are not that simple!
  - Partial incoherence (e.g. GPUs) [Wickerson et al. OOPSLA 2016]
  - Lazy coherence (e.g. TSO-CC) [Elver and Nagarajan HPCA 2014]
- CCI: Coherence-Consistency Interface



# Issue with Draft RISC-V MCM: Cumulativity

- Consider this litmus test variant (WRC):

- C11 atomics can specify memory orderings: REL = release, ACQ = acquire

Thread 0	Thread 1	Thread 2
St (x, 1, REL)	r0 = Ld (x, ACQ)	r1 = Ld (y, ACQ)
	St (y, 1, REL)	r2 = Ld (x, ACQ)
<b>Forbidden</b> by C11: r0 = 1, r1 = 1, r2 = 0		

- RISC-V lacked cumulative fences to enforce this ordering:

- (x5 and x6 contain addresses of x and y)

Core 0	Core 1	Core 2
sw x1, (x5)	lw x2, (x5)	lw x3, (x6)
	fence r, rw	fence r, rw
	fence rw, w	lw x4, (x5)
	sw x2, (x6)	
<b>Allowed</b> by draft RISC-V: x1 = 1, x2 = 1, x3 = 1, x4 = 0		



# Issue with Draft RISC-V MCM: Cumulativity

- Consider this litmus test variant (WRC):

- C11 atomics can specify memory orderings: REL = release, ACQ = acquire

Thread 0	Thread 1	Thread 2
St (x, 1, REL)	r0 = Ld (x, ACQ)	r1 = Ld (y, ACQ)
	St (y, 1, REL)	r2 = Ld (x, ACQ)
<b>Forbidden</b> by C11: r0 = 1, r1 = 1, r2 = 0		

- RISC-V lacked cumulative fences to enforce this ordering:

- (x5 and x6 contain addresses of x and y)

Core 0	Core 1	Core 2
sw x1, (x5)	lw x2, (x5)	lw x3, (x6)
	fence r, rw	fence r, rw
	fence rw, w	lw x4, (x5)
	sw x2, (x6)	
<b>Allowed</b> by draft RISC-V: x1 = 1, x2 = 1, x3 = 1, x4 = 0		



# Issue with Draft RISC-V MCM: Cumulativity

- Consider this litmus test variant (WRC):

- C11 atomics can specify memory orderings: REL = release, ACQ = acquire

Thread 0	Thread 1	Thread 2
St (x, 1, REL)	r0 = Ld (x, ACQ)	r1 = Ld (y, ACQ)
	St (y, 1, REL)	r2 = Ld (x, ACQ)
<b>Forbidden</b> by C11: r0 = 1, r1 = 1, r2 = 0		

- RISC-V lacked cumulative fences to enforce this ordering:

- (x5 and x6 contain addresses of x and y)

Core 0	Core 1	Core 2
sw x1, (x5)	lw x2, (x5)	lw x3, (x6)
	fence r, rw	fence r, rw
	fence rw, w	lw x4, (x5)
	sw x2, (x6)	
<b>Allowed</b> by draft RISC-V: x1 = 1, x2 = 1, x3 = 1, x4 = 0		



# Issue with Draft RISC-V MCM: Cumulativity

- Consider this litmus test variant (WRC):

- C11 atomics can specify memory orderings: REL = release, ACQ = acquire

Thread 0	Thread 1	Thread 2
St (x, 1, REL)	r0 = Ld (x, ACQ)	r1 = Ld (y, ACQ)
	St (y, 1, REL)	r2 = Ld (x, ACQ)
<b>Forbidden</b> by C11: r0 = 1, r1 = 1, r2 = 0		

- RISC-V lacked cumulative fences to enforce this ordering:

- (x5 and x6 contain addresses of x and y)

Core 0	Core 1	Core 2
sw x1, (x5)	lw x2, (x5)	lw x3, (x6)
	fence r, rw	fence r, rw
	fence rw, w	lw x4, (x5)
	sw x2, (x6)	
<b>Allowed</b> by draft RISC-V: x1 = 1, x2 = 1, x3 = 1, x4 = 0		



# Issue with Draft RISC-V MCM: Cumulativity

- Consider this litmus test variant (WRC):

- C11 atomics can specify memory orderings: REL = release, ACQ = acquire

Thread 0	Thread 1	Thread 2
St (x, 1, REL)	r0 = Ld (x, ACQ)	r1 = Ld (y, ACQ)
	St (y, 1, REL)	r2 = Ld (x, ACQ)
<b>Forbidden</b> by C11: r0 = 1, r1 = 1, r2 = 0		

- RISC-V lacked cumulative fences to enforce this ordering:

- (x5 and x6 contain addresses of x and y)

Core 0	Core 1	Core 2
sw x1, (x5)	lw x2, (x5)	lw x3, (x6)
	fence r, rw	fence r, rw
	fence rw, w	lw x4, (x5)
sw x2, (x6)		
<b>Allowed</b> by draft RISC-V: x1 = 1, x2 = 1, x3 = 1, x4 = 0		



# Issue with Draft RISC-V MCM: Cumulativity

- Consider this litmus test variant (WRC):

- C11 atomics can specify memory orderings: REL = release, ACQ = acquire

Thread 0	Thread 1	Thread 2
St (x, 1, REL)	r0 = Ld (x, ACQ)	r1 = Ld (y, ACQ)
	St (y, 1, REL)	r2 = Ld (x, ACQ)
<b>Forbidden</b> by C11: r0 = 1, r1 = 1, r2 = 0		

- RISC-V lacked cumulative fences to enforce this ordering:

- (x5 and x6 contain addresses of x and y)

Core 0	Core 1	Core 2
sw x1, (x5)	lw x2, (x5)	lw x3, (x6)
	fence r, rw	fence r, rw
	fence rw, w	lw x4, (x5)
	sw x2, (x6)	
<b>Allowed</b> by draft RISC-V: x1 = 1, x2 = 1, x3 = 1, x4 = 0		



# ARMv7/Power Trailing-Sync Counterexample

- Consider this litmus test variant (IRIW):

- Total order over all SC atomic accesses is required

Thread 0	Thread 1	Thread 2	Thread 3
St (x, 1, SC)	St (y, 1, SC)	r0 = Ld (x, ACQ)	r2 = Ld (y, ACQ)
		r1 = Ld (y, SC)	r3 = Ld (x, SC)
<b>Forbidden</b> by C11: r0 = 1, r1 = 0, r2 = 1, r3 = 0			

- With the trailing-sync mapping, this compiles to the following:

- Allowed on Power [Sarkar et al. PLDI 2011] and ARMv7 [Alglave et al. TOPLAS 2014]

Core 0	Core 1	Core 2	Core 3
str 1, [x]	str 1, [y]	ldr r1, [x]	ldr r3, [y]
		ctrlisb/ctrlisync	ctrlisb/ctrlisync
		ldr r2, [y]	ldr r4, [x]
<b>Allowed</b> by Power/ARMv7: r1 = 1, r2 = 0, r3 = 1, r4 = 0			



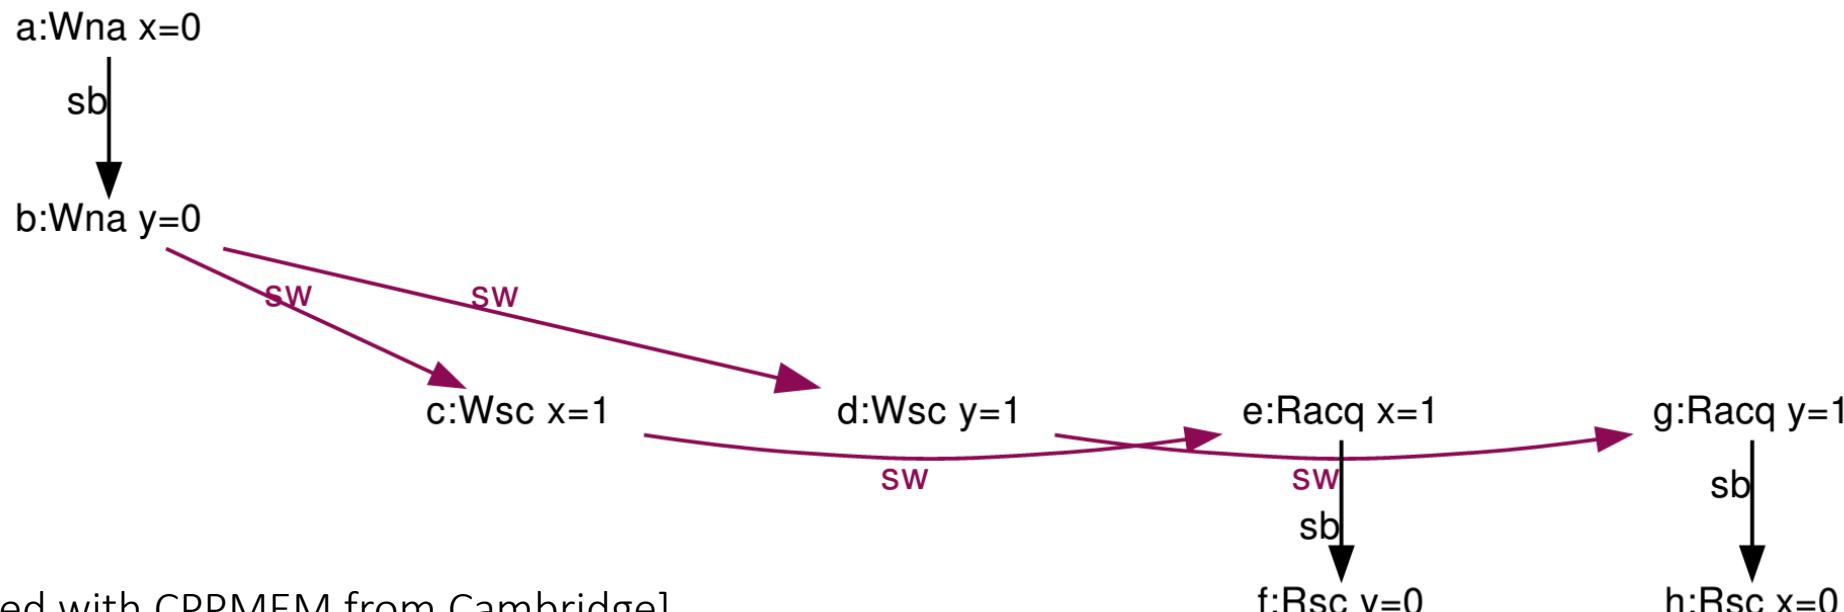
# ARMv7/Power Trailing-Sync Counterexample

- Consider this litmus test variant (IRIW):

- Total order over all SC atomic accesses is required

Thread 0	Thread 1	Thread 2	Thread 3
St (x, 1, SC)	St (y, 1, SC)	r0 = Ld (x, ACQ)	r2 = Ld (y, ACQ)
		r1 = Ld (y, SC)	r3 = Ld (x, SC)
<b>Forbidden</b> by C11: r0 = 1, r1 = 0, r2 = 1, r3 = 0			

- SC total order must respect happens-before i.e.  $(sb \cup sw) +$



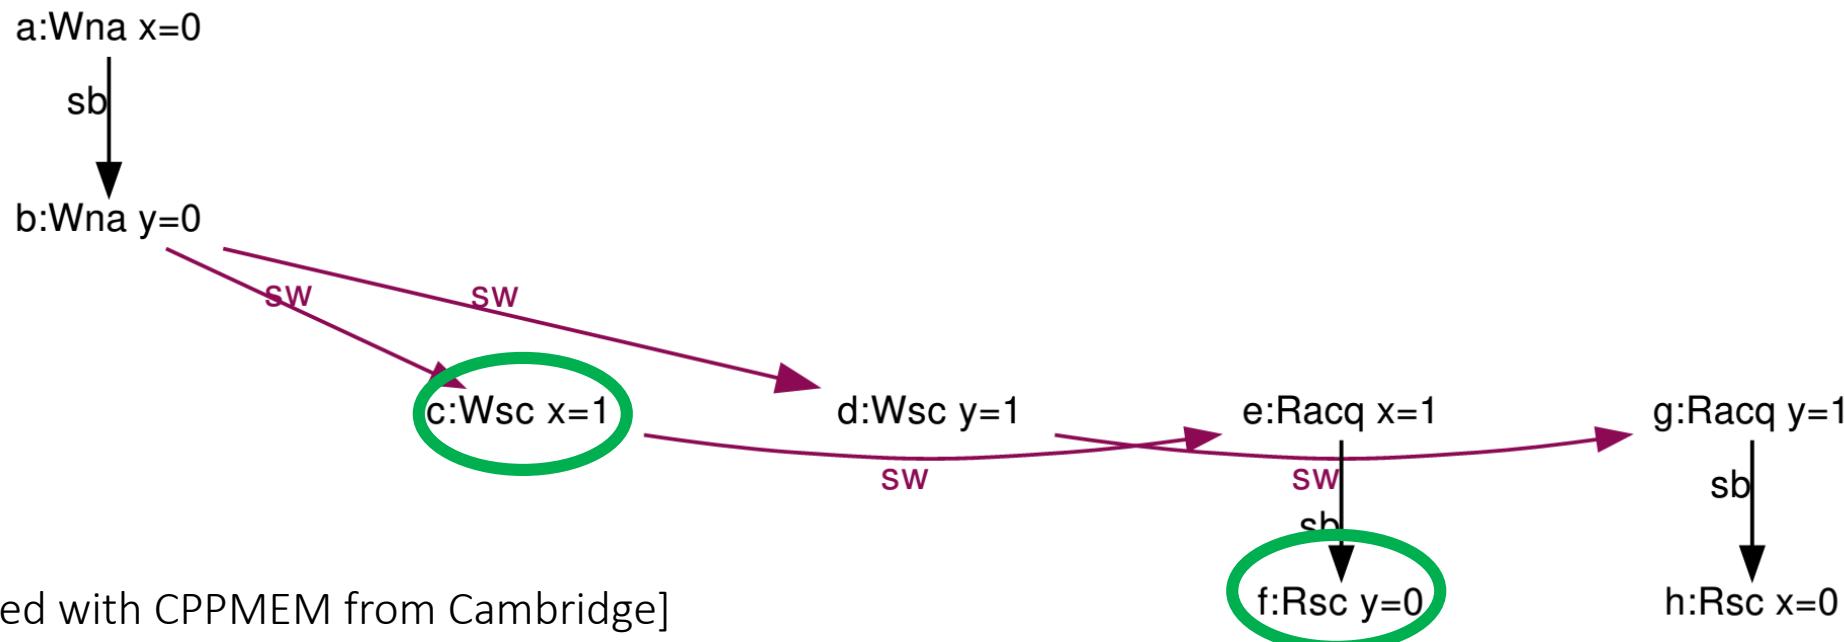
# ARMv7/Power Trailing-Sync Counterexample

- Consider this litmus test variant (IRIW):

- Total order over all SC atomic accesses is required

Thread 0	Thread 1	Thread 2	Thread 3
St (x, 1, SC)	St (y, 1, SC)	r0 = Ld (x, ACQ)	r2 = Ld (y, ACQ)
		r1 = Ld (y, SC)	r3 = Ld (x, SC)
<b>Forbidden</b> by C11: r0 = 1, r1 = 0, r2 = 1, r3 = 0			

- SC total order must respect happens-before i.e.  $(sb \cup sw) +$



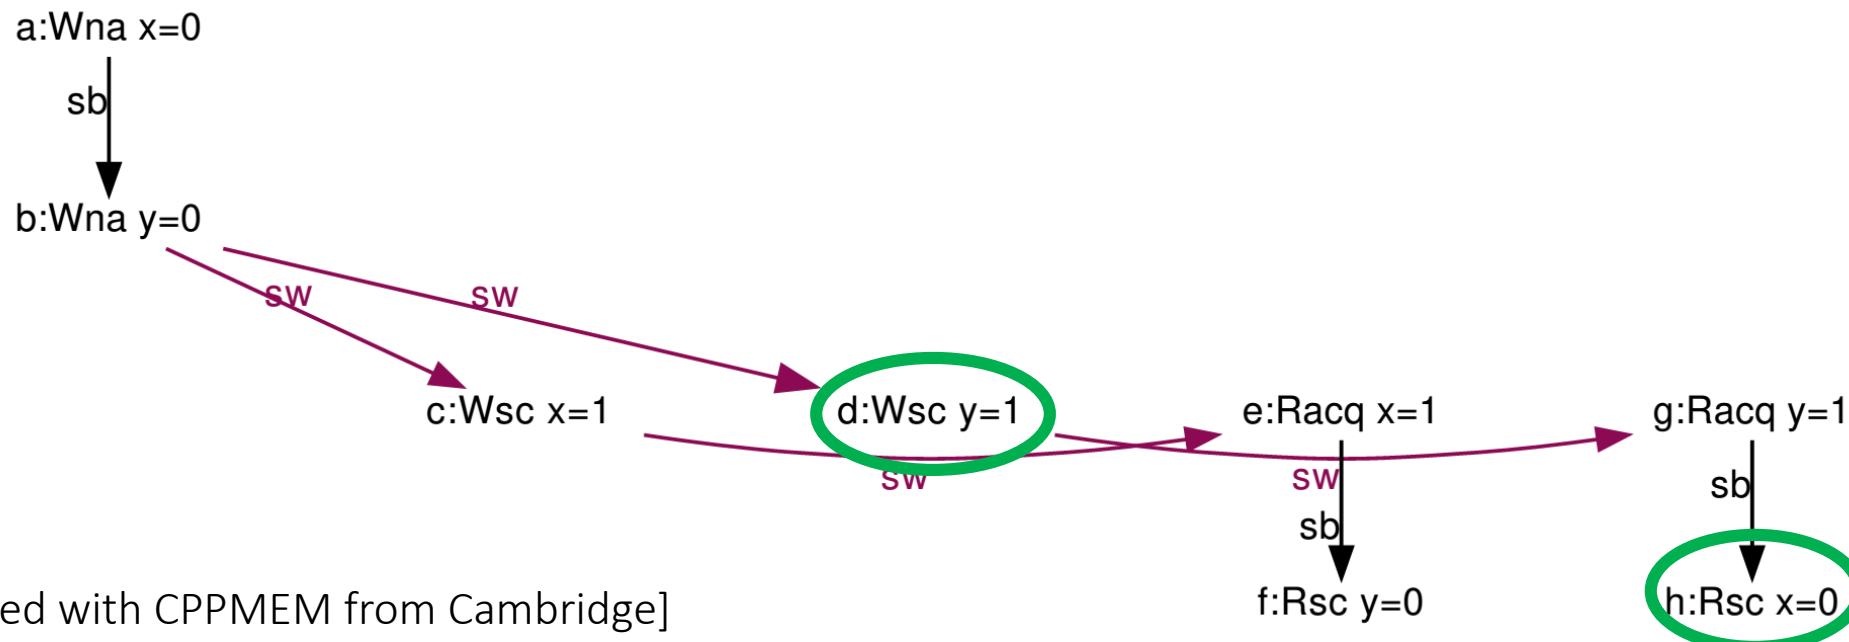
# ARMv7/Power Trailing-Sync Counterexample

- Consider this litmus test variant (IRIW):

- Total order over all SC atomic accesses is required

Thread 0	Thread 1	Thread 2	Thread 3
St (x, 1, SC)	St (y, 1, SC)	r0 = Ld (x, ACQ)	r2 = Ld (y, ACQ)
		r1 = Ld (y, SC)	r3 = Ld (x, SC)
<b>Forbidden</b> by C11: r0 = 1, r1 = 0, r2 = 1, r3 = 0			

- SC total order must respect happens-before i.e.  $(sb \cup sw) +$



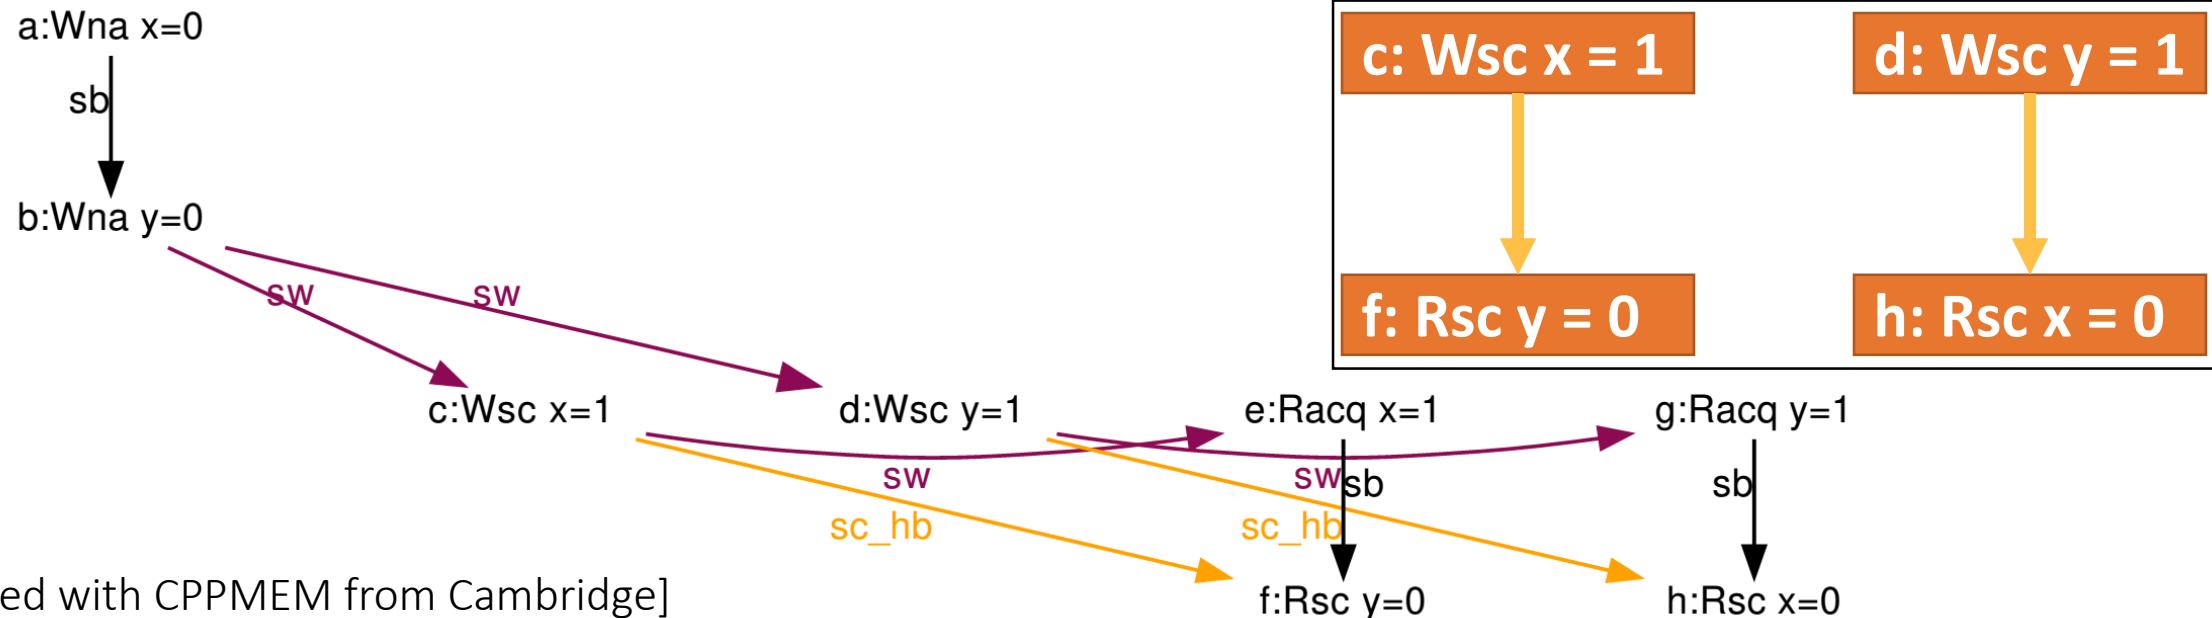
# ARMv7/Power Trailing-Sync Counterexample

- Consider this litmus test variant (IRIW):

- Total order over all SC atomic accesses is required

Thread 0	Thread 1	Thread 2	Thread 3
St (x, 1, SC)	St (y, 1, SC)	r0 = Ld (x, ACQ)	r2 = Ld (y, ACQ)
		r1 = Ld (y, SC)	r3 = Ld (x, SC)
<b>Forbidden</b> by C11: r0 = 1, r1 = 0, r2 = 1, r3 = 0			

- SC total order must respect happens-before i.e. (sb U sw) +



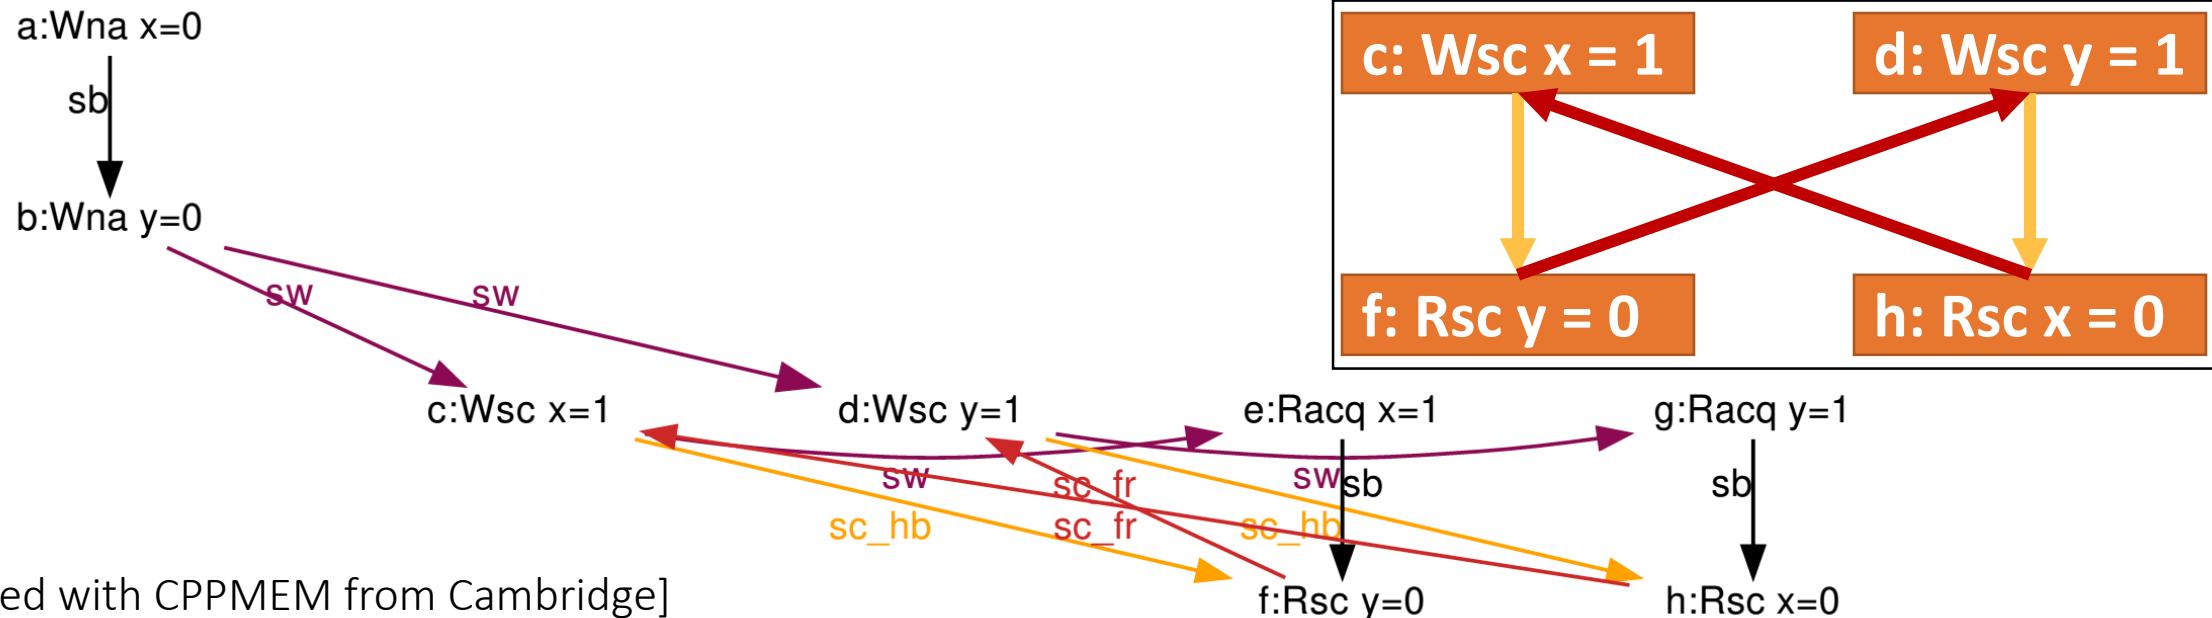
# ARMv7/Power Trailing-Sync Counterexample

- Consider this litmus test variant (IRIW):

- Total order over all SC atomic accesses is required

Thread 0	Thread 1	Thread 2	Thread 3
St (x, 1, SC)	St (y, 1, SC)	r0 = Ld (x, ACQ)	r2 = Ld (y, ACQ)
		r1 = Ld (y, SC)	r3 = Ld (x, SC)
<b>Forbidden</b> by C11: r0 = 1, r1 = 0, r2 = 1, r3 = 0			

- SC reads must be before later SC writes



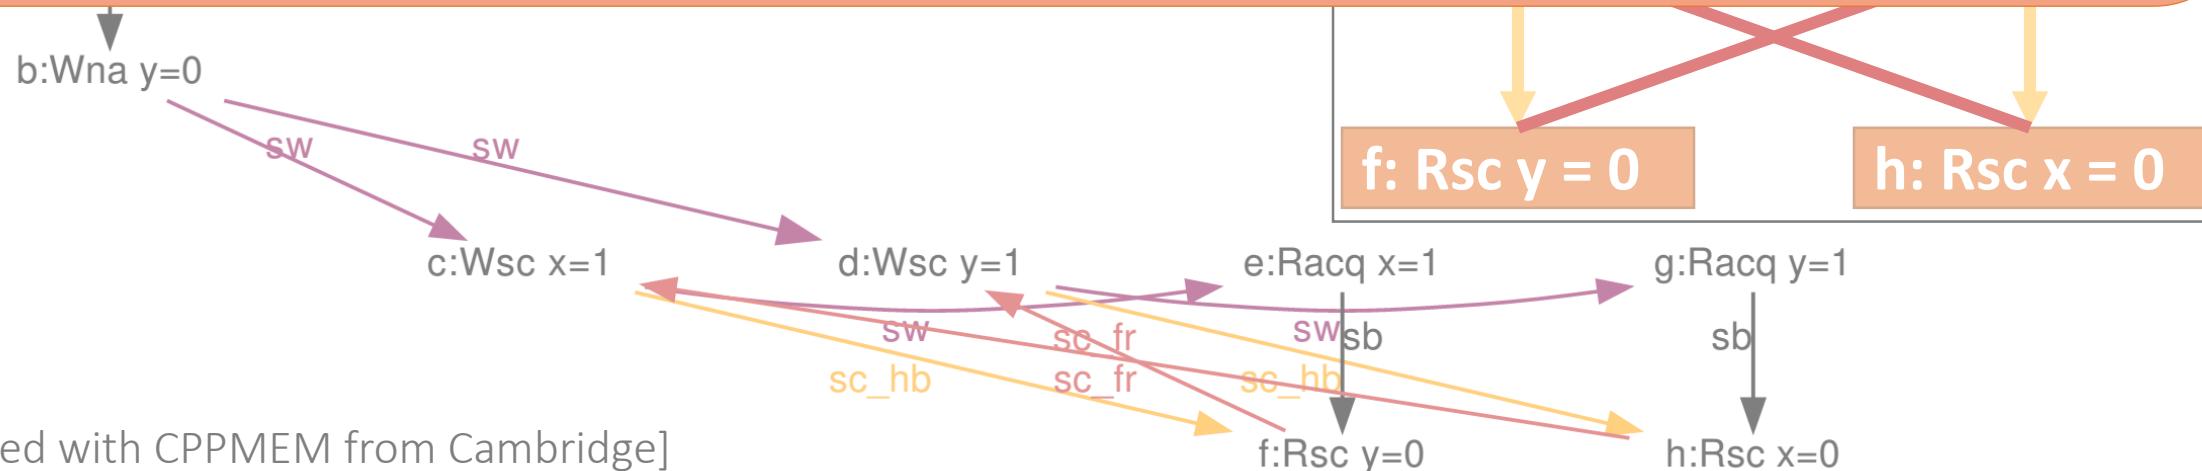
# ARMv7/Power Trailing-Sync Counterexample

- Consider this litmus test variant (IRIW):

- Total order over all SC atomic accesses is required

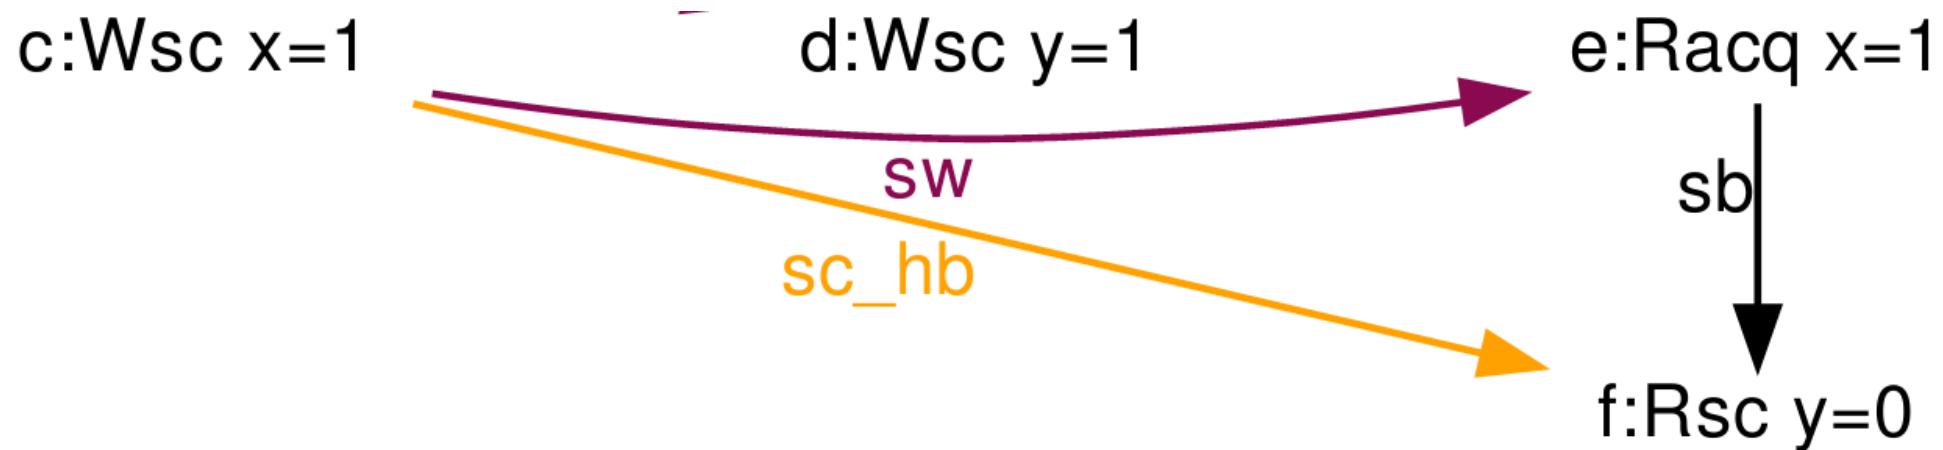
Thread 0 Thread 1 Thread 2 Thread 3

- Cycle in the SC order implies outcome is forbidden**
- But compiled code allows the behaviour!**



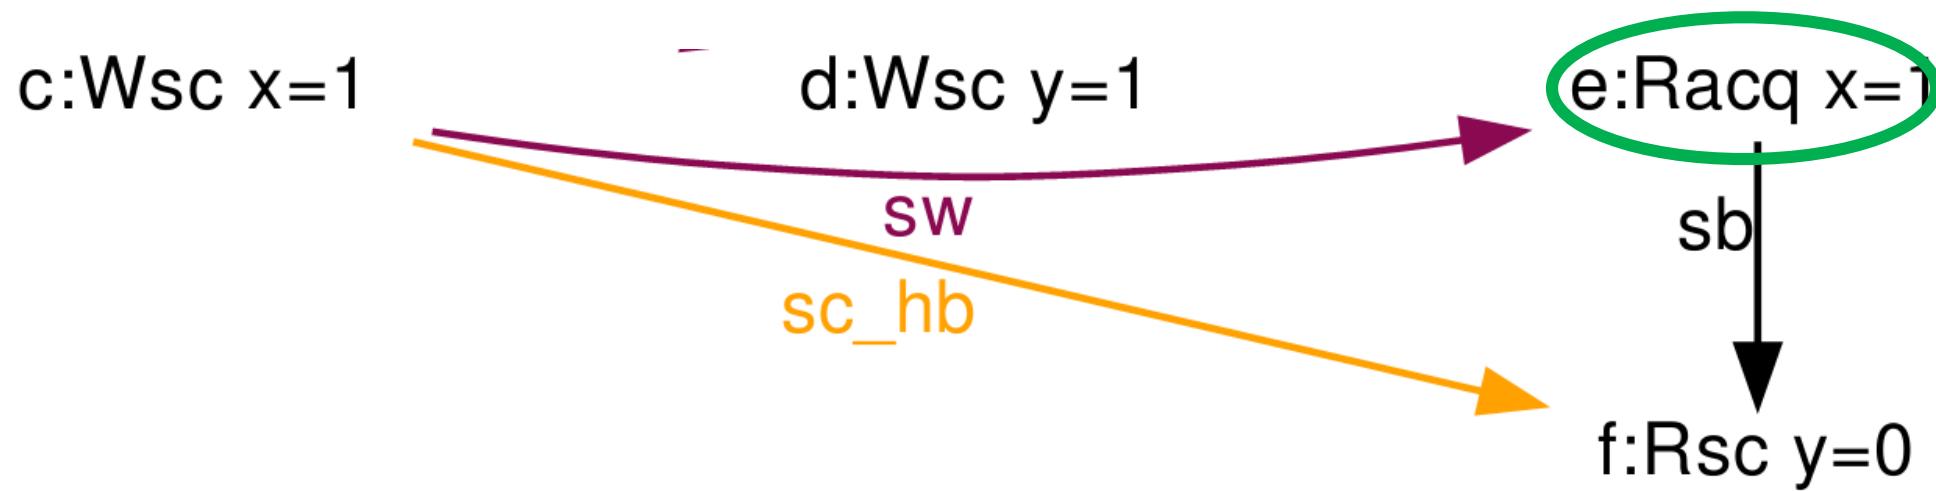
# What went wrong?

- It was thought that program order and coherence edges **directly between SC accesses** were all that needed enforcing [Batty et al. POPL 2012]
- But *hb* edges can arise between SC accesses through the transitive composition of edges to and from a non-SC **intermediate** access
- Occurs in IRIW counterexample:



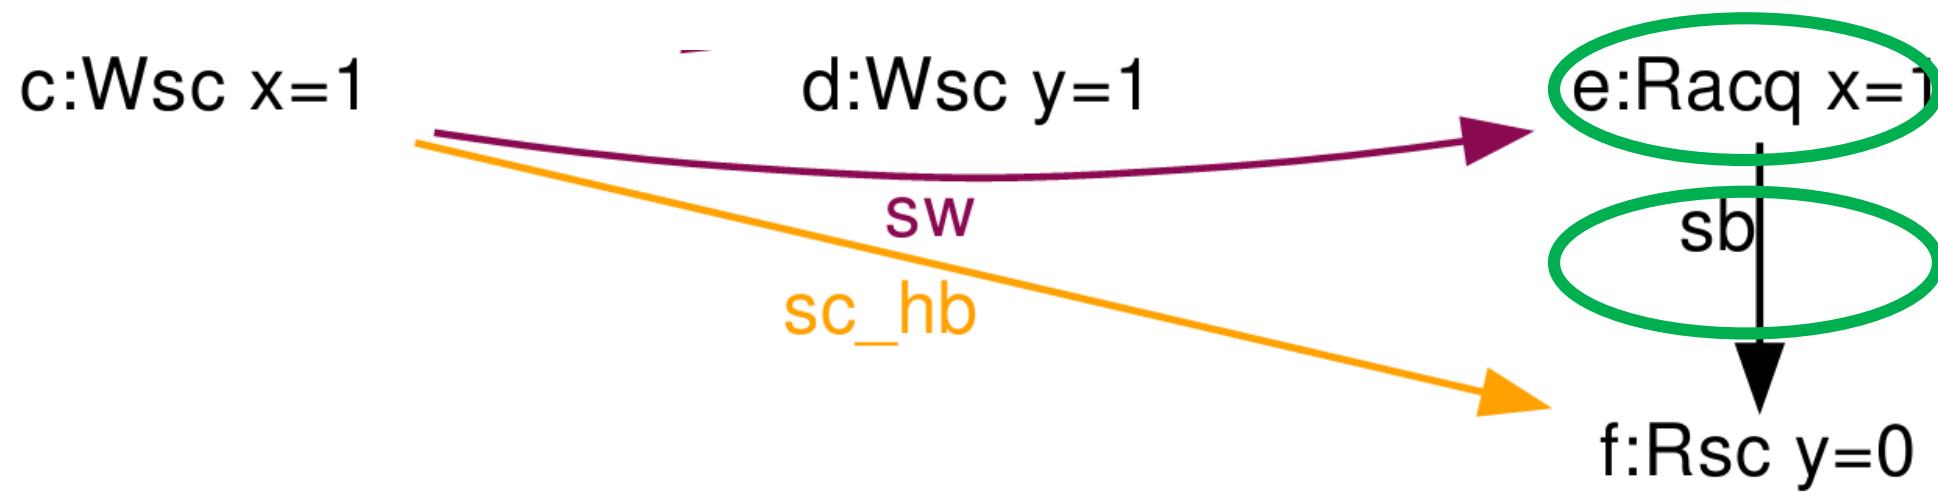
# What went wrong?

- It was thought that program order and coherence edges **directly between SC accesses** were all that needed enforcing [Batty et al. POPL 2012]
- But *hb* edges can arise between SC accesses through the transitive composition of edges to and from a non-SC **intermediate** access
- Occurs in IRIW counterexample:



# What went wrong?

- It was thought that program order and coherence edges **directly between SC accesses** were all that needed enforcing [Batty et al. POPL 2012]
- But *hb* edges can arise between SC accesses through the transitive composition of edges to and from a non-SC **intermediate** access
- Occurs in IRIW counterexample:



# Assumption Generation

- Need to restrict executions to those of litmus test
- Three classes of assumptions:
  - Memory initialization
    - Instr. mem and data mem
  - Register initialization
  - Value assumptions
    - Load value assumptions: loads return correct value (when they occur)
    - Final value assumptions: Required final values of memory are respected
- RTLCheck generates SystemVerilog Assumptions to constrain executions
  - Utilises user-provided **program mapping function**



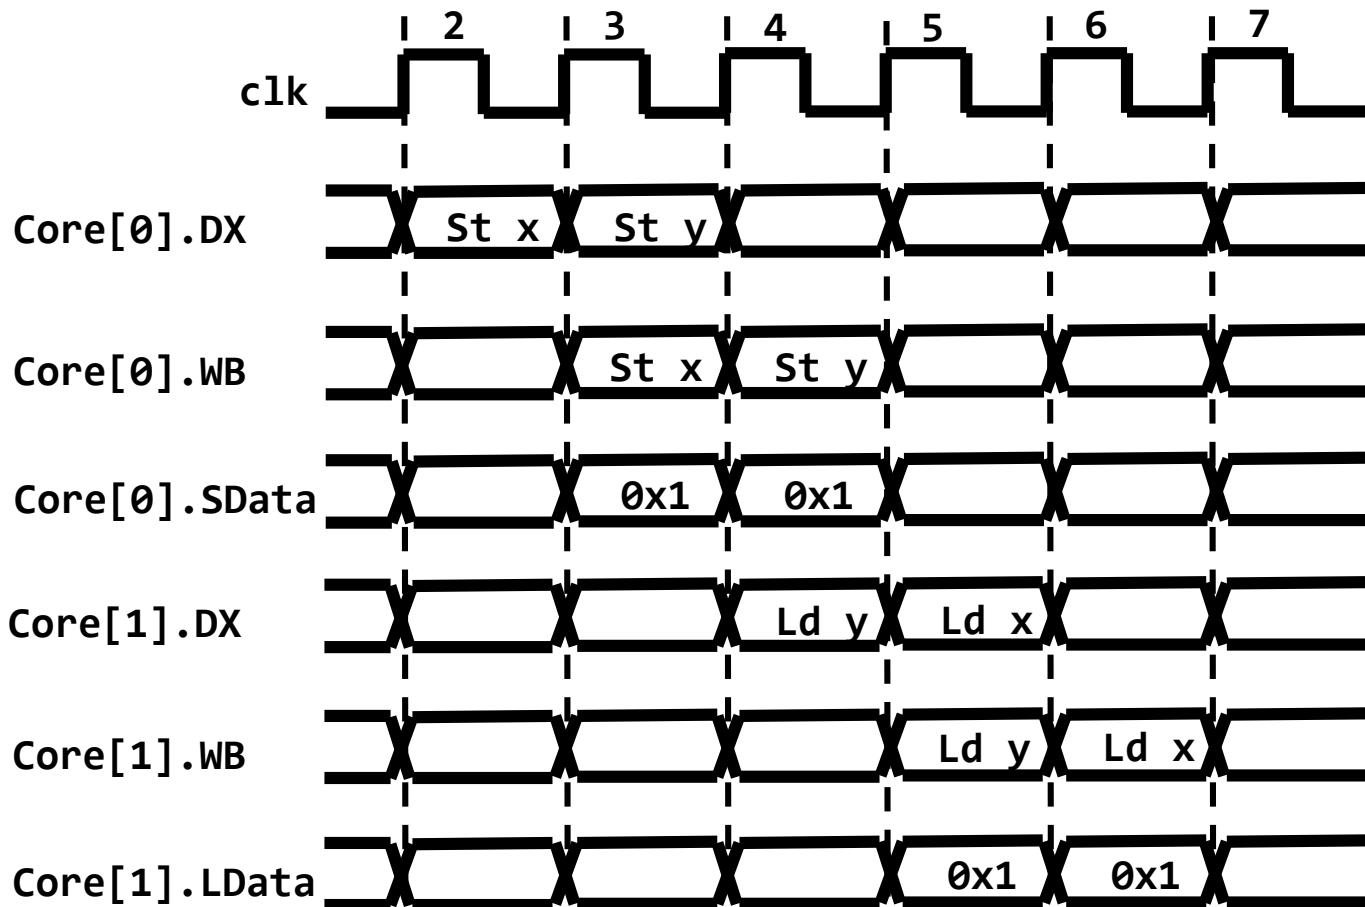
# Assumption Generation

- **Covering trace**: execution where assumption condition is enforced
  - Eg: execution where load of x returns 0
  - Must obey **all** assumptions
- **Covering final value assum. == finding forbidden execution!**
  - No covering trace => equivalent to verifying overall test!
- Quicker verification for some tests
  - Expect benefit to be largest for small designs



# The Benefits of Final Value Assumptions

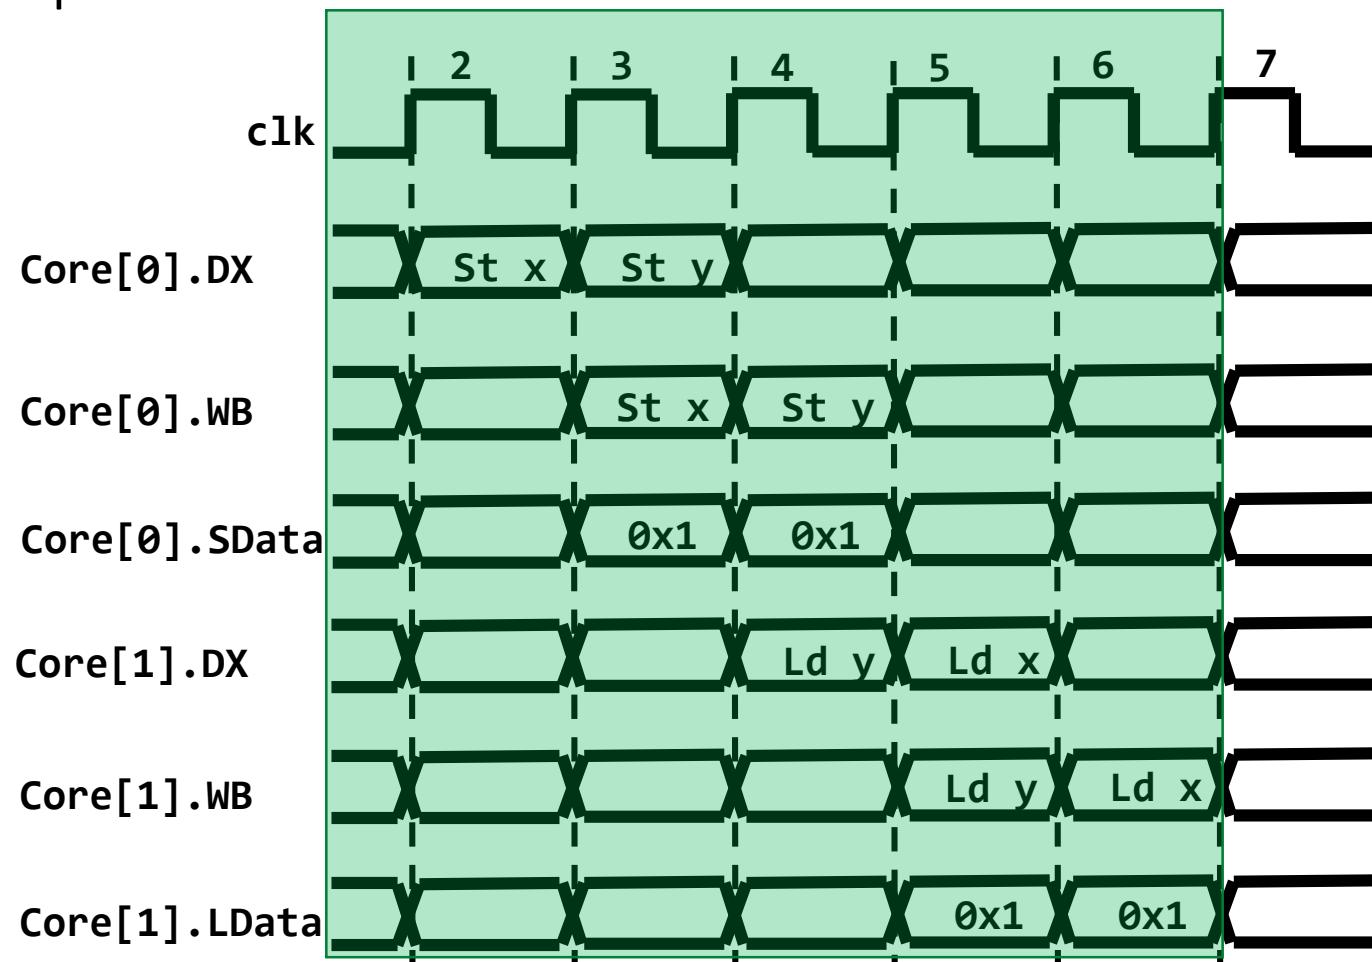
- Why generate final value assumptions if test has no final conditions?
- Answer: Covering traces can lead to faster verification
- These are traces where assumption condition occurs and can be enforced



# The Benefits of Final Value Assumptions

- Why generate final value assumptions if test has no final conditions?
- Answer: Covering traces can lead to faster verification
- These are traces where assumption condition occurs and can be enforced

Covering trace for final val  
assumption is complete  
execution of litmus test

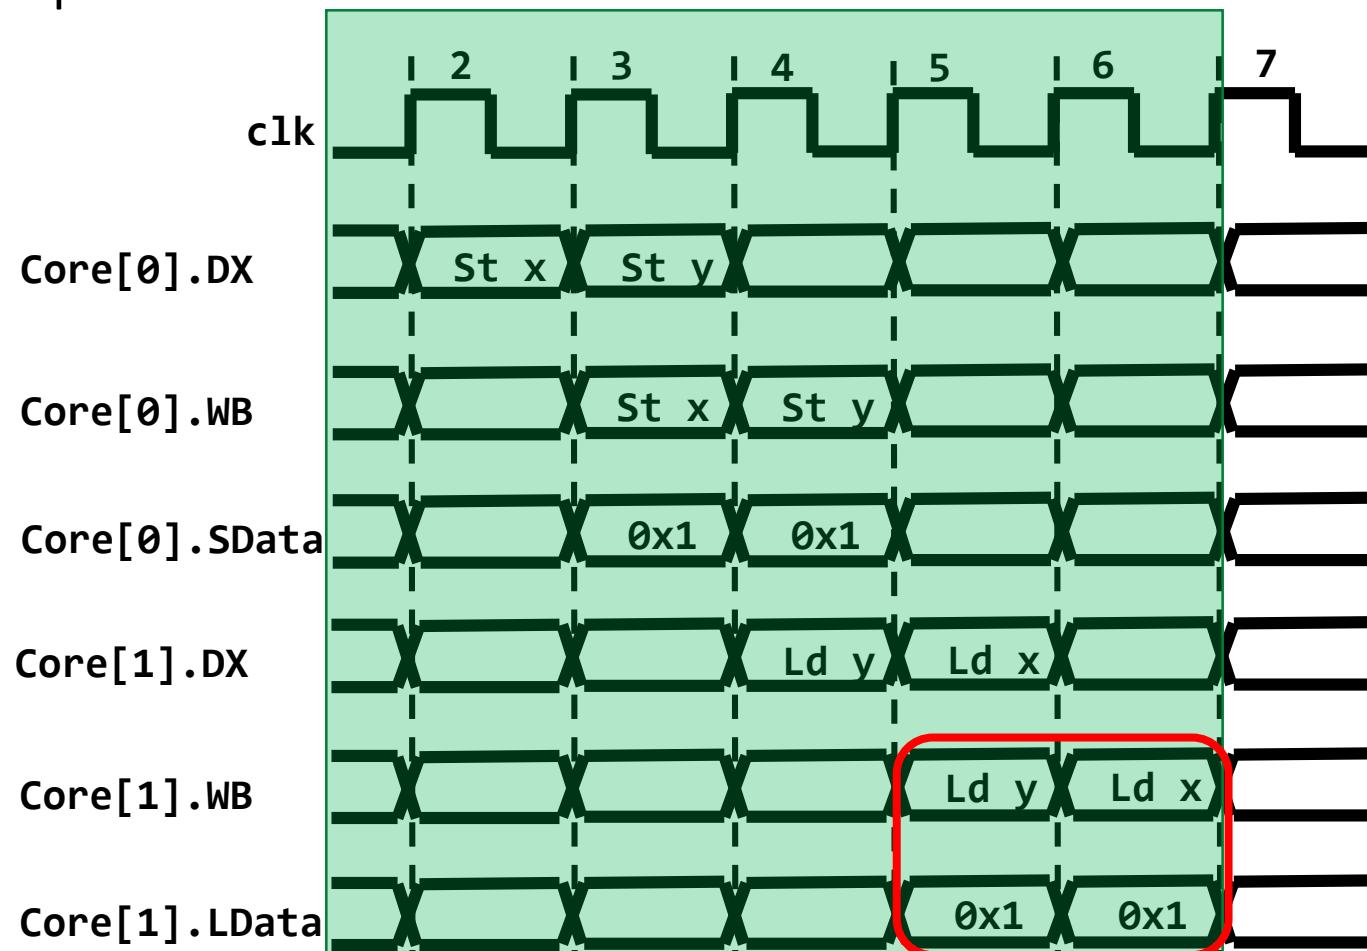


# The Benefits of Final Value Assumptions

- Why generate final value assumptions if test has no final conditions?
- Answer: Covering traces can lead to faster verification
- These are traces where assumption condition occurs and can be enforced

Covering trace for final val assumption is complete execution of litmus test

Covering trace must also obey other assumptions, including load val assumptions  
(For mp, Ld y = 1 and Ld x = 0)



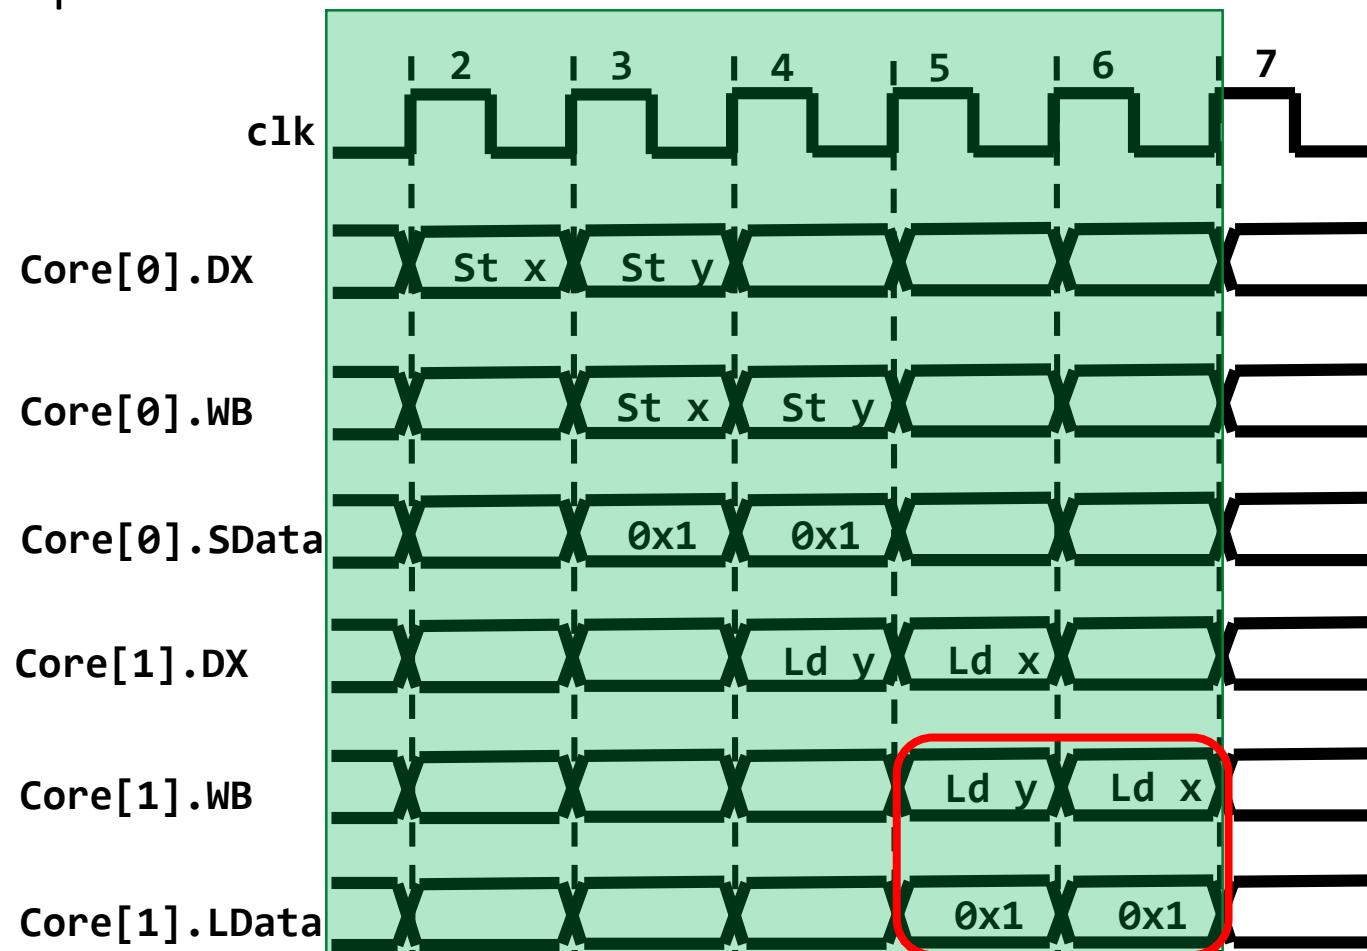
# The Benefits of Final Value Assumptions

- Why generate final value assumptions if test has no final conditions?
- Answer: Covering traces can lead to faster verification
- These are traces where assumption condition occurs and can be enforced

Covering trace for final val assumption is complete execution of litmus test

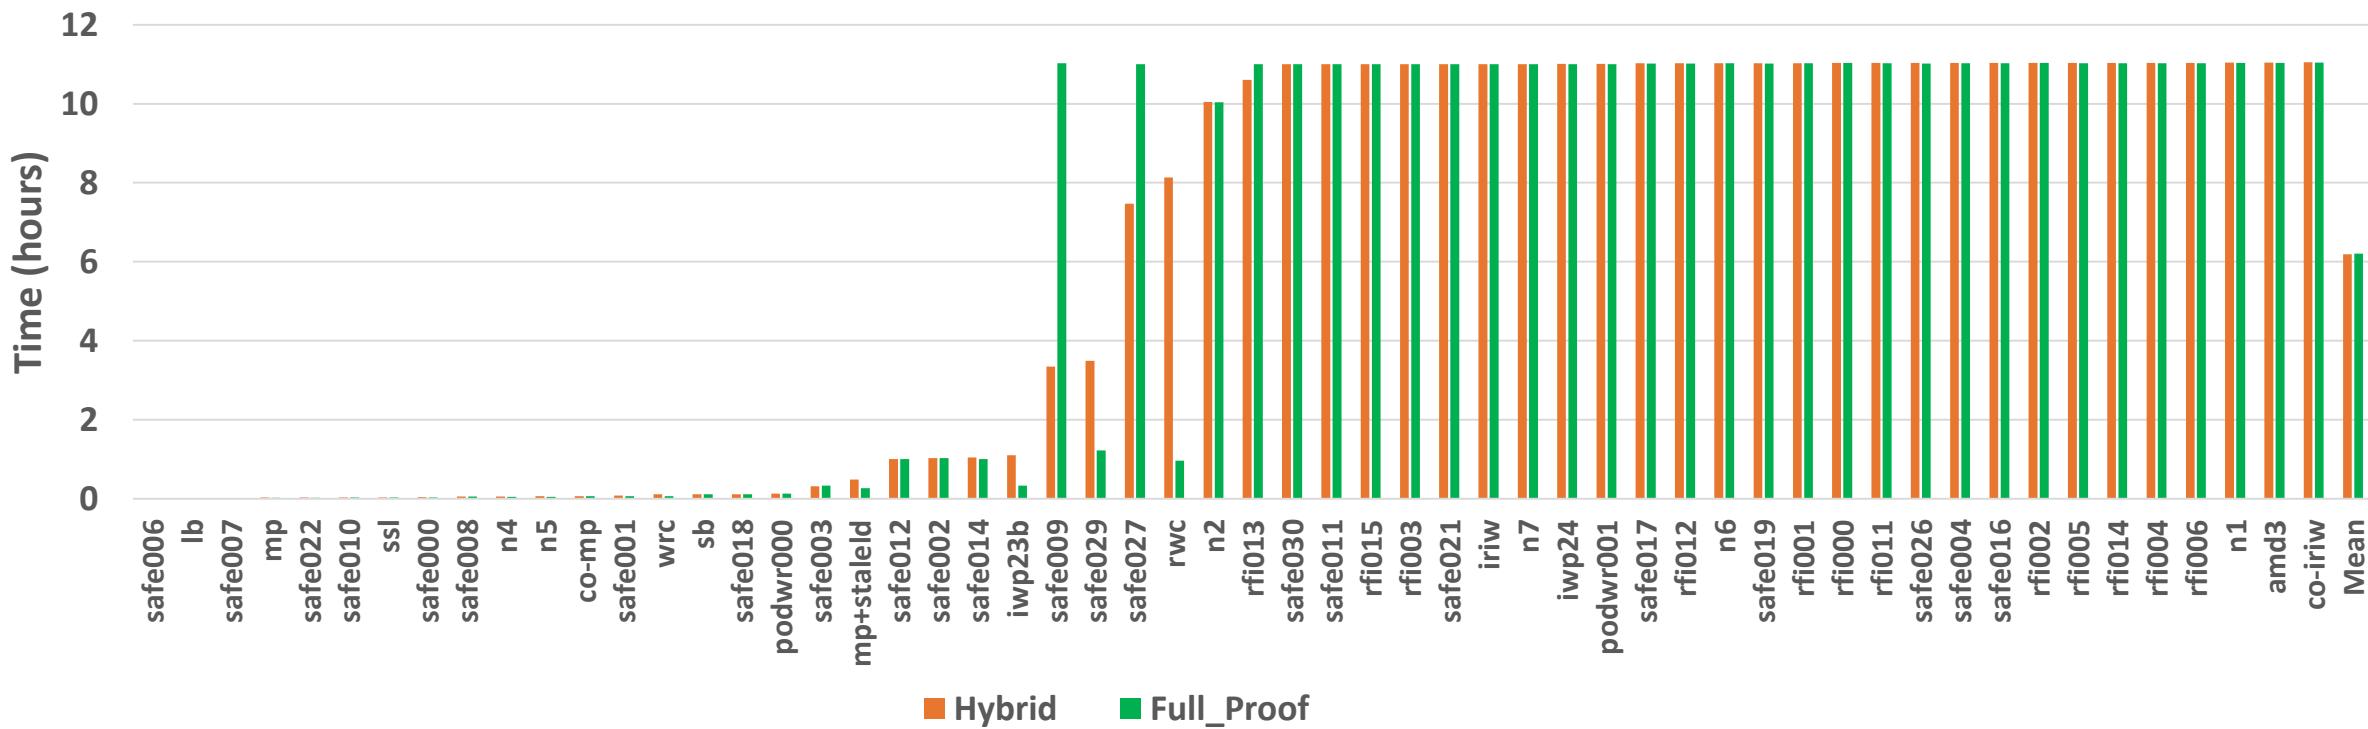
Covering trace must also obey other assumptions, including load val assumptions (For mp, Ld y = 1 and Ld x = 0)

Thus, covering trace for mp final val assumption (full execution with Ld y=1 and Ld x=0) is equivalent to finding forbidden execution of mp!



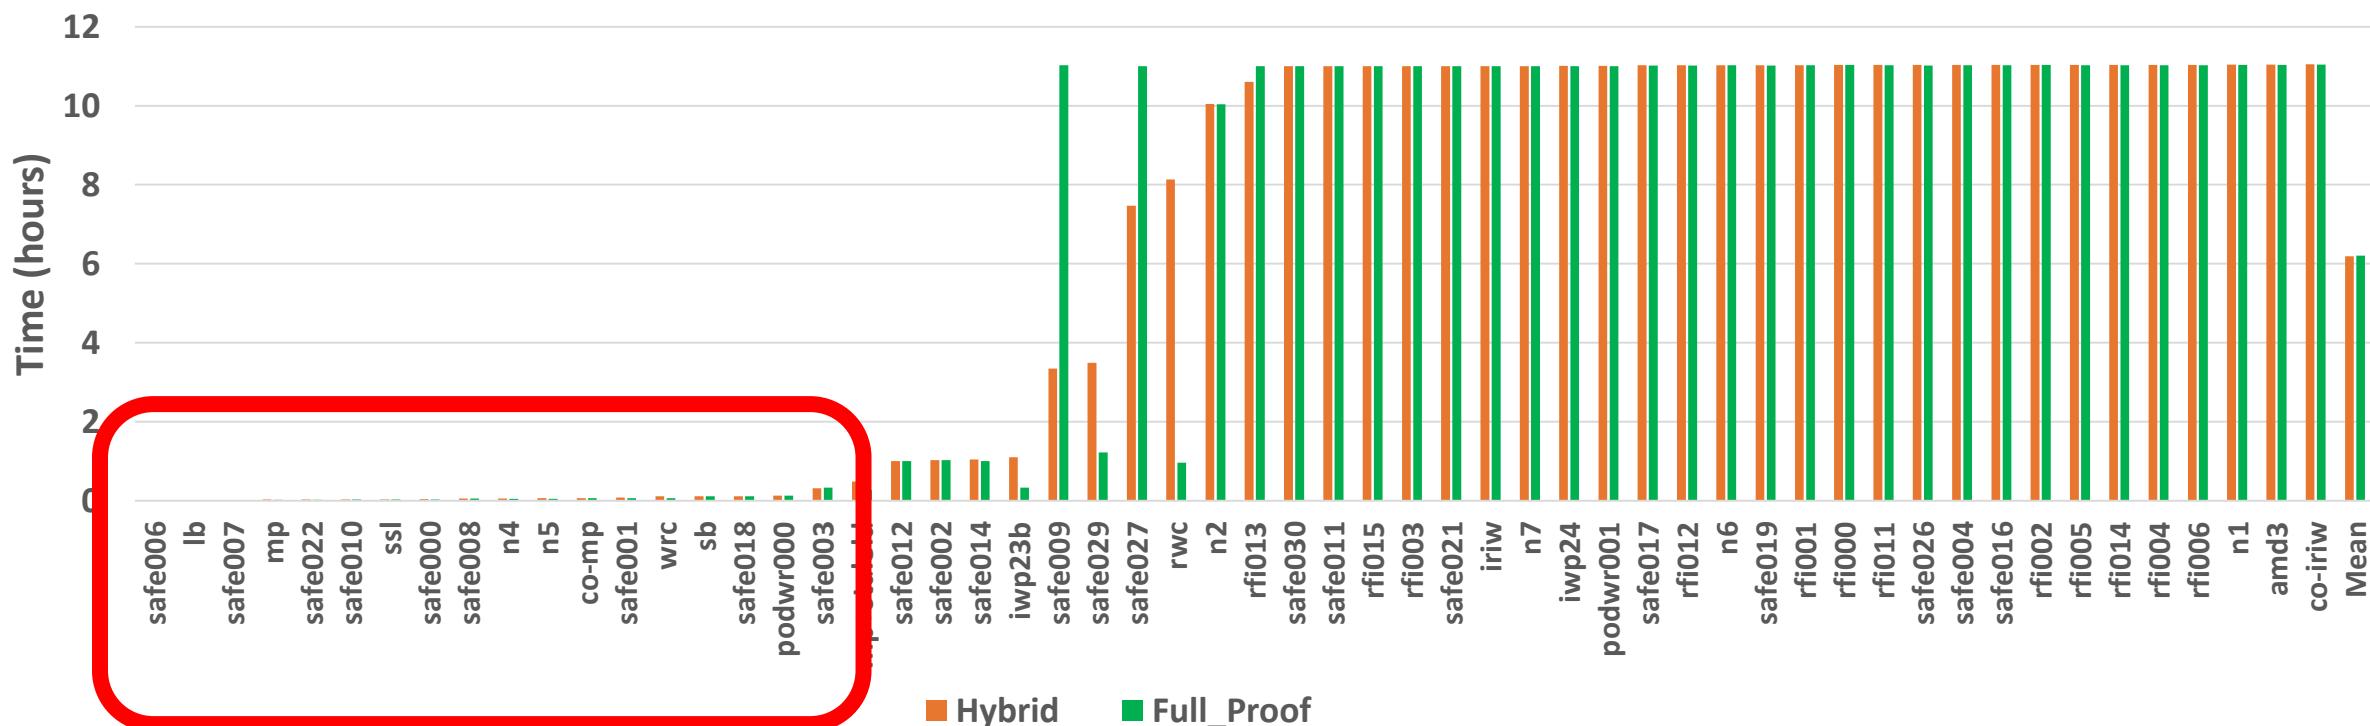
# Results: Time to Prove Properties

- Two configurations (**Hybrid** and **Full\_Proof**), avg. runtime 6.2 hrs
  - See paper for configuration details



# Results: Time to Prove Properties

- Two configurations (**Hybrid** and **Full\_Proof**), avg. runtime 6.2 hrs
  - See paper for configuration details

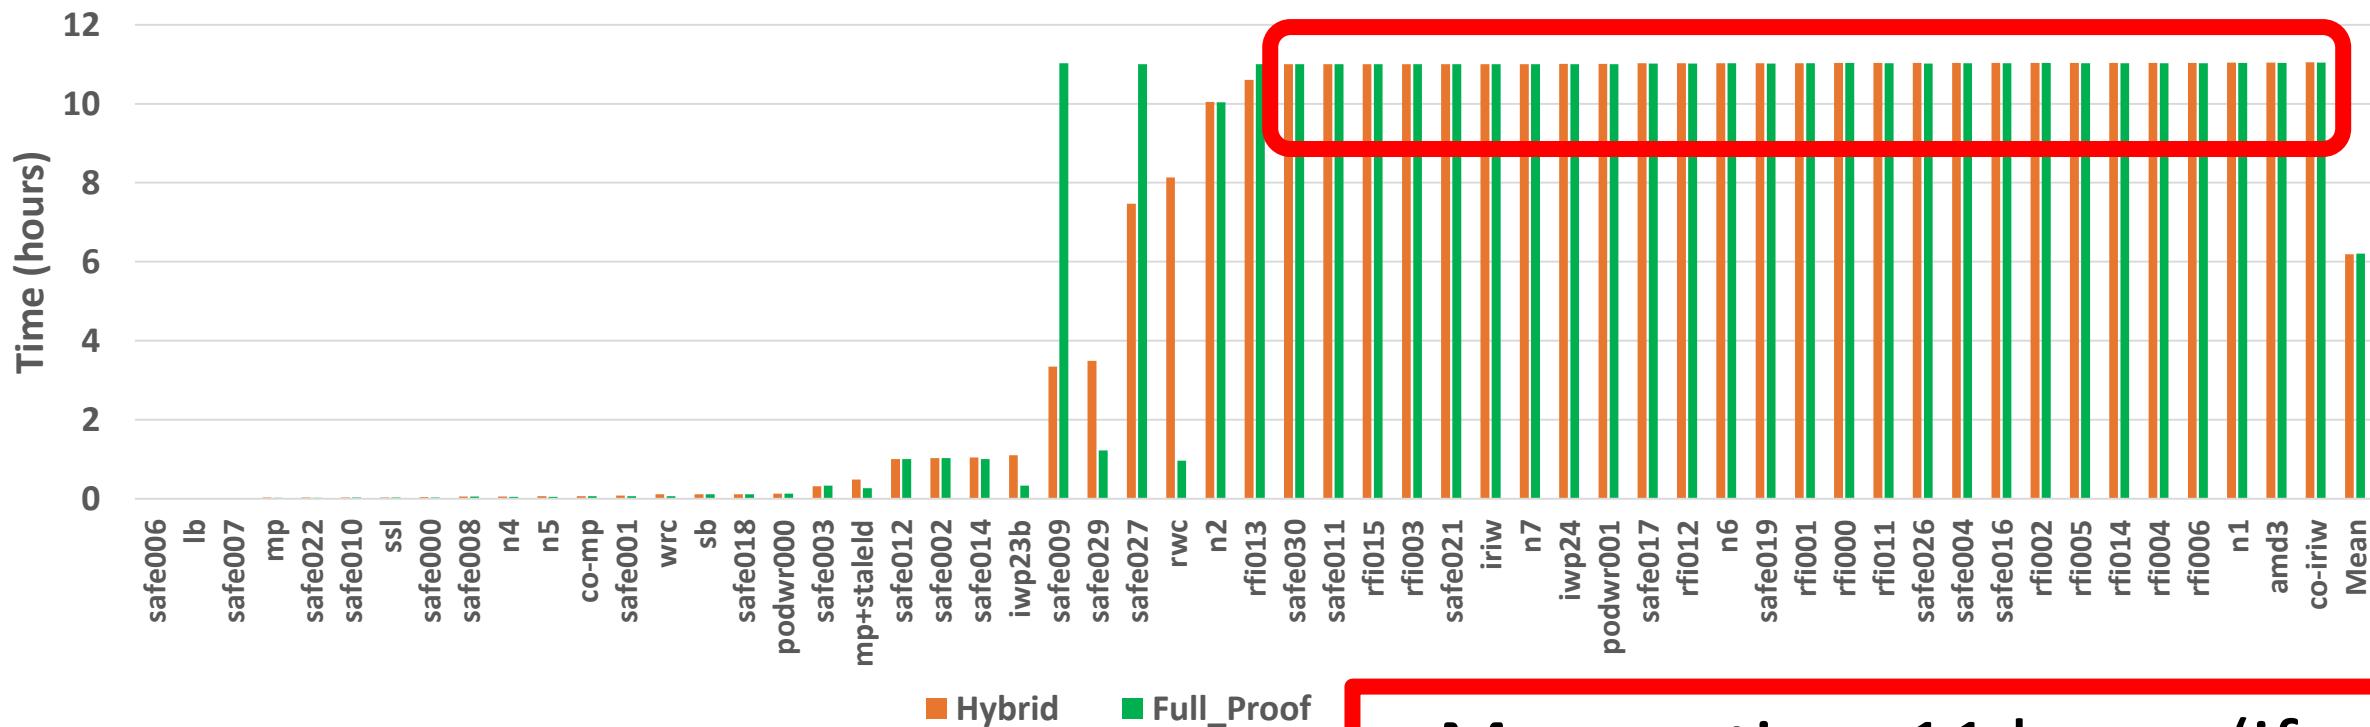


Complete quickly due to  
covering traces



# Results: Time to Prove Properties

- Two configurations (**Hybrid** and **Full\_Proof**), avg. runtime 6.2 hrs
  - See paper for configuration details

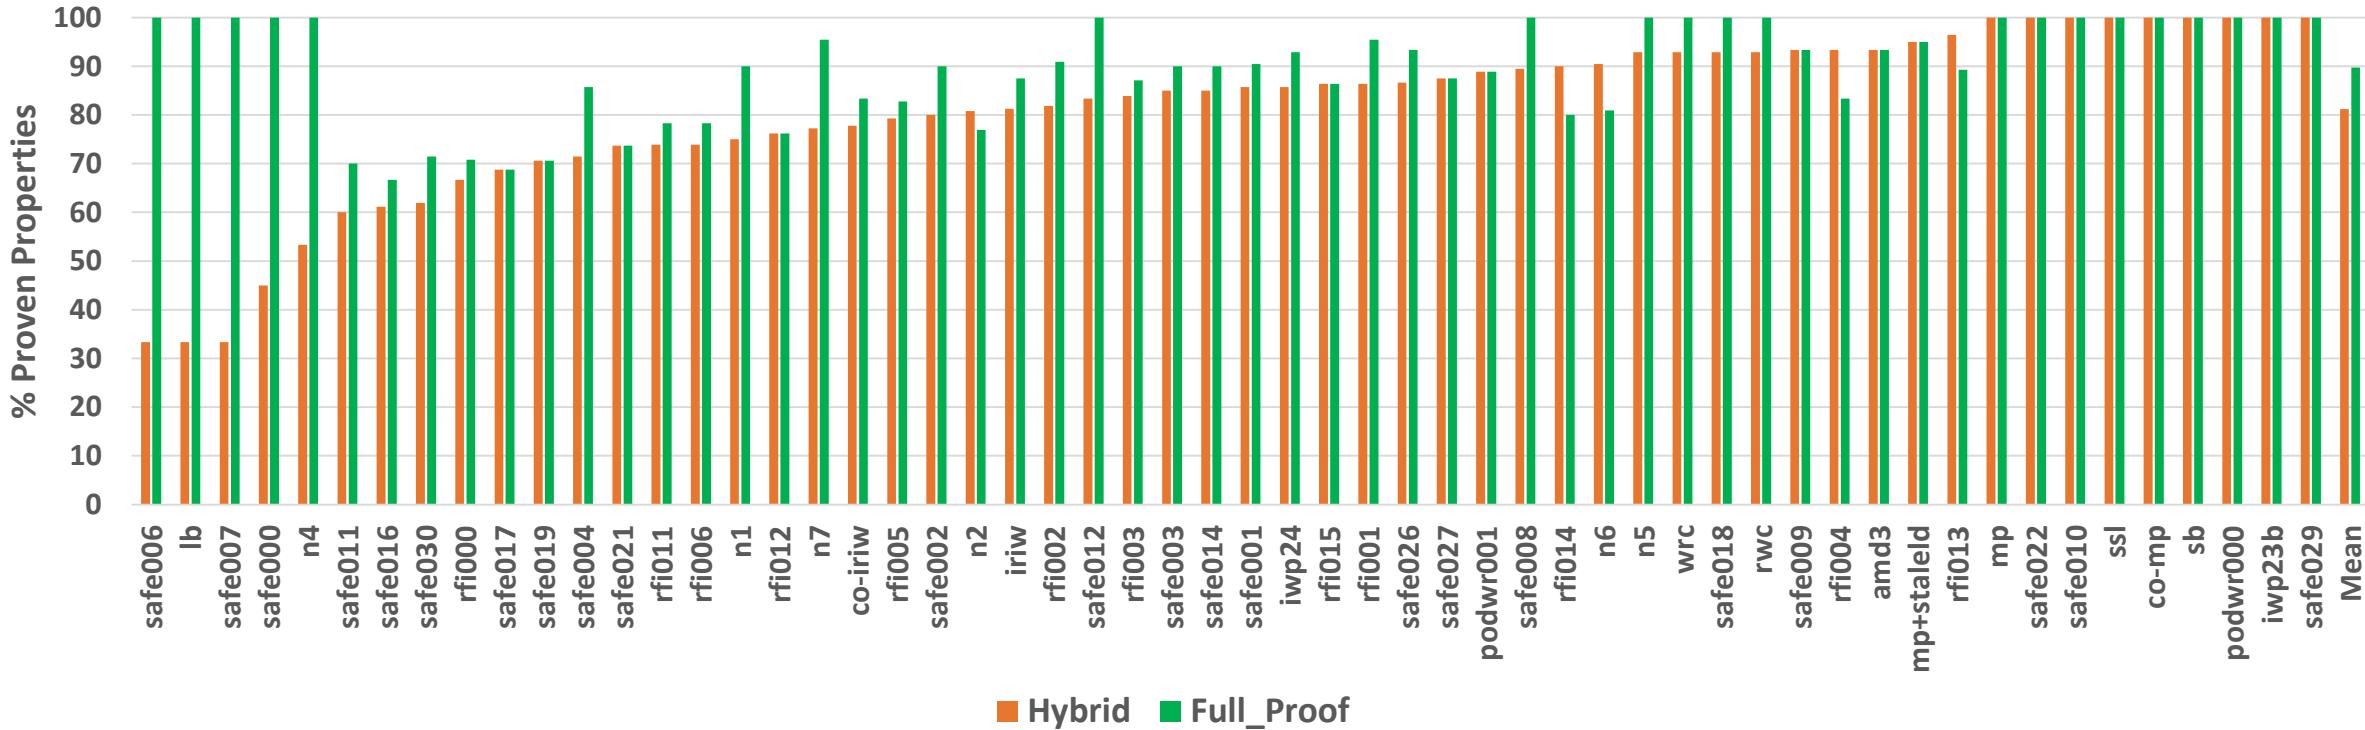


Max runtime 11 hours (if some properties unproven)



# Results: Proven Properties

- **Full\_Proof** generally better (90%/test) than **Hybrid** (81%/test)
- On average, **Full\_Proof** can prove more properties in same time



# Results: Proven Properties

- **Full\_Proof** generally better (90%/test) than **Hybrid** (81%/test)
- On average, **Full\_Proof** can prove more properties in same time

**Hybrid better for only a few tests**

