

Automated Formal Memory Consistency Verification of Hardware

Yatin A. Manerkar

Princeton University

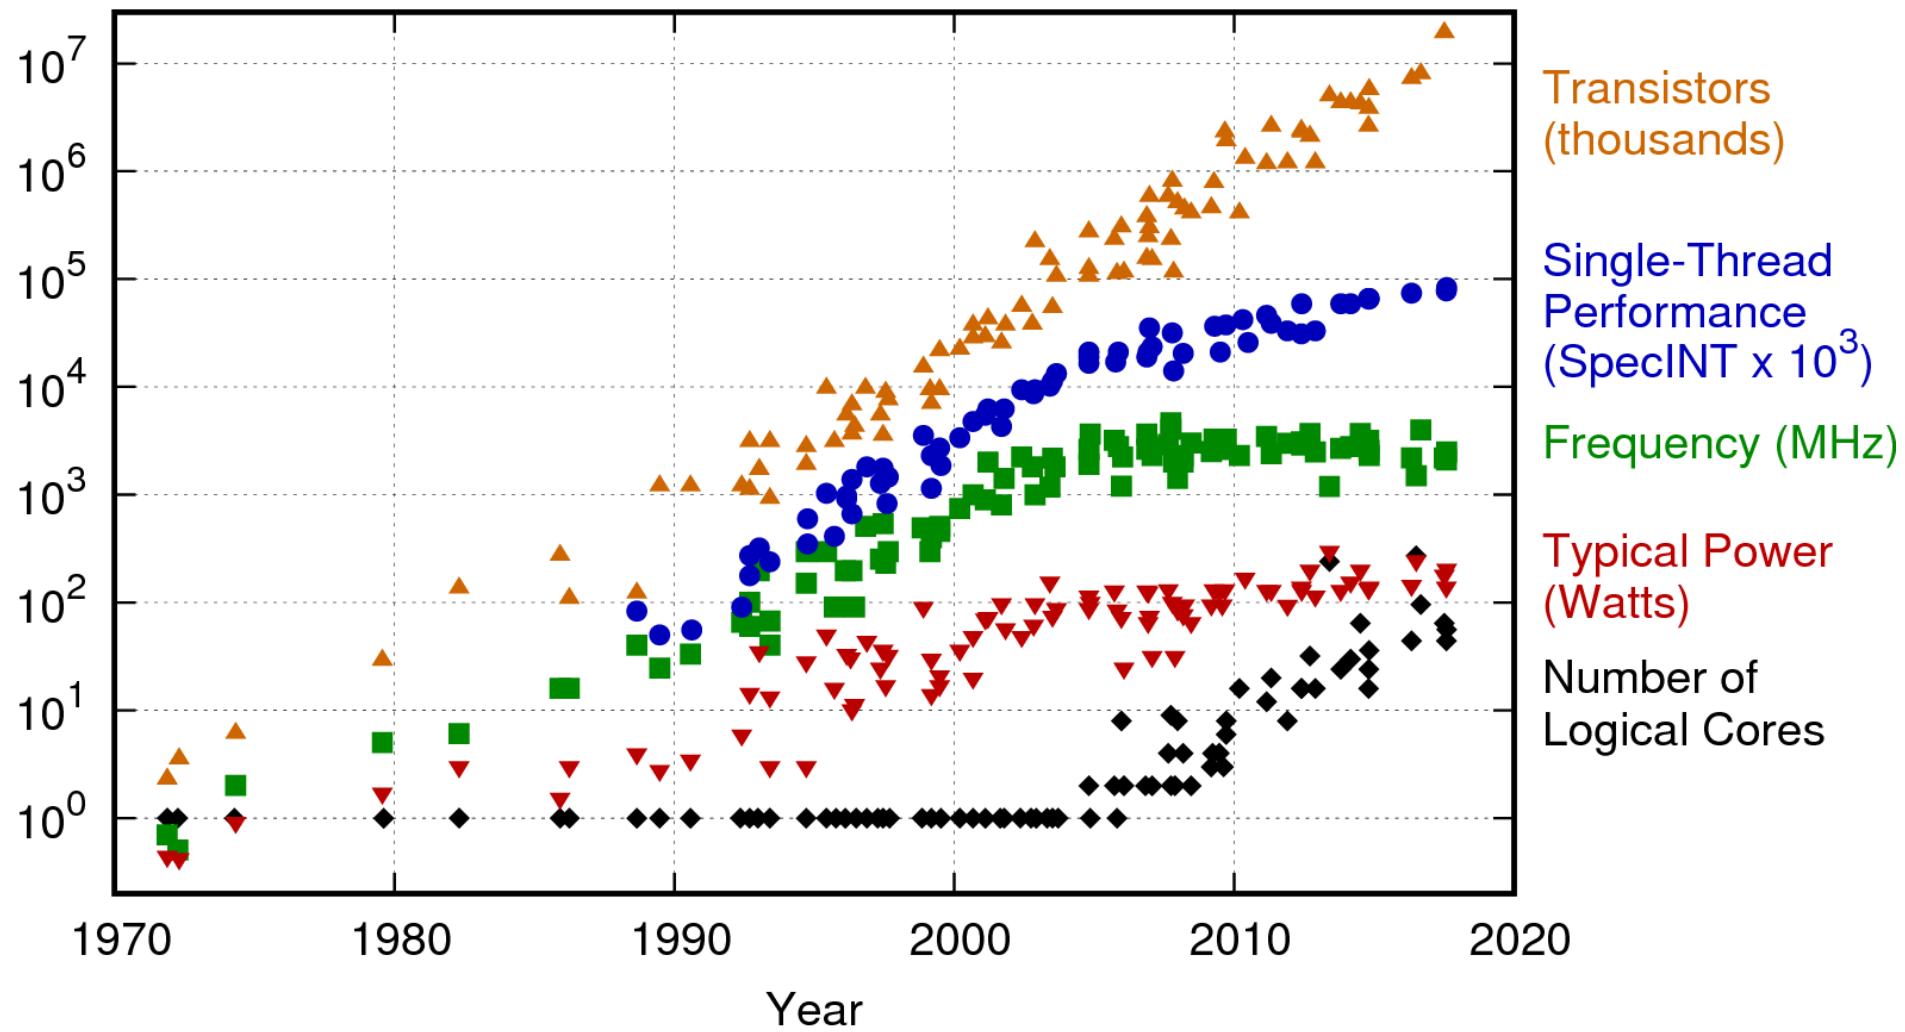
June 23rd, 2019

<http://www.cs.princeton.edu/~manerkar>



The Rise of Parallelism...

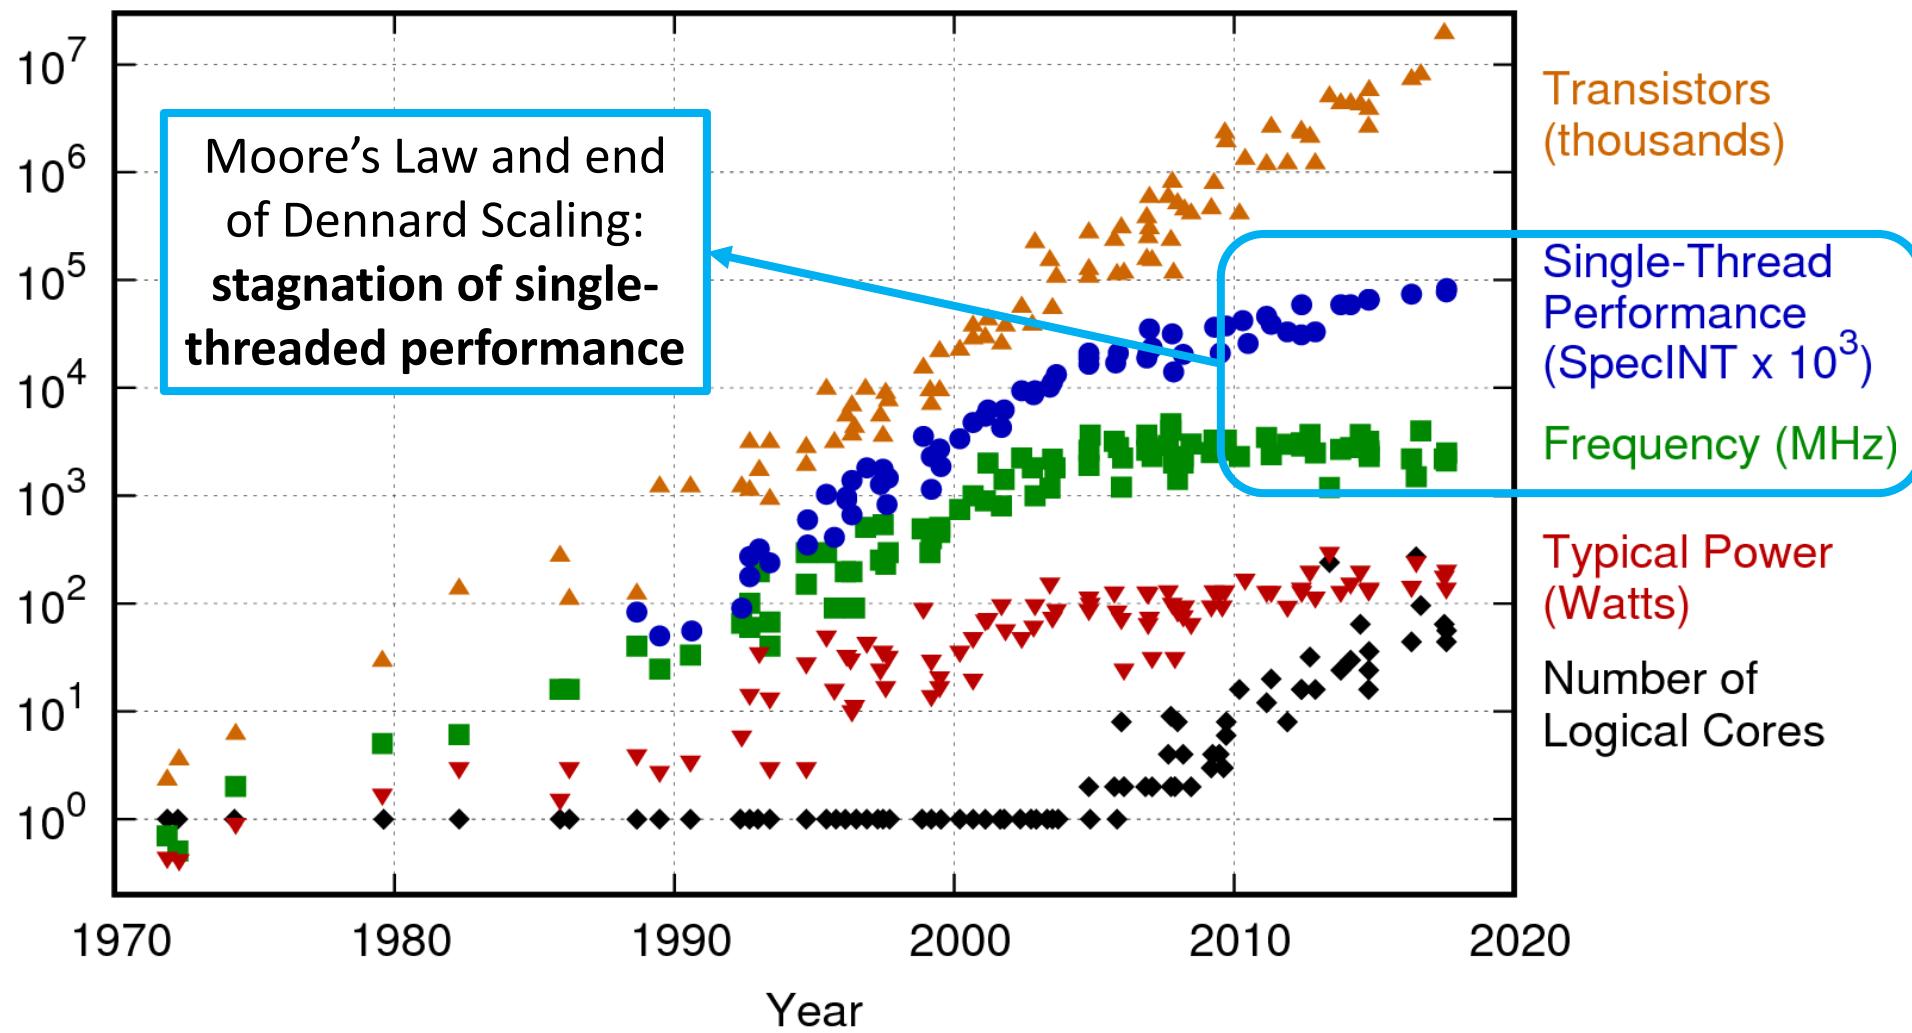
42 Years of Microprocessor Trend Data



Original data up to the year 2010 collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten
New plot and data collected for 2010-2017 by K. Rupp

The Rise of Parallelism...

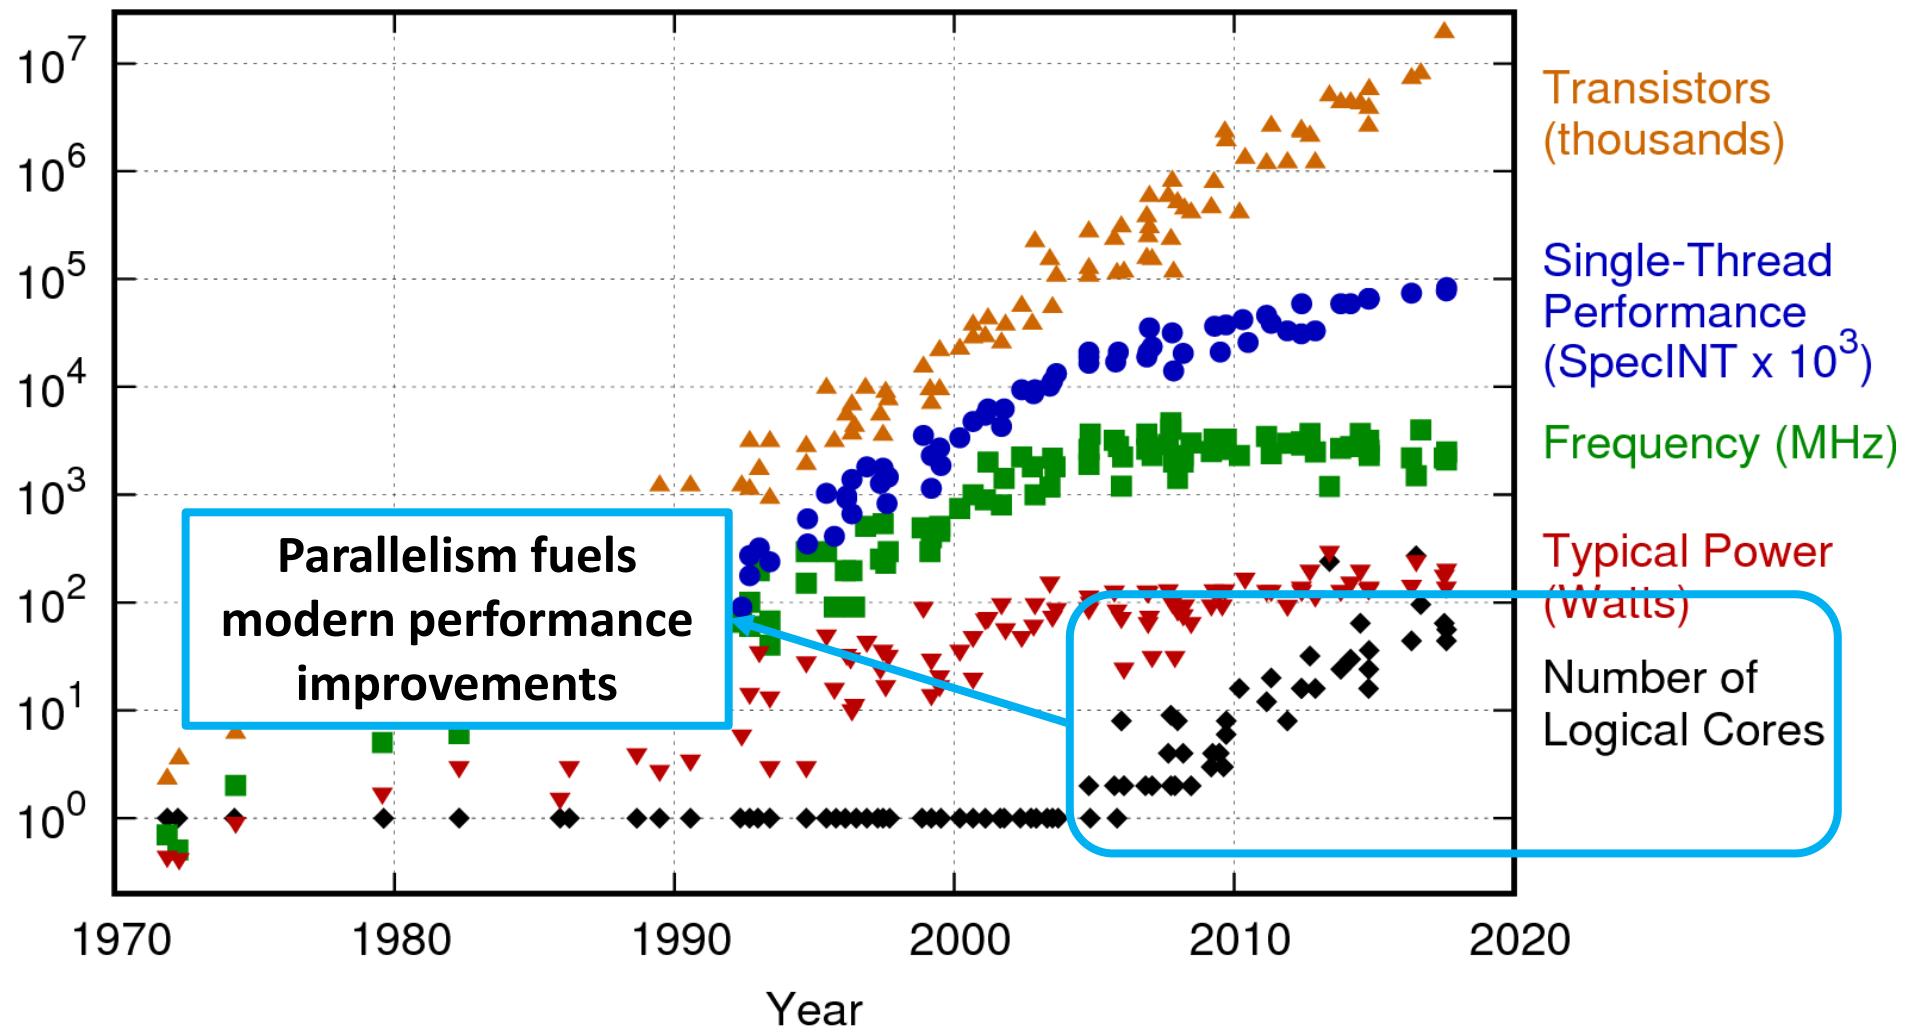
42 Years of Microprocessor Trend Data



Original data up to the year 2010 collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten
New plot and data collected for 2010-2017 by K. Rupp

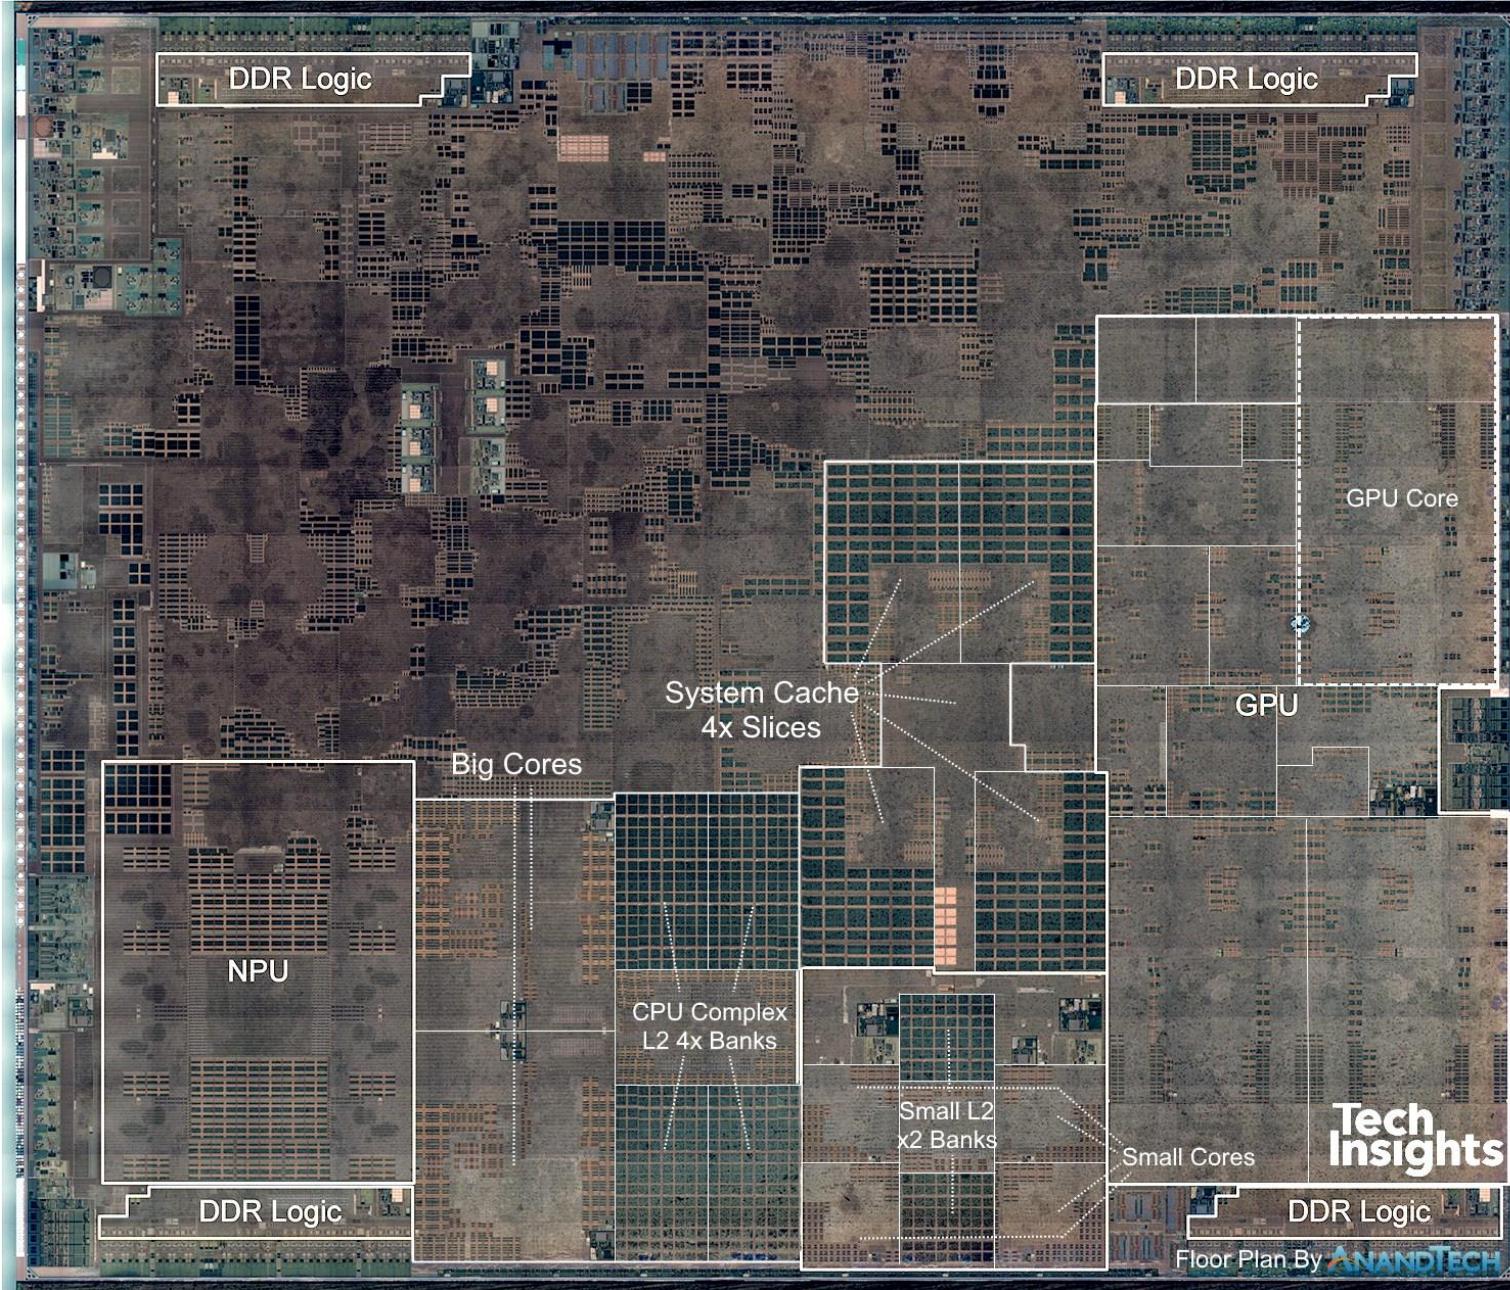
The Rise of Parallelism...

42 Years of Microprocessor Trend Data

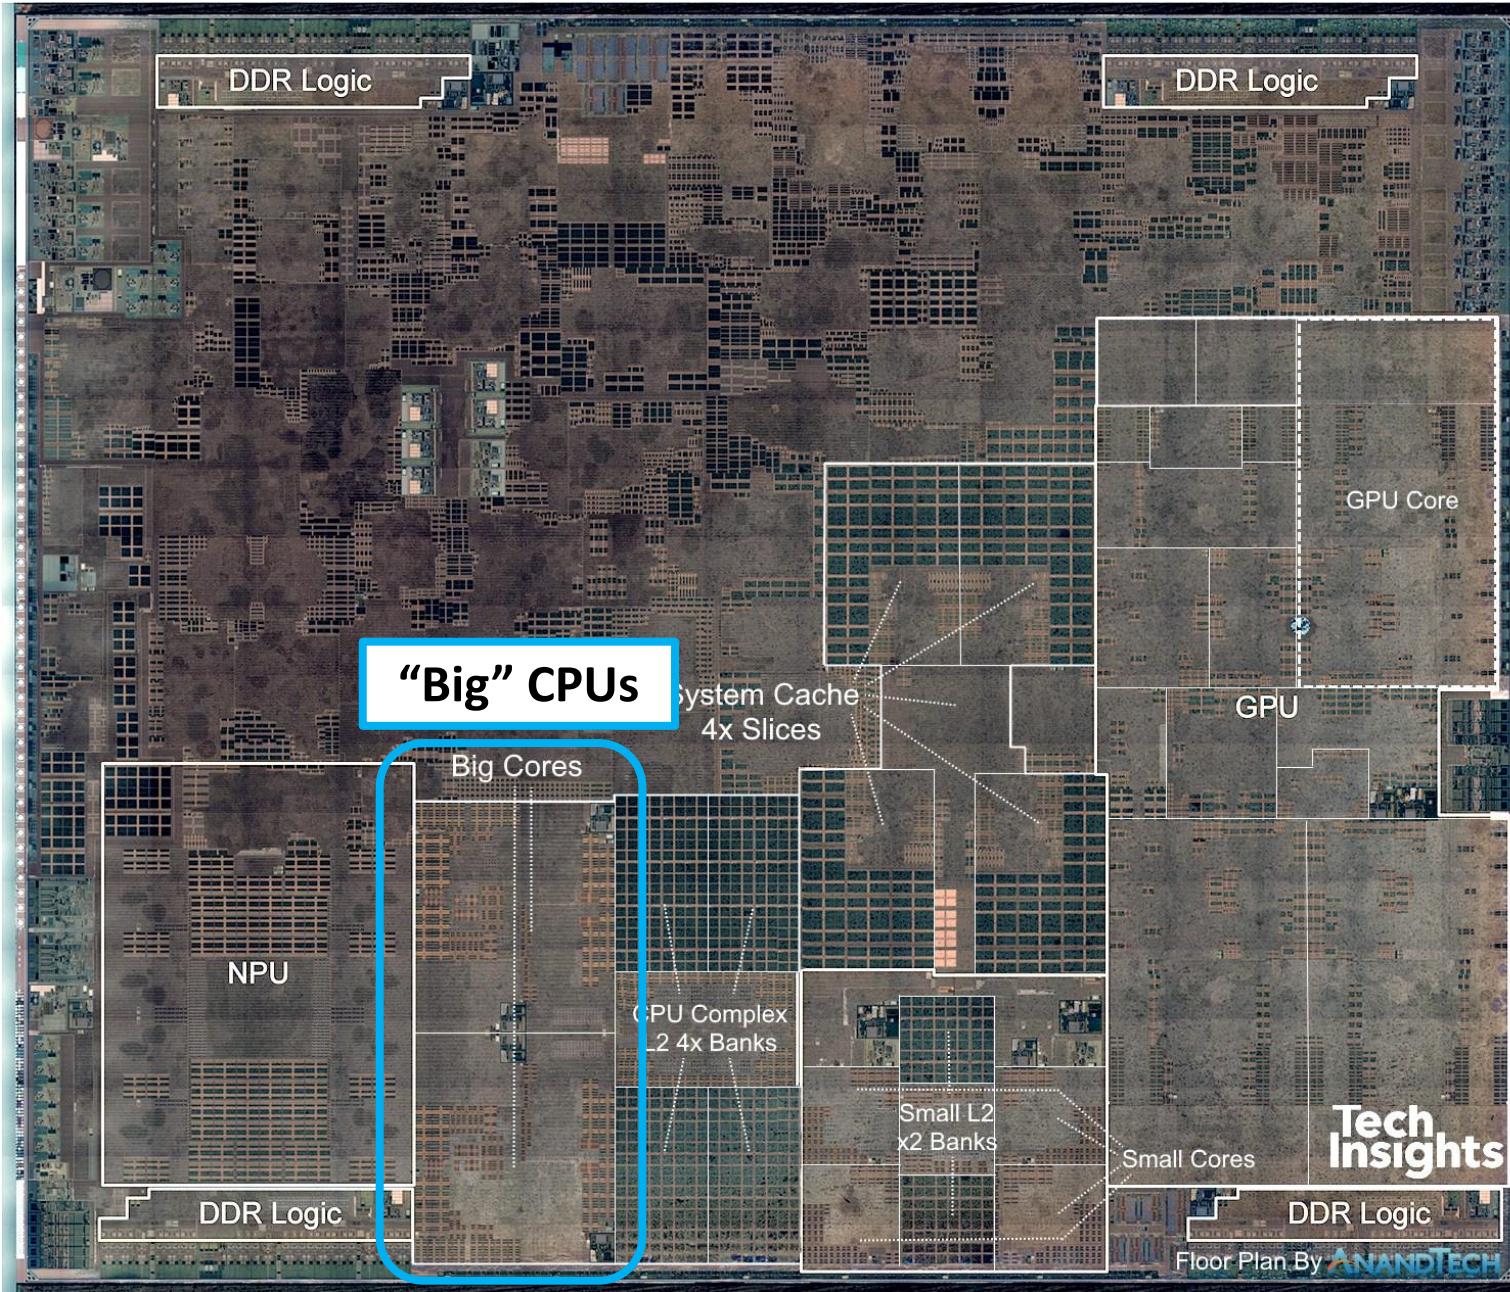


Original data up to the year 2010 collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten
New plot and data collected for 2010-2017 by K. Rupp

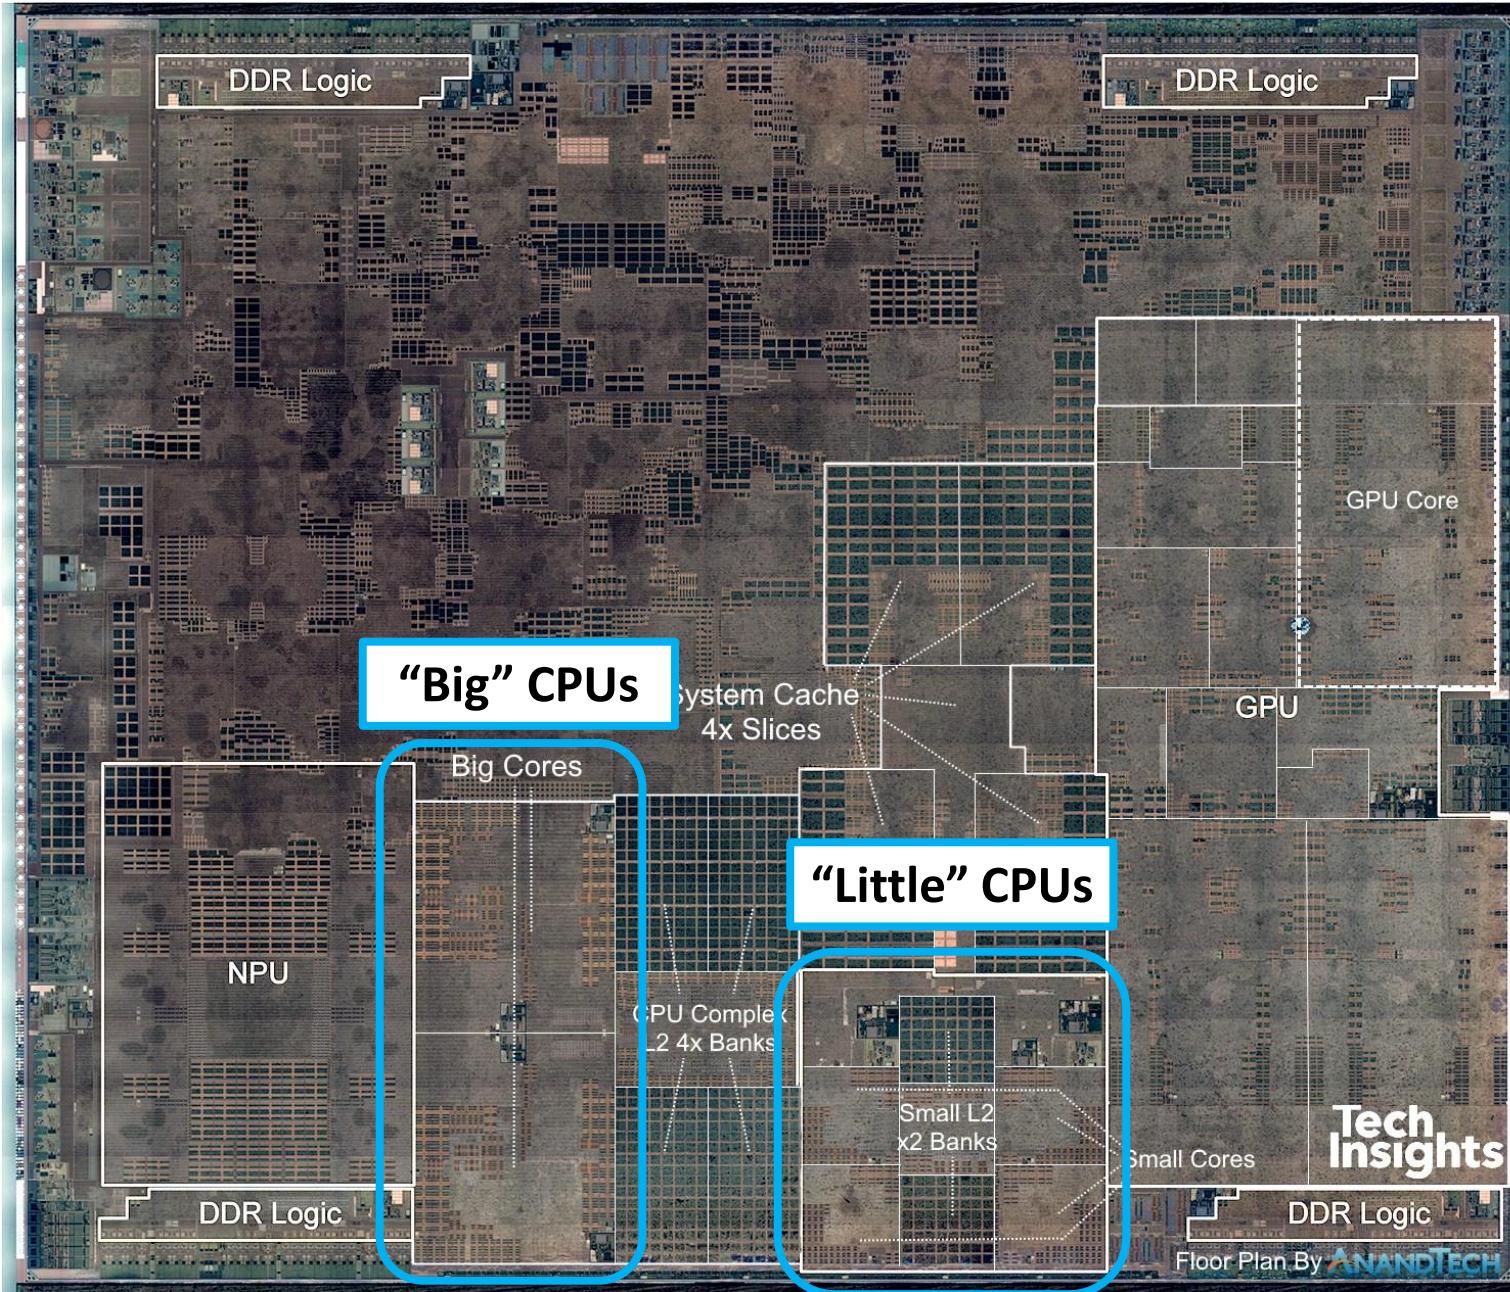
...and Heterogeneity (Example: Apple A12)



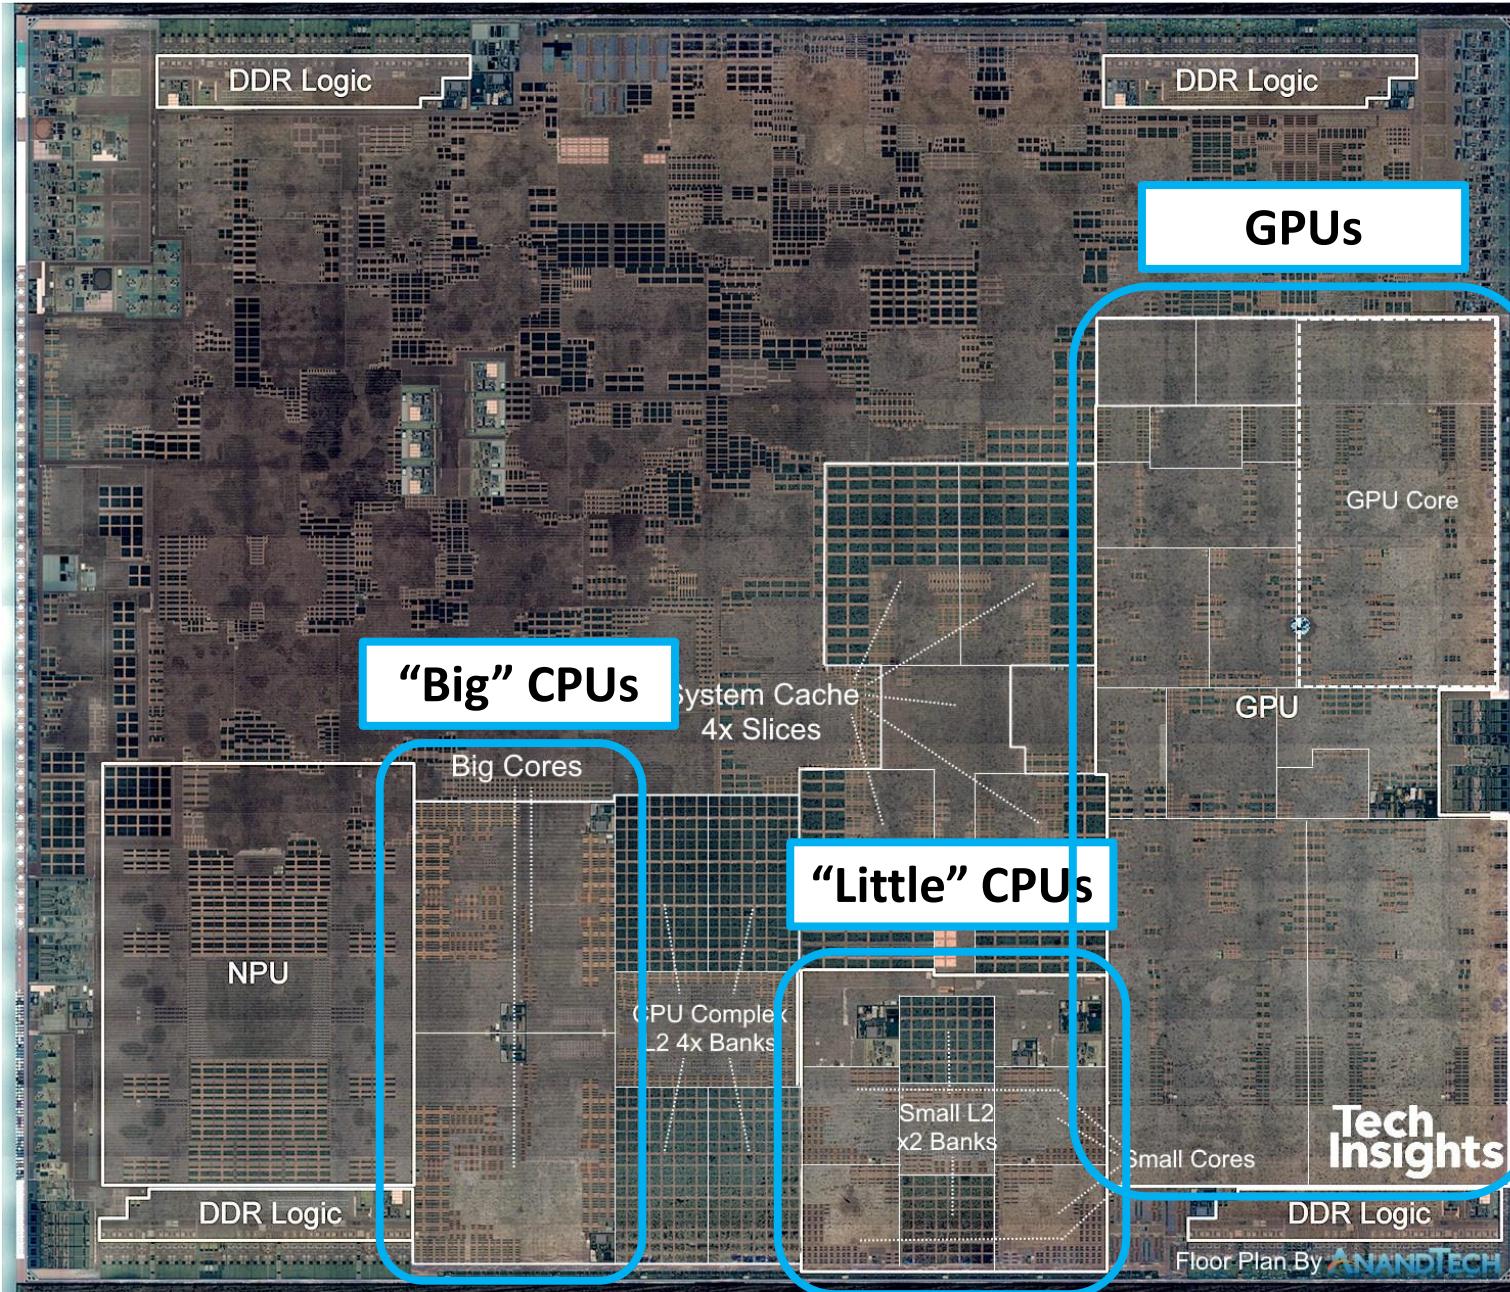
...and Heterogeneity (Example: Apple A12)



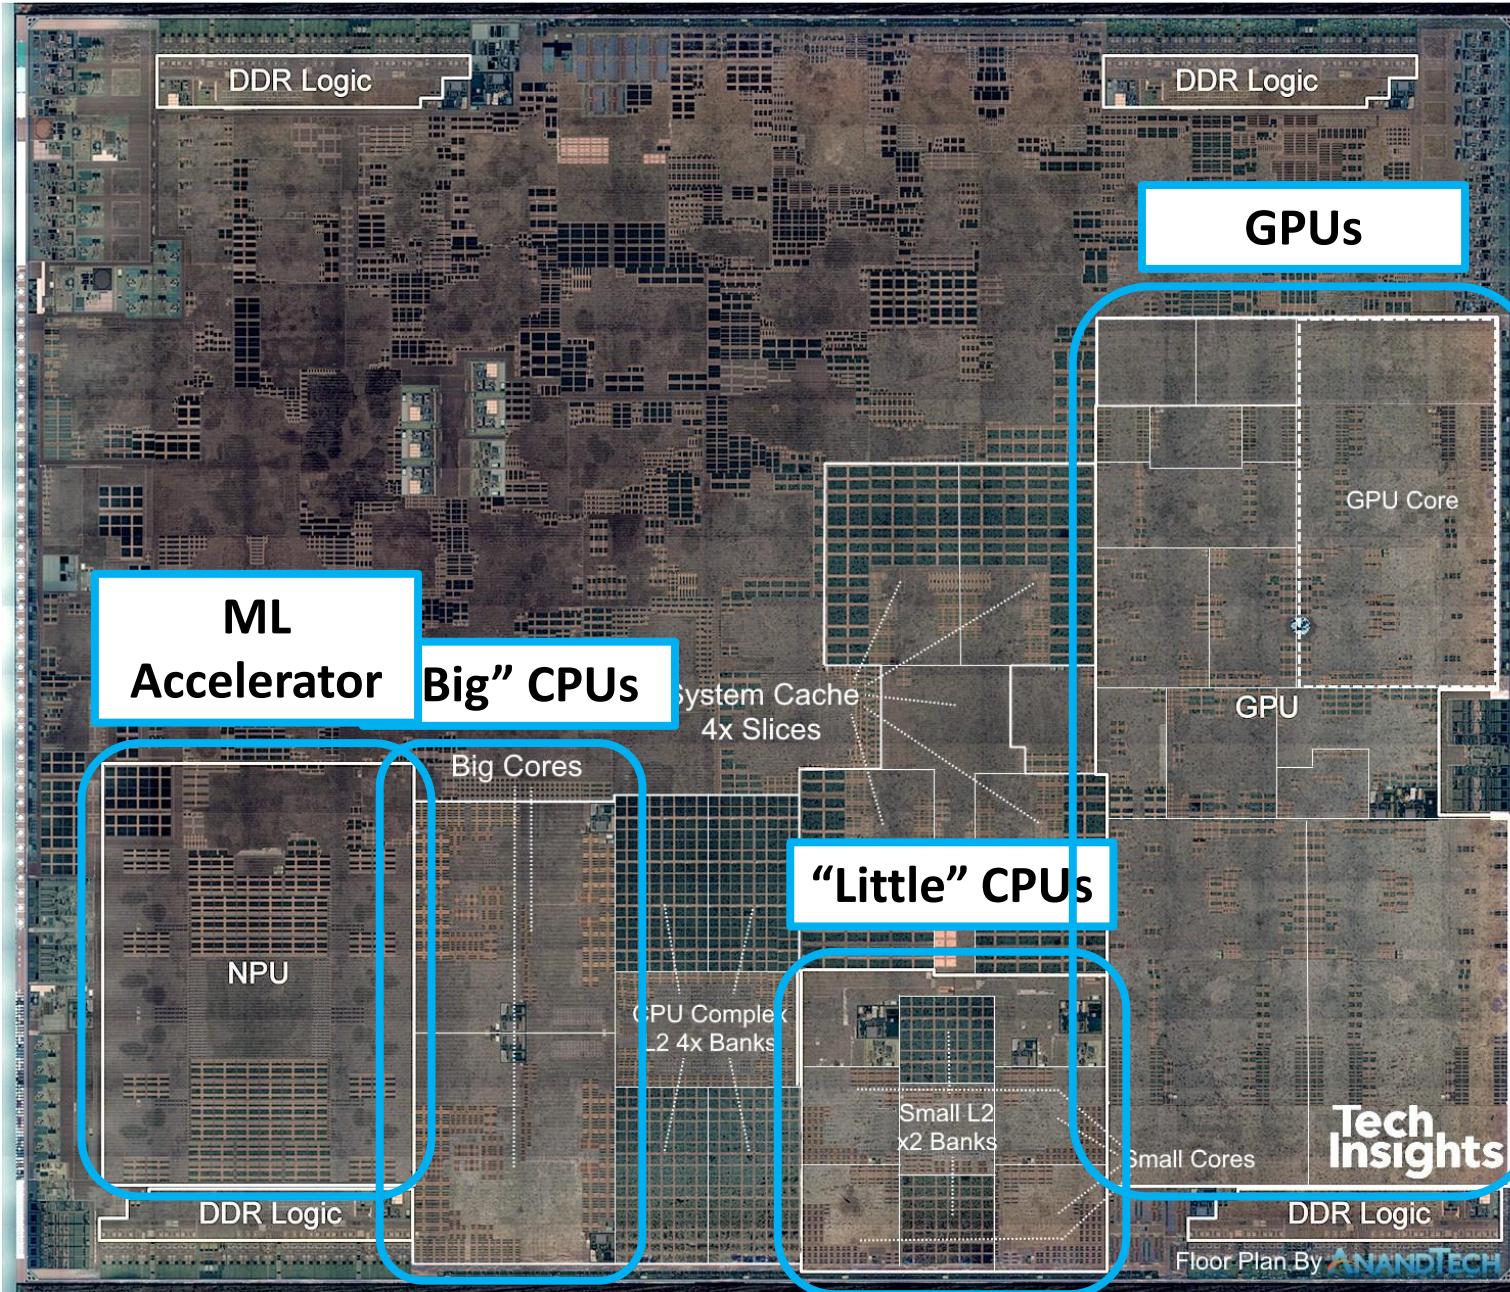
...and Heterogeneity (Example: Apple A12)



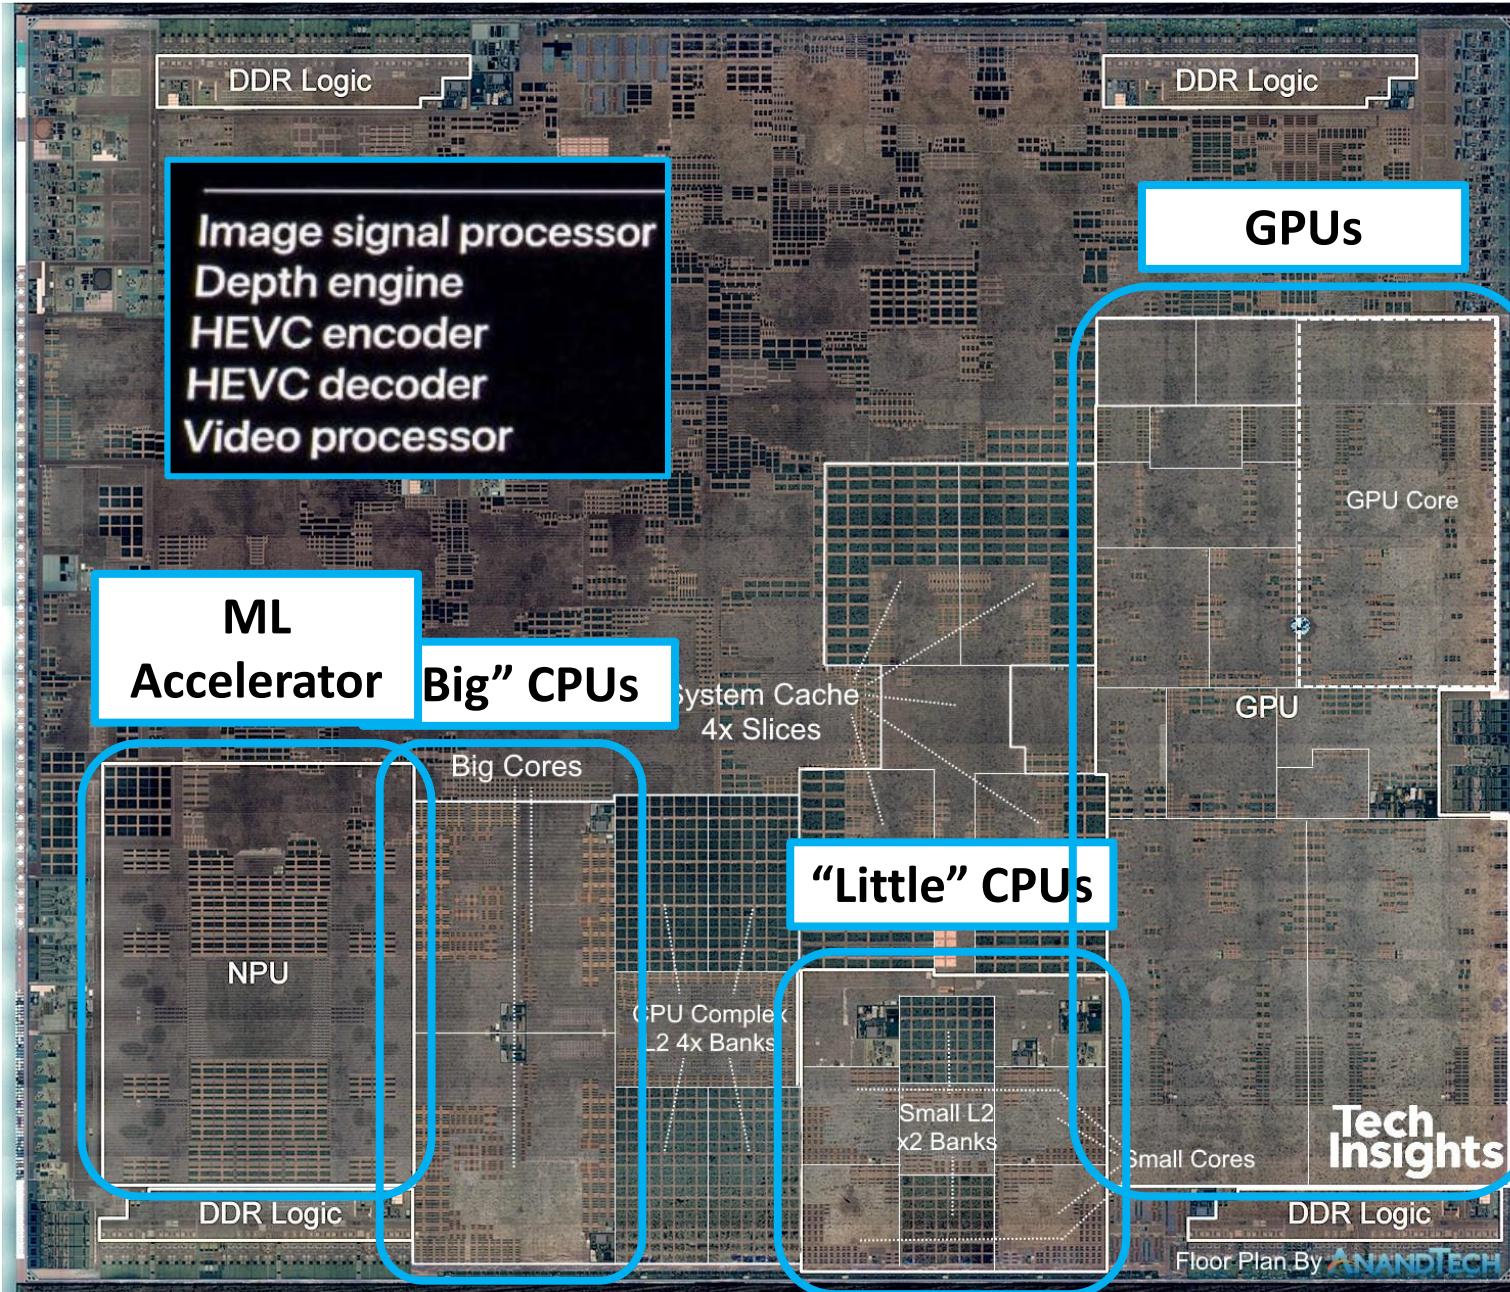
...and Heterogeneity (Example: Apple A12)



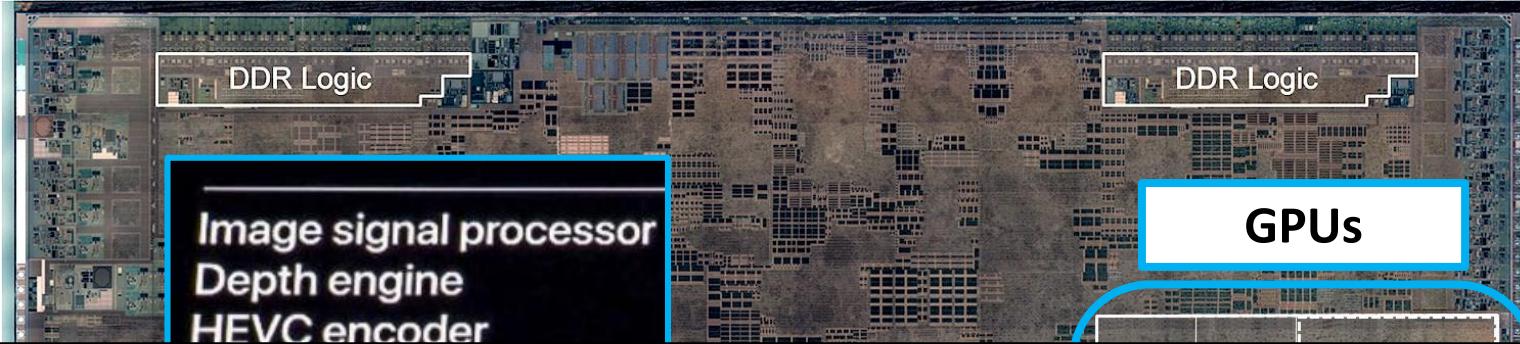
...and Heterogeneity (Example: Apple A12)



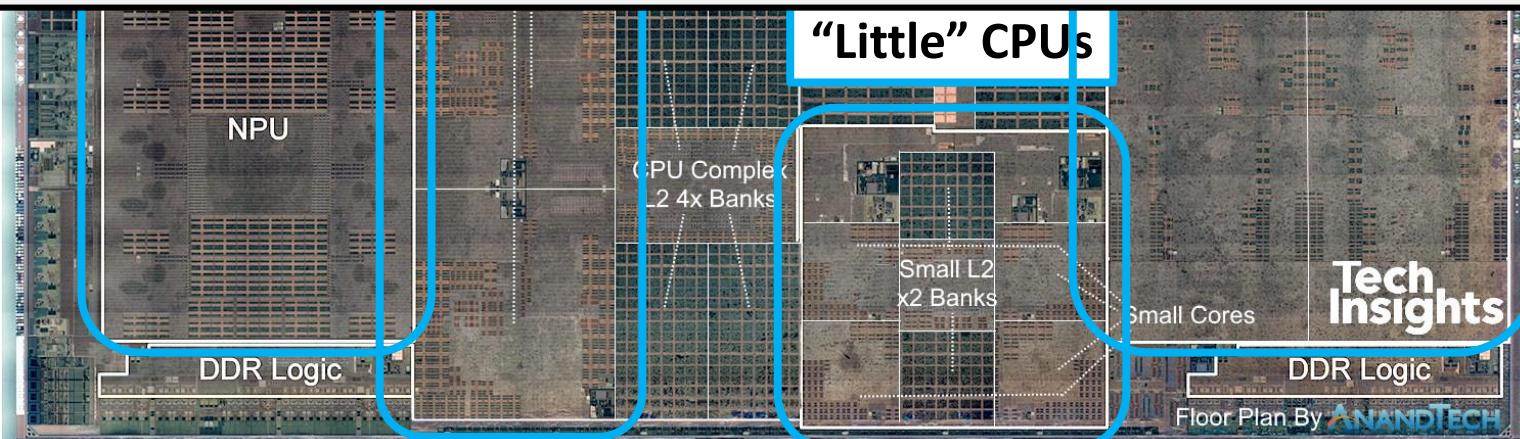
...and Heterogeneity (Example: Apple A12)



...and Heterogeneity (Example: Apple A12)



Parallel processors are hard to get right!
How can we formally verify parallel hardware?



Building a Formally Verified Processor

Formal Methods Expert



- Build proven-correct processor (e.g. Kami) or...

Building a Formally Verified Processor

Formal Methods Expert



REMS

- Build proven-correct processor (e.g. Kami) or...
- ...construct formal model of implementation and verify that (REMS)

Building a Formally Verified Processor

Formal Methods Expert



REMS

- Build proven-correct processor (e.g. Kami) or...
- ...construct formal model of implementation and verify that (REMS)
- **Formal methods expert carries most of the verification burden**

Building a Formally Verified Processor

Formal Methods Expert



REMS

- Build proven-correct processor (e.g. Kami) or...
- ...construct formal model of implementation and verify that (REMS)
- **Formal methods expert carries most of the verification burden**

Computer Architect



OpenPiton

- Experts on building processors
- Generally not much formal methods expertise
- **Can they share more of the verification burden?**

Building a Formally Verified Processor

Formal Methods Expert



Computer Architect



My work: Automated tools that enable engineers to formally verify their systems by themselves!

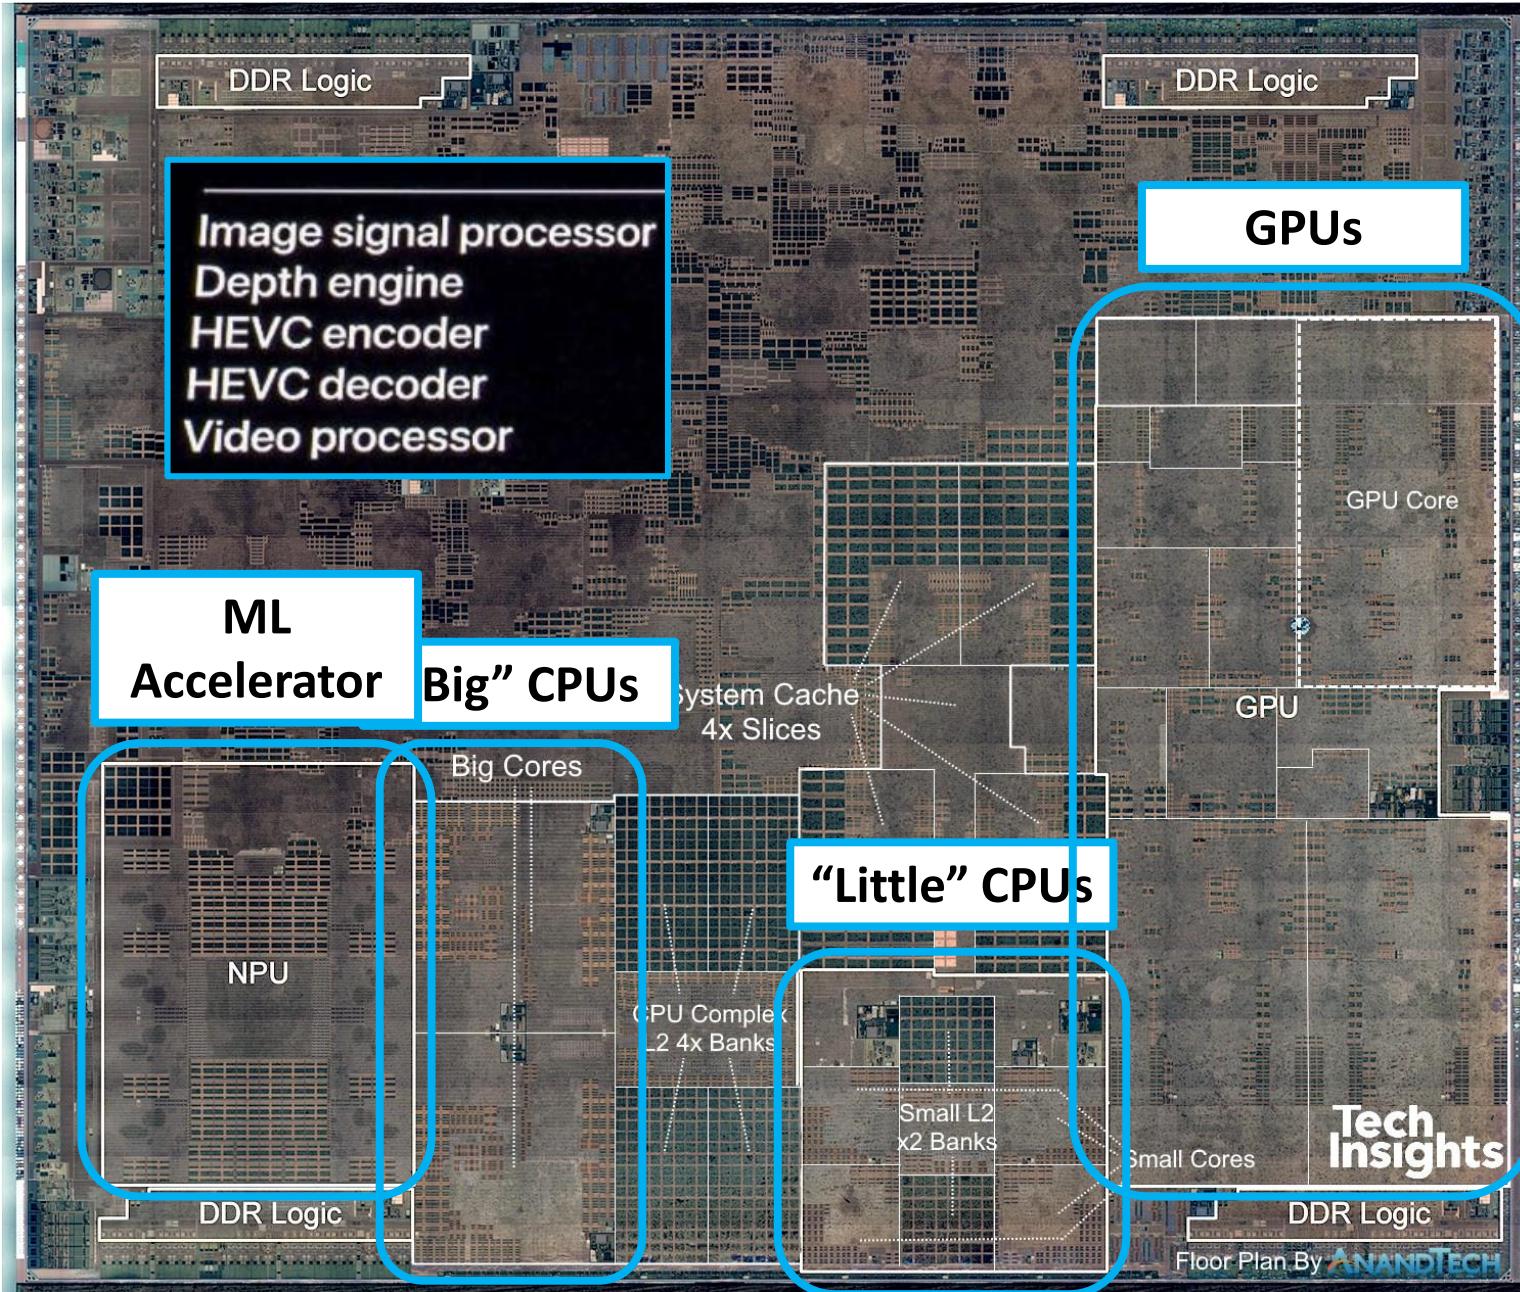
Case Study: Memory Consistency Verification

- Formal methods expert carries most of the verification burden
- Can they share more of the verification burden?

Talk Outline

- Overview
- Memory Consistency Background
- **PipeProof:** All-Program Microarchitectural MCM Verification
- **RTLCheck:** MCM Verification of Verilog RTL
- Expanding to other domains
- Conclusion

Processors Communicate via Shared Memory



What does this program print?

Thread 0	Thread 1
1 x = 1;	3 if (y == 1) print("Answer is:");
2 y = 1;	4 if (x == 1) print("42");



What does this program print?

Thread 0	Thread 1
① <code>x = 1;</code>	③ <code>if (y == 1)</code> <code>print("Answer is:");</code>
② <code>y = 1;</code>	④ <code>if (x == 1)</code> <code>print("42");</code>

Can it print “Answer is: 42”? **Yes**, eg:

① ② ③ ④



What does this program print?

Thread 0	Thread 1
1 x = 1;	3 if (y == 1) print("Answer is:");
2 y = 1;	4 if (x == 1) print("42");

Can it print “Answer is: 42”? **Yes**, eg:

1 2 3 4

How about just “42”? **Yes**, eg:

1 3 4 2



What does this program print?

Thread 0	Thread 1
① <code>x = 1;</code>	③ <code>if (y == 1) print("Answer is:");</code>
② <code>y = 1;</code>	④ <code>if (x == 1) print("42");</code>

Can it print “Answer is: 42”? **Yes**, eg:

1 2 3 4

How about just “42”? **Yes**, eg:

1 3 4 2

Could it print **nothing**? **Yes**, eg:

3 4 1 2



What does this program print?

Thread 0	Thread 1
① $x = 1;$	③ if ($y == 1$) print("Answer is:");
② $y = 1;$	④ if ($x == 1$) print("42");

Can it print “Answer is: 42”? **Yes**, eg:

1 2 3 4

How about just “42”? **Yes**, eg:

1 3 4 2

Could it print **nothing**? **Yes**, eg:

3 4 1 2

These executions obey **Sequential Consistency (SC)** [Lamport79], which requires that the results of the overall program correspond to some in-order interleaving of the statements from each individual thread.



What does this program print?

Thread 0	Thread 1
1 x = 1;	3 if (y == 1) print("Answer is:");
2 y = 1;	4 if (x == 1) print("42");

How about “Answer is:”?

2 1 3 4



What does this program print?

Thread 0	Thread 1
1 x = 1;	3 if (y == 1) print("Answer is:");
2 y = 1;	4 if (x == 1) print("42");

How about “Answer is:”?

It depends!

2 1 3 4



What does this program print?

Thread 0	Thread 1
1 x = 1;	3 if (y == 1) print("Answer is:");
2 y = 1;	4 if (x == 1) print("42");

How about “Answer is:”?

It depends!

2 1 3 4



NO!



What does this program print?

Thread 0	Thread 1
1 x = 1;	3 if (y == 1) print("Answer is:");
2 y = 1;	4 if (x == 1) print("42");

How about “Answer is:”?

It depends!



NO!



YES!

2 1 3 4
arm



What does this program print?

Thread 0

① x = 1;

Thread 1

③ if (y == 1)
print("Answer is:");

Most processors today implement “weak memory models” that relax orderings required by SC!



NO!

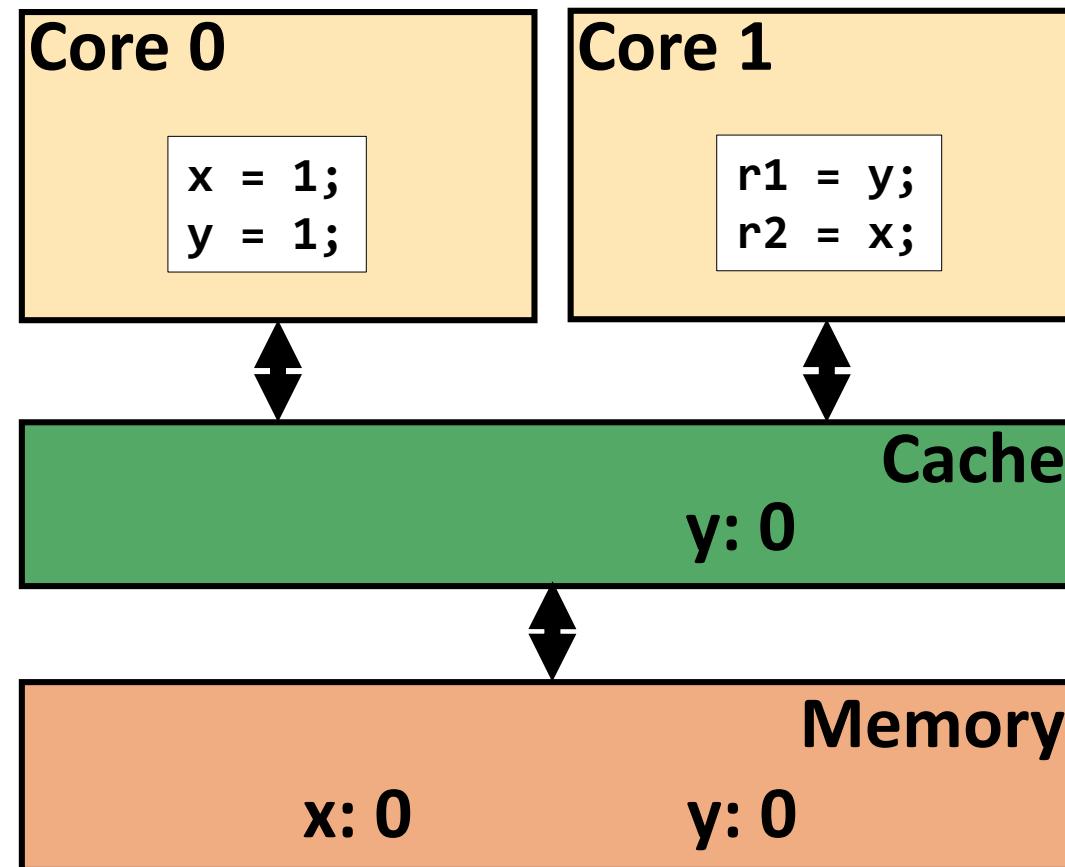
YES!



arm

Why reorder memory operations?

Answer: Performance!



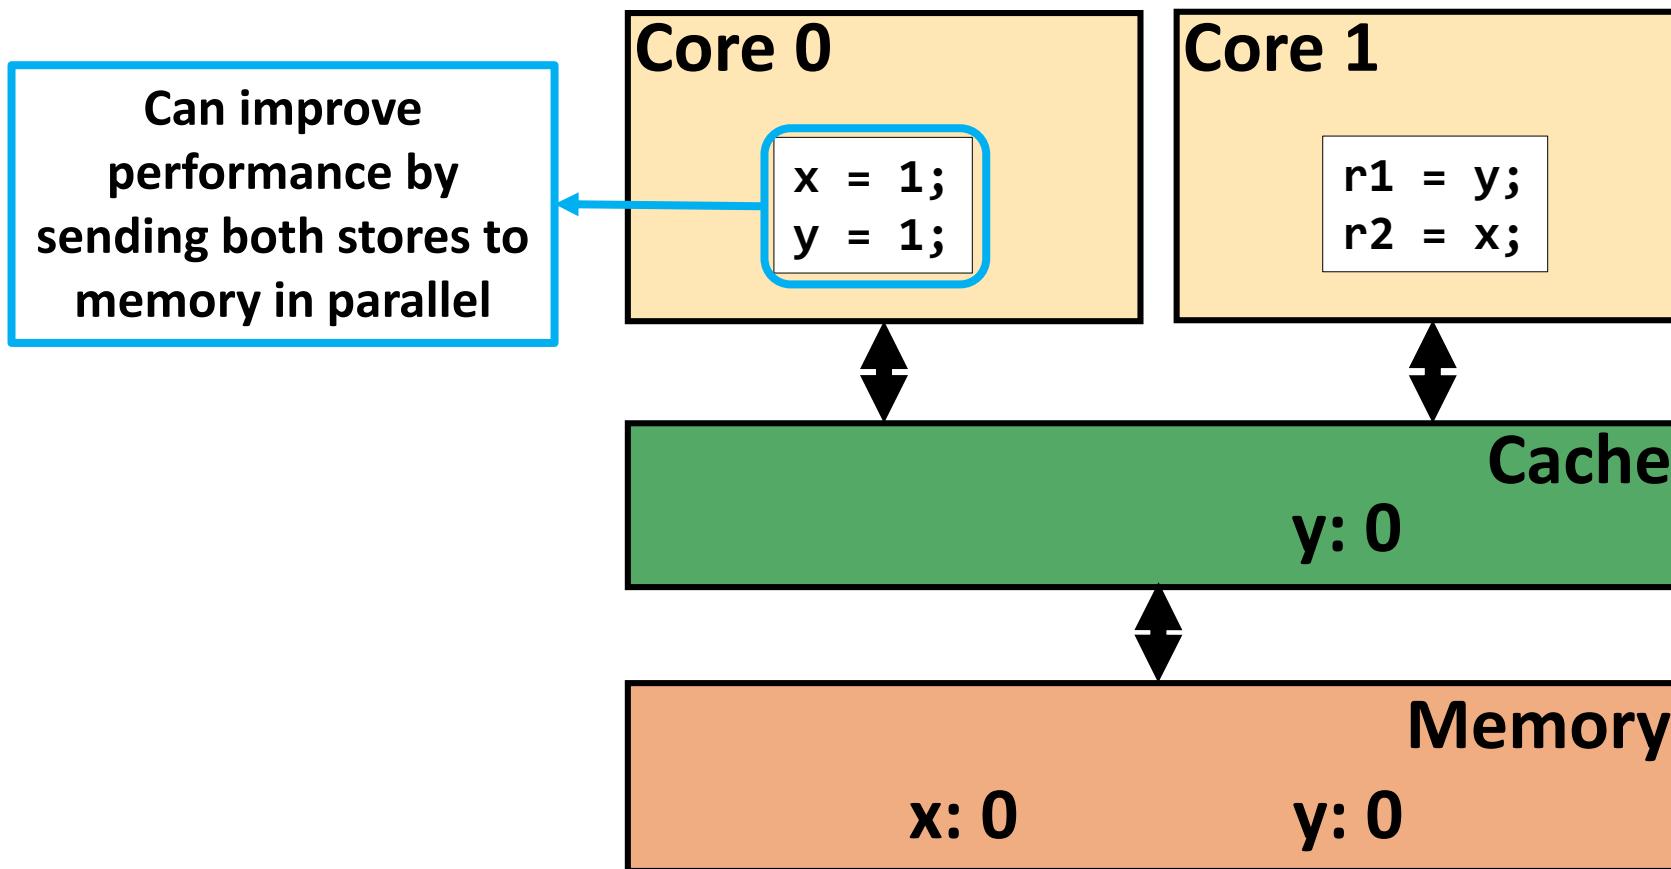
Message Passing (mp)

Core 0	Core 1
$x = 1;$	$r1 = y;$
$y = 1;$	$r2 = x;$
Can $r1=1$ and $r2=0$?	



Why reorder memory operations?

Answer: Performance!



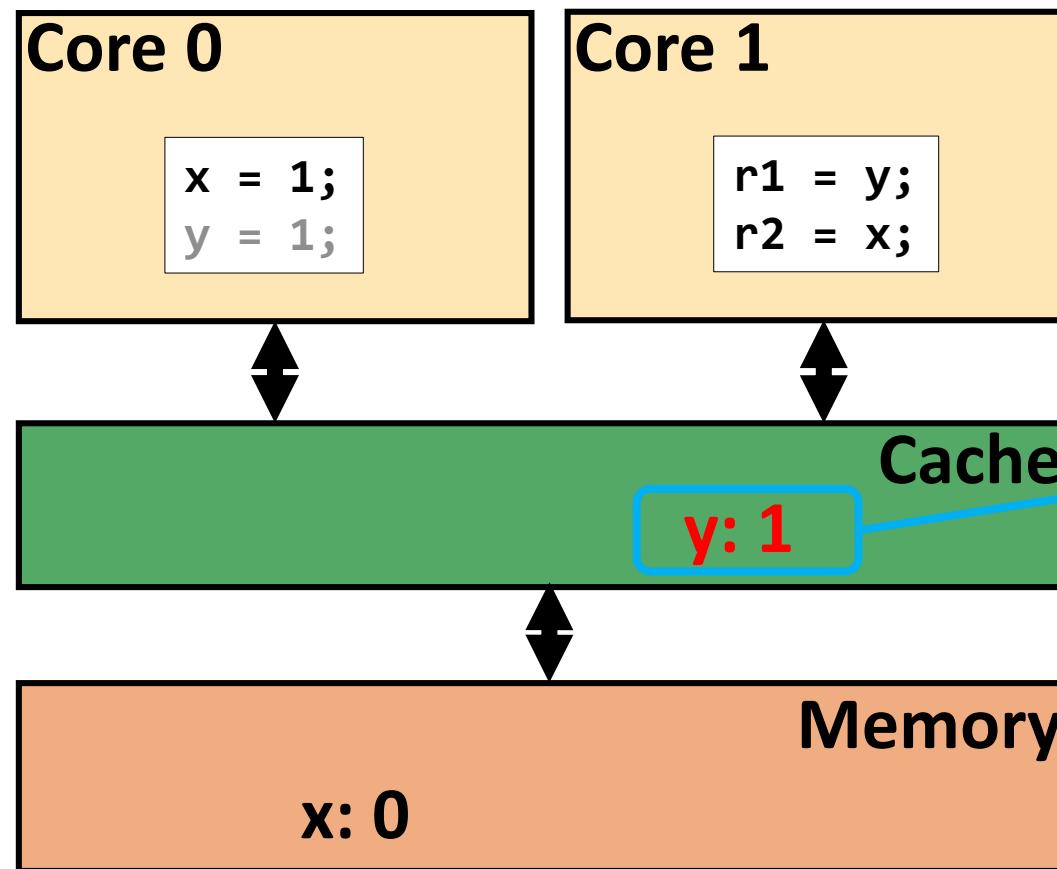
Message Passing (mp)

Core 0	Core 1
$x = 1;$	$r1 = y;$
$y = 1;$	$r2 = x;$
Can $r1=1$ and $r2=0$?	



Why reorder memory operations?

Answer: Performance!



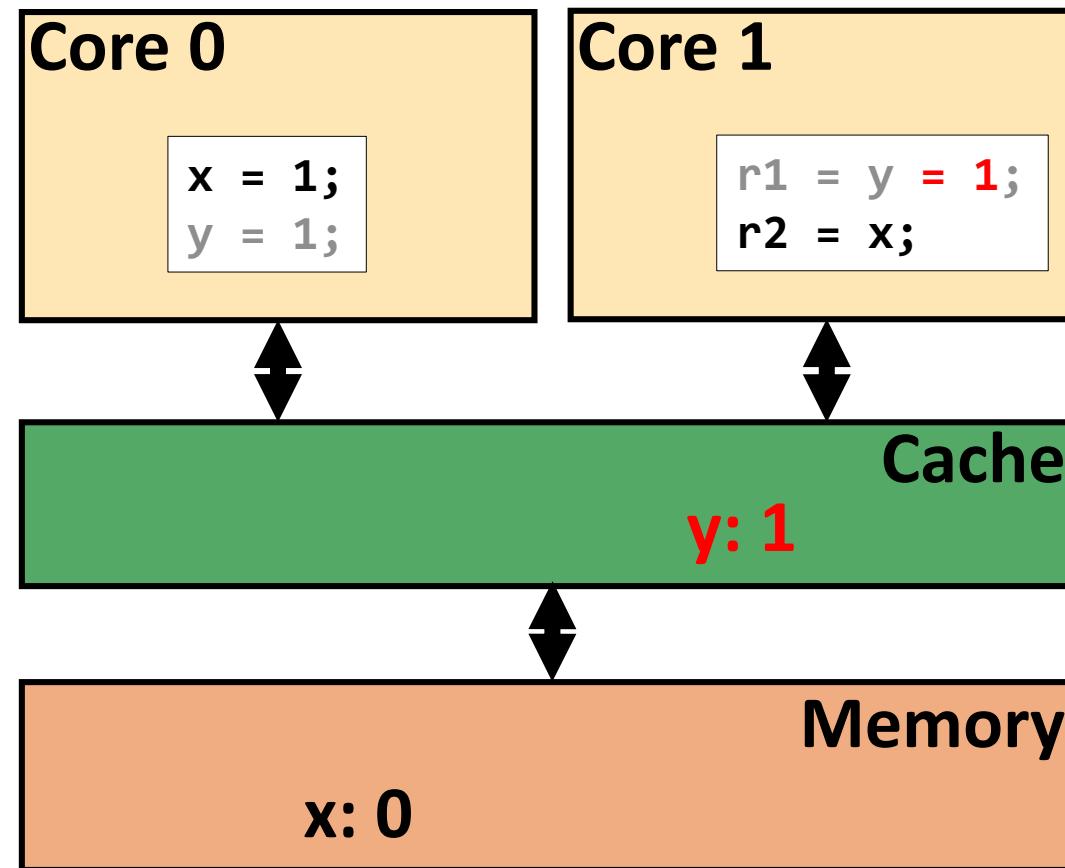
Message Passing (mp)

Core 0	Core 1
<code>x = 1;</code>	<code>r1 = y;</code>
<code>y = 1;</code>	<code>r2 = x;</code>
Can <code>r1=1</code> and <code>r2=0</code> ?	



Why reorder memory operations?

Answer: Performance!



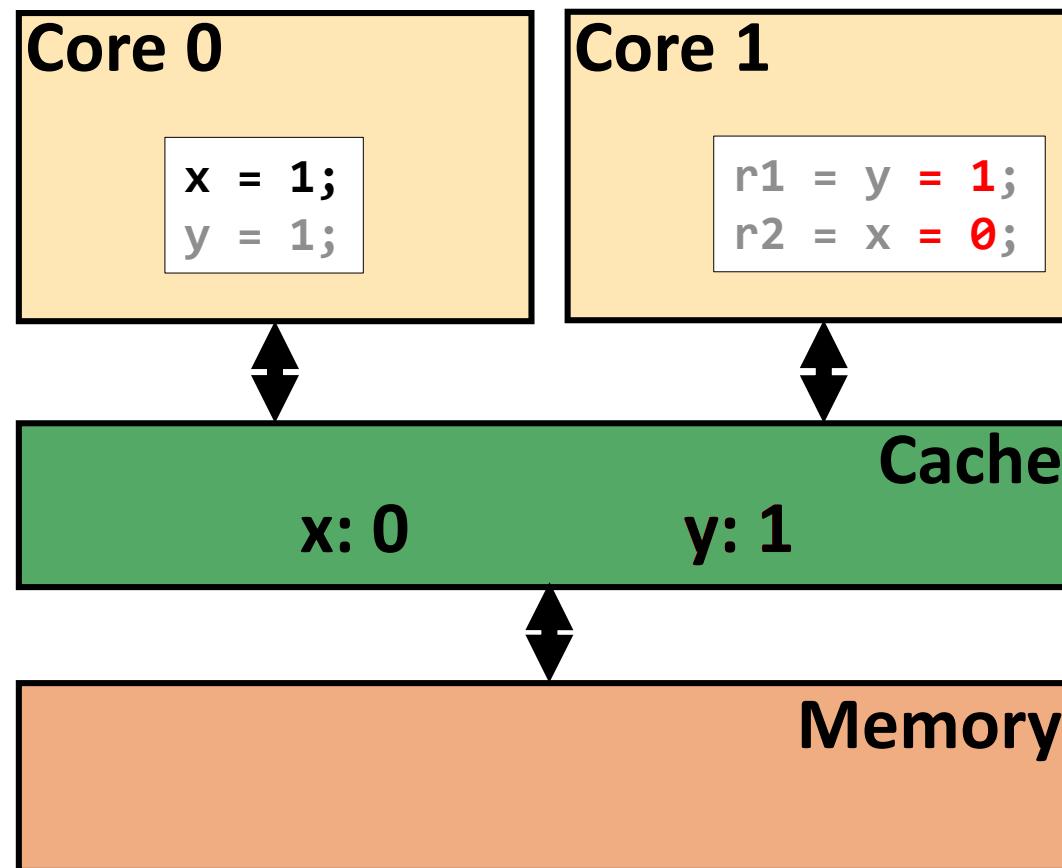
Message Passing (mp)

Core 0	Core 1
$x = 1;$	$r1 = y;$
$y = 1;$	$r2 = x;$
Can $r1=1$ and $r2=0$?	



Why reorder memory operations?

Answer: Performance!



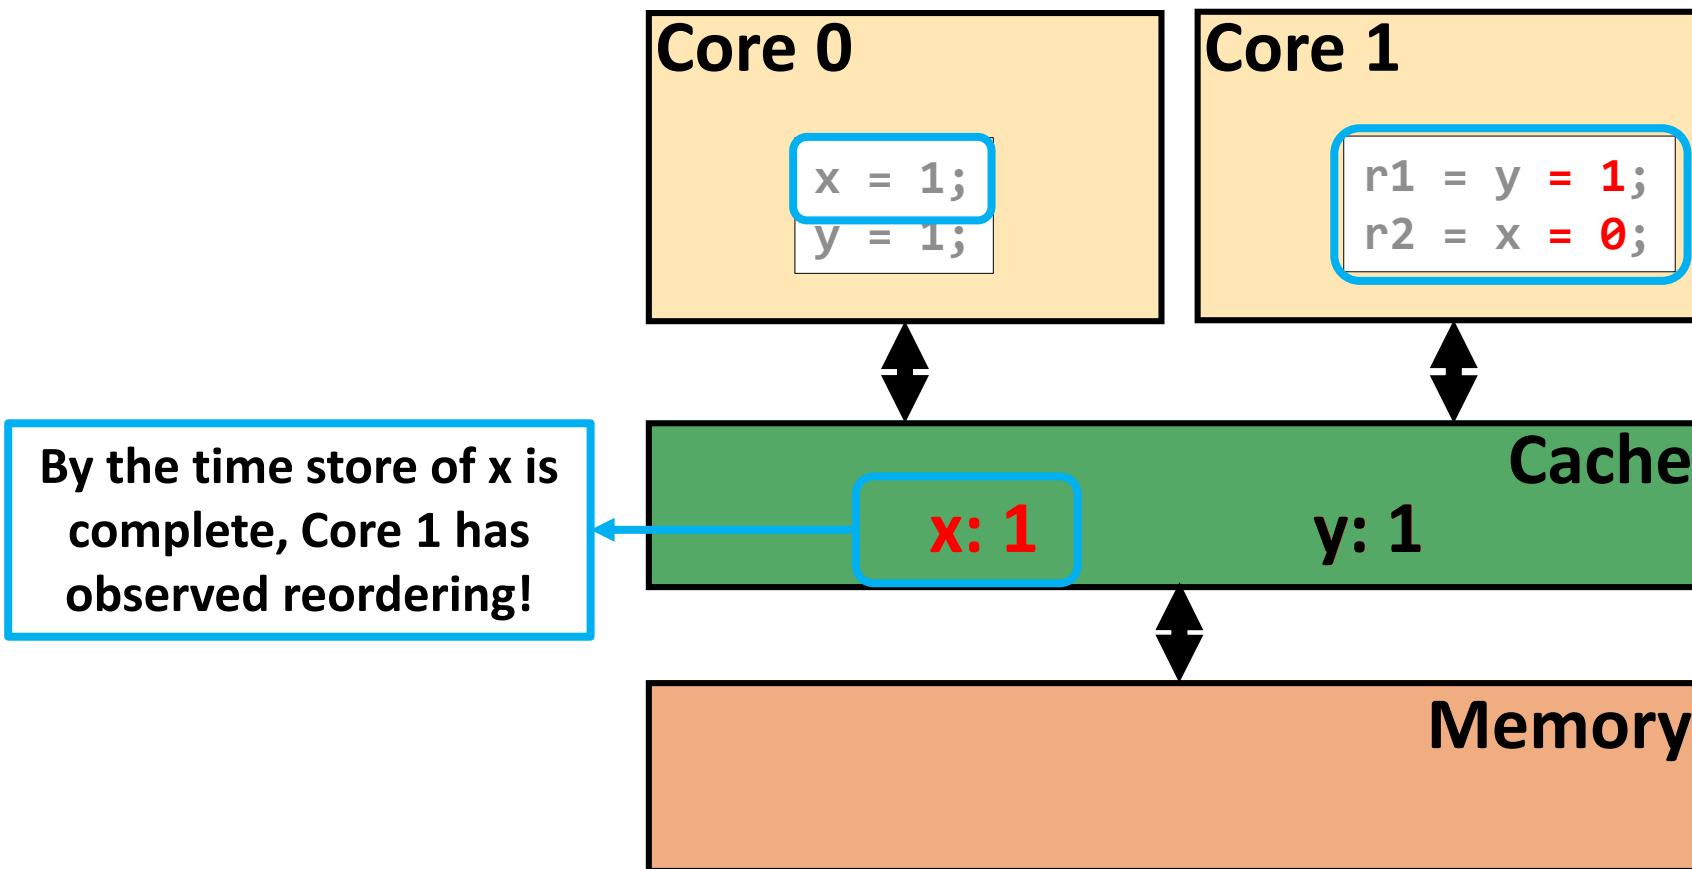
Message Passing (mp)

Core 0	Core 1
<code>x = 1;</code>	<code>r1 = y;</code>
<code>y = 1;</code>	<code>r2 = x;</code>
Can <code>r1=1</code> and <code>r2=0</code> ?	



Why reorder memory operations?

Answer: Performance!



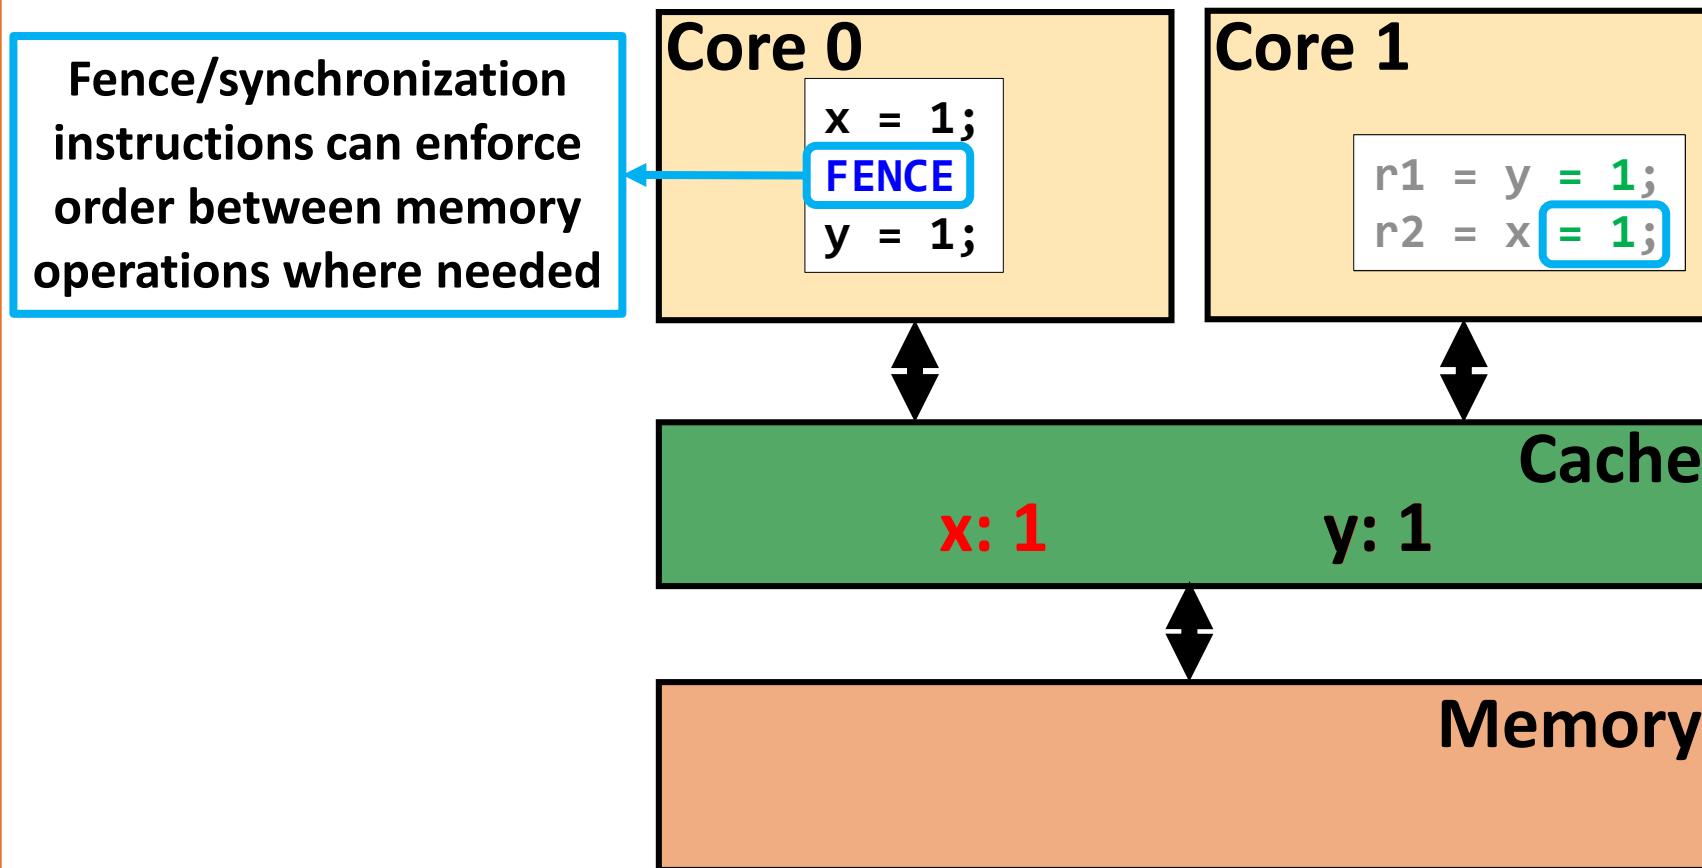
Message Passing (mp)

Core 0	Core 1
<code>x = 1;</code>	<code>r1 = y;</code>
<code>y = 1;</code>	<code>r2 = x;</code>
Can <code>r1=1</code> and <code>r2=0</code> ?	



Why reorder memory operations?

Answer: Performance!



Message Passing (mp)

Core 0	Core 1
<code>x = 1;</code>	<code>r1 = y;</code>
<code>y = 1;</code>	<code>r2 = x;</code>
Can <code>r1=1</code> and <code>r2=0</code> ?	



Memory Consistency Models (MCMs)

- Instruction sets (ISAs) represent hardware operations (add, ld, st, ...)
- MCMs similarly represent the orderings among hardware memory ops

Compiler

Hardware



Memory Consistency Models (MCMs)

- Instruction sets (ISAs) represent hardware operations (add, ld, st, ...)
- MCMs similarly represent the orderings among hardware memory ops

Where do I
need to add
fences?

Compiler

Hardware



Memory Consistency Models (MCMs)

- Instruction sets (ISAs) represent hardware operations (add, ld, st, ...)
- MCMs similarly represent the orderings among hardware memory ops

Where do I
need to add
fences?

Compiler

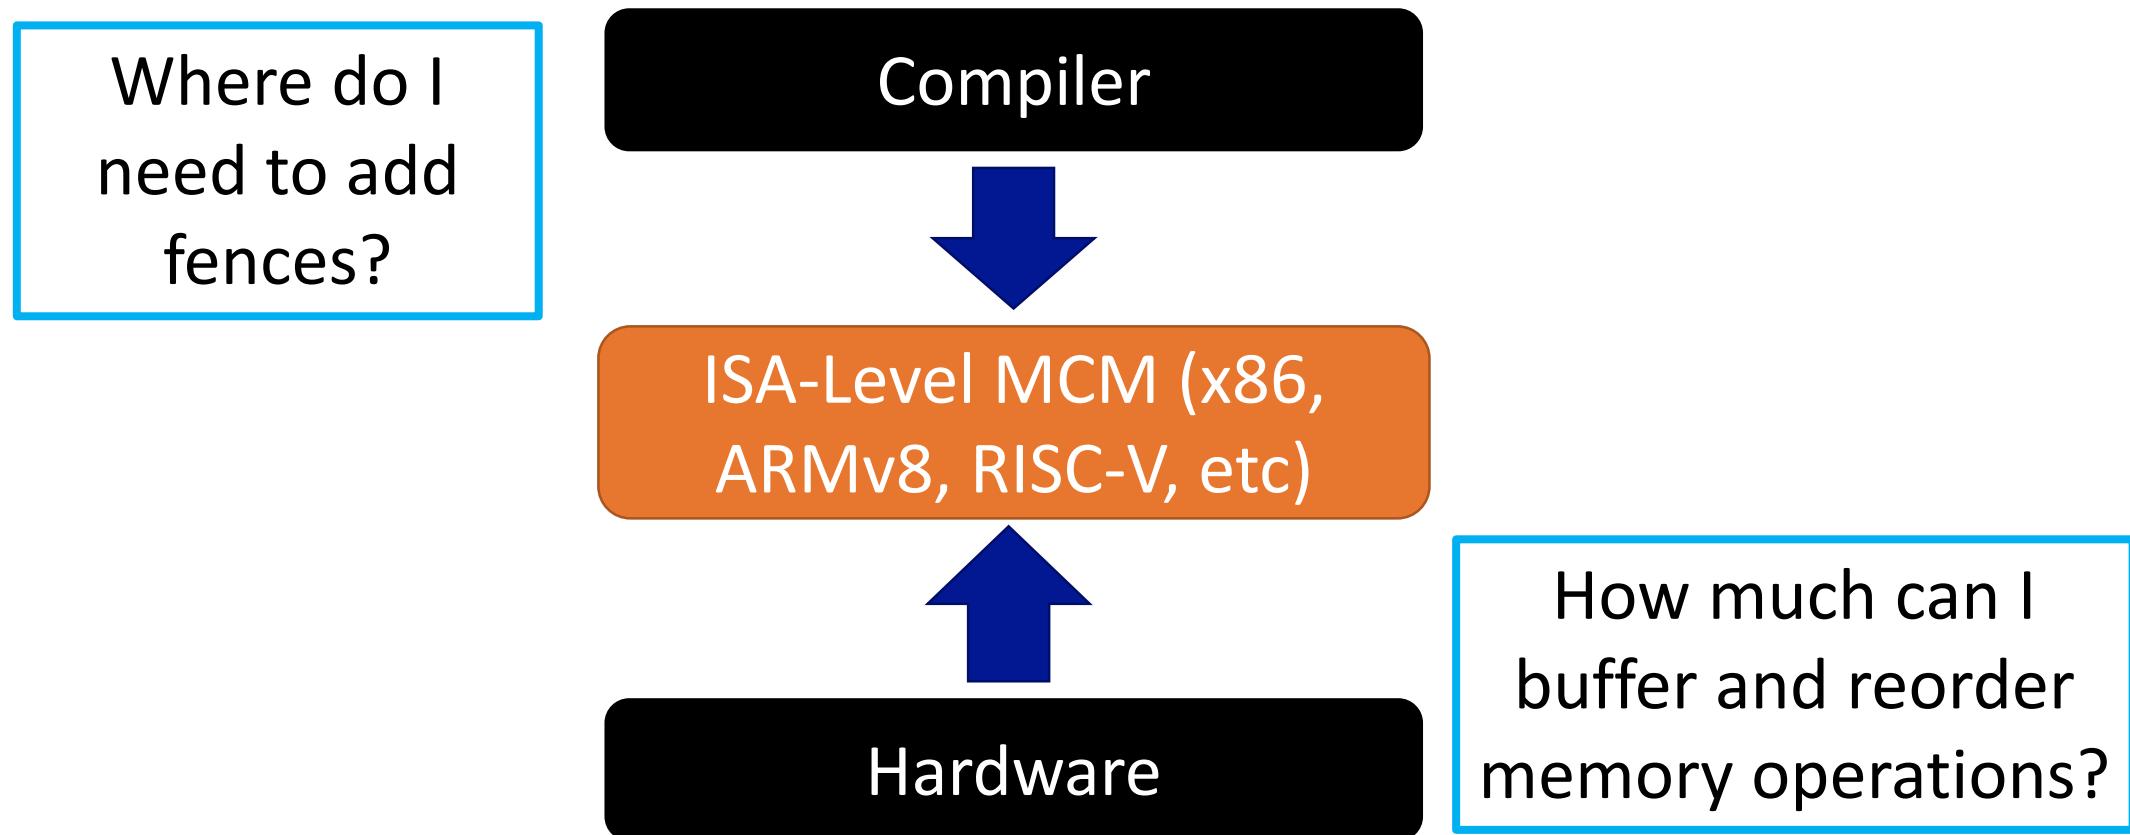
Hardware

How much can I
buffer and reorder
memory operations?



Memory Consistency Models (MCMs)

- Instruction sets (ISAs) represent hardware operations (add, ld, st, ...)
- MCMs similarly represent the orderings among hardware memory ops



Memory Consistency Models (MCMs)

- Instruction sets (ISAs) represent hardware operations (add, ld, st, ...)
- MCMs similarly represent the orderings among hardware memory ops

In a nutshell: MCMs specify what value will be returned when your program does a load!

ARMv8, RISC-V, etc)



Hardware

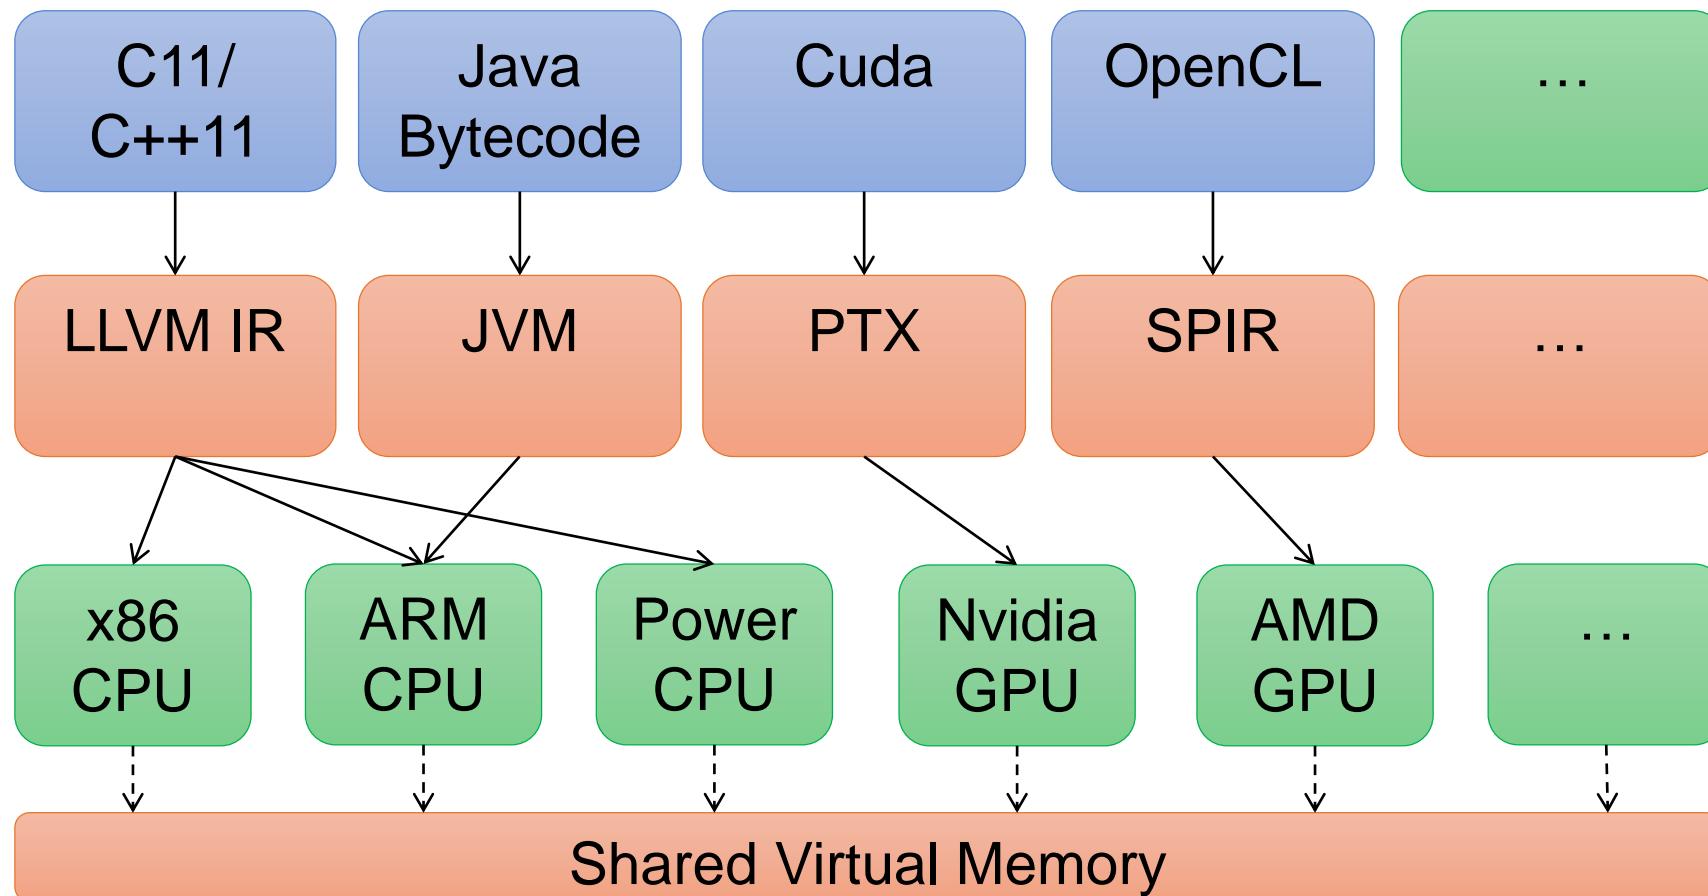
How much can I buffer and reorder memory operations?



Memory Consistency Models (MCMs)

Memory Consistency Models (MCMs)

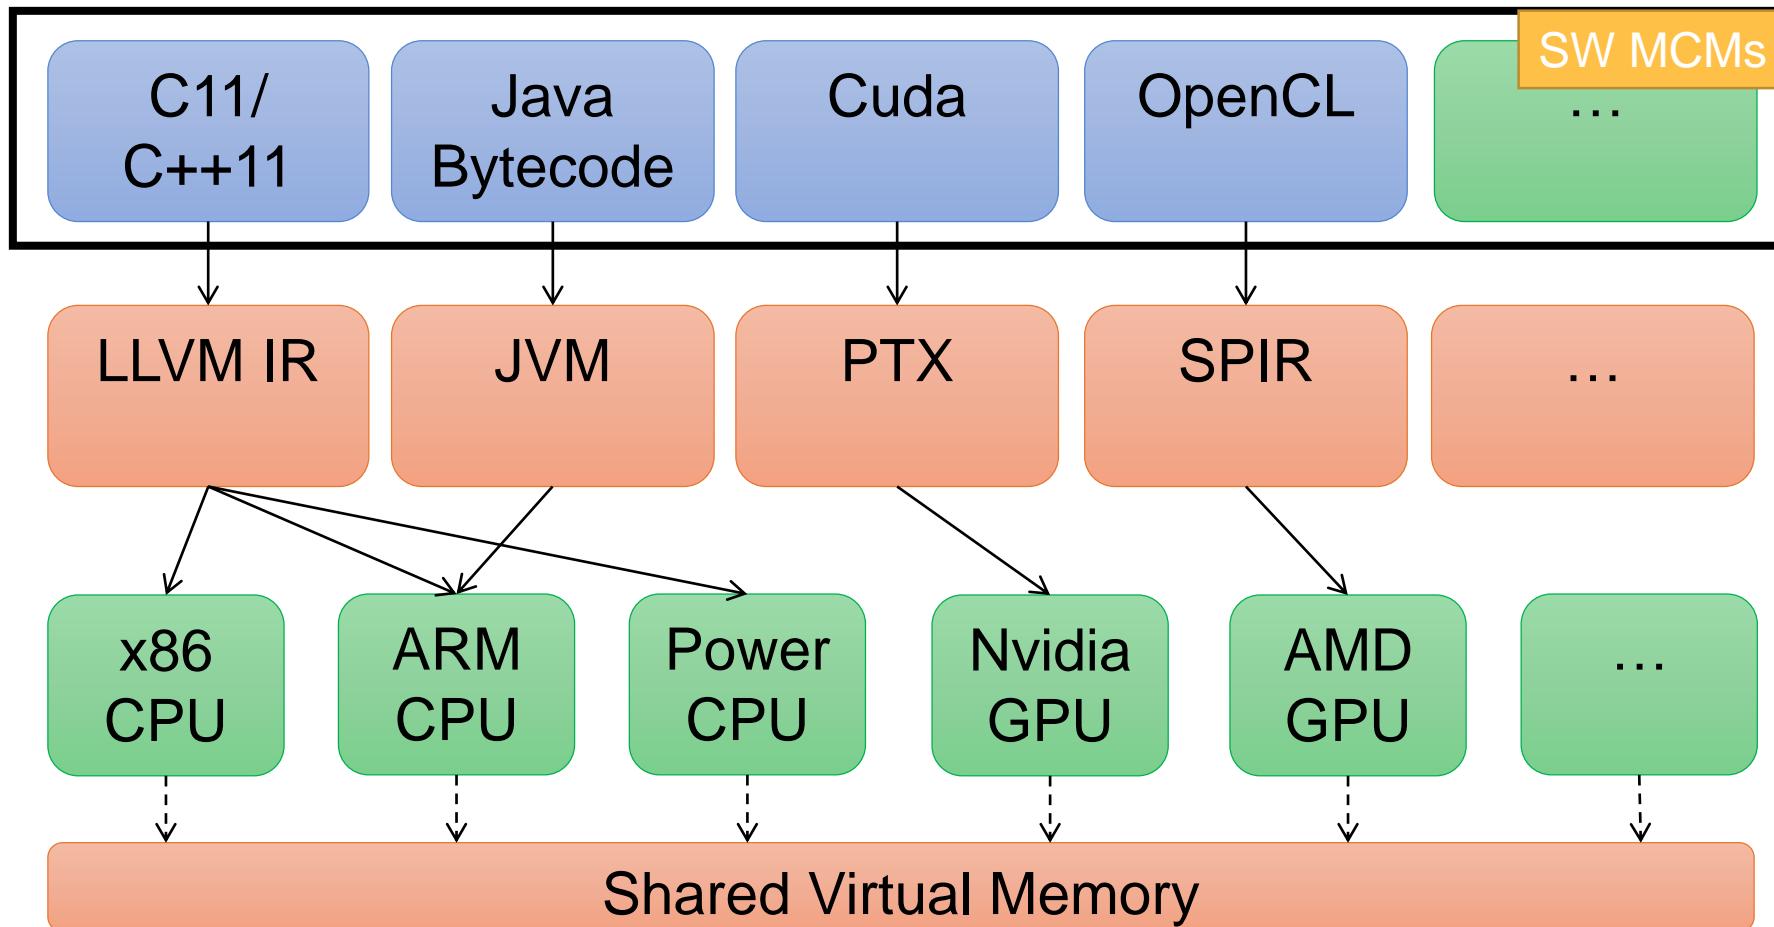
Specify rules and guarantees about the ordering and visibility of accesses to shared memory [Sorin et al., 2011].



Memory Consistency Models (MCMs)

Memory Consistency Models (MCMs)

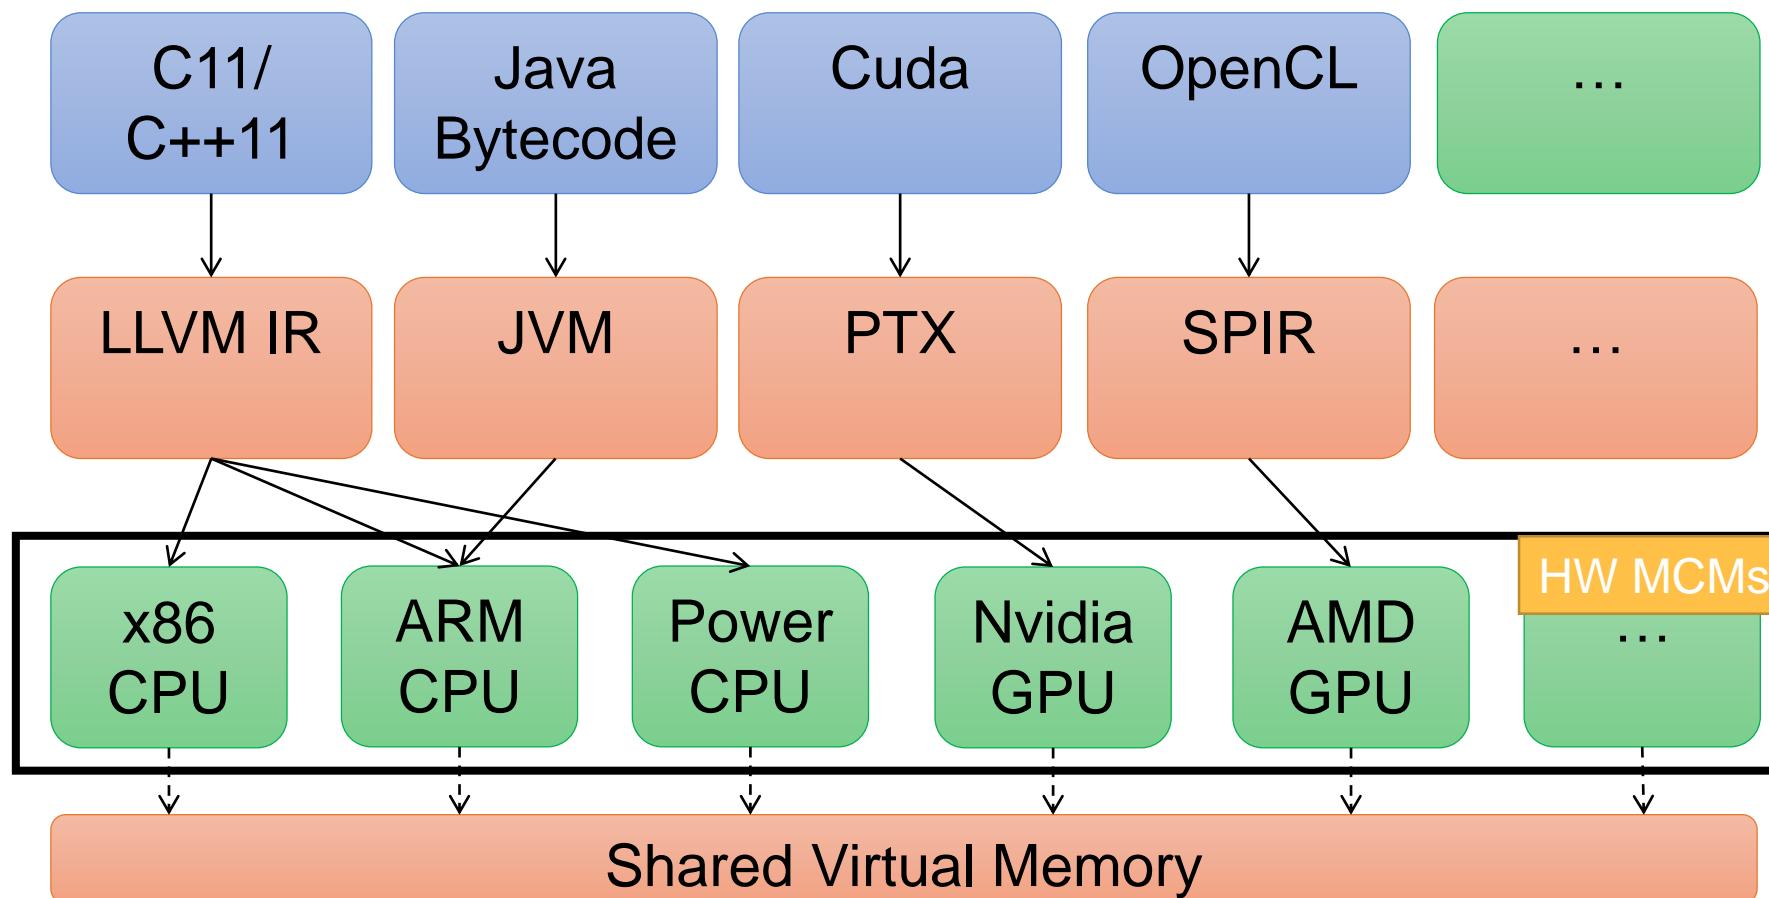
Specify rules and guarantees about the ordering and visibility of accesses to shared memory [Sorin et al., 2011].



Memory Consistency Models (MCMs)

Memory Consistency Models (MCMs)

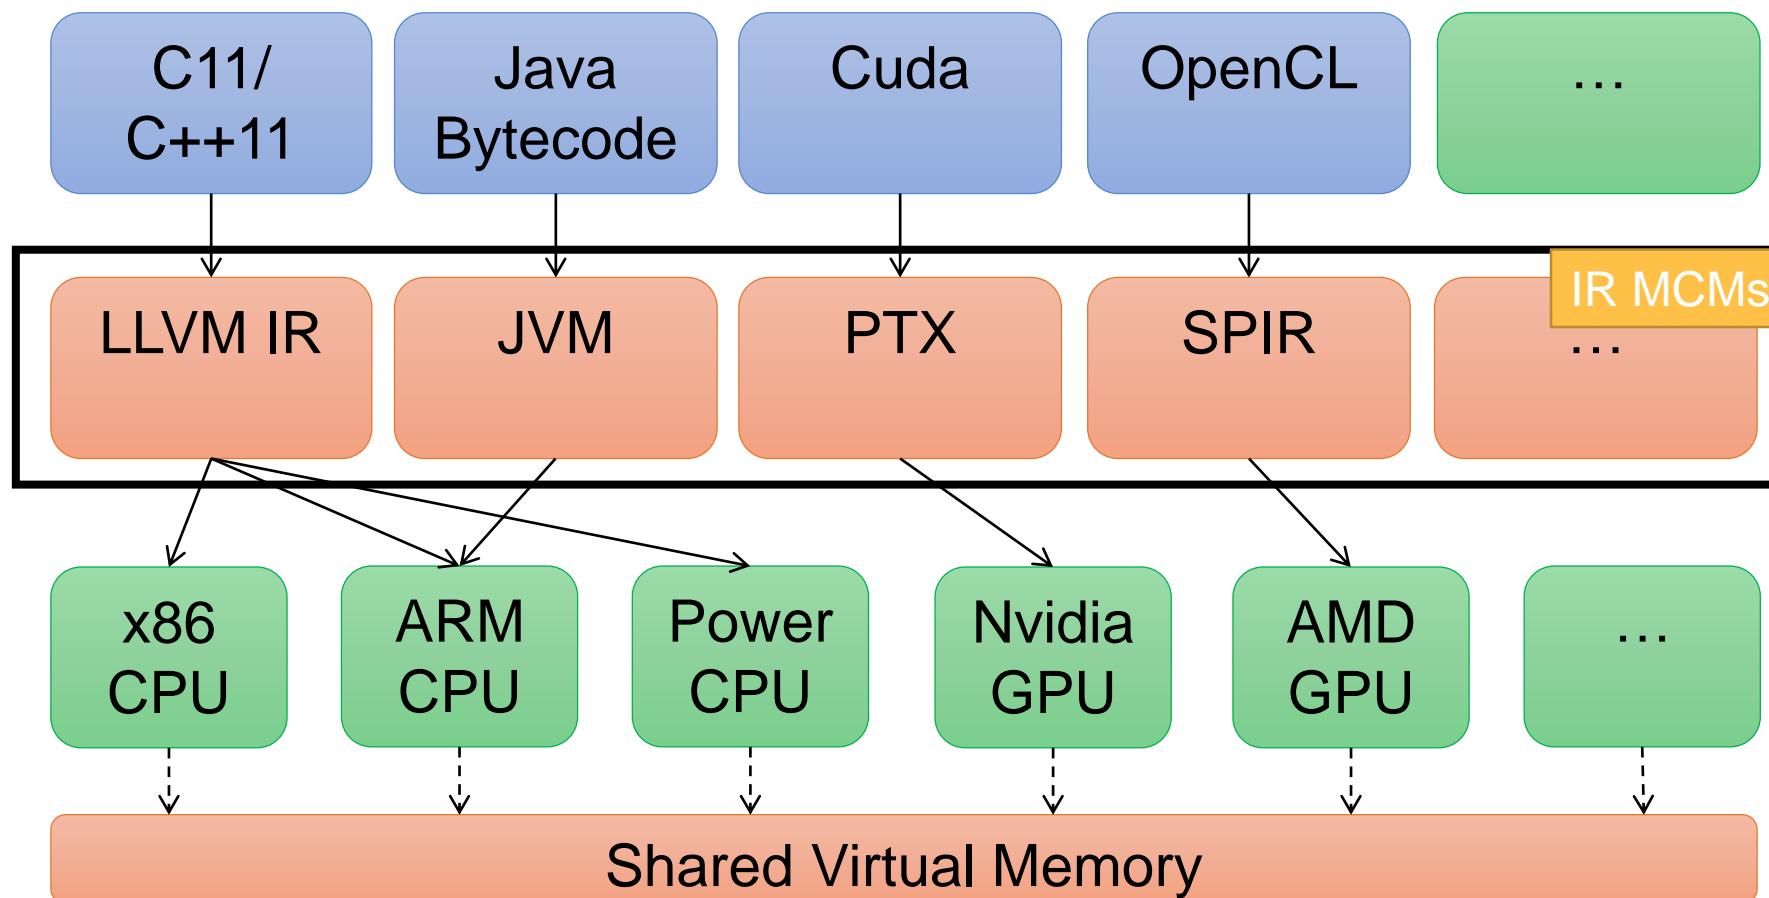
Specify rules and guarantees about the ordering and visibility of accesses to shared memory [Sorin et al., 2011].



Memory Consistency Models (MCMs)

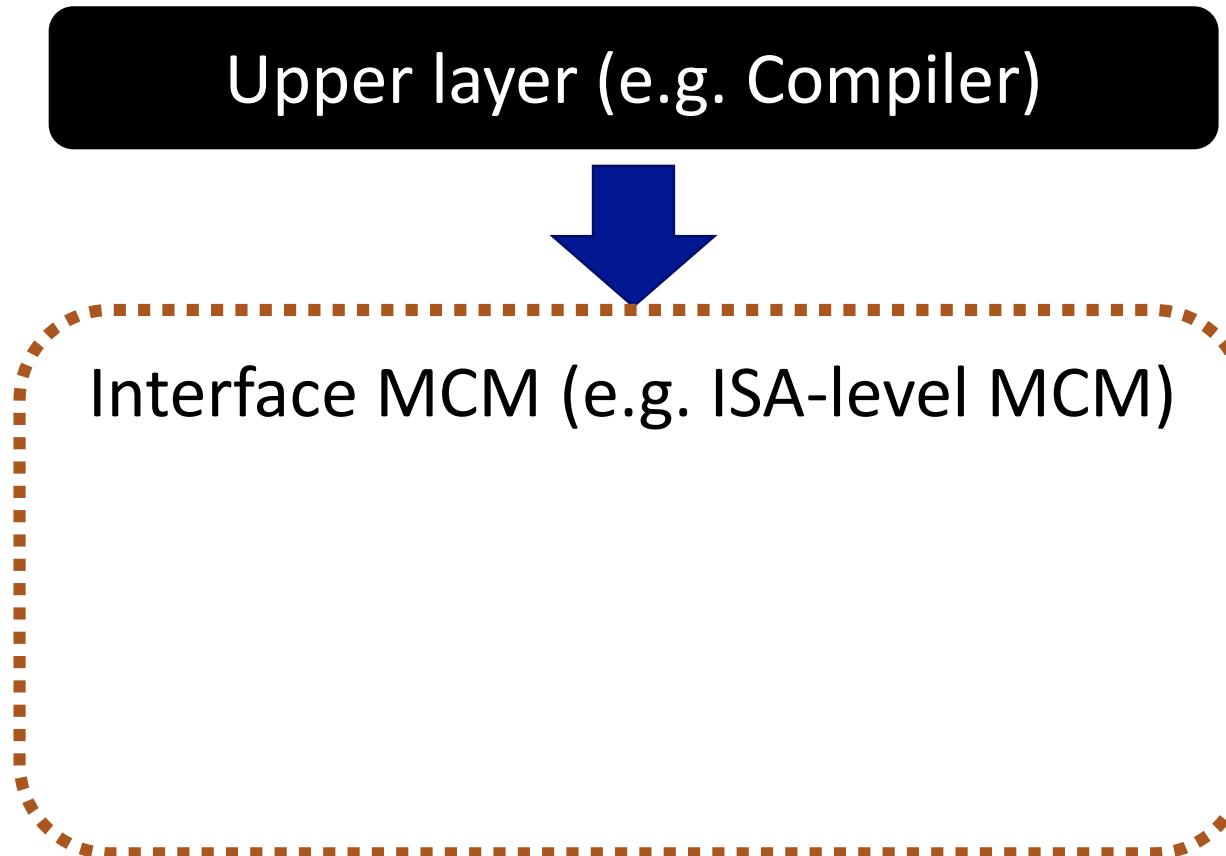
Memory Consistency Models (MCMs)

Specify rules and guarantees about the ordering and visibility of accesses to shared memory [Sorin et al., 2011].



The Need for MCM Verification

- MCMs are specified at interfaces between layers of the stack
 - Upper layers target MCM; **lower layers must maintain it for all programs!**

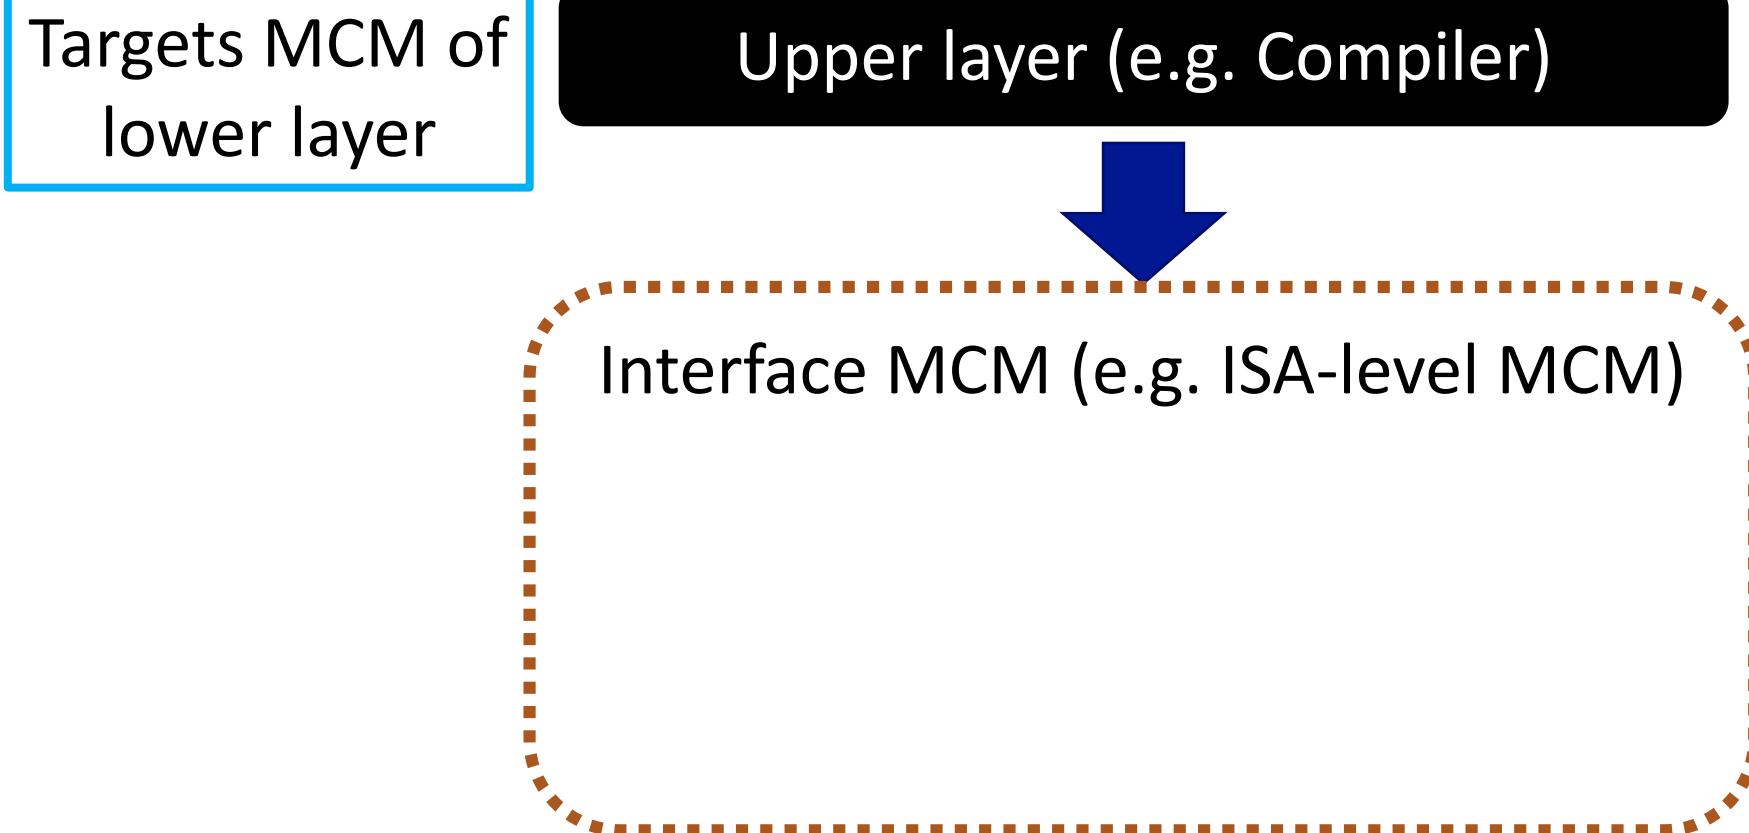


¹Microarchitecture is a component-level (e.g. caches, pipeline stages, store buffers) model of the hardware.



The Need for MCM Verification

- MCMs are specified at interfaces between layers of the stack
 - Upper layers target MCM; **lower layers must maintain it for all programs!**

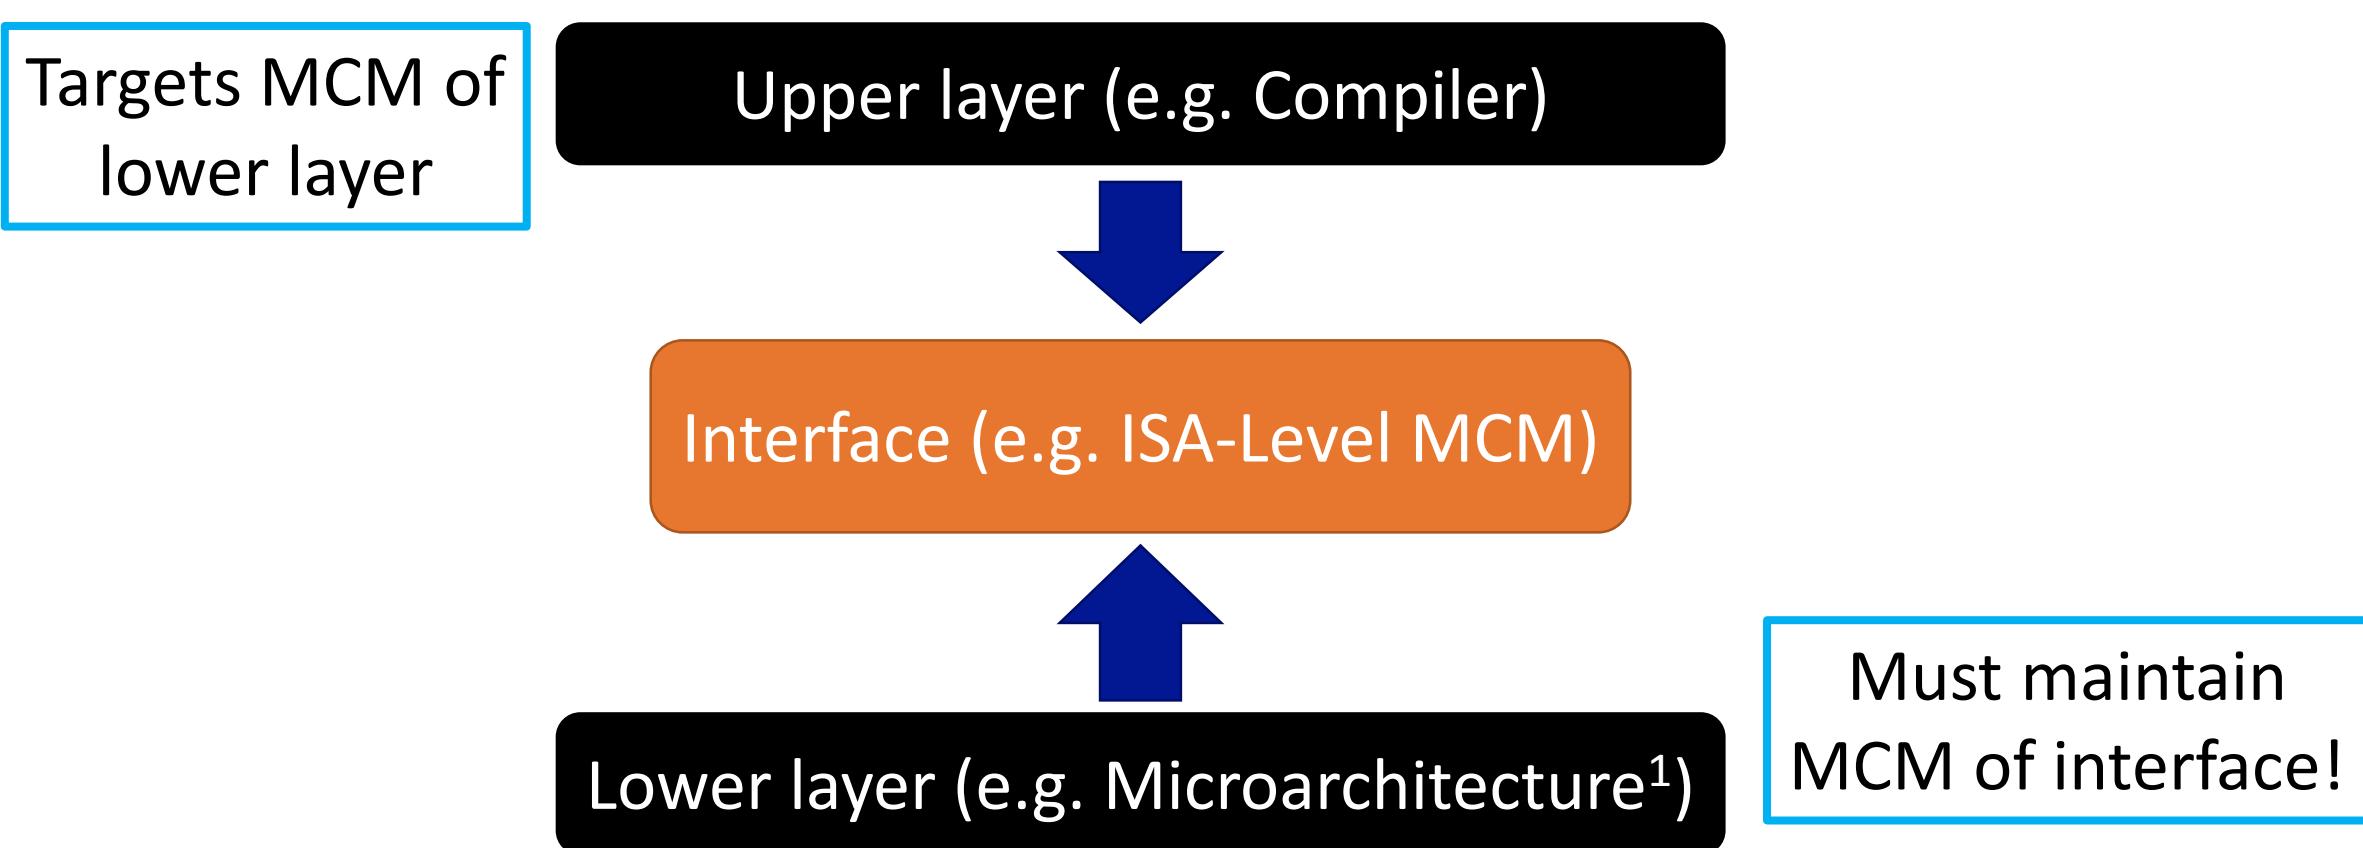


¹Microarchitecture is a component-level (e.g. caches, pipeline stages, store buffers) model of the hardware.



The Need for MCM Verification

- MCMs are specified at interfaces between layers of the stack
 - Upper layers target MCM; **lower layers must maintain it for all programs!**



¹Microarchitecture is a component-level (e.g. caches, pipeline stages, store buffers) model of the hardware.

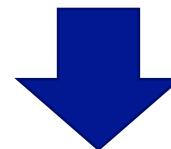


The Need for MCM Verification

- MCMs are specified at interfaces between layers of the stack
 - Upper layers target MCM; **lower layers must maintain it for all programs!**

Targets MCM of
lower layer

Upper layer (e.g. Compiler)



Lower layer (e.g. Microarchitecture¹)

Must maintain
MCM of interface!

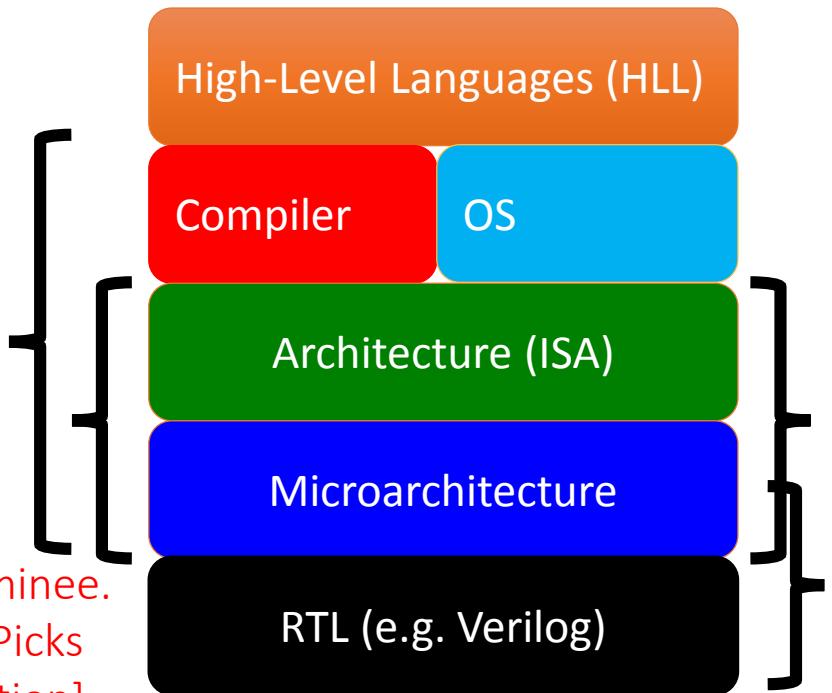
¹Microarchitecture is a component-level (e.g. caches, pipeline stages, store buffers) model of the hardware.



The Check Suite: Automated Tools For Verifying Memory Orderings and their Security Implications

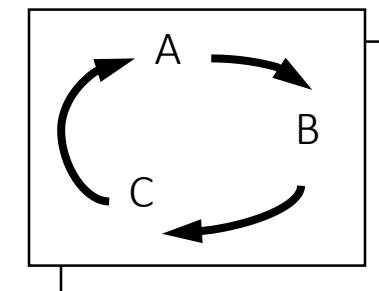
CheckMate
[Micro '18]
[IEEE Micro
Top Picks]

PipeProof
[Micro '18]
[Best Paper Nominee.
IEEE Micro Top Picks
Honorable Mention]



- Brackets on the right side group components into tool suites:
- Brackets group HLL, Compiler, OS, and Architecture (ISA) under **TriCheck** [ASPLOS '17] [IEEE MICRO Top Picks]
 - Brackets group HLL, Compiler, OS, and Architecture (ISA) under **COATCheck** [ASPLOS '16] [IEEE MICRO Top Picks]
 - Brackets group Microarchitecture and RTL under **PipeCheck** [Micro '14] [IEEE MICRO Top Picks]
 - Brackets group Microarchitecture and RTL under **CCICheck** [Micro '15] [Nominated for Best Paper Award]
 - Brackets group Microarchitecture and RTL under **RTLCheck** [Micro '17] [IEEE MICRO Top Picks Honorable Mention]

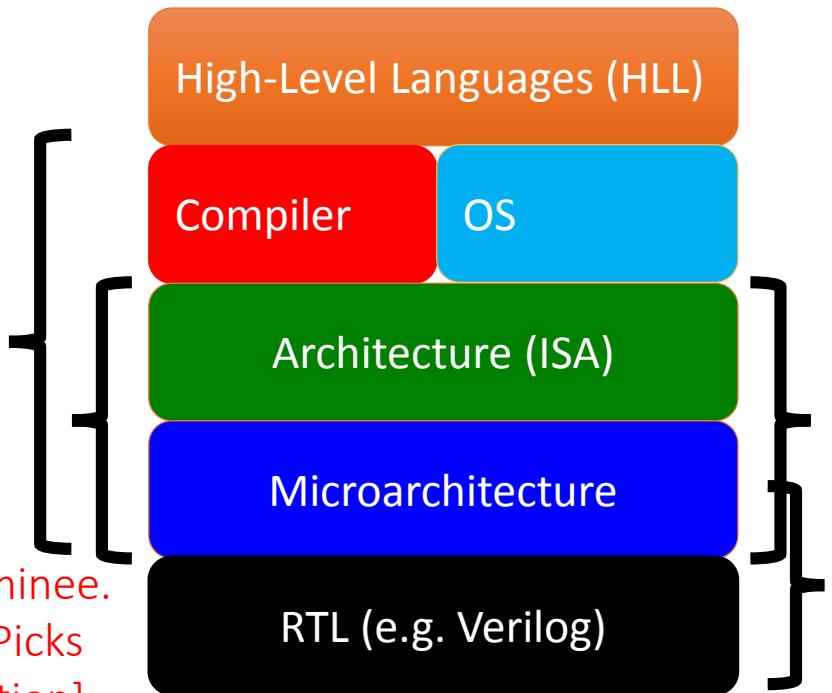
- Axiomatic specifications** -> **Happens-before graphs**
 - Cyclic => Impossible, Acyclic => Possible
- Model Checking space of graphs using **SMT solvers**
- Most tools **written in Gallina** => can be proven correct



The Check Suite: Automated Tools For Verifying Memory Orderings and their Security Implications

CheckMate
[Micro '18]
[IEEE Micro
Top Picks]

PipeProof
[Micro '18]
[Best Paper Nominee.
IEEE Micro Top Picks
Honorable Mention]



TriCheck [ASPLOS '17] [IEEE MICRO Top Picks]

COATCheck [ASPLOS '16] [IEEE MICRO Top Picks]

PipeCheck [Micro '14] [IEEE MICRO Top Picks]

CCICheck [Micro '15] [Nominated for Best Paper Award]

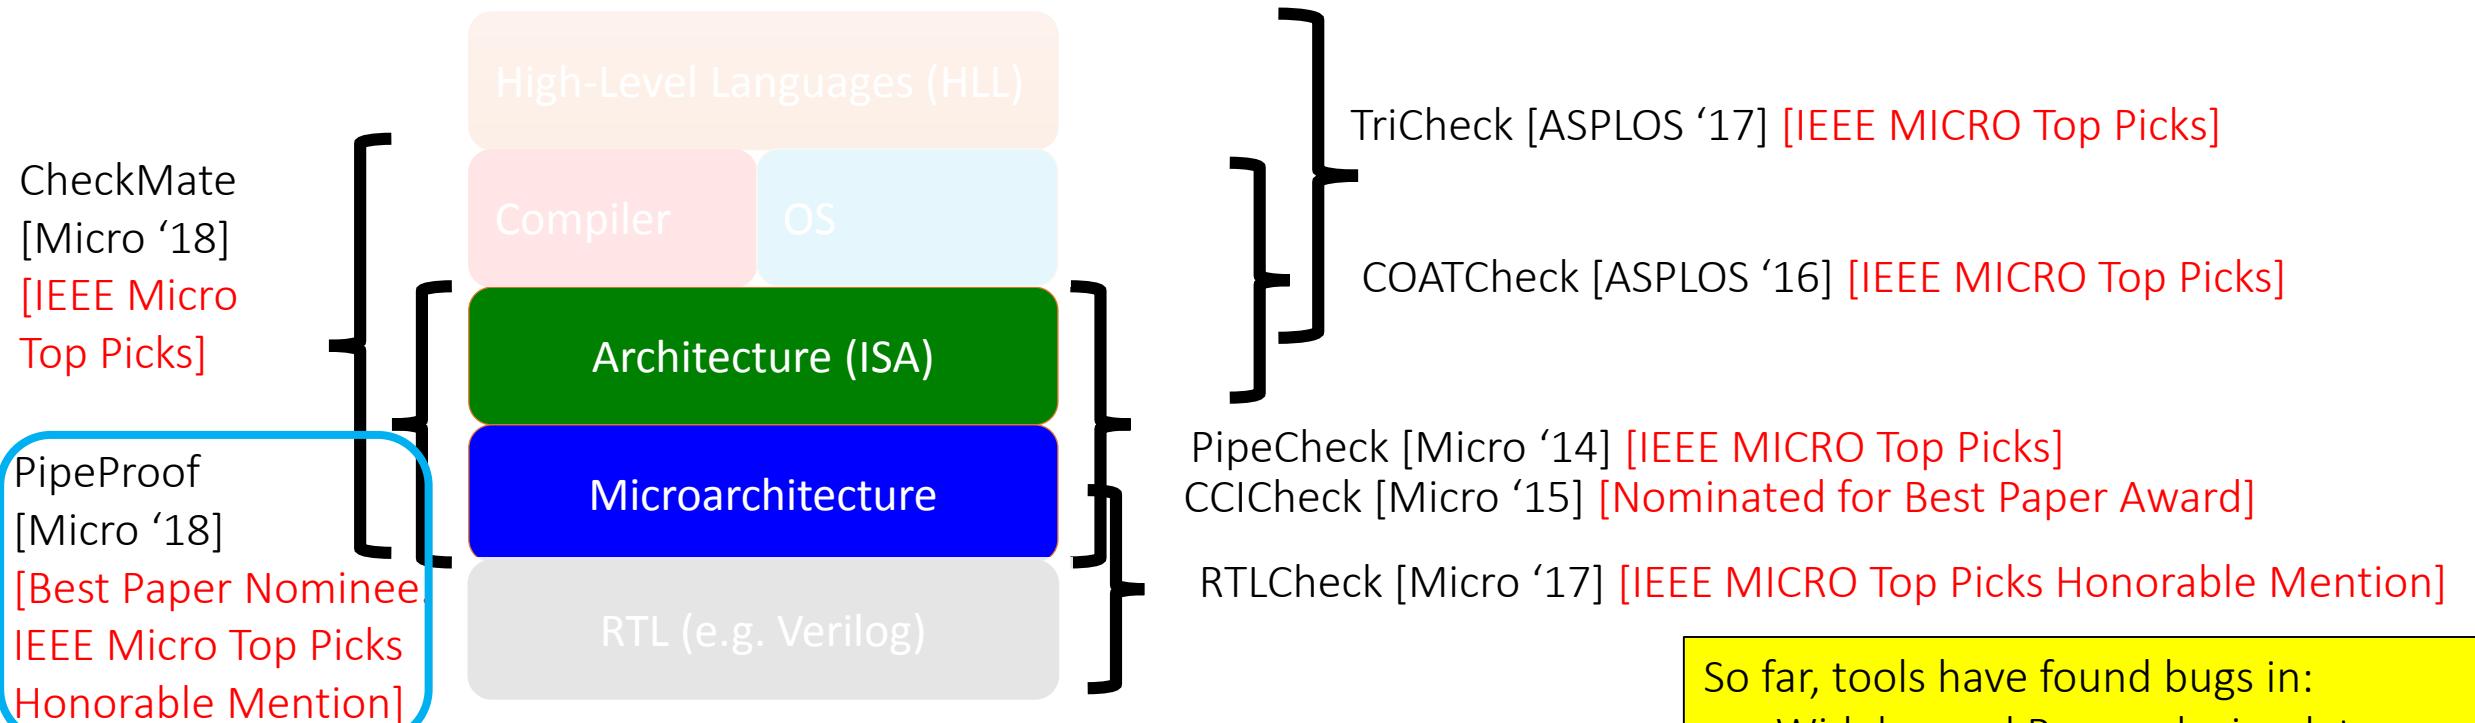
RTLCheck [Micro '17] [IEEE MICRO Top Picks Honorable Mention]

So far, tools have found bugs in:

- Widely-used Research simulator
- Cache coherence paper
- IBM XL C++ compiler (fixed in v13.1.5)
- In-design commercial processors
- RISC-V ISA specification
- Open-source RTL (Verilog)
- C++ 11 mem model
- SpectrePrime, MeltdownPrime

- **Axiomatic specifications -> Happens-before graphs**
 - Cyclic => Impossible, Acyclic => Possible
- Model Checking space of graphs using **SMT solvers**
- Most tools **written in Gallina** => can be proven correct

The Check Suite: Automated Tools For Verifying Memory Orderings and their Security Implications



- **Axiomatic specifications -> Happens-before graphs**
 - Cyclic => Impossible, Acyclic => Possible
- Model Checking space of graphs using **SMT solvers**
- Most tools **written in Gallina** => can be proven correct

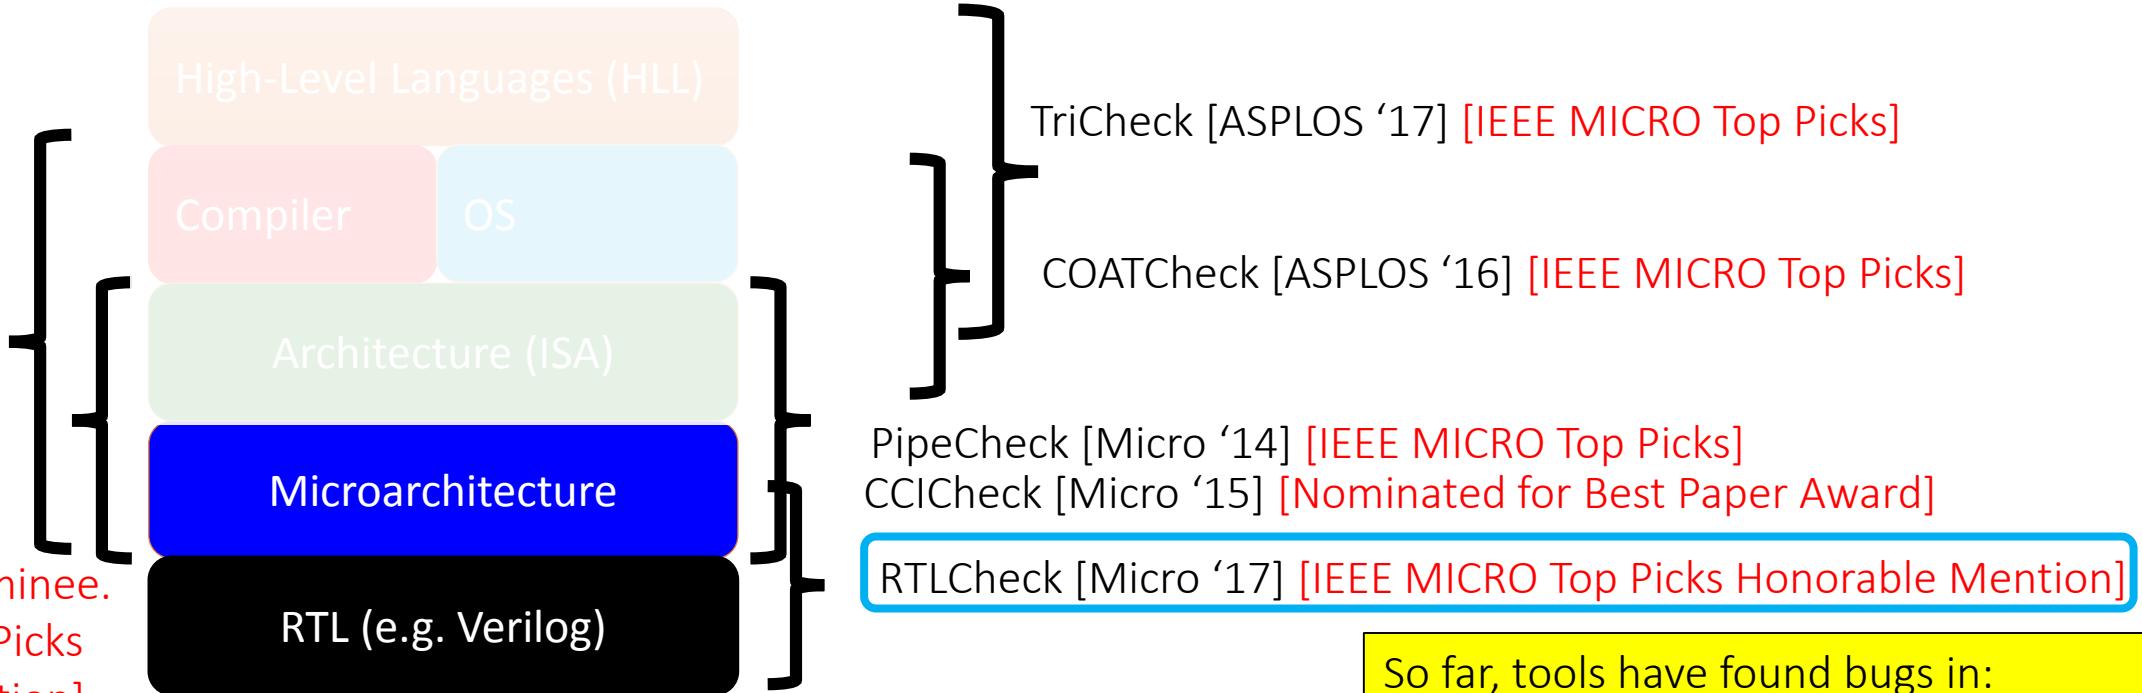
So far, tools have found bugs in:

- Widely-used Research simulator
- Cache coherence paper
- IBM XL C++ compiler (fixed in v13.1.5)
- In-design commercial processors
- RISC-V ISA specification
- Open-source RTL (Verilog)
- C++ 11 mem model
- SpectrePrime, MeltdownPrime

The Check Suite: Automated Tools For Verifying Memory Orderings and their Security Implications

CheckMate
[Micro '18]
[IEEE Micro
Top Picks]

PipeProof
[Micro '18]
[Best Paper Nominee.
IEEE Micro Top Picks
Honorable Mention]



- **Axiomatic specifications -> Happens-before graphs**
 - Cyclic => Impossible, Acyclic => Possible
- Model Checking space of graphs using **SMT solvers**
- Most tools **written in Gallina** => can be proven correct

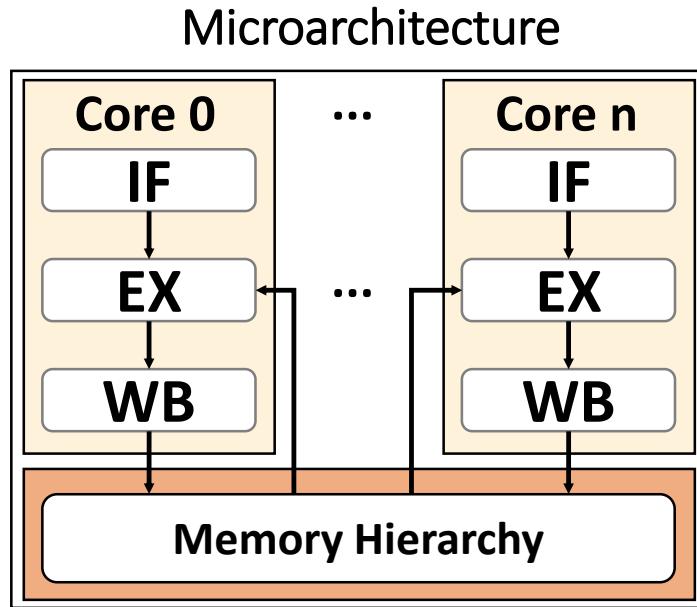
So far, tools have found bugs in:

- Widely-used Research simulator
- Cache coherence paper
- IBM XL C++ compiler (fixed in v13.1.5)
- In-design commercial processors
- RISC-V ISA specification
- Open-source RTL (Verilog)
- C++ 11 mem model
- SpectrePrime, MeltdownPrime

Talk Outline

- Overview and Motivation
- Memory Consistency Background
- **PipeProof:** All-Program Microarchitectural MCM Verification
- **RTLCheck:** MCM Verification of Verilog RTL
- Expanding to other domains
- Conclusion

Microarchitectural MCM Verification



SC/TSO/RISC-V MCM?

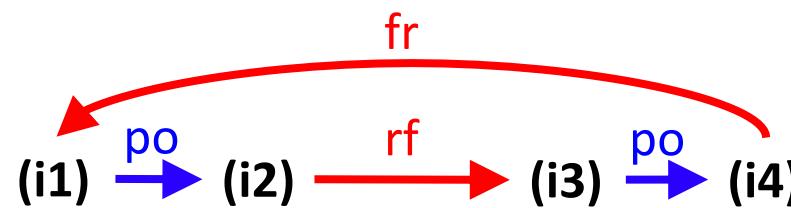
- PipeProof proves that a microarchitecture respects its ISA MCM
 - For all possible programs!
- How do we formally specify
 - ISA-level MCMs?
 - Microarchitectural orderings?



ISA-Level MCM Specifications

- MCMs often defined using relational patterns
 - [Shasha and Snir TOPLAS 1988] [Alglave et al. TOPLAS 2014]
- ISA-level executions are graphs
 - **nodes:** instructions, **edges:** ISA-level relations
- Eg: SC is *acyclic*($po \cup co \cup rf \cup fr$)

Message passing (mp) litmus test	
Core 0	Core 1
(i1) $x = 1;$	(i3) $r1 = y;$
(i2) $y = 1;$	(i4) $r2 = x;$
SC Forbids: $r1 = 1, r2 = 0$	



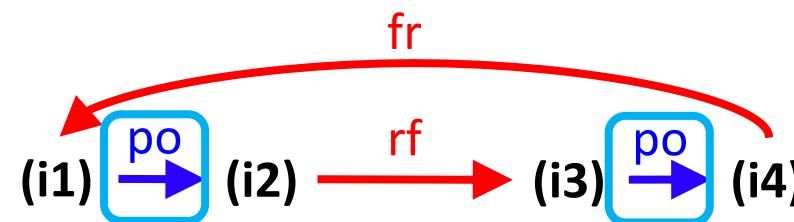
Legend:
po = Program order
co = coherence order
rf = reads-from
fr = from-reads

- Formal specifications of ISA + HLL MCMs in recent years
 - x86 [Owens et al. TPHOLS2009], ARM [Pulte et al. POPL2018], C11 [Batty et al. POPL 2011], ...
- Automated formal tools e.g. **herd** [Alglave et al. TOPLAS 2014]
 - Can formally analyse small test programs against these models

ISA-Level MCM Specifications

- MCMs often defined using relational patterns
 - [Shasha and Snir TOPLAS 1988] [Alglave et al. TOPLAS 2014]
- ISA-level executions are graphs
 - **nodes:** instructions, **edges:** ISA-level relations
- Eg: SC is *acyclic(po \cup co \cup rf \cup fr)*

Message passing (mp) litmus test	
Core 0	Core 1
(i1) x = 1;	(i3) r1 = y;
(i2) y = 1;	(i4) r2 = x;
SC Forbids: r1 = 1, r2 = 0	



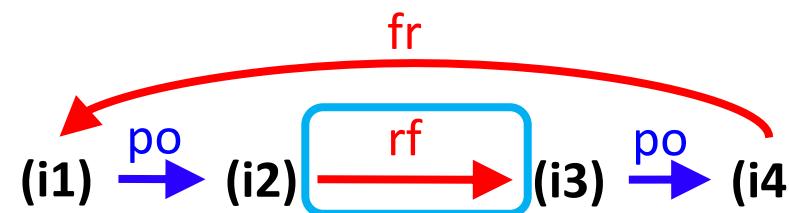
Legend:
po = Program order
co = coherence order
rf = reads-from
fr = from-reads

- Formal specifications of ISA + HLL MCMs in recent years
 - x86 [Owens et al. TPHOLS2009], ARM [Pulte et al. POPL2018], C11 [Batty et al. POPL 2011], ...
- Automated formal tools e.g. **herd** [Alglave et al. TOPLAS 2014]
 - Can formally analyse small test programs against these models

ISA-Level MCM Specifications

- MCMs often defined using relational patterns
 - [Shasha and Snir TOPLAS 1988] [Alglave et al. TOPLAS 2014]
- ISA-level executions are graphs
 - **nodes:** instructions, **edges:** ISA-level relations
- Eg: SC is *acyclic*($po \cup co \cup rf \cup fr$)

Message passing (mp) litmus test	
Core 0	Core 1
(i1) $x = 1;$	(i3) $r1 = y;$
(i2) $y = 1;$	(i4) $r2 = x;$
SC Forbids: $r1 = 1, r2 = 0$	



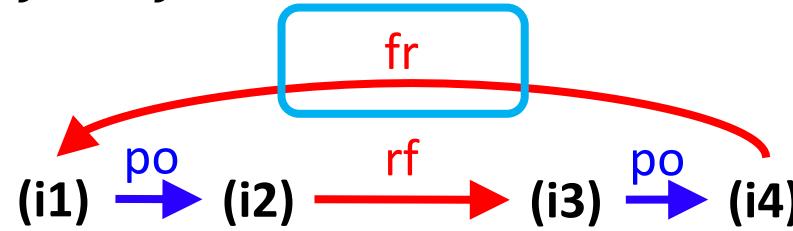
Legend:
po = Program order
co = coherence order
rf = reads-from
fr = from-reads

- Formal specifications of ISA + HLL MCMs in recent years
 - x86 [Owens et al. TPHOLS2009], ARM [Pulte et al. POPL2018], C11 [Batty et al. POPL 2011], ...
- Automated formal tools e.g. **herd** [Alglave et al. TOPLAS 2014]
 - Can formally analyse small test programs against these models

ISA-Level MCM Specifications

- MCMs often defined using relational patterns
 - [Shasha and Snir TOPLAS 1988] [Alglave et al. TOPLAS 2014]
- ISA-level executions are graphs
 - **nodes:** instructions, **edges:** ISA-level relations
- Eg: SC is *acyclic*($po \cup co \cup rf \cup fr$)

Message passing (mp) litmus test	
Core 0	Core 1
(i1) $x = 1;$	(i3) $r1 = y;$
(i2) $y = 1;$	(i4) $r2 = x;$
SC Forbids: $r1 = 1, r2 = 0$	



Legend:
po = Program order
co = coherence order
rf = reads-from
fr = from-reads

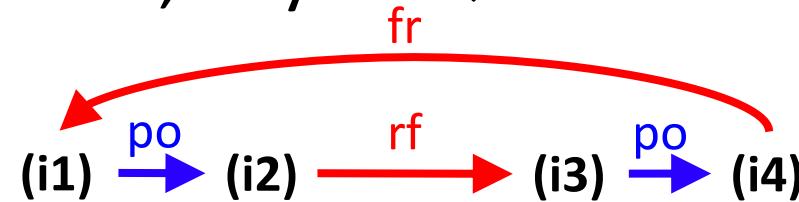
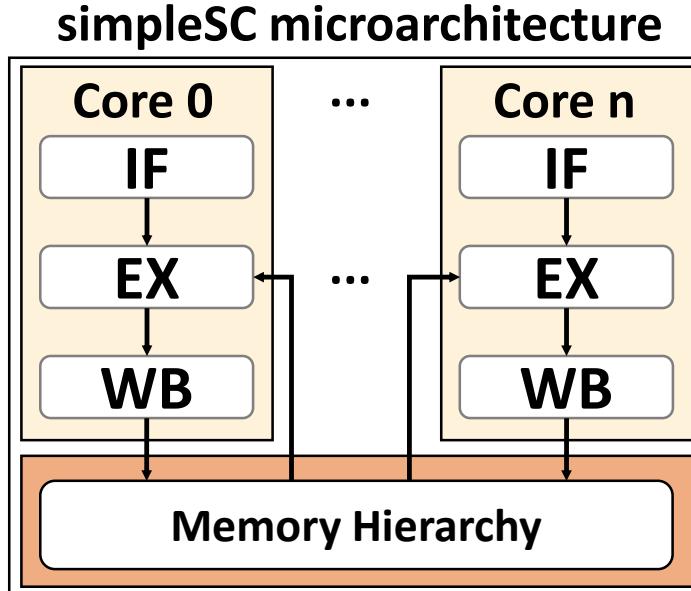
- Formal specifications of ISA + HLL MCMs in recent years
 - x86 [Owens et al. TPHOLS2009], ARM [Pulte et al. POPL2018], C11 [Batty et al. POPL 2011], ...
- Automated formal tools e.g. **herd** [Alglave et al. TOPLAS 2014]
 - Can formally analyse small test programs against these models

Microarchitectural Happens-Before (μ hb) Graphs

- Developed by PipeCheck [Lustig et al. MICRO 2014]
- Microarchitecture performs instrs. in stages
- Microarchitectural executions are μ hb graphs
 - **Nodes:** instr. sub-events, **edges:** happens-before relationships
- Cyclic μ hb graph → **unobservable**, Acyclic → **observable**

Message passing (mp) litmus test	
Core 0	Core 1
(i1) $x = 1;$	(i3) $r1 = y;$
(i2) $y = 1;$	(i4) $r2 = x;$

SC Forbids: $r1 = 1, r2 = 0$



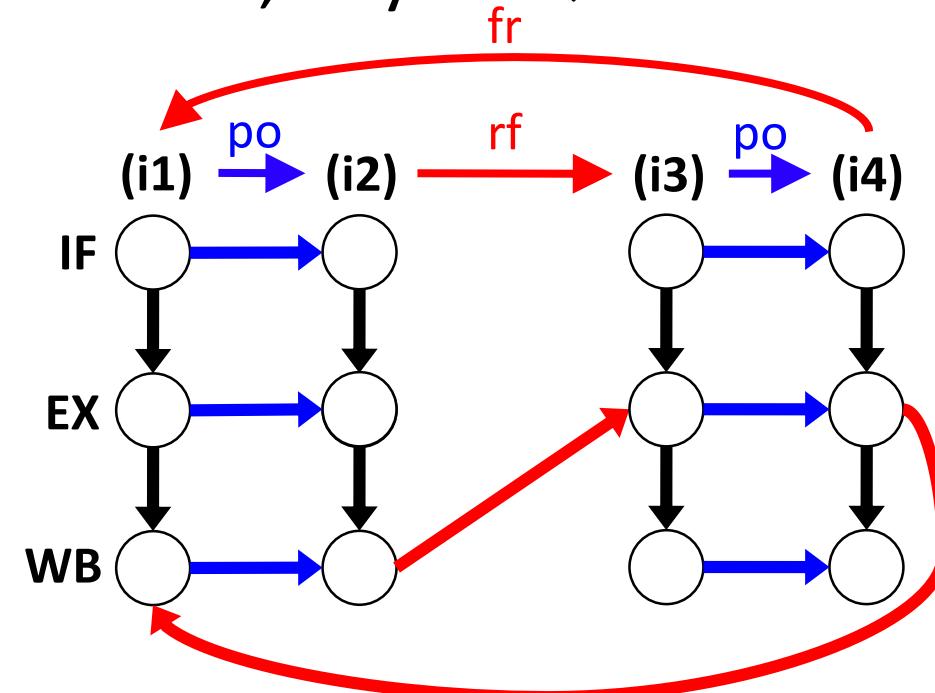
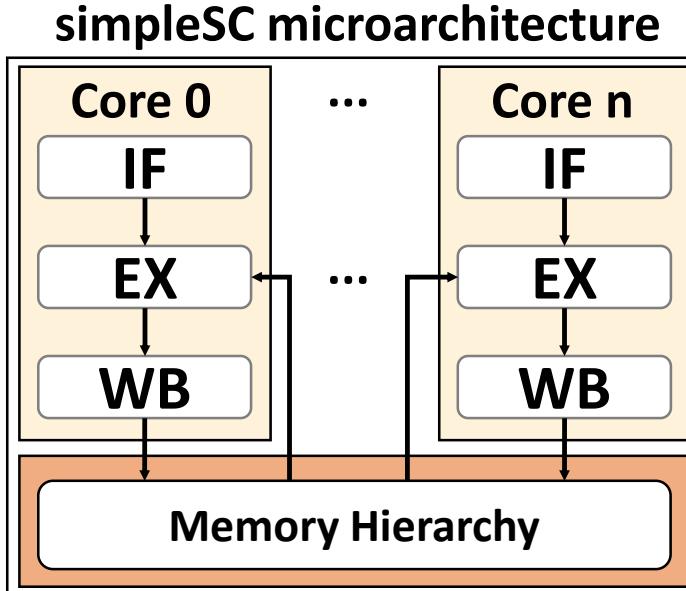
Legend:
IF = Fetch
EX = Execute
WB = Writeback

Microarchitectural Happens-Before (μ hb) Graphs

- Developed by PipeCheck [Lustig et al. MICRO 2014]
- Microarchitecture performs instrs. in stages
- Microarchitectural executions are μ hb graphs
 - **Nodes:** instr. sub-events, **edges:** happens-before relationships
- Cyclic μ hb graph → **unobservable**, Acyclic → **observable**

Message passing (mp) litmus test	
Core 0	Core 1
(i1) $x = 1;$	(i3) $r1 = y;$
(i2) $y = 1;$	(i4) $r2 = x;$

SC Forbids: $r1 = 1, r2 = 0$

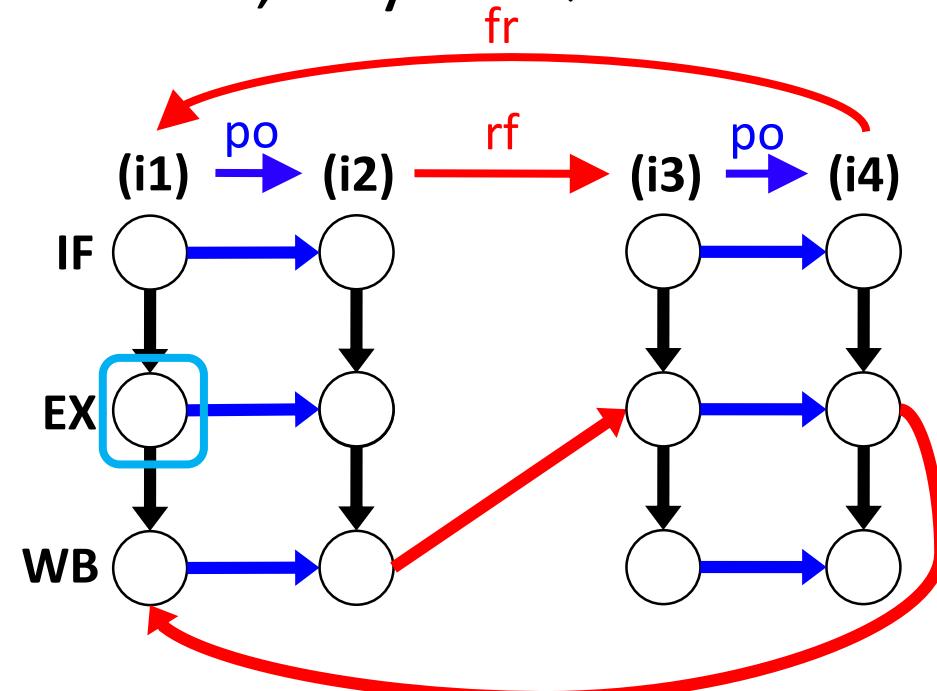
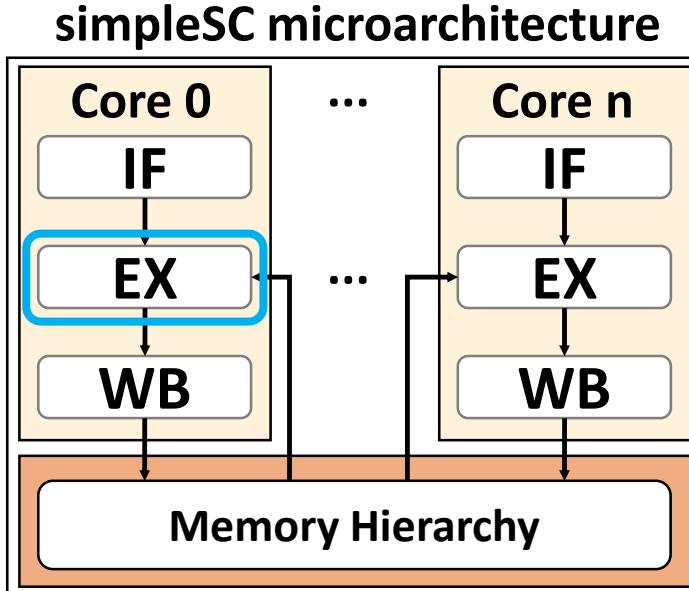


Microarchitectural Happens-Before (μ hb) Graphs

- Developed by PipeCheck [Lustig et al. MICRO 2014]
- Microarchitecture performs instrs. in stages
- Microarchitectural executions are μ hb graphs
 - **Nodes:** instr. sub-events, **edges:** happens-before relationships
- Cyclic μ hb graph → **unobservable**, Acyclic → **observable**

Message passing (mp) litmus test	
Core 0	Core 1
(i1) $x = 1;$	(i3) $r1 = y;$
(i2) $y = 1;$	(i4) $r2 = x;$

SC Forbids: $r1 = 1, r2 = 0$

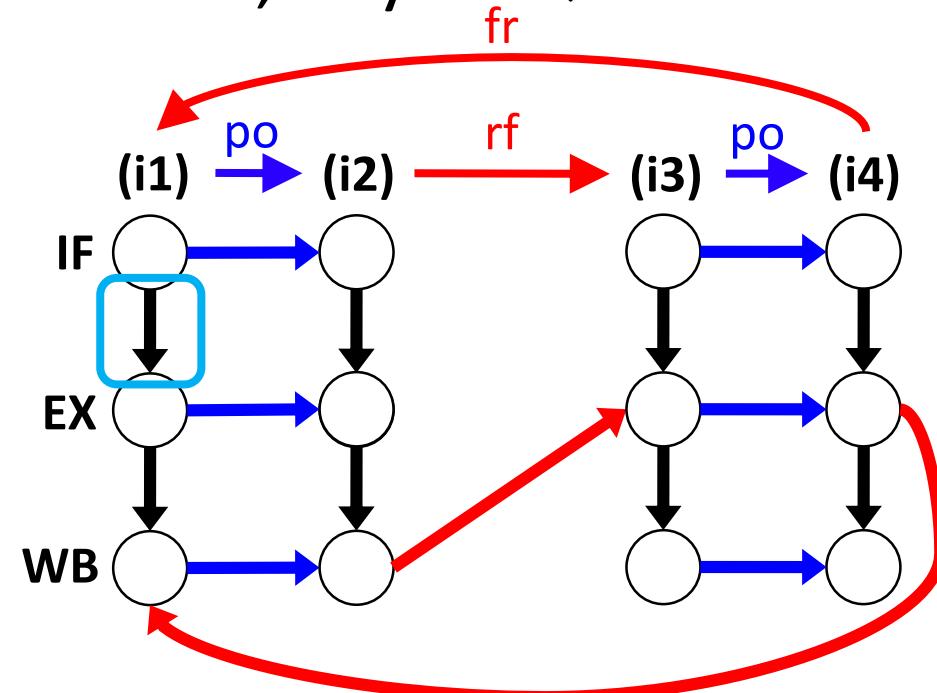
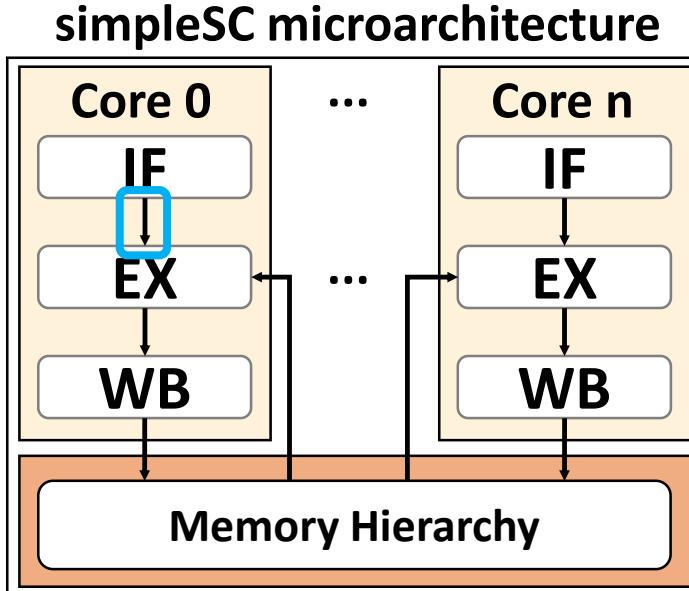


Microarchitectural Happens-Before (μ hb) Graphs

- Developed by PipeCheck [Lustig et al. MICRO 2014]
- Microarchitecture performs instrs. in stages
- Microarchitectural executions are μ hb graphs
 - **Nodes:** instr. sub-events, **edges:** happens-before relationships
- Cyclic μ hb graph → **unobservable**, Acyclic → **observable**

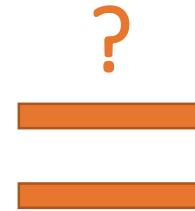
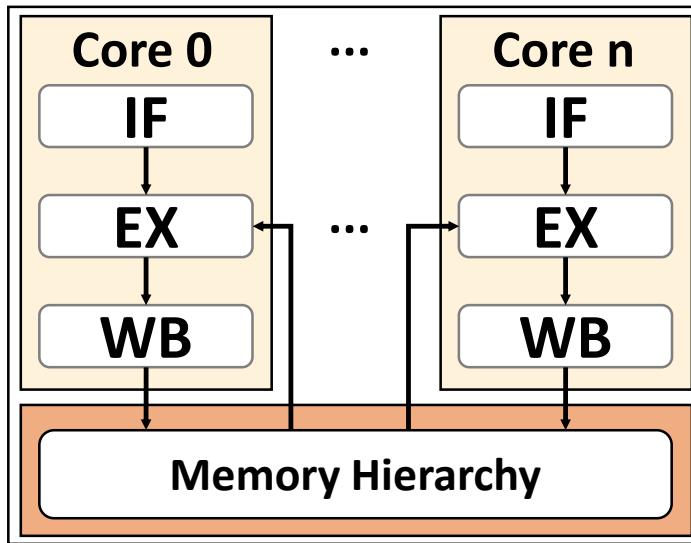
Message passing (mp) litmus test	
Core 0	Core 1
(i1) $x = 1;$	(i3) $r1 = y;$
(i2) $y = 1;$	(i4) $r2 = x;$

SC Forbids: $r1 = 1, r2 = 0$



Microarchitectural MCM Verification

Microarchitecture



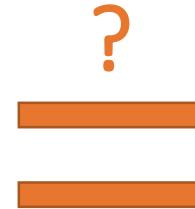
SC/TSO/RISC-V MCM?

Microarchitectural MCM Verification

Microarchitecture Specification in μ Spec DSL

```
Axiom "PO_Fetch":  
forall microops "i1",  
forall microops "i2",  
SameCore i1 i2 /\ ProgramOrder i1 i2 =>  
  AddEdge ((i1, Fetch), (i2, Fetch), "PO").
```

```
Axiom "Execute_stage_is_in_order":  
forall microops "i1",  
...
```



SC/TSO/RISC-V MCM?

- μ Spec DSL [Lustig et al. ASPLOS 2016] is similar to first-order logic (FOL)
 - `forall`, `exists`, AND ($/\$), OR (\vee), NOT (\sim), implication (\Rightarrow)
 - Has built-in predicates which take memory operations as input
 - e.g. `ProgramOrder i j` where `i` and `j` are loads/stores
 - Predicates can reference nodes and edges (μ hb edges closed under transitivity)
 - e.g. `EdgeExists ((i1, Fetch), (i2, Fetch))`



PipeProof: Automated All-Program MCM Verif.

High-Level Languages (HLL)

Compiler

Instruction Set (ISA)

Microarchitecture

Processor RTL (Verilog)

- PipeProof verifies that a microarchitecture correctly respects its ISA MCM across all possible programs
 - Early-stage **design-time verification** (i.e. before RTL)

Microarch. and
ISA MCM Specs

Aux. Inputs
(e.g. Mappings)

PipeProof

All-Program MCM
Correctness Proof!

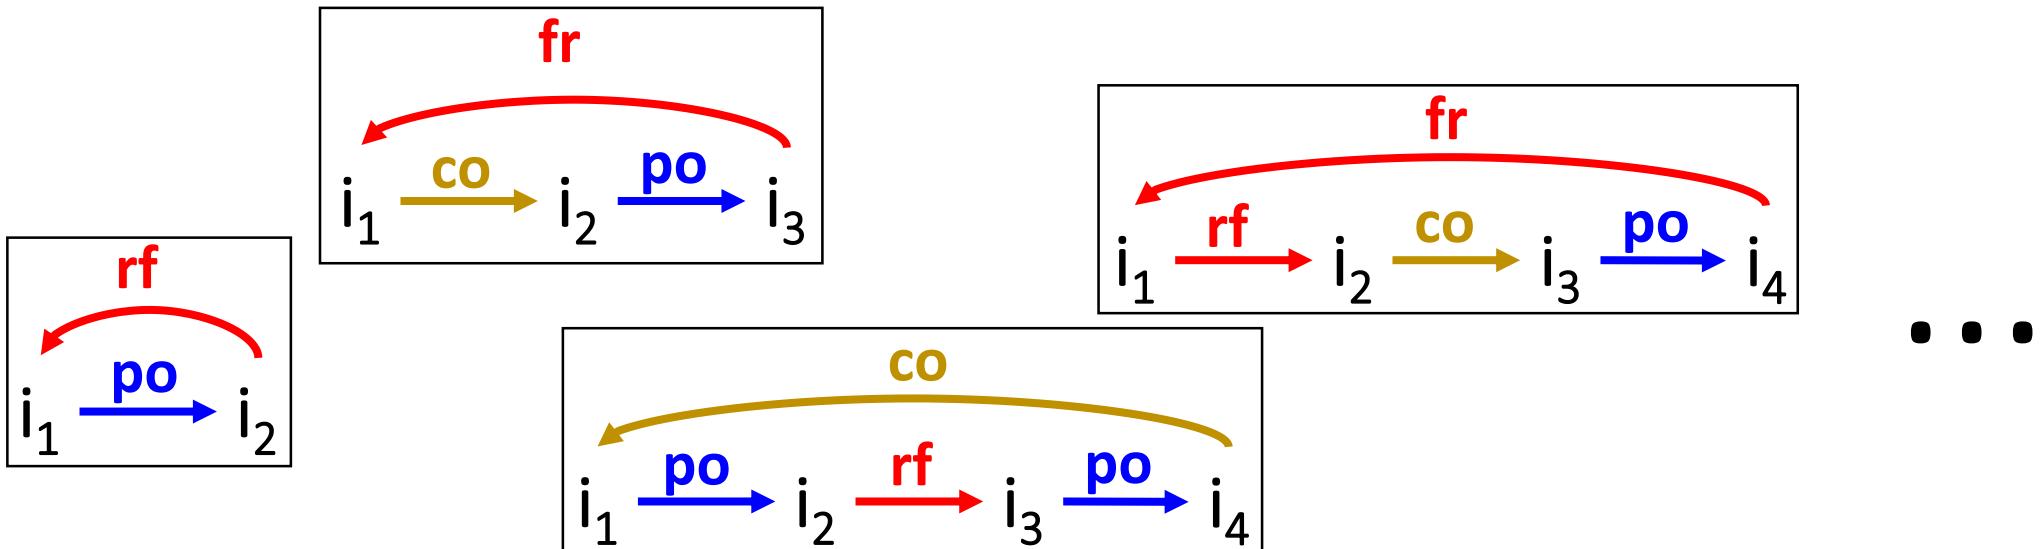
Verifying Across All Possible Programs

- Are all forbidden programs microarchitecturally unobservable?
 - If so, then microarchitecture is correct
- **Infinite** number of forbidden programs
 - E.g.: For SC, must check all possibilities of $cyclic(po \cup co \cup rf \cup fr)$
- Prove using **abstractions and induction**
 - Based on Counterexample-guided abstraction refinement [Clarke et al. CAV 2000]



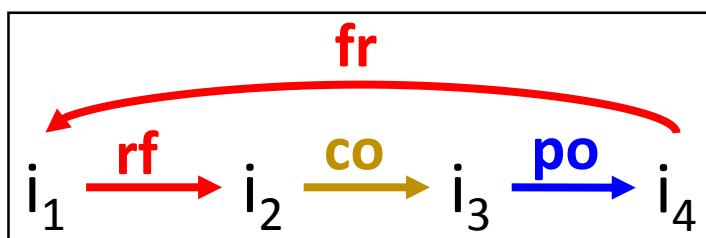
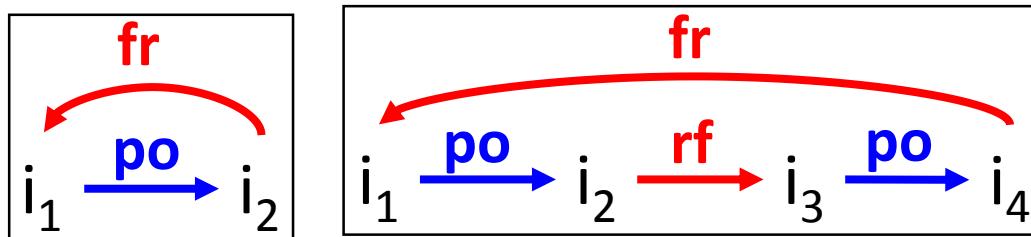
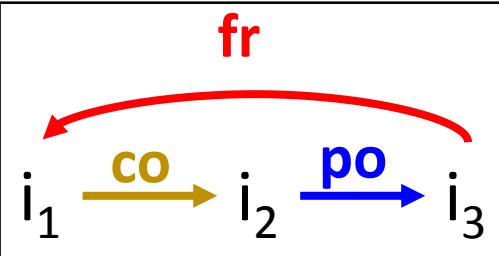
Verifying Across All Possible Programs

- Are all forbidden programs microarchitecturally unobservable?
 - If so, then microarchitecture is correct
- **Infinite** number of forbidden programs
 - E.g.: For SC, must check all possibilities of $cyclic(po \cup co \cup rf \cup fr)$
- Prove using **abstractions and induction**
 - Based on Counterexample-guided abstraction refinement [Clarke et al. CAV 2000]



The Transitive Chain (TC) Abstraction

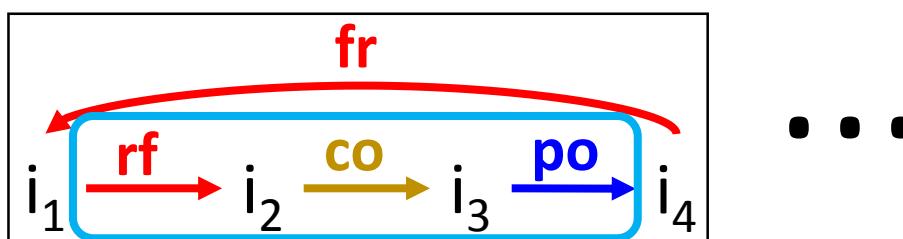
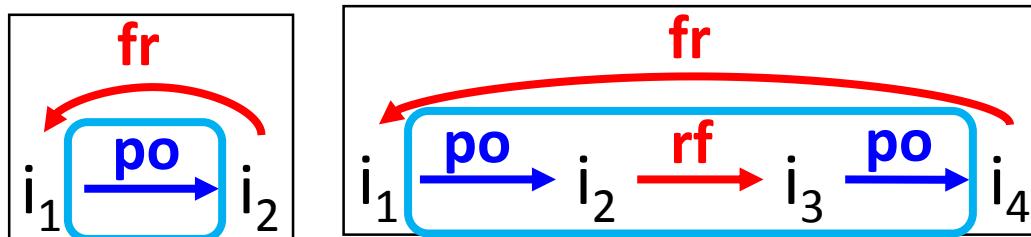
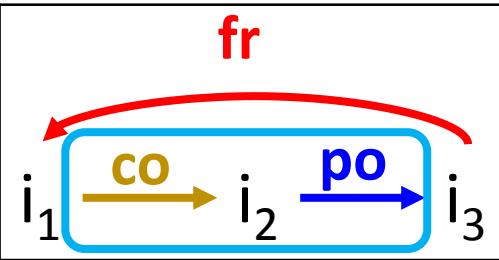
All non-unary cycles containing **fr**
(Infinite set)



• • •

The Transitive Chain (TC) Abstraction

All non-unary cycles containing **fr**
(Infinite set)

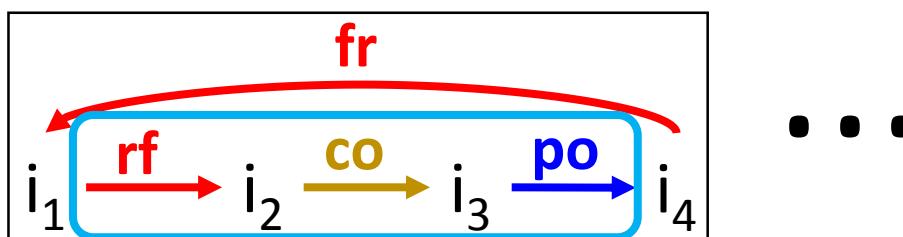
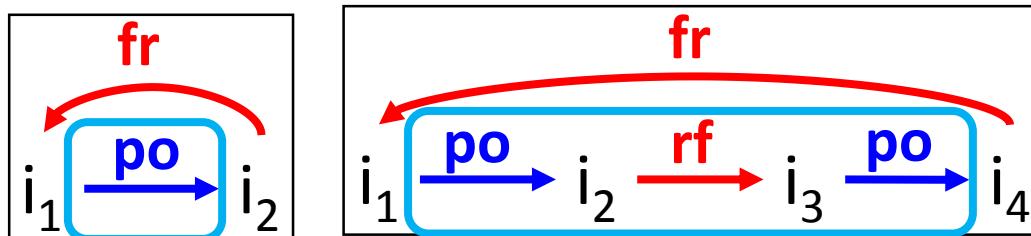
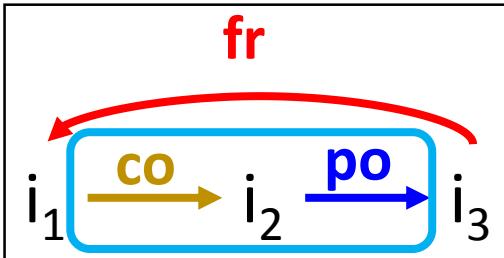


• • •

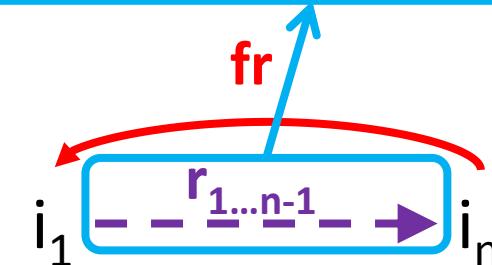
**Cycle = Transitive Chain (sequence)
+ Loopback edge (fr)**

The Transitive Chain (TC) Abstraction

All non-unary cycles containing fr
(Infinite set)



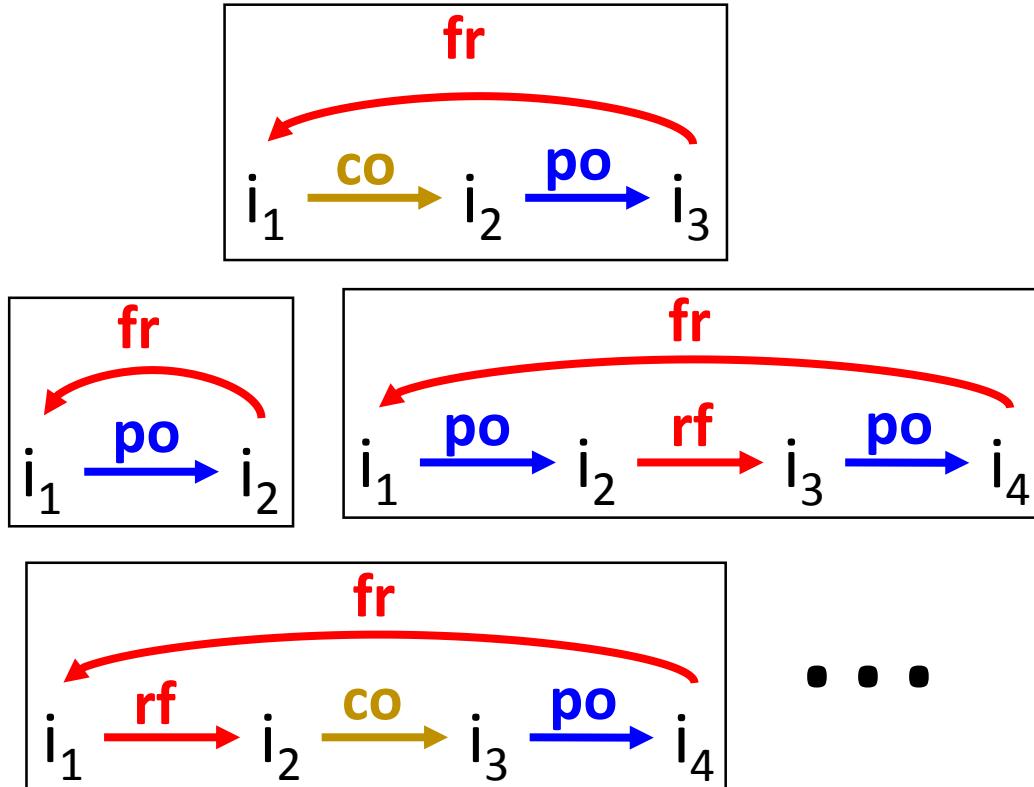
Transitive chain (sequence)
of ISA-level edges



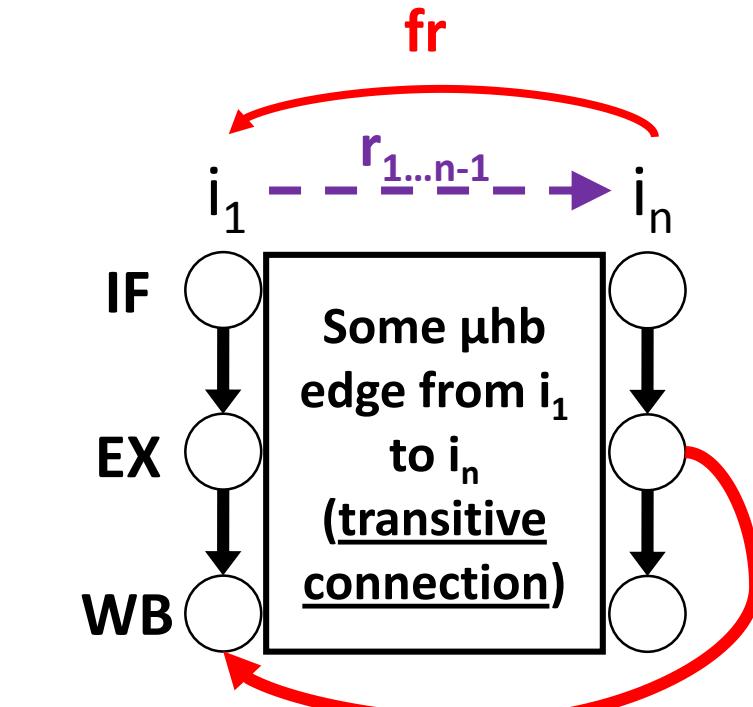
**Cycle = Transitive Chain (sequence)
+ Loopback edge (fr)**

The Transitive Chain (TC) Abstraction

All non-unary cycles containing fr
(Infinite set)



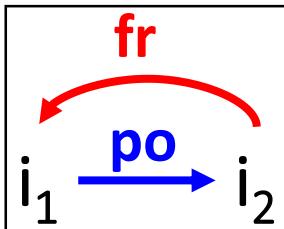
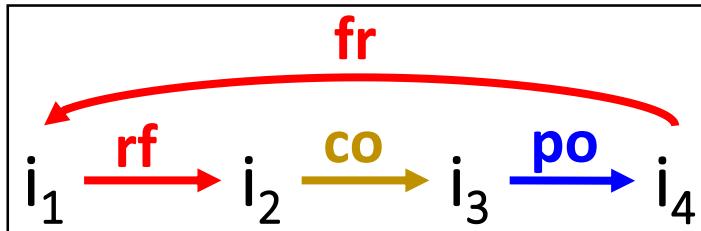
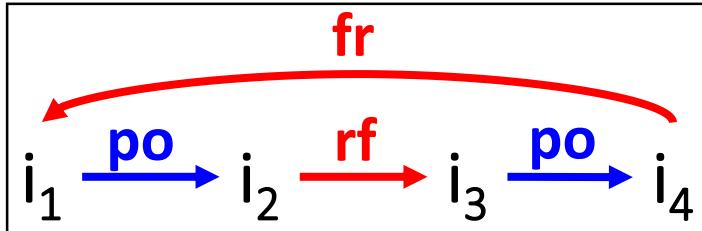
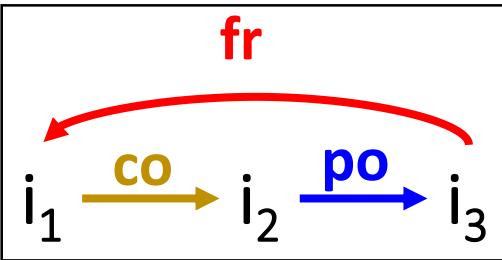
Cycle = Transitive Chain (sequence)
+ Loopback edge (fr)



ISA-level transitive chain =>
Microarch. level transitive connection

The Transitive Chain (TC) Abstraction

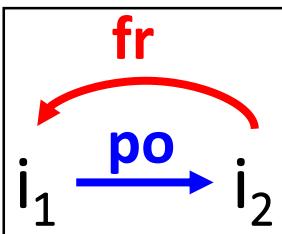
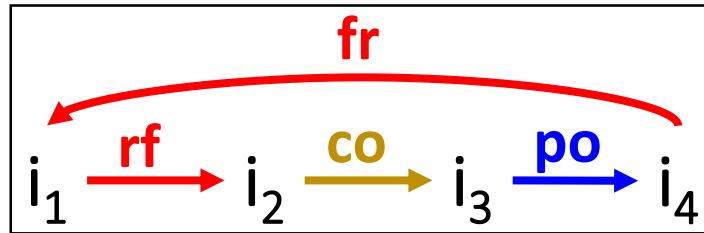
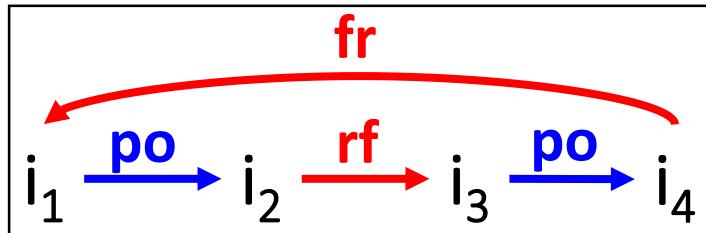
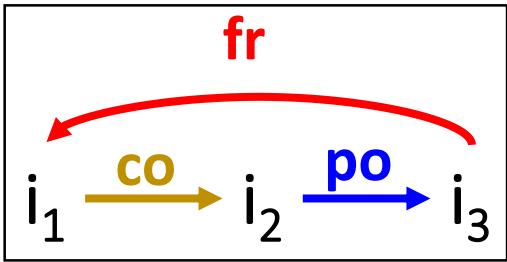
Infinite!



• • •

The Transitive Chain (TC) Abstraction

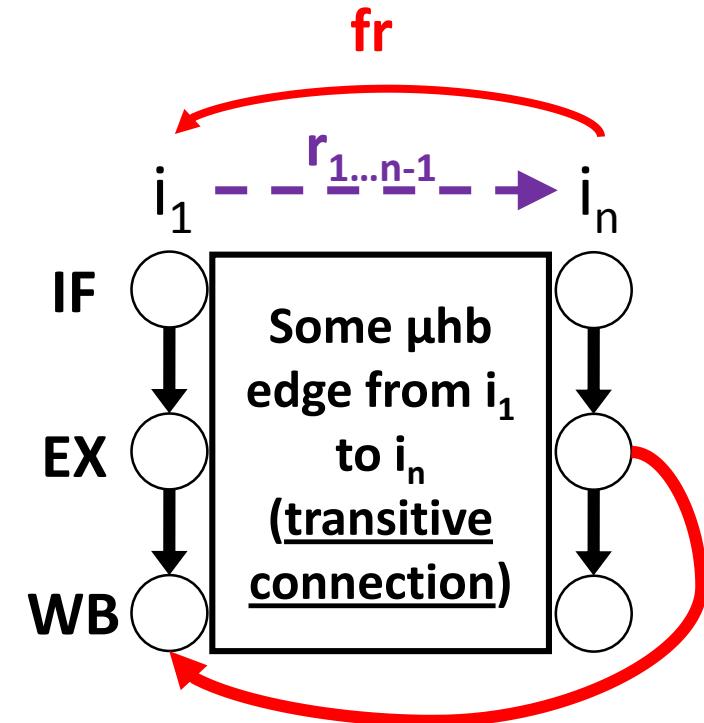
Infinite!



Using
TC Abstraction



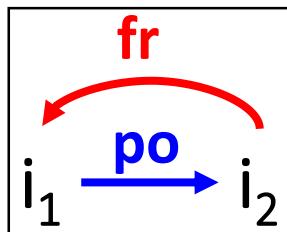
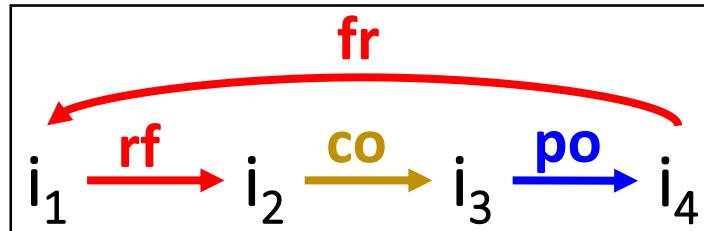
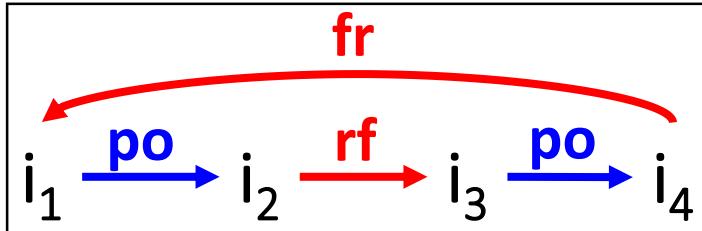
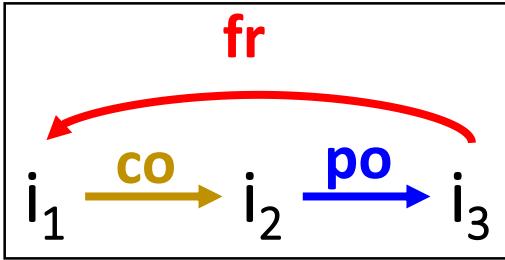
Finite!



3 x 3 = 9 possible
transitive connections
from i_1 to i_n

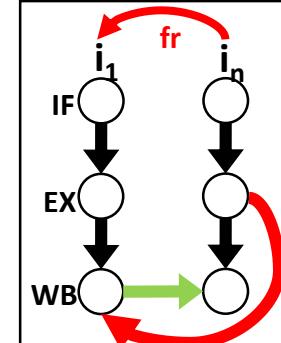
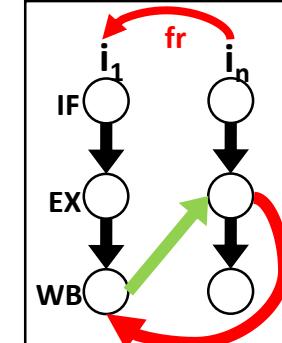
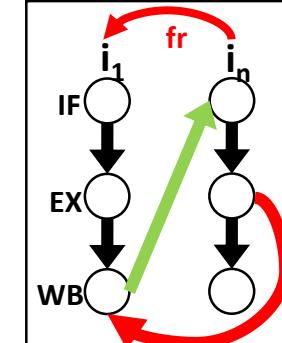
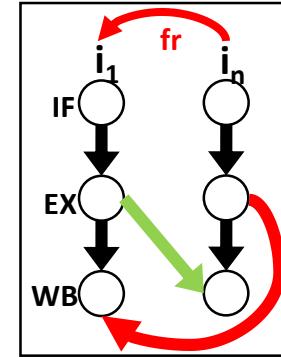
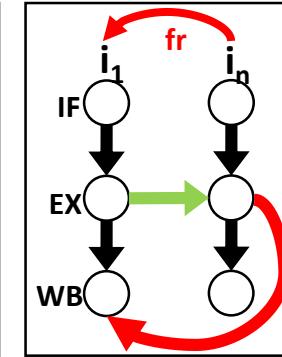
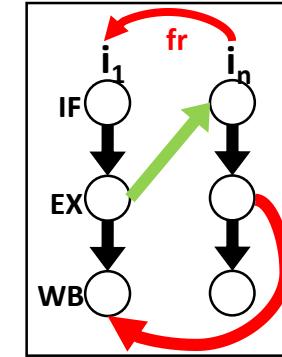
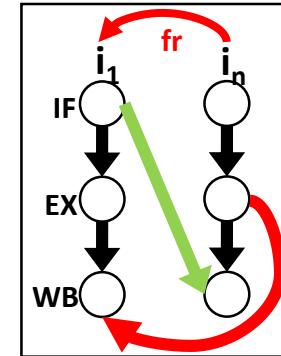
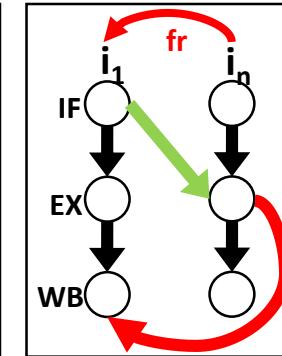
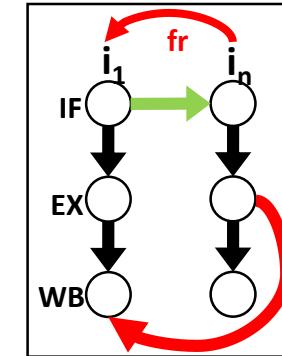
The Transitive Chain (TC) Abstraction

Infinite!



...

Finite!

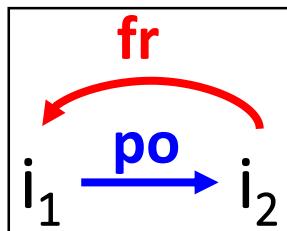
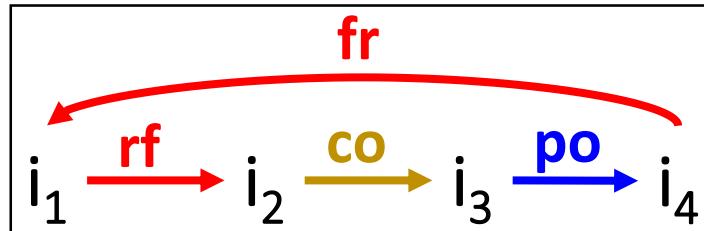
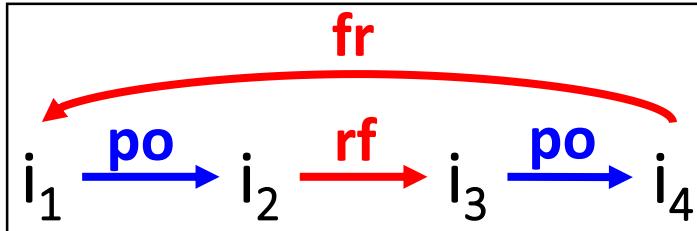
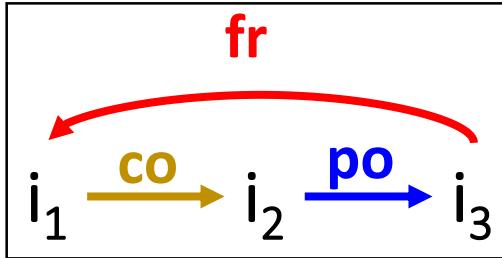


Using
TC Abstraction



The Transitive Chain (TC) Abstraction

Infinite!



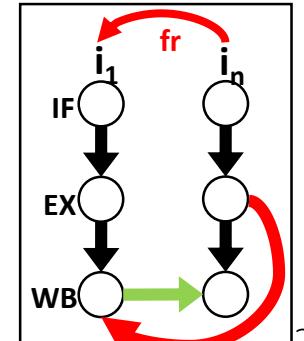
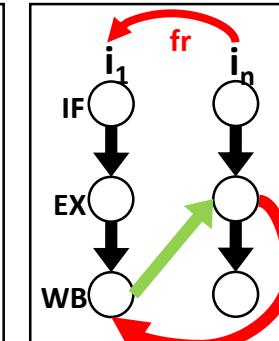
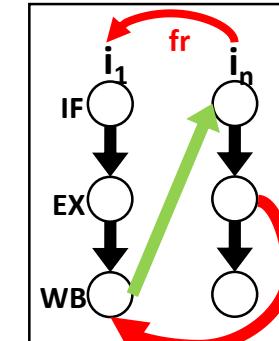
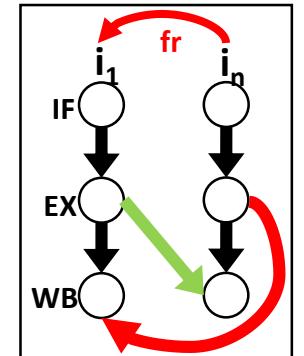
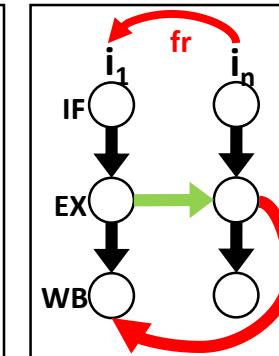
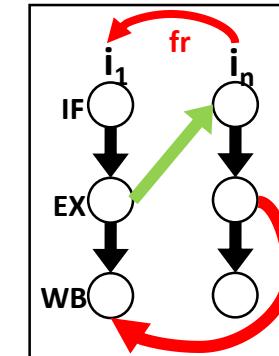
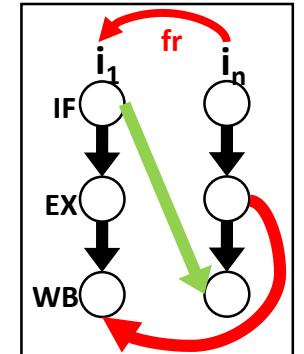
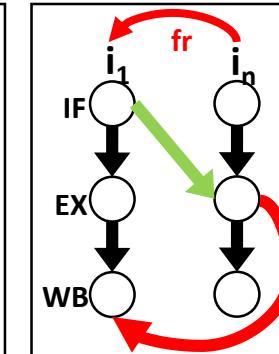
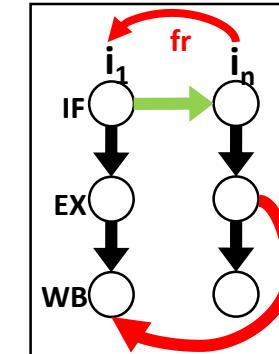
...

Abstraction soundness
automatically verified
as a supporting proof!

Using
TC Abstraction

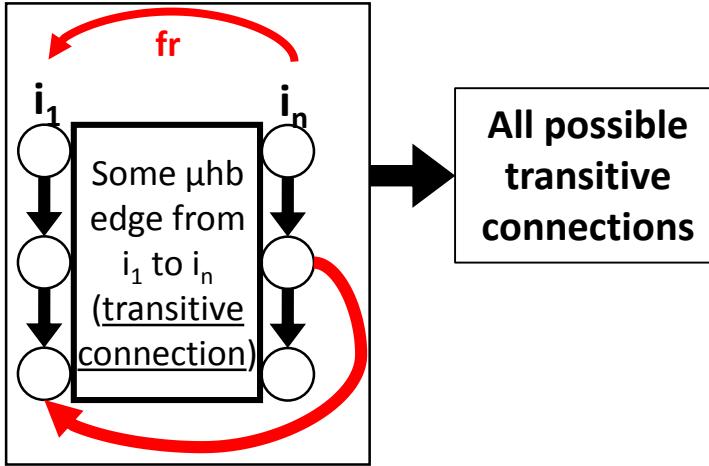


Finite!

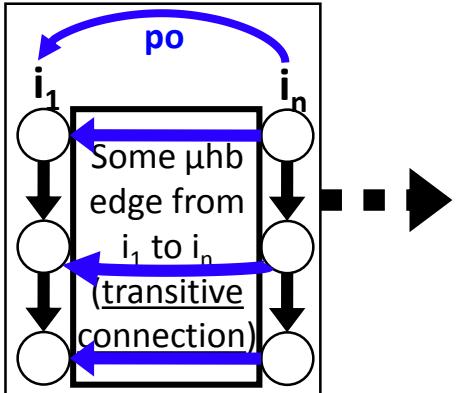


Microarchitectural Correctness Proof

Cycles containing *fr*



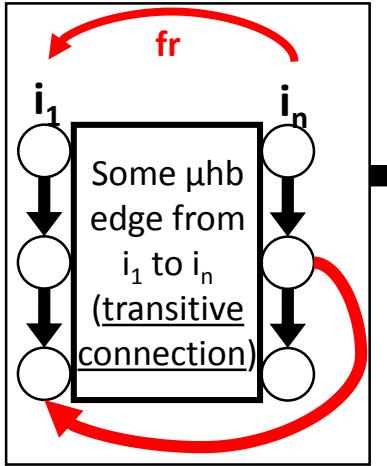
Cycles containing *po*



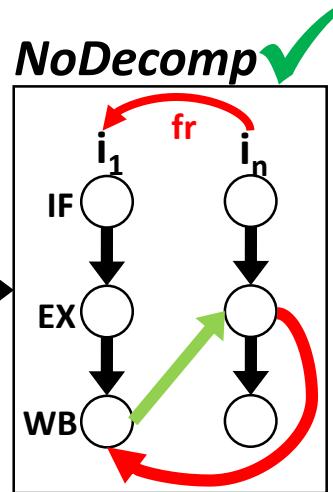
Other ISA-level
cycles...

Microarchitectural Correctness Proof

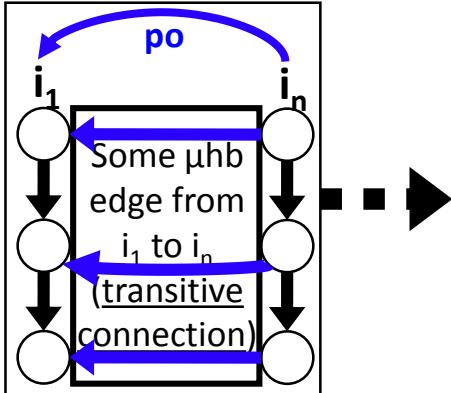
Cycles containing *fr*



All possible transitive connections



Cycles containing *po*

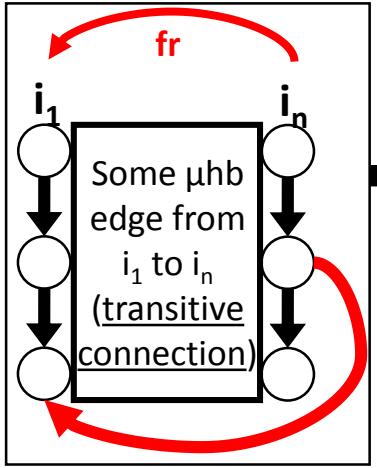


Other ISA-level cycles...

Other transitive connections...

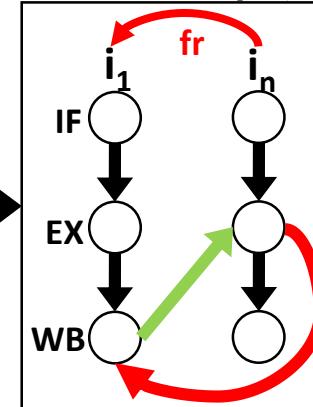
Microarchitectural Correctness Proof

Cycles containing *fr*

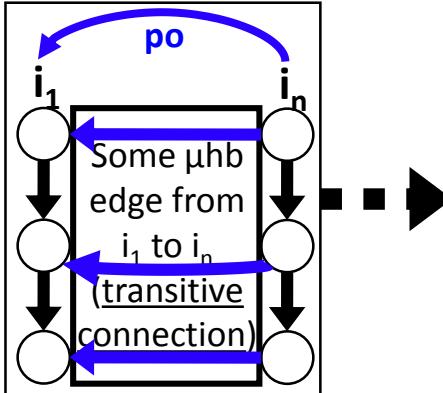


All possible transitive connections

NoDecomp ✓

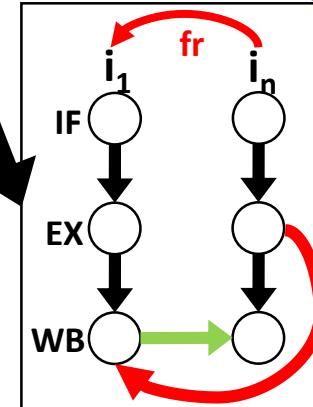


Cycles containing *po*



Other ISA-level cycles...

AbsCounterX?

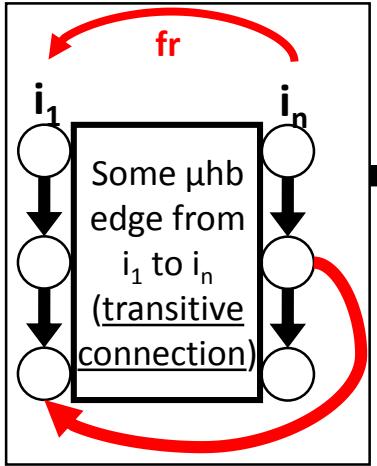


Acyclic graph with **transitive connection** =>
Abstract Counterexample (i.e. possible bug)

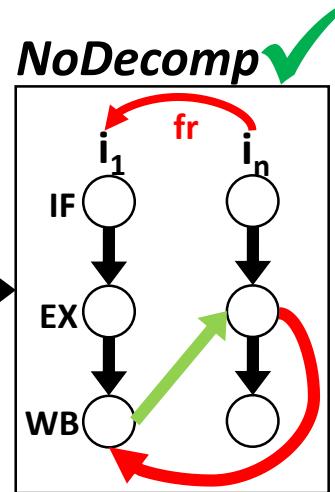
Other transitive connections...

Microarchitectural Correctness Proof

Cycles containing *fr*

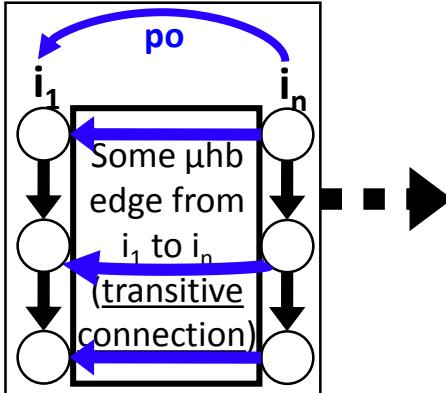


All possible transitive connections



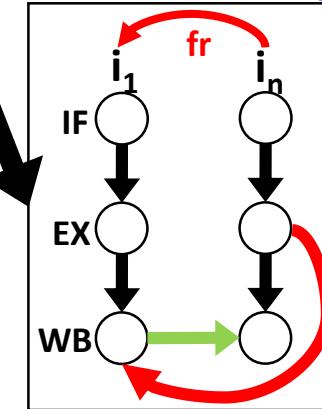
NoDecomp ✓
Transitive connection (green edge) may represent one or multiple ISA-level edges

Cycles containing *po*



Other ISA-level cycles...

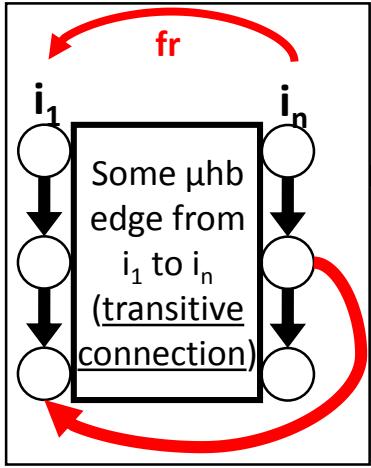
AbsCounterX?



Other transitive connections...

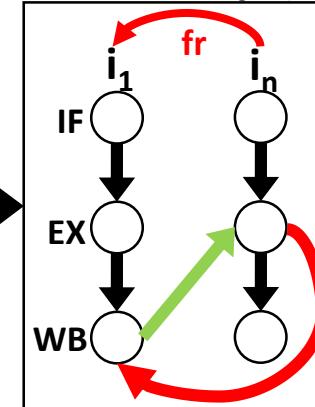
Microarchitectural Correctness Proof

Cycles containing *fr*



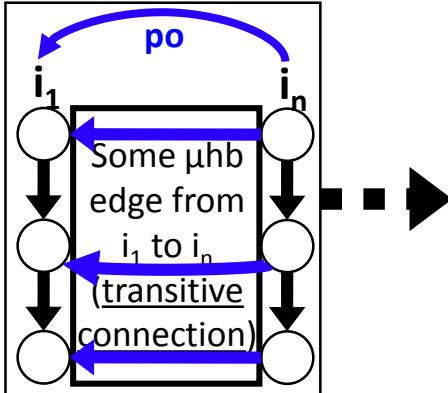
All possible transitive connections

NoDecomp ✓



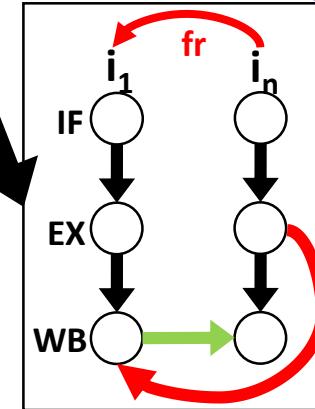
Transitive connection (green edge) may represent one or multiple ISA-level edges

Cycles containing *po*



Other ISA-level cycles...

AbsCounterX?



Try to Concretize (Replace transitive connection with one ISA-level edge)

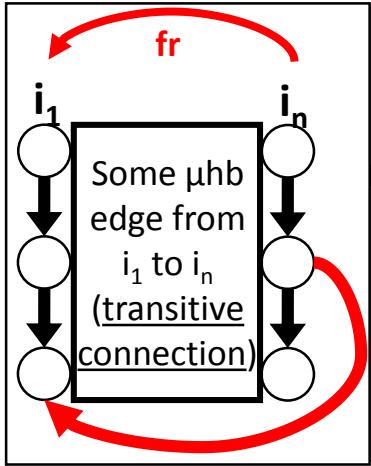
Observable

Other transitive connections...

Microarch Buggy,
Return Counterexample

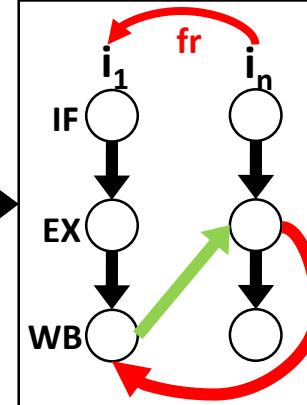
Microarchitectural Correctness Proof

Cycles containing *fr*



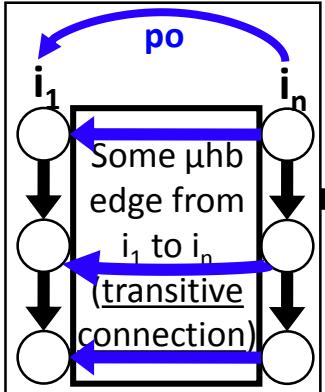
All possible transitive connections

NoDecomp ✓



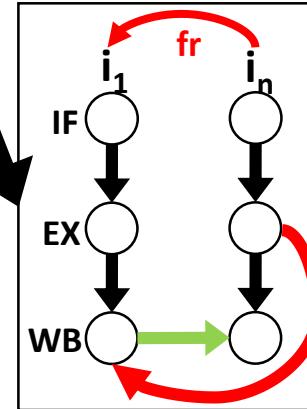
Transitive connection (green edge) may represent one or multiple ISA-level edges

Cycles containing *po*



Other ISA-level cycles...

AbsCounterX?



Try to Concretize (Replace transitive connection with one ISA-level edge)

Unobs.

Consider all Decompositions (Inductively break down Transitive Chain)

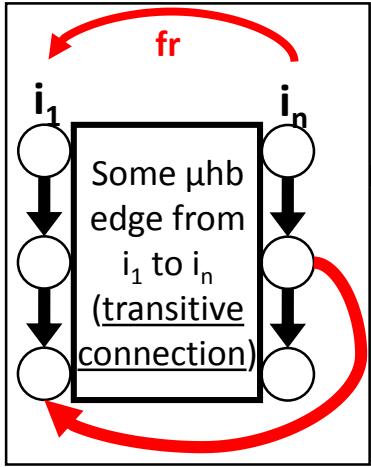
Observable

Other transitive connections...

Microarch Buggy, Return Counterexample

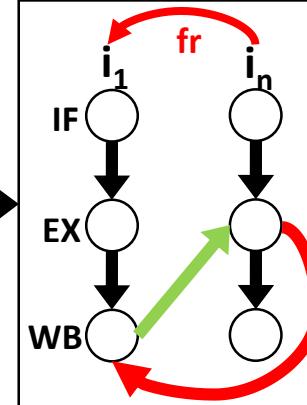
Microarchitectural Correctness Proof

Cycles containing *fr*



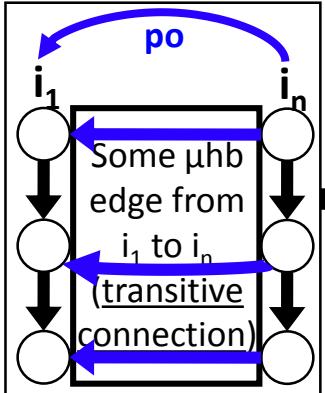
All possible transitive connections

NoDecomp ✓



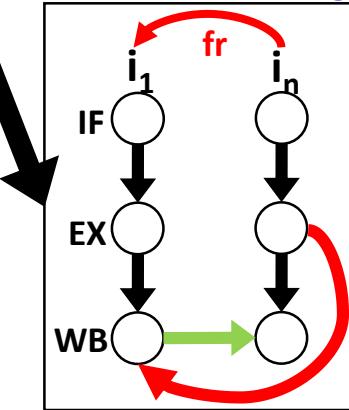
Transitive connection (green edge) may represent one or multiple ISA-level edges

Cycles containing *po*



"Refinement Loop"

AbsCounterX?



Try to Concretize (Replace transitive connection with one ISA-level edge)

Unobs.

Consider all Decompositions (Inductively break down Transitive Chain)

Observable

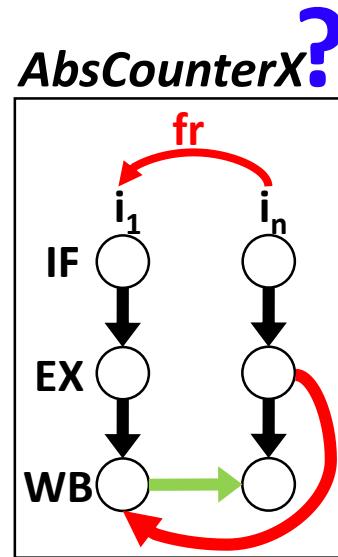
Other ISA-level cycles...

Other transitive connections...

Microarch Buggy, Return Counterexample

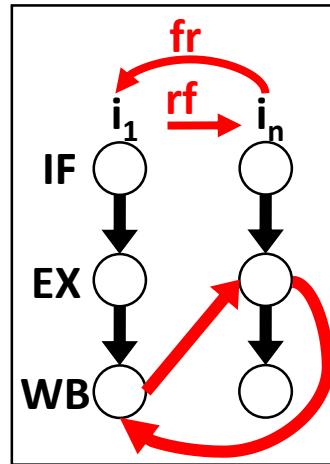
Refinement Loop: Concretization

- Replaces transitive connection with a single ISA-level edge
 - All concretizations must be unobservable
 - Observable concretizations are counterexamples (bugs)



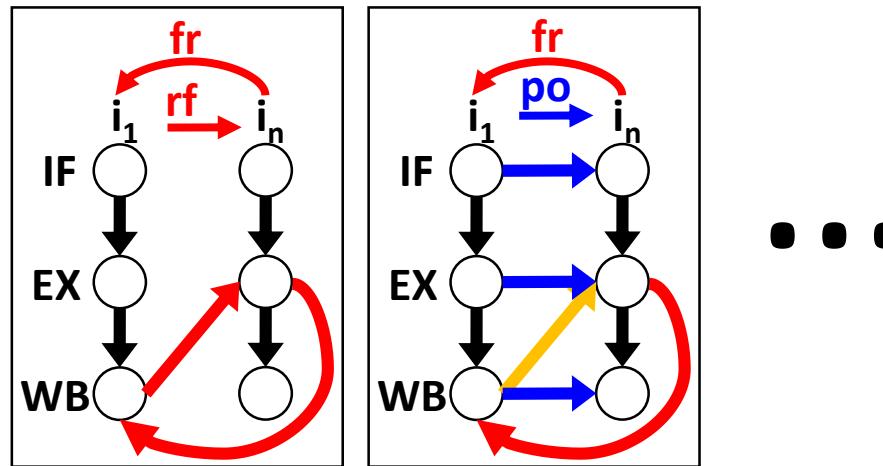
Refinement Loop: Concretization

- Replaces transitive connection with a single ISA-level edge
 - All concretizations must be unobservable
 - Observable concretizations are counterexamples (bugs)



Refinement Loop: Concretization

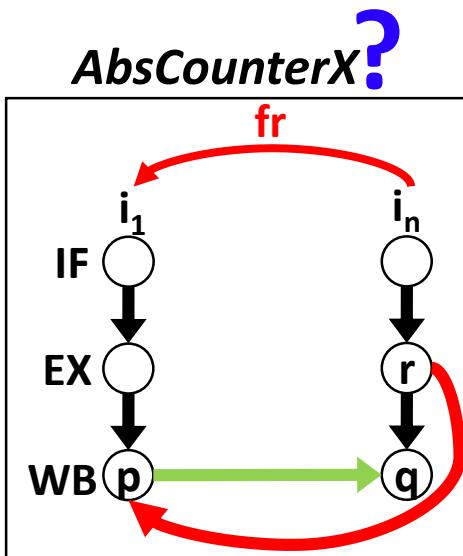
- Replaces transitive connection with a single ISA-level edge
 - All concretizations must be unobservable
 - Observable concretizations are counterexamples (bugs)



Refinement Loop: Decomposition

- Inductively break down transitive chain
 - Additional constraints may be enough to make execution unobservable

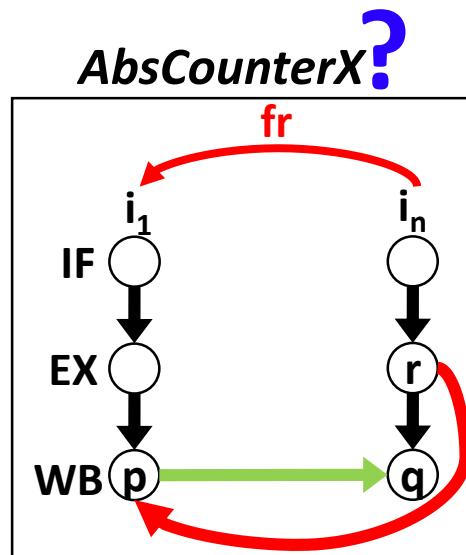
```
factorial(n) = factorial(n-1) * n
```



Refinement Loop: Decomposition

- Inductively break down transitive chain
 - Additional constraints may be enough to make execution unobservable

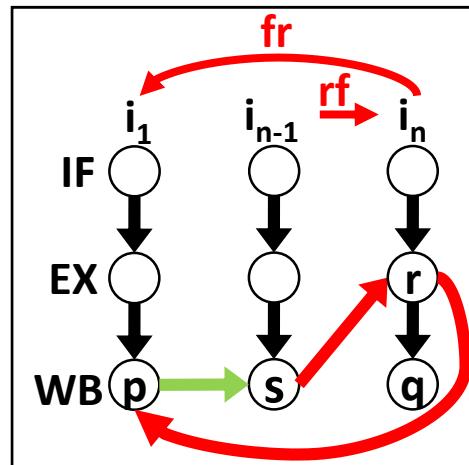
```
factorial(n)      = factorial(n-1) * n  
                   ;           ;  
Chain of length n = Chain of length n-1 + "Peeled-off" edge
```



Refinement Loop: Decomposition

- Inductively break down transitive chain
 - Additional constraints may be enough to make execution unobservable

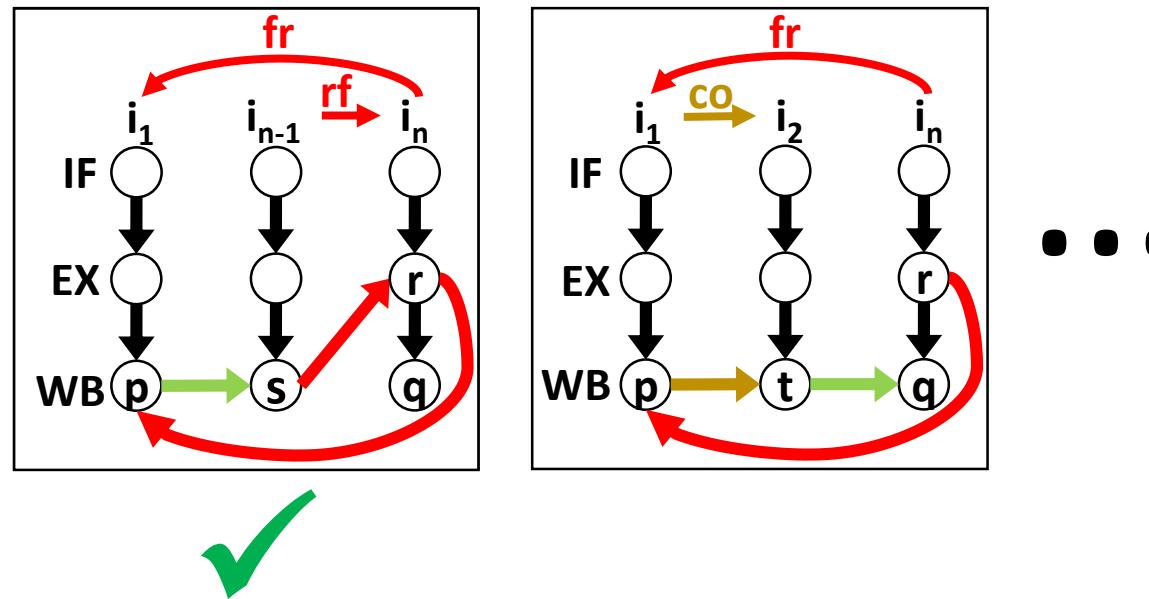
```
factorial(n) = factorial(n-1) * n  
; ; ;  
Chain of length n = Chain of length n-1 + "Peeled-off" edge
```



Refinement Loop: Decomposition

- Inductively break down transitive chain
 - Additional constraints may be enough to make execution unobservable

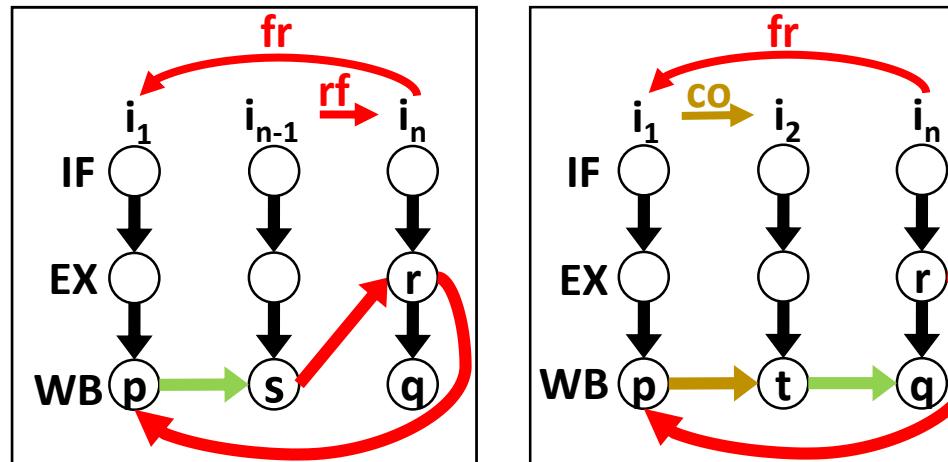
```
factorial(n)      = factorial(n-1) * n  
                   ;           ;  
Chain of length n = Chain of length n-1 + "Peeled-off" edge
```



Refinement Loop: Decomposition

- Inductively break down transitive chain
 - Additional constraints may be enough to make execution unobservable

```
factorial(n)      = factorial(n-1) * n  
                   ;           ;  
Chain of length n = Chain of length n-1 + "Peeled-off" edge
```



If decomposition is abstract
counterexample, repeat concretization
and decomposition!

Results

- Ran PipeProof on simpleSC (SC) and simpleTSO (TSO¹) μarches
 - 3-stage in-order pipelines
- TSO verification made feasible by optimizations
 - Explicitly checking all decompositions => **case explosion**
 - **Covering Sets Optimization** (eliminate redundant transitive connections)
 - **Memoization** (eliminate previously checked ISA-level cycles)

	simpleSC	simpleSC (w/ Covering Sets + Memoization)
Total Time	225.9 sec	19.1 sec
	simpleTSO	simpleTSO (w/ Covering Sets + Memoization)
Total Time	Timeout	2449.7 sec (\approx 41 mins)

¹TSO (Total Store Order) is the MCM of Intel x86 processors. It relaxes Store->Load ordering.

PipeProof Takeaways

- First Ever Automated **All-Program** Microarchitectural MCM Verification
 - Designers get both **completeness** and **automation** of verification
 - Engineers can verify microarchitectures themselves, before RTL is written!
- Based on techniques from formal methods (**CEGAR**) [Clarke et al. CAV 2000]
- **Transitive Chain (TC) Abstraction** models infinite set of executions
- Accolades:
 - Nominated for Best Paper at MICRO 2018
 - “Honorable Mention” in 2018 IEEE Micro *Top Picks of Comp. Arch. Conferences*

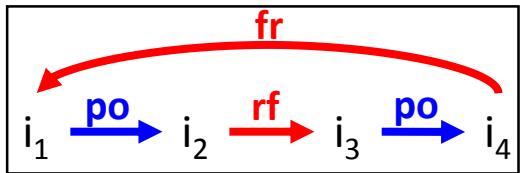
Talk Outline

- Overview and Motivation
- Memory Consistency Background
- **PipeProof:** All-Program Microarchitectural MCM Verification
- **RTLCheck:** MCM Verification of Verilog RTL
- Expanding to other domains
- Conclusion

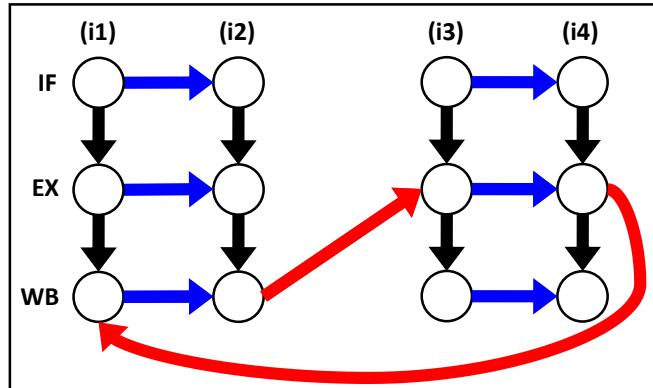


What if I want to verify RTL (Verilog)?

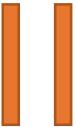
ISA-Level MCM



Microarchitectural Orderings



acyclic (po U co U rf U fr)

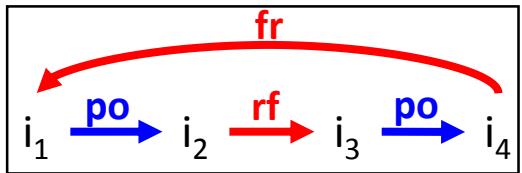


Axiom "PO_Fetch":
forall microop "i1", "i2",
SameCore i1 i2 /\ ProgramOrder i1 i2 =>
AddEdge ((i1, IF), (i2, IF)).
...

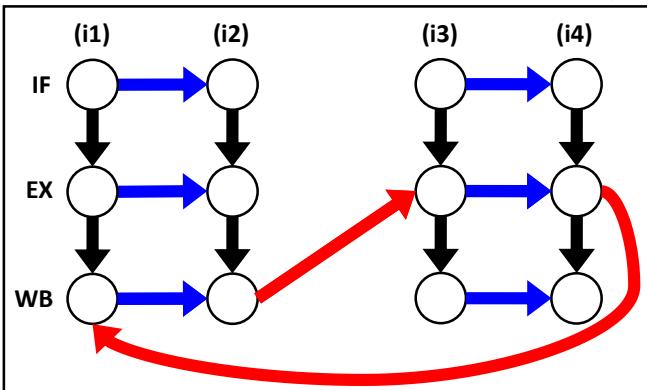
Verified with
PipeProof

What if I want to verify RTL (Verilog)?

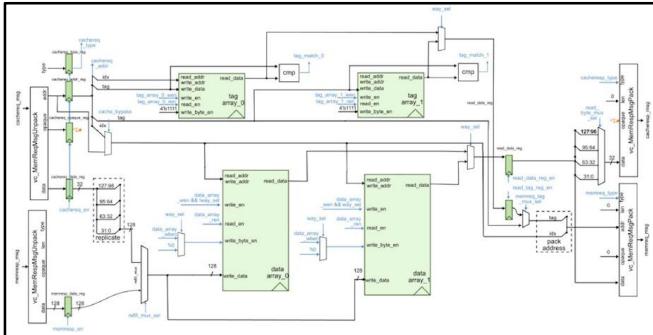
ISA-Level MCM



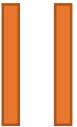
Microarchitectural Orderings



RTL implementation (Verilog)



acyclic ($po \cup co \cup rf \cup fr$)



Axiom "PO_Fetch":
forall microop "i1", "i2",
SameCore i1 i2 /\ ProgramOrder i1 i2 =>
AddEdge ((i1, IF), (i2, IF)).
...

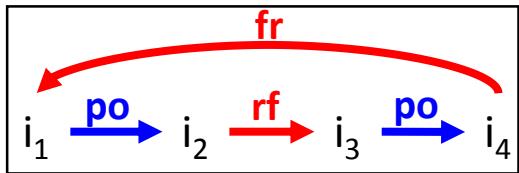


Verified with
PipeProof

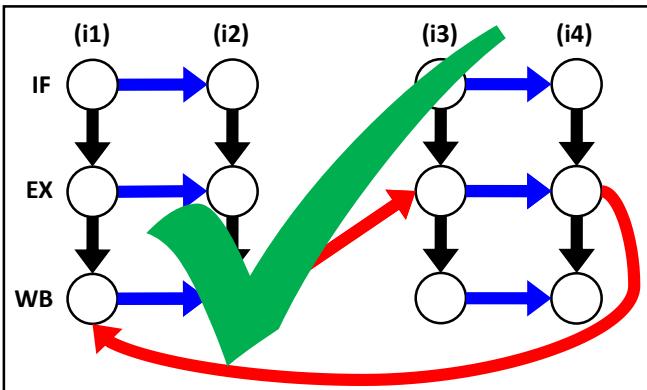


What if I want to verify RTL (Verilog)?

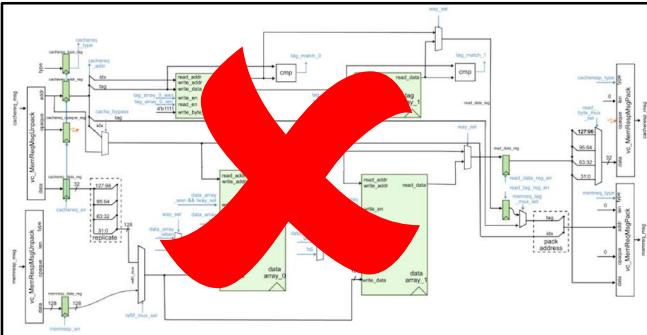
ISA-Level MCM



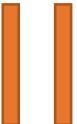
Microarchitectural Orderings



RTL implementation (Verilog)



acyclic ($po \cup co \cup rf \cup fr$)



Axiom "PO_Fetch":
forall microop "i1", "i2",
SameCore i1 i2 /\ ProgramOrder i1 i2 =>
AddEdge ((i1, IF), (i2, IF)).
...

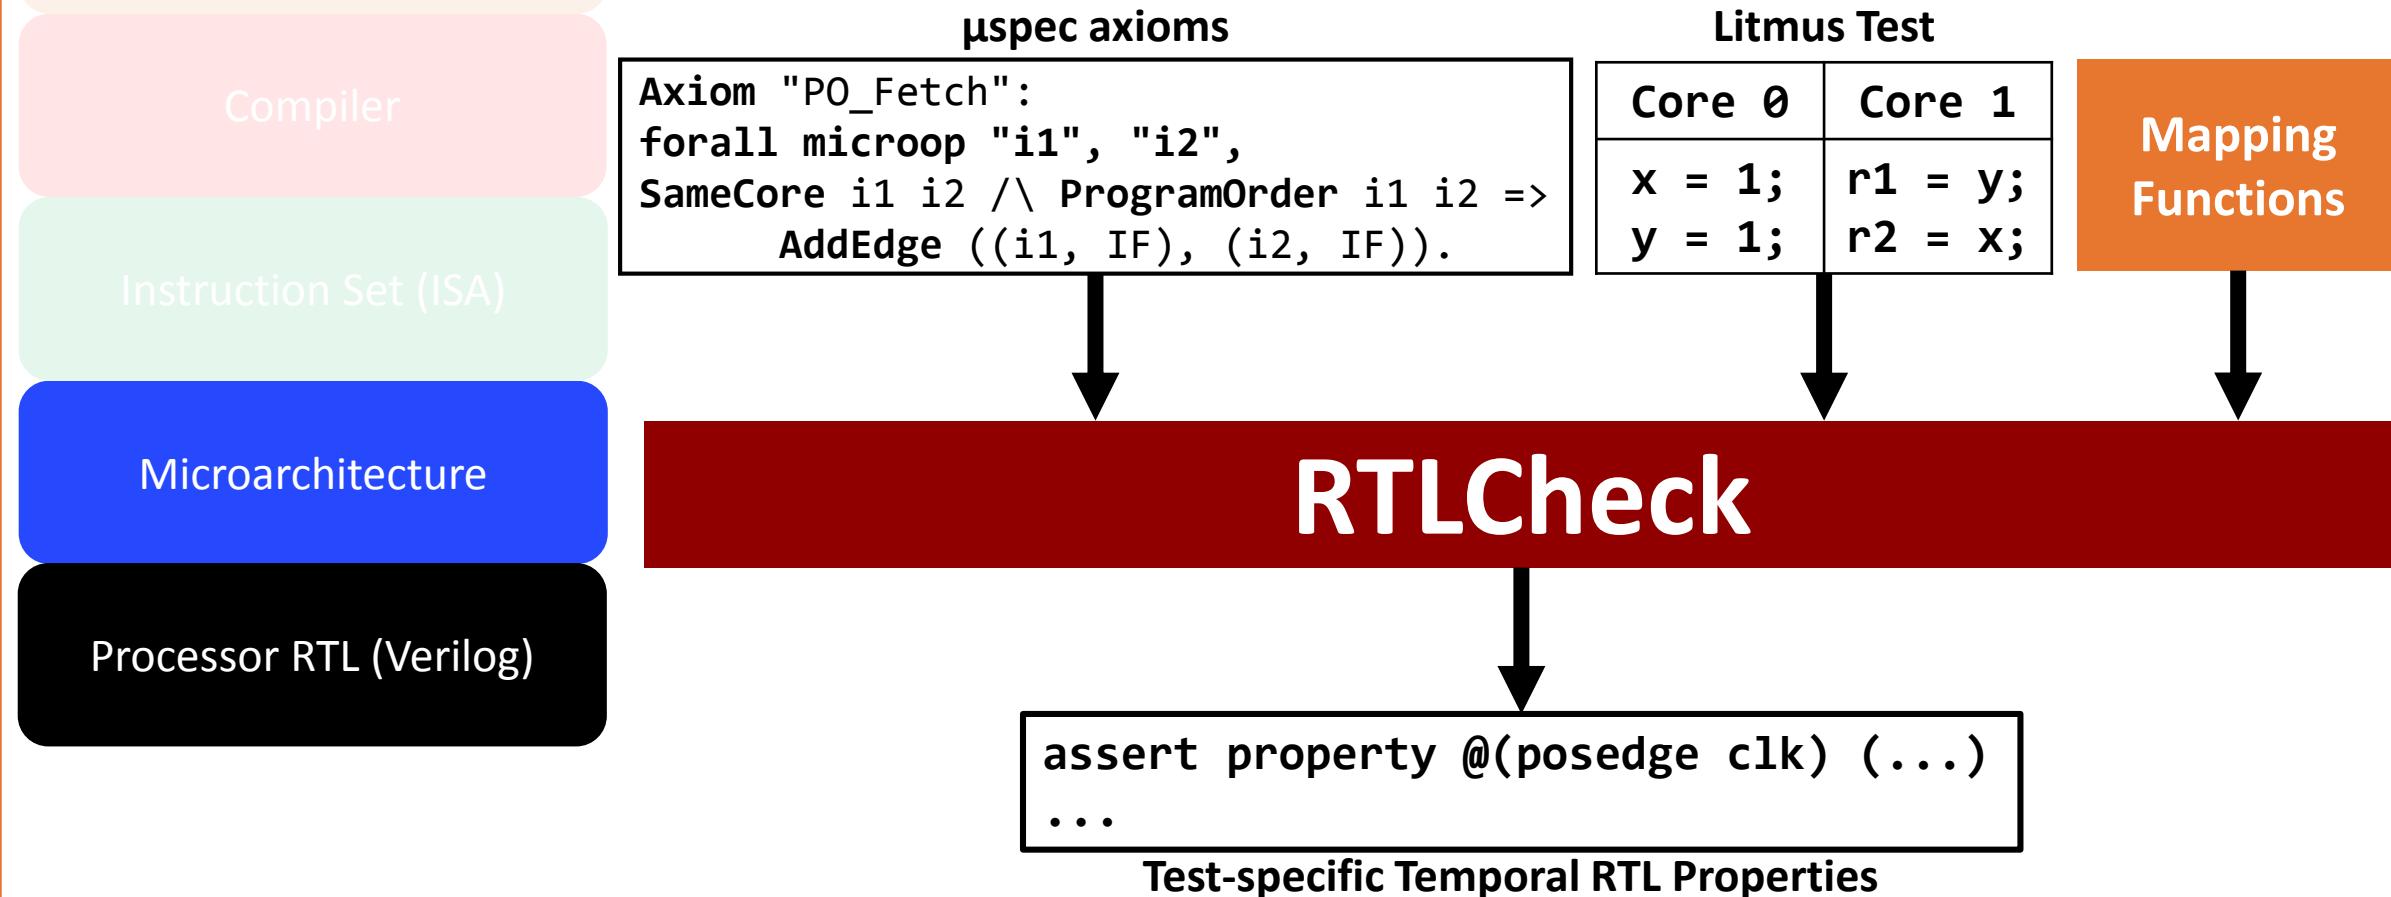


Verified with
PipeProof



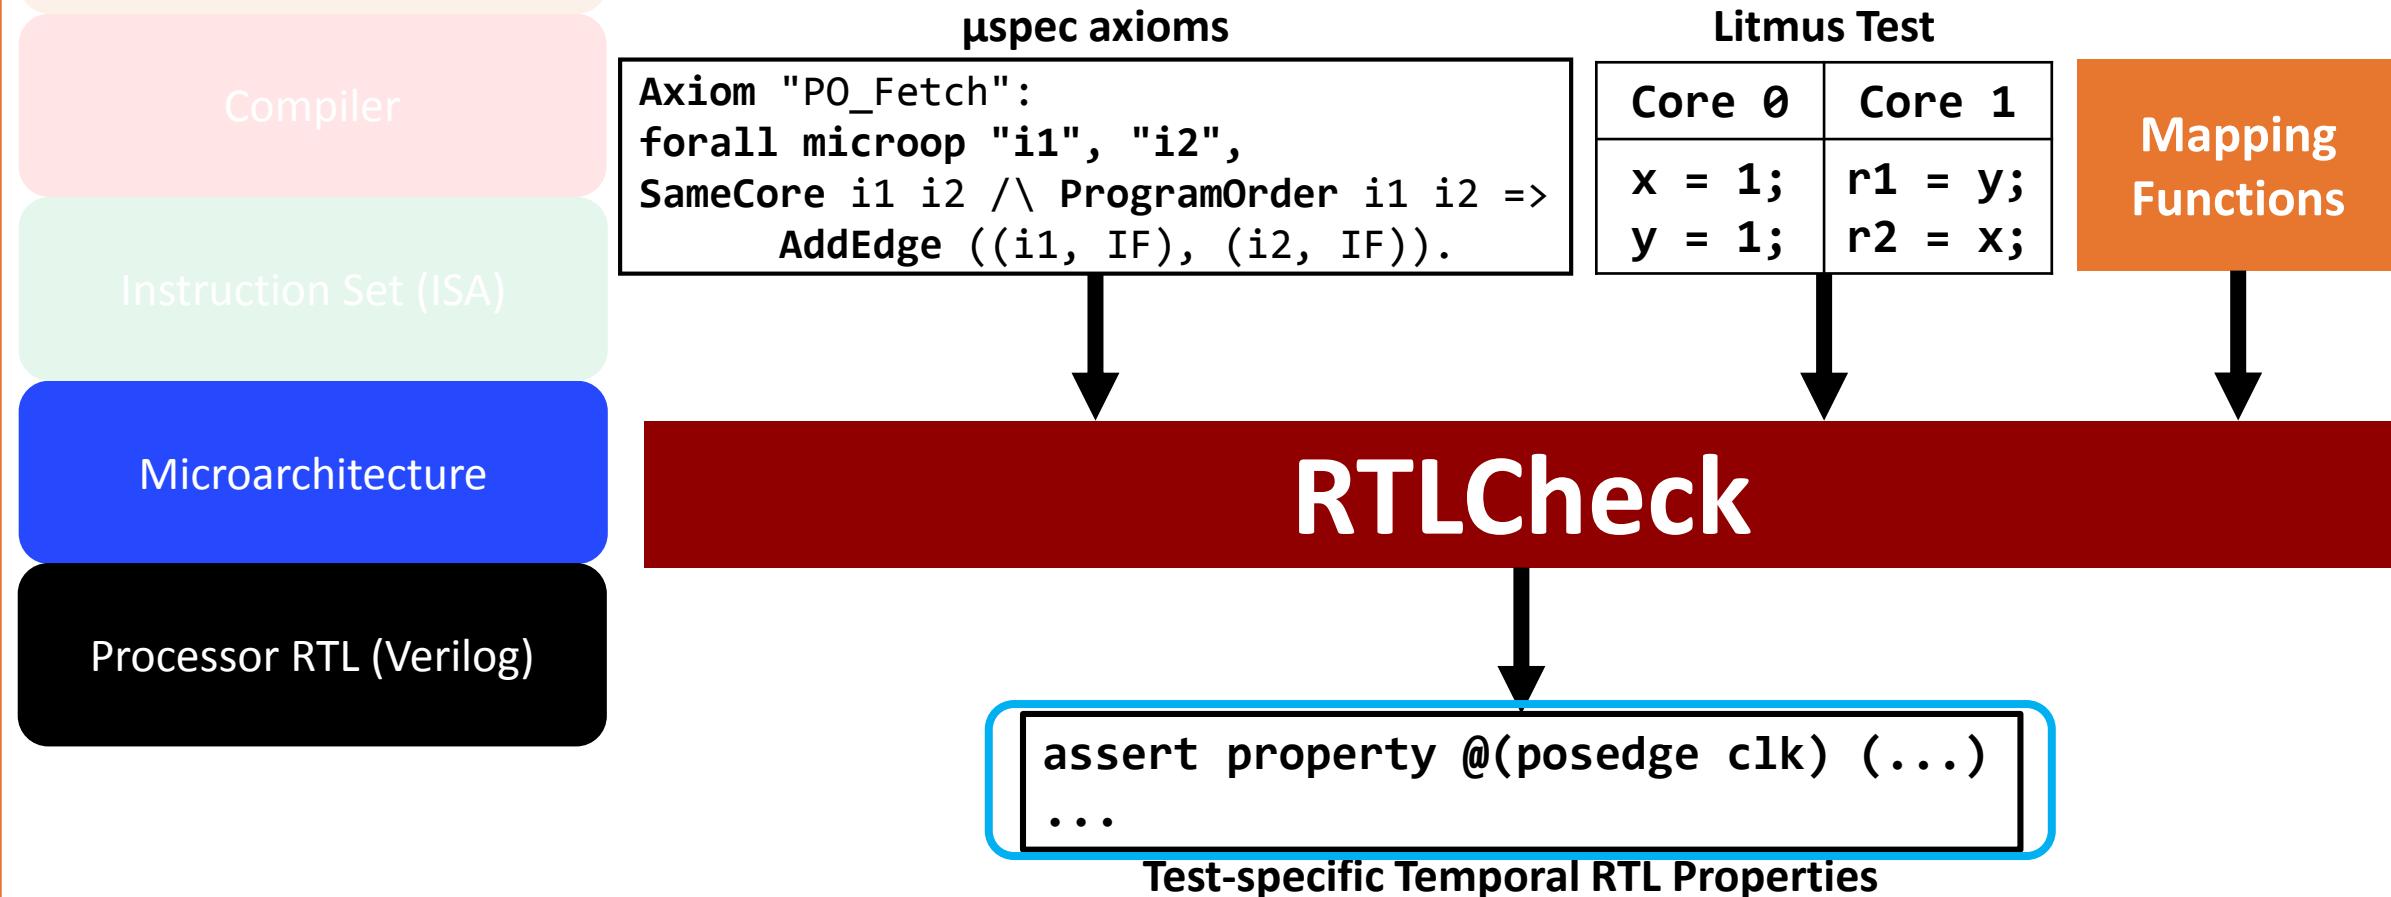
RTLCheck: Checking RTL Consistency Orderings

- RTLCheck enables **automated** checking of Verilog RTL against μ spec axioms for litmus test suites



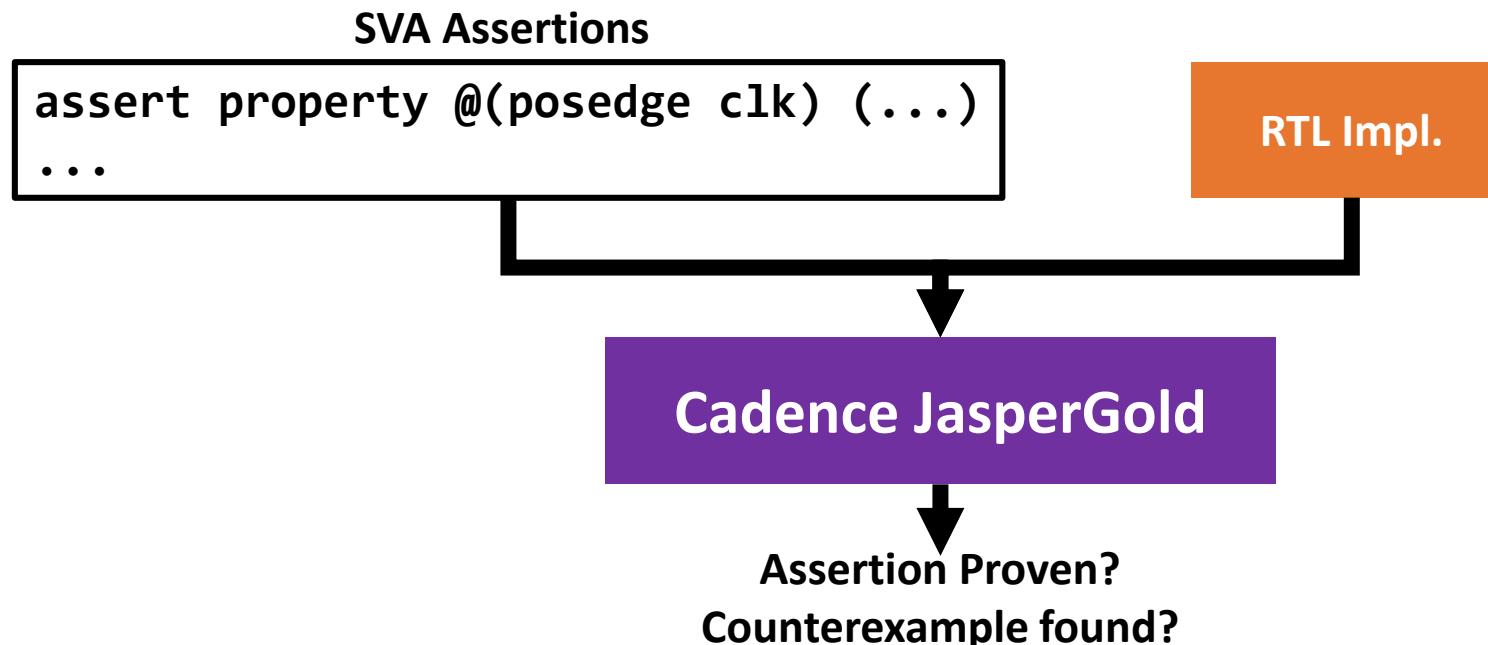
RTLCheck: Checking RTL Consistency Orderings

- RTLCheck enables **automated** checking of Verilog RTL against μ spec axioms for litmus test suites



SystemVerilog Assertions (SVA)

- SVA: Industry standard for RTL verification, e.g.: ARM [Reid et al. CAV 2016]
 - Based on Linear Temporal Logic (LTL) with regular operators
- Commercial tools (e.g. JasperGold) can formally verify SVA assertions
- **Translating μ spec to SVA => RTL MCM verification using industry flows**
- **But it's not that simple!**



Meaning can be Lost in Translation!

小心地滑

(Caution: Slippery Floor)



Meaning can be Lost in Translation!

小心地滑
(Caution: Slippery Floor)



[Image: Barbara Younger]

[Inspiration: Tae Jun Ham]



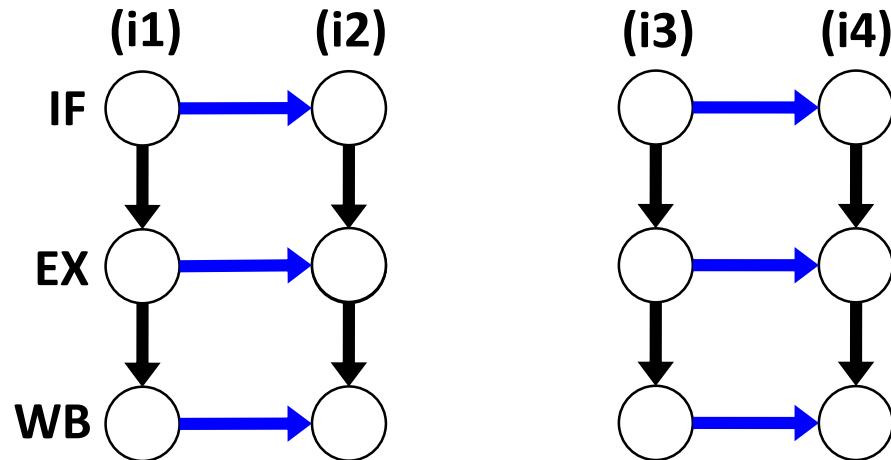
The μ spec/SVA Mismatch

- Tricky to translate μ spec to SVA while maintaining μ spec semantics
- **SVA Verifiers (JasperGold) don't implement full SVA spec!**
 - Causes further complications
- **Example: Outcome Filtering**
 - Filtering litmus test executions to those that have particular values for loads



Outcome Filtering with Execution as a Single Unit

- In this case, outcome filtering is **easy and efficient**
- Know load values, so can draw (**red**) edges based on these values
 - Example: i4 reads 0 => i4 must read mem before write i1

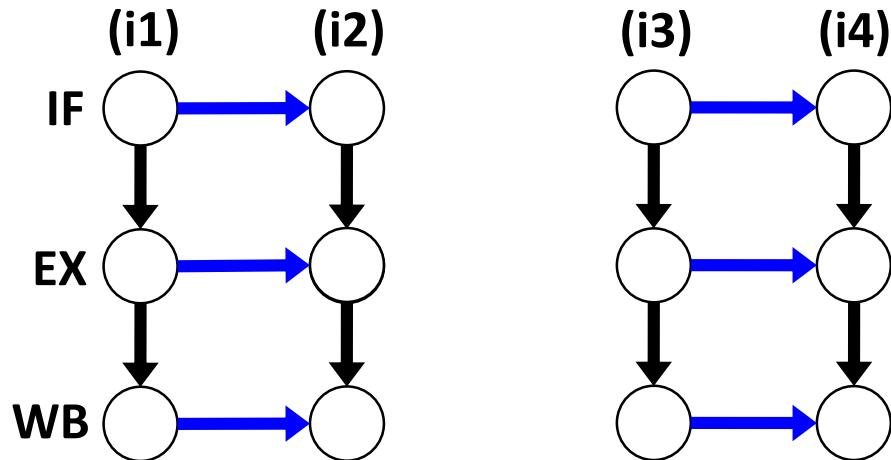


mp litmus test	
Core 0	Core 1
(i1) x = 1; (i2) y = 1;	(i3) r1 = y; (i4) r2 = x;



Outcome Filtering with Execution as a Single Unit

- In this case, outcome filtering is **easy and efficient**
- Know load values, so can draw (**red**) edges based on these values
 - Example: i4 reads 0 => i4 must read mem before write i1

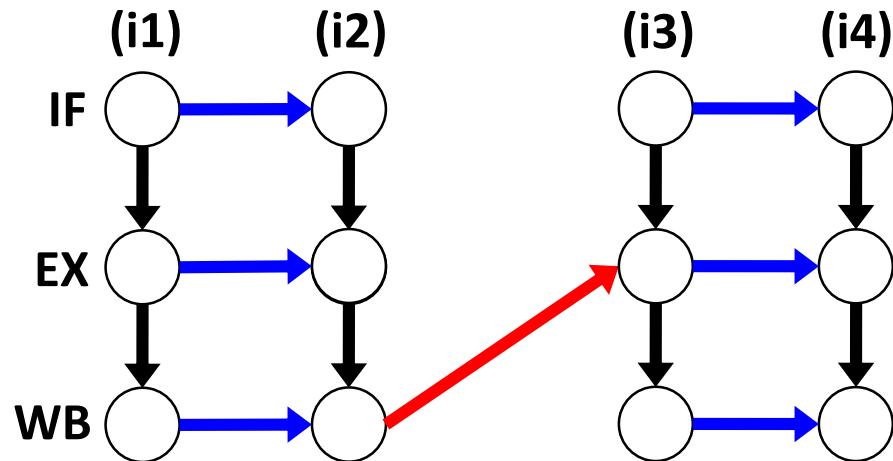


mp litmus test	
Core 0	Core 1
(i1) x = 1;	(i3) r1 = y;
(i2) y = 1;	(i4) r2 = x;
SC Forbids: r1 = 1, r2 = 0	



Outcome Filtering with Execution as a Single Unit

- In this case, outcome filtering is **easy and efficient**
- Know load values, so can draw (**red**) edges based on these values
 - Example: i4 reads 0 => i4 must read mem before write i1

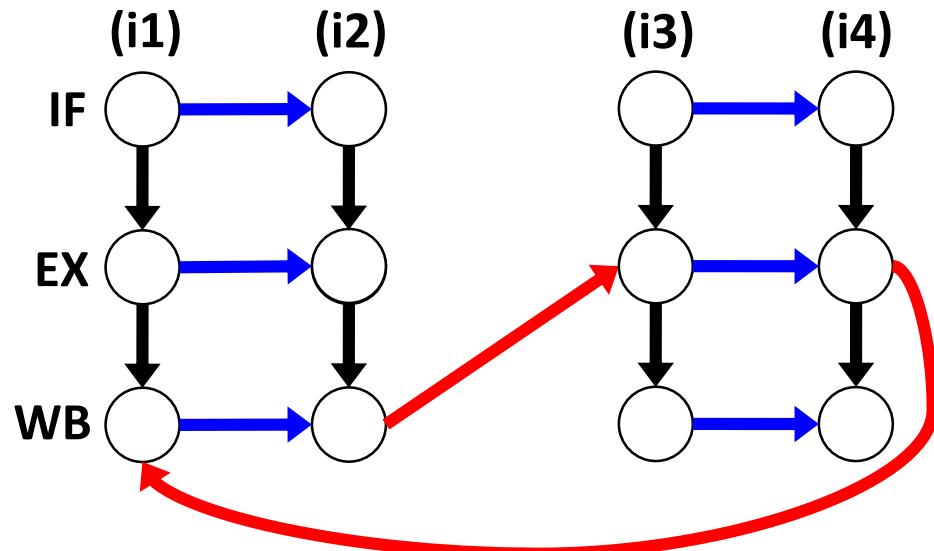


mp litmus test	
Core 0	Core 1
(i1) x = 1;	(i3) r1 = v;
(i2) y = 1;	(i4) r2 = x;
SC Forbids: r1 = 1, r2 = 0	



Outcome Filtering with Execution as a Single Unit

- In this case, outcome filtering is **easy and efficient**
- Know load values, so can draw (**red**) edges based on these values
 - Example: i4 reads 0 => i4 must read mem before write i1



mp litmus test	
Core 0	Core 1
(i1) x = 1;	(i3) r1 = y;
(i2) y = 1;	(i4) r2 = x;
SC Forbids: r1 = 1, r2 = 0	



Outcome Filtering with Temporal Logic

```
assume property (a); // e.g. Load i4 returns 0  
assert property (b); // e.g. i4 reads mem before write i1  
  
//The above is equivalent to...  
assert property ((always a) implies (always b));
```

- In temporal logic syntax ($G = \text{always}$, $F = \text{eventually}$), this becomes:
$$G a \rightarrow G b = (\sim(G a)) \vee G b = (F \sim a) \vee G b$$
- Assumptions introduce liveness: expensive to check! [Cerny et al. 2010]
- SVA verifiers **approximate**: only check assumptions until current state
 - This results in a property which is easier to check...
 - ...but **makes outcome filtering impossible** with such verifiers!
- **RTLCheck Solution:** Generate properties that handle all test outcomes

RTLCheck Takeaways

- First **automated RTL MCM verification** for litmus test suites
 - Engineers can check MCM properties of their RTL themselves
 - Compatible with existing industry flows and tools
- Novel algorithms to translate μ spec **axioms** to **temporal** SVA properties
 - Ongoing work: Formalise mismatch between μ spec and SVA
- **Discovered bug** in memory implementation of RISC-V V-scale processor
- Accolades:
 - “Honorable Mention” in 2017 IEEE Micro *Top Picks of Comp. Arch. Conferences*

Talk Outline

- Overview and Motivation
- Background on MCM Specification and Verification
- **PipeProof:** All-Program Microarchitectural MCM Verification
- **RTLCheck:** MCM Verification of Verilog RTL
- Expanding to other domains
- Conclusion



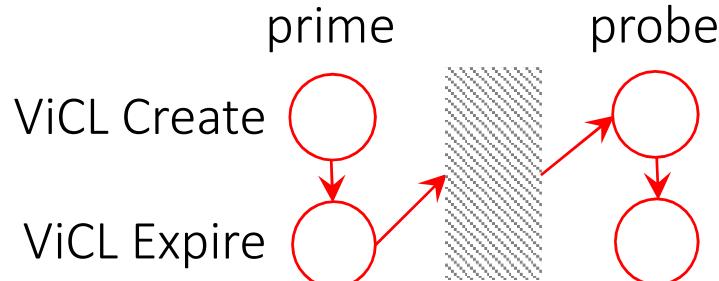
Security Analysis with CheckMate [Trippel et al. MICRO 2018]

- Work by another member of our research group (Caroline Trippel)
- Her key insight: **μhb graphs can be used for reasoning about security!**

Microarchitecture + OS Specification in Alloy

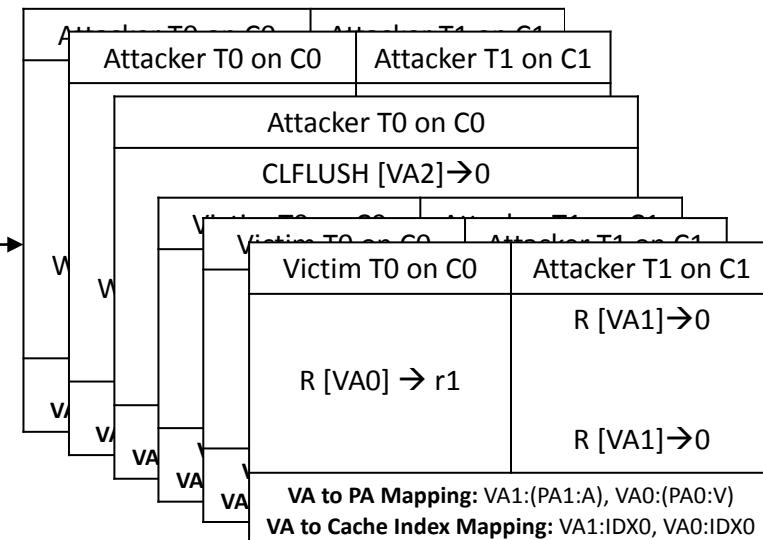
```
fact Program_Order_Fetch {  
    all disj e0, e1 : Event |  
    ProgramOrder[e0, e1] =>  
    EdgeExists[e0, Fetch, e1, Fetch, uhb_inter]  
}  
  
fact In_Order_Decode {  
    all disj e0, e1 : Event |  
    EdgeExists[e0, Fetch, e1, Fetch, uhb_inter] =>  
    EdgeExists[e0, Decode, e1, Decode, uhb_inter]  
}
```

Exploit Pattern Specification



CheckMate
Hardware Exploit
Prog. Synthesis

Exploits synthesized from μhb analysis



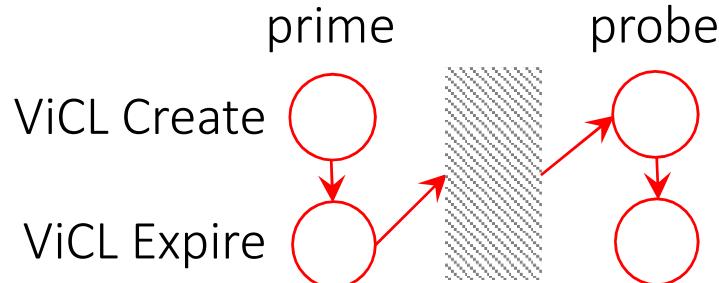
Security Analysis with CheckMate [Trippel et al. MICRO 2018]

- Work by another member of our research group (Caroline Trippel)
- Her key insight: **μhb graphs can be used for reasoning about security!**

Microarchitecture + OS Specification in Alloy

```
fact Program_Order_Fetch {  
    all disj e0, e1 : Event |  
    ProgramOrder[e0, e1] =>  
    EdgeExists[e0, Fetch, e1, Fetch, uhb_inter]  
}  
  
fact In_Order_Decode {  
    all disj e0, e1 : Event |  
    EdgeExists[e0, Fetch, e1, Fetch, uhb_inter] =>  
    EdgeExists[e0, Decode, e1, Decode, uhb_inter]  
}
```

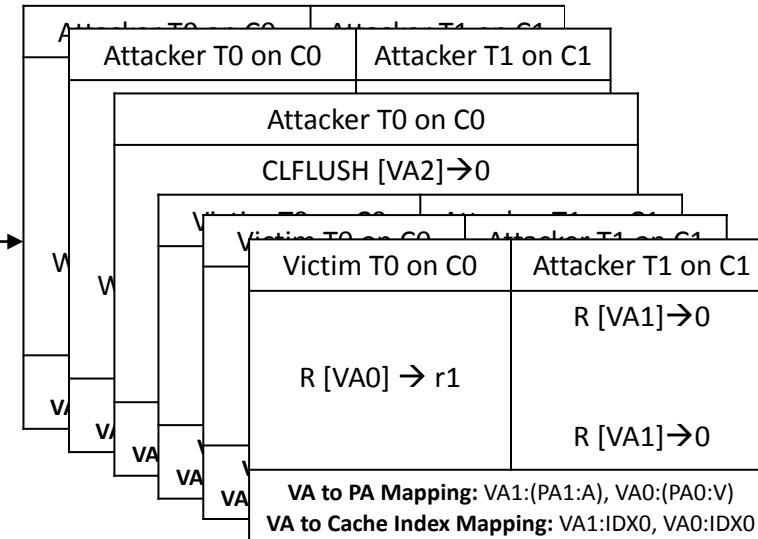
Exploit Pattern Specification



CheckMate
Hardware Exploit
Prog. Synthesis

Includes new exploits!
(SpectrePrime, MeltdownPrime)

Exploits synthesized from μhb analysis



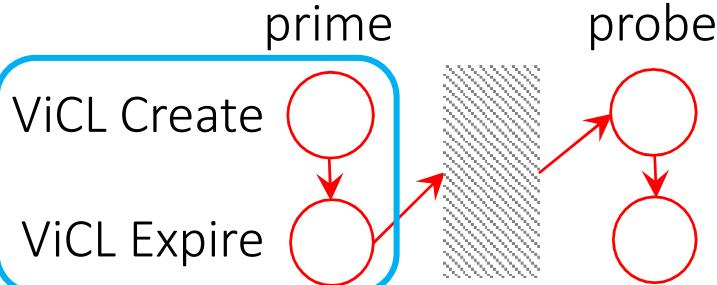
Security Analysis with CheckMate [Trippel et al. MICRO 2018]

- Work by another member of our research group (Caroline Trippel)
- Her key insight: **μhb graphs can be used for reasoning about security!**

Microarchitecture + OS Specification in Alloy

```
fact Program_Order_Fetch {  
    all disj e0, e1 : Event |  
    ProgramOrder[e0, e1] =>  
    EdgeExists[e0, Fetch, e1, Fetch, uhb_inter]  
}  
  
fact In_Order_Decode {  
    all disj e0, e1 : Event |  
    EdgeExists[e0, Fetch, e1, Fetch, uhb_inter] =>  
    EdgeExists[e0, Decode, e1, Decode, uhb_inter]  
}
```

Exploit Pattern Specification

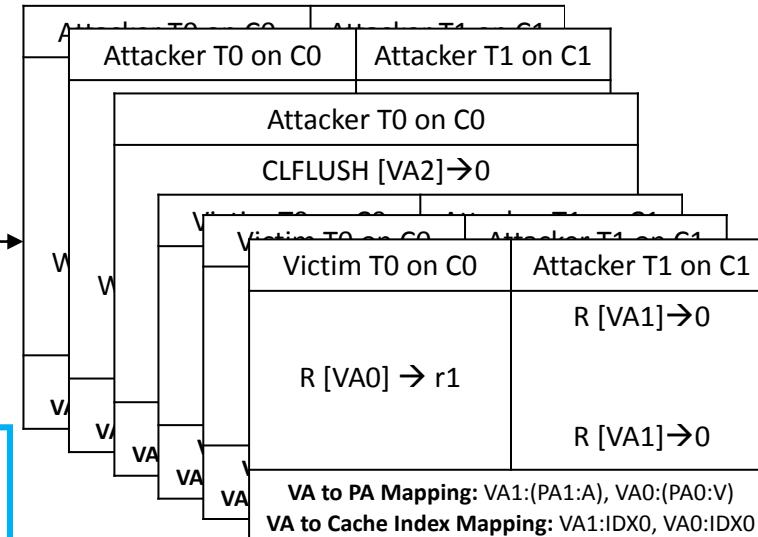


CheckMate
Hardware Exploit
Prog. Synthesis

ViCL abstraction [Manerkar et al. MICRO 2015] used to model cache behaviour

Includes new exploits!
(SpectrePrime, MeltdownPrime)

Exploits synthesized from μhb analysis



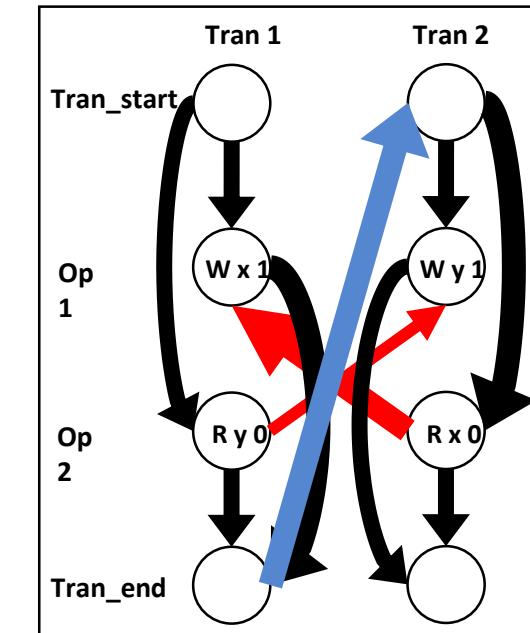
Ongoing Work: Verifying Distributed Systems

- Joint work with Themis Melissaris
- Distributed systems have some similarities to shared-memory systems
 - Distributed protocols (e.g. Paxos) similar to **cache coherence protocols**
 - Replicated data store consistency models similar to **MCMs**



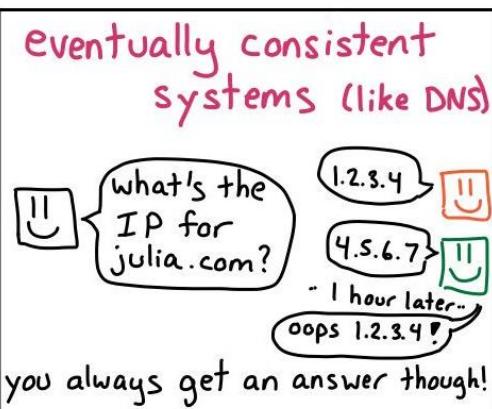
Ongoing Work: Verifying Distributed Systems

- Joint work with Themis Melissaris
- Distributed systems have some similarities to shared-memory systems
 - Distributed protocols (e.g. Paxos) similar to **cache coherence protocols**
 - Replicated data store consistency models similar to **MCMs**

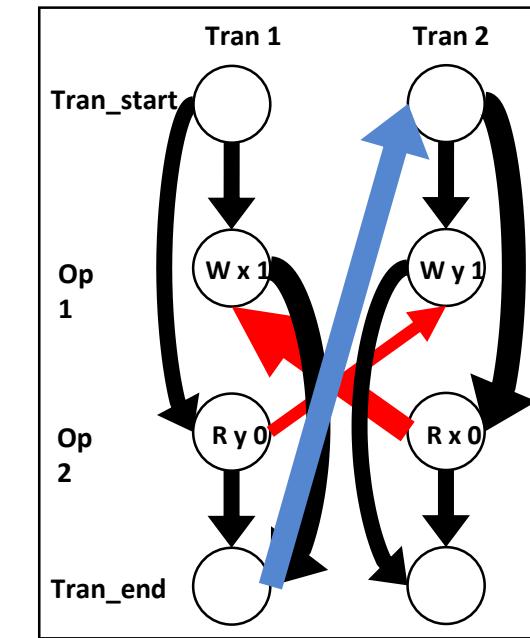


Ongoing Work: Verifying Distributed Systems

- Joint work with Themis Melissaris
- Distributed systems have some similarities to shared-memory systems
 - Distributed protocols (e.g. Paxos) similar to **cache coherence protocols**
 - Replicated data store consistency models similar to **MCMs**
- Also have features with no shared-memory analogue!
 - Correctness in the presence of node failures
 - Eventual consistency [Vogels CACM 2009]



[Cartoon by Julia Evans]



Talk Outline

- Overview and Motivation
- Background on MCM Specification and Verification
- **PipeProof:** All-Program Microarchitectural MCM Verification
- **RTLCheck:** MCM Verification of Verilog RTL
- Expanding to other domains
- Conclusion



Conclusions

- **Complexity of computing hardware is increasing**
 - Ubiquitous parallelism and increased heterogeneity
- **Automated formal verification helps engineers handle this complexity**
 - Give engineers the ability to formally verify their systems themselves
 - **PipeProof**: Automated All-Program Microarchitectural MCM Verification
 - **RTLCheck**: Per-Program MCM Verification of RTL Designs
- **Techniques for MCM analysis applicable to other domains**
 - e.g. Security [Trippel et al. MICRO 2018] and distributed systems

Collaborators



Margaret Martonosi



**Daniel Lustig
(NVIDIA)**



Aarti Gupta



**Michael Pelluaer
(NVIDIA)**



Caroline Trippel



Sharad Malik



Hongce Zhang

Automated Formal Memory Consistency Verification of Hardware

Yatin A. Manerkar

Princeton University

June 23rd, 2019

<http://www.cs.princeton.edu/~manerkar>



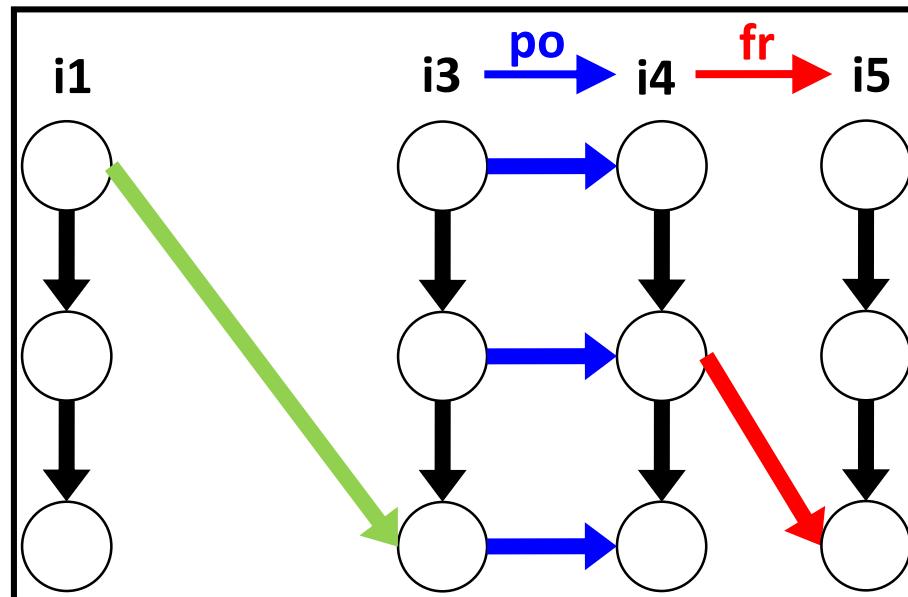
Backup Slides



Chain Invariants

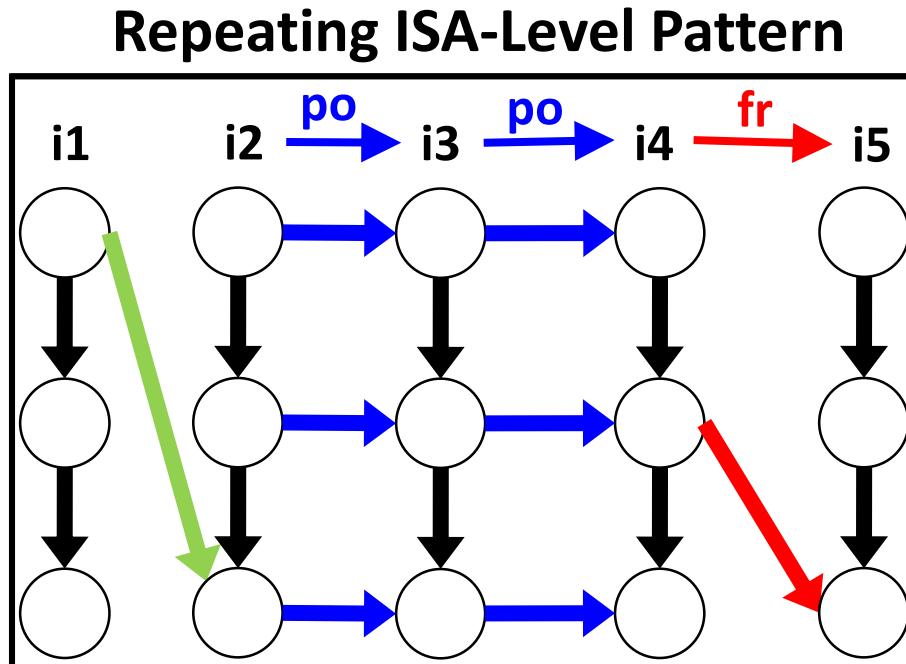
- Abstractly represent repeated ISA-level patterns
- Sometimes needed for refinement loop to terminate
- **Inductively proven by PipeProof before their use in proof algorithms**
- Example: checking for edge from i1 to i5 (TC abstraction support proof)

Abstract Counterexample



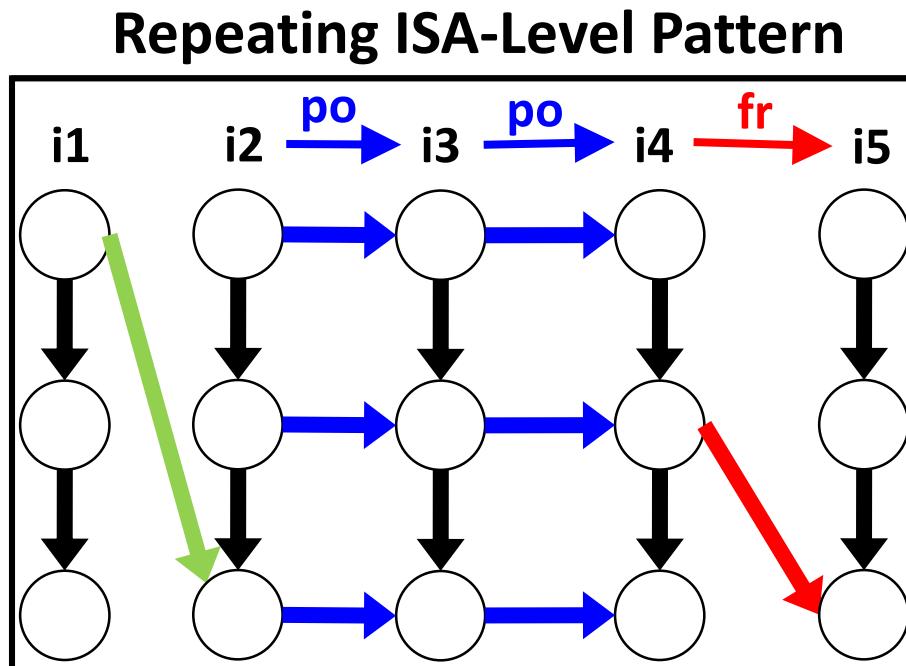
Chain Invariants

- Abstractly represent repeated ISA-level patterns
- Sometimes needed for refinement loop to terminate
- **Inductively proven by PipeProof before their use in proof algorithms**
- Example: checking for edge from i1 to i5 (TC abstraction support proof)



Chain Invariants

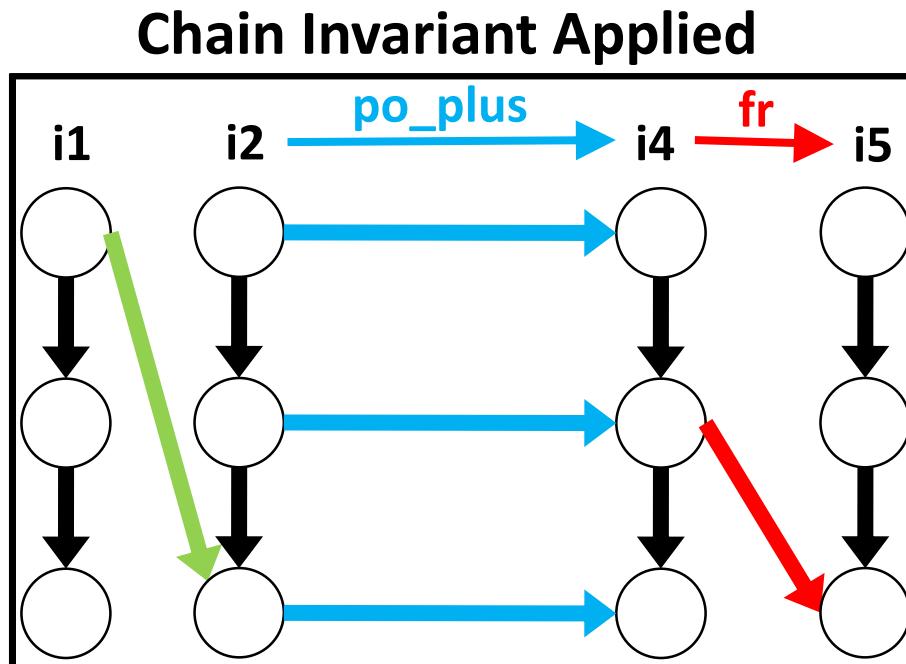
- Abstractly represent repeated ISA-level patterns
- Sometimes needed for refinement loop to terminate
- **Inductively proven by PipeProof before their use in proof algorithms**
- Example: checking for edge from i1 to i5 (TC abstraction support proof)



Can continue decomposing in this way forever!

Chain Invariants

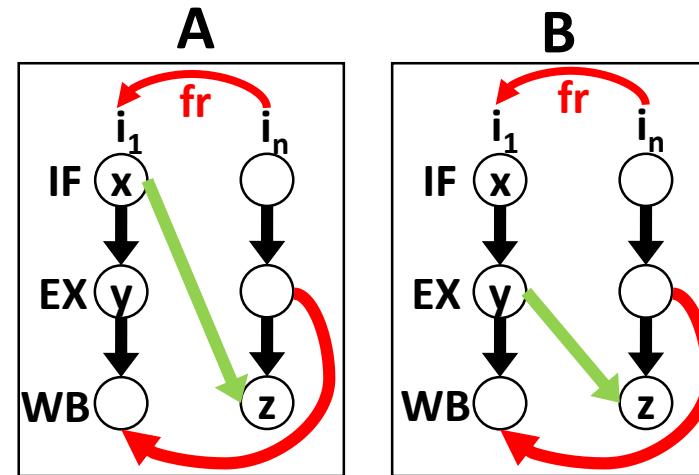
- Abstractly represent repeated ISA-level patterns
- Sometimes needed for refinement loop to terminate
- **Inductively proven by PipeProof before their use in proof algorithms**
- Example: checking for edge from i1 to i5 (TC abstraction support proof)



-**po_plus** = arbitrary number of repetitions of **po**
-Next edge peeled off will be something other than **po**

Covering Sets Optimization

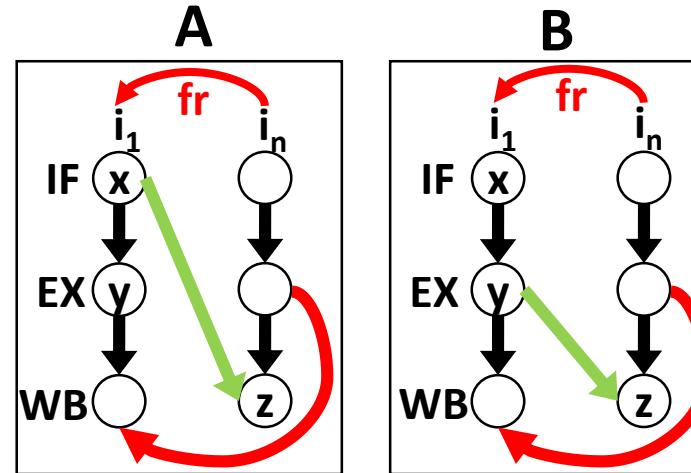
- Must verify across all possible transitive connections
- Each decomposition creates a new set of transitive connections
 - Can quickly lead to a case explosion
- The Covering Sets Optimization eliminates redundant transitive connections



Covering Sets Optimization

- Must verify across all possible transitive connections
- Each decomposition creates a new set of transitive connections
 - Can quickly lead to a case explosion
- The Covering Sets Optimization eliminates redundant transitive connections

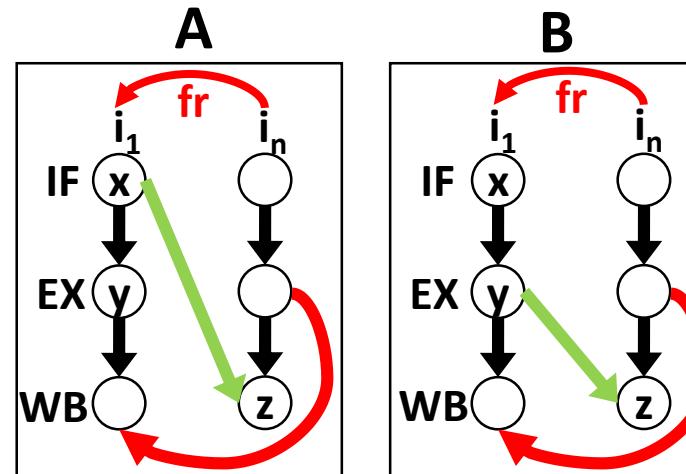
Graph A has an edge from $x \rightarrow z$ (tran conn.)



Covering Sets Optimization

- Must verify across all possible transitive connections
- Each decomposition creates a new set of transitive connections
 - Can quickly lead to a case explosion
- The Covering Sets Optimization eliminates redundant transitive connections

Graph A has an edge from $x \rightarrow z$ (tran conn.)



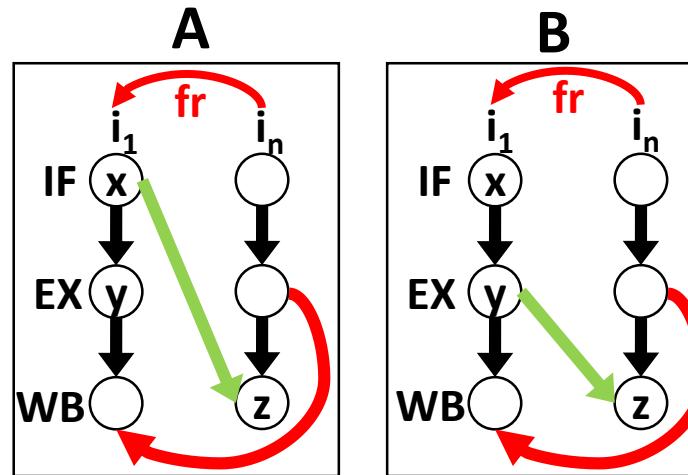
Graph B has edges from $y \rightarrow z$ (tran conn.) and $x \rightarrow z$ (by transitivity)



Covering Sets Optimization

- Must verify across all possible transitive connections
- Each decomposition creates a new set of transitive connections
 - Can quickly lead to a case explosion
- The Covering Sets Optimization eliminates redundant transitive connections

Graph A has an edge from $x \rightarrow z$ (tran conn.)



Graph B has edges from $y \rightarrow z$ (tran conn.) and $x \rightarrow z$ (by transitivity)

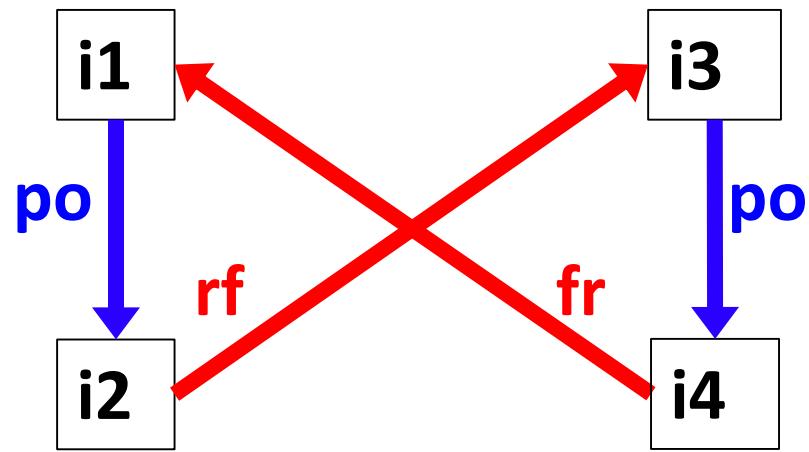
Correctness of A => Correctness of B (since B contains A's tran conn.)

Checking B explicitly is redundant!



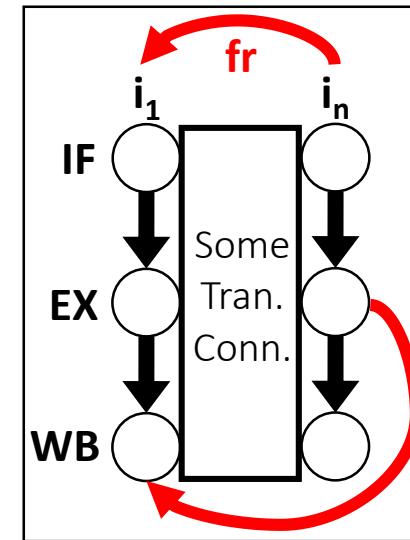
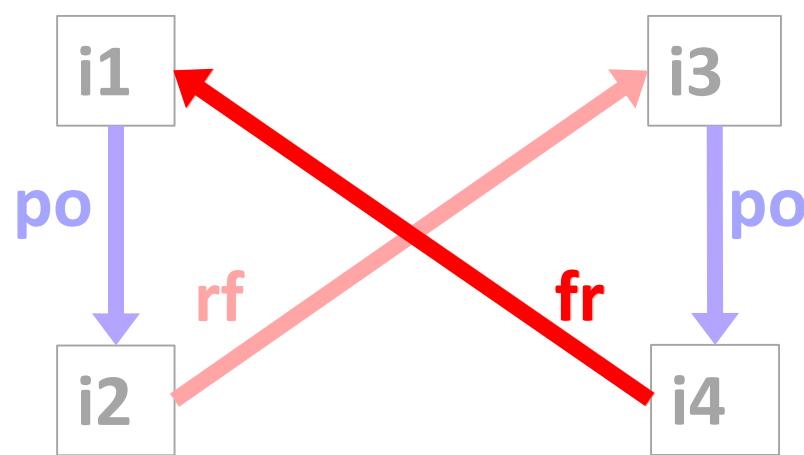
Memoization Optimization

- Base PipeProof algorithm examines some cycles multiple times
- Memoization eliminates redundant checks of cycles that have already been verified



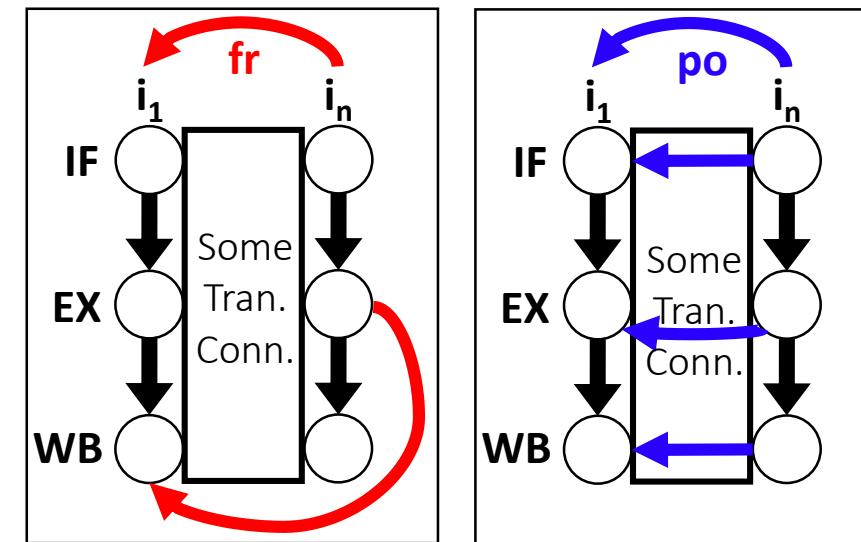
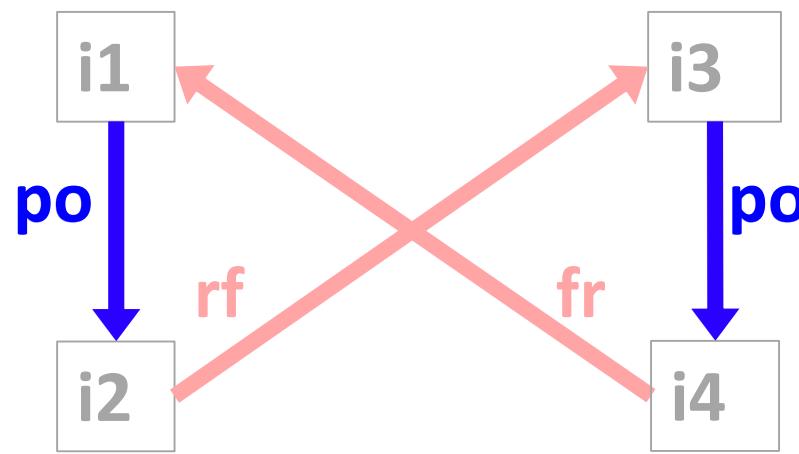
Memoization Optimization

- Base PipeProof algorithm examines some cycles multiple times
- Memoization eliminates redundant checks of cycles that have already been verified



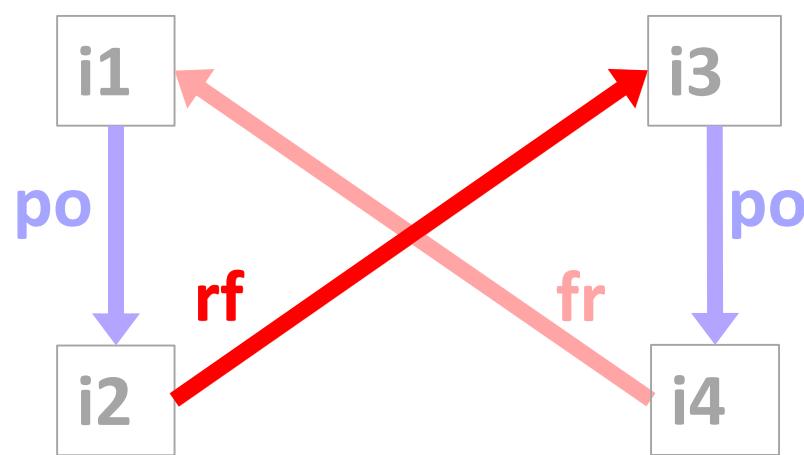
Memoization Optimization

- Base PipeProof algorithm examines some cycles multiple times
- Memoization eliminates redundant checks of cycles that have already been verified

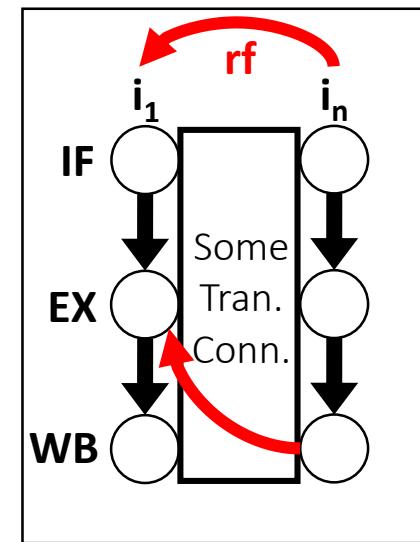
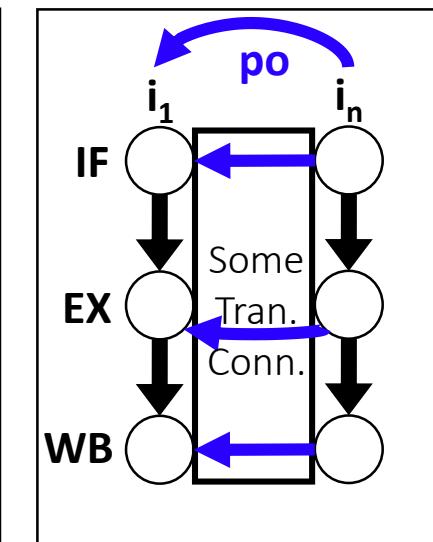
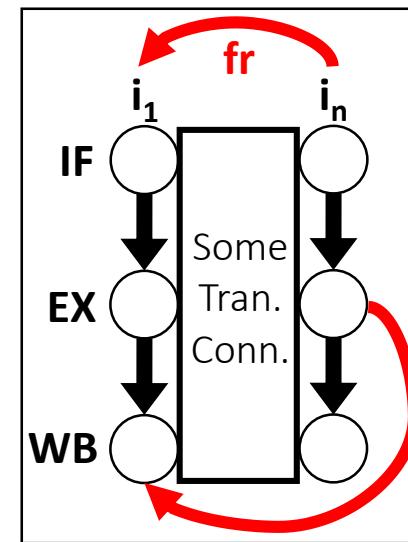


Memoization Optimization

- Base PipeProof algorithm examines some cycles multiple times
- Memoization eliminates redundant checks of cycles that have already been verified

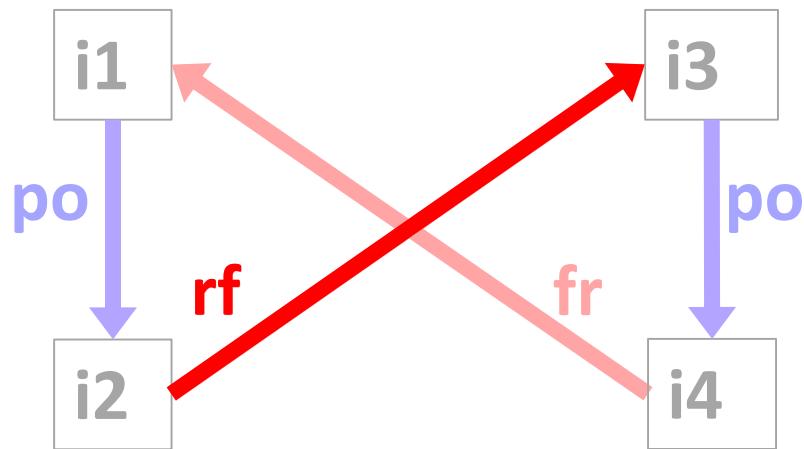


Same cycle is checked 3 times!

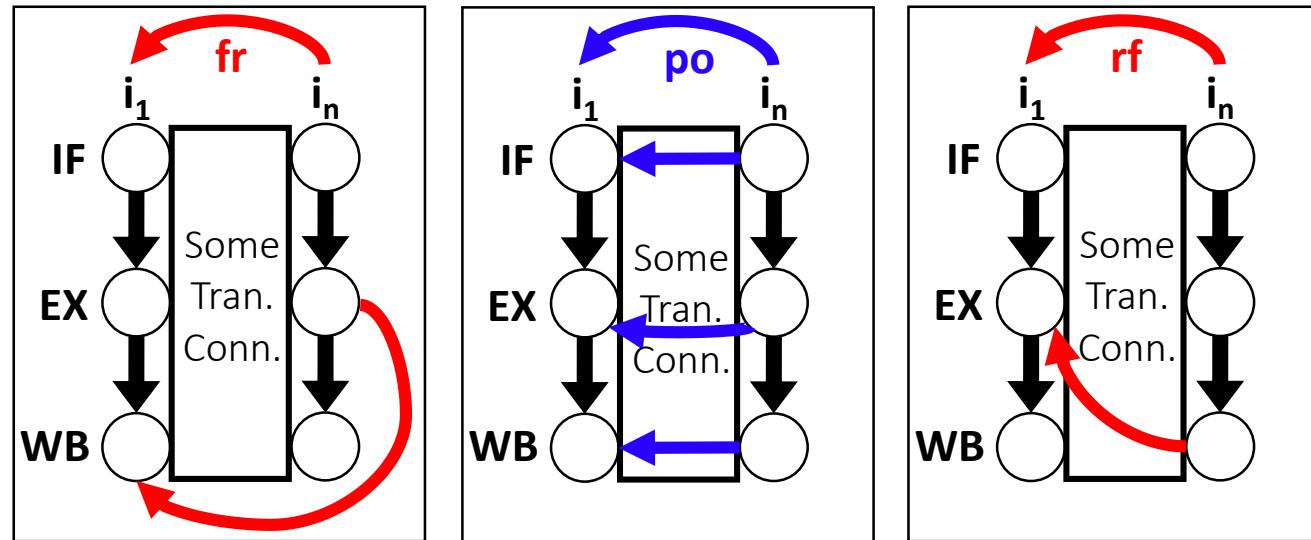


Memoization Optimization

- Base PipeProof algorithm examines some cycles multiple times
- Memoization eliminates redundant checks of cycles that have already been verified



Same cycle is checked 3 times!

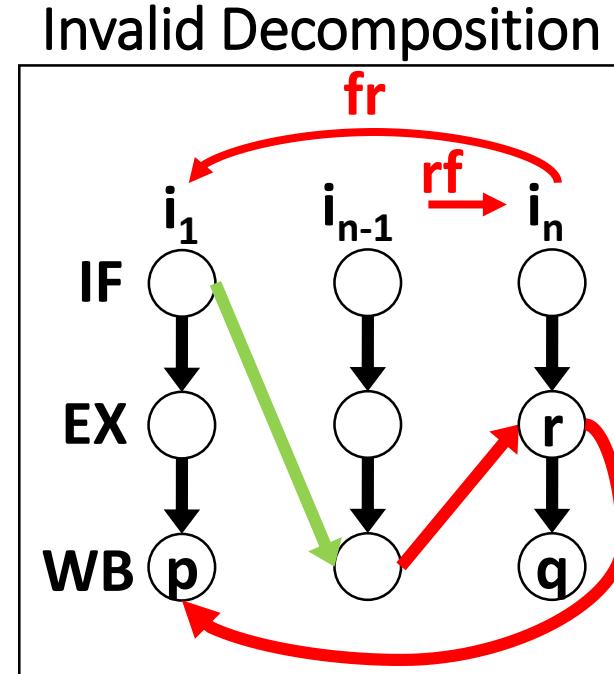
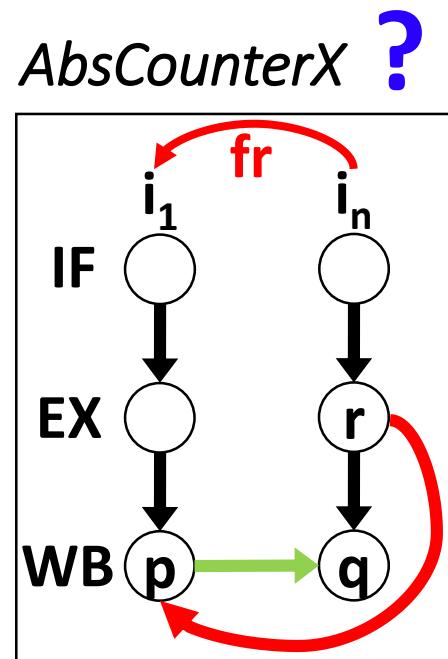


Procedure: If all ISA-level cycles containing edge r_i have been checked,
do not peel off r_i edges when checking subsequent cycles



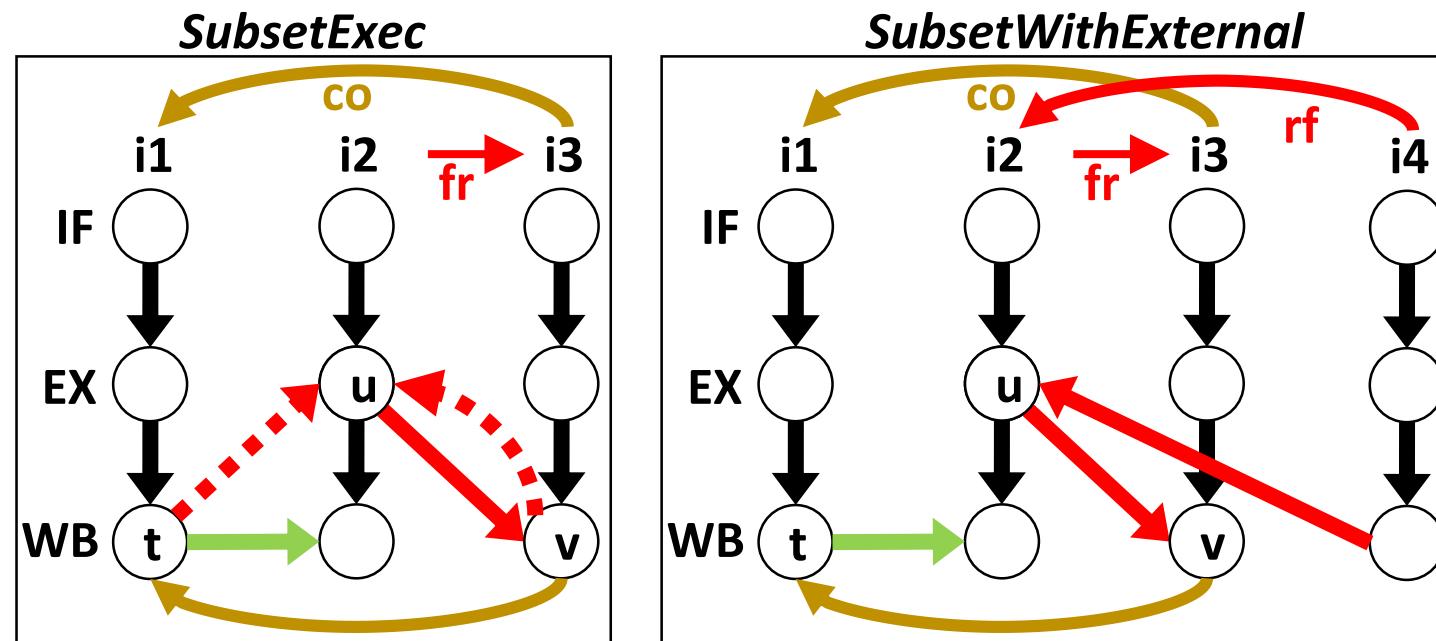
Filtering Invalid Decompositions

- When decomposing a transitive connection, the decomposition should guarantee the transitive connections of its parent abstract cexes.
- Decompositions that do not do this are invalid and filtered out

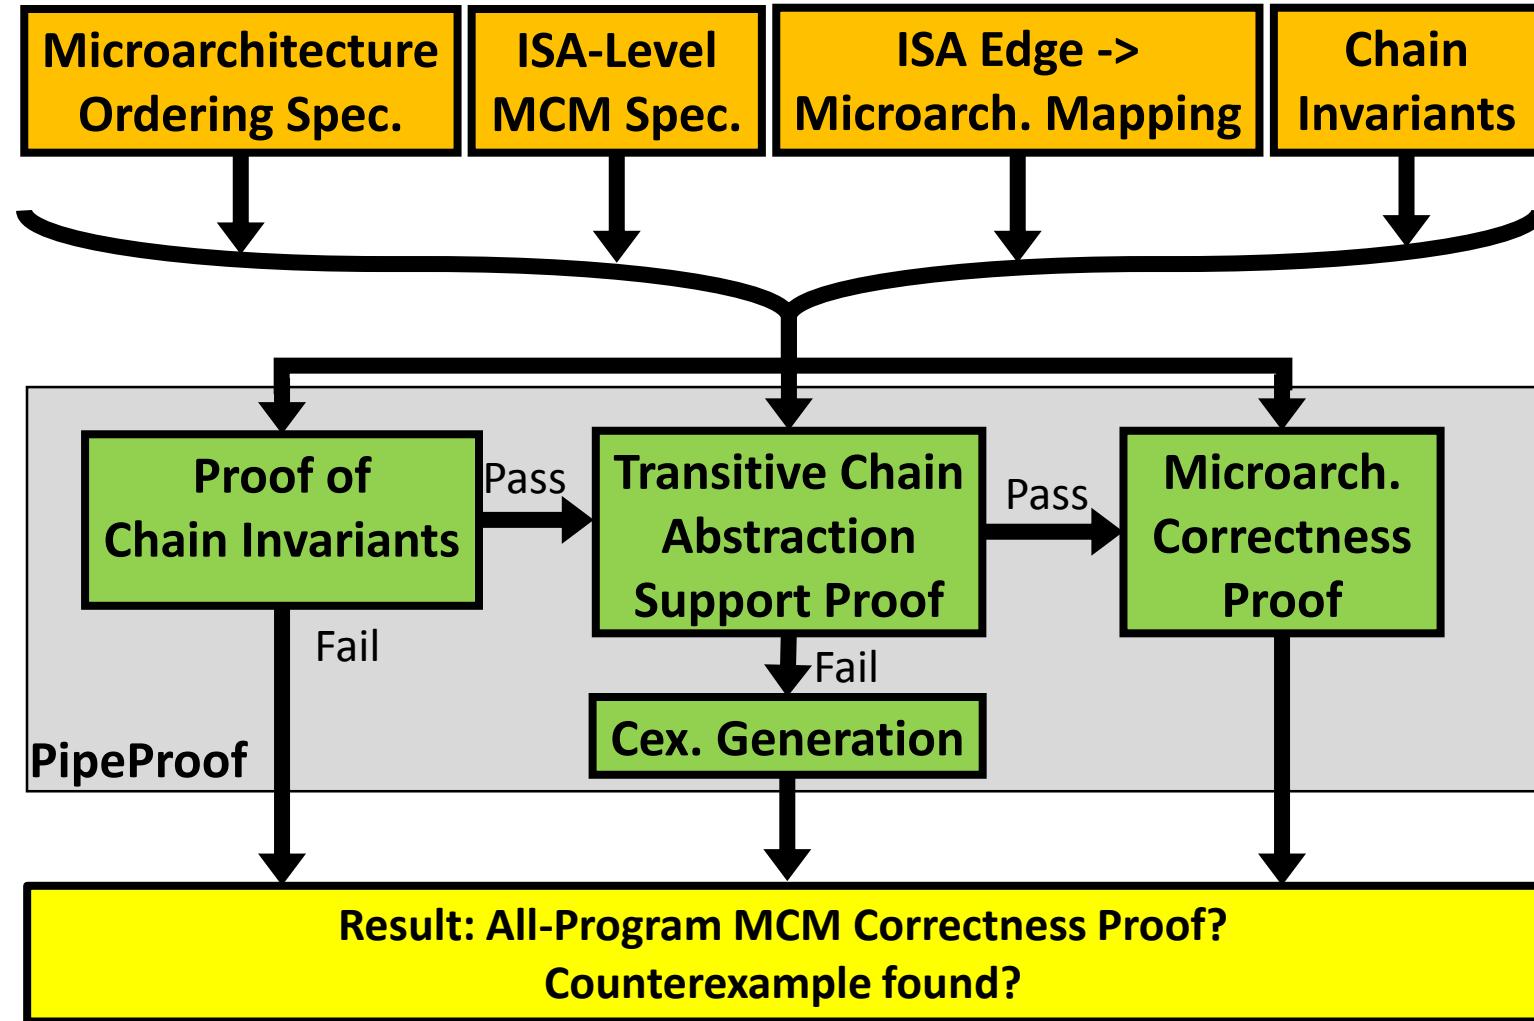


The Adequate Model Over-Approximation

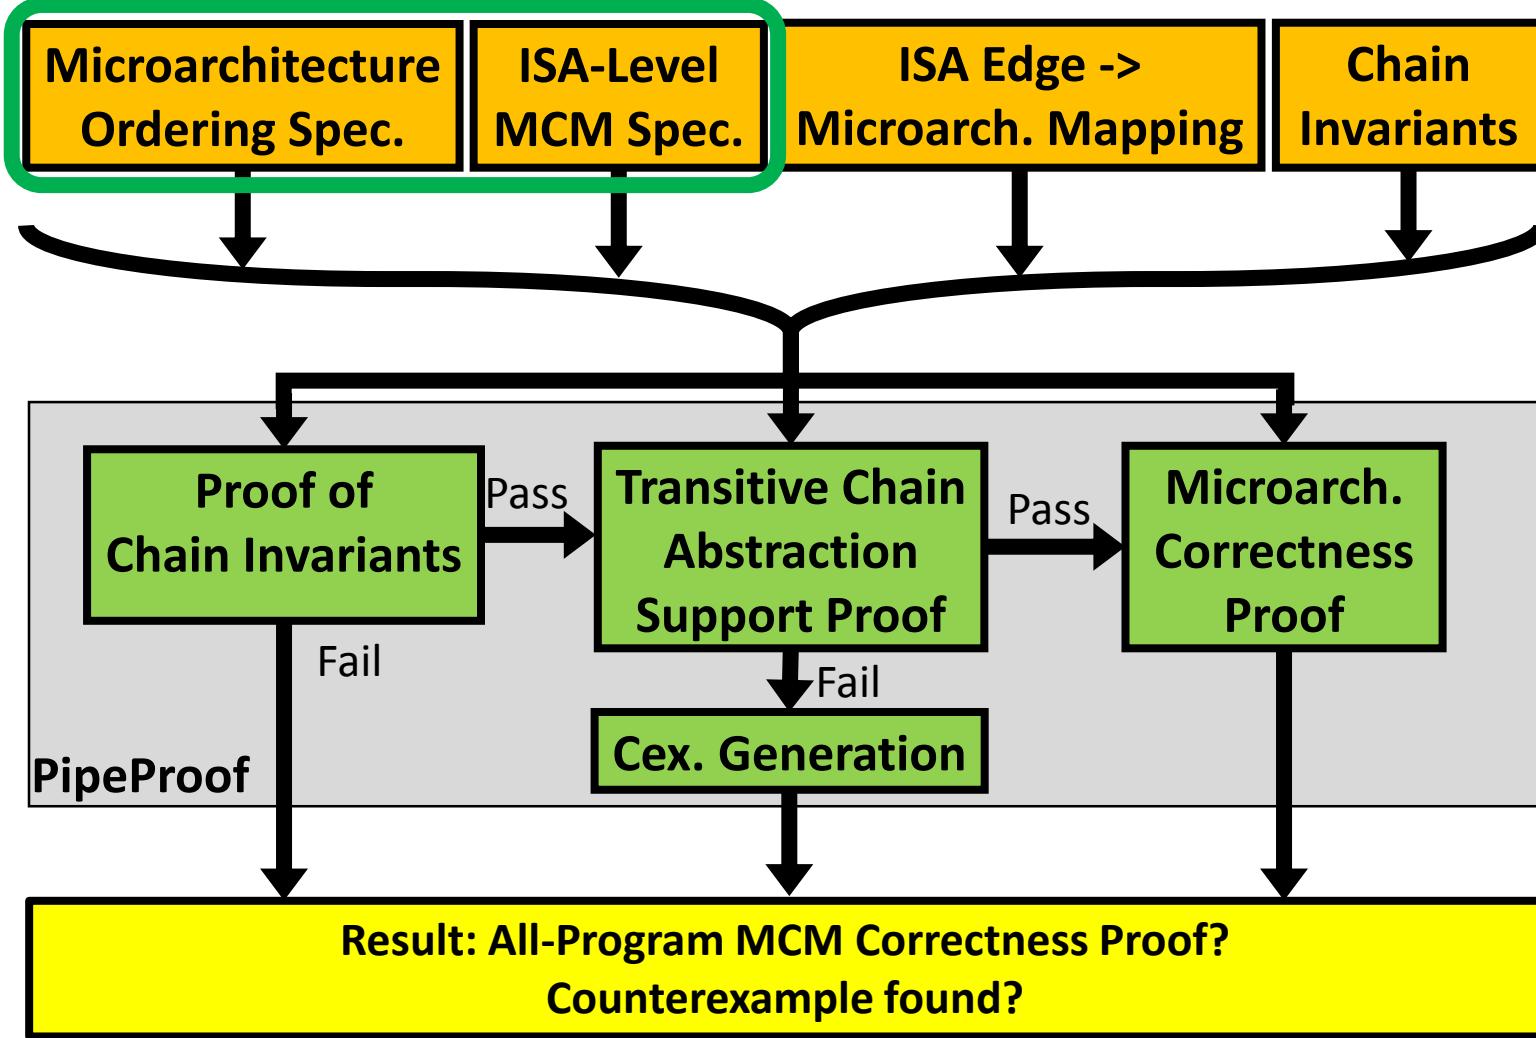
- Addition of an instruction can make unobservable execution observable!
- Need to work with over-approximation of microarchitectural constraints
- PipeProof sets all **exists** clauses to true as its over-approximation



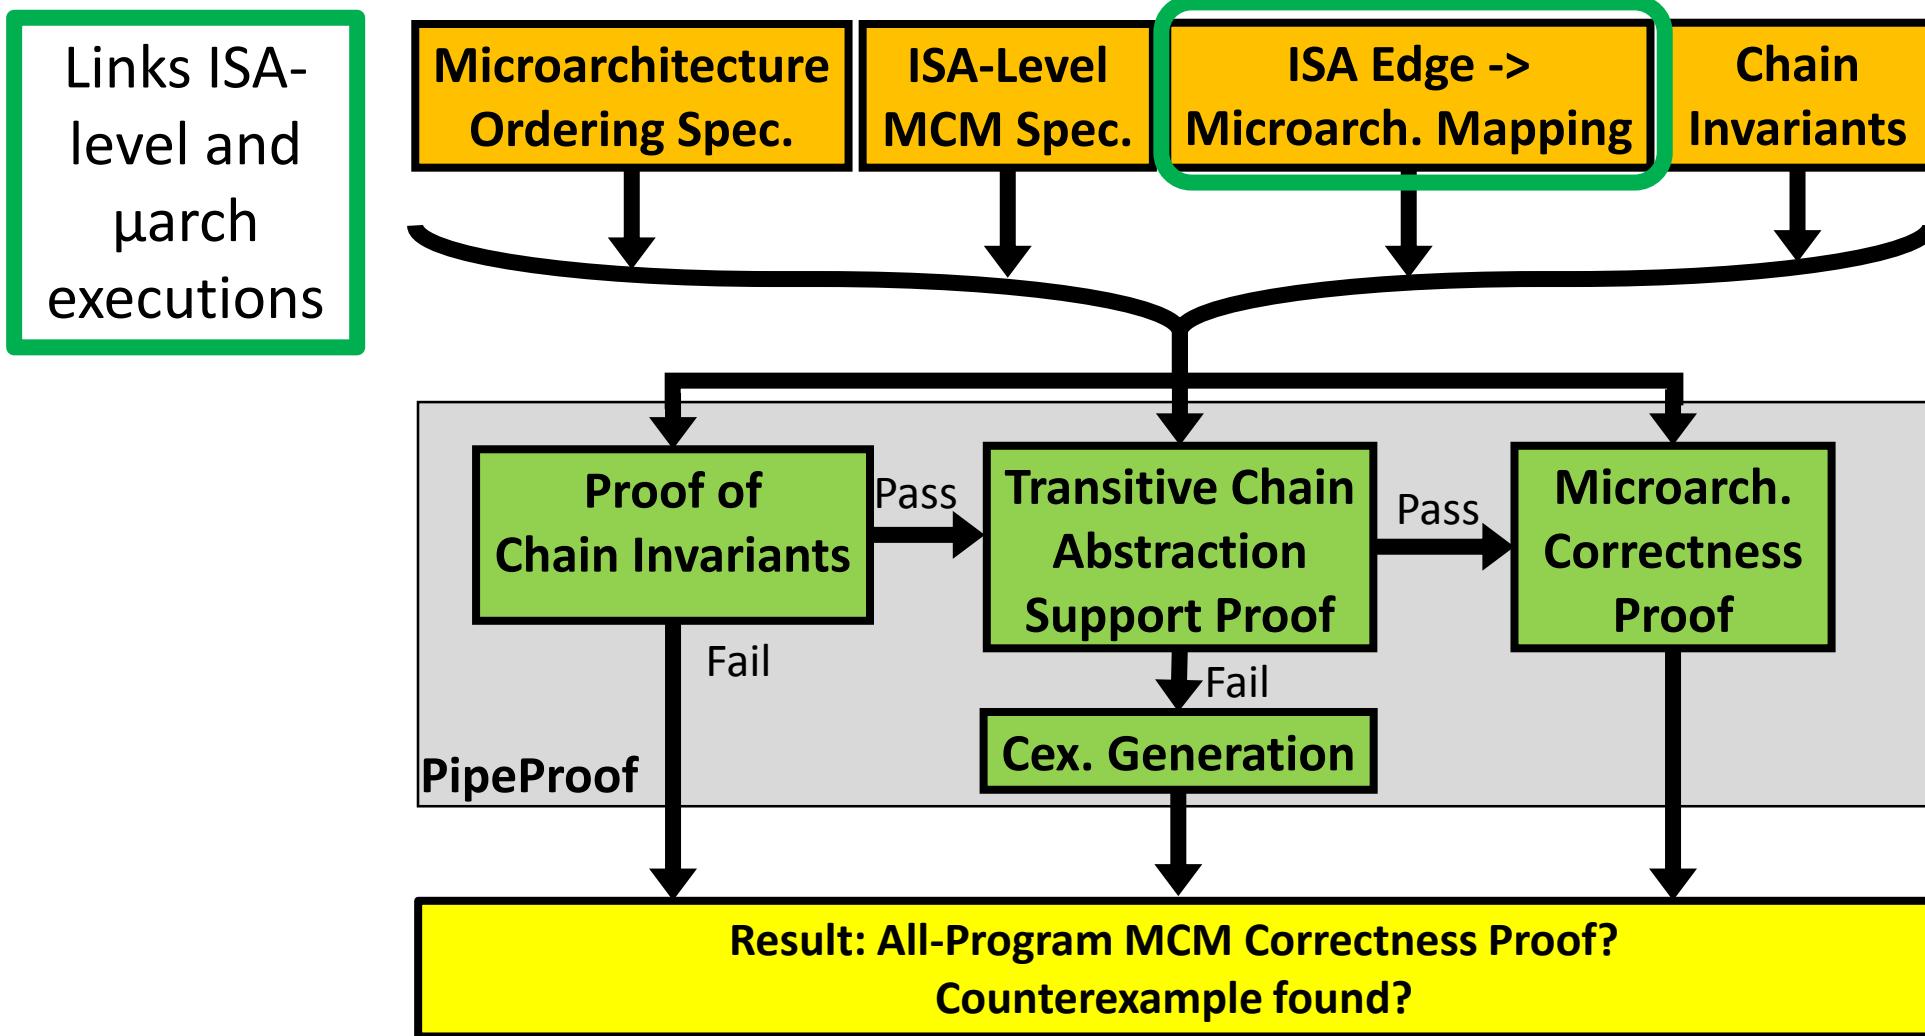
PipeProof Block Diagram



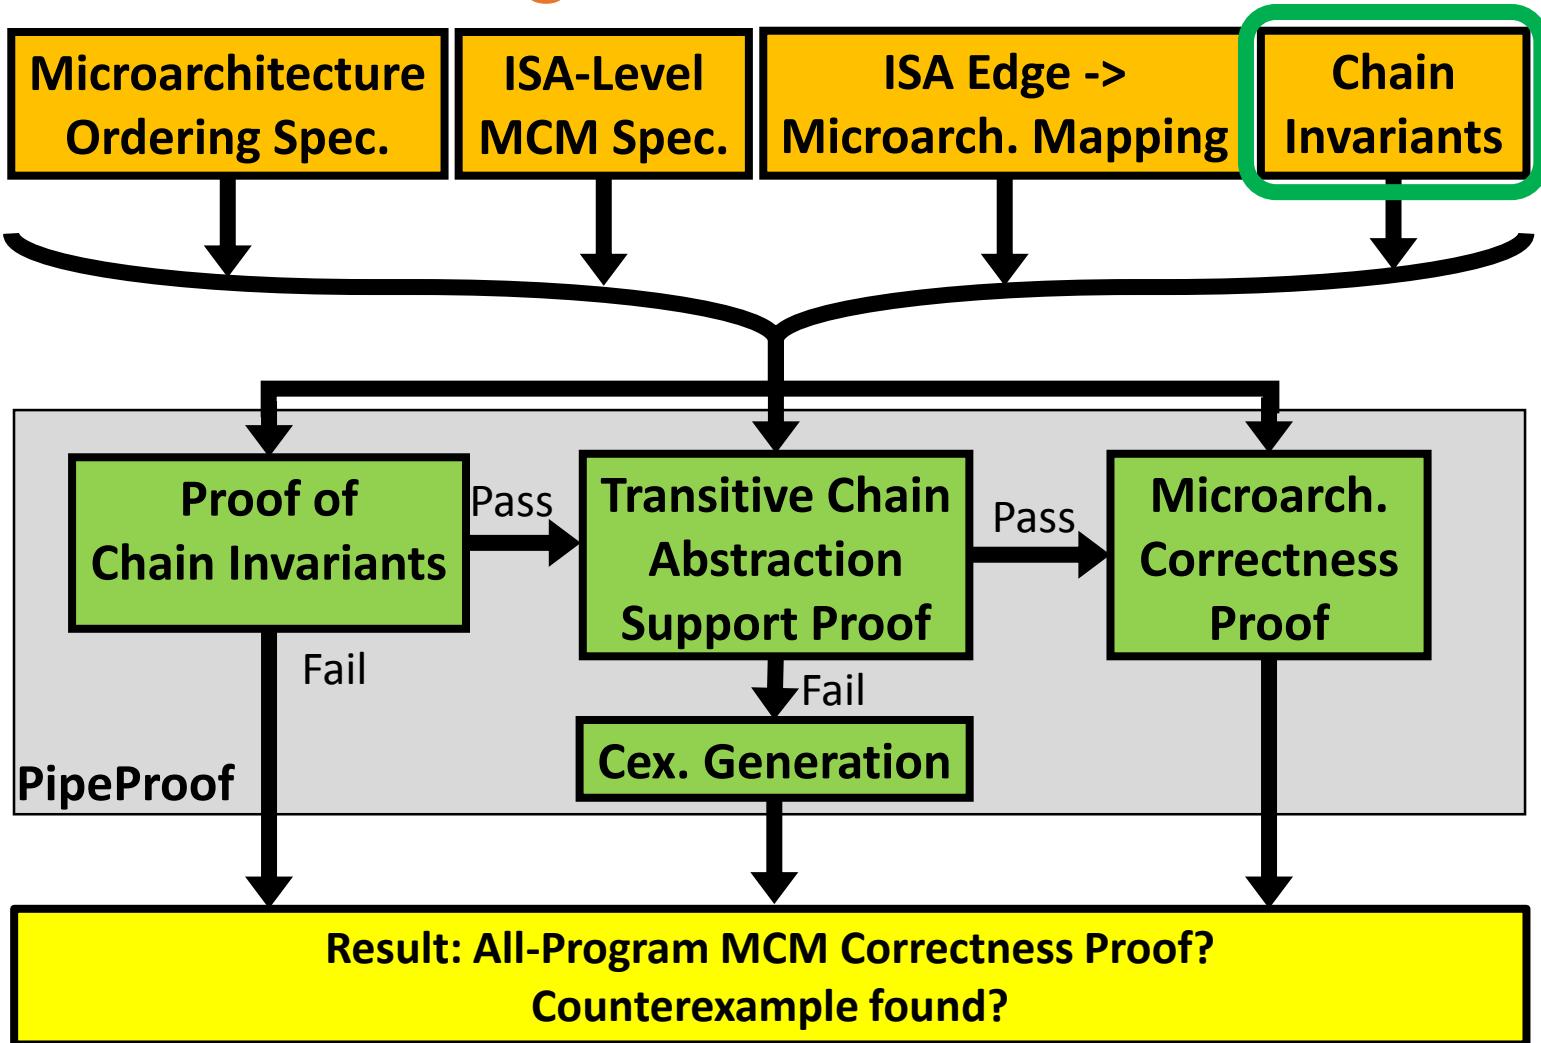
PipeProof Block Diagram



PipeProof Block Diagram



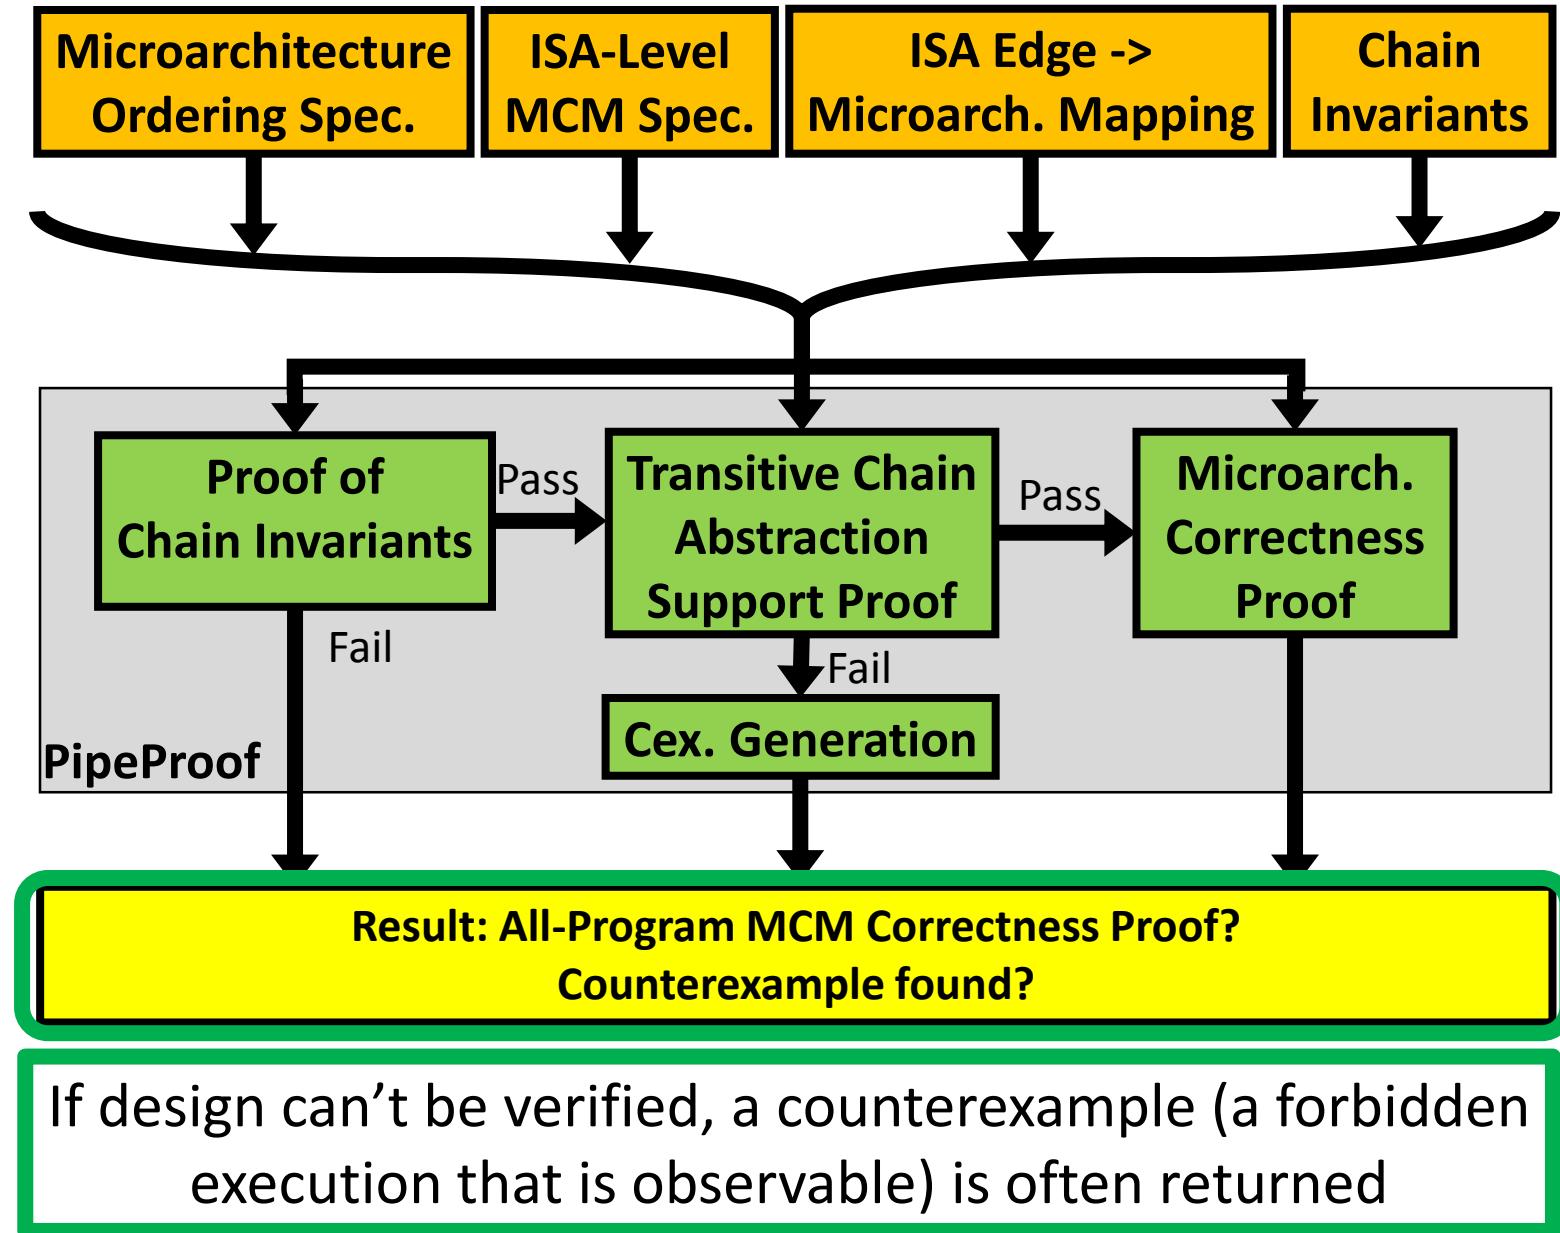
PipeProof Block Diagram



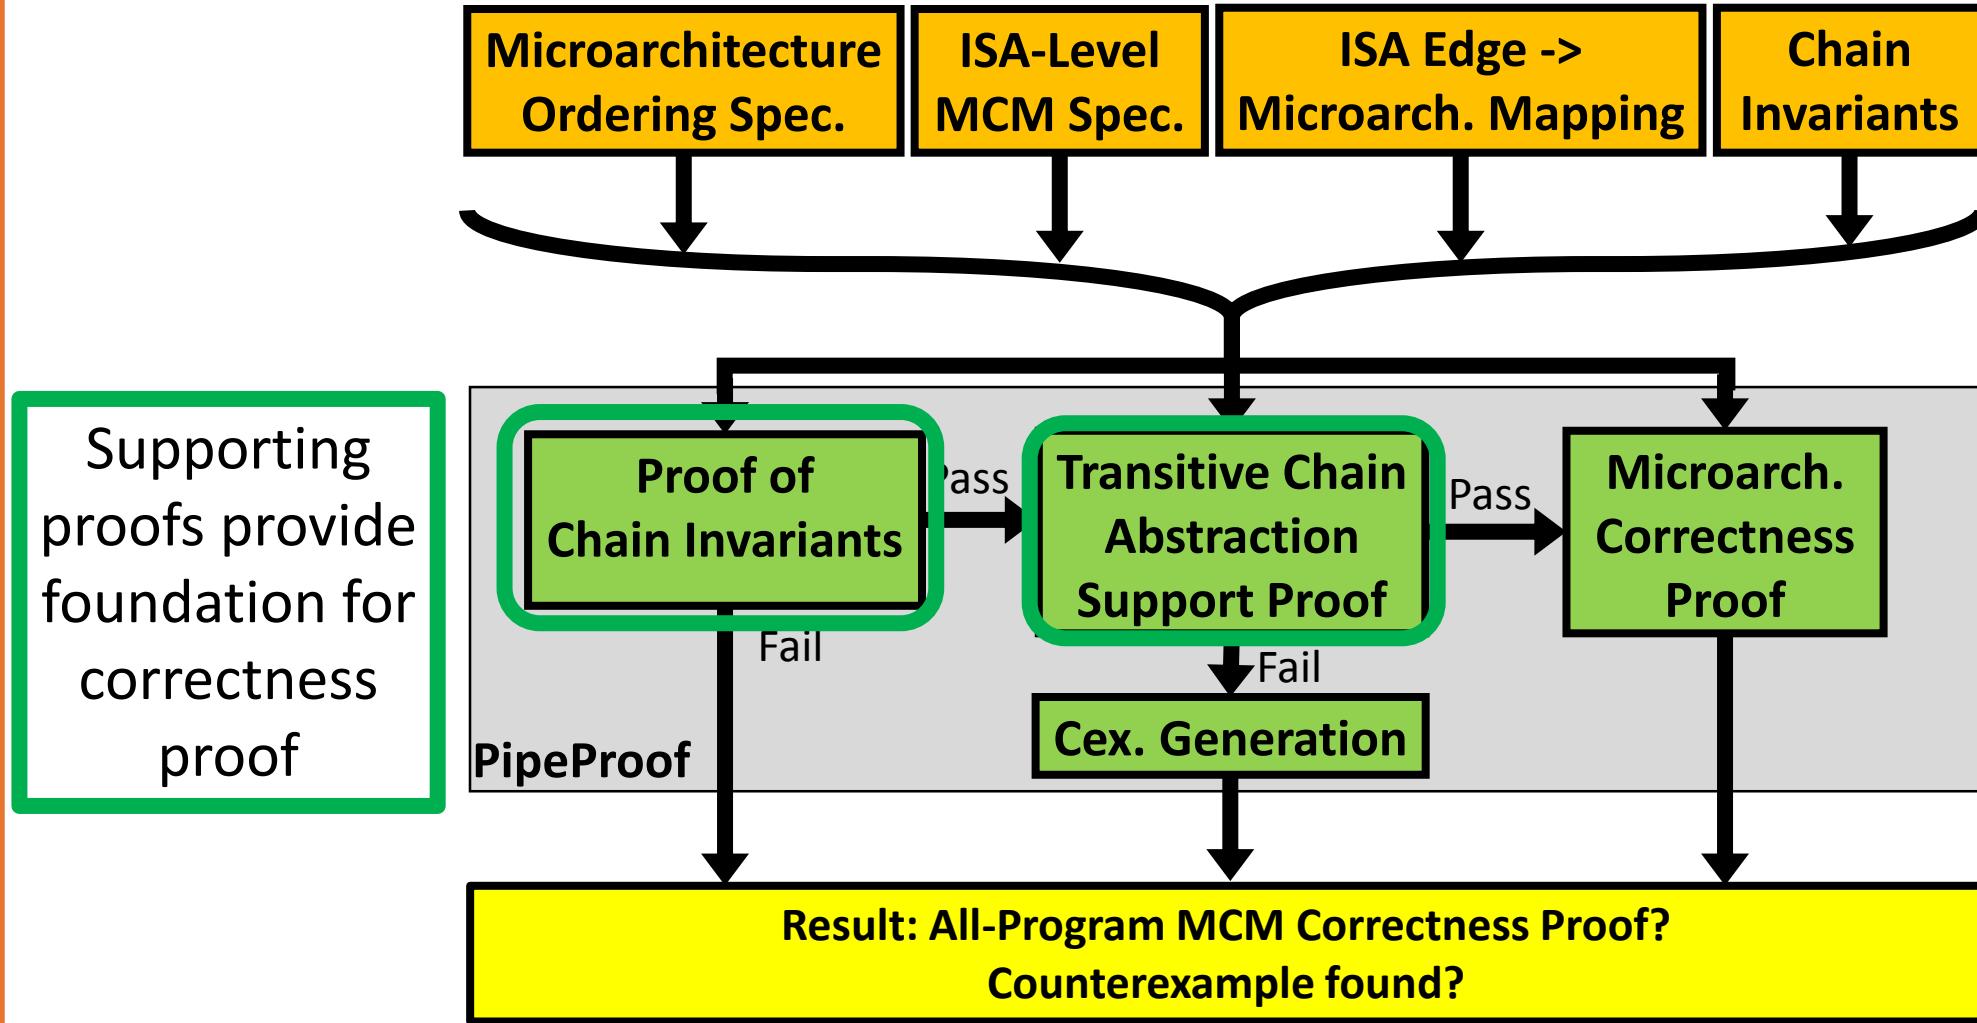
Represent repeated ISA-level patterns



PipeProof Block Diagram

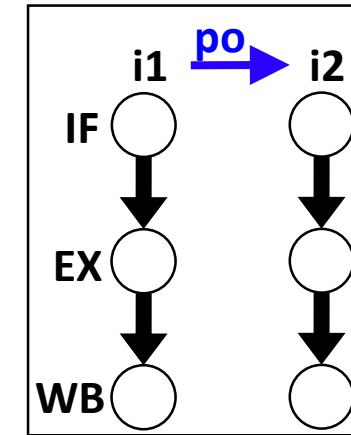


PipeProof Block Diagram



Mapping ISA-Level Edges to Microarchitecture

- Translate each edge in ISA-level cycle to microarchitectural constraints
- Do so with user-provided **Mapping Axioms**
- Example: Mapping of *po* edges



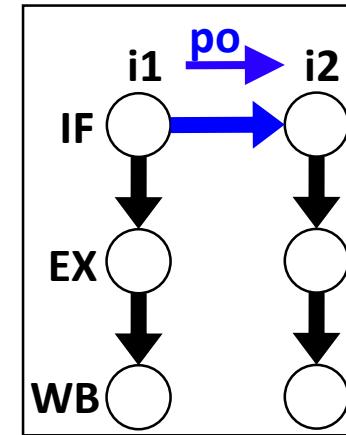
Axiom "Mapping_po":

```
forall microop "i",
forall microop "j",
(HasDependency po i j =>
  AddEdge ((i, Fetch), (j, Fetch), "po_arch", "blue")).
```



Mapping ISA-Level Edges to Microarchitecture

- Translate each edge in ISA-level cycle to microarchitectural constraints
- Do so with user-provided **Mapping Axioms**
- Example: Mapping of *po* edges



Axiom "Mapping_po":

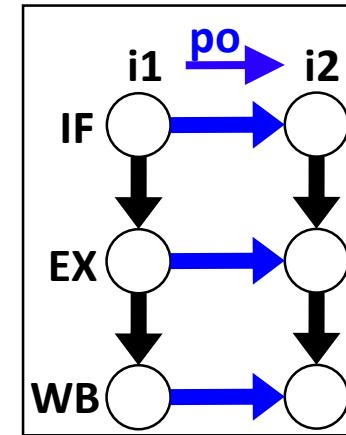
```
forall microop "i",
forall microop "j",
(HasDependency po i j =>
    AddEdge ((i, Fetch), (j, Fetch), "po_arch", "blue").
```



Mapping ISA-Level Edges to Microarchitecture

- Translate each edge in ISA-level cycle to microarchitectural constraints
- Do so with user-provided **Mapping Axioms**
- Example: Mapping of *po* edges

Blue edges between EX and WB stages added by other FIFO axioms (refer to μ spec file)



Axiom "Mapping_po":

```
forall microop "i",
forall microop "j",
(HasDependency po i j =>
    AddEdge ((i, Fetch), (j, Fetch), "po_arch", "blue")).
```

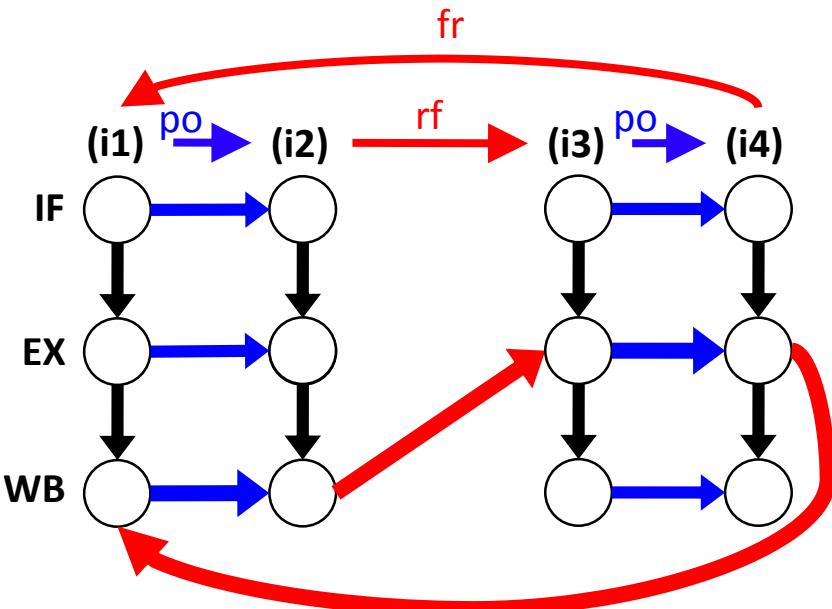


Can “litmus tests” provide complete coverage?

- Open question as to whether a set of litmus tests is **complete**

mp Litmus Test

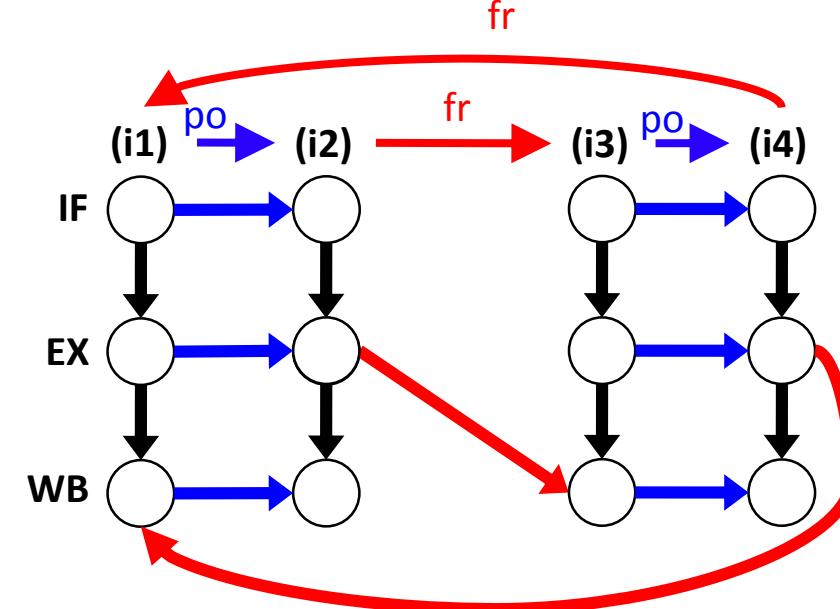
Core 0	Core 1
$x = 1;$ $y = 1;$	$r1 = y;$ $r2 = x;$
Forbid: $r1 = 1, r2 = 0$	



Cyclic => Still unobservable

sb Litmus Test

Core 0	Core 1
$x = 1;$ $r1 = y;$	$y = 1;$ $r2 = x;$
Forbid: $r1 = 0, r2 = 0$	



Acyclic => BUG!

Can “litmus tests” provide complete coverage?

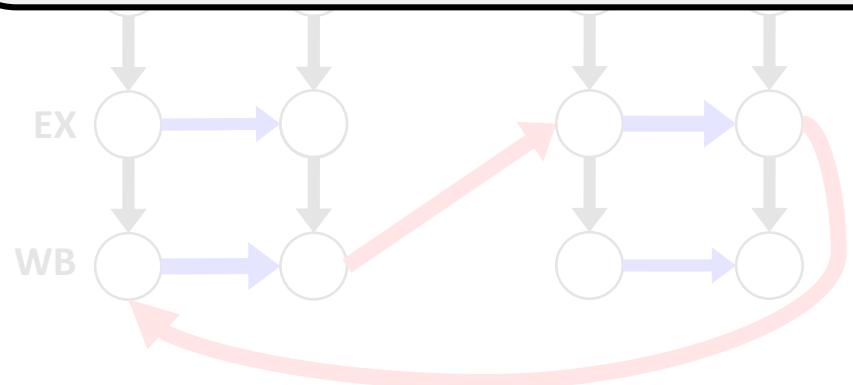
- Open question as to whether a set of litmus tests is **complete**

mp Litmus Test	
Core 0	Core 1
x = 1;	r1 = y;

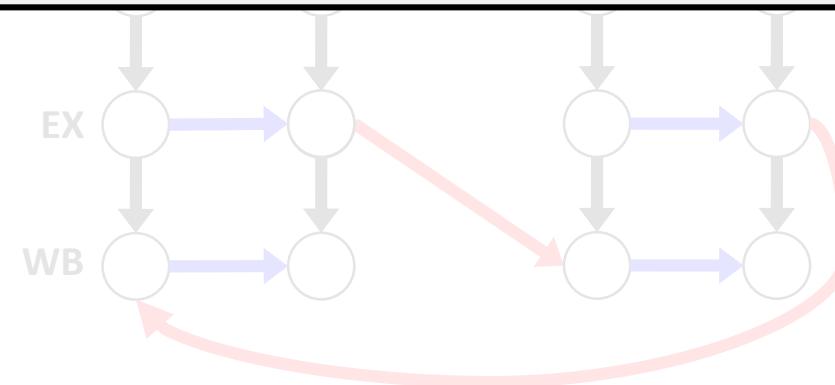
sb Litmus Test	
Core 0	Core 1
x = 1;	y = 1;

Different tests catch different bugs!

To catch all bugs, must verify across all programs!



Cyclic => Still unobservable



Acyclic => BUG!



Solution: Load Value Constraints

- Don't filter based on outcome
 - Translate all possible outcomes
- Tag each case with appropriate **load value constraints**
 - reflect the data constraints required for edge(s)
- Ongoing work: Precisely formalise the μ spec/SVA mismatch
 - How much is fundamental? How much is due to SVA verifier approximation?

Core 0	Core 1
(i1) $x = 1;$	(i3) $r1 = y;$
(i2) $y = 1;$	(i4) $r2 = x;$
SC Forbids: $r1 = 1, r2 = 0$	

Axiom "*Read_Values*":

Every load either reads **BeforeAllWrites** OR reads **FromLatestWrite**

Property to check:

`mapNode(Ld x → St x, Ld x == 0) or mapNode(St x → Ld x, Ld x == 1);`



Solution: Load Value Constraints

- Don't filter based on outcome
 - Translate all possible outcomes
- Tag each case with appropriate **load value constraints**
 - reflect the data constraints required for edge(s)
- Ongoing work: Precisely formalise the μ spec/SVA mismatch
 - How much is fundamental? How much is due to SVA verifier approximation?

Core 0	Core 1
(i1) $x = 1;$	(i3) $r1 = y;$
(i2) $y = 1;$	(i4) $r2 = x;$
SC Forbids: $r1 = 1, r2 = 0$	

Axiom "*Read_Values*":

Every load either reads **BeforeAllWrites** OR reads **FromLatestWrite**

Property to check:

mapNode($Ld\ x \rightarrow St\ x, Ld\ x == 0$) or mapNode($St\ x \rightarrow Ld\ x, Ld\ x == 1$);



Solution: Load Value Constraints

- Don't filter based on outcome
 - Translate all possible outcomes
- Tag each case with appropriate **load value constraints**
 - reflect the data constraints required for edge(s)
- Ongoing work: Precisely formalise the μ spec/SVA mismatch
 - How much is fundamental? How much is due to SVA verifier approximation?

Core 0	Core 1
(i1) $x = 1;$	(i3) $r1 = y;$
(i2) $y = 1;$	(i4) $r2 = x;$
SC Forbids: $r1 = 1, r2 = 0$	

Axiom "*Read_Values*":

Every load either reads **BeforeAllWrites** OR reads **FromLatestWrite**

Property to check:

`mapNode(Ld x → St x, Ld x == 0)` or `mapNode(St x → Ld x, Ld x == 1)`;



Solution: Load Value Constraints

- Don't filter based on outcome
 - Translate all possible outcomes
- Tag each case with appropriate **load value constraints**
 - reflect the data constraints required for edge(s)
- Ongoing work: Precisely formalise the μ spec/SVA mismatch
 - How much is fundamental? How much is due to SVA verifier approximation?

Core 0	Core 1
(i1) $x = 1;$	(i3) $r1 = y;$
(i2) $y = 1;$	(i4) $r2 = x;$
SC Forbids: $r1 = 1, r2 = 0$	

Axiom "*Read_Values*":

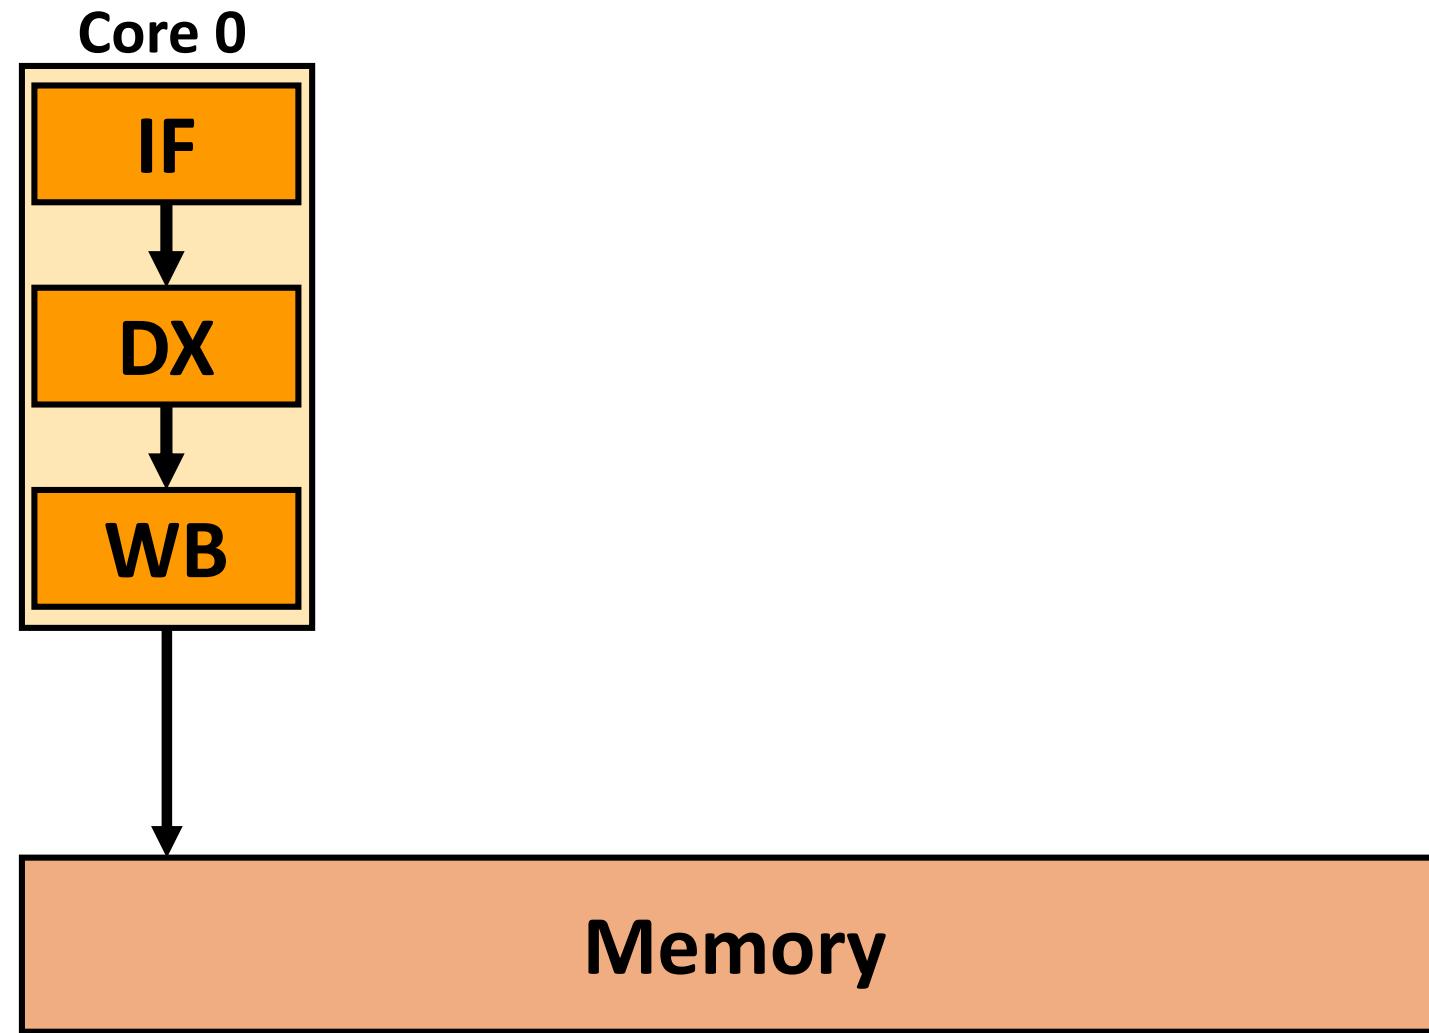
Every load either reads **BeforeAllWrites** OR reads **FromLatestWrite**

Property to check:

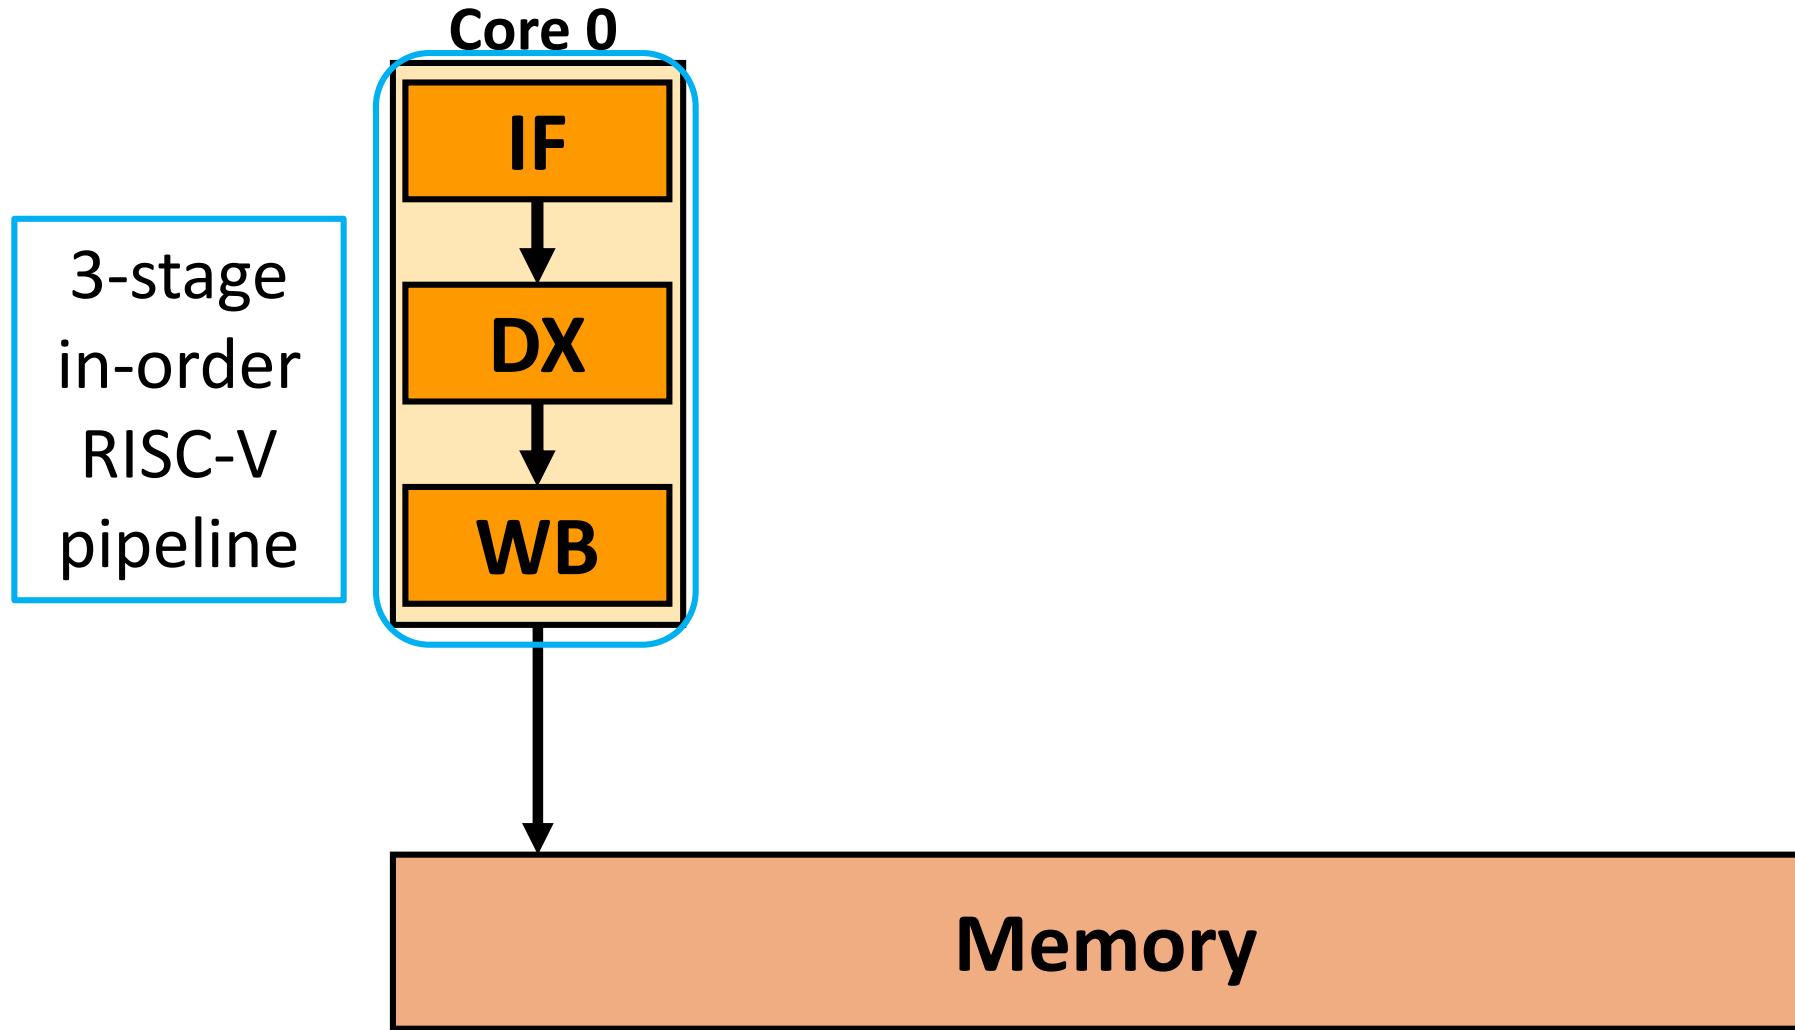
mapNode($Ld\ x \rightarrow St\ x, Ld\ x == 0$) or mapNode($St\ x \rightarrow Ld\ x, Ld\ x == 1$);



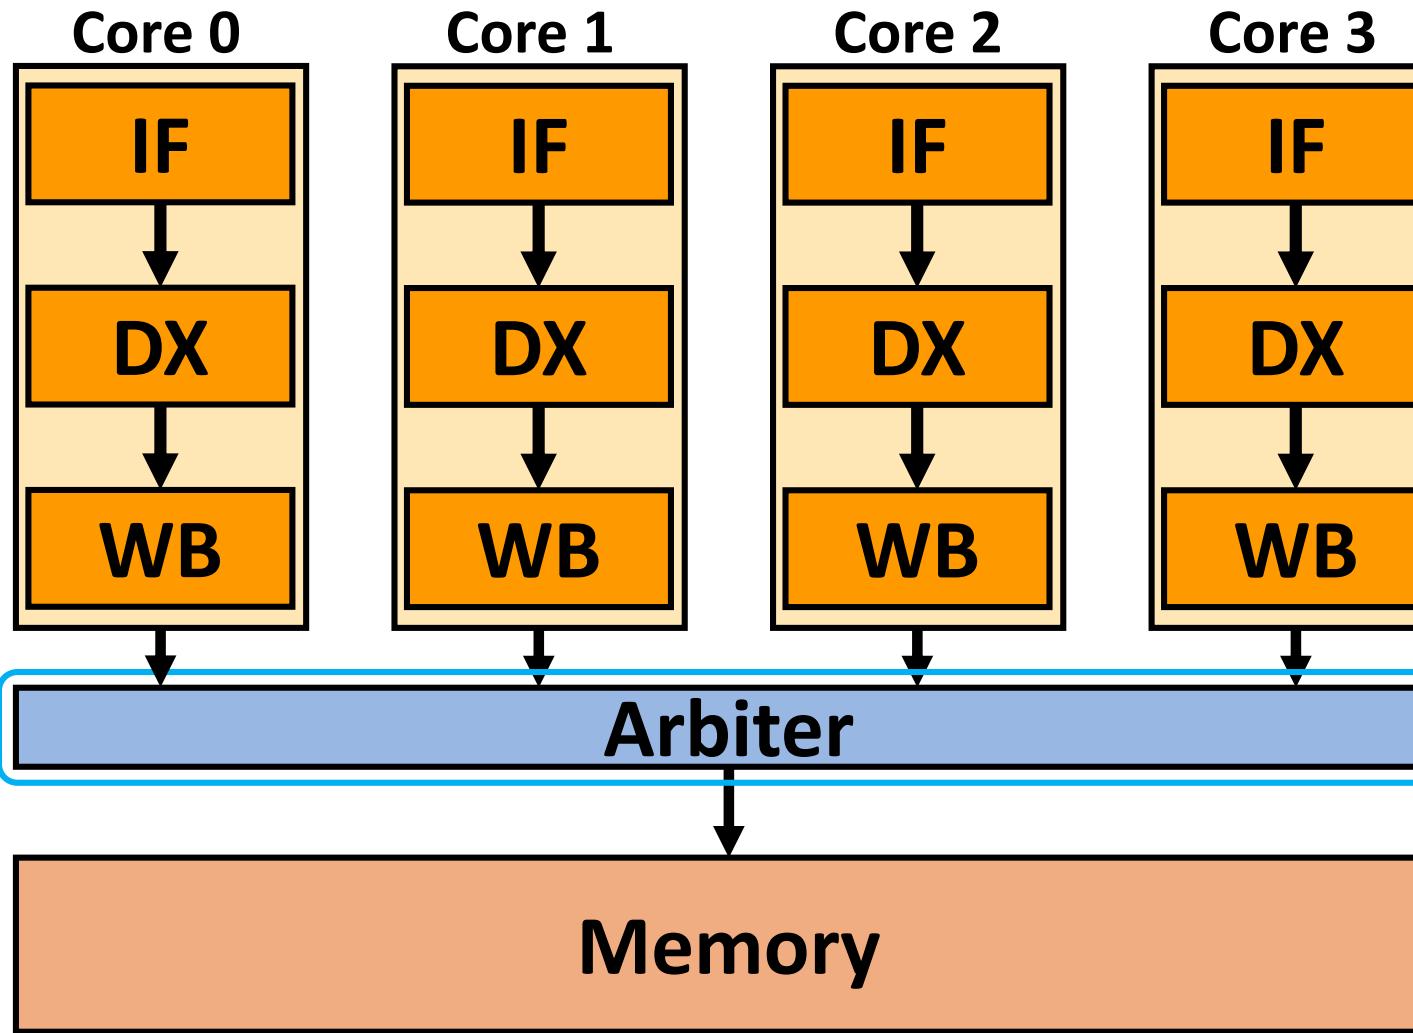
Multi-V-scale: a Multicore Case Study



Multi-V-scale: a Multicore Case Study

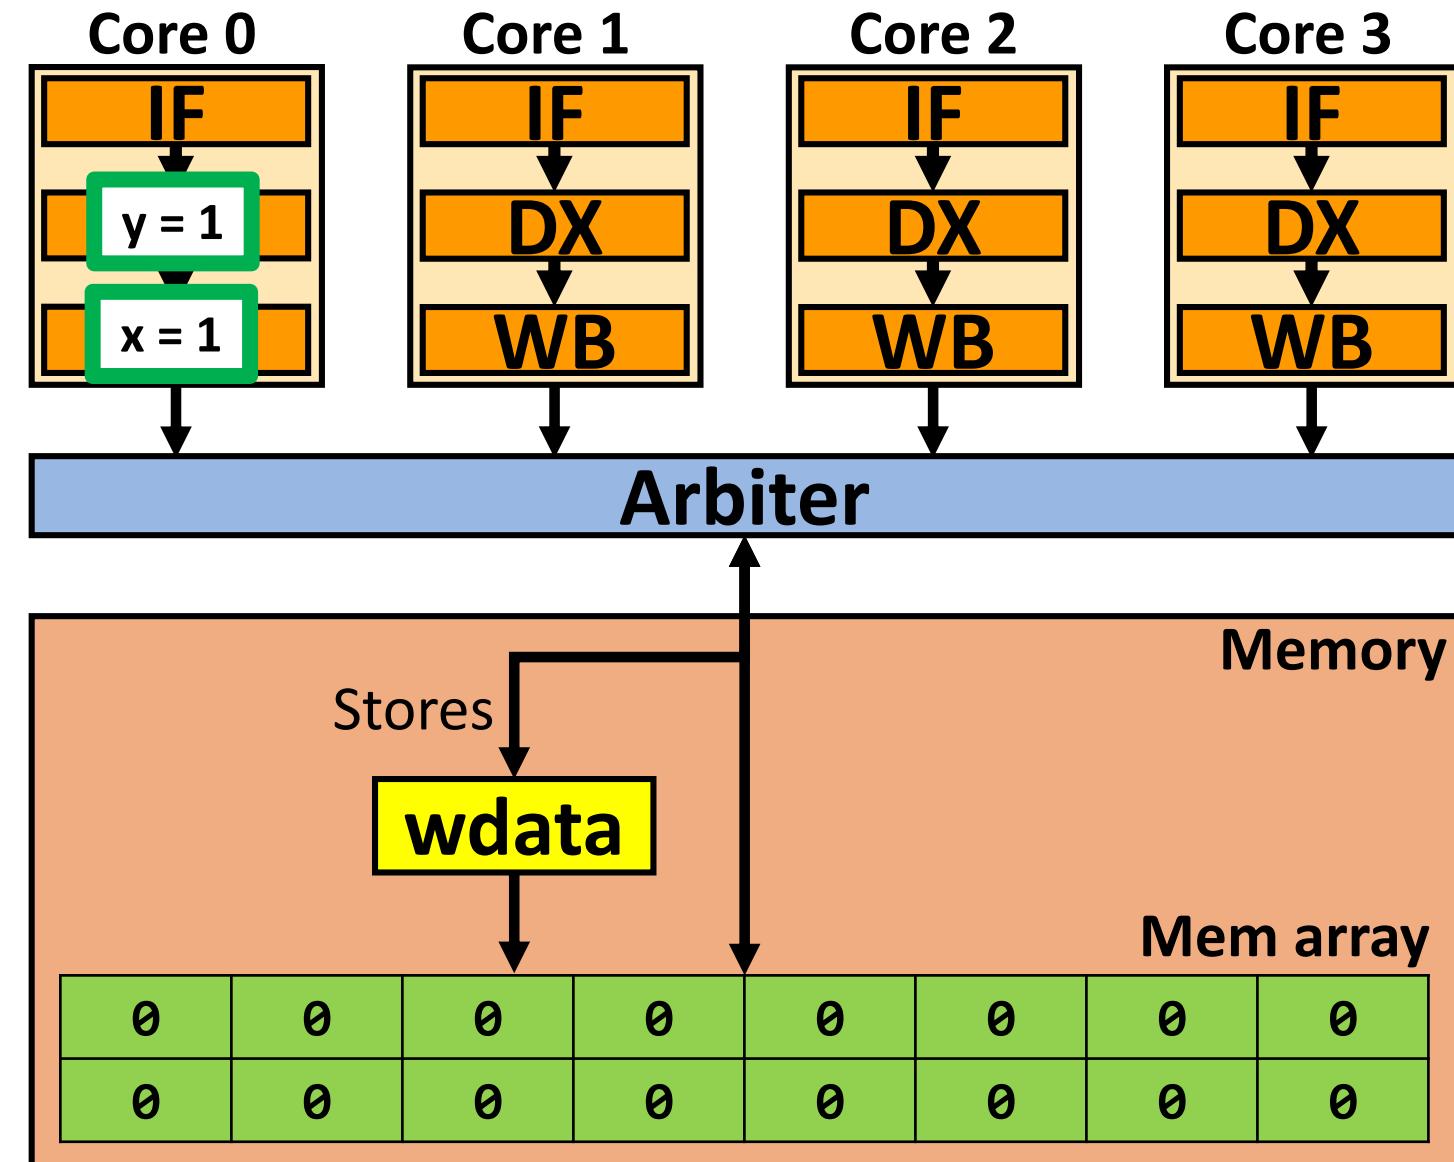


Multi-V-scale: a Multicore Case Study



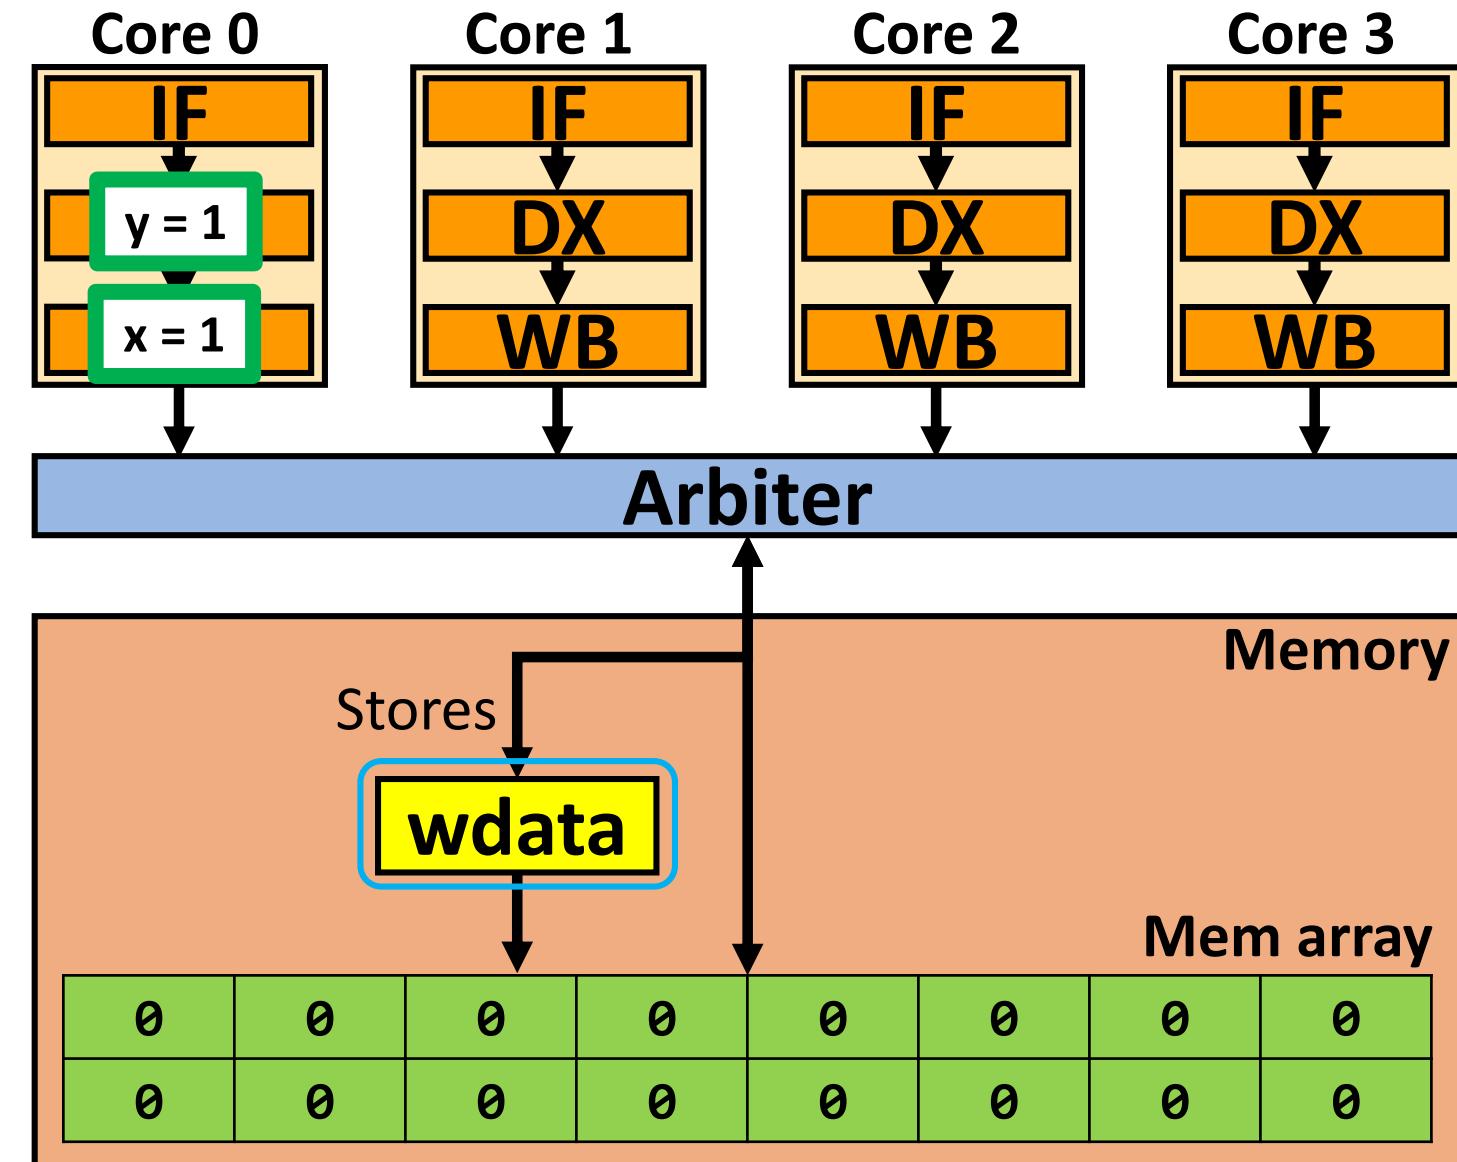
Bug Discovered in V-scale Mem. Implementation

- When two stores are sent to memory in successive cycles, first of two stores is dropped by memory!
- Bug would occur even in single-core V-scale
- Fixed bug by eliminating intermediate wdata reg



Bug Discovered in V-scale Mem. Implementation

- When two stores are sent to memory in successive cycles, first of two stores is dropped by memory!
- Bug would occur even in single-core V-scale
- Fixed bug by eliminating intermediate wdata reg



Bug Discovered in V-scale Mem. Implementation

- When two stores are sent to memory in successive cycles, first of two stores is dropped by memory!
- Bug would occur even in single-core V-scale
- Fixed bug by eliminating intermediate wdata reg

